

# CLKSCREW

Exposing the Perils of Security-Oblivious Energy Management

Adrian Tang, Simha Sethumadhavan, Salvatore Stolfo  
*Columbia University*

# \$> whoami

---

**Adrian Tang** — Ph.D. candidate @ Columbia University

Member of

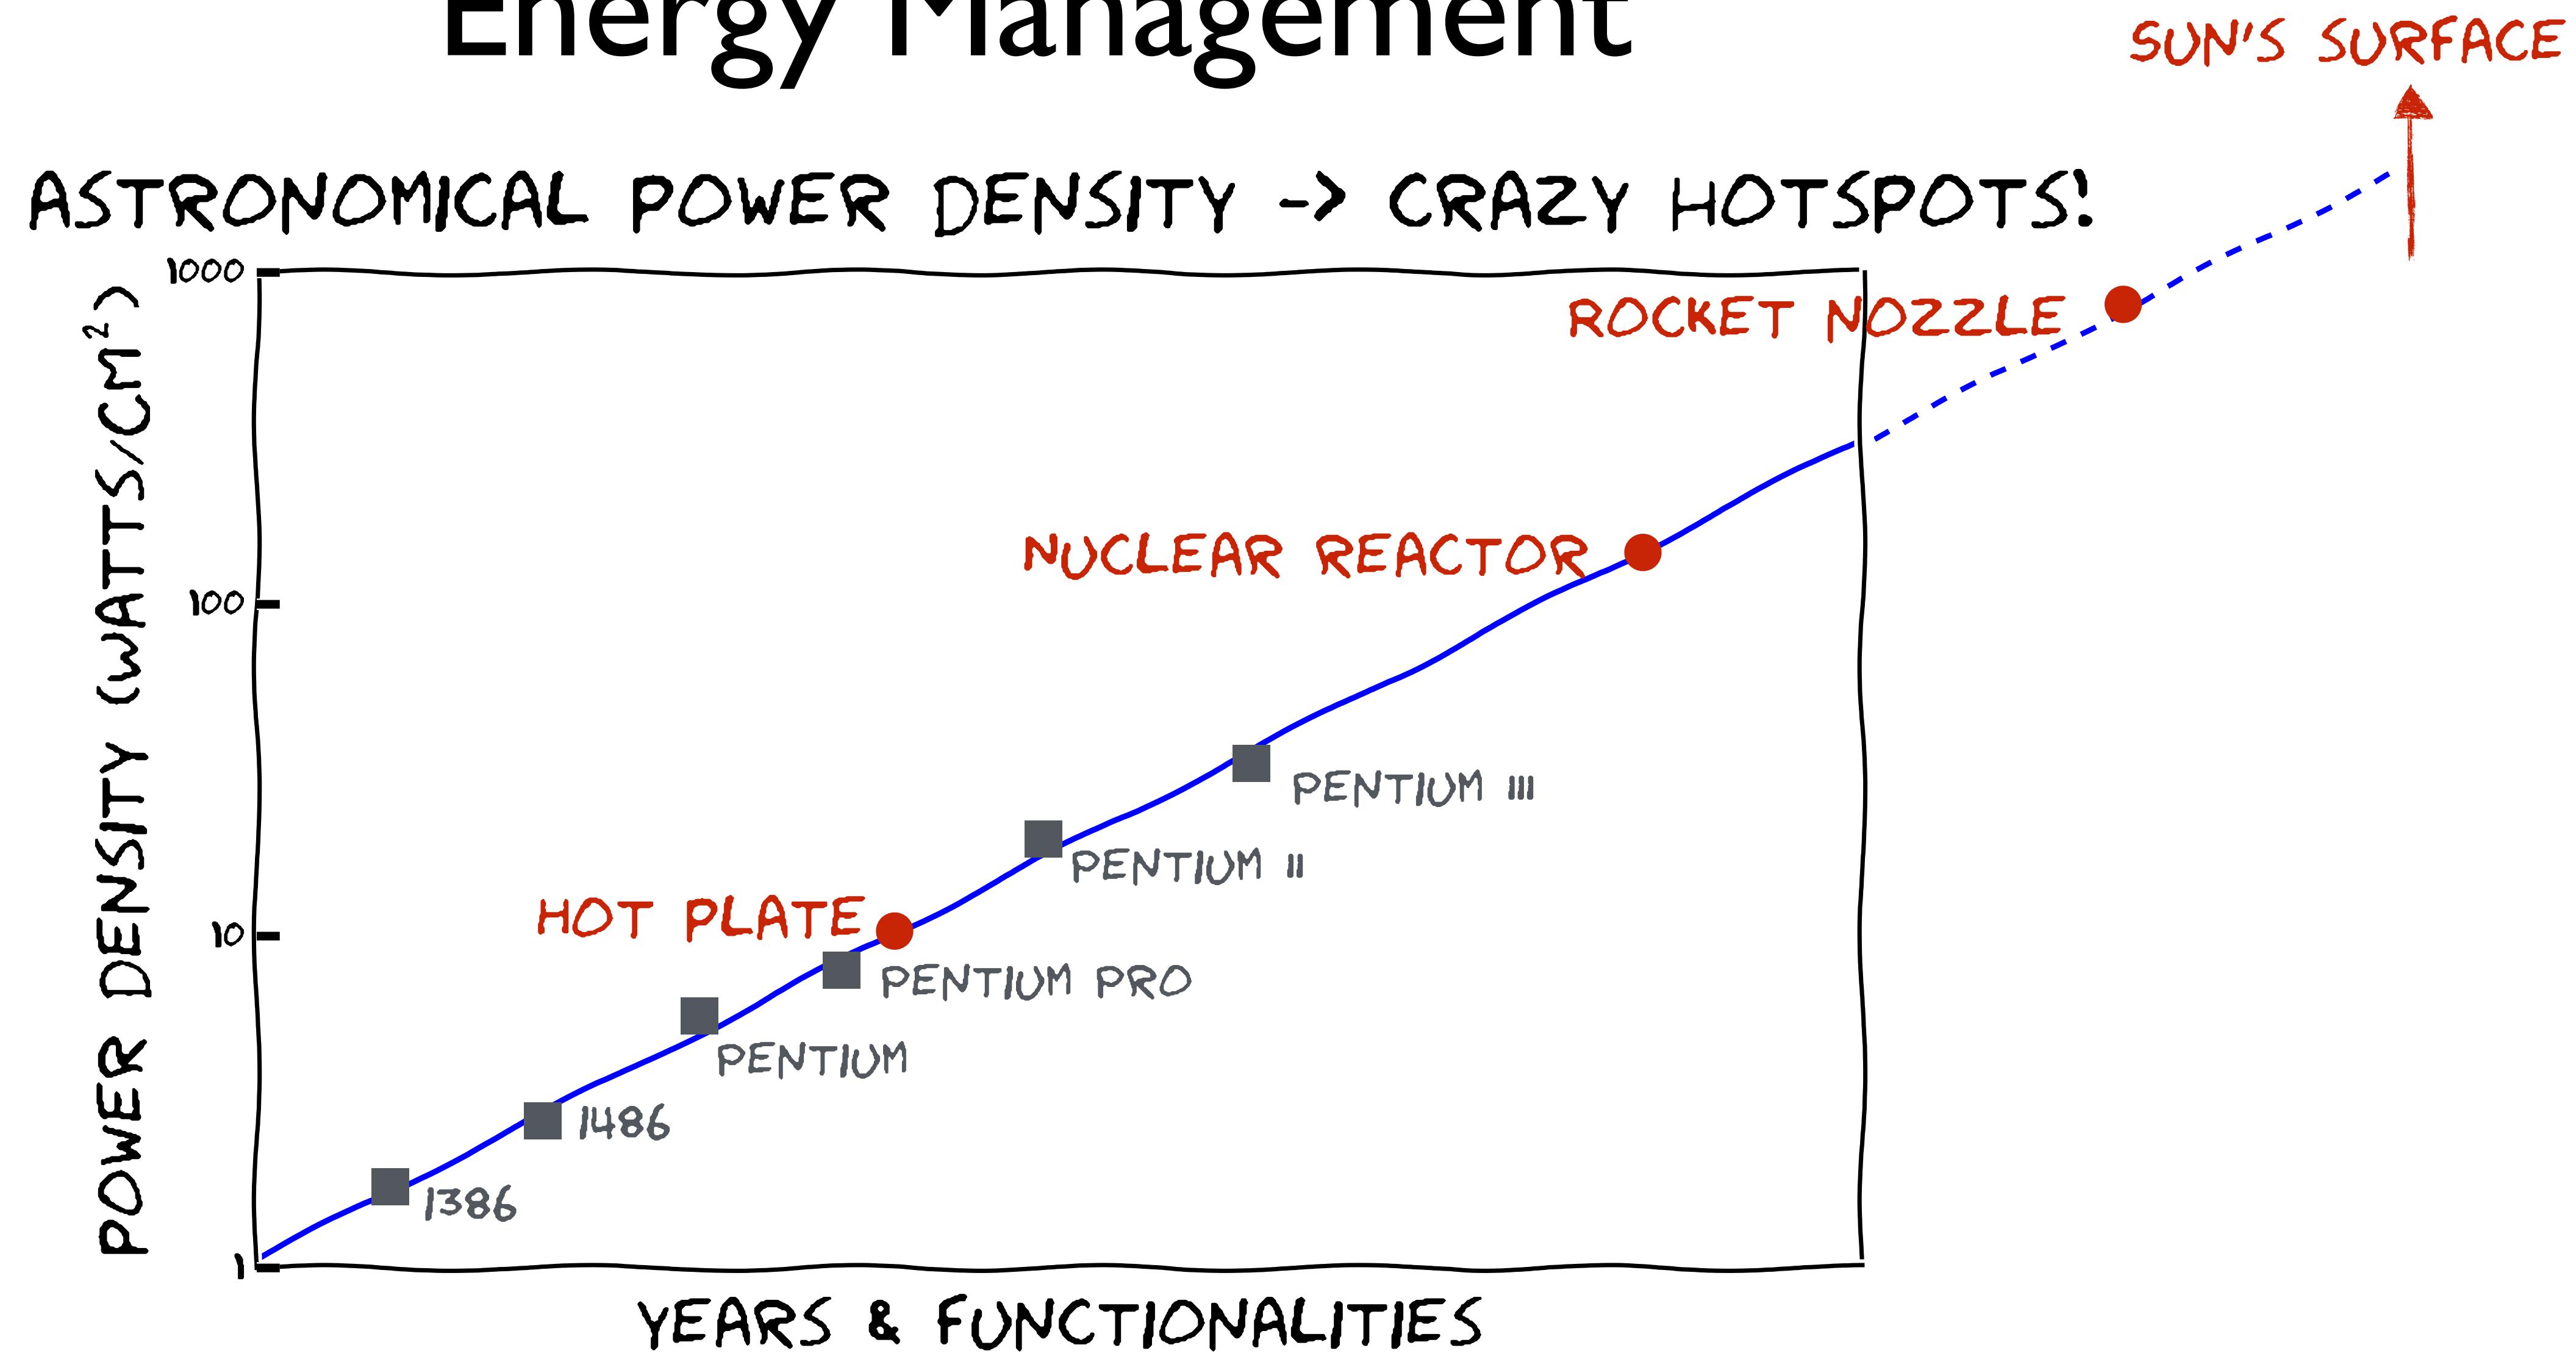
- Intrusion Detection Systems Labs (IDS)
- Computer Architecture and Security Technologies Lab (CASTL)

Interests

- Reverse engineering
- Bug hunting
- Malware analysis

Security research mainly focusing on hardware-software interfaces

# Today's systems cannot exist without Energy Management



# Today's systems cannot exist without Energy Management

## Industry

Snapdragon 820 Consumes 30% Less Power<sup>1</sup>

Power Consumption Trend  
Normalized Real Life Usage

Enhanced Tuning/Overclocking on 4th Gen  
Intel® Core™ Processors

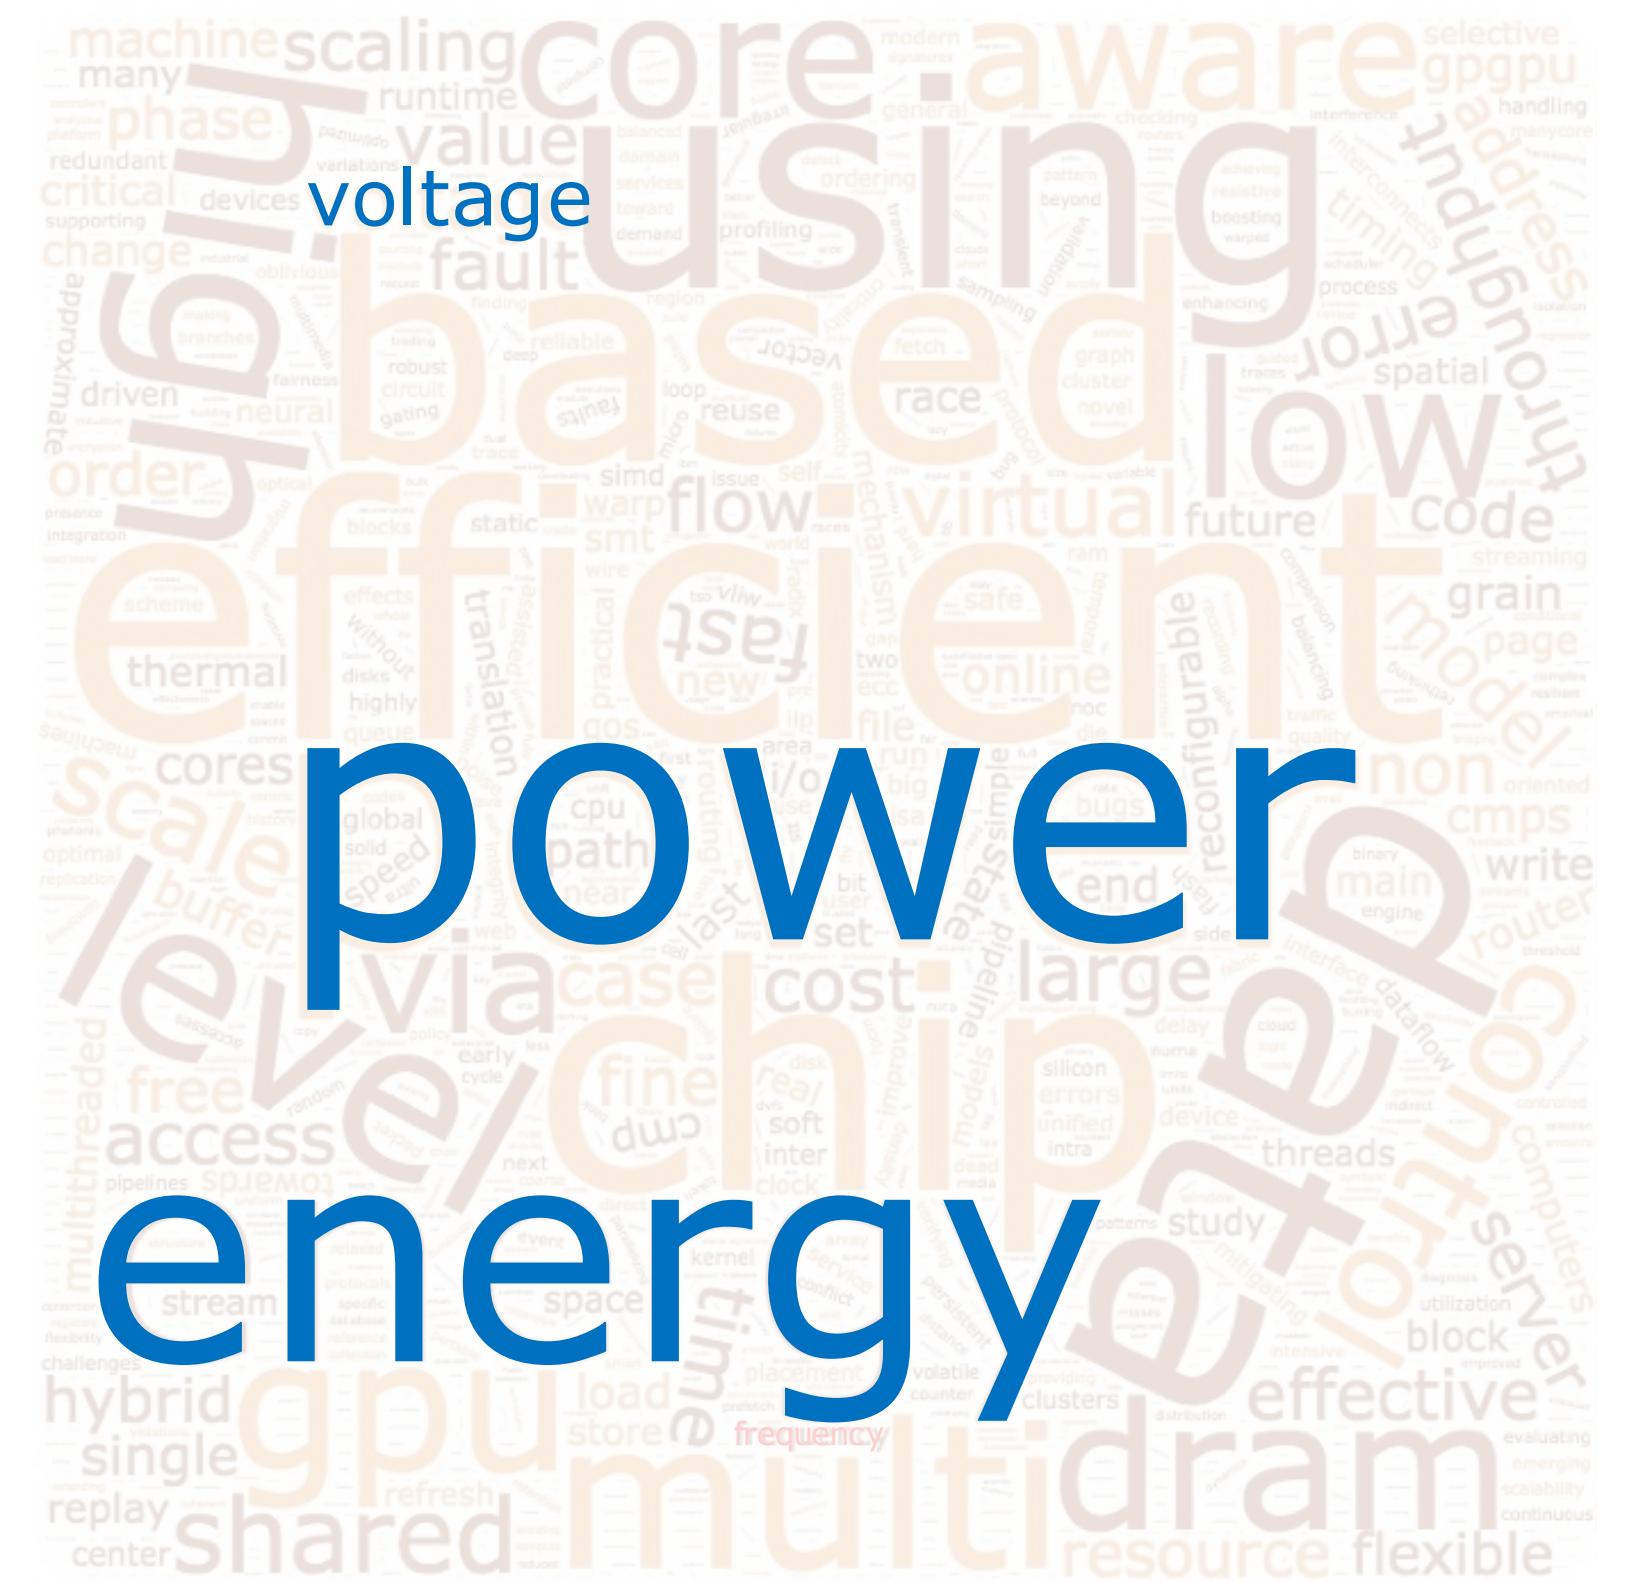
Increase core ratios via Turbo

New Power Management Capabilities

- Per Core P-States (PCPS)
  - Allows cores to run at individual frequency/voltage
- Energy Efficient Turbo Mode (EET)
  - Core throughput / stall behavior monitored
  - Core frequency is increased only if it is energy efficient
- Uncore Voltage/Frequency Scaling (UFS)
  - Nehalem: Core could turbo up, Uncore at fixed frequency
  - Sandy Bridge: Core and Uncore turbo up/down together
  - Haswell: Each Core & Uncore treated independently
    - Core-Bound Applications: Drive Core frequency higher without needing to increase Uncore
    - LLC/Memory-Bound Applications: Drive Uncore frequency higher without burning core power

Core #	Independent Frequencies
2	Low
4	High
6	High
8	Very High
10	Medium
R	Medium

## Academia



Source: Word-cloud from ISCA, ASPLOS, MICRO, HPCA (2000 - 2016)

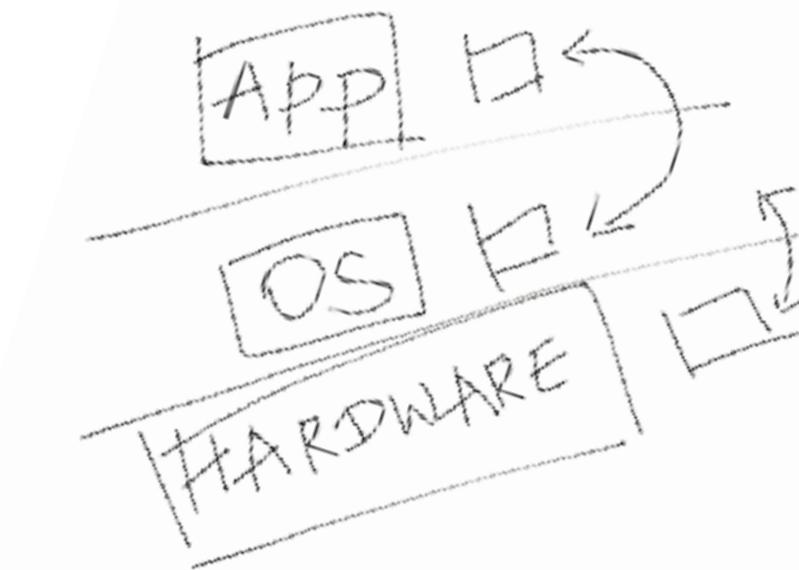


Essential



Pervasive

Today's systems cannot ~~exist without~~ stay secure with  
**Energy Management**



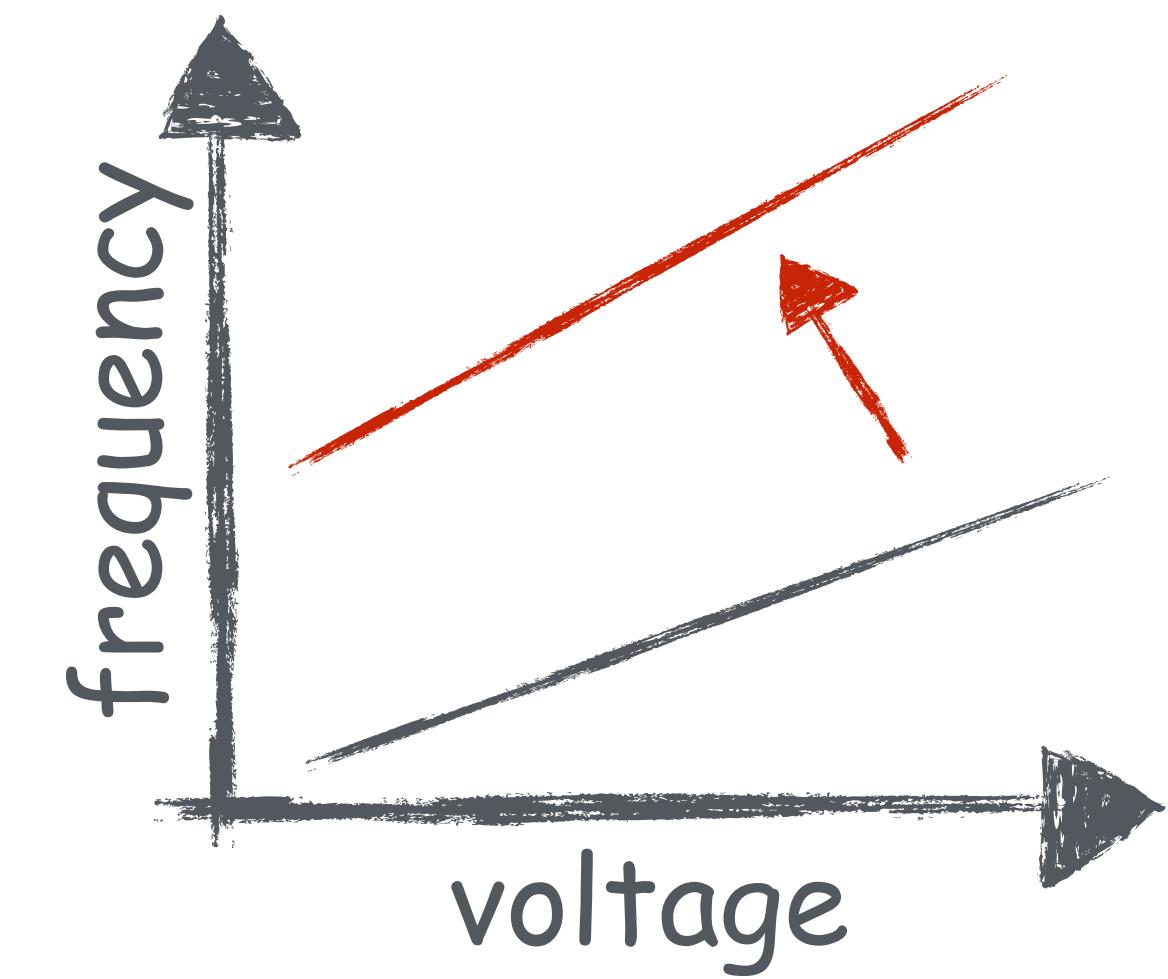
Complicated

# Exploiting software interfaces to Energy Management

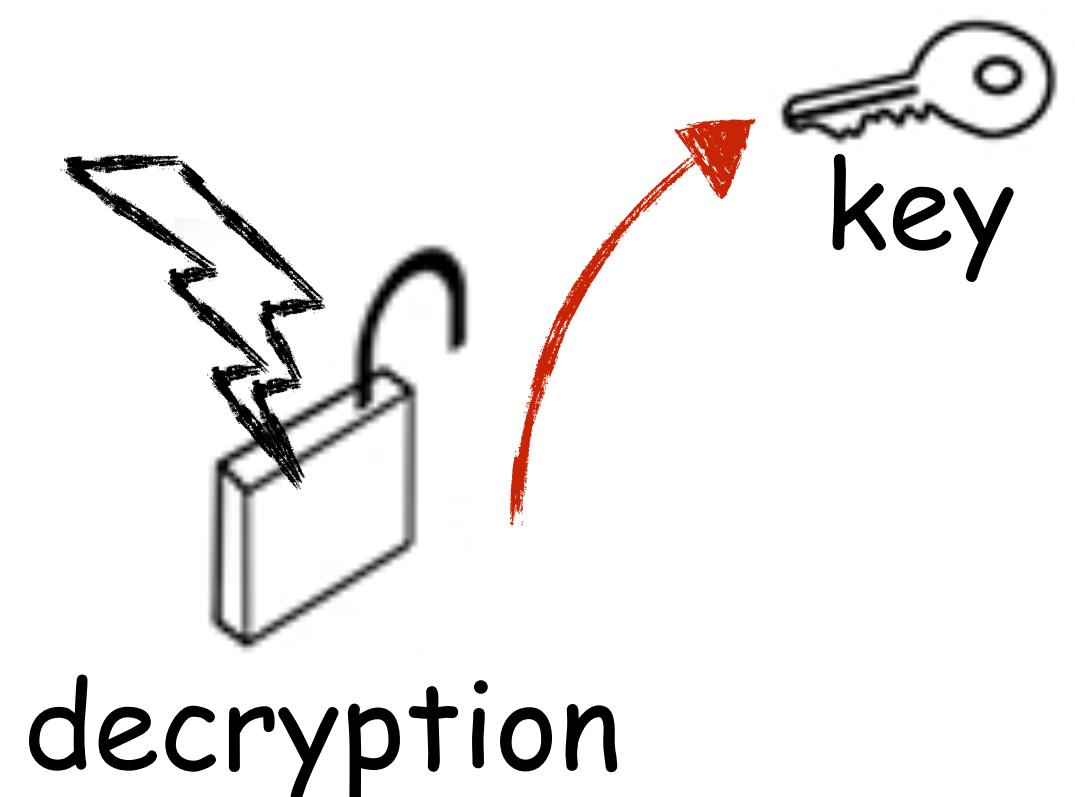
Software-based  
attacker



Stretch  
operational limits



Induce faults



# Exploiting software interfaces to Energy Management

Software-based  
attacker



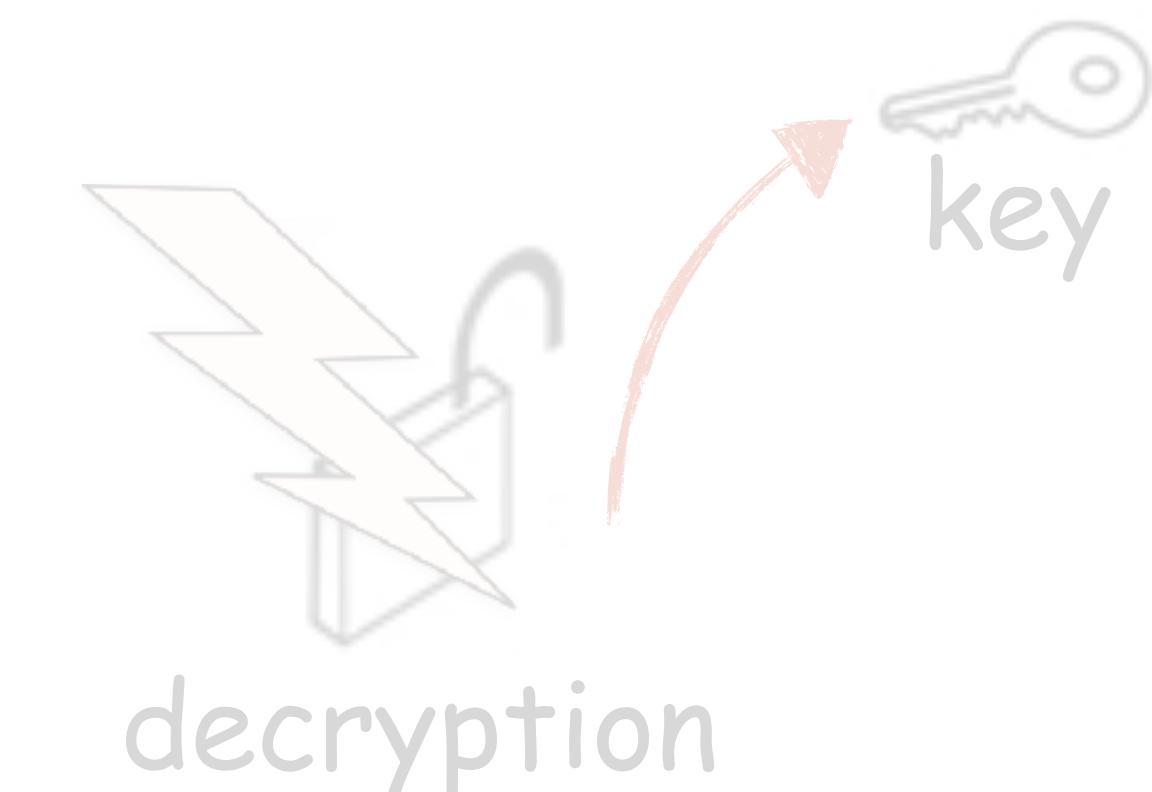
Stretch

## Traditional fault attacks

- Need physical proximity
- Need separate equipment
- Soldering, crocodile clips, wire, etc



Induce faults



# **CLKSCREW**: Exposing the perils of security-oblivious Energy Management

**New attack vector** that exploits energy management

**Practical attack** on trusted computing on ARM devices

**Impacts** hundreds of millions of deployed devices

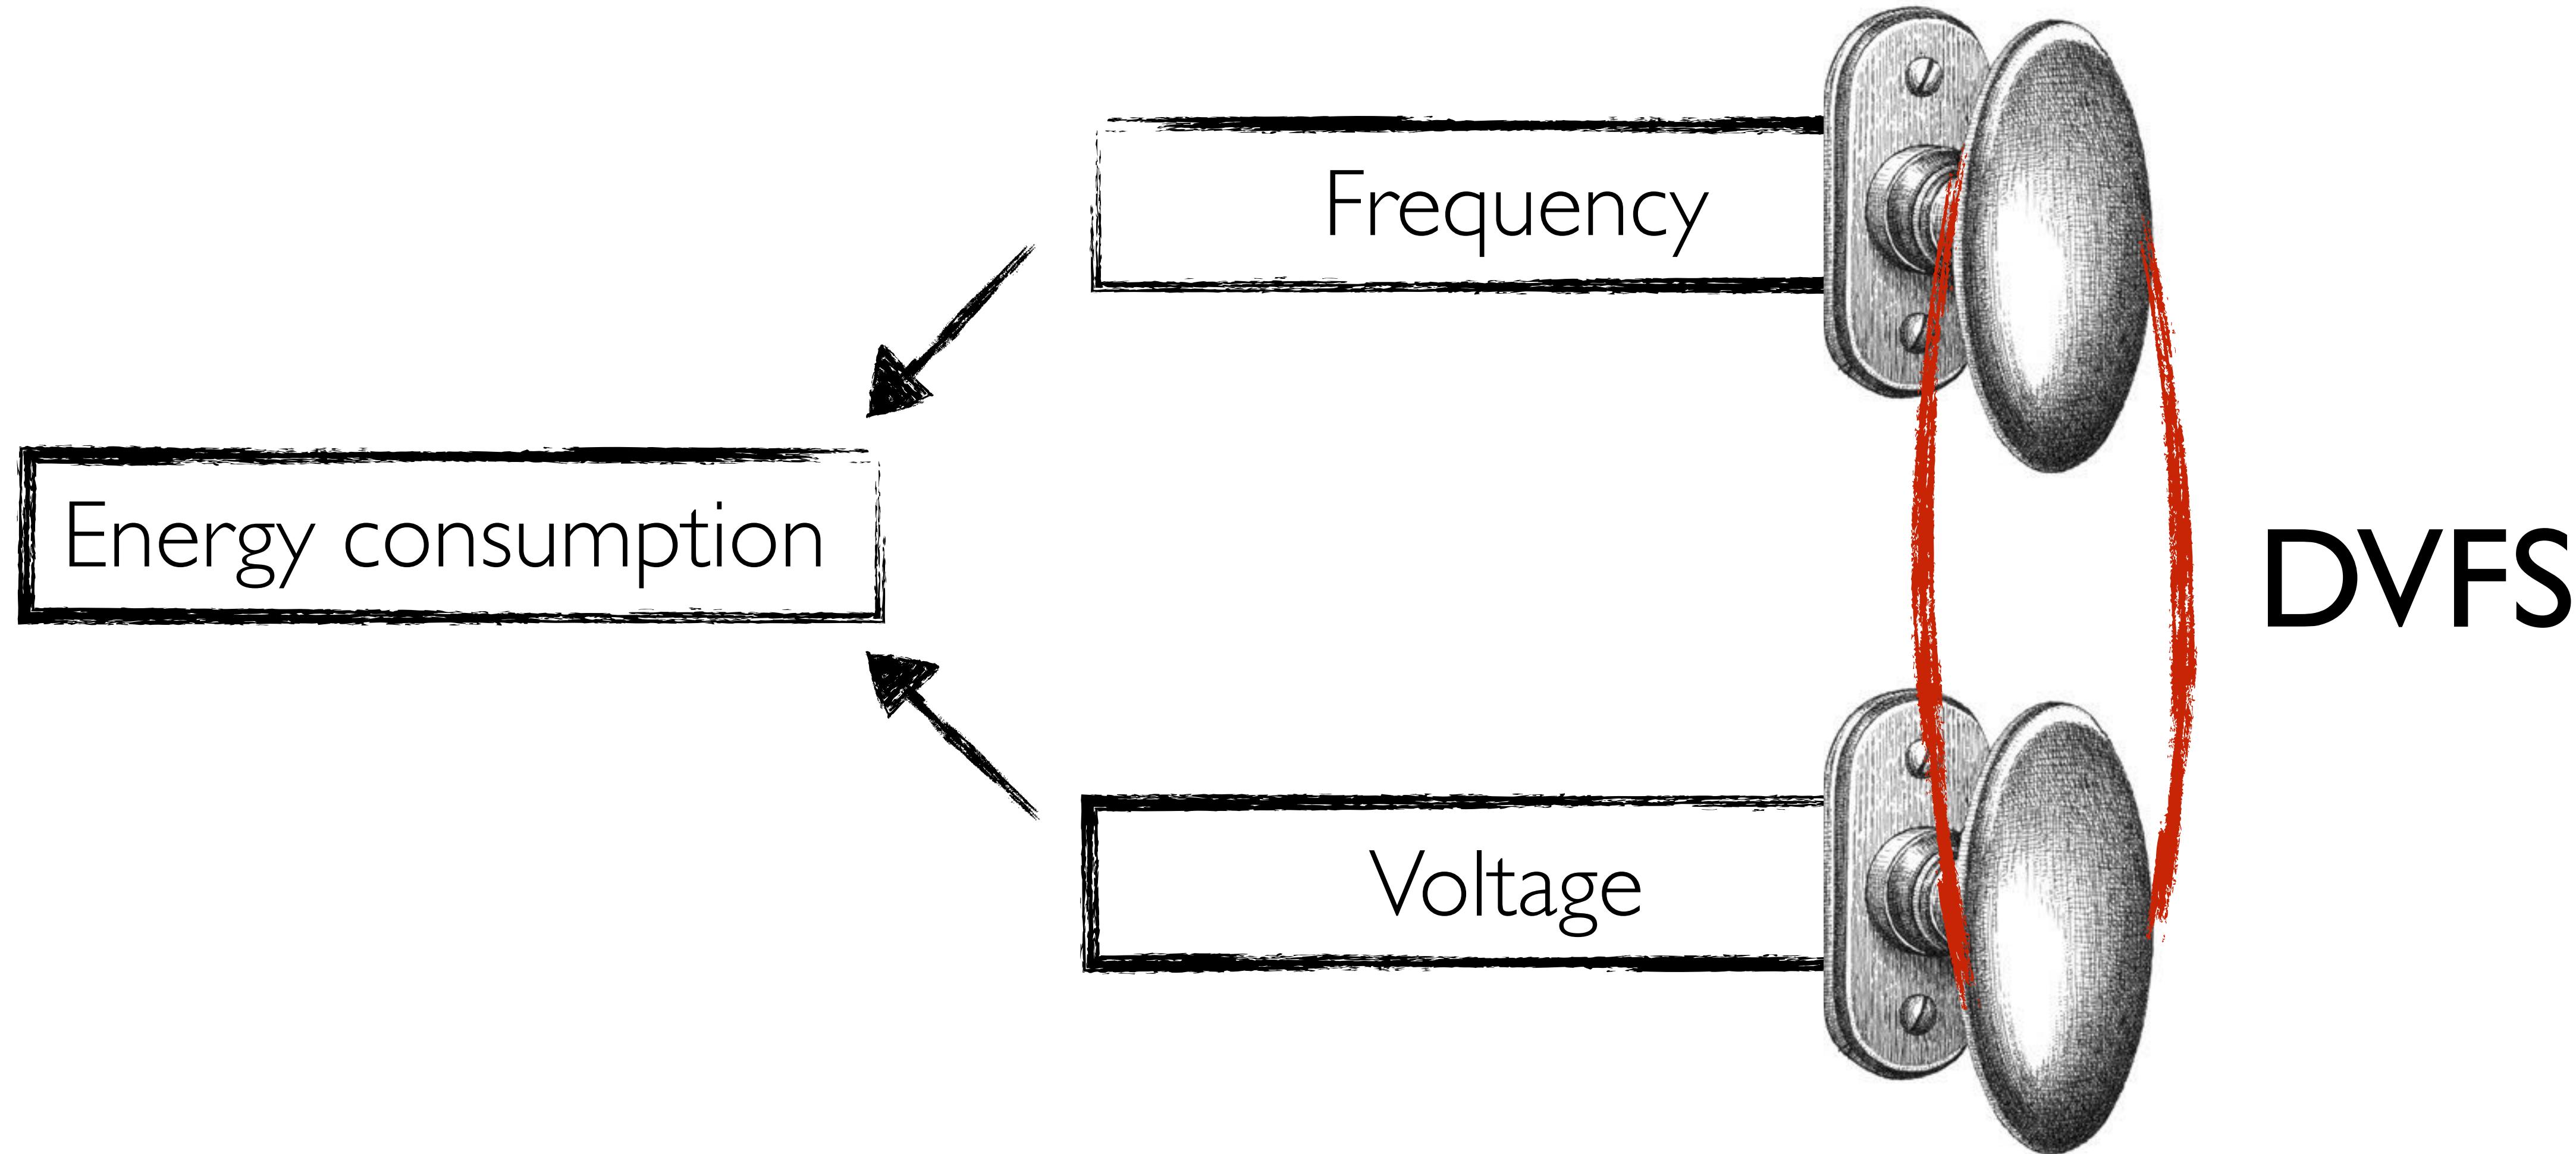
**Lessons** for future energy management designs to be security-conscious

# Outline

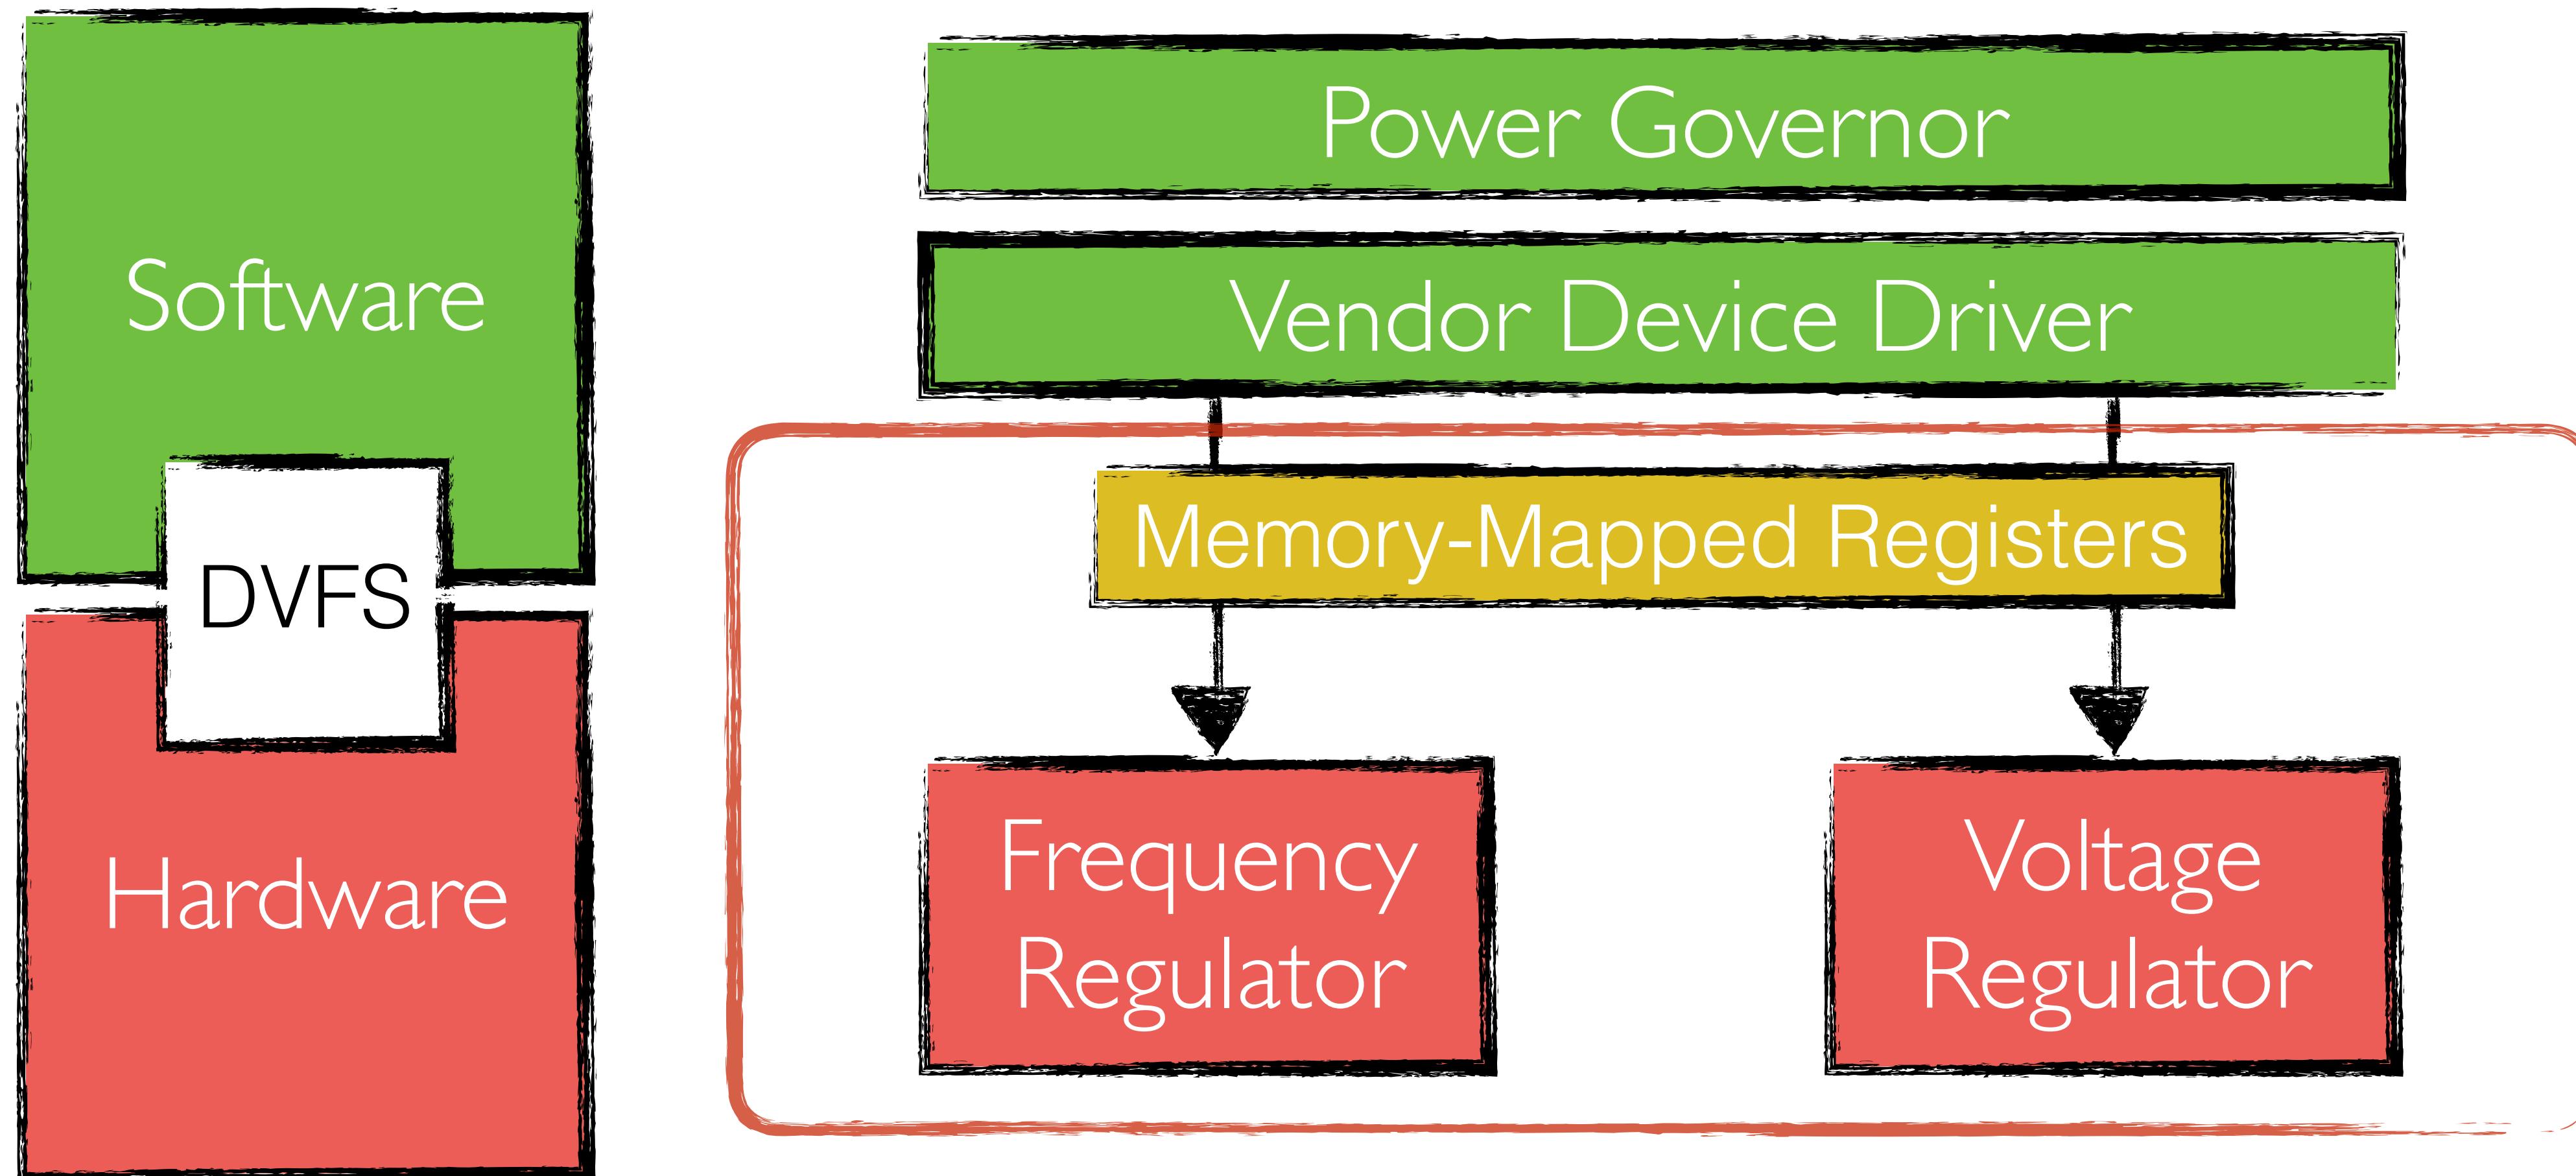
---

- I. DVFS and Deep Dive into Hardware Regulators
- II. The CLKSCREW Attack
- III. Trustzone Attack I: Secret AES Key Inference
- IV. Trustzone Attack 2:Tricking RSA Signature Validation
- V. Concluding Remarks

# Dynamic Voltage and Frequency Scaling (DVFS)



# Hardware & Software Support for DVFS



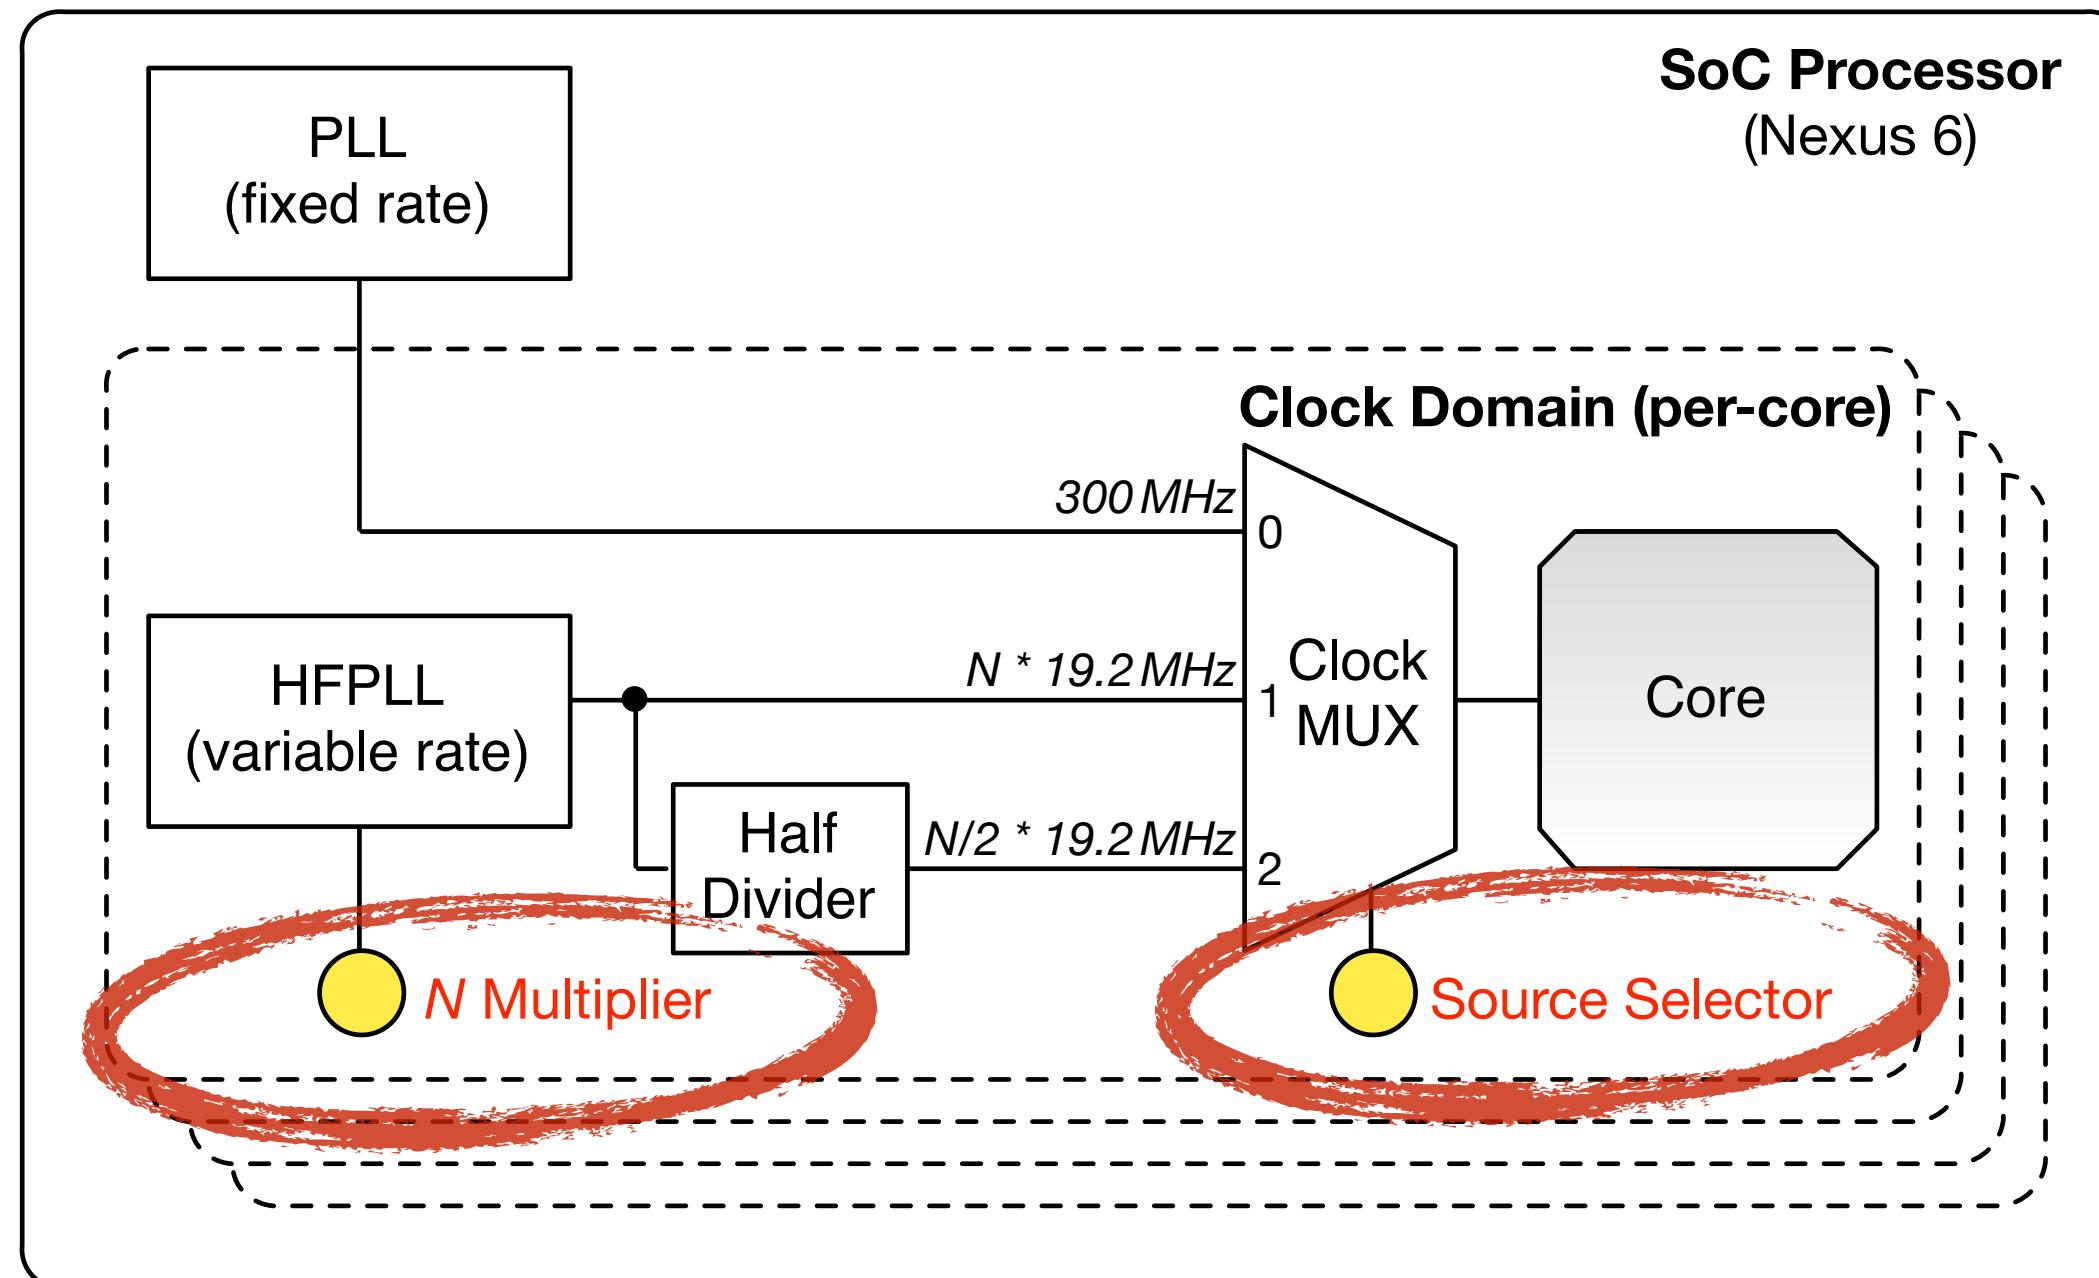
# Our Target

Nexus 6 - ARMv7  
Quad-core 2.6GHz  
Snapdragon Krait SoC

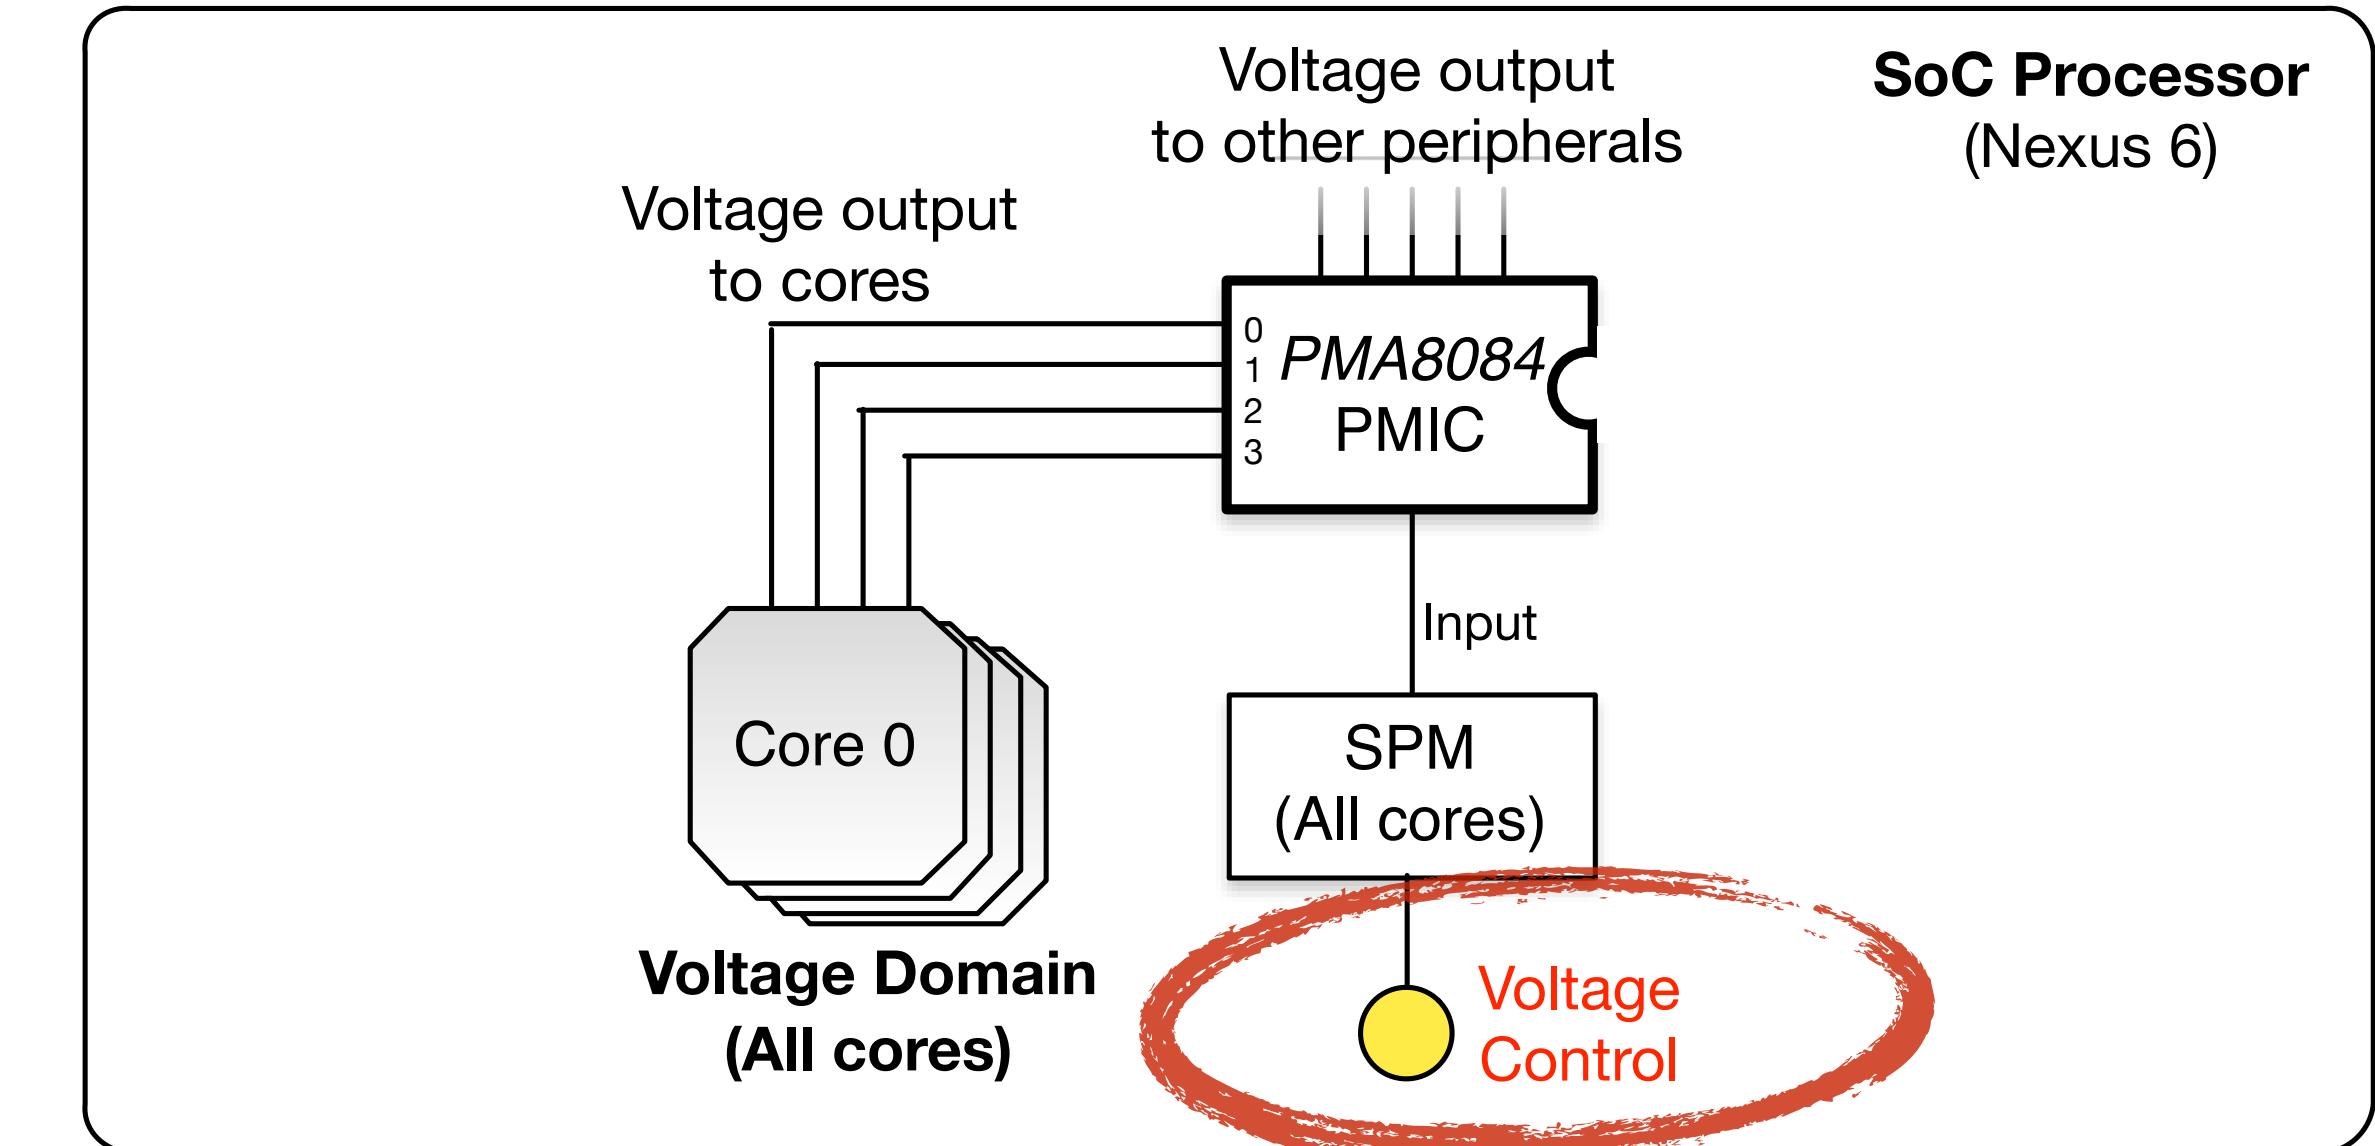


# Nexus 6 - HW Regulators and SW Interfaces

## Frequency regulators



## Voltage regulators

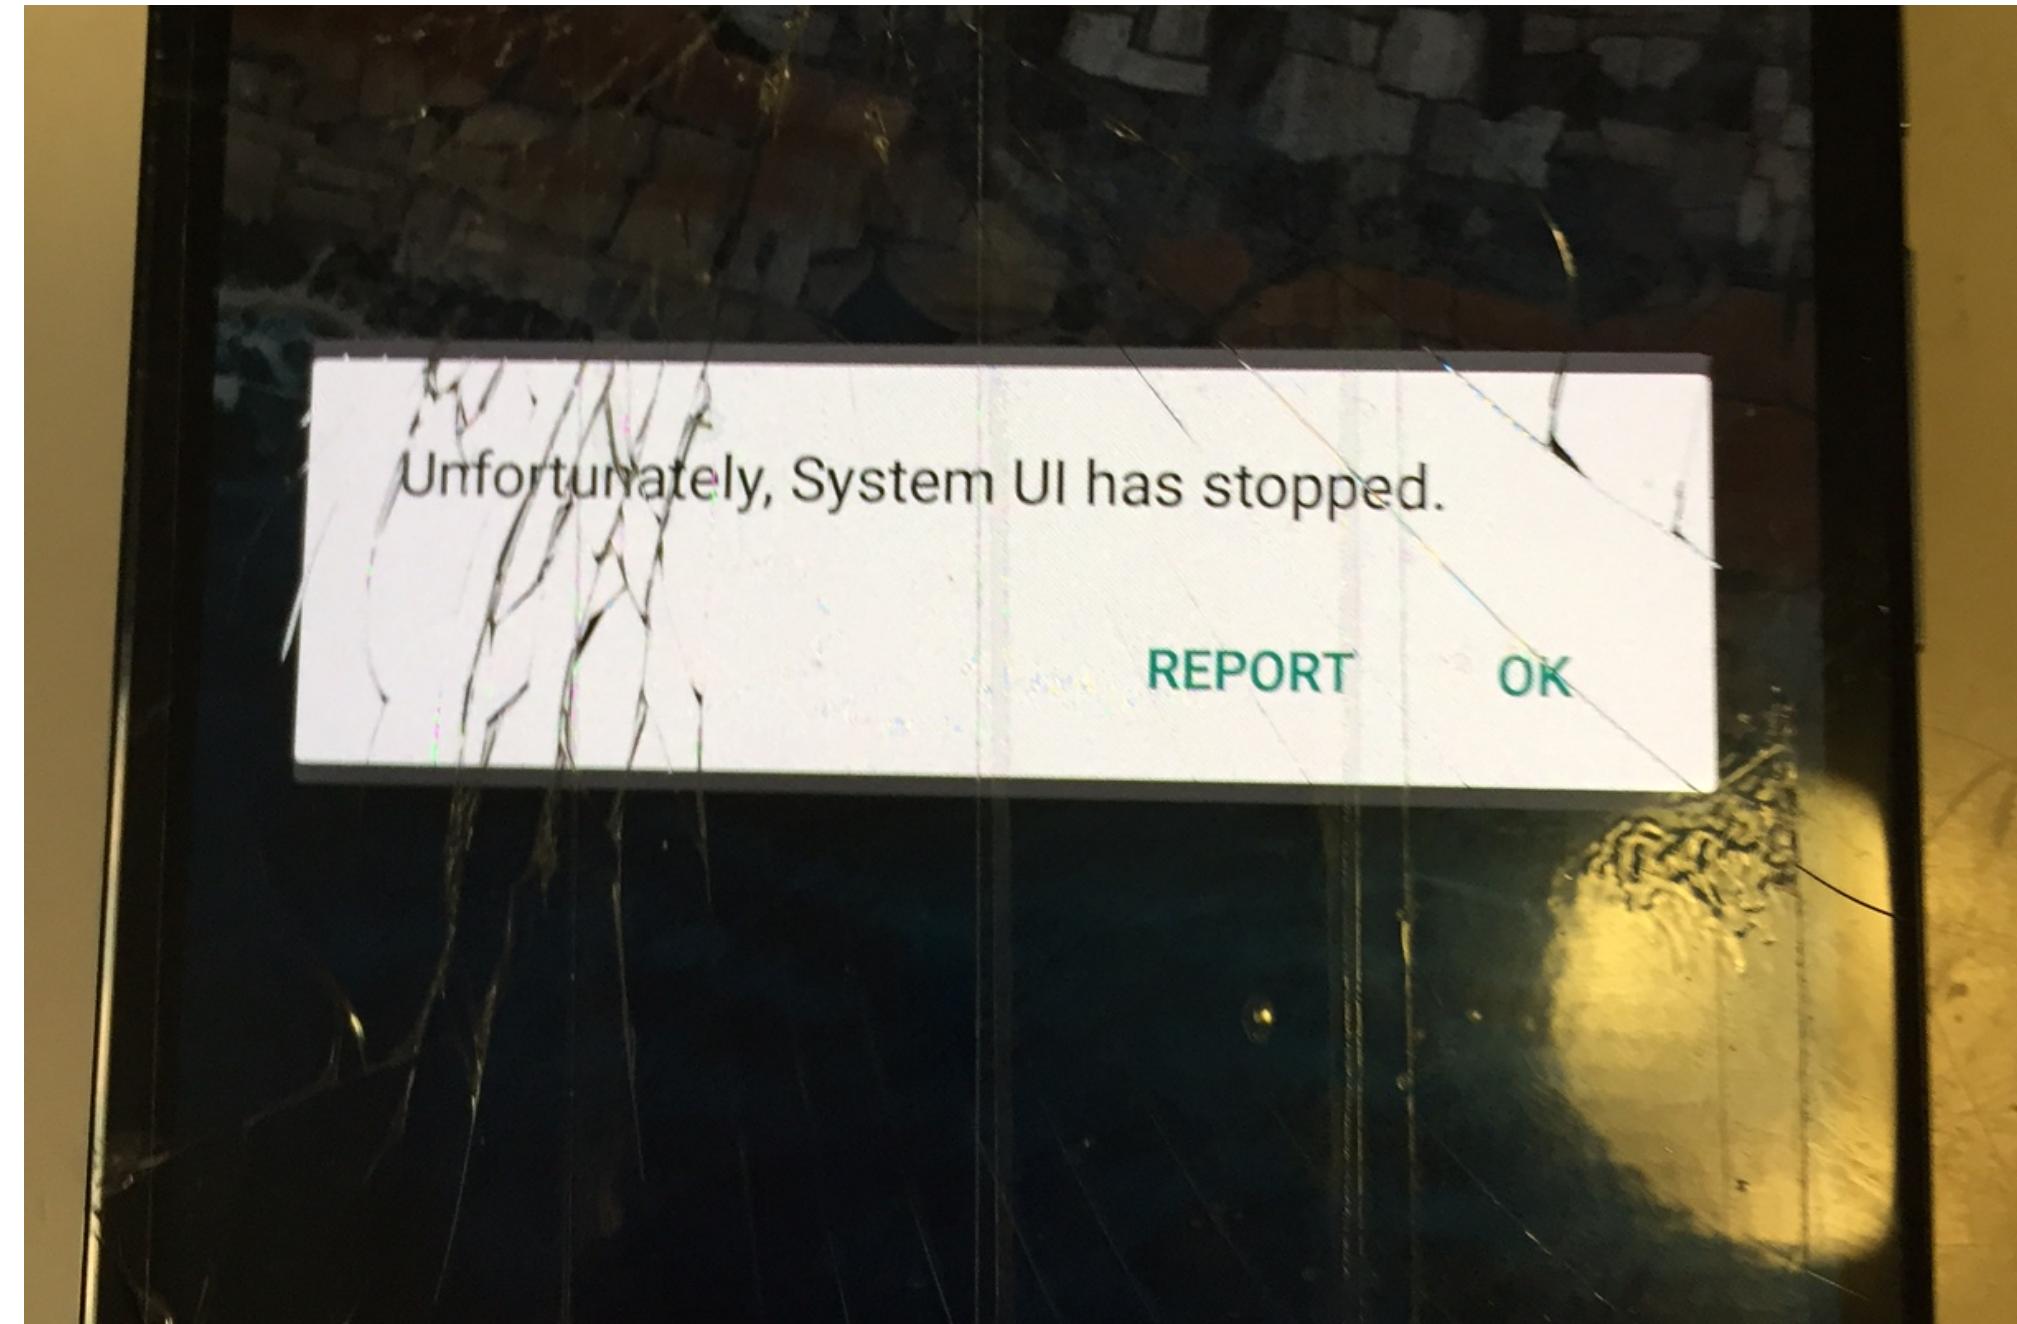


Operating frequency and voltage can be configured  
via memory-mapped registers from software



[https://github.com/0x0atang/clkscrew/blob/master/faultmin\\_SD805/glitch\\_sd805.c](https://github.com/0x0atang/clkscrew/blob/master/faultmin_SD805/glitch_sd805.c)

# First sign of trouble...



Curious crashes

```
root@shamu:/ # logcat | grep "fault"
    Fatal signal 11 (SIGSEGV), code 1, fault addr 0x0 in tid 507 (uprimebenchmark)
    signal 11 (SIGSEGV), code 1 (SEGV_MAPERR), fault addr 0x0
    Fatal signal 11 (SIGSEGV), code 2, fault addr 0xa0950e08 in tid 5422 (RenderThread)
    signal 11 (SIGSEGV), code 2 (SEGV_ACCERR), fault addr 0xa0950e08
    Fatal signal 4 (SIGILL), code 1, fault addr 0xa93f5d40 in tid 4953 (RenderThread)
    signal 4 (SIGILL), code 1 (ILL_ILLOPC), fault addr 0xa93f5d40
    signal 6 (SIGABRT), code -6 (SI_TKILL), fault addr -----
```

A photograph of a silver laptop sitting on a bright yellow desk stand. The laptop screen displays a terminal window with green text. A black power cord runs from the laptop to a white electrical outlet on a white wall. To the right of the laptop is a white refrigerator. The floor is made of light-colored tiles.

# Thinking out of the box...

Temperature matters

Do hardware regulators impose limits  
to frequency/voltage changes?

# Freq / Voltage Operating Point Pairs (OPPs)

---

- ✓ Unintended computing behaviors
- ✓ Software-controlled frequency and voltage settings
- ✓ Verify frequency and voltage settings are indeed properly configured

```
Frequency: cat /d/clk/krait0_clk/measure  
Voltage: cat /d/regulator/krait0/voltage
```

# Freq / Voltage Operating Point Pairs (OPPs)



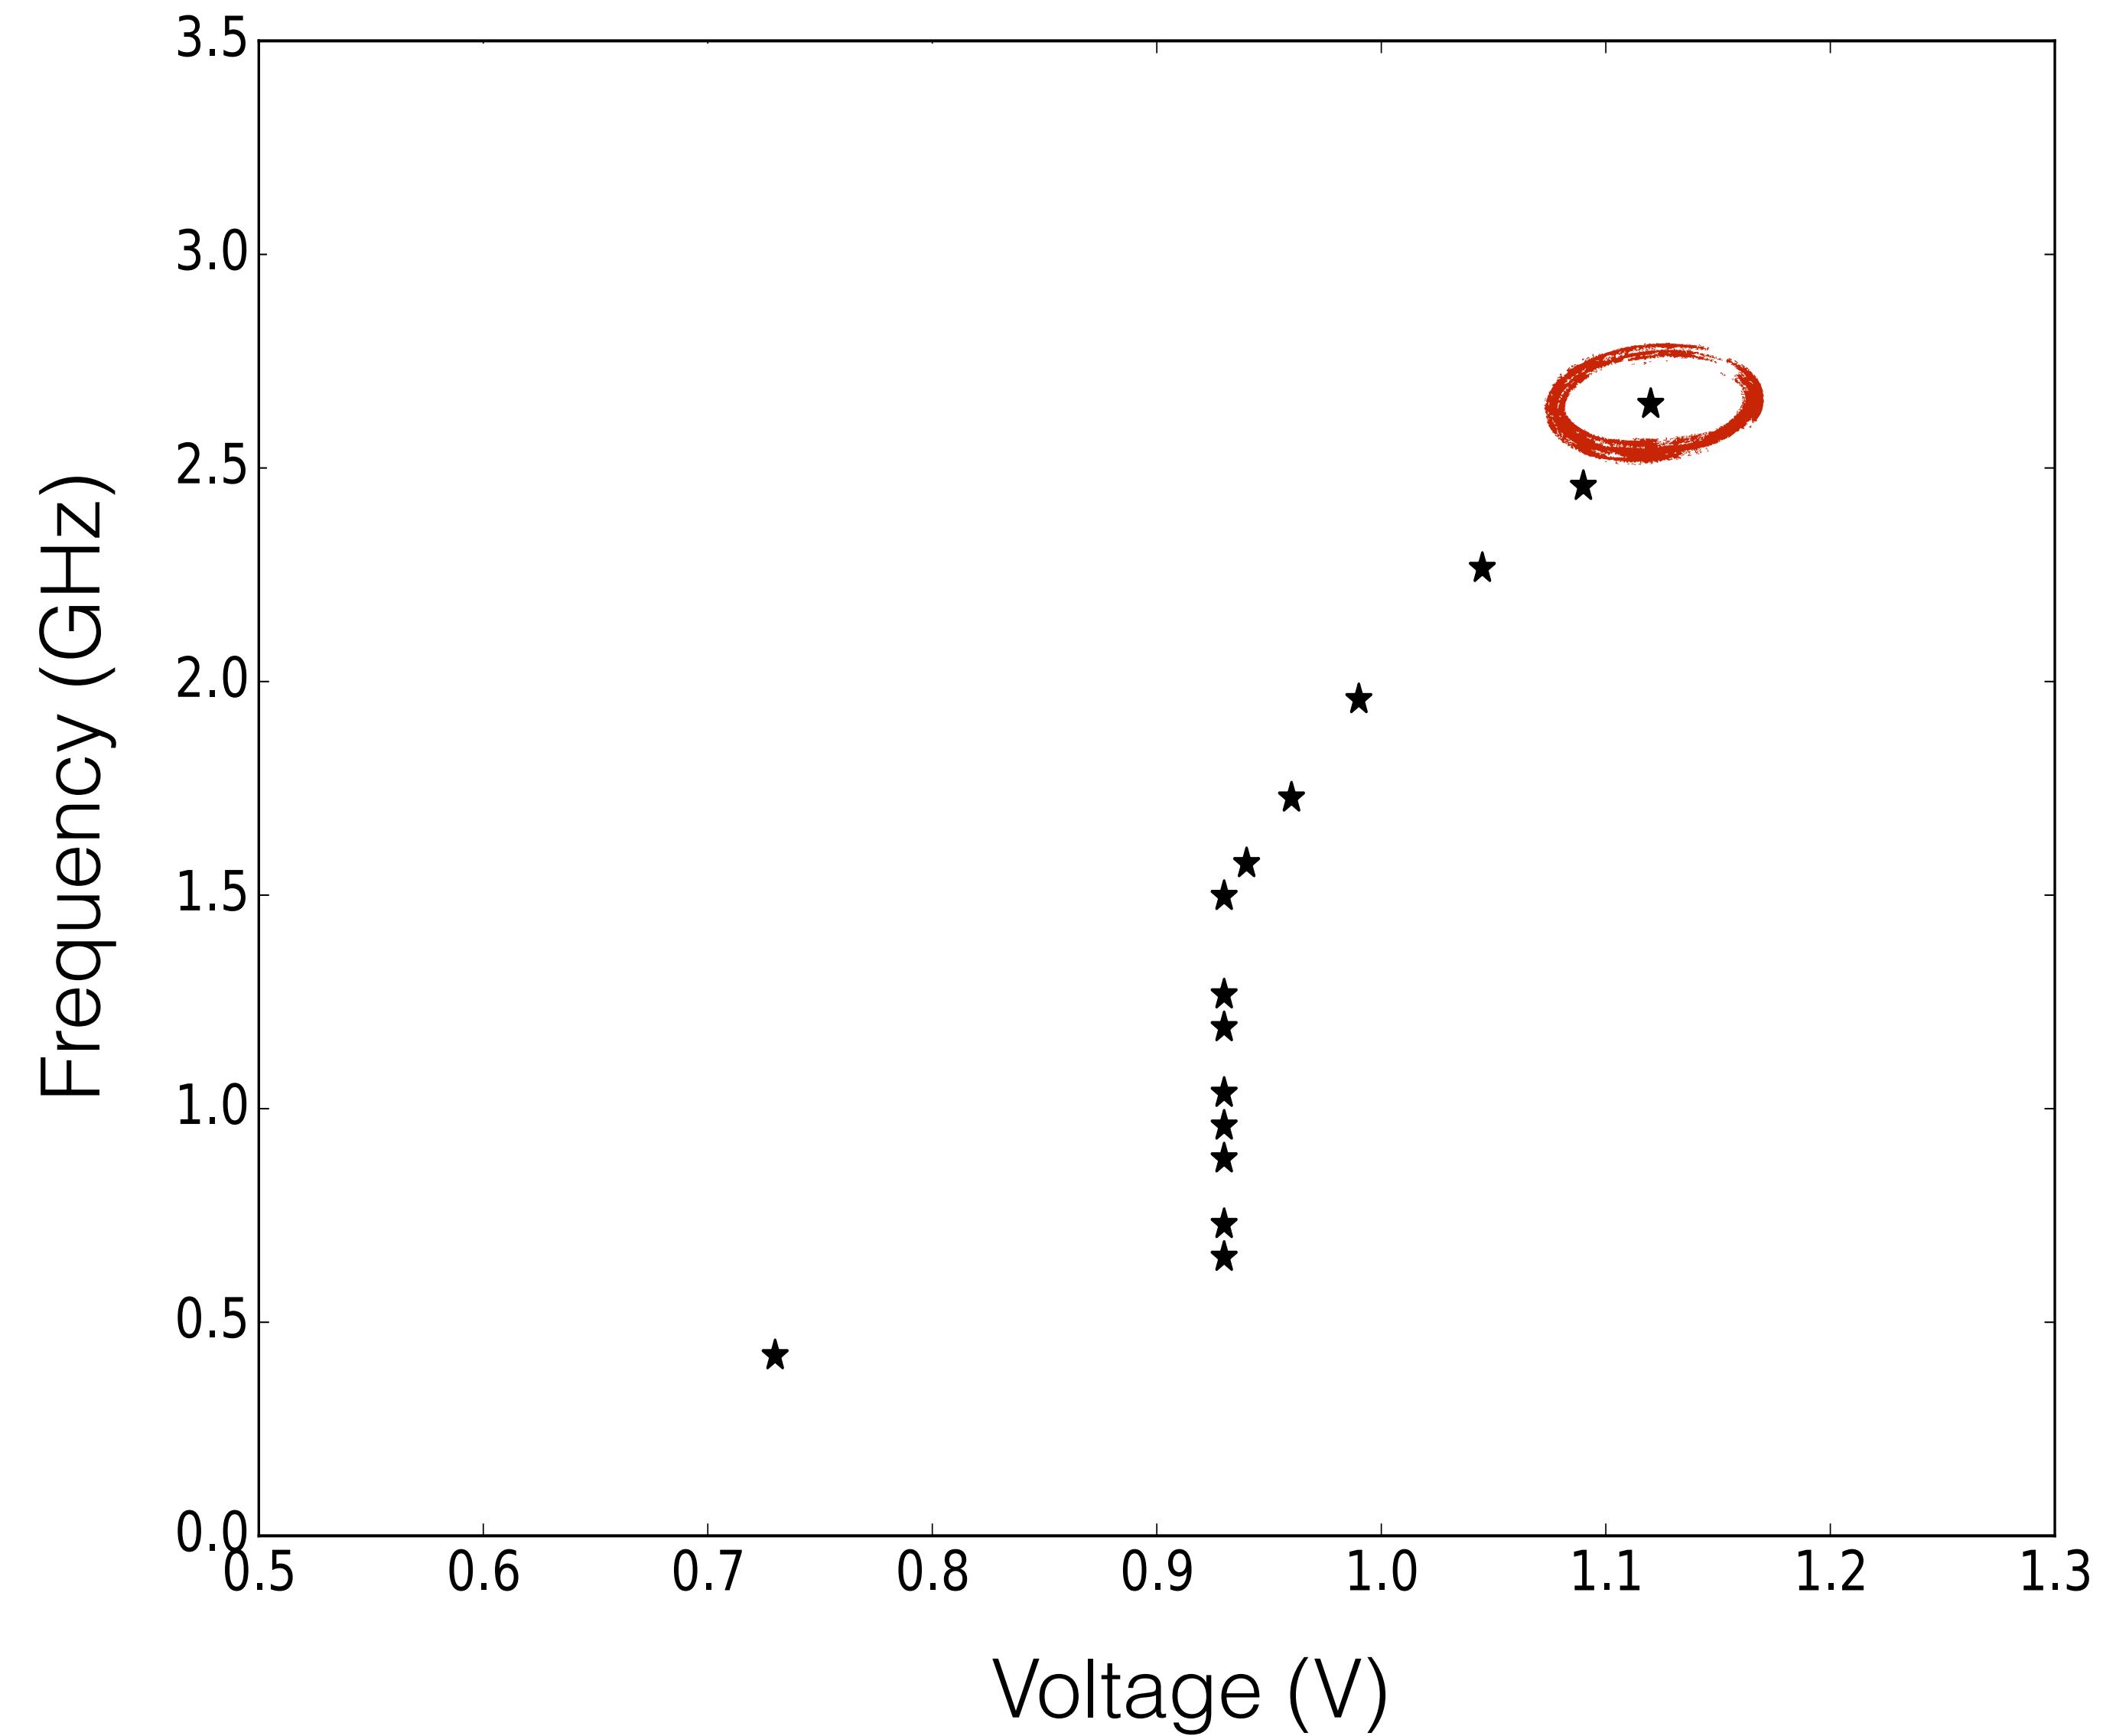
Nexus 6

★★★★★ 791 2

- Android v5(Lollipop)
- Turbo Charging
- 5.94 inch QHD A
- 2.7 GHz Process

Legend:

★★★ Vendor-recommended



# Frequency / Voltage Operating Point Pairs (OPPs)

No safeguard hardware limits

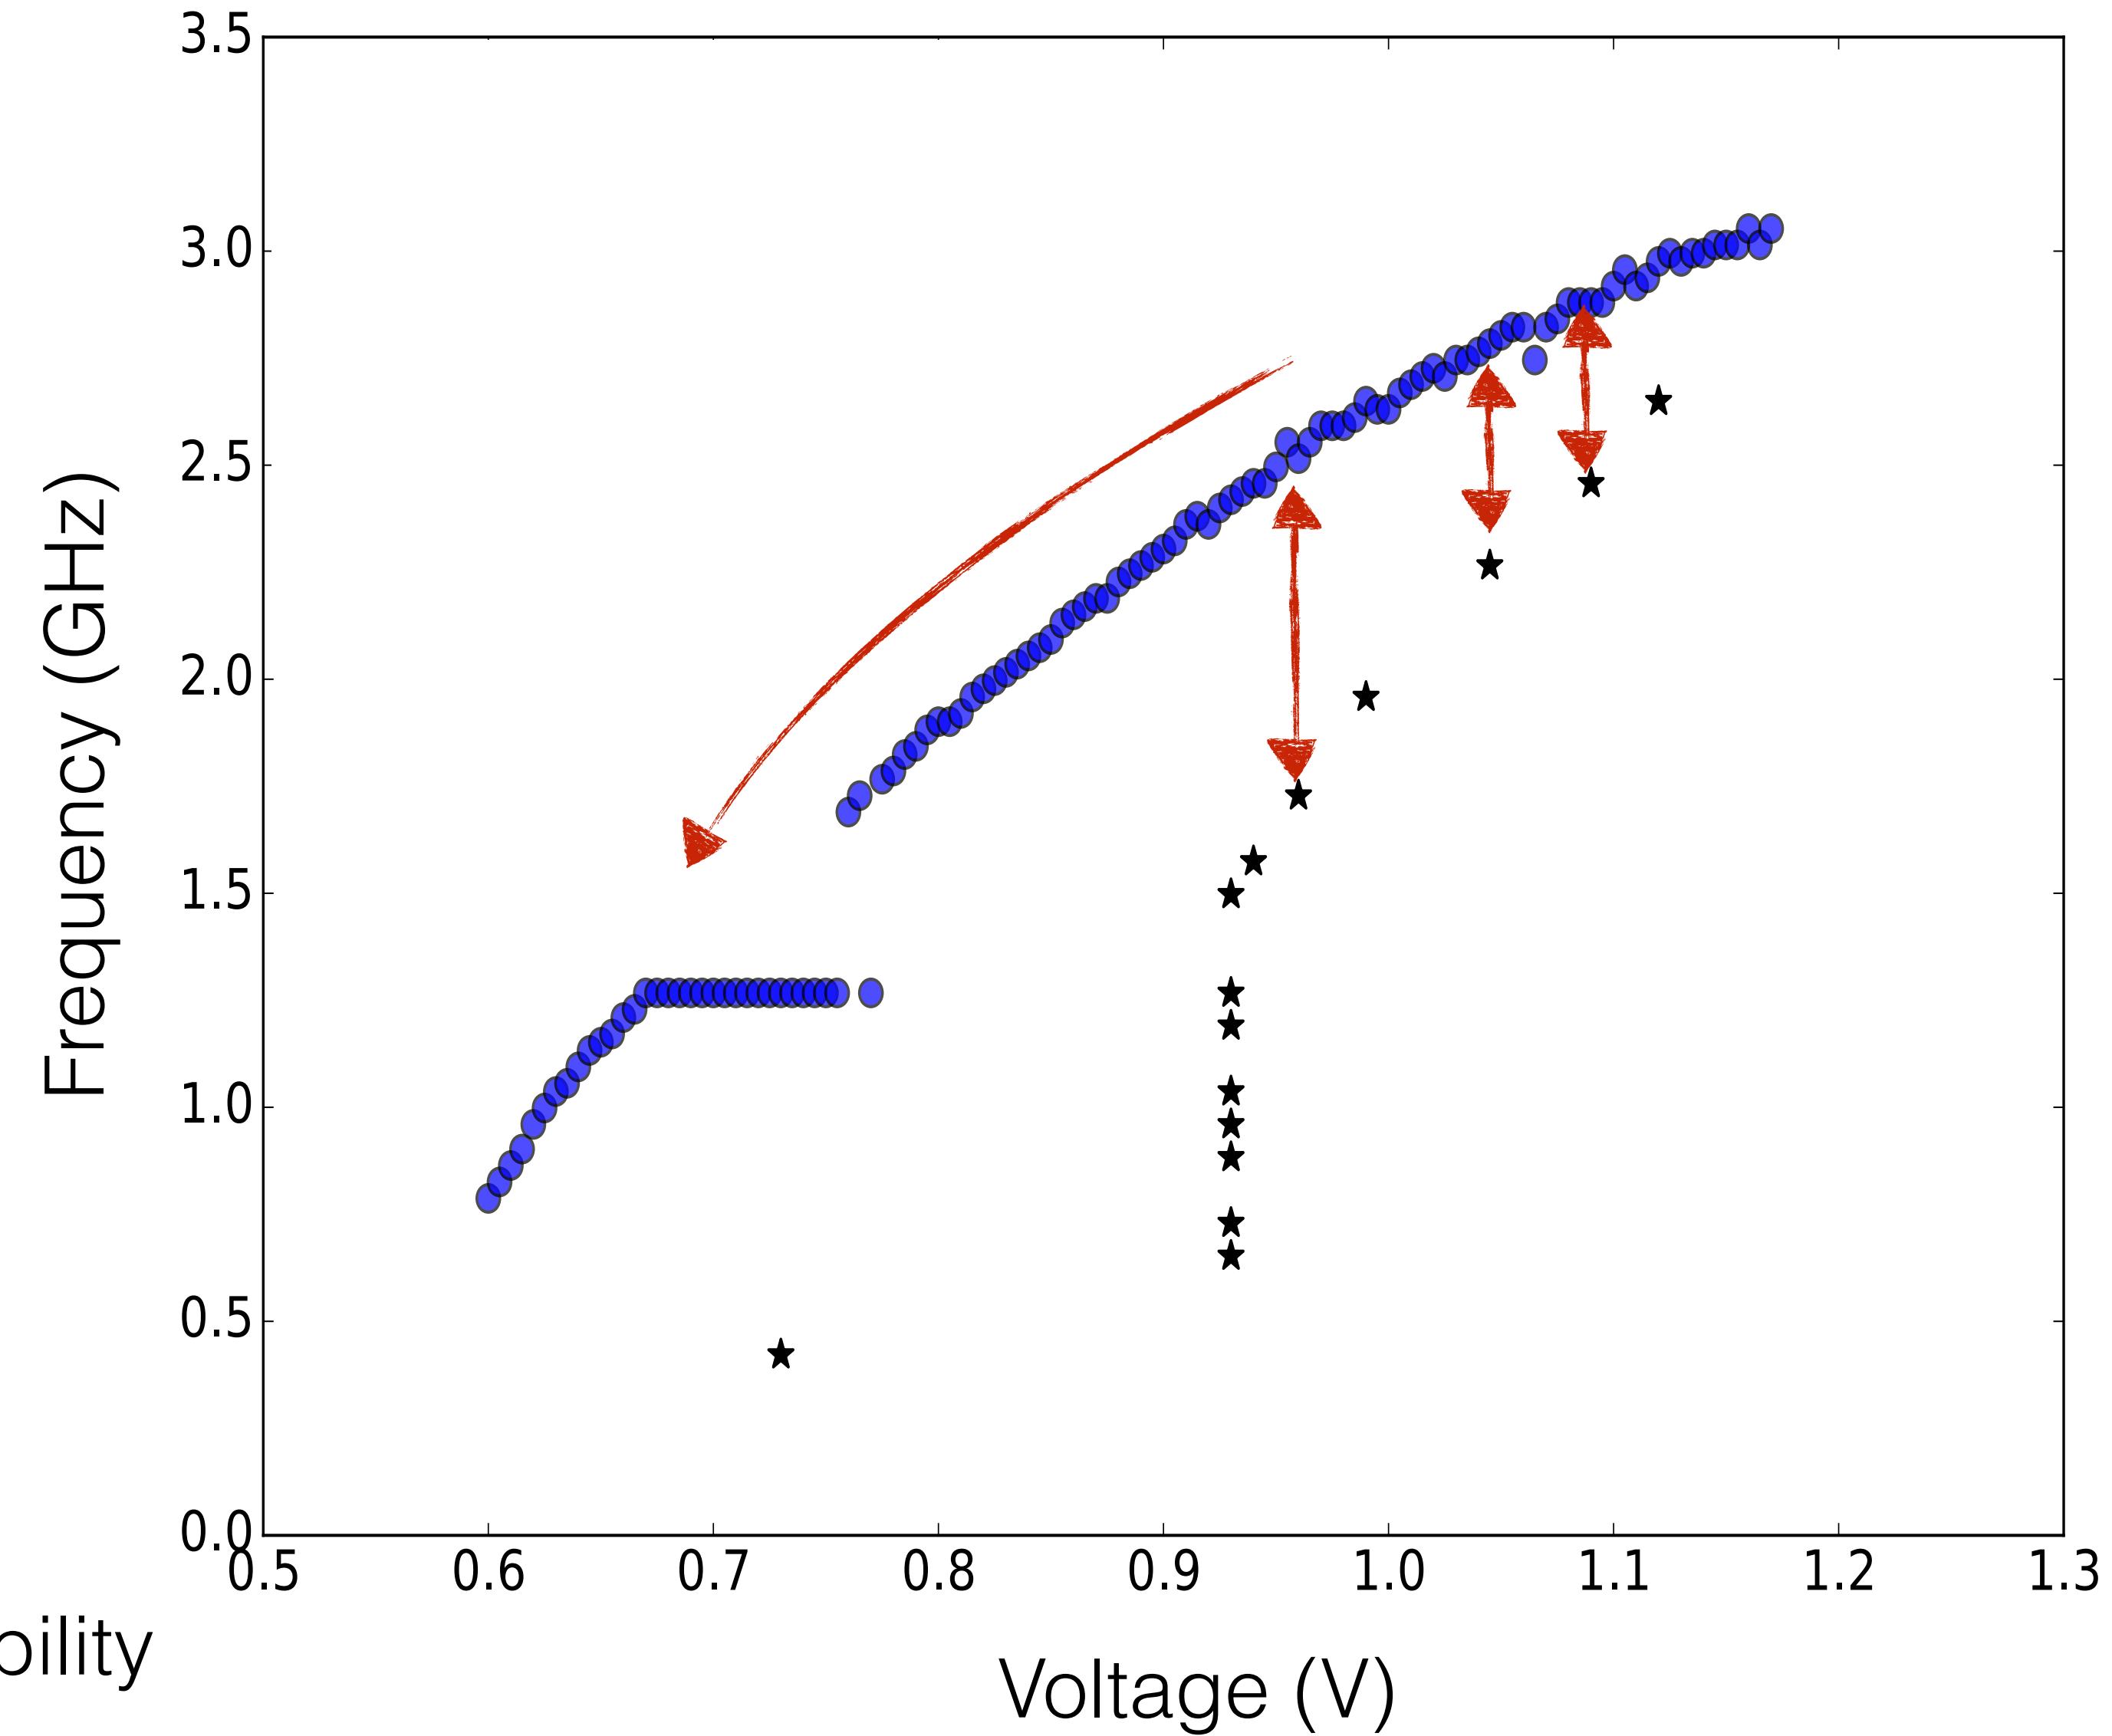
Lower voltage →

Lower minimum required  
frequency to induce instability

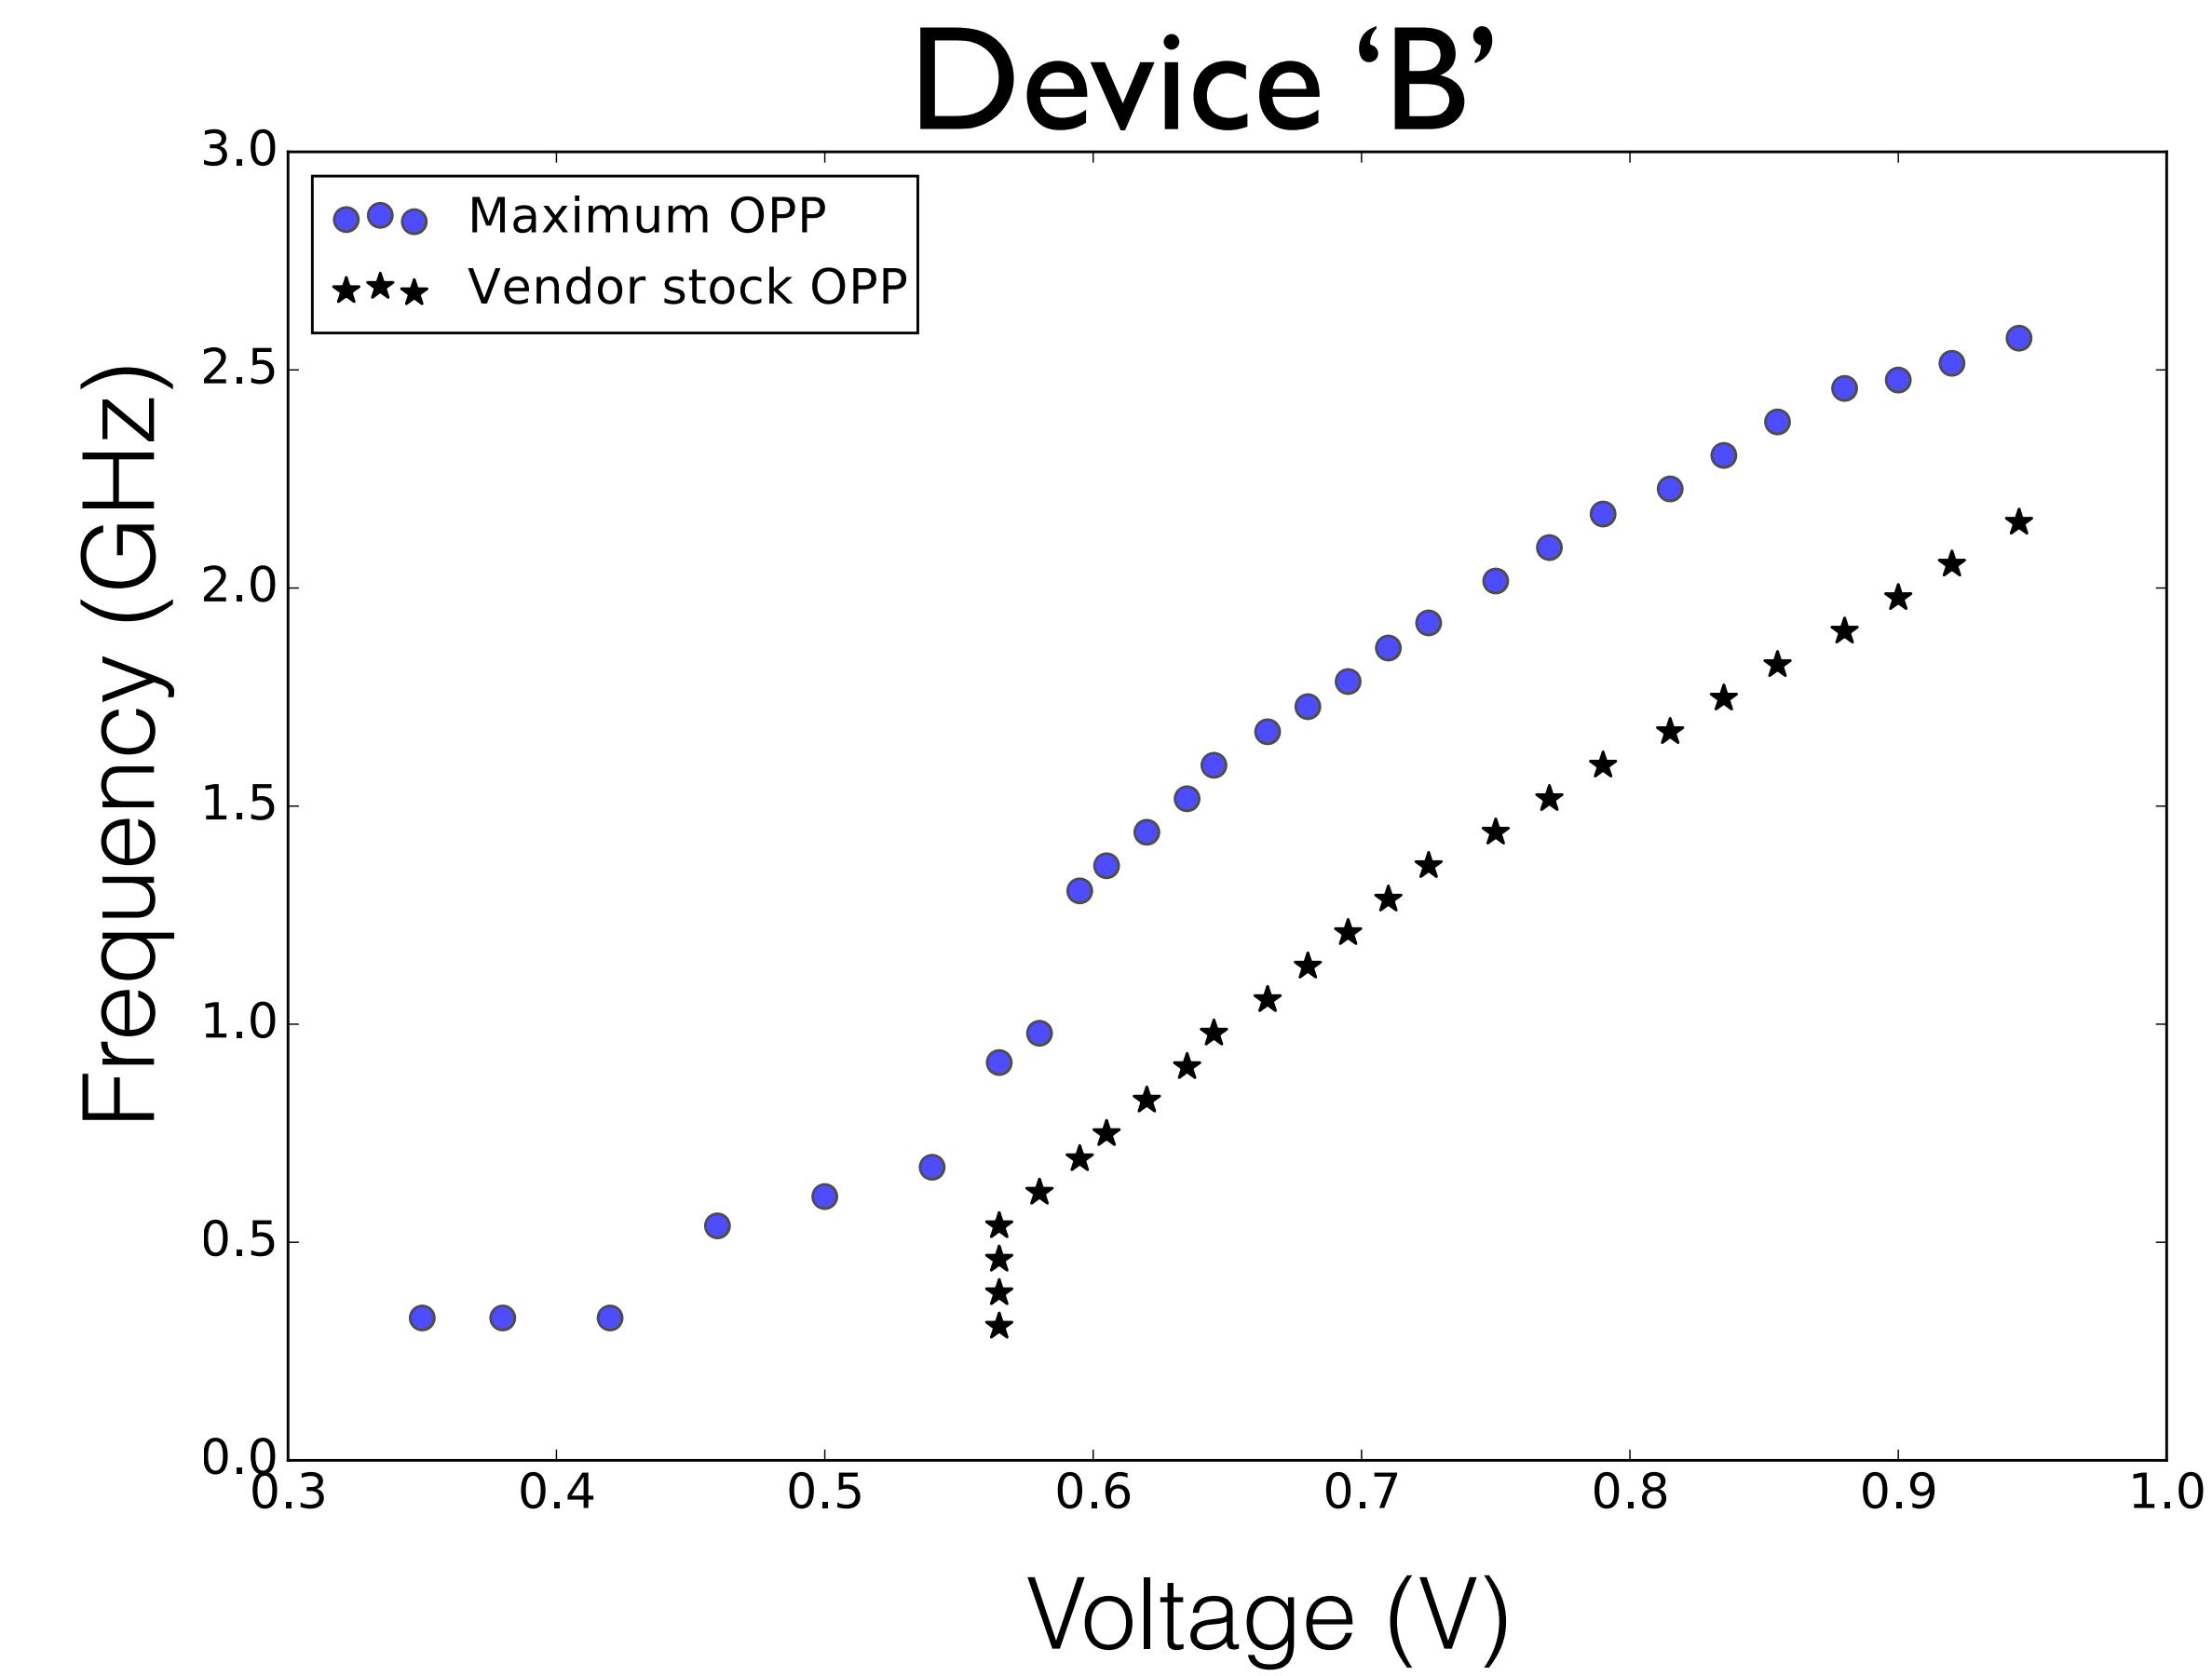
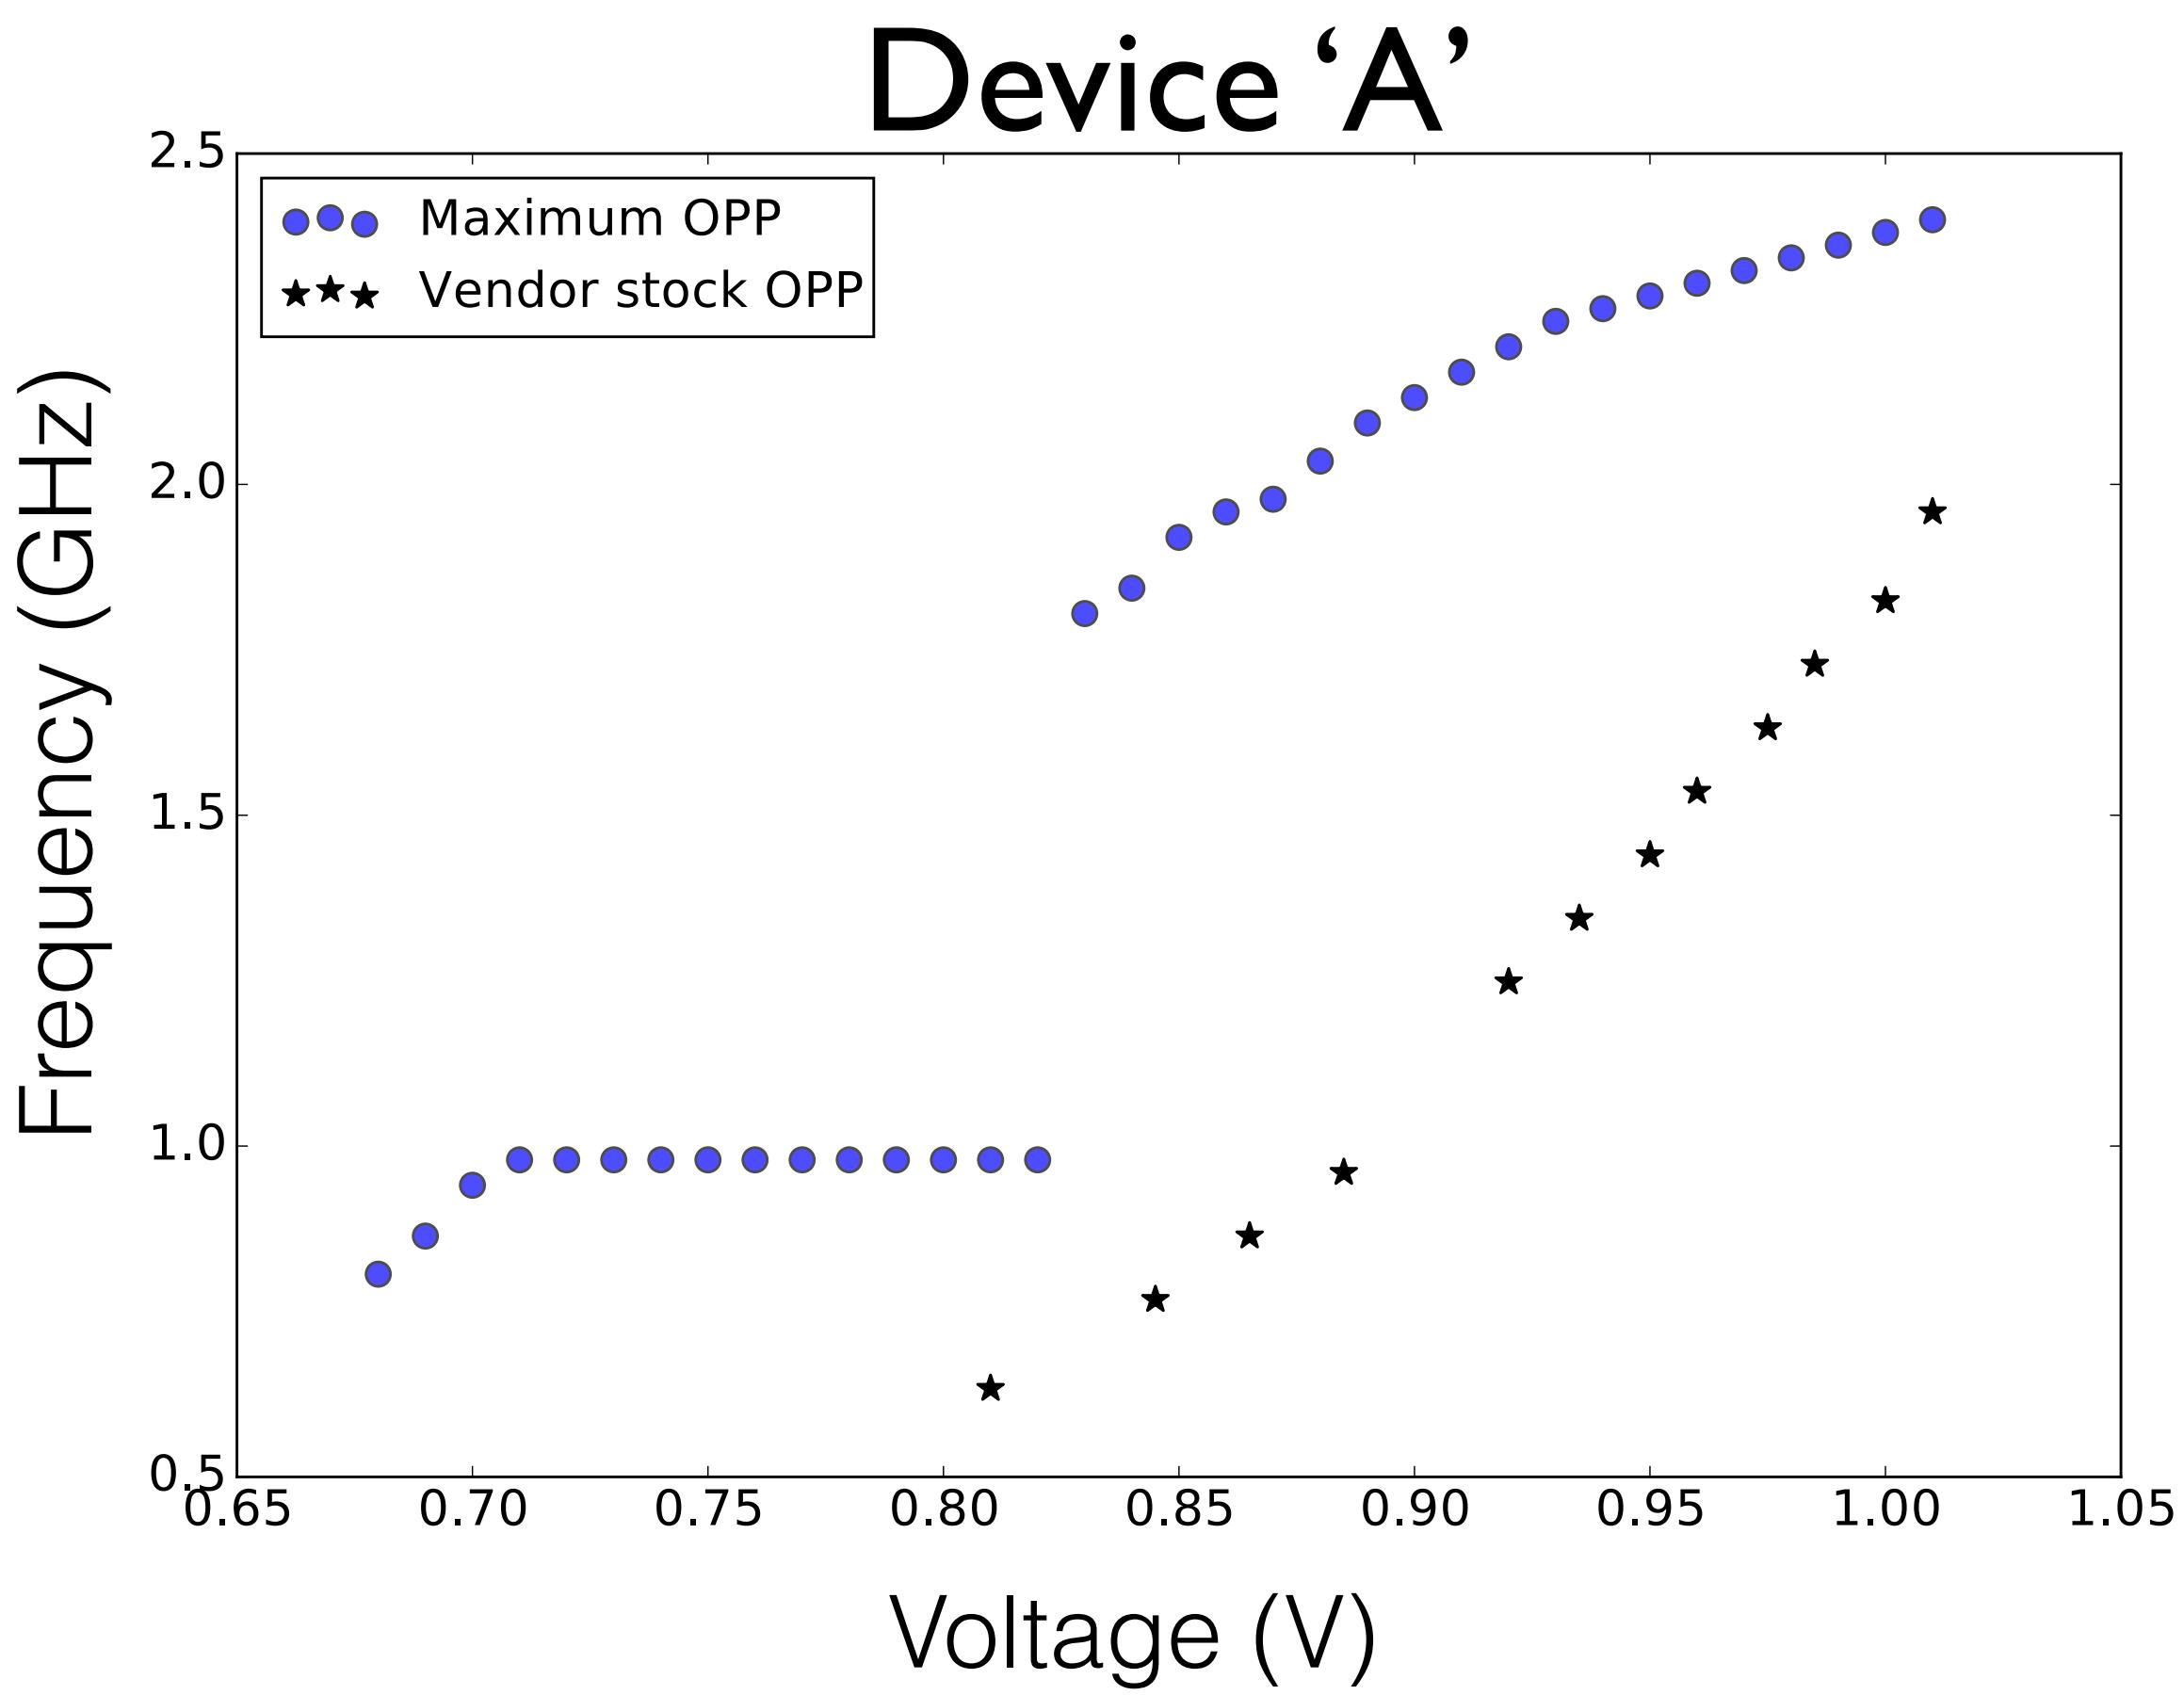
Legend:

★★★ Vendor-recommended

●●● Max OPP reached before instability



# Freq / Voltage Operating Point Pairs (OPPs)

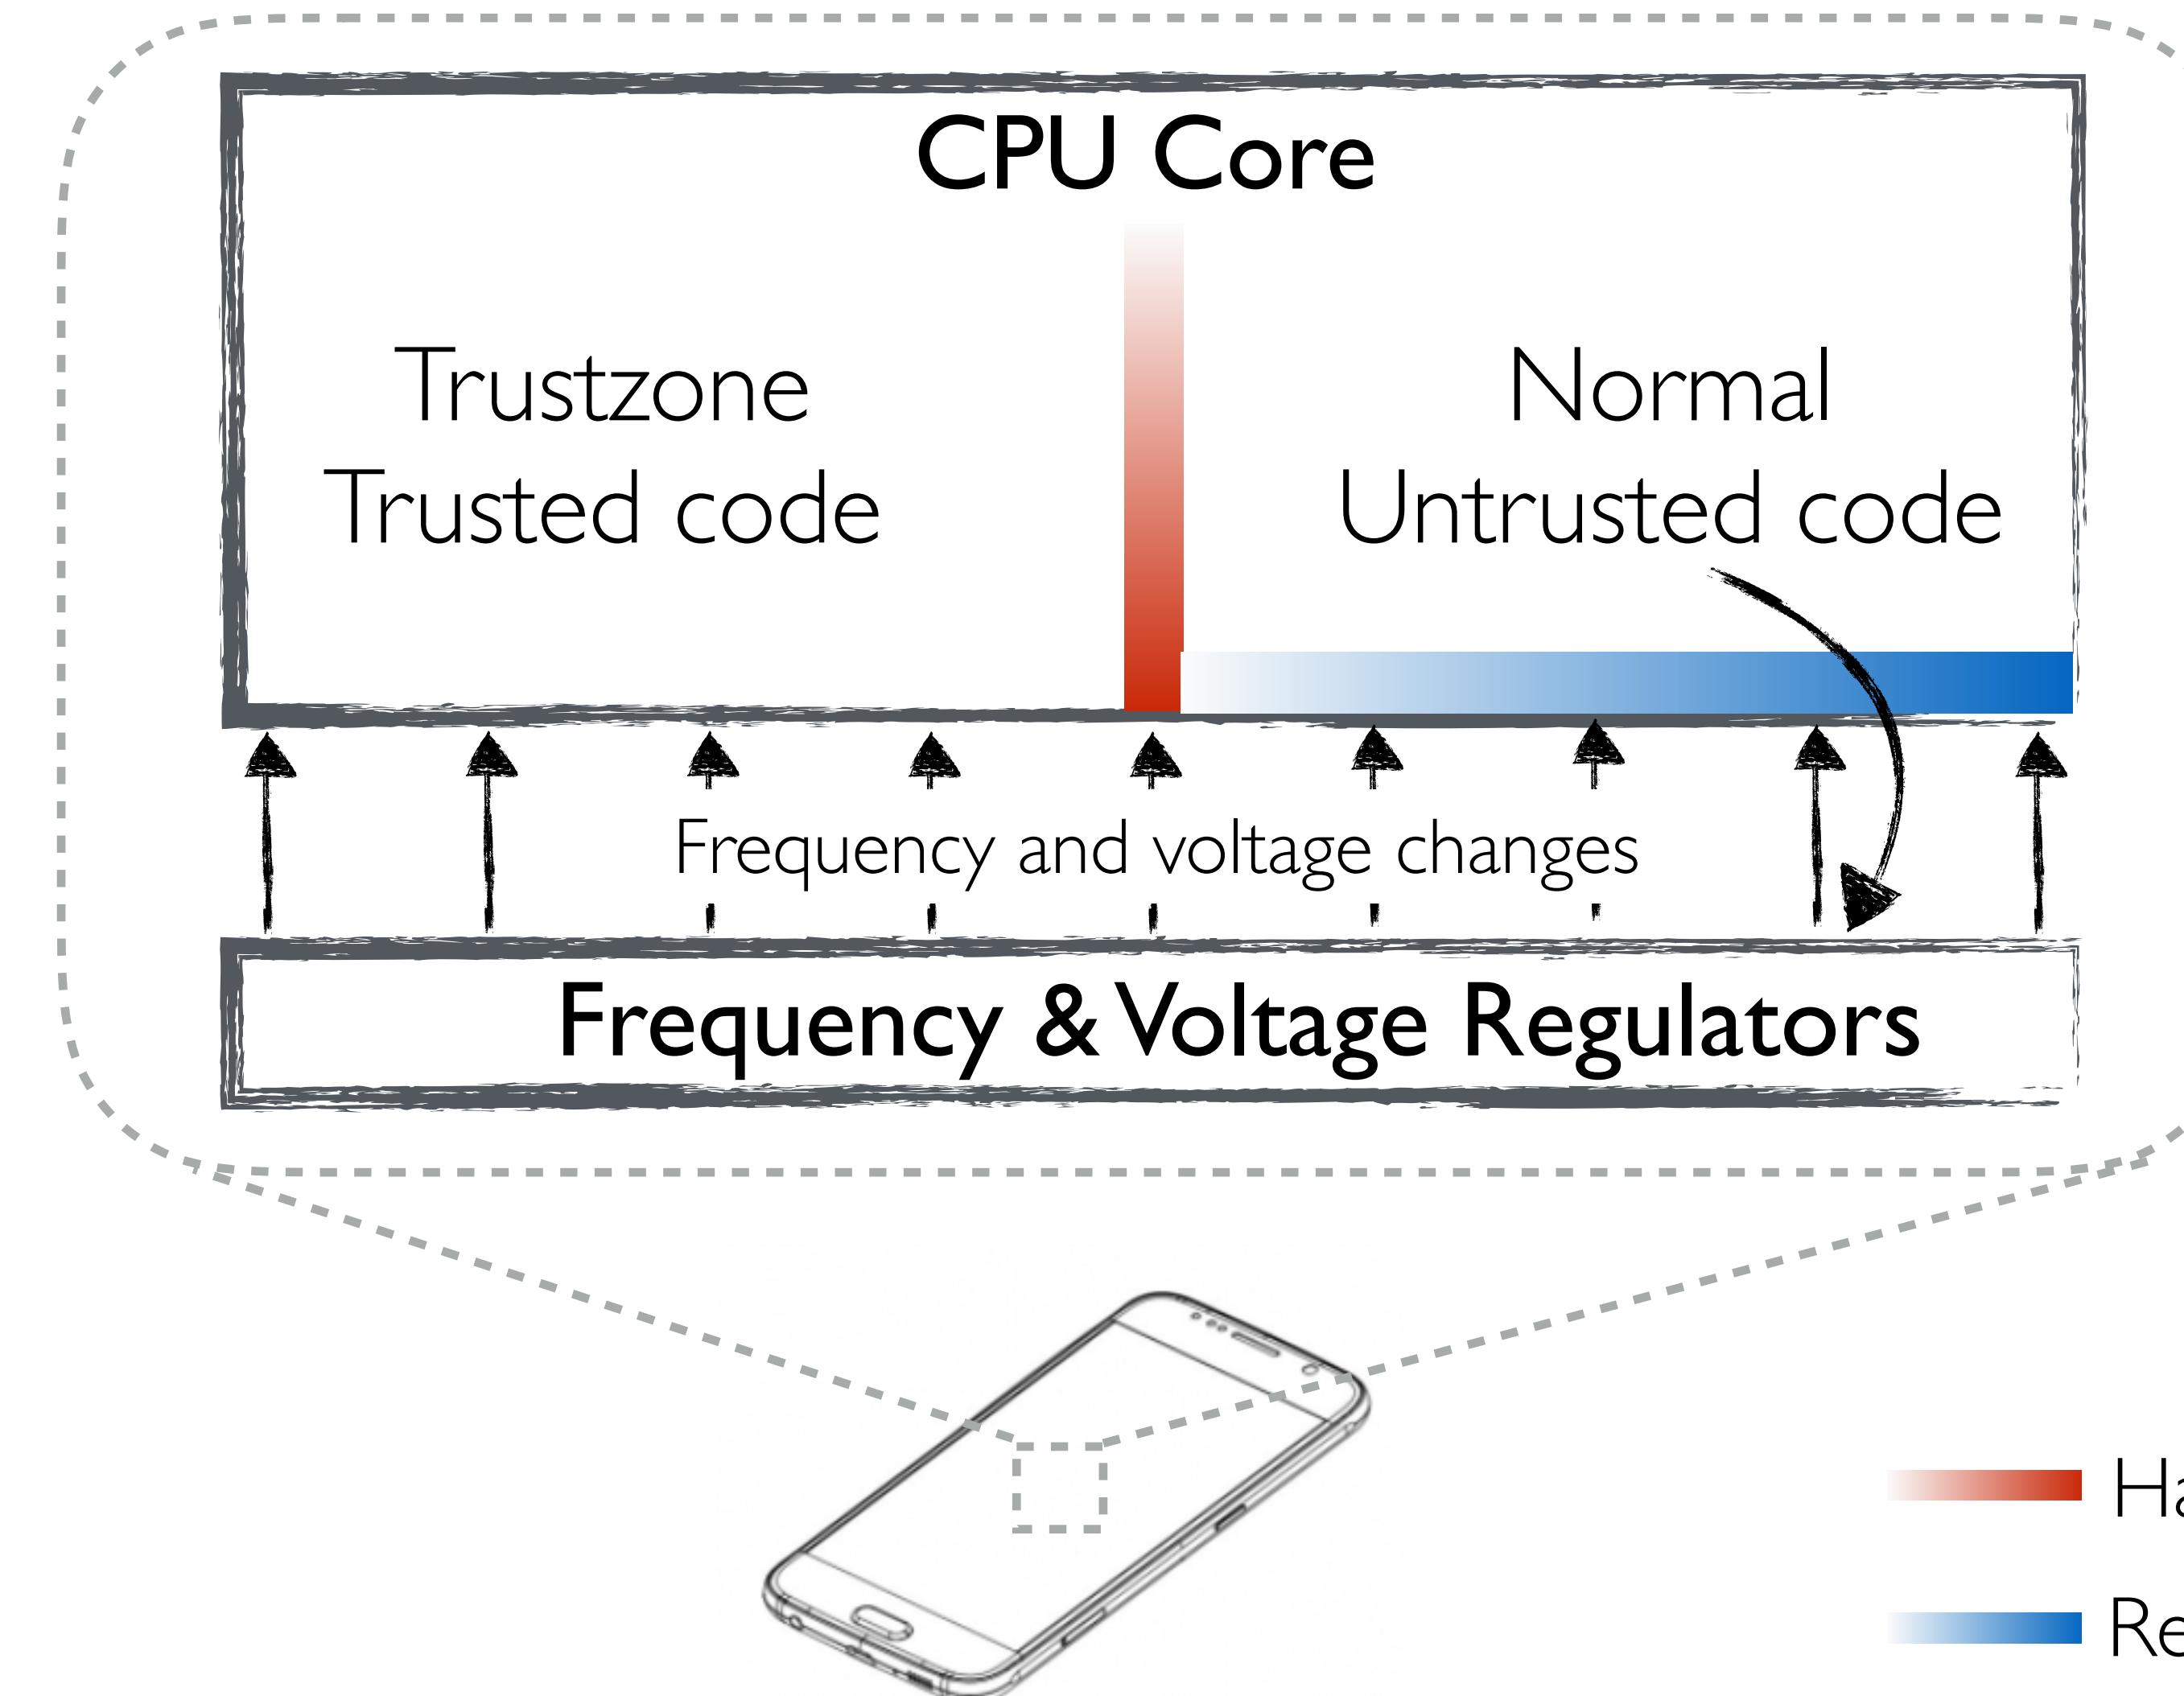


Do regulators operate across security boundaries?



Trusted Execution Environments (TEE)

# Is DVFS Trustzone-Aware? No!



# Outline

---

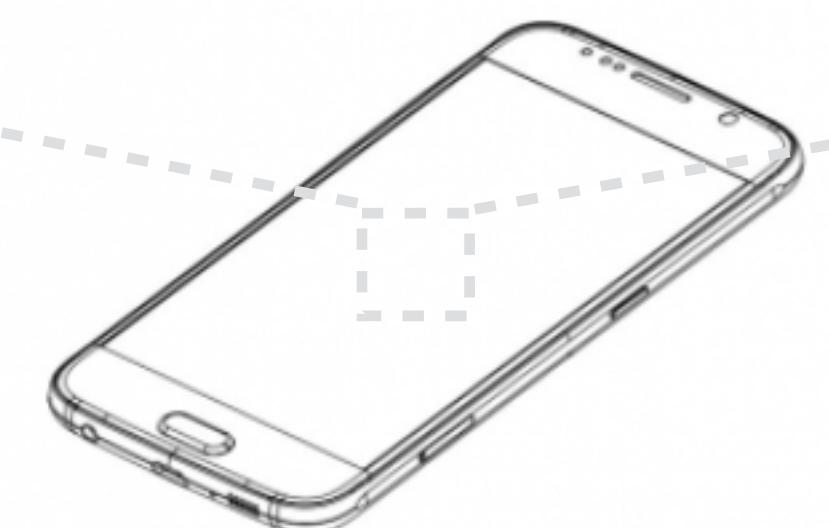
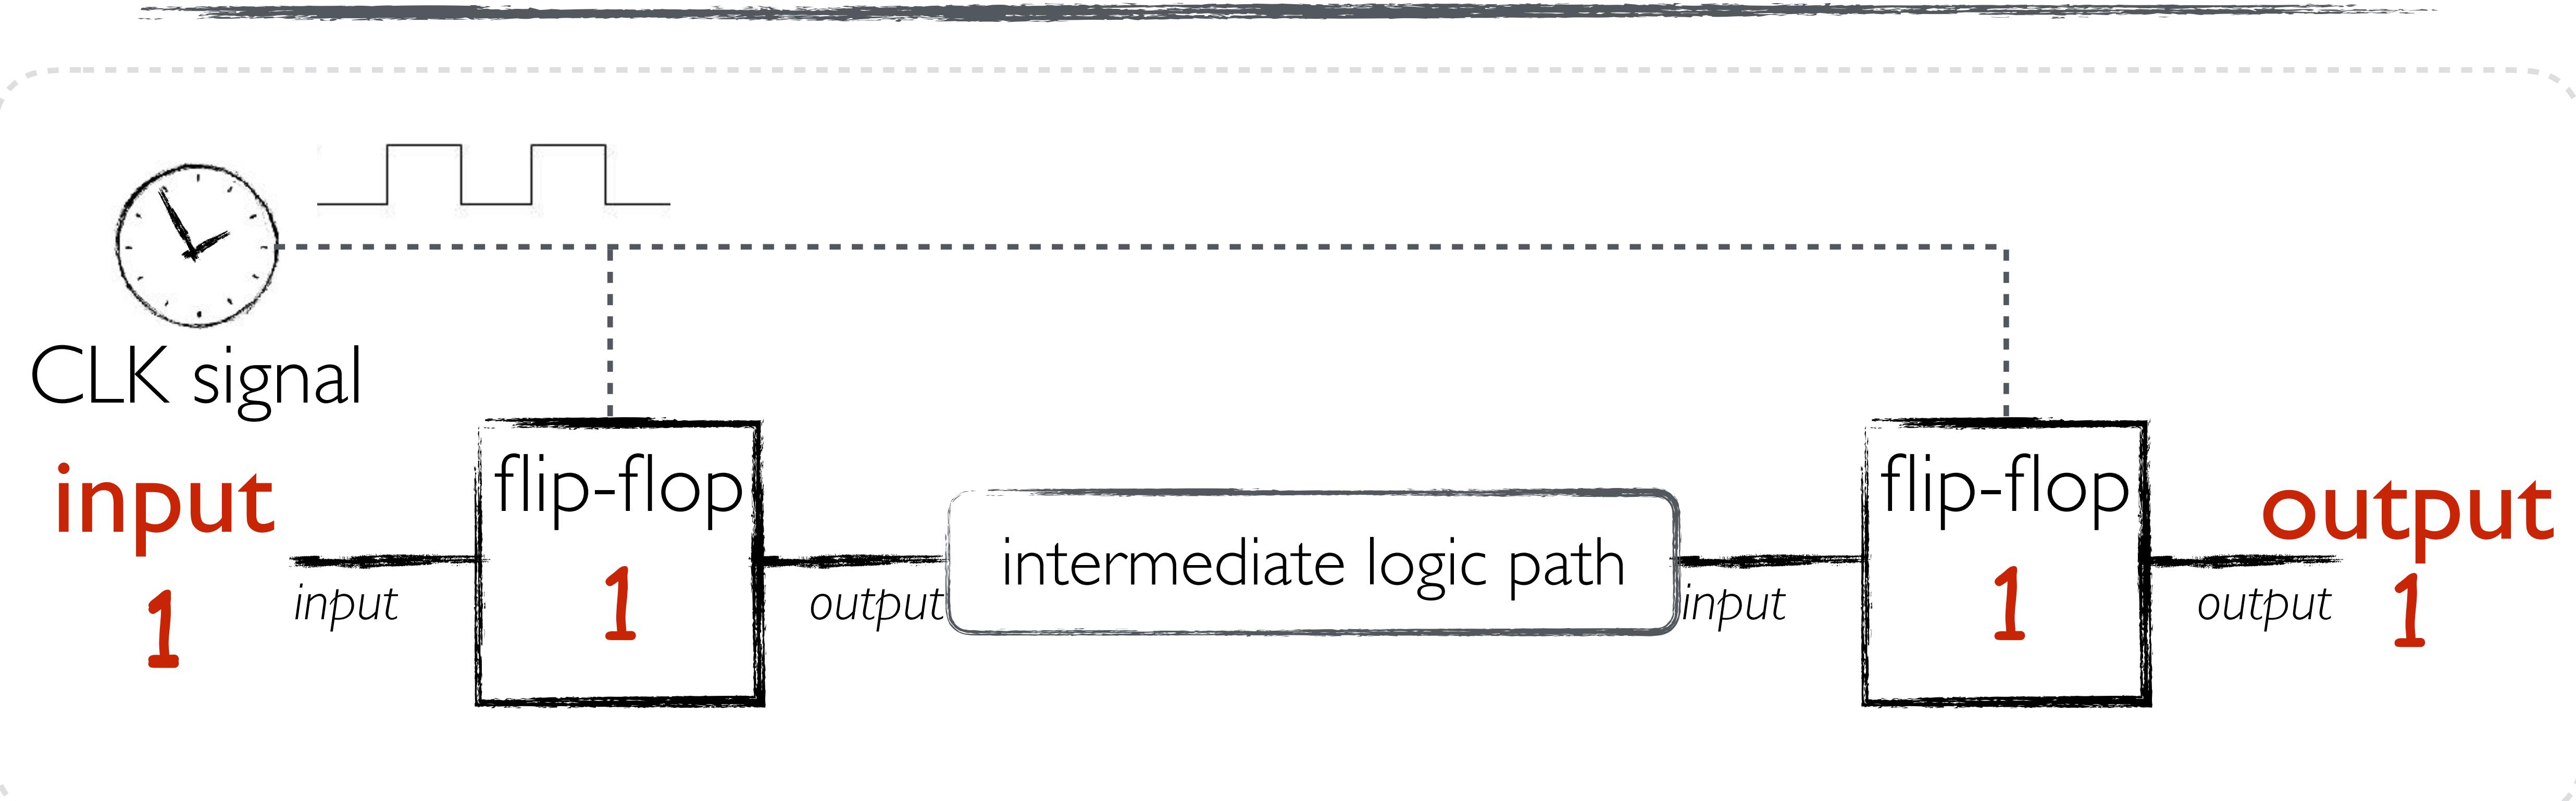
- I. DVFS and Deep Dive into Hardware Regulators
- II. The CLKSCREW Attack**
- III. Trustzone Attack I: Secret AES Key Inference
- IV. Trustzone Attack 2:Tricking RSA Signature Validation
- V. Concluding Remarks

Can we attack Trustzone code execution  
using software-only control of the regulators?

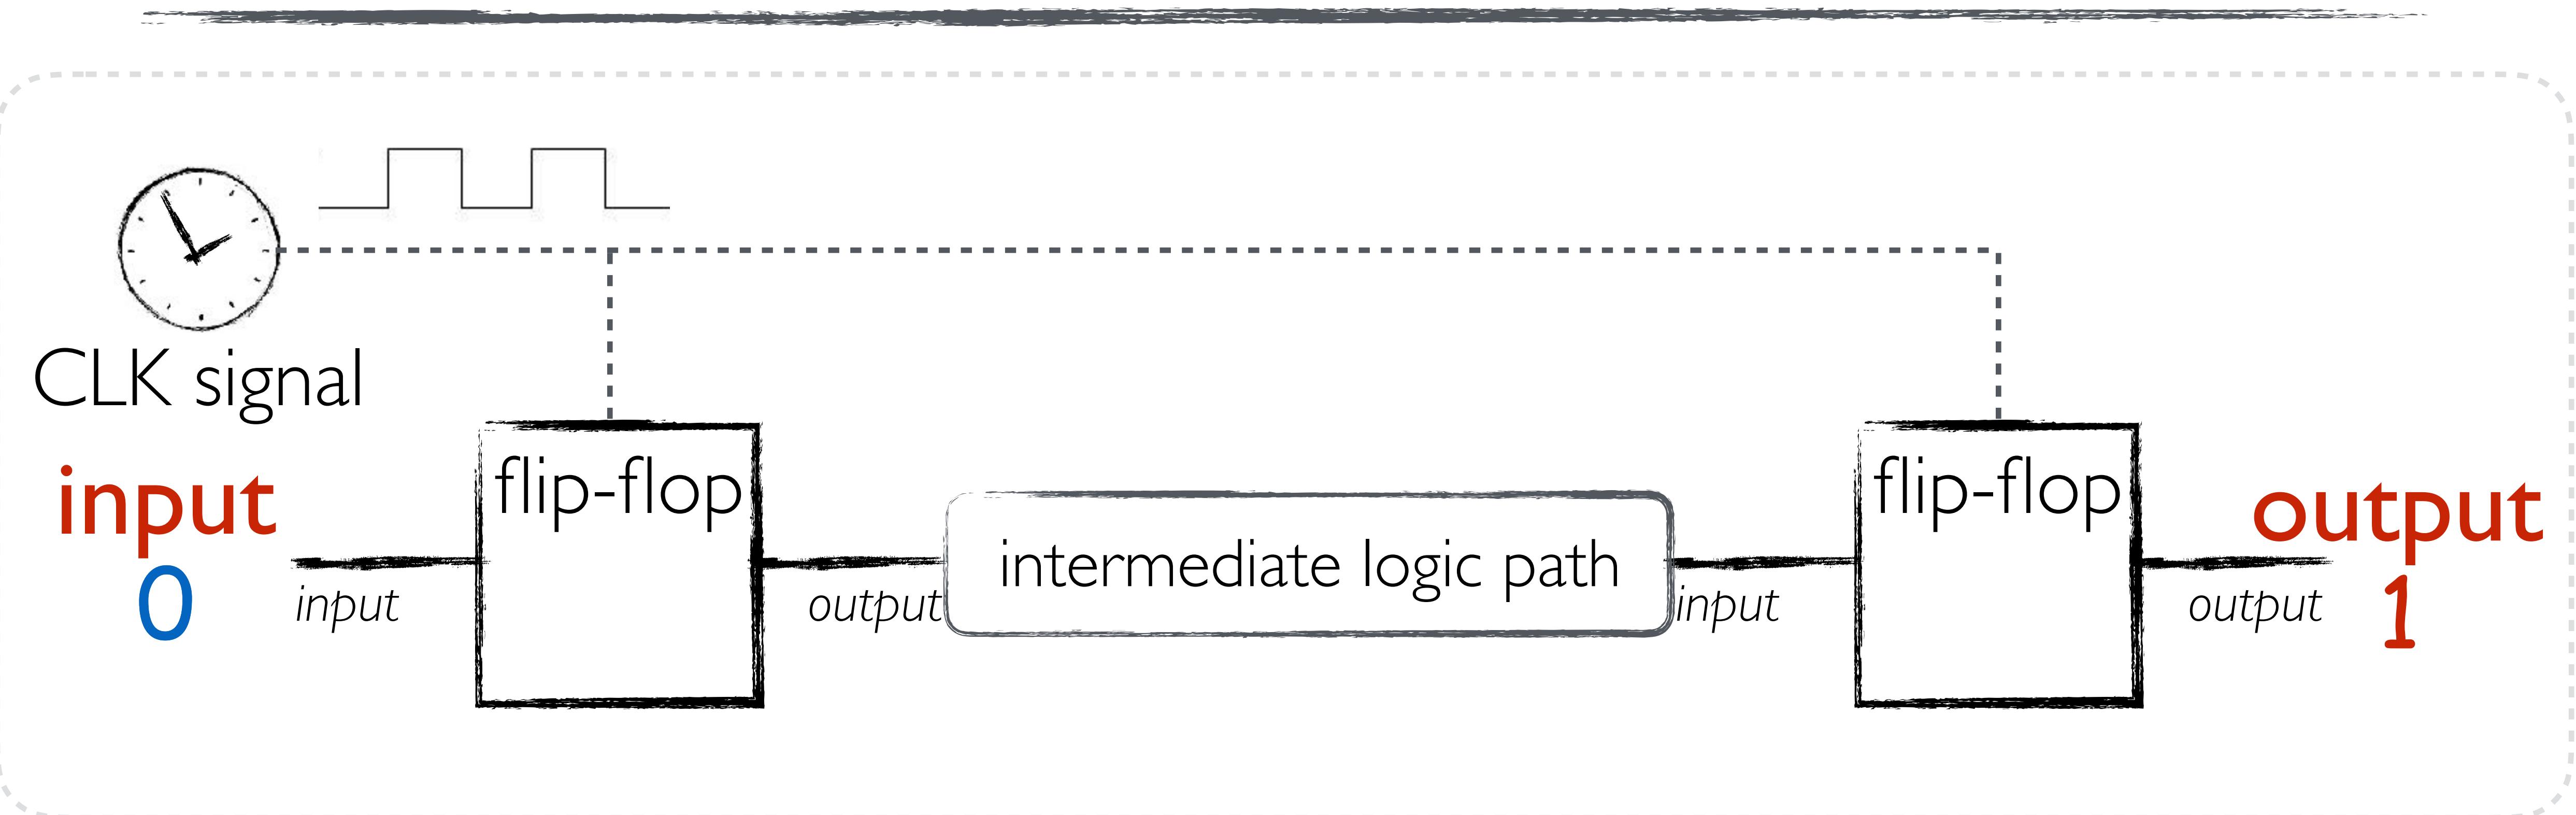
# Induce timing faults

confidentiality  
integrity  
~~availability~~

# How do faults occur (due to over-raising frequency)?



# How do faults occur (due to over-raising frequency)?



How dangerous are faults  
induced by software-based overclocking/undervolting?

# Faults induced by software-based overclocking

Influence **control flow**

Expected: Authentication fails

```
root@shamu:/ # /data/local/tmp/ubench 2
-----
[+] Verifying operation...
[-] Auth *FAIL*
[+] [MAIN]: Exiting and cleaning up

root@shamu:/ # /data/local/tmp/ubench 2
-----
[+] Verifying operation...
[-] Auth *PASS*
[-] Spinning... Ctrl-C to terminate.
```

Core I (Victim)

Runtime Fault Attack: Induce authentication to pass

```
root@shamu:/ # echo 1 > /sys/powerplay/status
root@shamu:/ #
```

Core 2 (Attacker)

# Faults induced by software-based overclocking

Influence **data** flow

Expected: Computation should return (0, 1, 2)

```
root@shamu:/ # /data/local/tmp/compute 3
[+] Performing computations
[-] *PASS*
[-] Iter1: Expected:0 Got:0.0
[-] Iter2: Expected:1 Got:1.0
[-] Iter3: Expected:2 Got:2.0
root@shamu:/ # /data/local/tmp/compute 3
[+] Performing computations
[-] *FAIL*
[-] Iter1: Expected:0 Got:nan
[-] Iter2: Expected:1 Got:1.0
[-] Iter3: Expected:2 Got:2.0
root@shamu:/ # 
```

Core 1 (Victim)

Runtime Fault Attack: Corrupt result to (nan, 1, 2)

```
root@shamu:/ # echo 1 > /sys/powerplay/status
root@shamu:/ # 
```

Core 2 (Attacker)

# CLKSCREW Challenges & Solutions

---

#1: Regulator operating limits

#2: Self-containment within same device

#3: Noisy complex OS environment

#4: Precise timing

#5: Fine-grained timing resolution

# CLKSCREW Challenges & Solutions

#1: Regulator operating limits

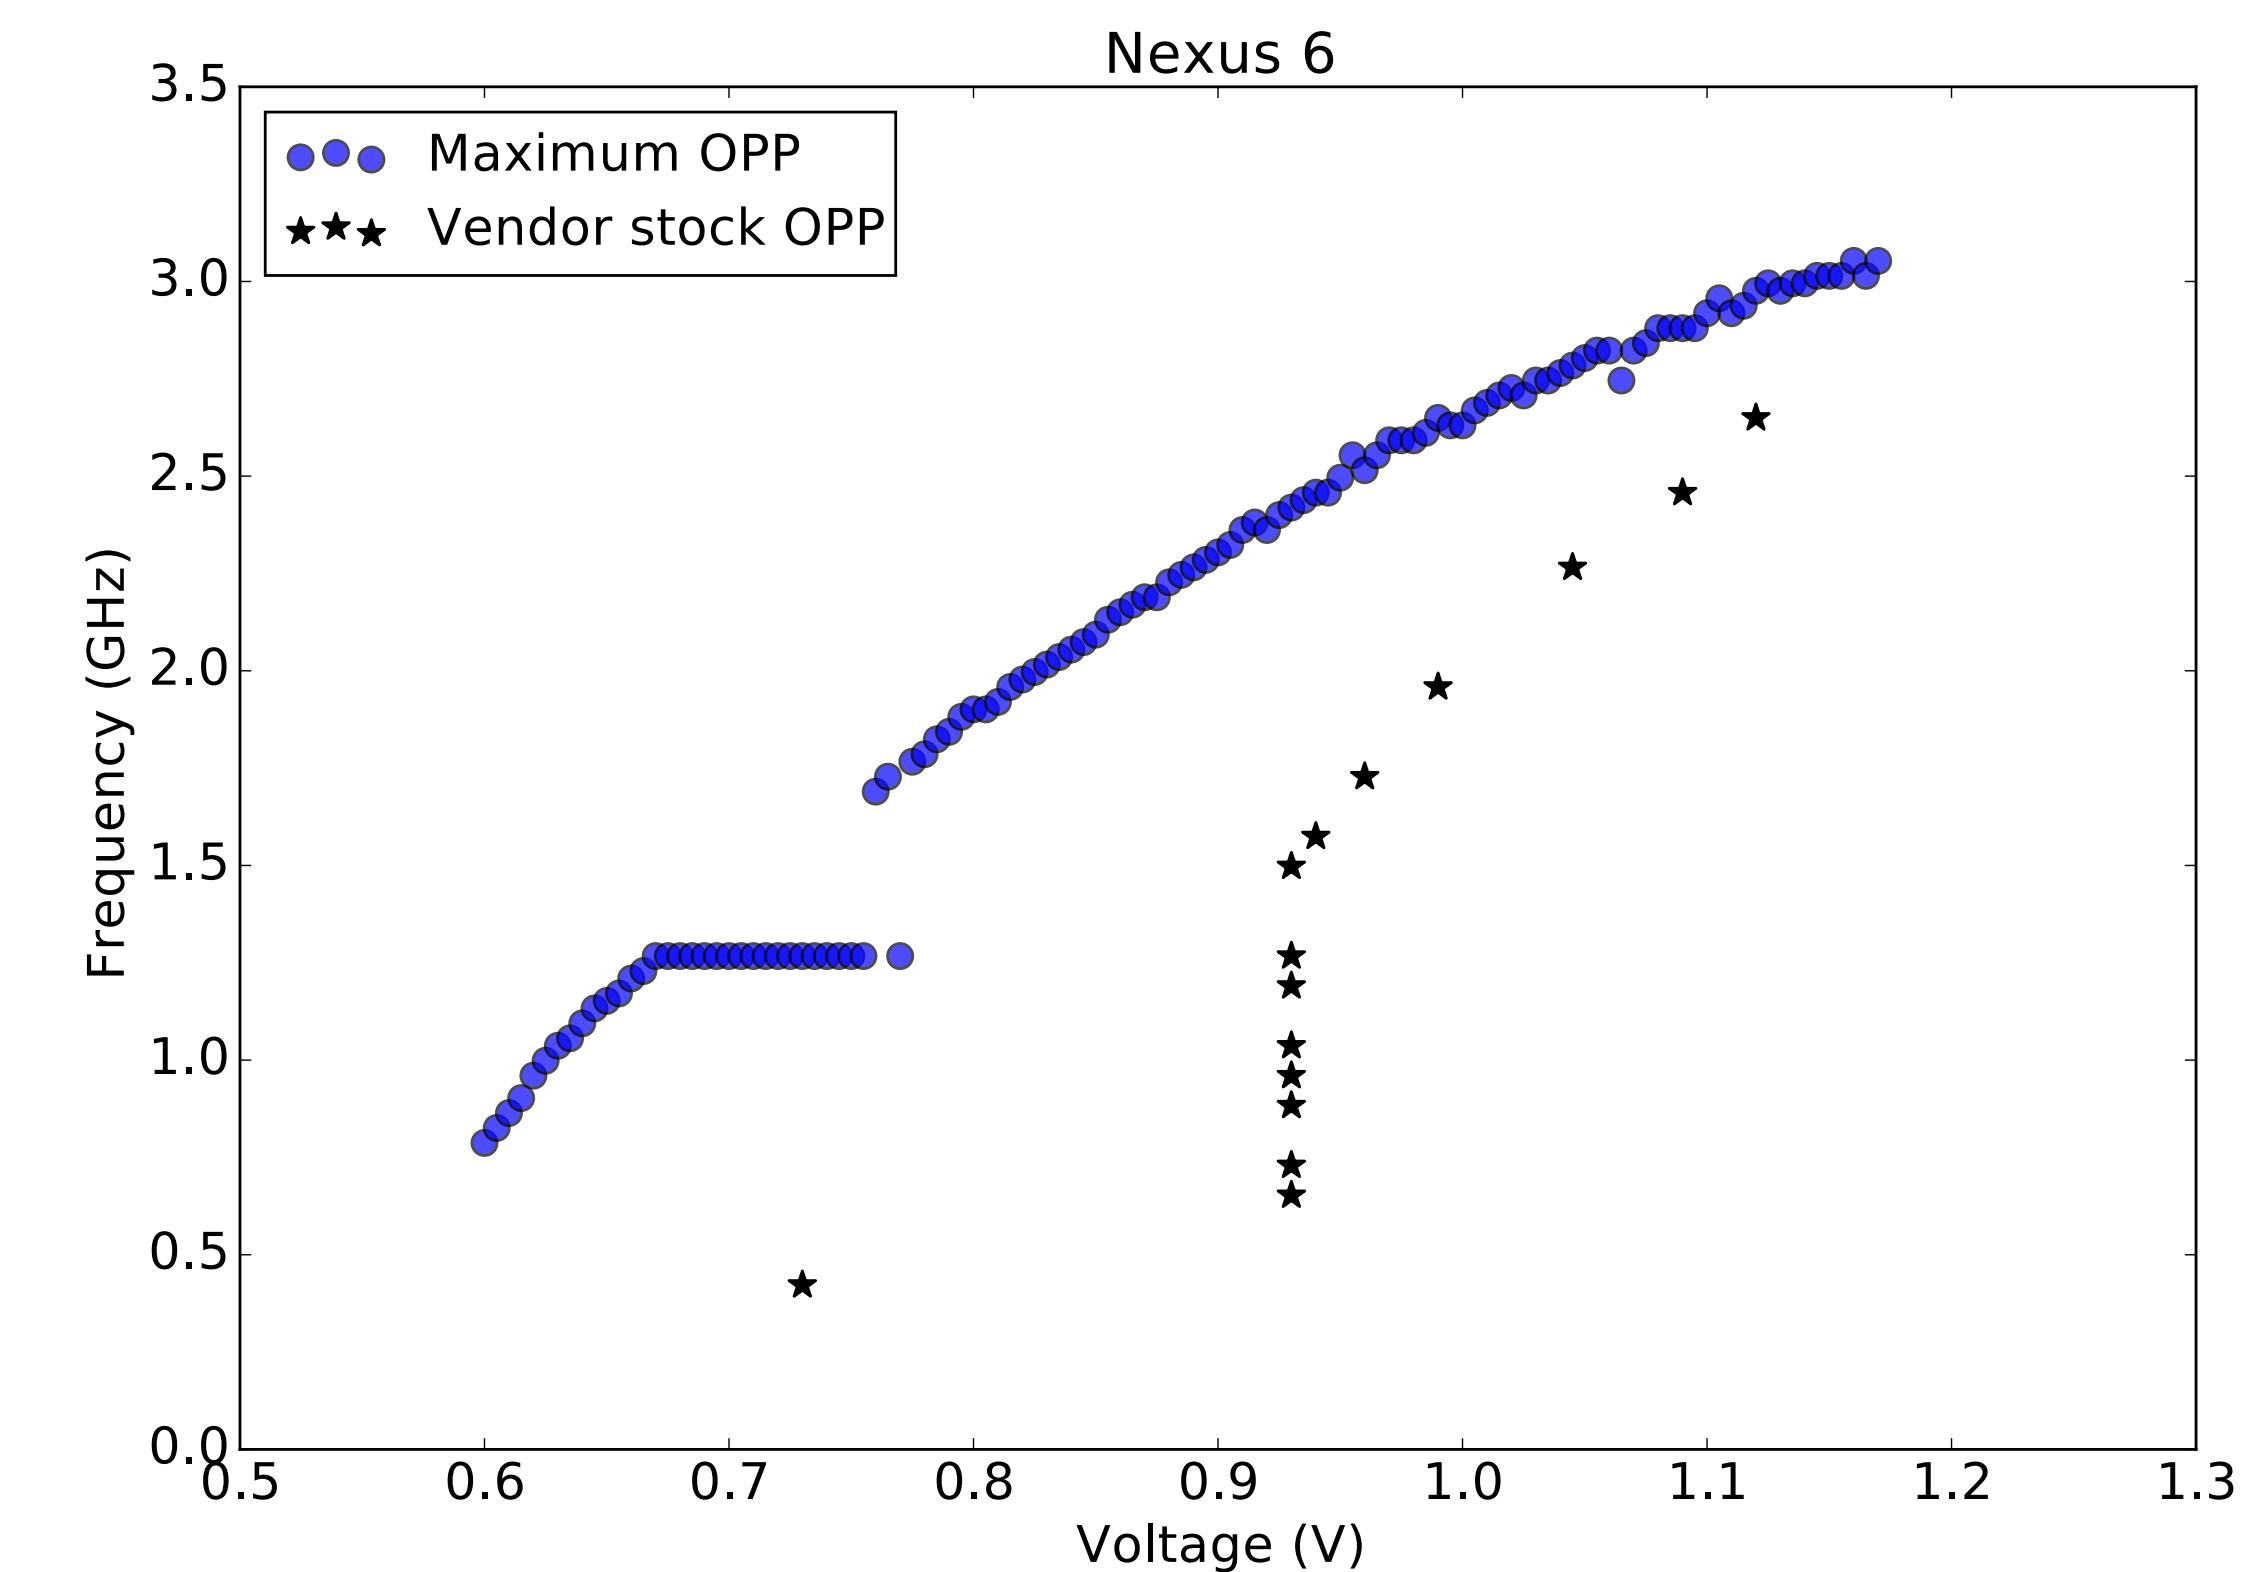
#2: Self-containment within same device

#3: Noisy complex OS environment

#4: Precise timing

#5: Fine-grained timing resolution

Addressed earlier in DVFS regulators



# CLKSCREW Challenges & Solutions

#1: Regulator operating limits

#2: Self-containment within same device

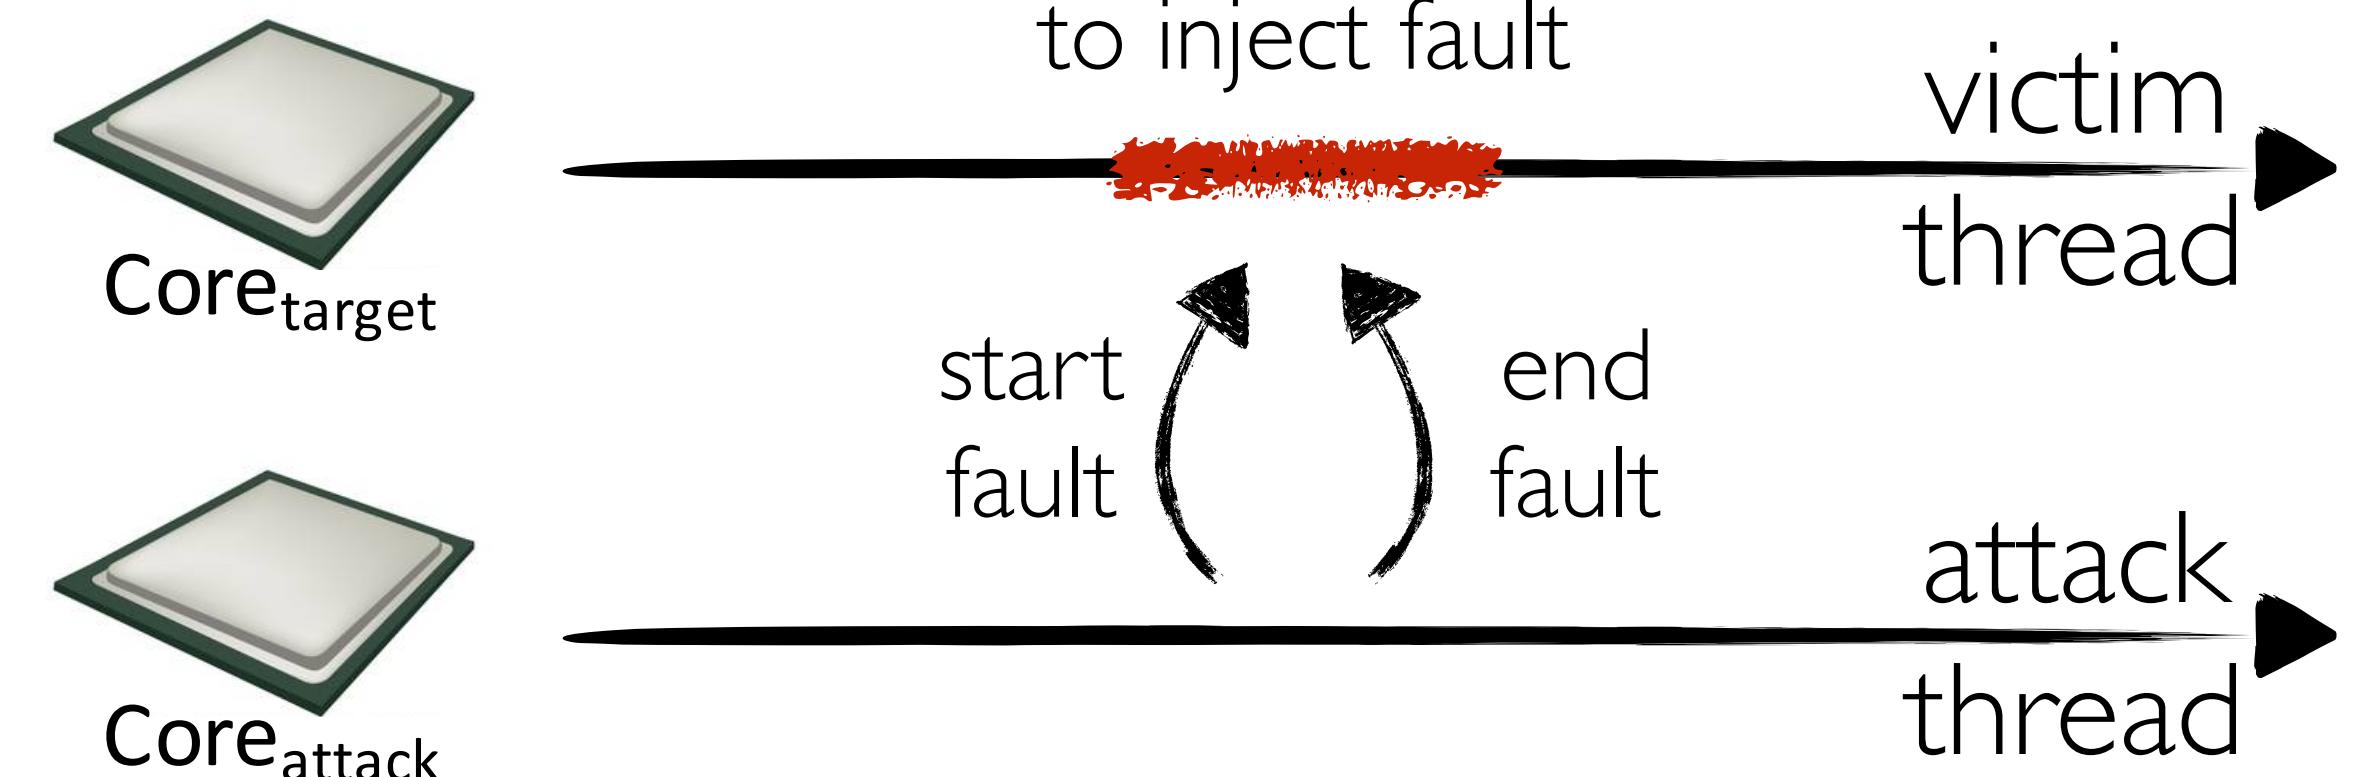
#3: Noisy complex OS environment

#4: Precise timing

#5: Fine-grained timing resolution

Cores have different frequency regulators

Core pinning



# CLKSCREW Challenges & Solutions

#1: Regulator operating limits

#2: Self-containment within same device

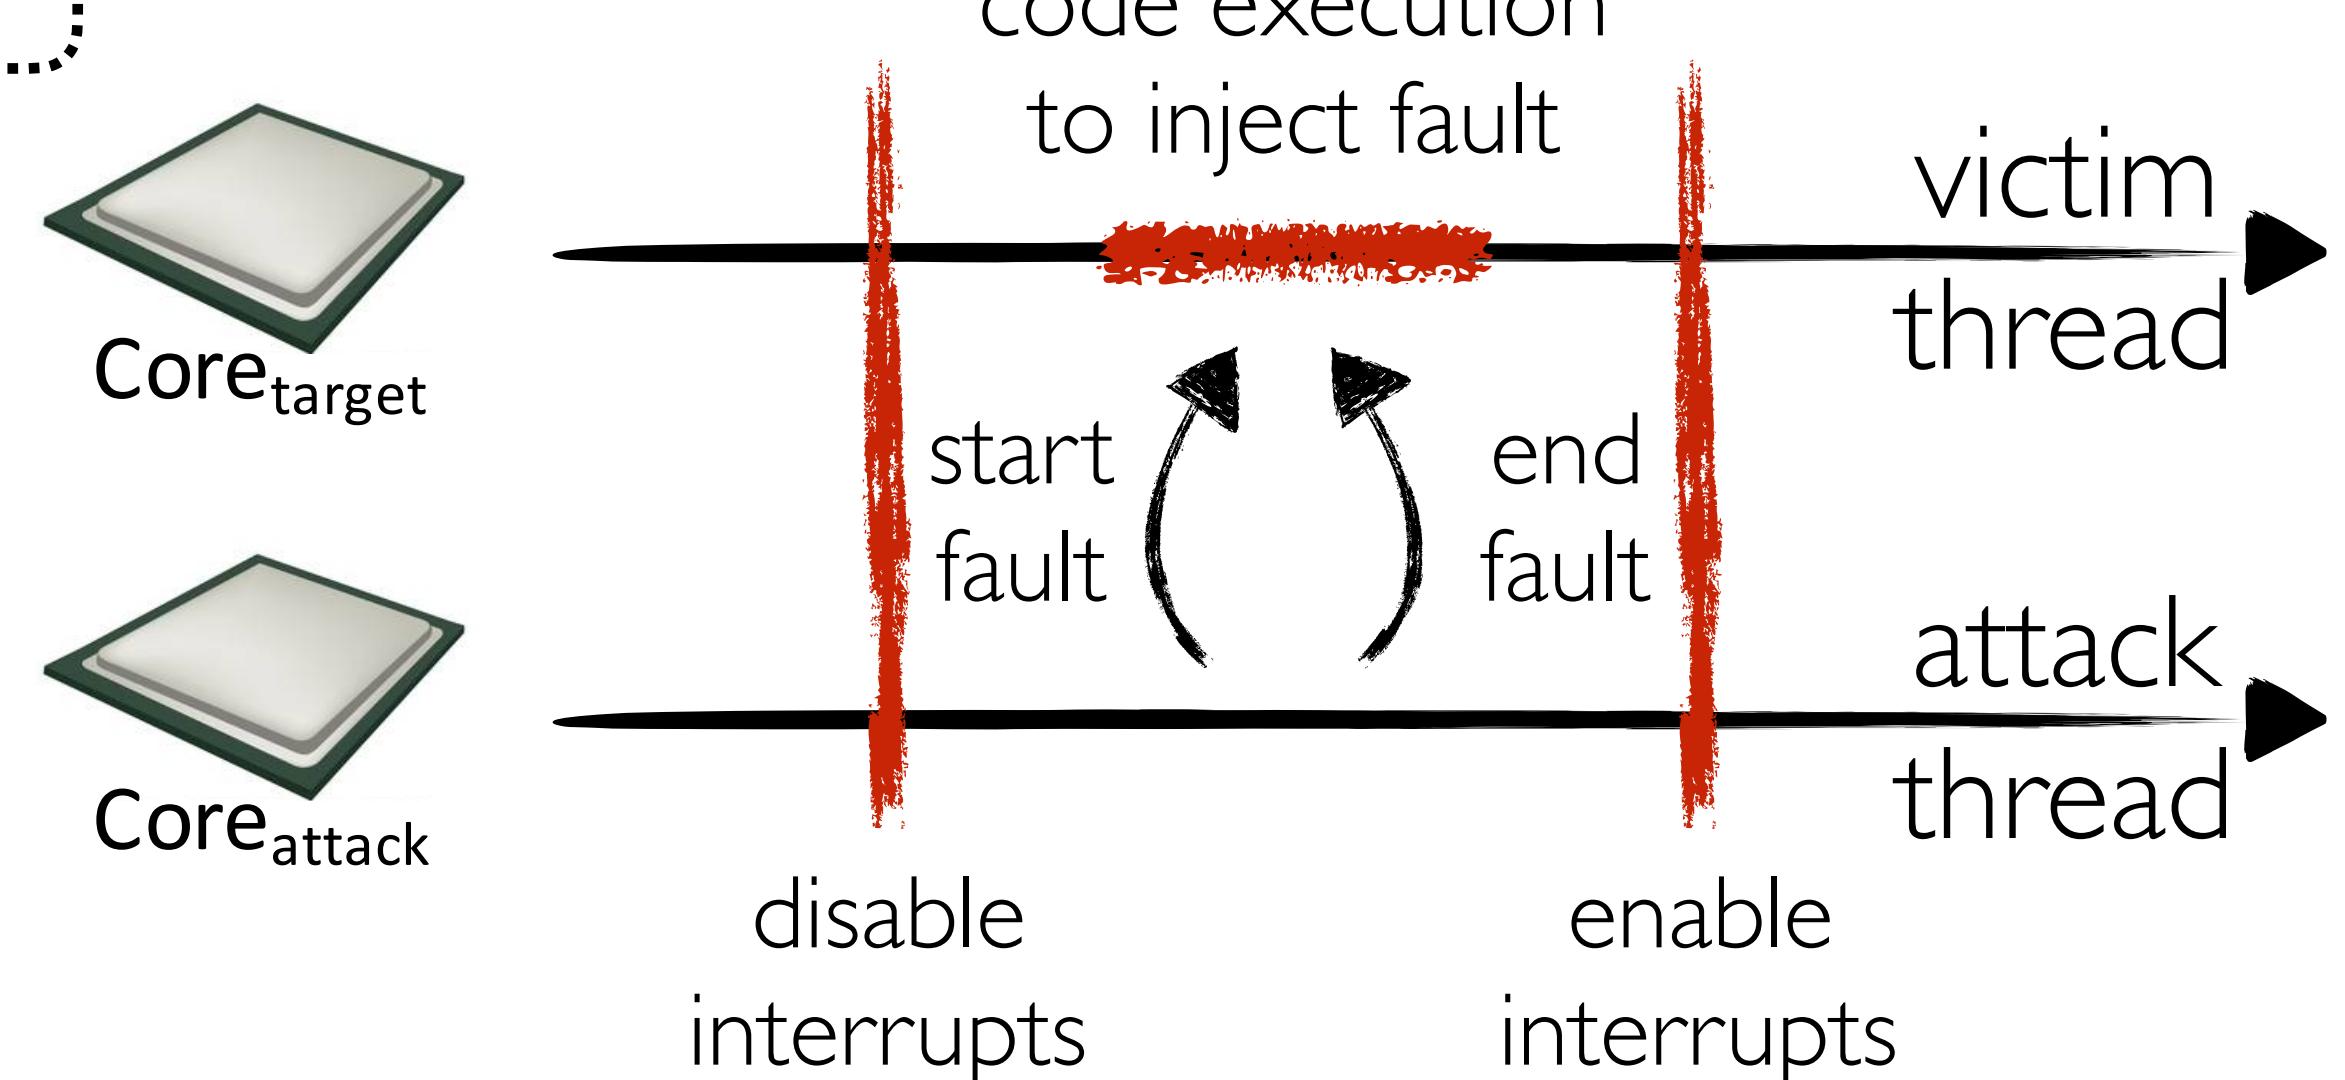
#3: Noisy complex OS environment

#4: Precise timing

#5: Fine-grained timing resolution

Core pinning

Disable interrupts during attack



# CLKSCREW Challenges & Solutions

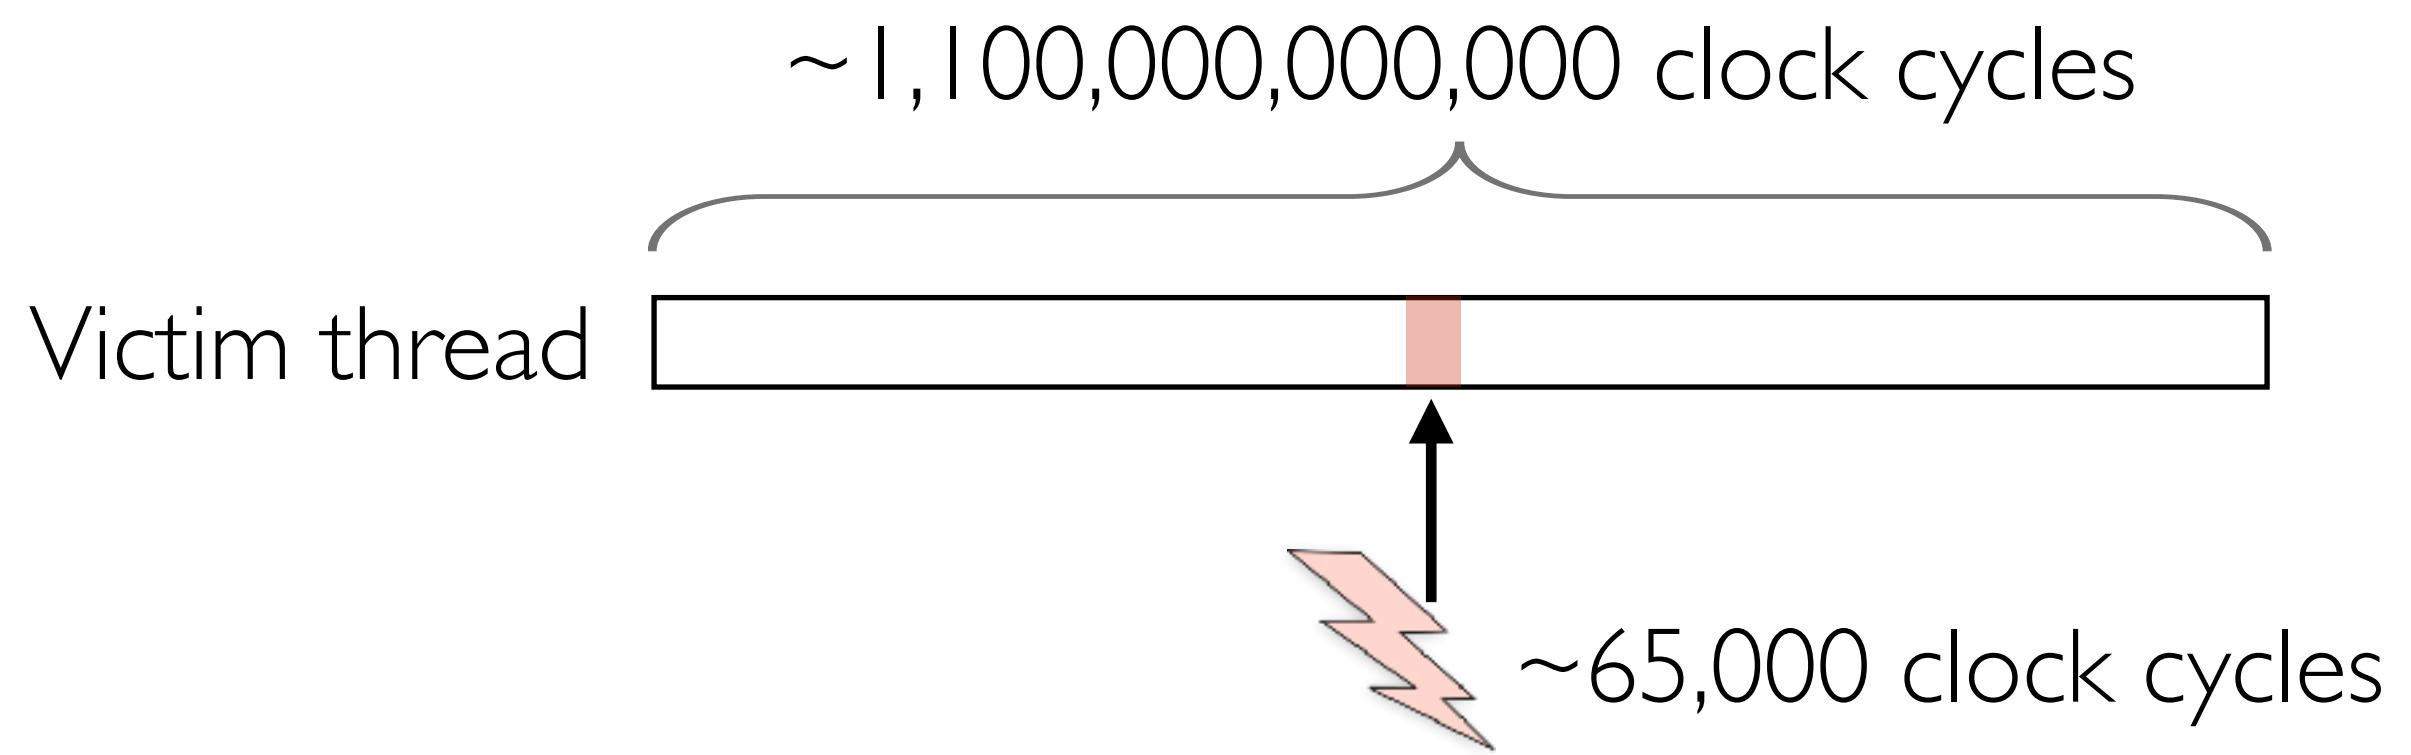
#1: Regulator operating limits

#2: Self-containment within same device

#3: Noisy complex OS environment

#4: Precise timing

#5: Fine-grained timing resolution



```
asm volatile("1: subs %0, %0, #1 \n"
            "    bhi 1b \n":::"r" (loops));
```

High-precision timing loops in attack architecture

Cache-based execution timing profiling

# Outline

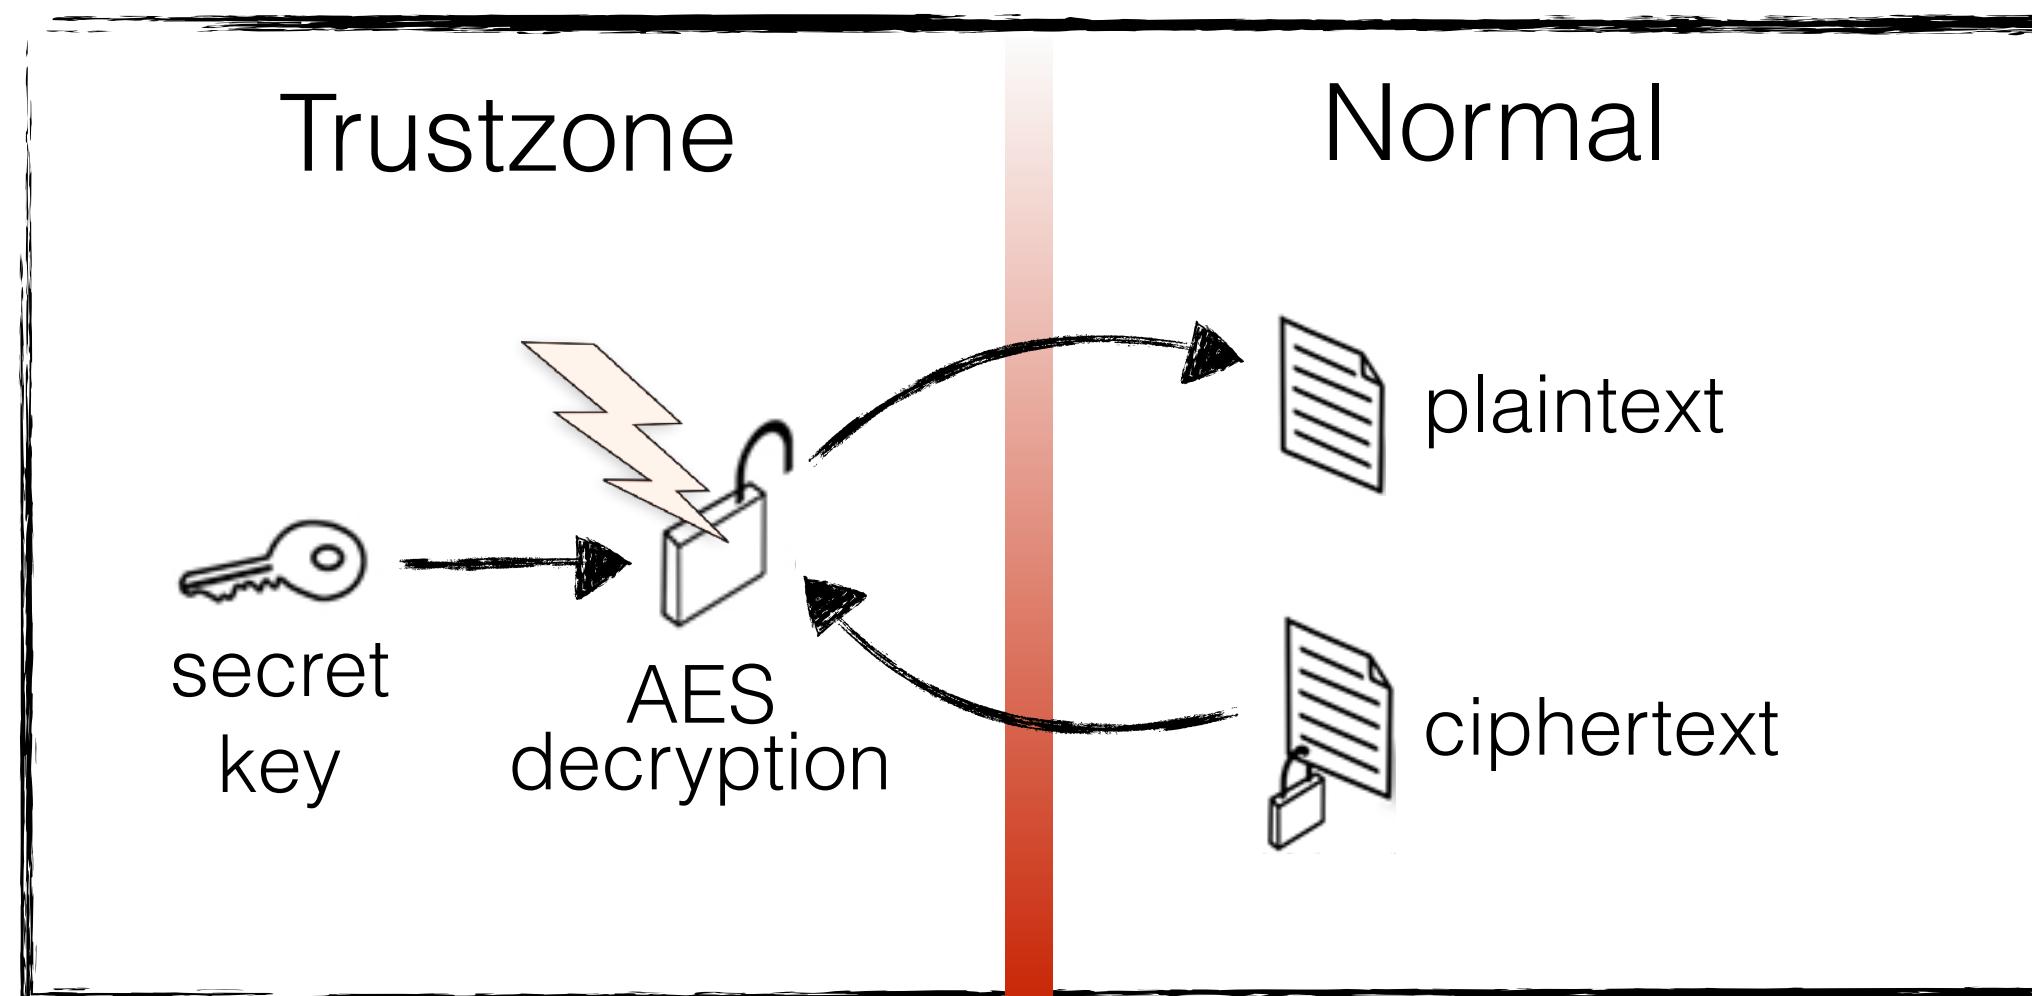
---

- I. DVFS and Deep Dive into Hardware Regulators
- II. The CLKSCREW Attack
- III. Trustzone Attack I: Secret AES Key Inference**
- IV. Trustzone Attack 2:Tricking RSA Signature Validation
- V. Concluding Remarks

# Subverting Trustzone Isolation with CLKSCREW

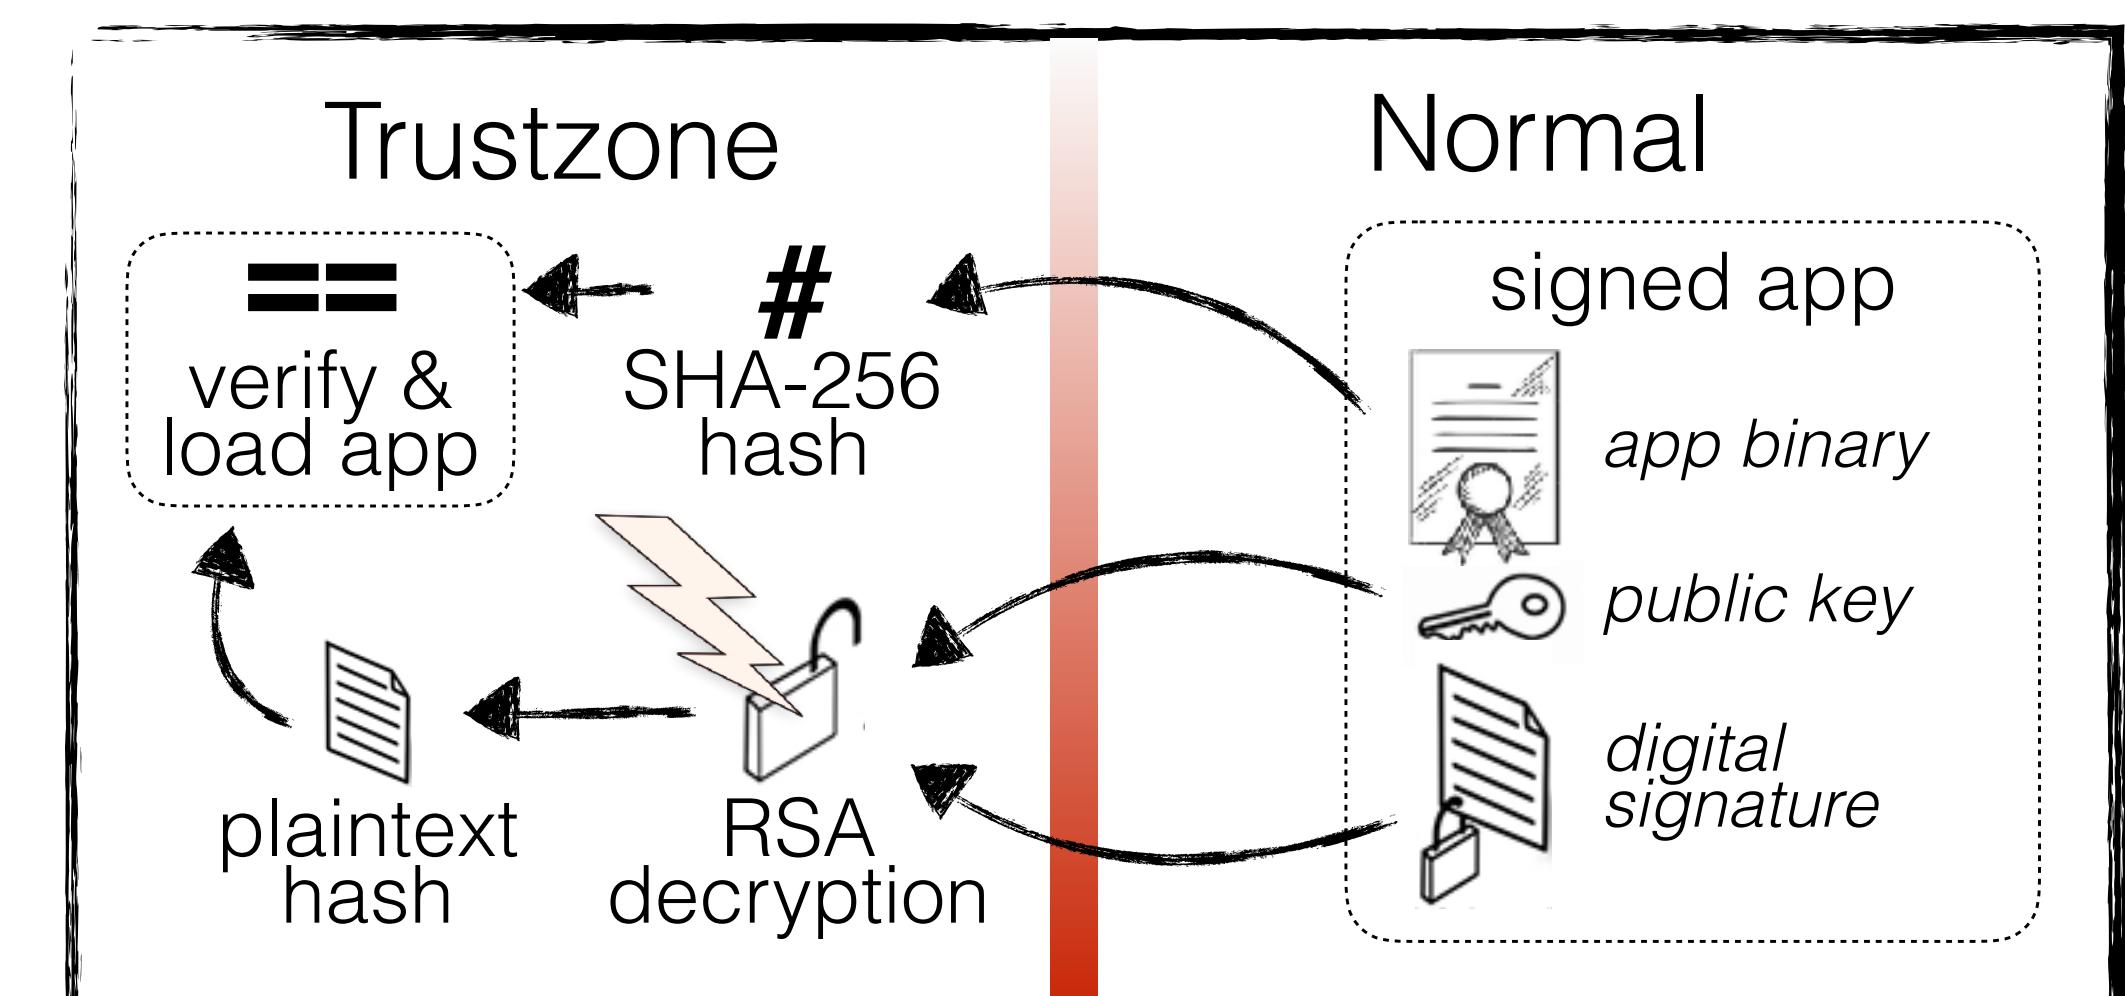
## Confidentiality Attack

infer secret AES key stored  
within Trustzone



## Integrity Attack

load self-signed app into  
Trustzone



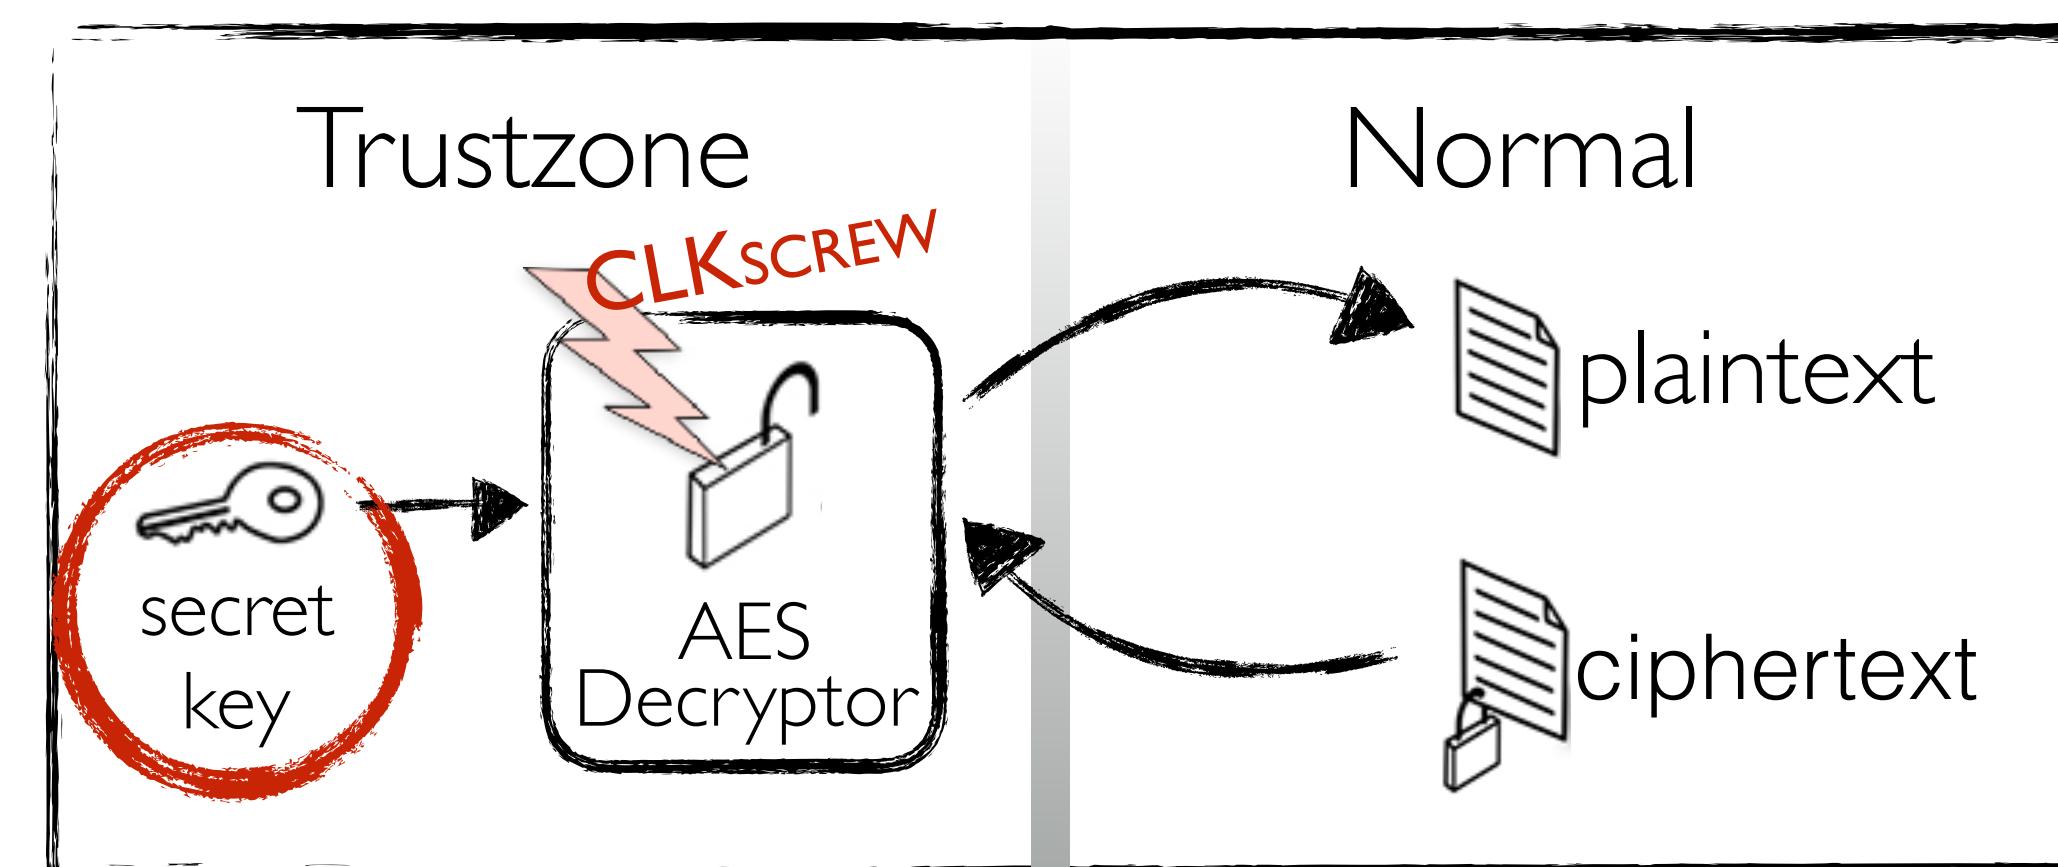
# Key Inference Attack: Threat Model

Victim app: AES decryption app executing in Trustzone

Attacker's goal: Get secret AES key from outside Trustzone

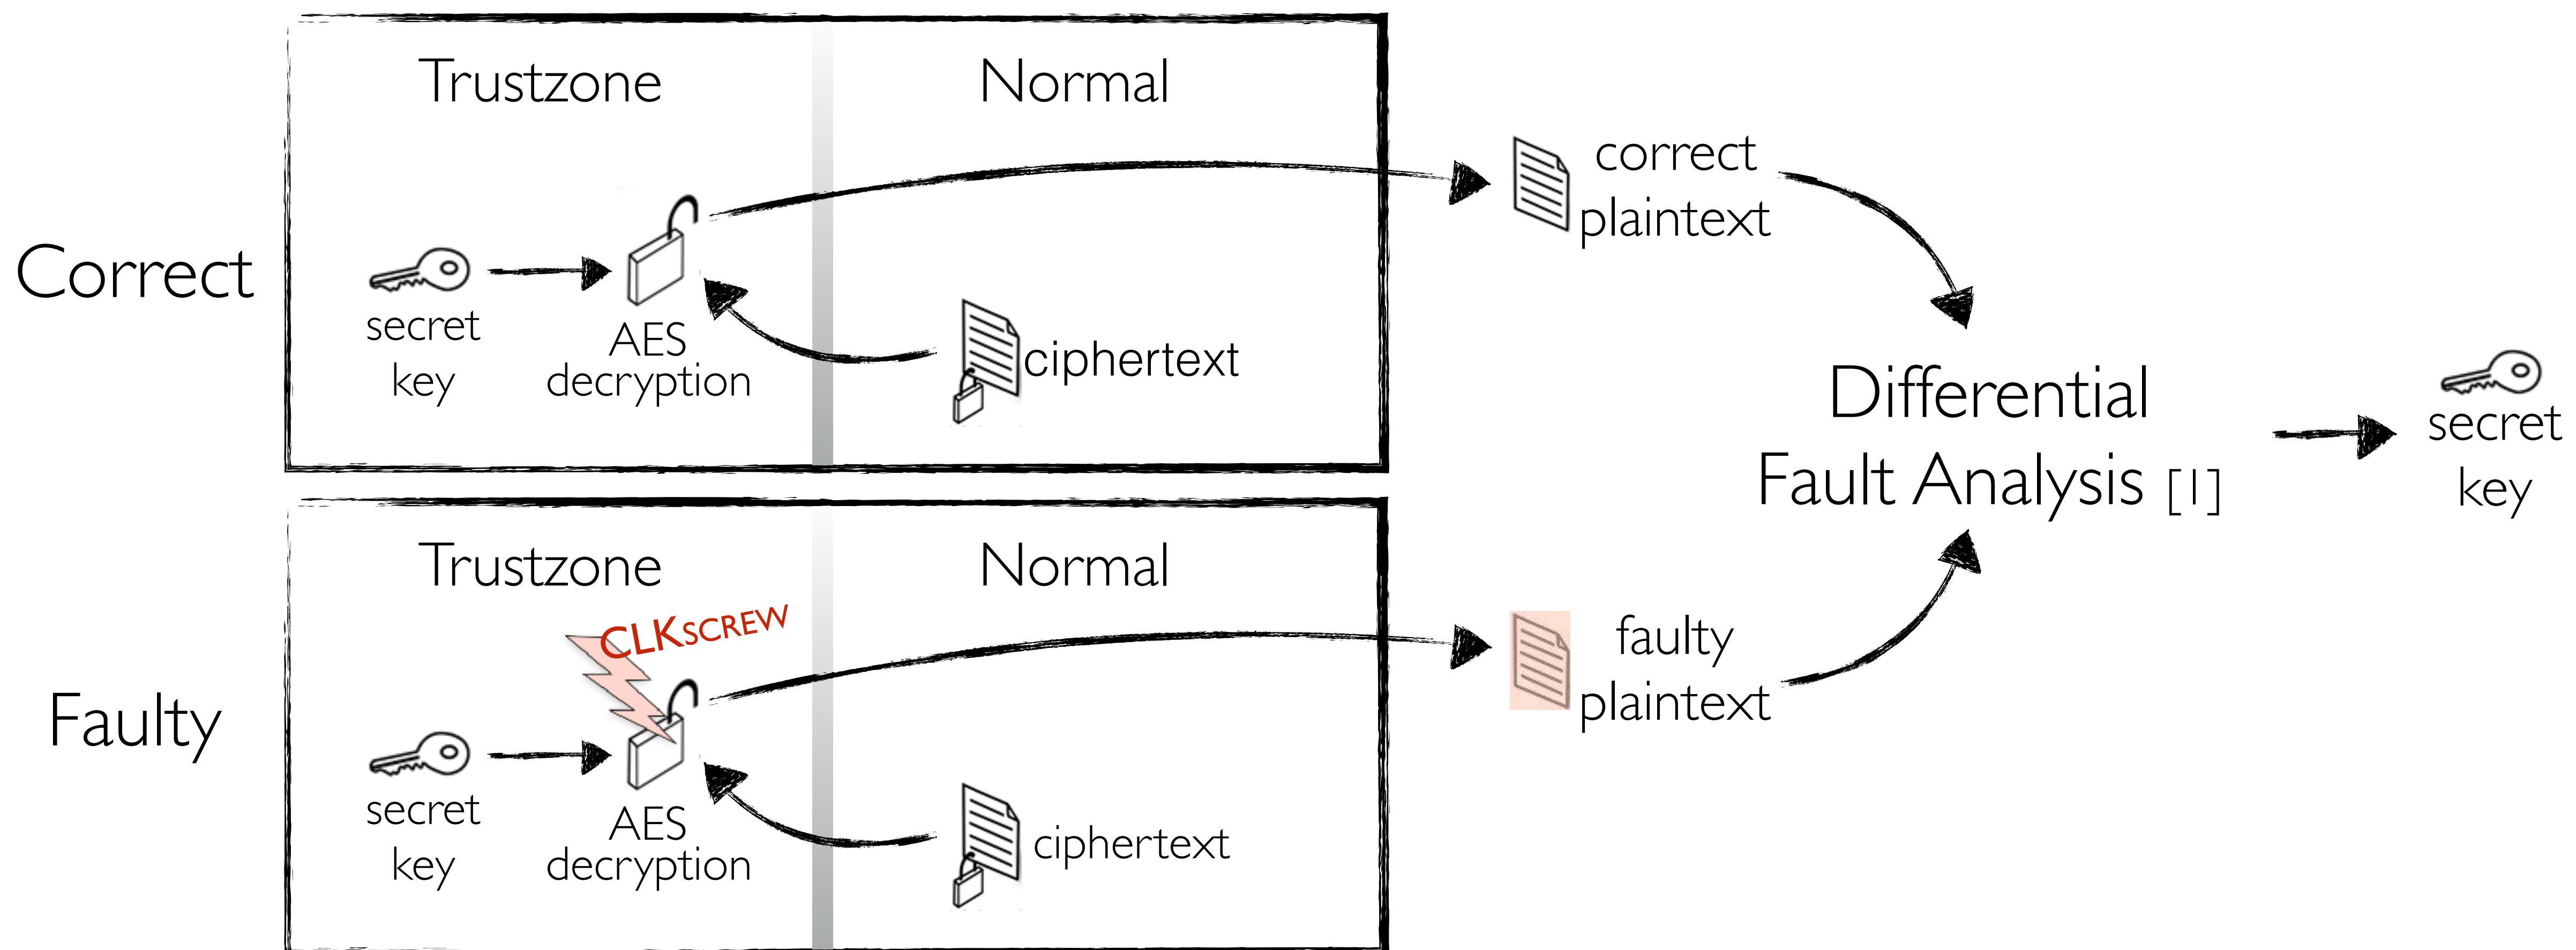
Attacker's capabilities:

- 1) Can repeatedly invoke the decryption app
- 2) Has software access to hardware regulators



# Key Inference Attack: Summary

Idea: Induce a fault during the AES decryption  
Infer key from a pair of correct and faulty plaintext



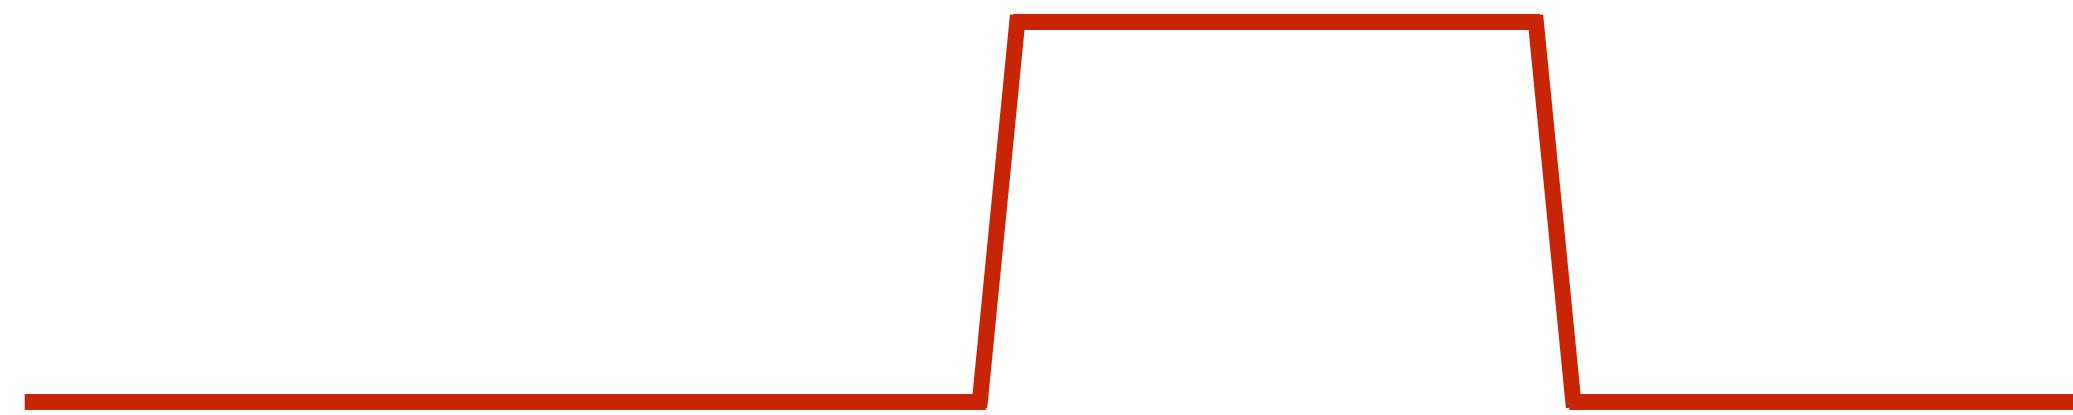
# Key Inference Attack: CLKSCREW Parameters

Base voltage: 1.055V

High frequency: 3.69GHz

Low frequency: 2.61GHz

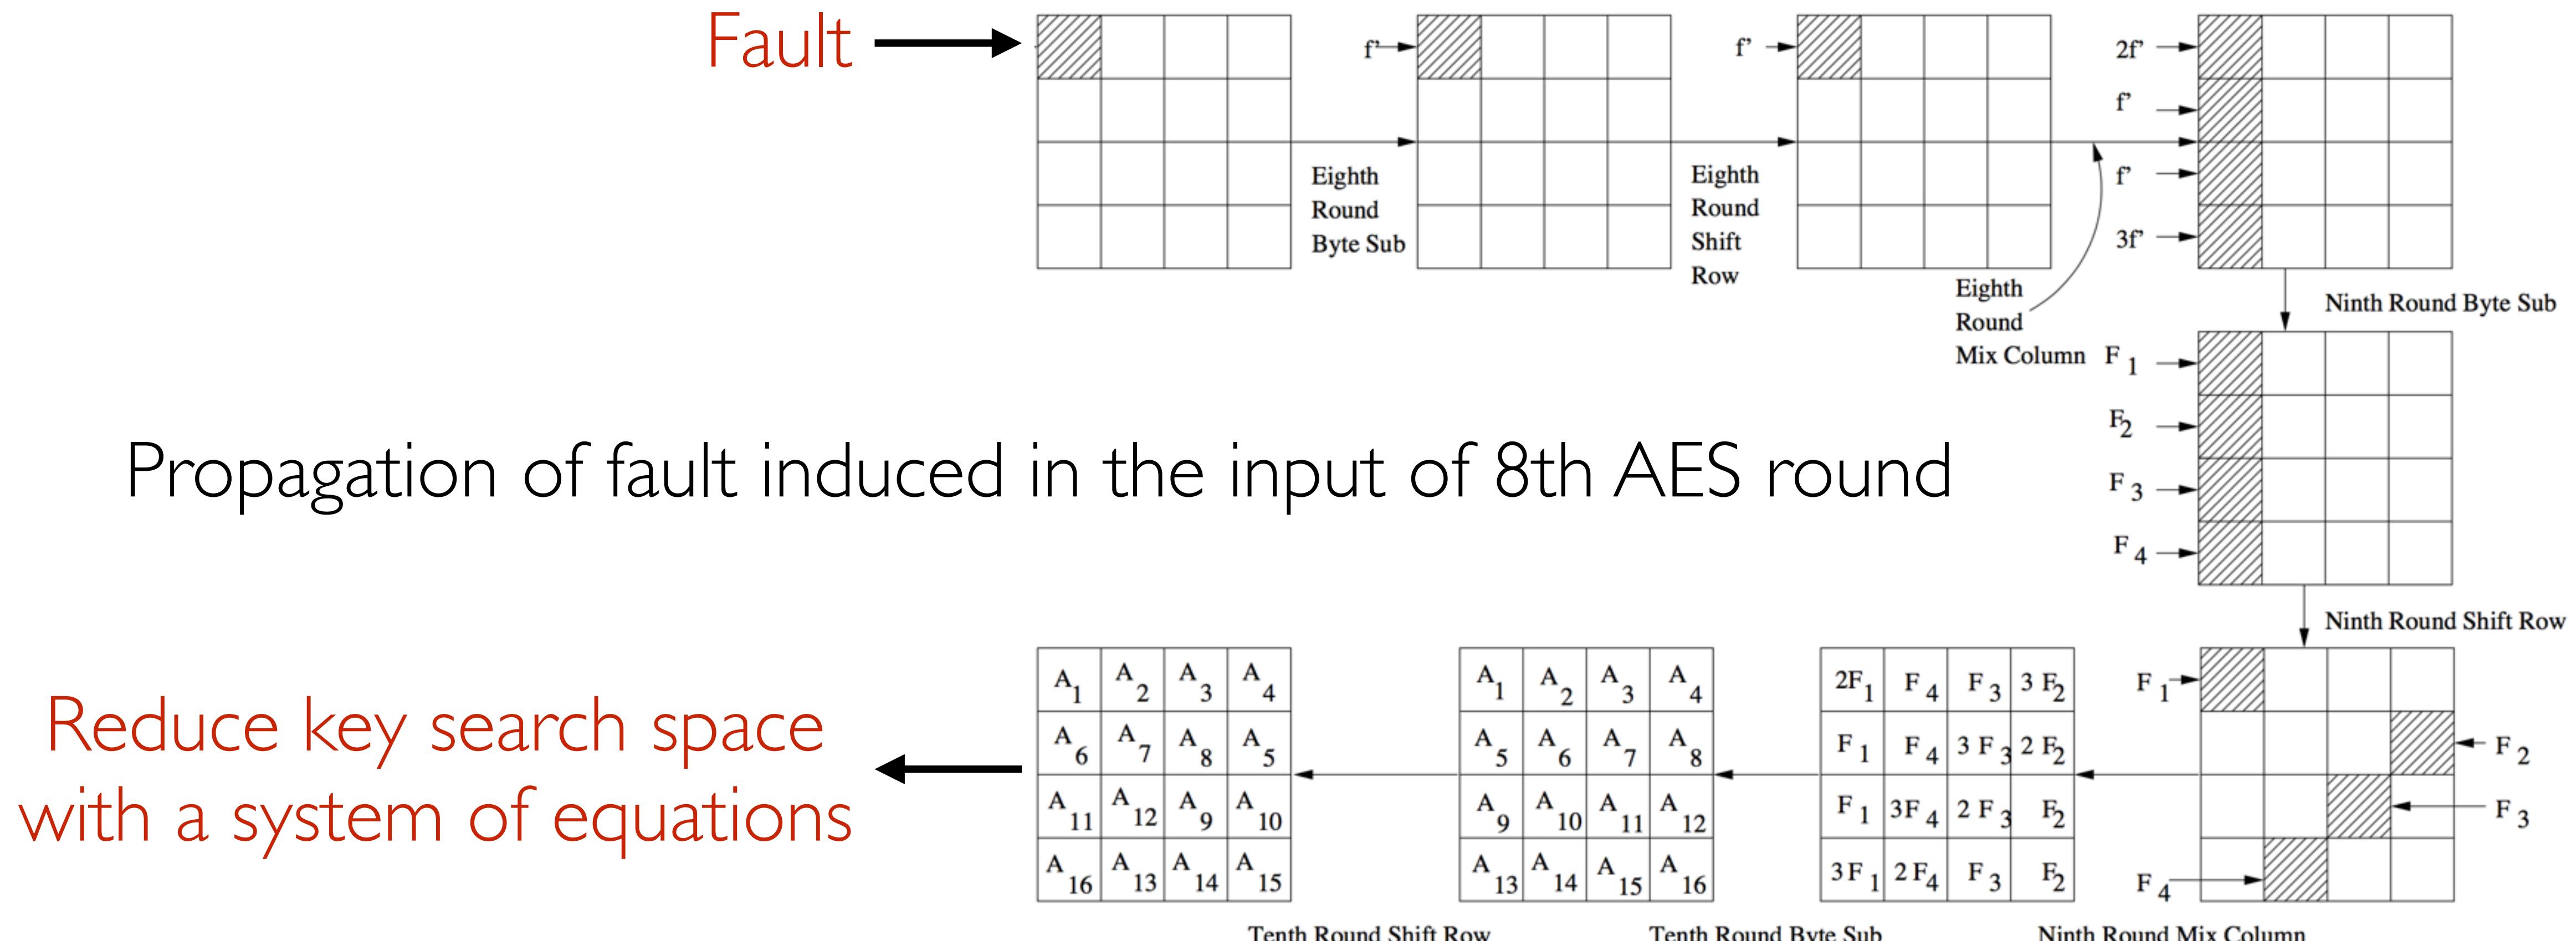
Fault injection duration:  
680 no-op loops (~39 µsec)



Differential Fault Analysis needs CLKSCREW to deliver a one-byte fault to the 7th AES round

# Key Inference Attack: Differential Fault Analysis

Check out code at: <https://github.com/Daeinardfa-aes>



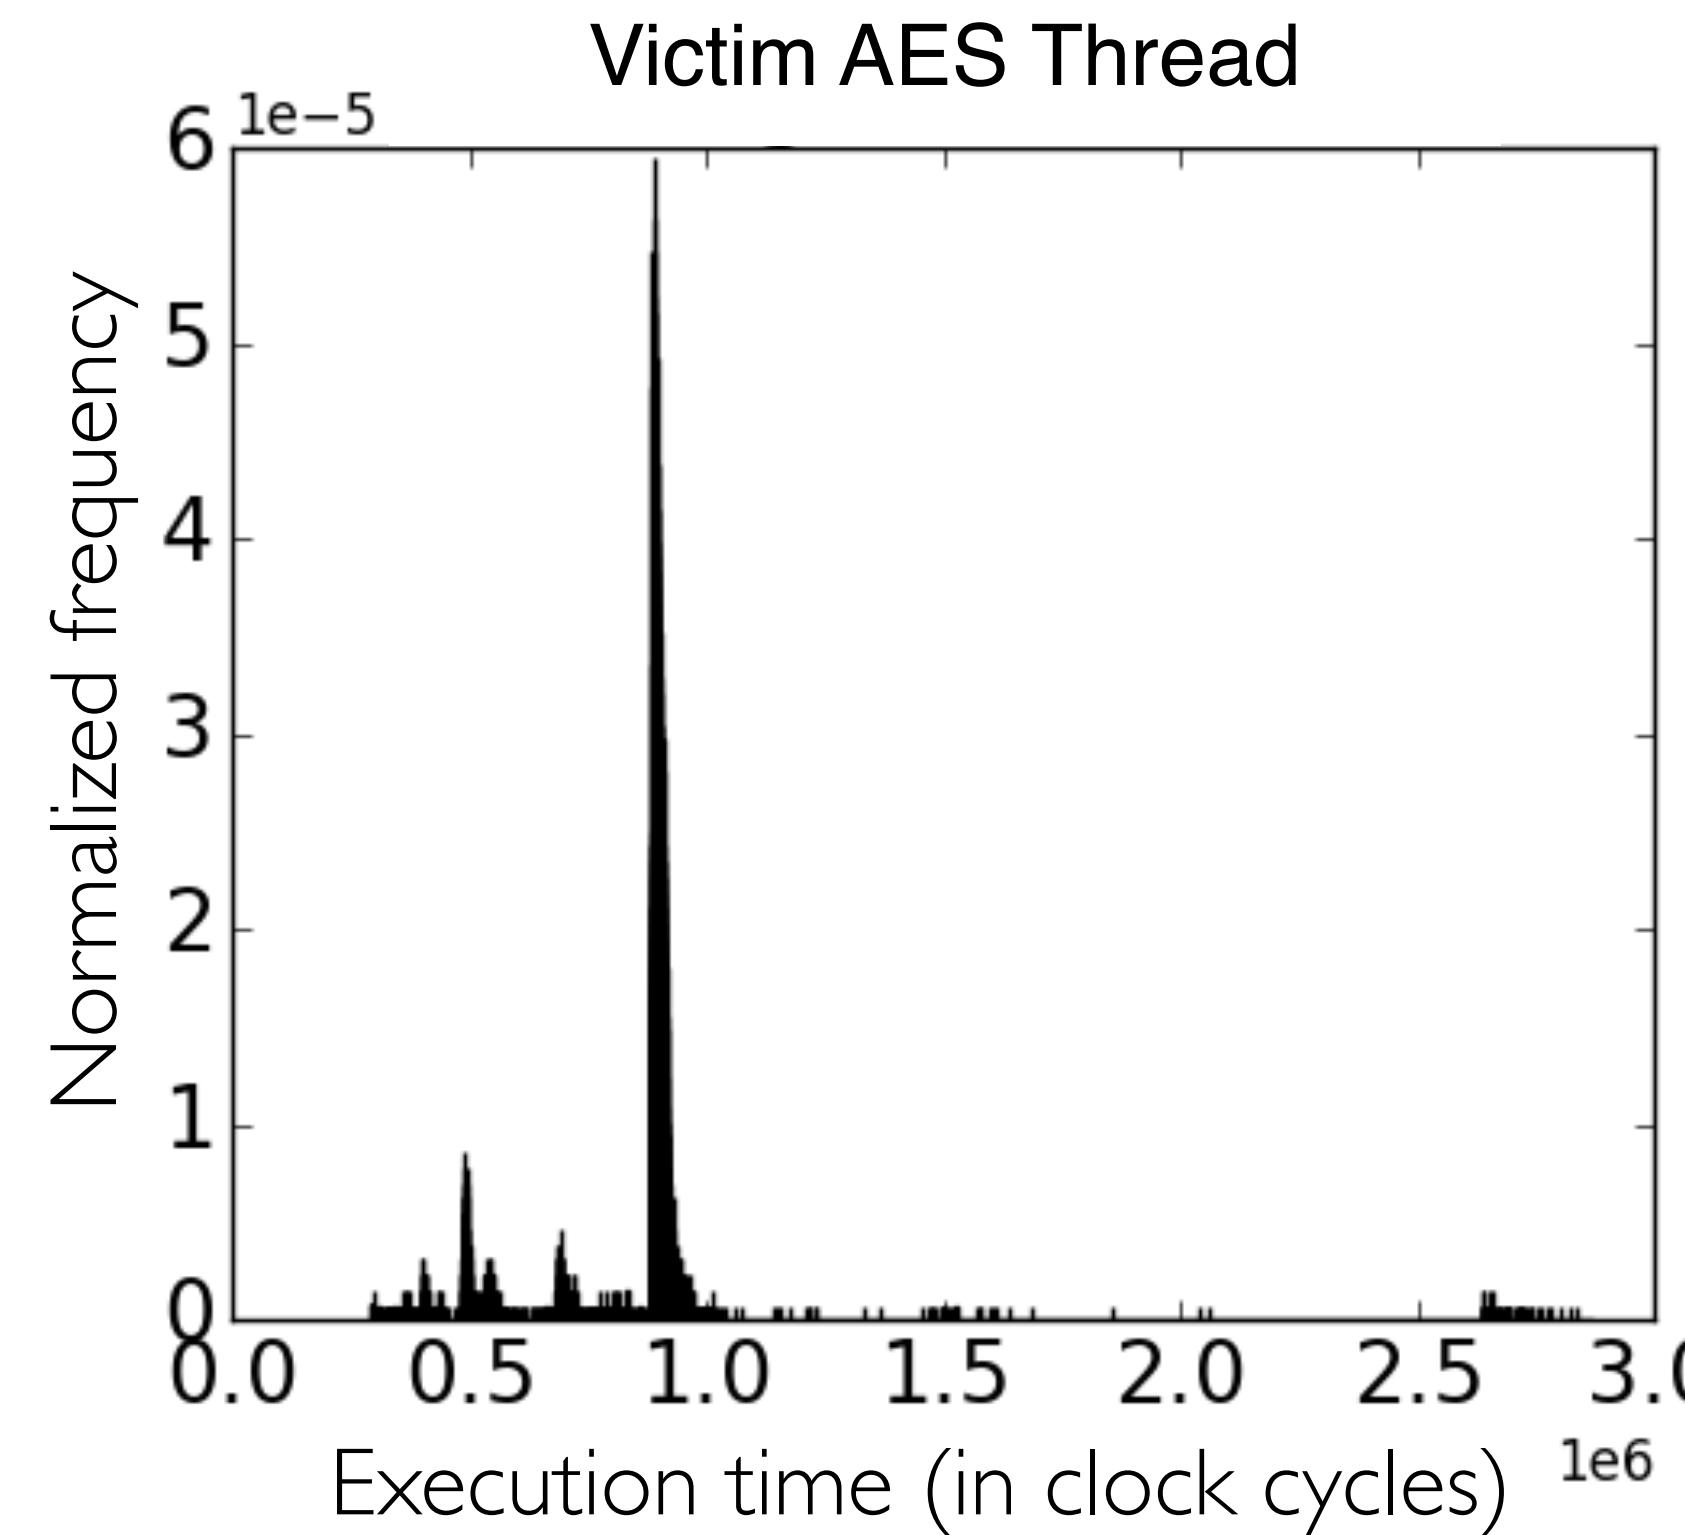
# Key Inference Attack: Timing Profiling

---

Execution timing of Trustzone code can be profiled with hardware cycle counters that are accessible outside of Trustzone

# Key Inference Attack: Timing Profiling

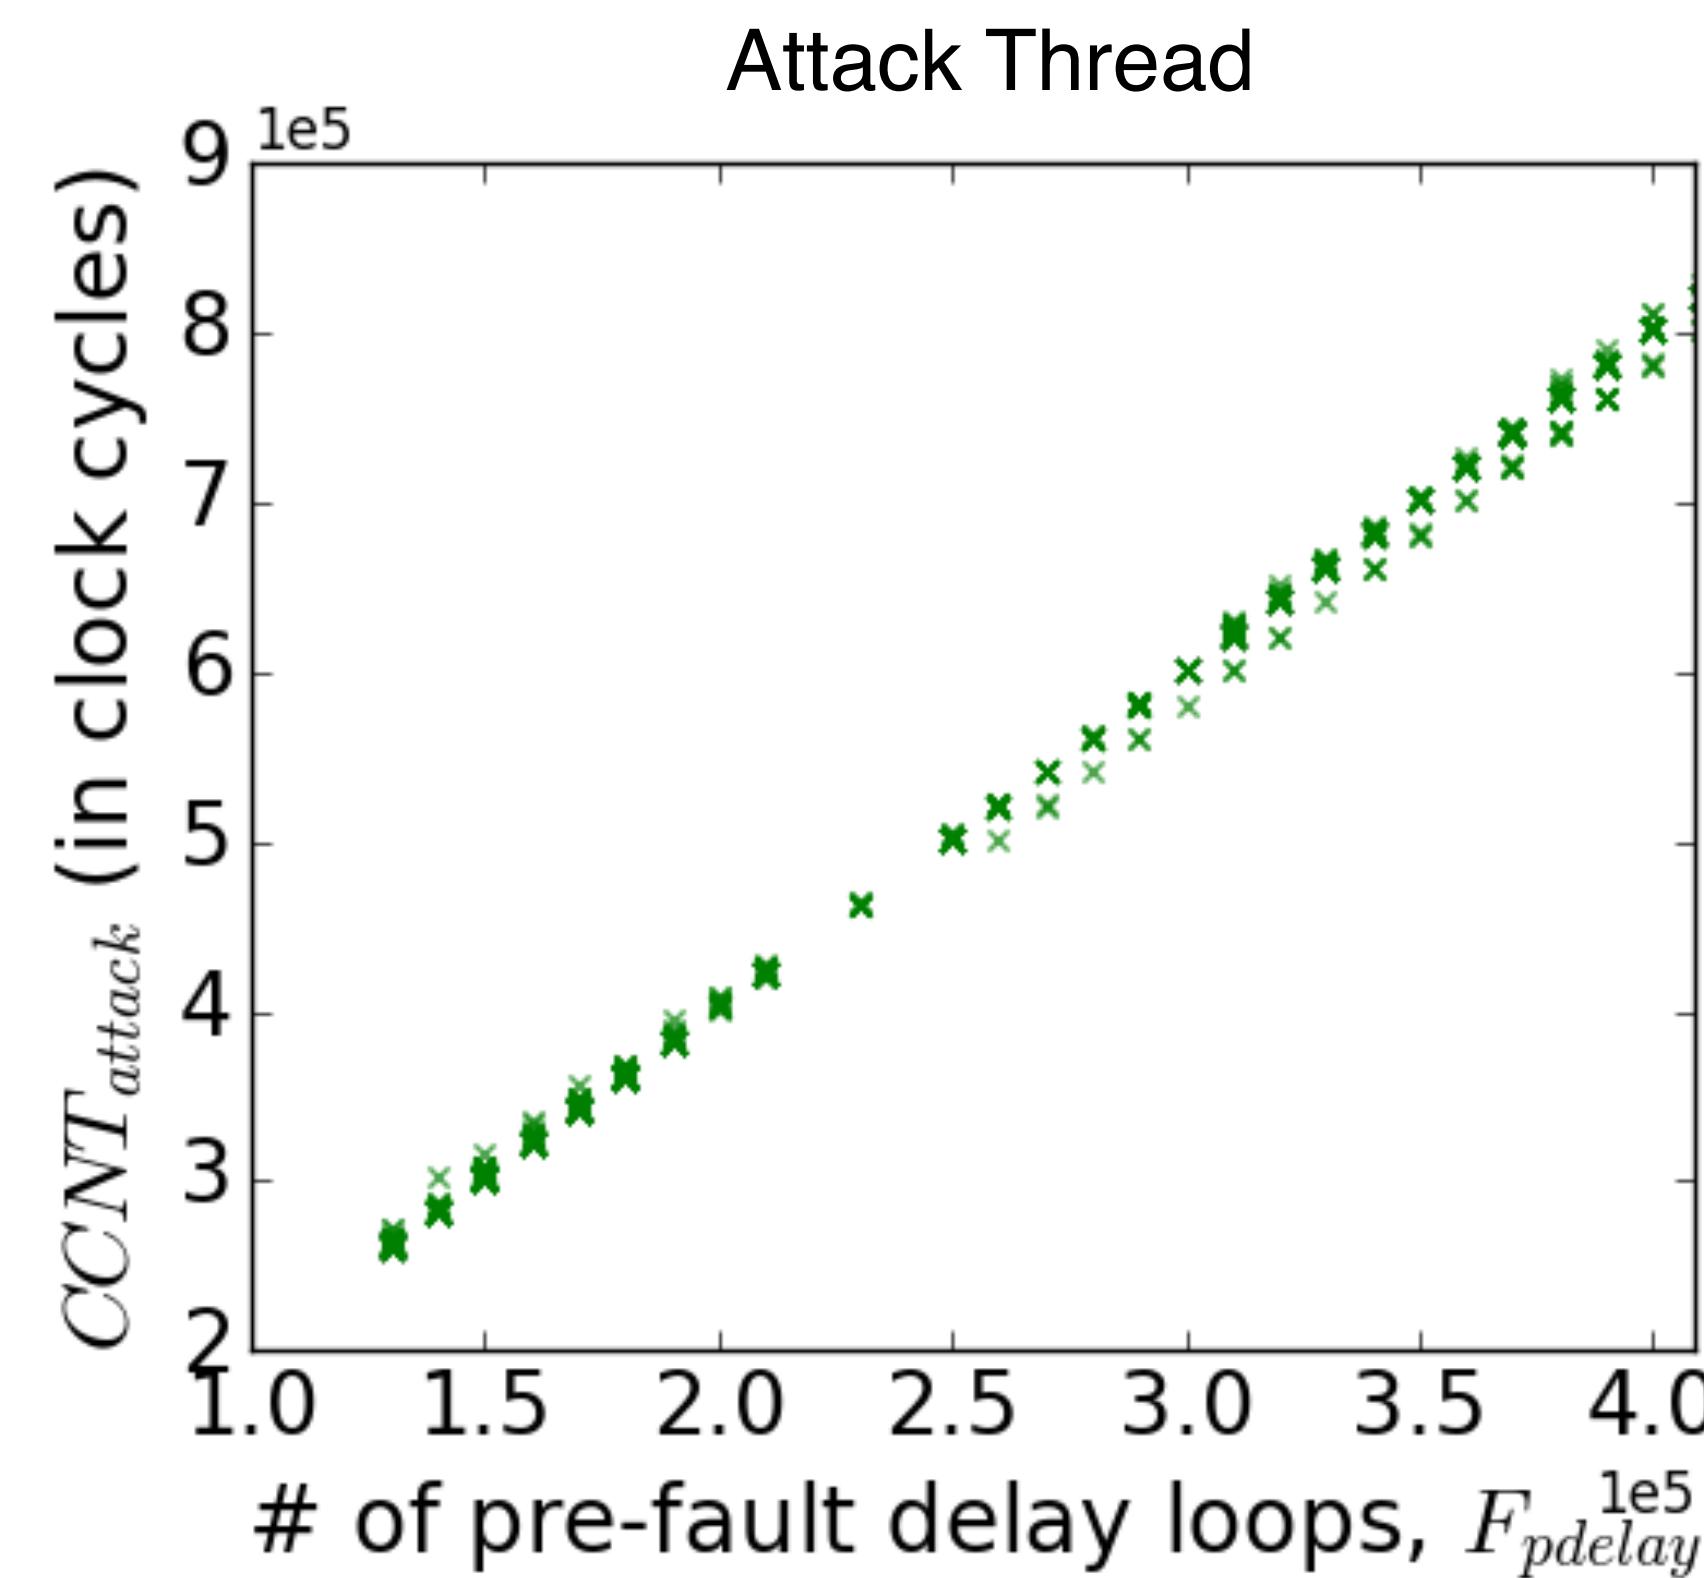
How varied is the execution timing of the victim decryption app?



Not too much variability in terms of execution time

# Key Inference Attack: Timing Profiling

Can we effectively control the timing of the fault delivery with no-op loops?



Number of no-op loops is a good proxy to control timing of fault delivery

# Key Inference Attack: Fault Model

---

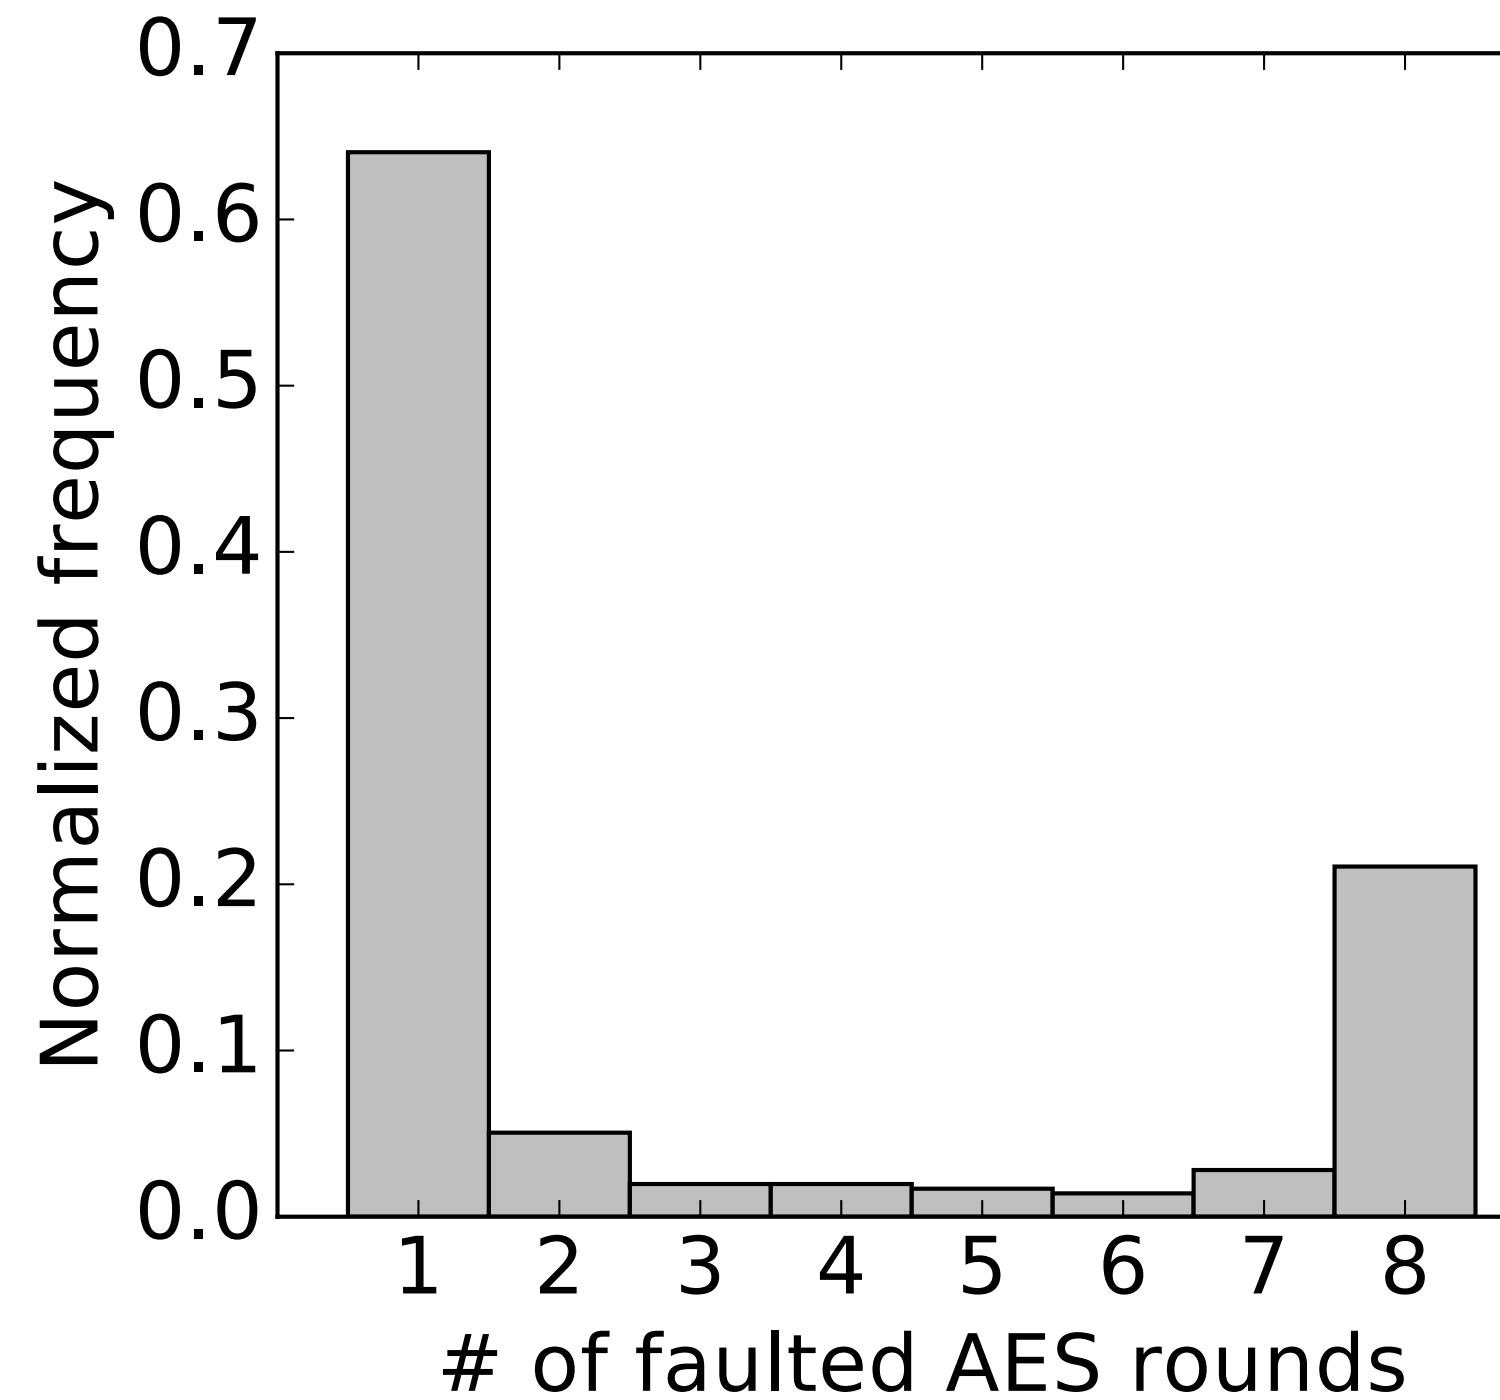
Our fault model requires our attack to inject fault

Exactly one AES round at the 7th round

Corruption of exactly one byte

# Key Inference Attack: Fault Model

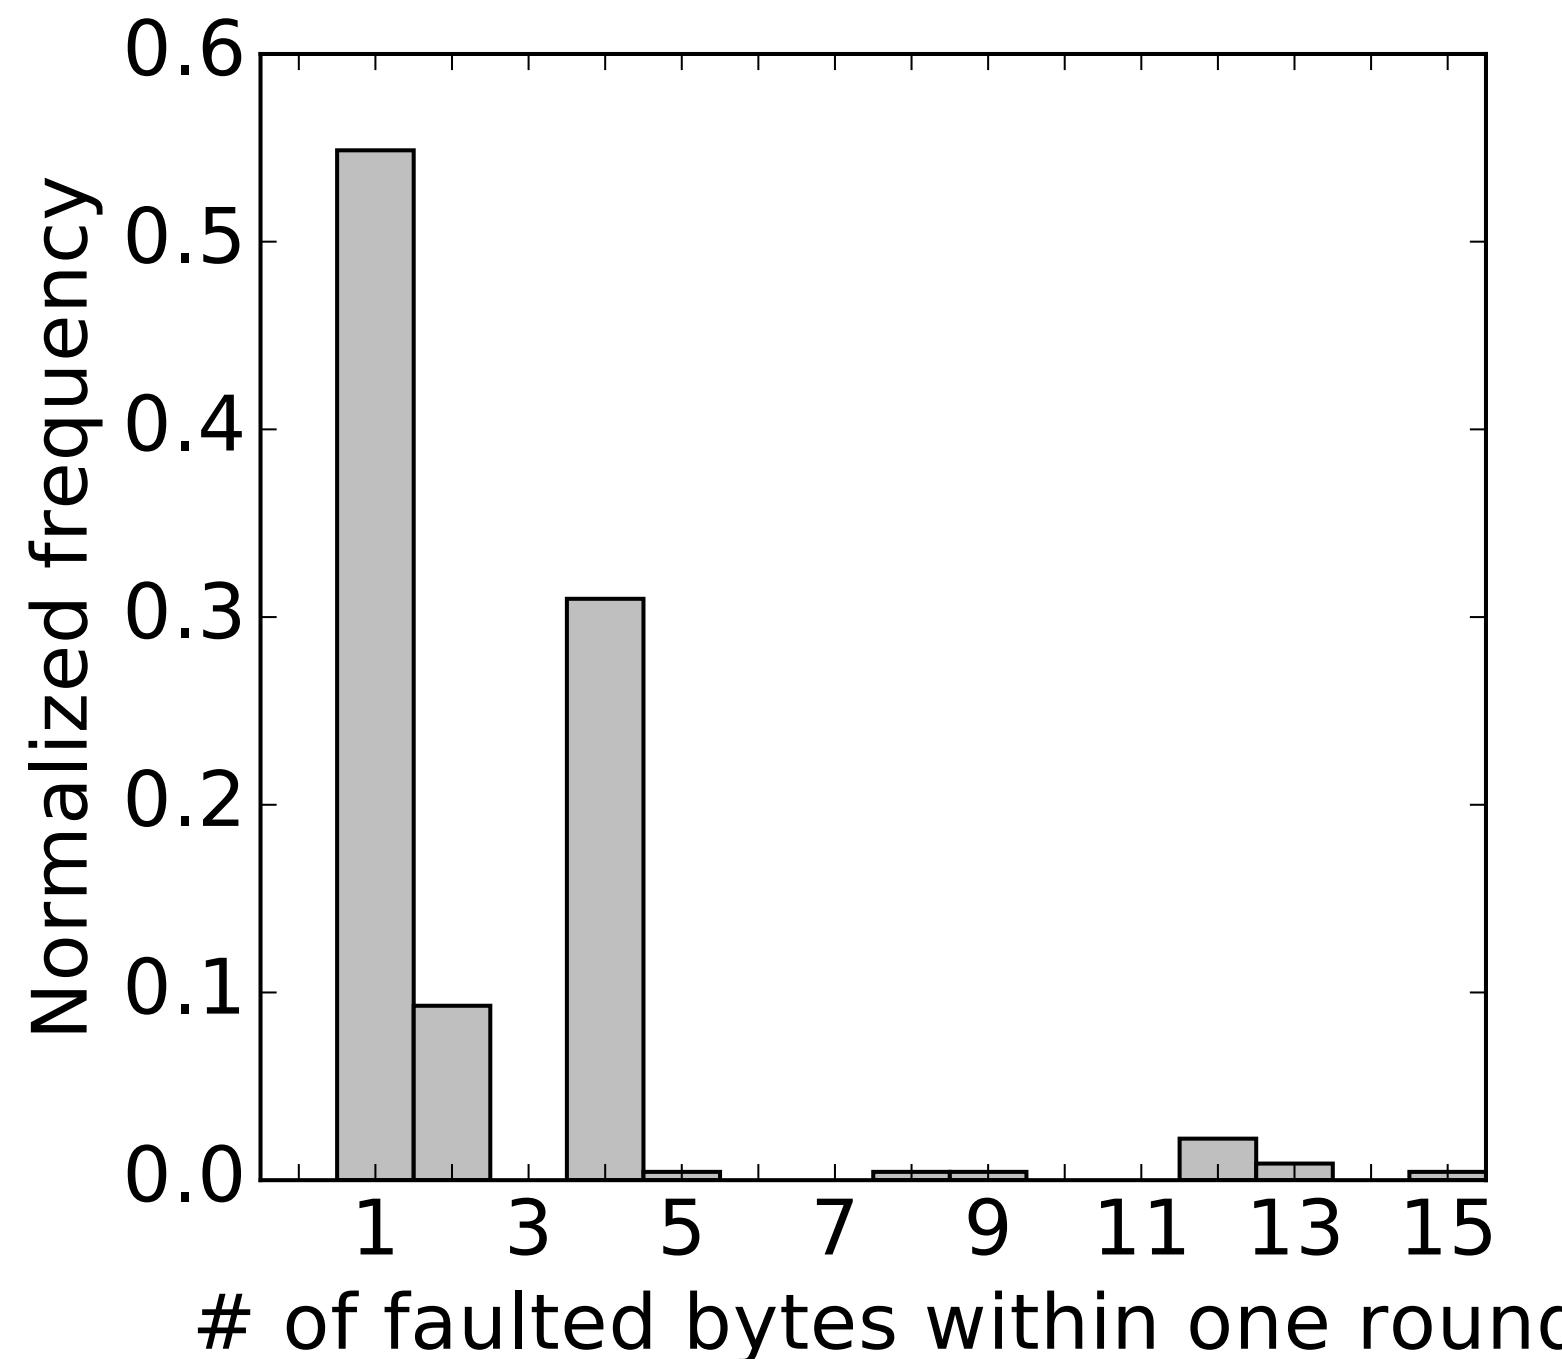
**Precision:** How likely can we inject fault in exactly one AES round?



More than 60% of the resulting faults are precise enough to corrupt exactly one AES round

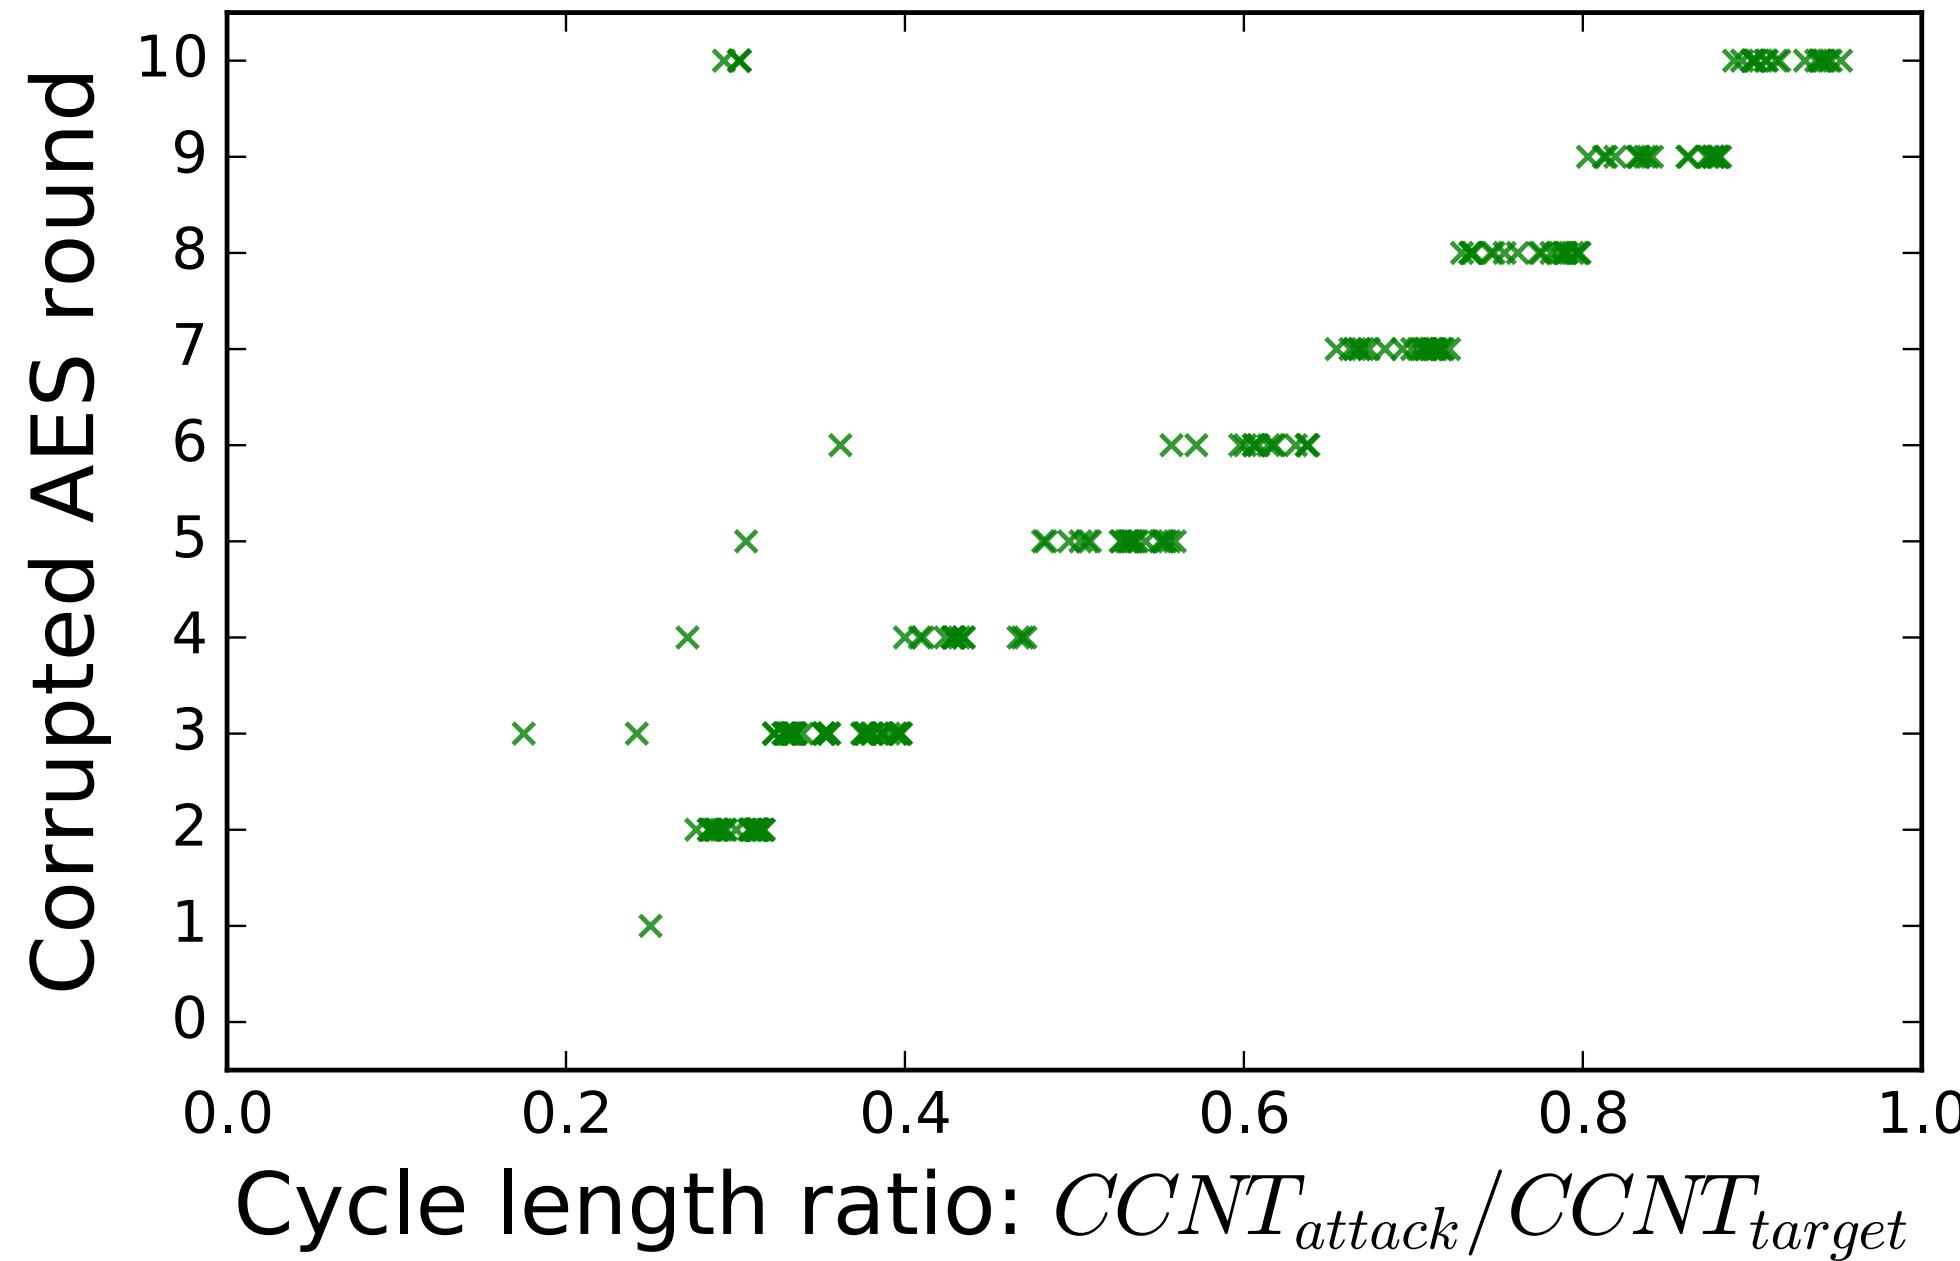
# Key Inference Attack: Fault Model

**Transience:** How likely can we corrupt exactly one byte?



Out of the above faults that affect one AES round, **more than half** are transient enough to corrupt **exactly one byte**

# Key Inference Attack: Results



Controlling  $F_{pdelay}$  allows us to precisely time the delivery of the fault to the targeted AES round

Statistics:

- ~20 faulting attempts to induce one-byte fault to desired AES round.
- ~12 min on a 2.7GHz quad core CPU to generate 3650 key hypotheses

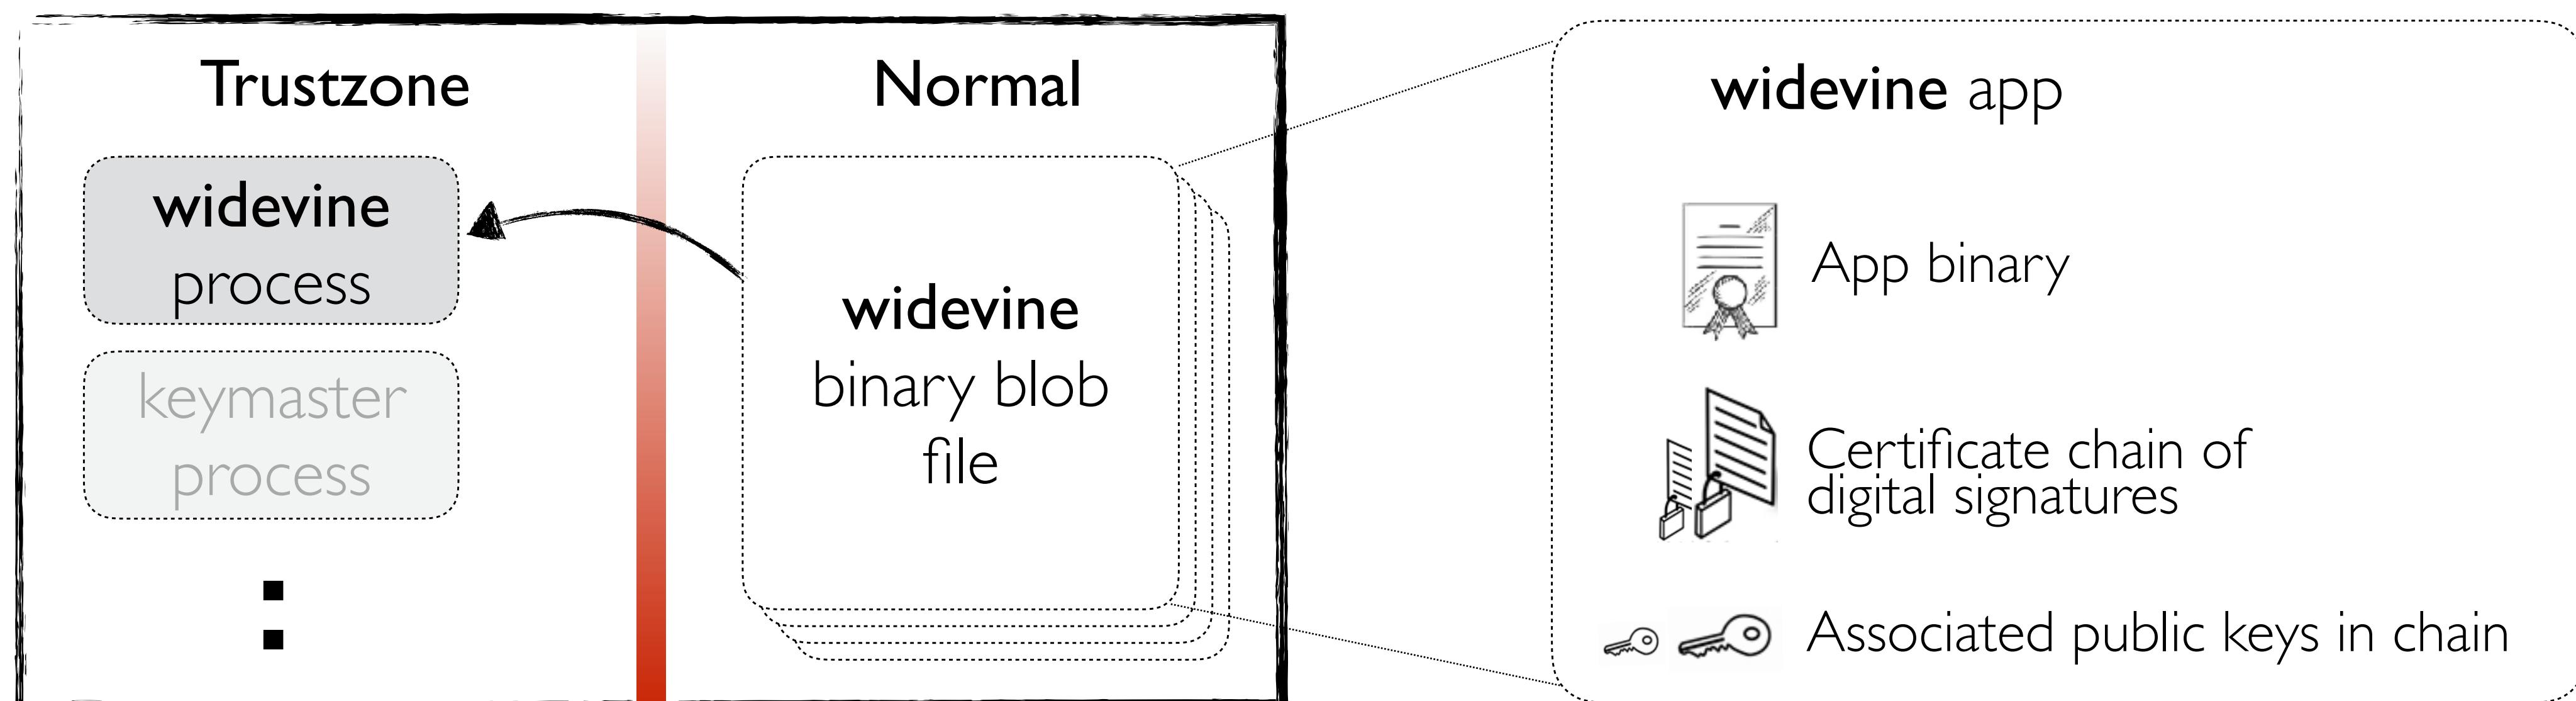
# Outline

---

- I. DVFS and Deep Dive into Hardware Regulators
- II. The CLKSCREW Attack
- III. Trustzone Attack I: Secret AES Key Inference
- IV. Trustzone Attack 2:Tricking RSA Signature Validation**
- V. Concluding Remarks

# Real-world Apps in Trustzone

- Apps running in Trustzone are building blocks for security
  - Eg: widevine (DRM), keymaster (hardware-backed key storage)
- Trustzone apps loaded from binary blob files at runtime
- Trustzone OS checks for a valid RSA signature before loading app



# RSA Signature Attack: Threat Model

---

To attack: RSA signature chain verification routine in Trustzone

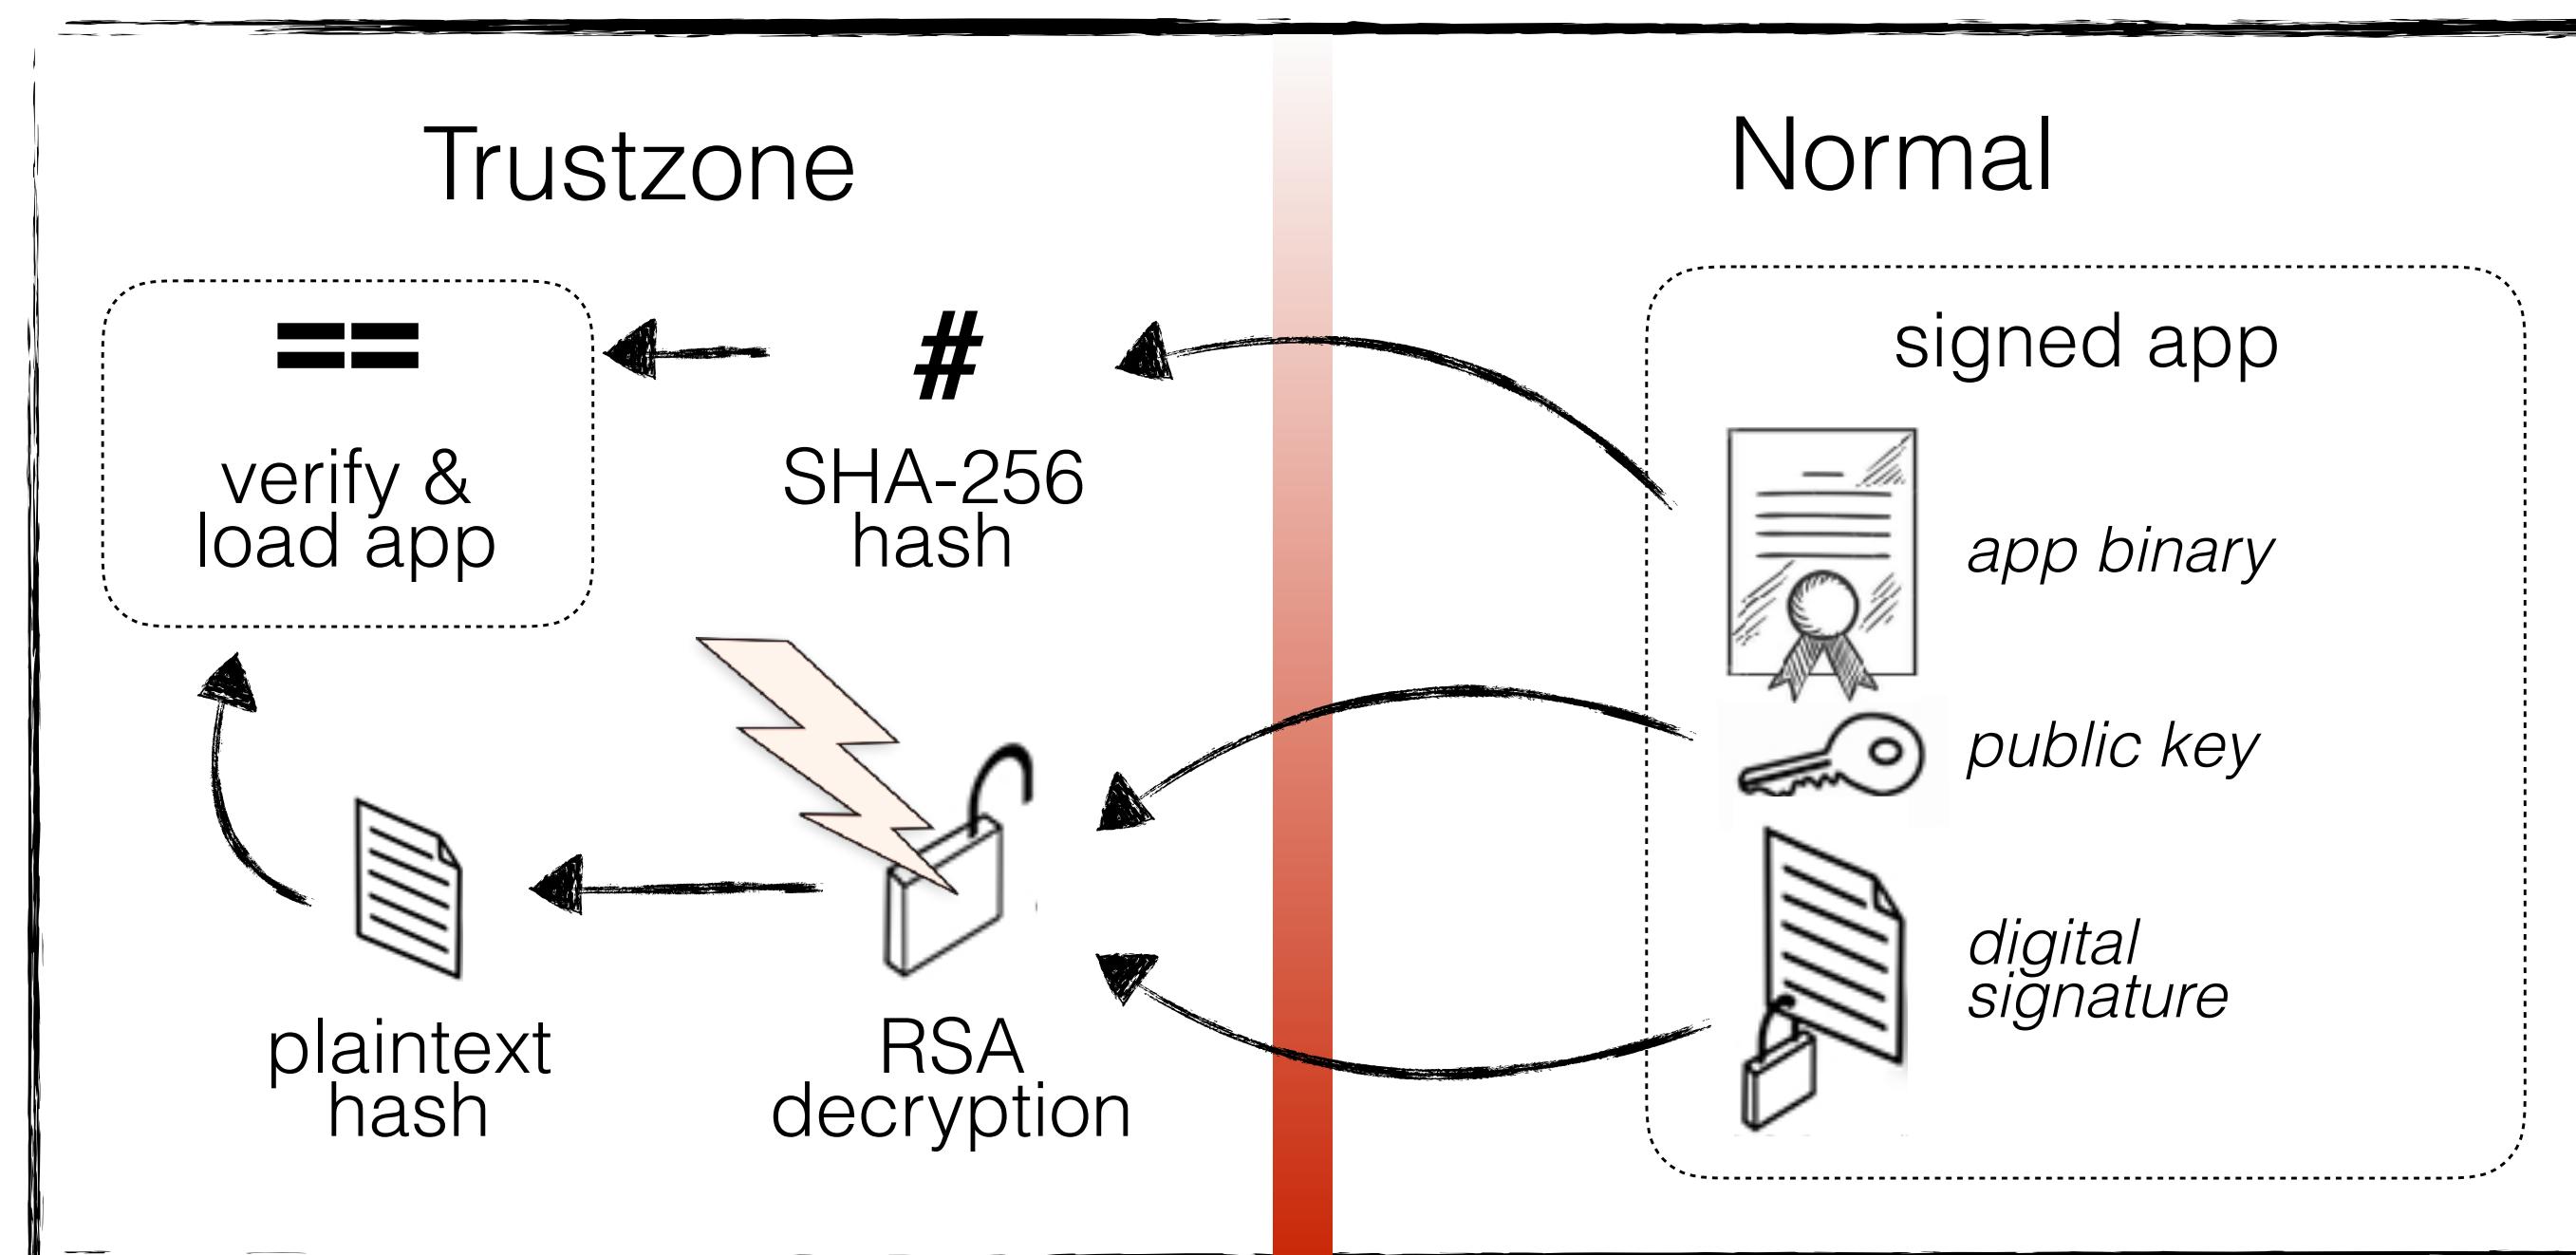
Attacker's goal: Trick routine into accepting a self-signed binary

Attacker's capabilities:

- 1) Can repeatedly invoke Trustzone to load app
- 2) Has software access to hardware regulators
- 3) Know when the app has been loaded successfully

# RSA Signature Attack: Summary

- Idea:
- Self-sign an app binary and invoke the app loading
  - Inject fault during signature verification
  - Corrupt RSA modulus used at runtime



# Attack Exploration and Formulation

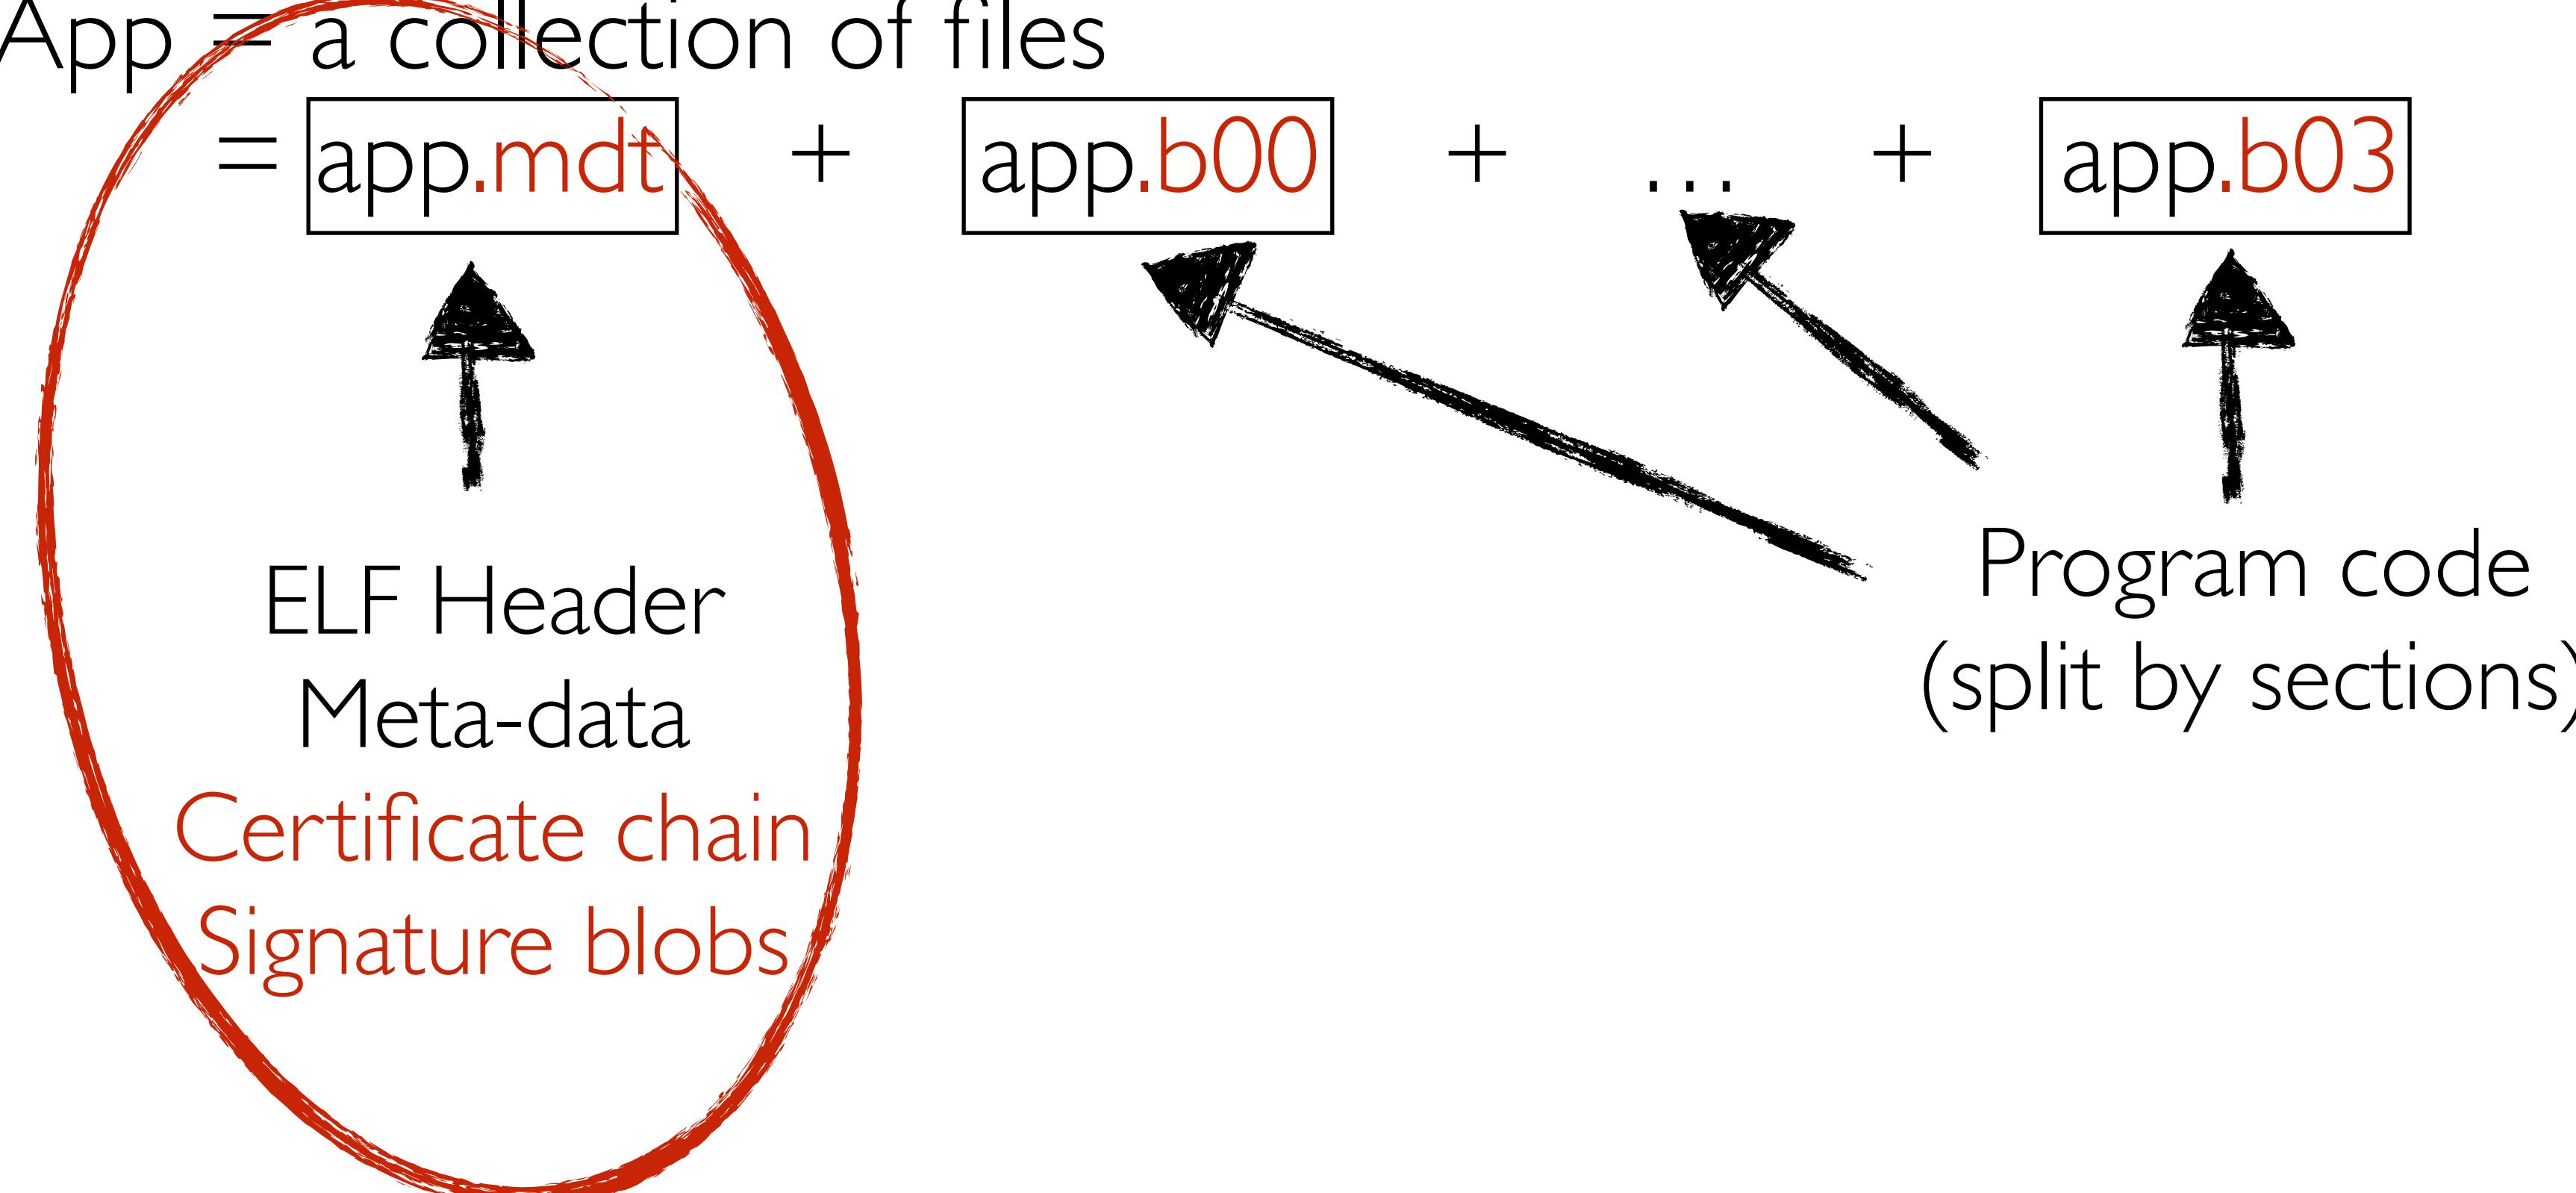
---

**Trustzone Apps:** How to craft self-signed app binary files?

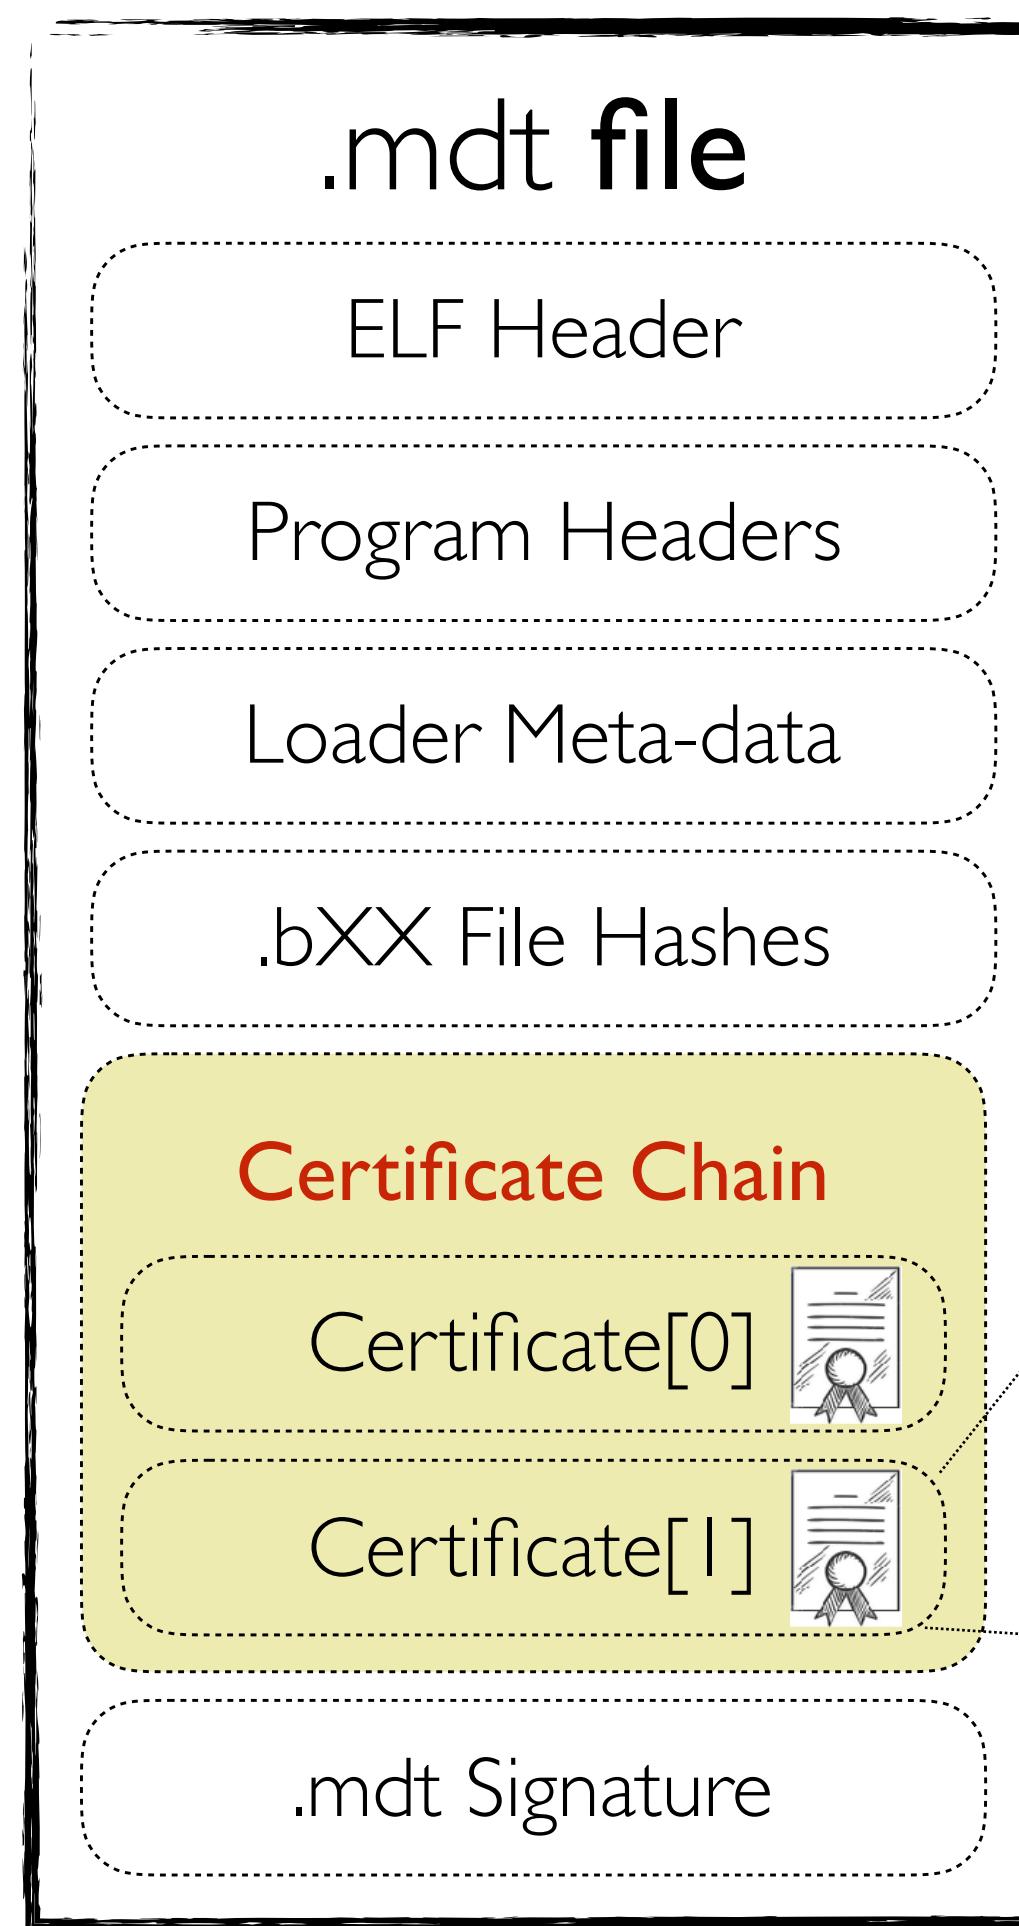
**Trustzone OS:** How, where and when to inject the CLKSCREW fault?

# Trustzone Apps: Format of Binary Files

- Trustzone firmware updates contain the Trustzone app binary files
- App = a collection of files



# Trustzone Apps: Format of Binary Files



Certificate stored in proprietary Motorola HAB (“High Assurance Boot”) binary format

## Actual Certificate[1] in <widevine.mdt>

**Subj:** O=Motorola Inc, OU=Motorola PKI, CN=CSF CA 637-I

**Issuer:** O=Motorola Inc, OU=Motorola PKI, CN=APP 637-I-2; ...

**Public exponent, e:** 0x10001

**Modulus, N:** c44dc735f6682a261a0b8545a62dd13df4c646a5ed...

**Signature:** 3cc1961f0d833a6197bd5537ee3f7d784dcf5dfb83b0...



Code: [https://github.com/0x0atang/clkscrew/blob/master/pycrypto/parse\\_mdt\\_certs.py](https://github.com/0x0atang/clkscrew/blob/master/pycrypto/parse_mdt_certs.py)

### Credits:

- Gal Beniamini: <http://bits-please.blogspot.com/2016/04/exploring-qualcomms-secure-execution.html>
- Tal Aloni: <http://vm1.duckdns.org/Public/Qualcomm-Secure-Boot/Qualcomm-Secure-Boot.htm>

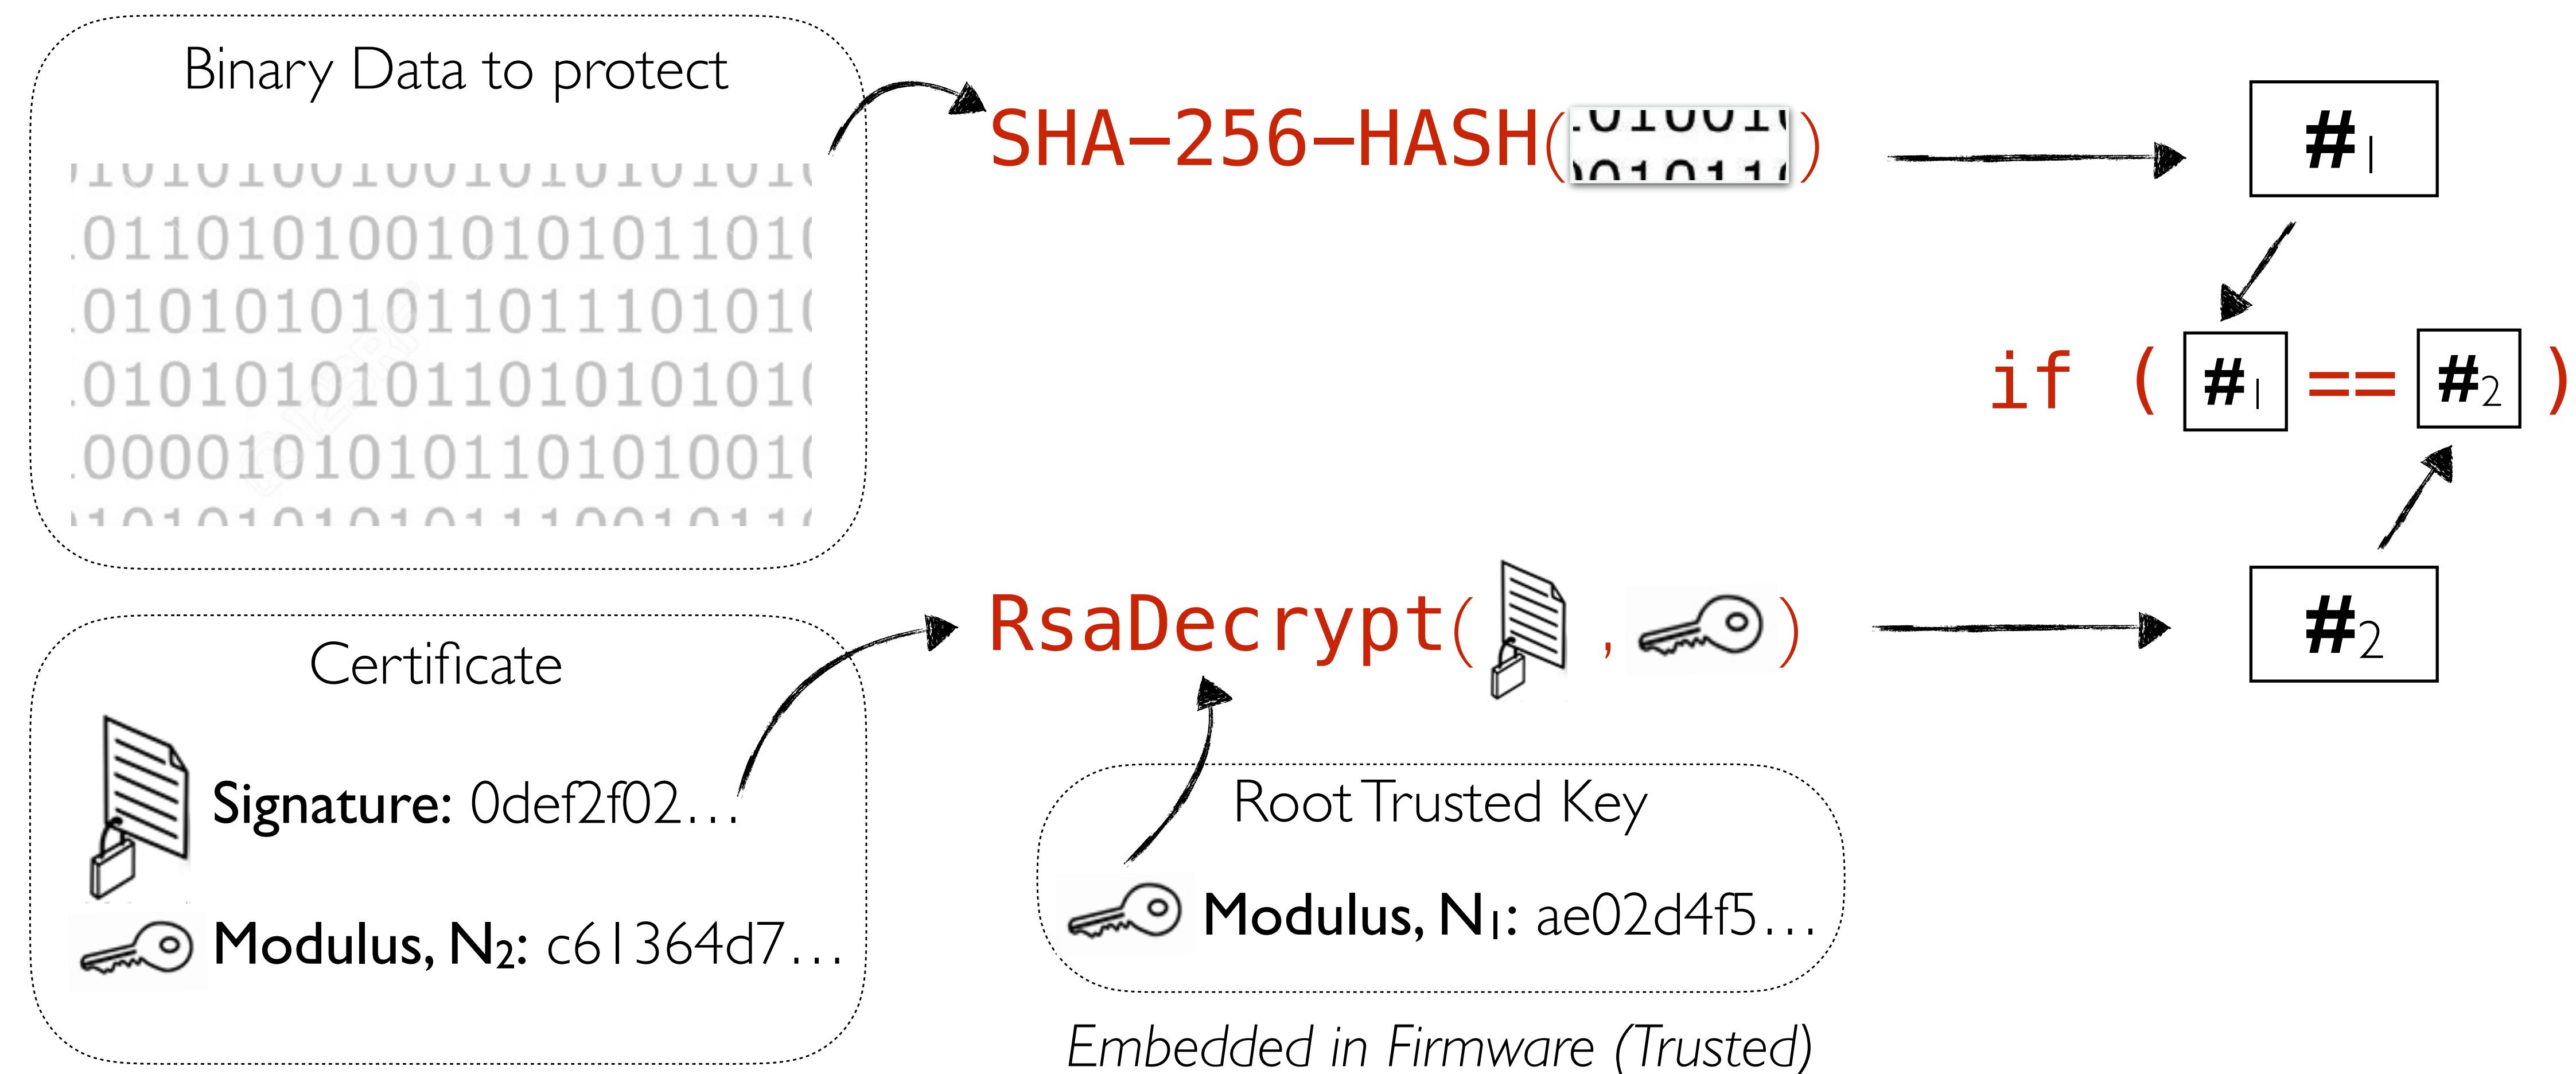
# Attack Exploration and Formulation

---

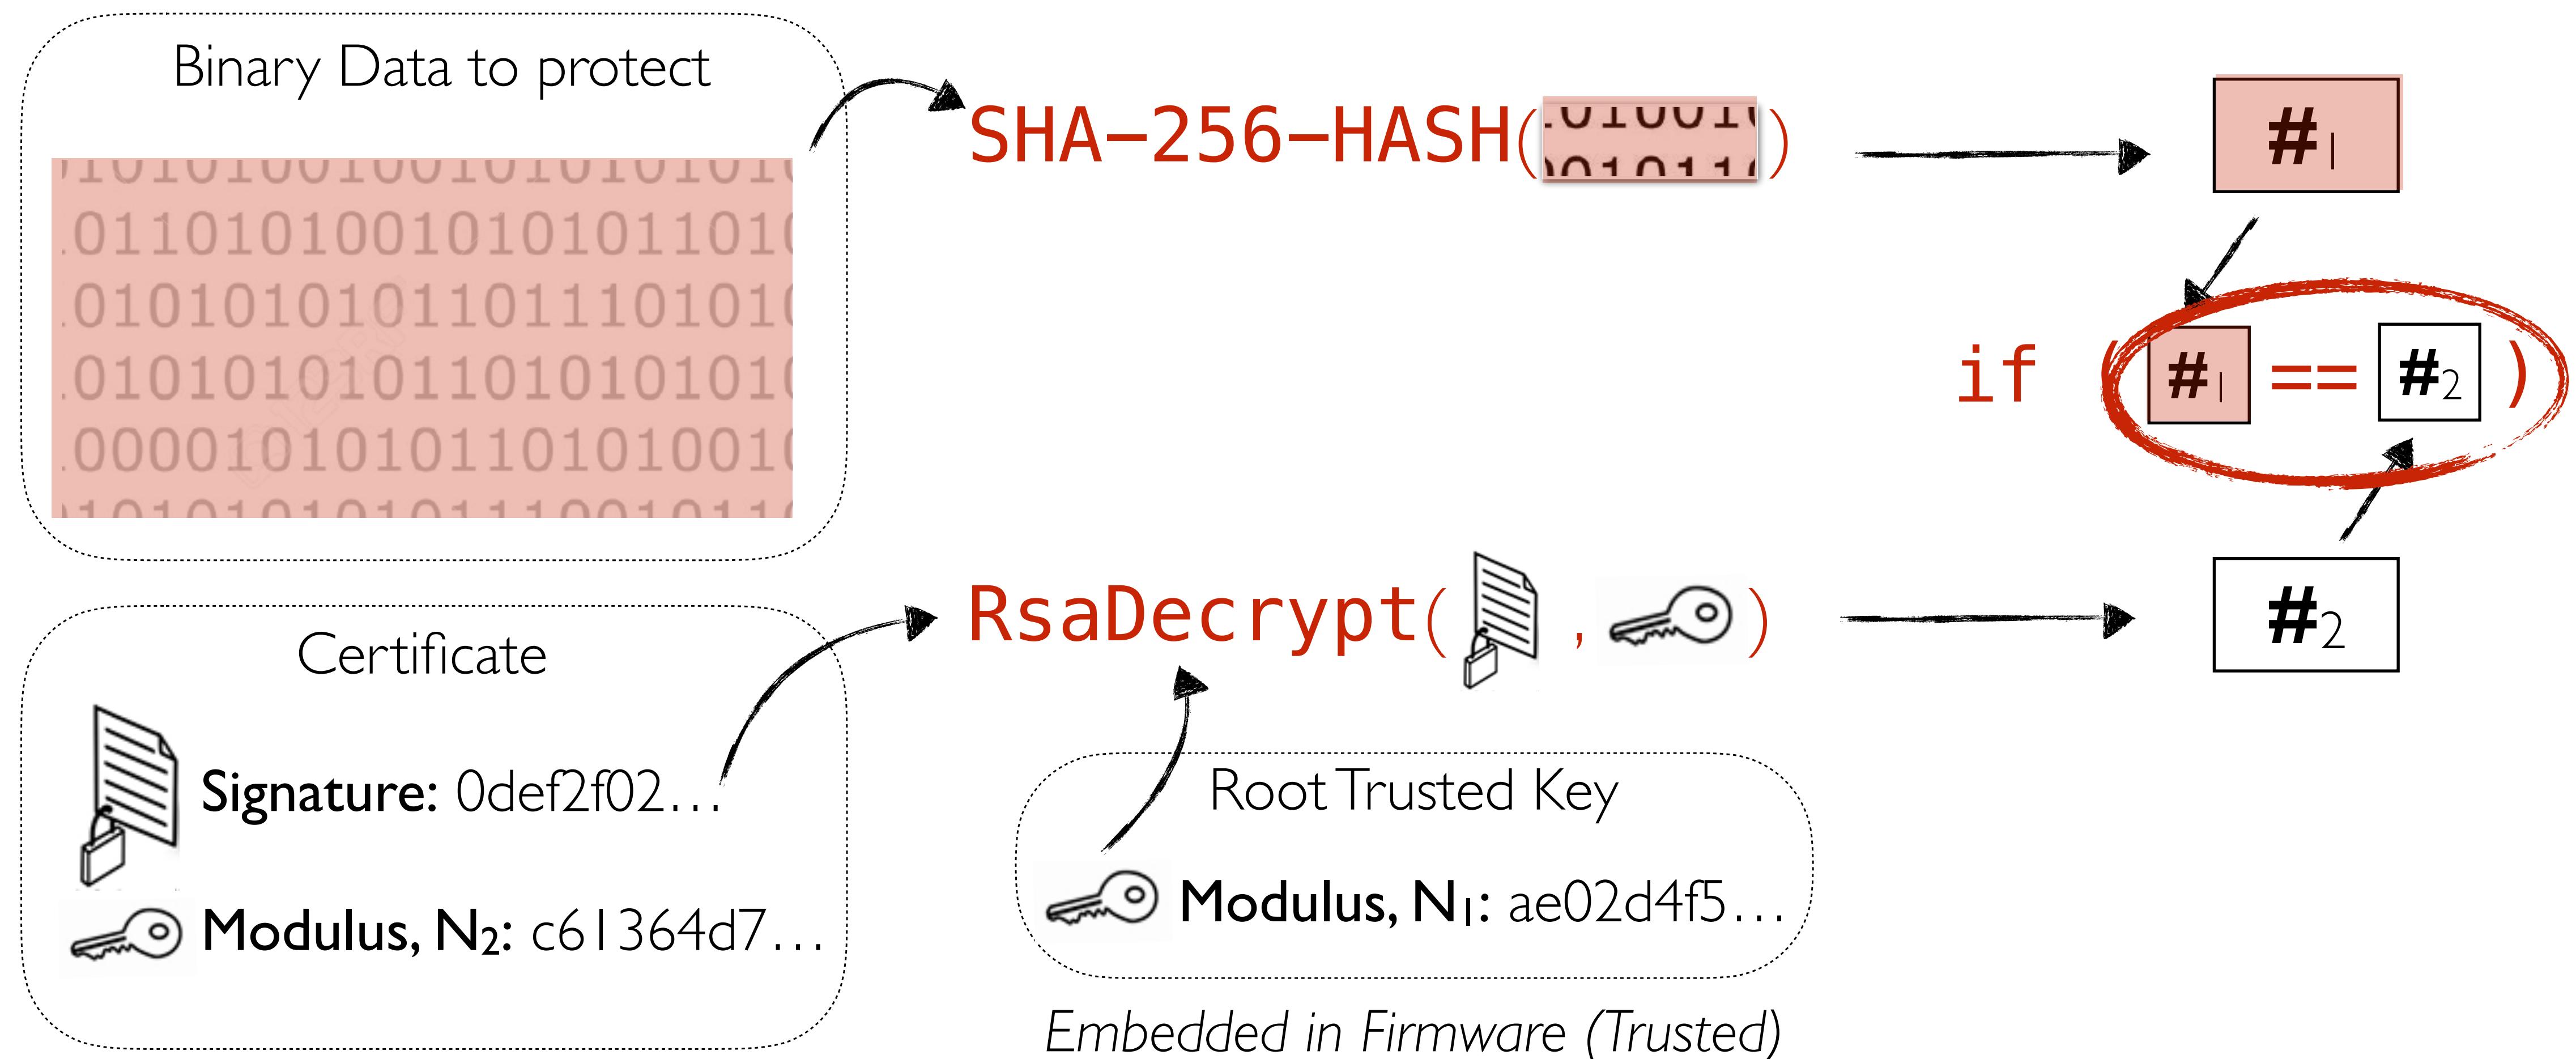
**Trustzone Apps:** How to craft self-signed app binary files?

**Trustzone OS:** How, where and when to inject the CLKSCREW fault?

# Quick Review: How signatures are verified

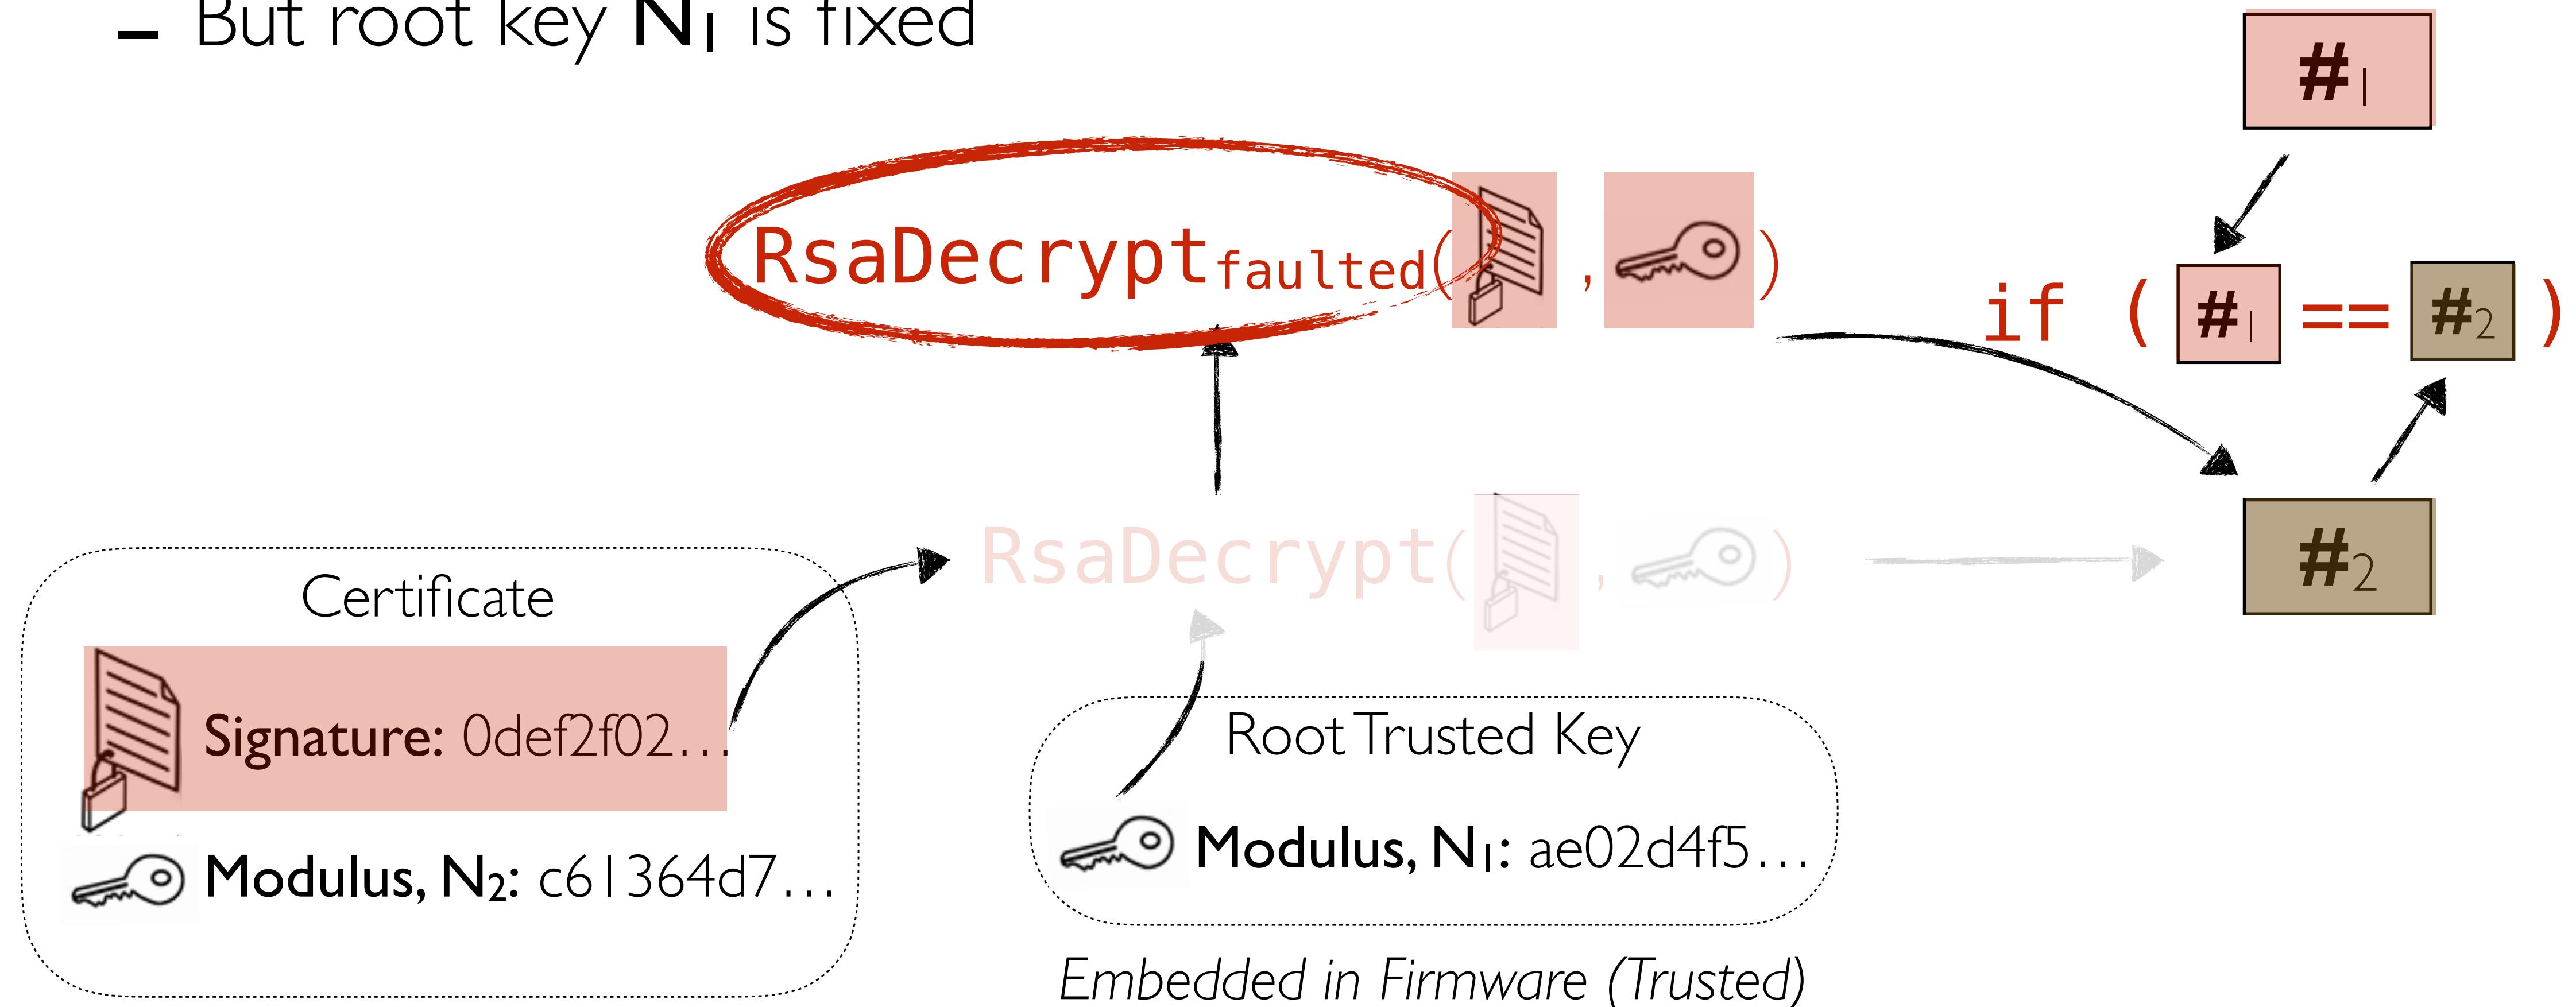


# What if we modify the binary data (Our App)?



# Why do we need to inject fault at runtime?

- We need  $\#_1 == \#_2$
- But root key  $N_1$  is fixed



# Digging deep into the firmware

```
LOAD:FE824E9C QWORD_FE824E9C DCB V ; DATA XREF: SUB_FE84C9/47D940  
LOAD:FE824E9C  
LOAD:FE824EA0 SRK_modulus DCB OxAE, 2, 0xD4, 0xF5, 0xC5, 0x6F, 0x3C, 0xE0, 0xE3  
LOAD:FE824EA0 DCB OxE7, 0x72, 0xE6, 0x7F, 0x52, 0x2E, 0xE8, 0xBF, 0xCE  
LOAD:FE824EA0 DCB Ox3A, 0xC4, 0x24, 0x2B, 0xCD, 0x5F, 0x1A, 0x96, 0x72  
LOAD:FE824EA0 DCB Ox47, 0x62, 0x5D, 0xC6, 0x4B, 0x71, 0xF4, 0xD0, 0xD4  
LOAD:FE824EA0 DCB 0xF, 0x9E, 0x3B, 0x4D, 0x44, 0x21, 0xB9, 0x5F, 0x3E  
LOAD:FE824EA0 DCB OxAO, 0xE, 0x58, 0xEC, 0xEC, 0x3E, 0xD4, 0xC7, 0x3A  
LOAD:FE824EA0 DCB OxD5, 0x58, 0xF7, 0x3C, 0x9C, 0x2A, 0xC4, 0xE, 0xD7  
LOAD:FE824EA0 DCB Ox82, 0x23, 0x7A, 0xB6, 2, 0xA, 0xA6, 0xB6, 0xFE, 0xB5  
LOAD:FE824EA0 DCB OxB7, 0x44, 0x2B, 0x44, 0x8B, 0xD, 0x57, 0x10, 0x42  
LOAD:FE824EA0 DCB OxC4, 0x2E, 0x4F, 0x72, 0xF6, 0xC9, 0x88, 0x7B, 0xF8  
LOAD:FE824EA0 DCB OxB2, 0xB5, 0x58, 0x7C, 0x43, 0x7A, 0xA0, 0x4D, 0xC7  
LOAD:FE824EA0 DCB OxA8, 0x3E, 0x3A, 0xD5, 0xC0, 0x7C, 0xA3, 0xEA, 0x30  
LOAD:FE824EA0 DCB OxA4, 0x6F, 0x5B, 6, 0x4E, 0xB9, 0x9F, 0xE6, 0xF0  
LOAD:FE824EA0 DCB 0xFD, 0xAF, 0xB9, 0xBC, 0x62, 0x6E, 0x74, 0x97, 0x5B  
LOAD:FE824EA0 DCB 7, 0x13, 0x8E, 0x39, 0xD2, 0x13, 0xBD, 0x7C, 0x6C  
LOAD:FE824EA0 DCB OxDC, 0xF2, 0x8, 0xE4, 0x50, 0xDD, 0x60, 0x3C, 0x57  
LOAD:FE824EA0 DCB Ox77, 0xD0, 0x5B, 0x99, 0xC, 0xD0, 0xCD, 0xAE, 0x42  
LOAD:FE824EA0 DCB OxBA, 0x2F, 0x28, 0xA5, 0xCE, 0x1D, 0x13, 0xCE, 0xCB  
LOAD:FE824EA0 DCB OxB2, 0x68, 0x38, 0xD2, 0xF5, 0xEC, 0xDC, 0xF1, 0xB9  
LOAD:FE824EA0 DCB OxF9, 0xFF, 0x56, 0x27, 0x47, 0x66, 0xDE, 0x59, 0x4F  
LOAD:FE824EA0 DCB OxCO, 0xA5, 0xF4, 0x48, 0xEC, 0xB7, 0x61, 0xBA, 0xF1  
LOAD:FE824EA0 DCB Ox82, 0xB3, 0x96, 0x2A, 0xC1, 0x19, 0x12, 0xD3, 0xE4  
LOAD:FE824EA0 DCB Ox2F, 0xFE, 0x91, 0xBF, 0xFE, 0xFD, 0x8C, 0xE, 0x9F  
LOAD:FE824EA0 DCB Ox40, 0x3A, 0xE1, 0xD5, 0xEB, 0xBF, 0x92, 0x6E, 0xCD  
LOAD:FE824EA0 DCB Ox3B, 0xA9, 0x17, 0x3E, 0x41, 0xAF, 0x39, 0x28, 0xAC  
LOAD:FE824EA0 DCB Ox55, 0x5A, 2, 1, 0x81, 0x7C, 0x69, 0x14, 5, 0x3F  
LOAD:FE824EA0 DCB Ox1D, 0x27, 0x2D, 0x32, 0xAD, 0x8F, 0x85, 0xD6, 0x77  
LOAD:FE824EA0 DCB Ox65, 0xC3, 0xF0, 0x7E, 0xB1, 0x92, 0x76, 0x86, 0xC7  
LOAD:FE824EA0 DCB Ox54, 0xA3  
LOAD:FE824FA0 DCB OxDA ;  
LOAD:FE824FA1 DCB Ox7C ; |  
END-PROG
```

Root Trusted Key

Modulus, N<sub>1</sub>: ae02d4f5...

Embedded in Firmware (Trusted)

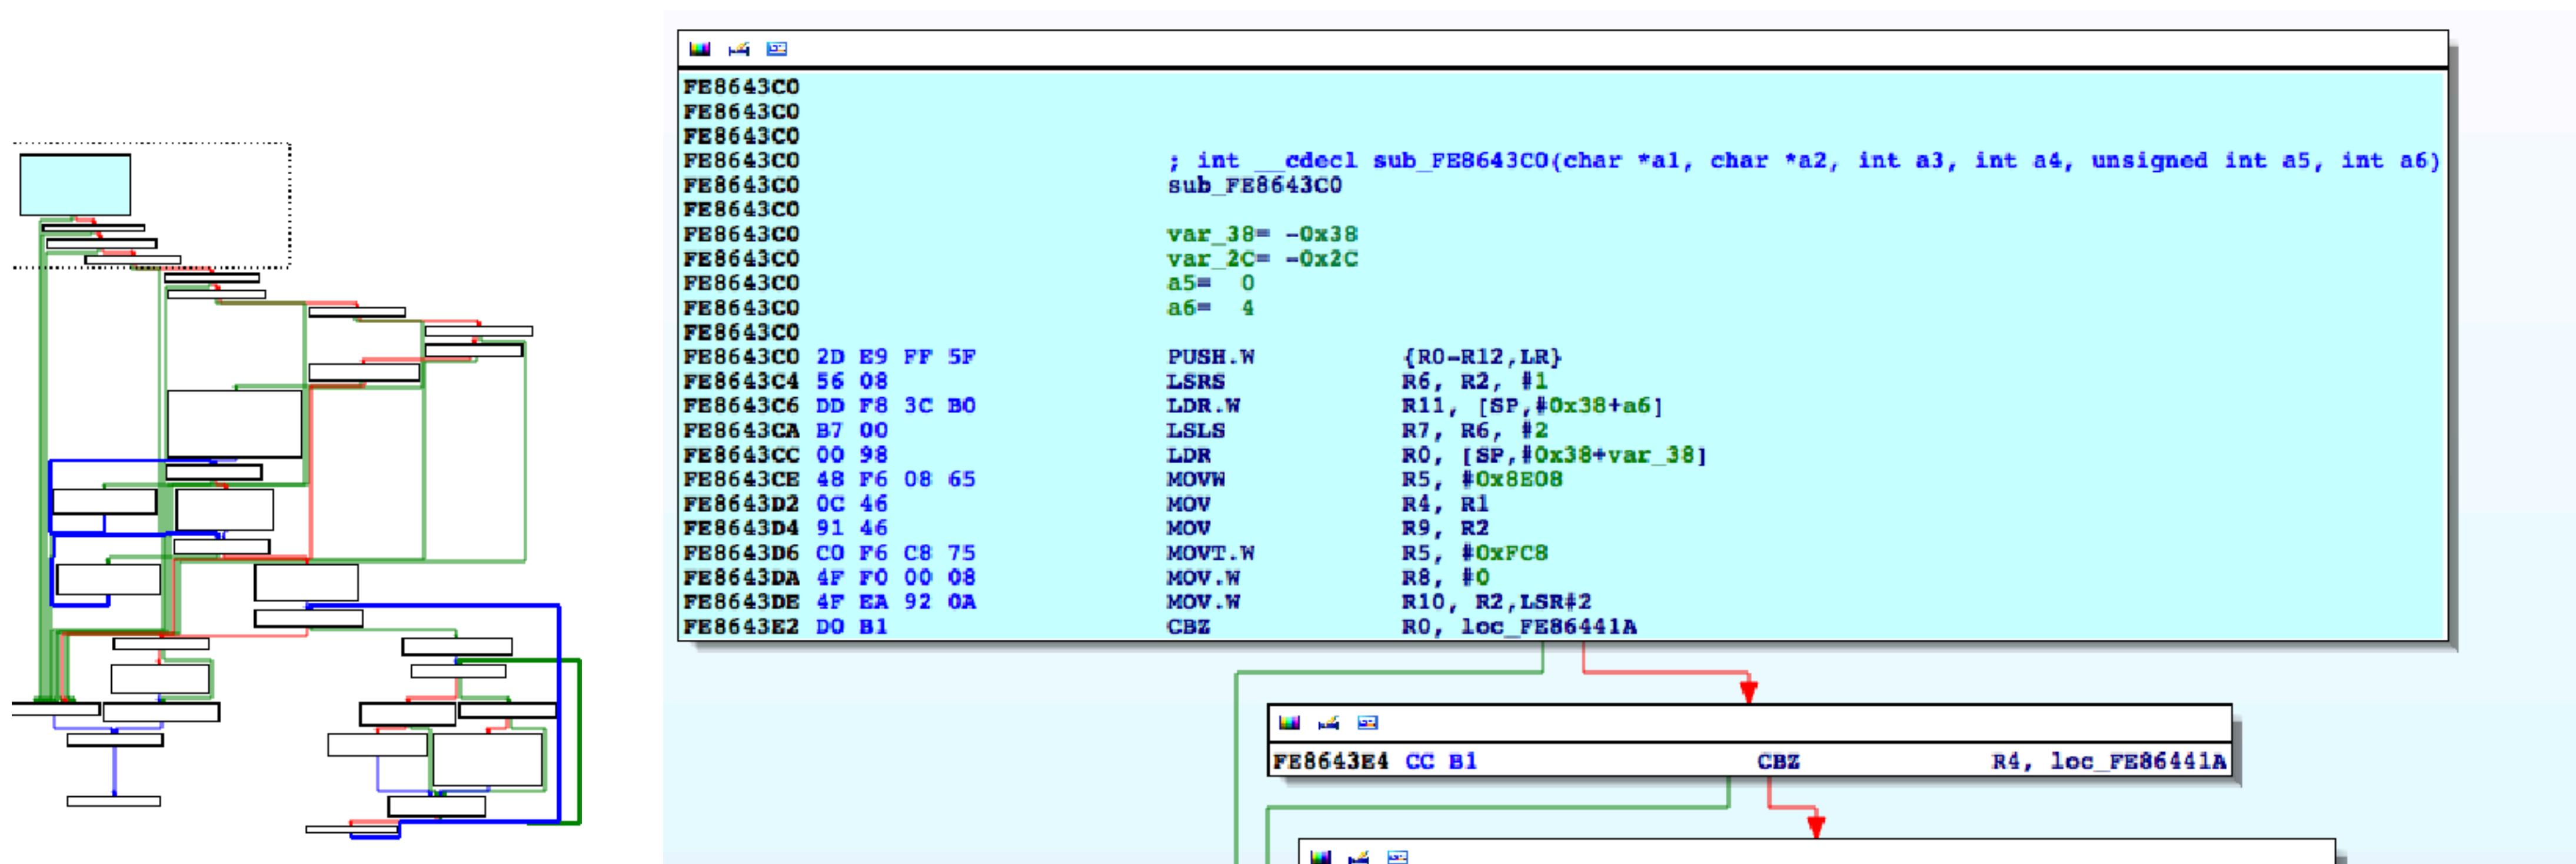
Super Root Key (SRK) modulus  
2048-bit

RSA-2048

# Digging deep into the firmware

signature      modulus      public exponent  
(0x10001)

RsaDecrypt( S , N , e ) =  $S^e \bmod N$  = decrypted hash



# Digging deep into the firmware

---

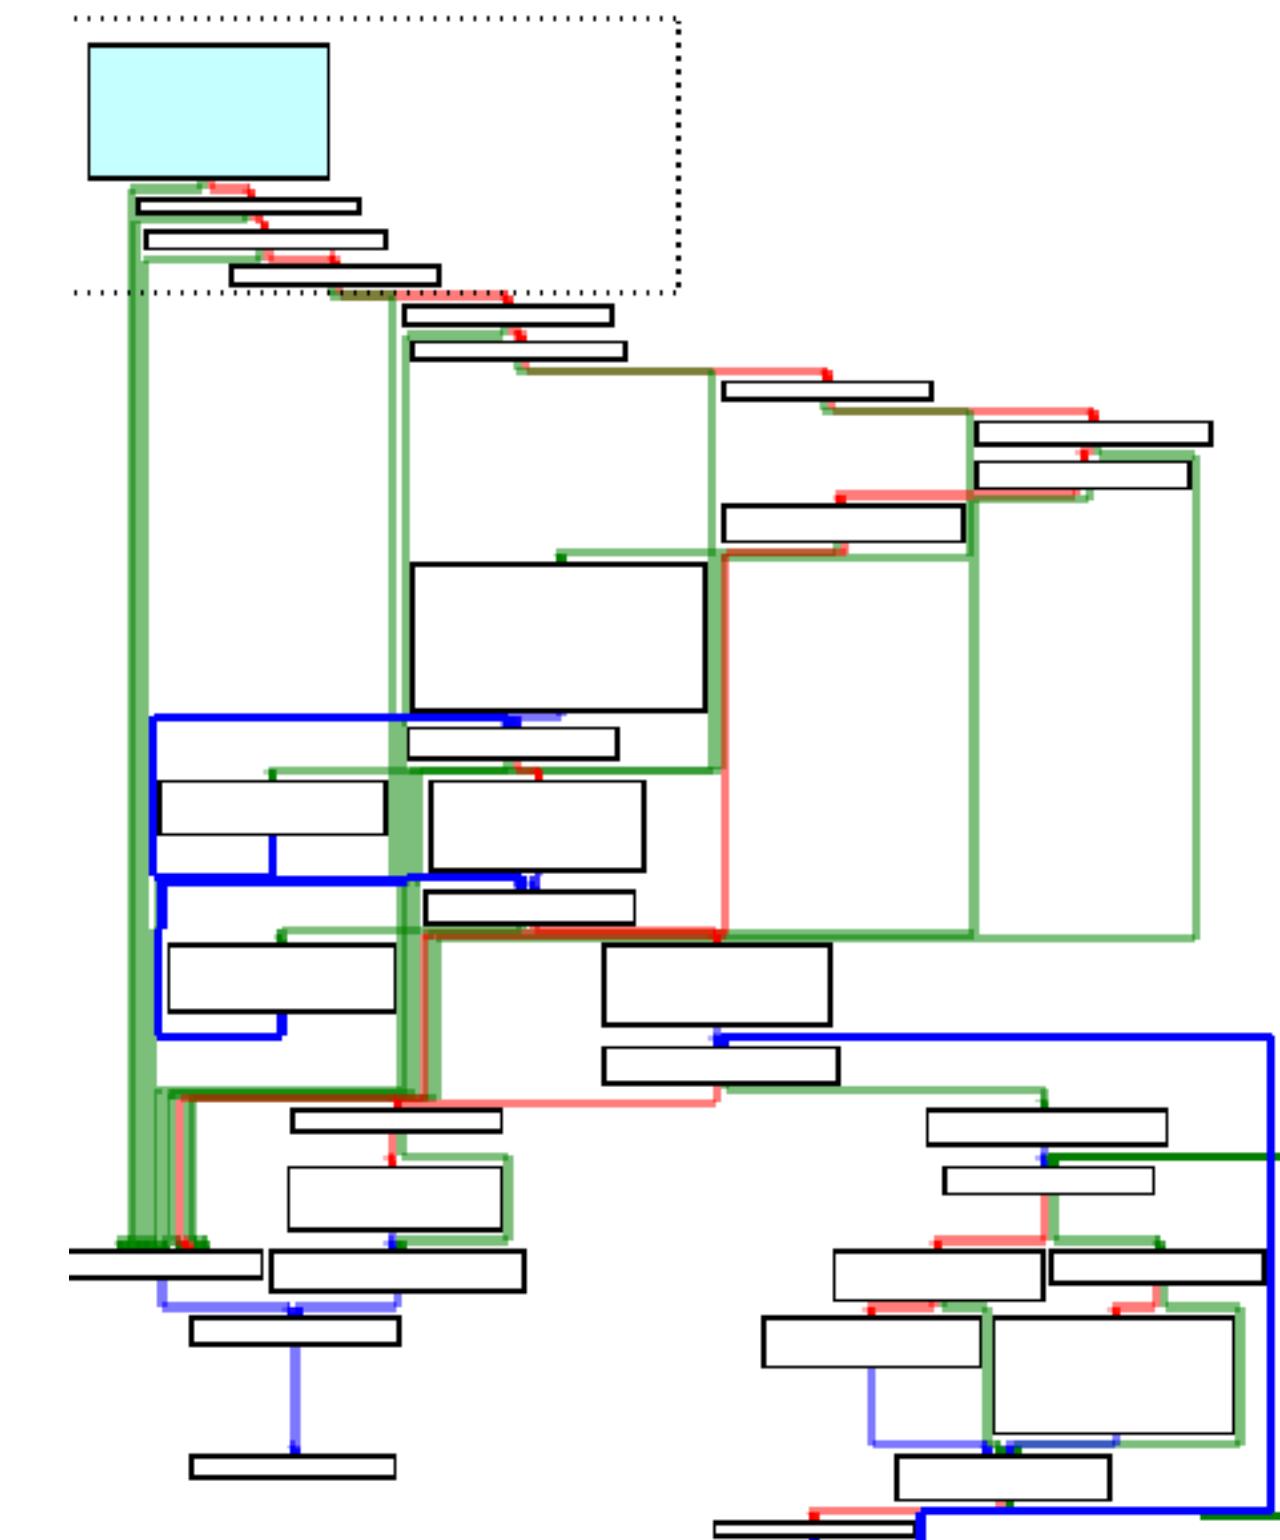
```

1: procedure DECRYPTSIG( $S, e, N$ )
2:    $r \leftarrow 2^{2048}$ 
3:    $R \leftarrow r^2 \bmod N$ 
4:    $N_{rev} \leftarrow \text{FLIPENDIANNESS}(N)$ 
5:    $r^{-1} \leftarrow \text{MODINVERSE}(r, N_{rev})$ 
6:    $\text{found\_first\_one\_bit} \leftarrow \text{false}$ 
7:   for  $i \in \{\text{bitlen}(e) - 1 \dots 0\}$  do
8:     if  $\text{found\_first\_one\_bit}$  then
9:        $x \leftarrow \text{MONTMULT}(x, x, N_{rev}, r^{-1})$ 
10:      if  $e[i] == 1$  then
11:         $x \leftarrow \text{MONTMULT}(x, a, N_{rev}, r^{-1})$ 
12:      end if
13:      else if  $e[i] == 1$  then
14:         $S_{rev} \leftarrow \text{FLIPENDIANNESS}(S)$ 
15:         $x \leftarrow \text{MONTMULT}(S_{rev}, R, N_{rev}, r^{-1})$ 
16:         $a \leftarrow x$ 
17:         $\text{found\_first\_one\_bit} \leftarrow \text{true}$ 
18:      end if
19:    end for
20:     $x \leftarrow \text{MONTMULT}(x, 1, N_{rev}, r^{-1})$ 
21:     $H \leftarrow \text{FLIPENDIANNESS}(x)$ 
22:    return  $H$ 
23: end procedure

```

---

RsaDecrypt(  $S$  ,  $N$  ,  $e$  )



signature	modulus	public exponent
$S$	$N$	(0x10001)

# Digging deep into the firmware

---

```
1: procedure DECRYPTSIG(S, e, N)
2:   r  $\leftarrow 2^{2048}$ 
3:   R  $\leftarrow r^2 \bmod N$ 
4:   Nrev  $\leftarrow \text{FLIPENDIANNESS}(N)$ 
5:   r-1  $\leftarrow \text{MODINVERSE}(r, N_{rev})$ 
6:   found_first_one_bit  $\leftarrow \text{false}$ 
7:   for i  $\in \{bitlen(e) - 1 \dots 0\} do
8:     if found_first_one_bit then
9:       x  $\leftarrow \text{MONTMULT}(x, x, N_{rev}, r^{-1})$ 
10:      if e[i] == 1 then
11:        x  $\leftarrow \text{MONTMULT}(x, a, N_{rev}, r^{-1})$ 
12:      end if
13:      else if e[i] == 1 then
14:        Srev  $\leftarrow \text{FLIPENDIANNESS}(S)$ 
15:        x  $\leftarrow \text{MONTMULT}(S_{rev}, R, N_{rev}, r^{-1})$ 
16:        a  $\leftarrow x$ 
17:        found_first_one_bit  $\leftarrow \text{true}$ 
18:      end if
19:    end for
20:    x  $\leftarrow \text{MONTMULT}(x, 1, N_{rev}, r^{-1})$ 
21:    H  $\leftarrow \text{FLIPENDIANNESS}(x)$ 
22:    return H
23: end procedure$ 
```

---

RsaDecrypt( *S* , *N* , *e* )

## Reverse engineering approaches:

- Statically via IDA
- Emulation using angr - <http://angr.io/>
- Dynamic code instrumentation on Trustzone code on actual phones  
(more details in future!)

# Digging deep into the firmware

```
1: procedure DECRYPTSIG( $S, e, N$ )
2:    $r \leftarrow 2^{2048}$ 
3:    $R \leftarrow r^2 \bmod N$ 
4:    $N_{rev} \leftarrow \text{FLIPENDIАНNESS}(N)$ 
5:    $r^{-1} \leftarrow \text{MODINVERSE}(r, N_{rev})$ 
6:    $\text{found\_first\_one\_bit} \leftarrow \text{false}$ 
7:   for  $i \in \{\text{bitlen}(e) - 1 \dots 0\}$  do
8:     if  $\text{found\_first\_one\_bit}$  then
9:        $x \leftarrow \text{MONTMULT}(x, x, N_{rev}, r^{-1})$ 
10:      if  $e[i] == 1$  then
11:         $x \leftarrow \text{MONTMULT}(x, a, N_{rev}, r^{-1})$ 
12:      end if
13:      else if  $e[i] == 1$  then
14:         $S_{rev} \leftarrow \text{FLIPENDIАНNESS}(S)$ 
15:         $x \leftarrow \text{MONTMULT}(S_{rev}, R, N_{rev}, r^{-1})$ 
16:         $a \leftarrow x$ 
17:         $\text{found\_first\_one\_bit} \leftarrow \text{true}$ 
18:      end if
19:    end for
20:     $x \leftarrow \text{MONTMULT}(x, 1, N_{rev}, r^{-1})$ 
21:     $H \leftarrow \text{FLIPENDIАНNESS}(x)$ 
22:    return  $H$ 
23: end procedure
```

**RsaDecrypt(  $S$  ,  $N$  ,  $e$  )**

- Computes modular exponentiation:  $S^e \bmod N$
- Implemented with an efficient form of multiplication called **Montgomery Multiplication**<sup>[I]</sup>

$$\text{MONTMULT}(x, y, N, r^{-1}) \leftarrow x \cdot y \cdot r^{-1} \bmod N$$

- Uses a memory-intensive function that reverses memory buffers

$$S_{rev} \leftarrow \text{FLIPENDIАНNESS}(S)$$

[I] KOC, C. K. High-speed RSA implementation. Tech. rep., Technical Report, RSA Laboratories, 1994.

# Digging deep into the firmware

```
1: procedure DECRYPTSIG( $S, e, N$ )
2:    $r \leftarrow 2^{2048}$ 
3:    $R \leftarrow r^2 \bmod N$ 
4:    $N_{rev} \leftarrow \text{FLIPENDIANCESS}(N)$  (circled)
5:    $r^{-1} \leftarrow \text{MODINVERSE}(r, N_{rev})$ 
6:    $\text{found\_first\_one\_bit} \leftarrow \text{false}$ 
7:   for  $i \in \{\text{bitlen}(e) - 1 \dots 0\}$  do
8:     if  $\text{found\_first\_one\_bit}$  then
9:        $x \leftarrow \text{MONTMULT}(x, x, N_{rev}, r^{-1})$ 
10:      if  $e[i] == 1$  then
11:         $x \leftarrow \text{MONTMULT}(x, a, N_{rev}, r^{-1})$ 
12:      end if
13:      else if  $e[i] == 1$  then
14:         $S_{rev} \leftarrow \text{FLIPENDIANCESS}(S)$ 
15:         $x \leftarrow \text{MONTMULT}(S_{rev}, R, N_{rev}, r^{-1})$ 
16:         $a \leftarrow x$ 
17:         $\text{found\_first\_one\_bit} \leftarrow \text{true}$ 
18:      end if
19:    end for
20:     $x \leftarrow \text{MONTMULT}(x, 1, N_{rev}, r^{-1})$ 
21:     $H \leftarrow \text{FLIPENDIANCESS}(x)$ 
22:    return  $H$ 
23: end procedure
```

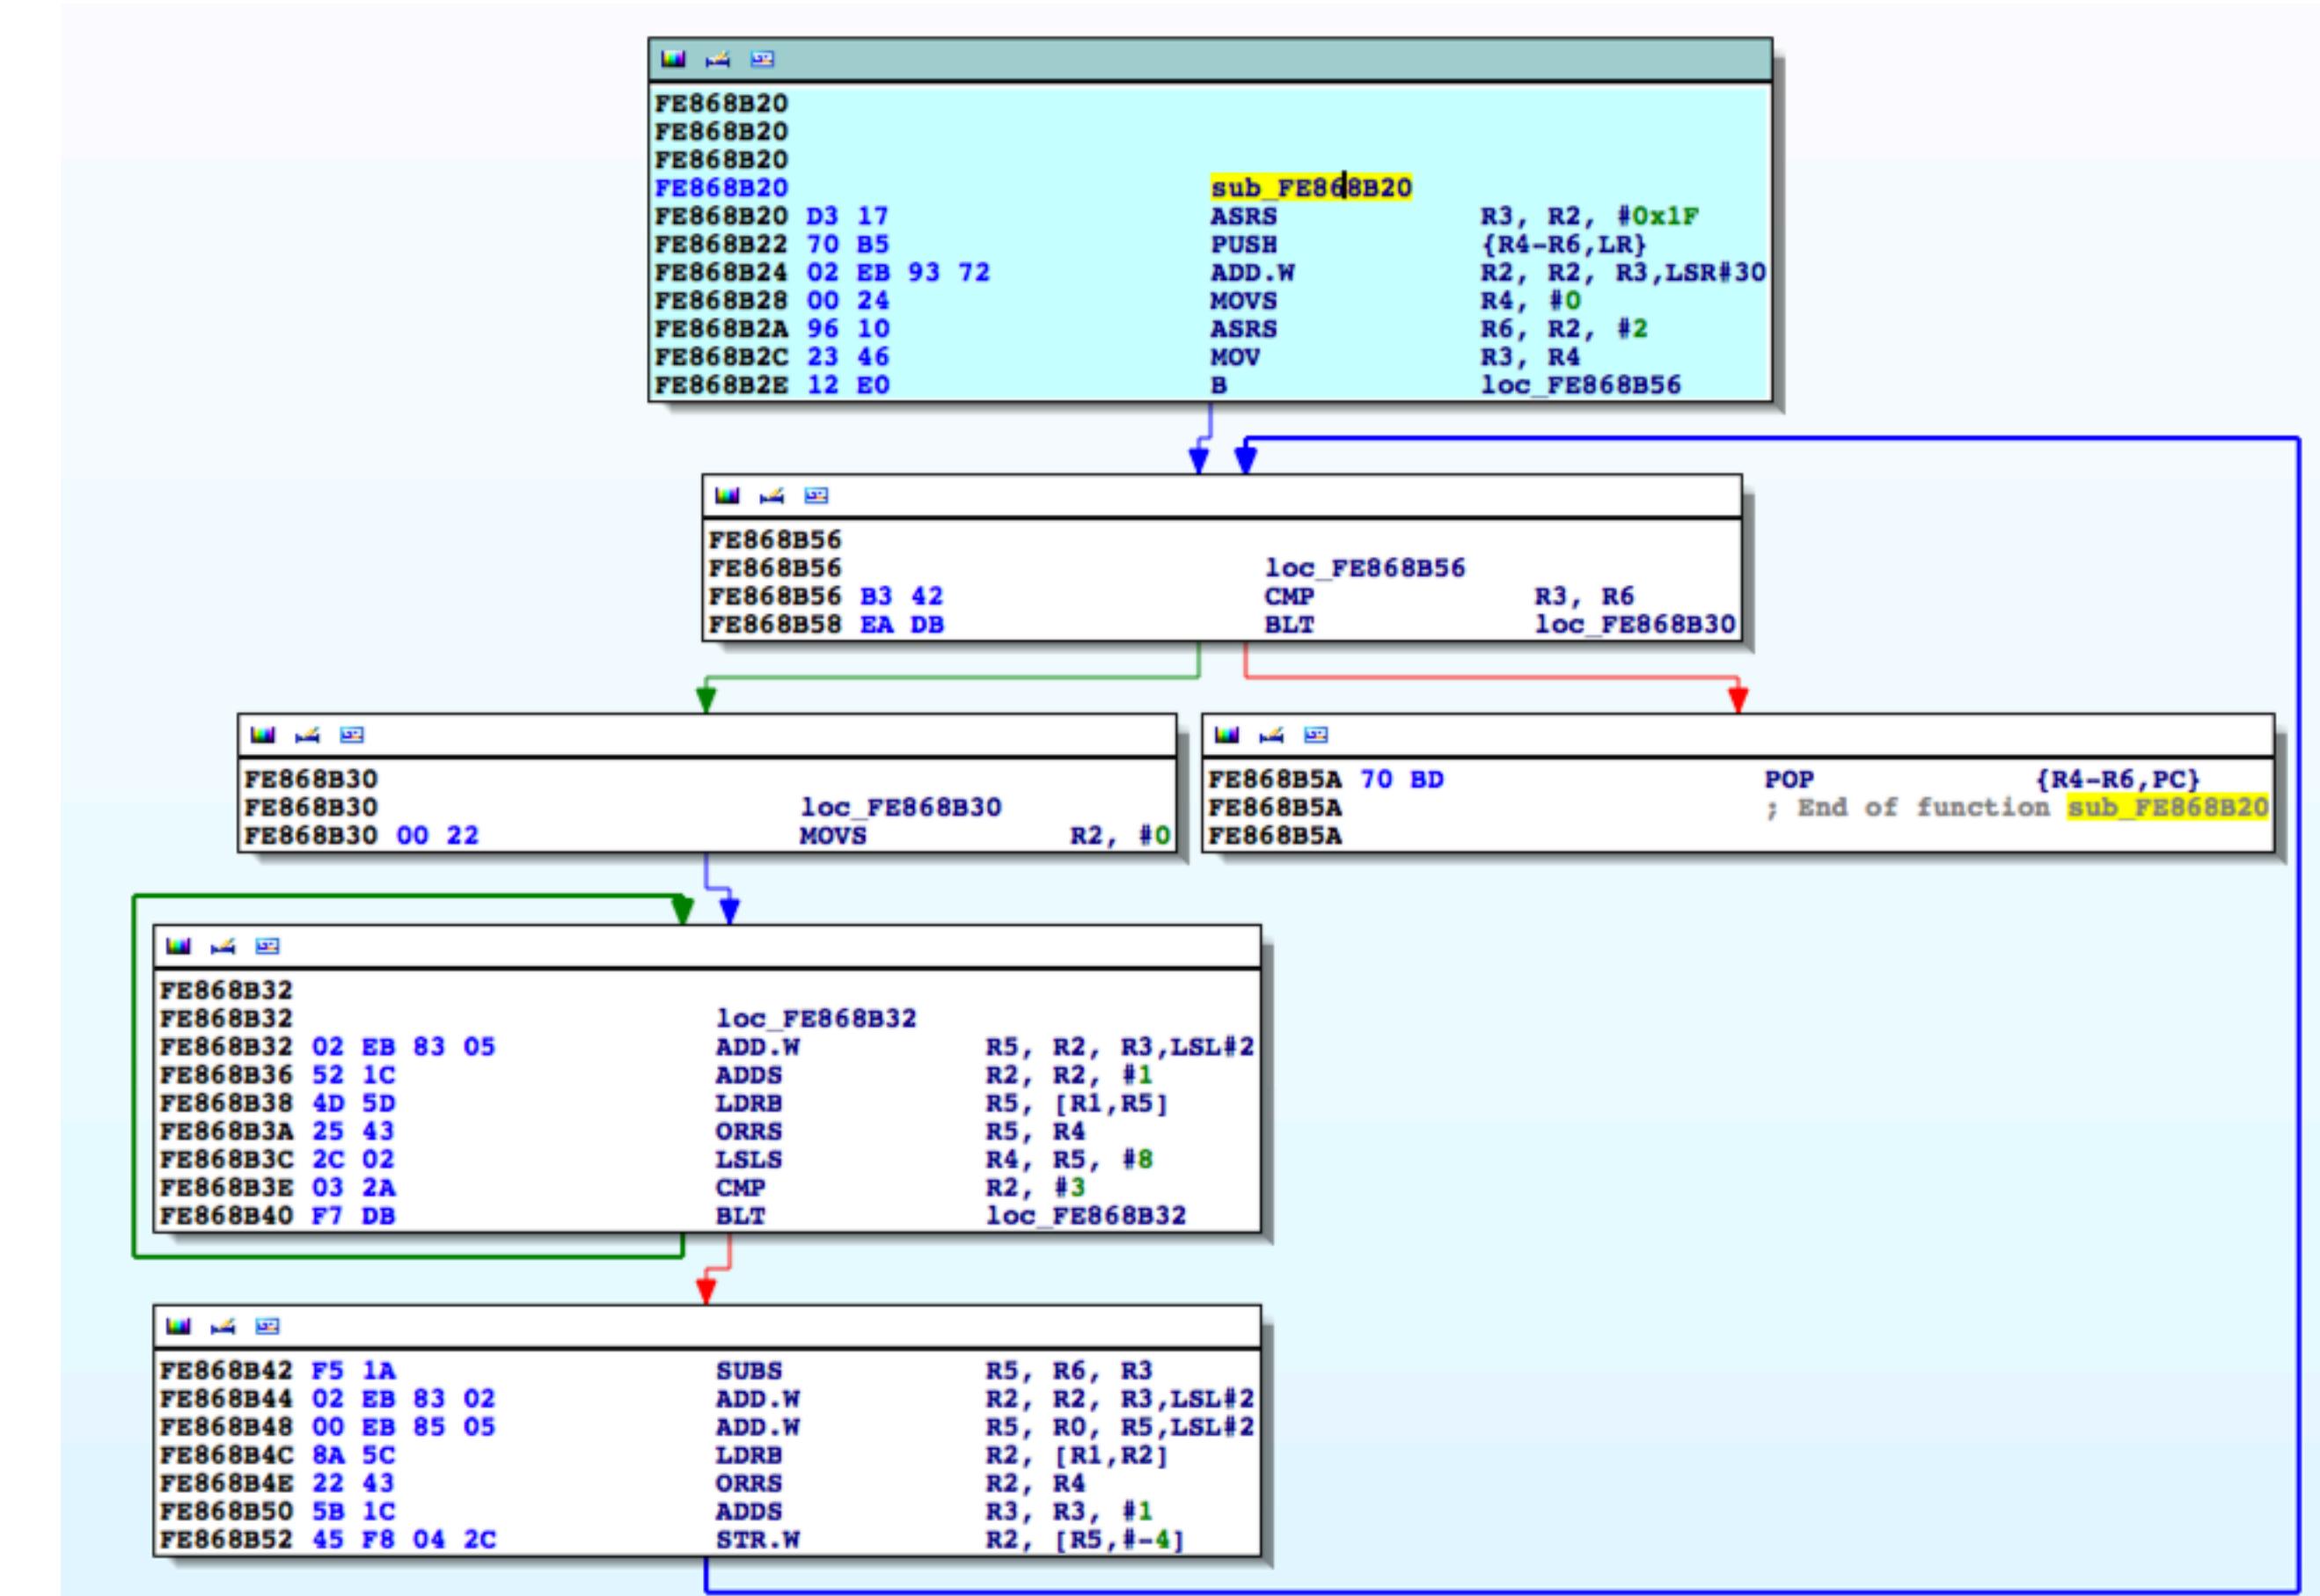
signature      modulus      public exponent  
(0x10001)

Where to inject the runtime fault?

- $N_{A,rev} \xleftarrow{\text{fault}} \text{FLIPENDIANCESS}(N)$
- Implemented with an efficient form of multiplication called **Montgomery Multiplication**  
 $\text{MONTMULT}(x, y, N, r^{-1}) \leftarrow x \cdot y \cdot r^{-1} \bmod N$
- Uses a memory-intensive function that reverses memory buffers  
 $S_{rev} \leftarrow \text{FLIPENDIANCESS}(S)$

# Corrupting FlipEndianness with Runtime Fault

```
1: procedure FLIPENDIANCESS(src)
2:   d ← 0
3:   dst ← {0}
4:   for i ∈ {0 .. len(src)/4 – 1} do
5:     for j ∈ {0 .. 2} do
6:       d ← (src[i * 4 + j] | d) ≪ 8
7:     end for
8:     d ← src[i * 4 + 3] | d
9:     k ← len(src) – i * 4 – 4
10:    dst[k .. k + 3] ← d
11:  end for
12:  return dst
13: end procedure
```



# Corrupting FlipEndianness with Runtime Fault

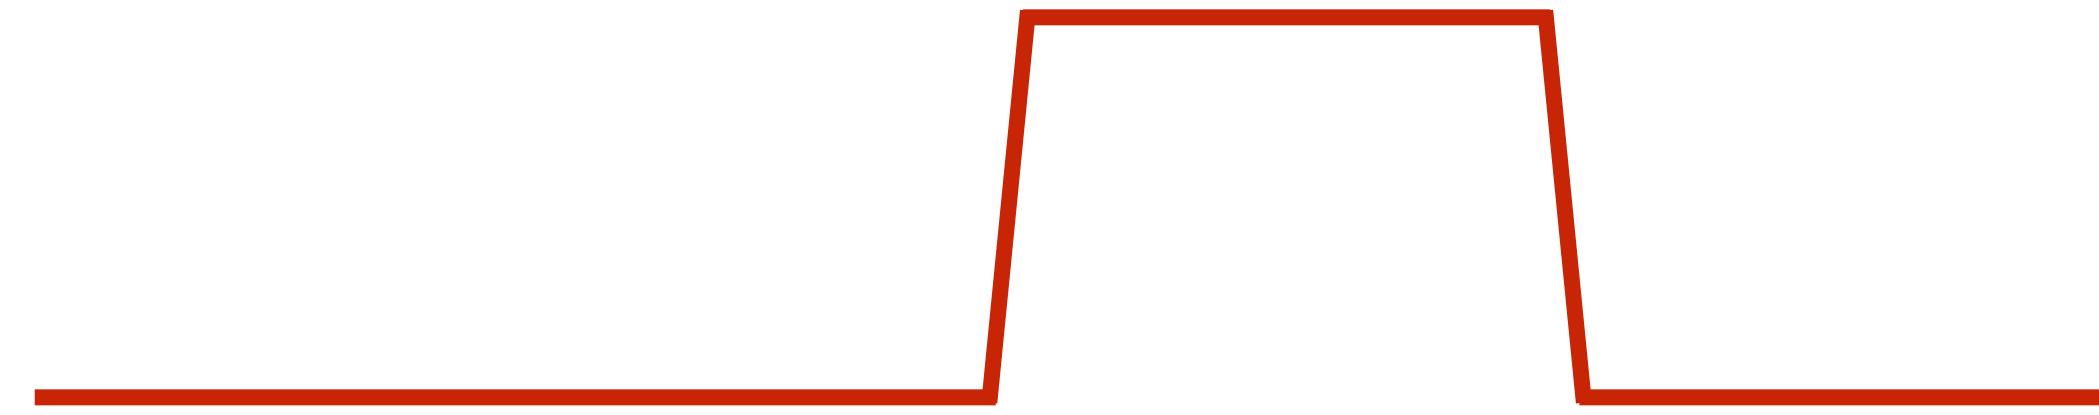
Base voltage: 1.055V

High frequency: 4.10GHz

Low frequency: 2.68GHz

Fault injection duration:

5 no-op loops (~0.287 μsec)



Expected modulus: ... bc099b4a ...

Faulty modulus used: ... bc09**4a**4a ...



Demo: <https://asciinema.org/a/5vvn3s9nzula930xui1z7tg65>



Code: [https://github.com/0x0atang/clkscrew/blob/master/faultmin\\_SD805/](https://github.com/0x0atang/clkscrew/blob/master/faultmin_SD805/)

# Digging deep into the firmware

```
1: procedure DECRYPTSIG( $S, e, N$ )
2:    $r \leftarrow 2^{2048}$ 
3:    $R \leftarrow r^2 \bmod N$ 
4:    $N_{rev} \leftarrow \text{FLIPENDIANCESS}(N)$ 
5:    $r^{-1} \leftarrow \text{MODINVERSE}(r, N_{rev})$ 
6:    $\text{found\_first\_one\_bit} \leftarrow \text{false}$ 
7:   for  $i \in \{\text{bitlen}(e) - 1 \dots 0\}$  do
8:     if  $\text{found\_first\_one\_bit}$  then
9:        $x \leftarrow \text{MONTMULT}(x, x, N_{rev}, r^{-1})$ 
10:      if  $e[i] == 1$  then
11:         $x \leftarrow \text{MONTMULT}(x, a, N_{rev}, r^{-1})$ 
12:      end if
13:      else if  $e[i] == 1$  then
14:         $S_{rev} \leftarrow \text{FLIPENDIANCESS}(S)$ 
15:         $x \leftarrow \text{MONTMULT}(S_{rev}, R, N_{rev}, r^{-1})$ 
16:         $a \leftarrow x$ 
17:         $\text{found\_first\_one\_bit} \leftarrow \text{true}$ 
18:      end if
19:    end for
20:     $x \leftarrow \text{MONTMULT}(x, 1, N_{rev}, r^{-1})$ 
21:     $H \leftarrow \text{FLIPENDIANCESS}(x)$ 
22:    return  $H$ 
23: end procedure
```

signature      modulus      public exponent  
(0x10001)

Where to inject the runtime fault?

$N_{A,rev} \xrightarrow{\text{fault}} \text{FLIPENDIANCESS}(N)$

– Implemented with an efficient form of multiplication called **Montgomery Multiplication**<sup>[I]</sup>

How to craft attack signature  $S_A'$ ?

$\text{DECRYPTSIG}(S_A', e, N) \xrightarrow{\text{fault}} H(C_A)$

# How to craft attack signature?

```
1: procedure DECRYPTSIG( $S, e, N$ )
2:    $r \leftarrow 2^{2048}$ 
3:    $R \leftarrow r^2 \bmod N$ 
4:    $N_{rev} \leftarrow \text{FLIPENDIANCESS}(N)$ 
5:    $r^{-1} \leftarrow \text{MODINVERSE}(r, N_{rev})$ 
6:    $\text{found\_first\_one\_bit} \leftarrow \text{false}$ 
7:   for  $i \in \{\text{bitlen}(e) - 1 \dots 0\}$  do
8:     if  $\text{found\_first\_one\_bit}$  then
9:        $x \leftarrow \text{MONTMULT}(x, x, N_{rev}, r^{-1})$ 
10:      if  $e[i] == 1$  then
11:         $x \leftarrow \text{MONTMULT}(x, a, N_{rev}, r^{-1})$ 
12:      end if
13:      else if  $e[i] == 1$  then
14:         $S_{rev} \leftarrow \text{FLIPENDIANCESS}(S)$ 
15:         $x \leftarrow \text{MONTMULT}(S_{rev}, R, N_{rev}, r^{-1})$ 
16:         $a \leftarrow x$ 
17:         $\text{found\_first\_one\_bit} \leftarrow \text{true}$ 
18:      end if
19:    end for
20:     $x \leftarrow \text{MONTMULT}(x, 1, N_{rev}, r^{-1})$ 
21:     $H \leftarrow \text{FLIPENDIANCESS}(x)$ 
22:    return  $H$ 
23: end procedure
```

Trickier than expected!!

Line 3:  $R \leftarrow r^2 \bmod N$

Line 4:  $N_{A,rev} \xleftarrow{\text{fault}} \text{FLIPENDIANCESS}(N)$

Line 15:  $\text{MONTMULT}(S_A', r^2 \bmod N, N_A, r_A^{-1})$

$$= S_A' \cdot (r^2 \bmod N) \cdot r_A^{-1} \bmod N_A$$

different moduli N's used

(More cryptanalysis in white paper...)



<https://github.com/0x0atang/clkscrew/blob/master/pycrypto/pycrypto.py>

# When (during execution) to inject fault?

---

- DecryptSig() is invoked *4 times* when verifying an app
  - 1) SRK.modulus => CERT[0]
  - 2) CERT[0].modulus => cert chain meta-data
  - 3) CERT[0].modulus => CERT[1]
  - 4) CERT[1].modulus => .mdt file hashes
- We need a way to profile when invocation (4) executes within Trustzone
- **Attack Enabler:** Memory accesses from outside Trustzone can evict cache lines used by Trustzone code



Side-channel-based  
cache profiling

# When (during execution) to inject fault?

---

- Instruction-cache **Prime+Probe** profiling more reliable than data-cache ones
  - More info on side-channel-based profiling attacks on ARM [1, 2, 3]
- I-Cache profiling not as convenient as D-Cache profiling
  - Instead of using memory read operations
  - Need to execute instructions at memory address congruent to cache sets we are monitoring
  - Create a JIT compiler — Given a list of cache sets to monitor
  - Allocate a large chunk of executable memory
  - Chain relative BR branch instructions in addresses congruent to monitored cache sets

[1] GRUSS, D., SPREITZER, R., AND MANGARD, S. Cache Template Attacks: Automating Attacks on Inclusive Last-Level Caches. USENIX 2015

[2] LIPP, M., GRUSS, D., SPREITZER, R., MAURICE, C., AND MANGARD, S. Armageddon: Cache attacks on mobile devices. USENIX 2016.

[3] ZHANG, X., XIAO, Y., AND ZHANG, Y. Return-Oriented Flush- Reload Side Channels on ARM and Their Implications for Android Devices. CCS 2016.

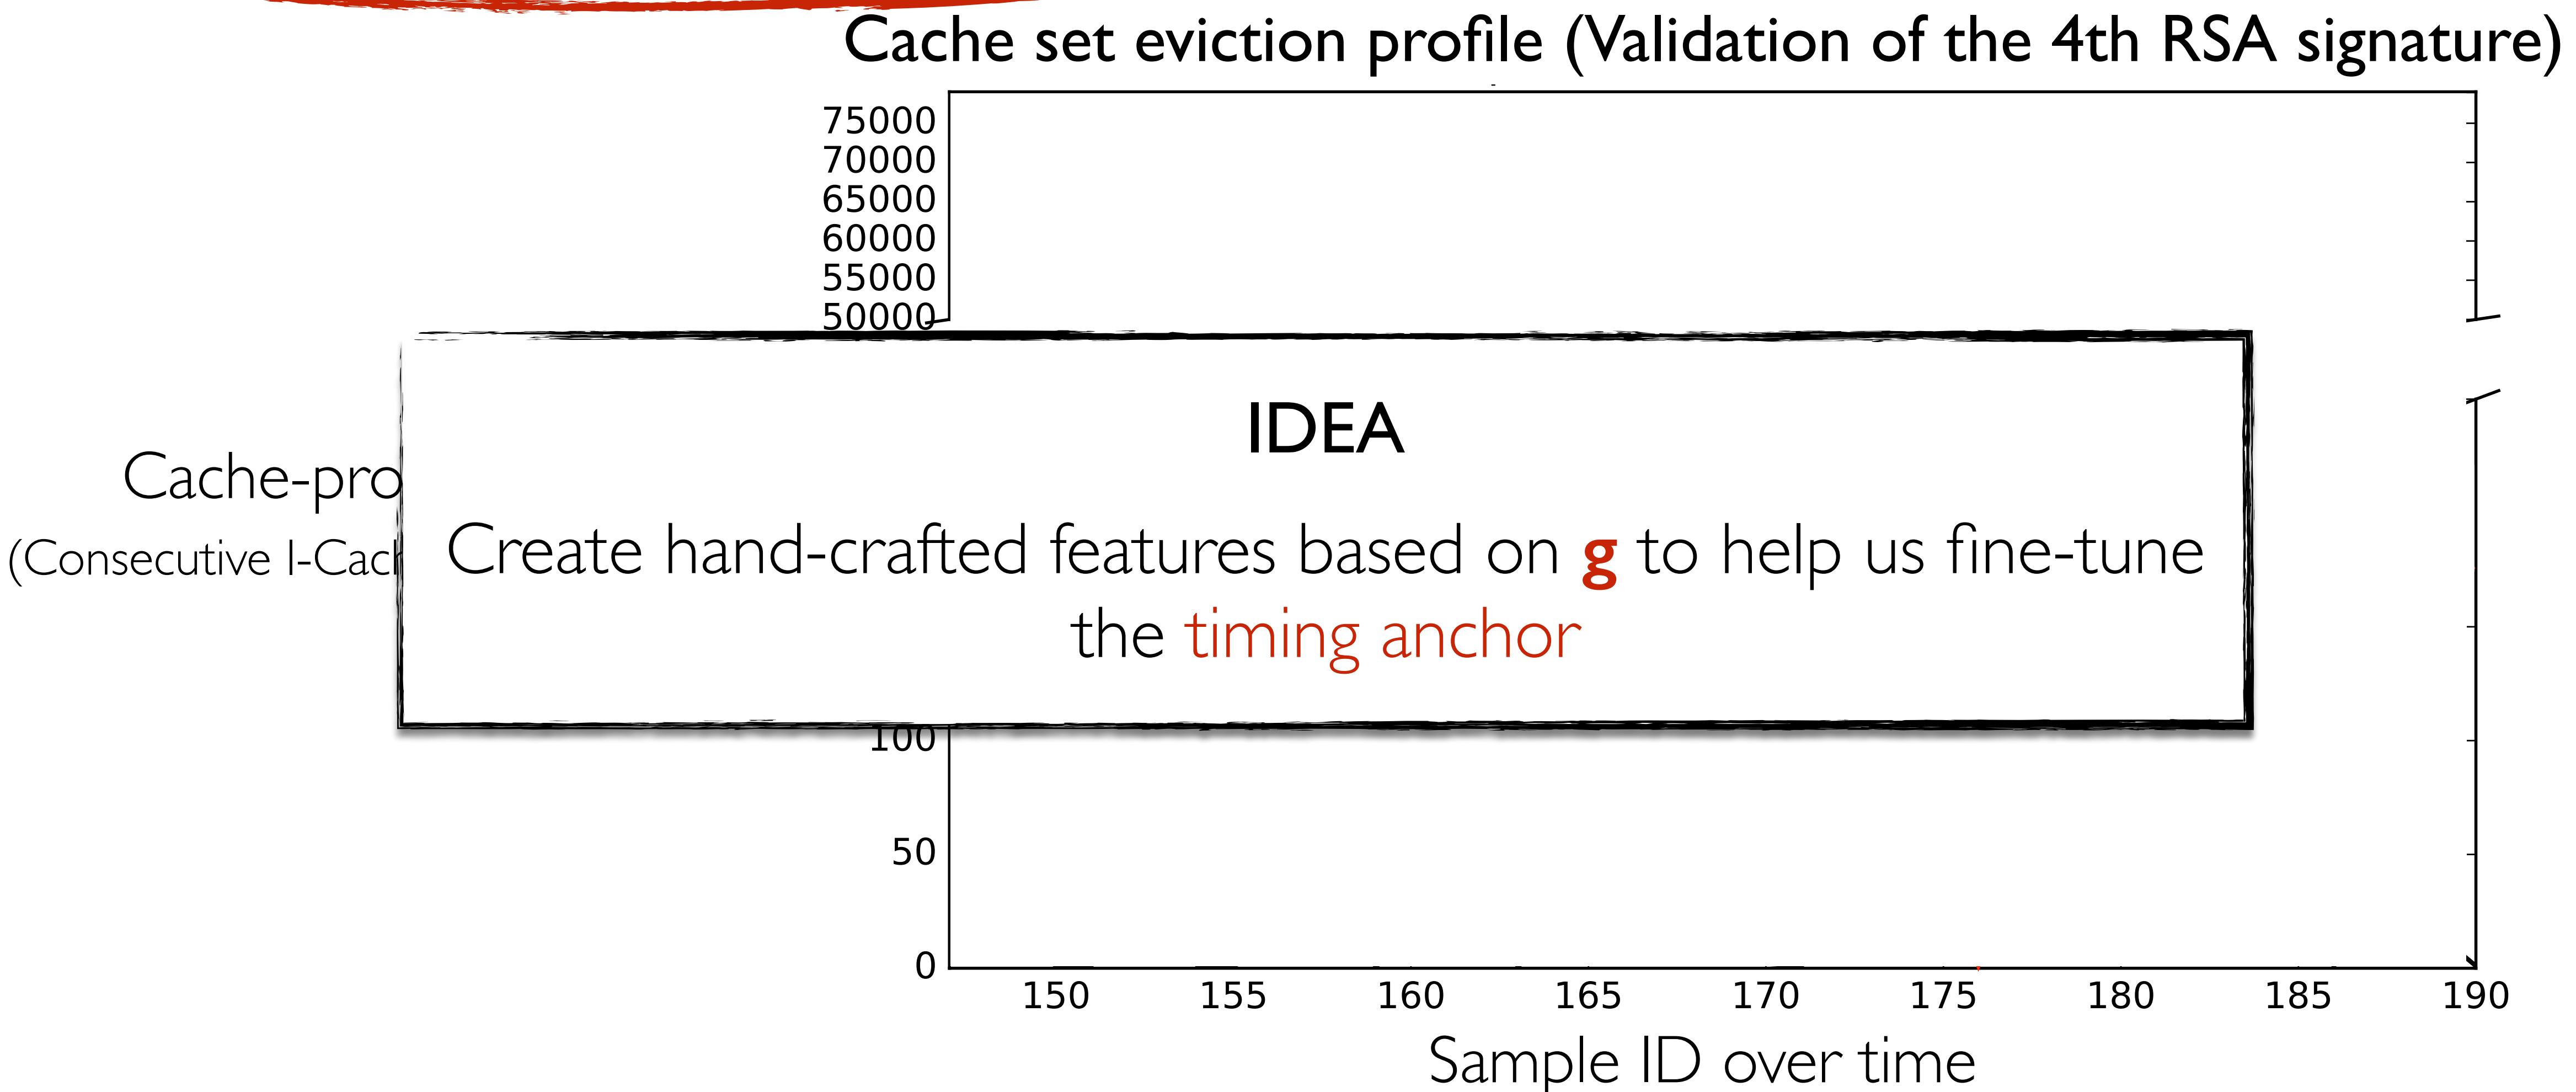
# When (during execution) to inject fault?

---

- Sketch of I-Cache profiling
  - Pick a few code areas before the target victim code to monitor
  - Monitor for I-Cache eviction for these cache sets simultaneously (We monitor 4 sets)
  - Say **E** is the event when all these cache sets are found to be evicted
  - Track the next time **E** happens
  - Use an incrementing counter (as a high-precision timer) to track the duration between consecutive **E**'s
  - Call this duration between consecutive E's — **g**
  - Time-series **g** => a fine-grained proxy of Trustcode code execution

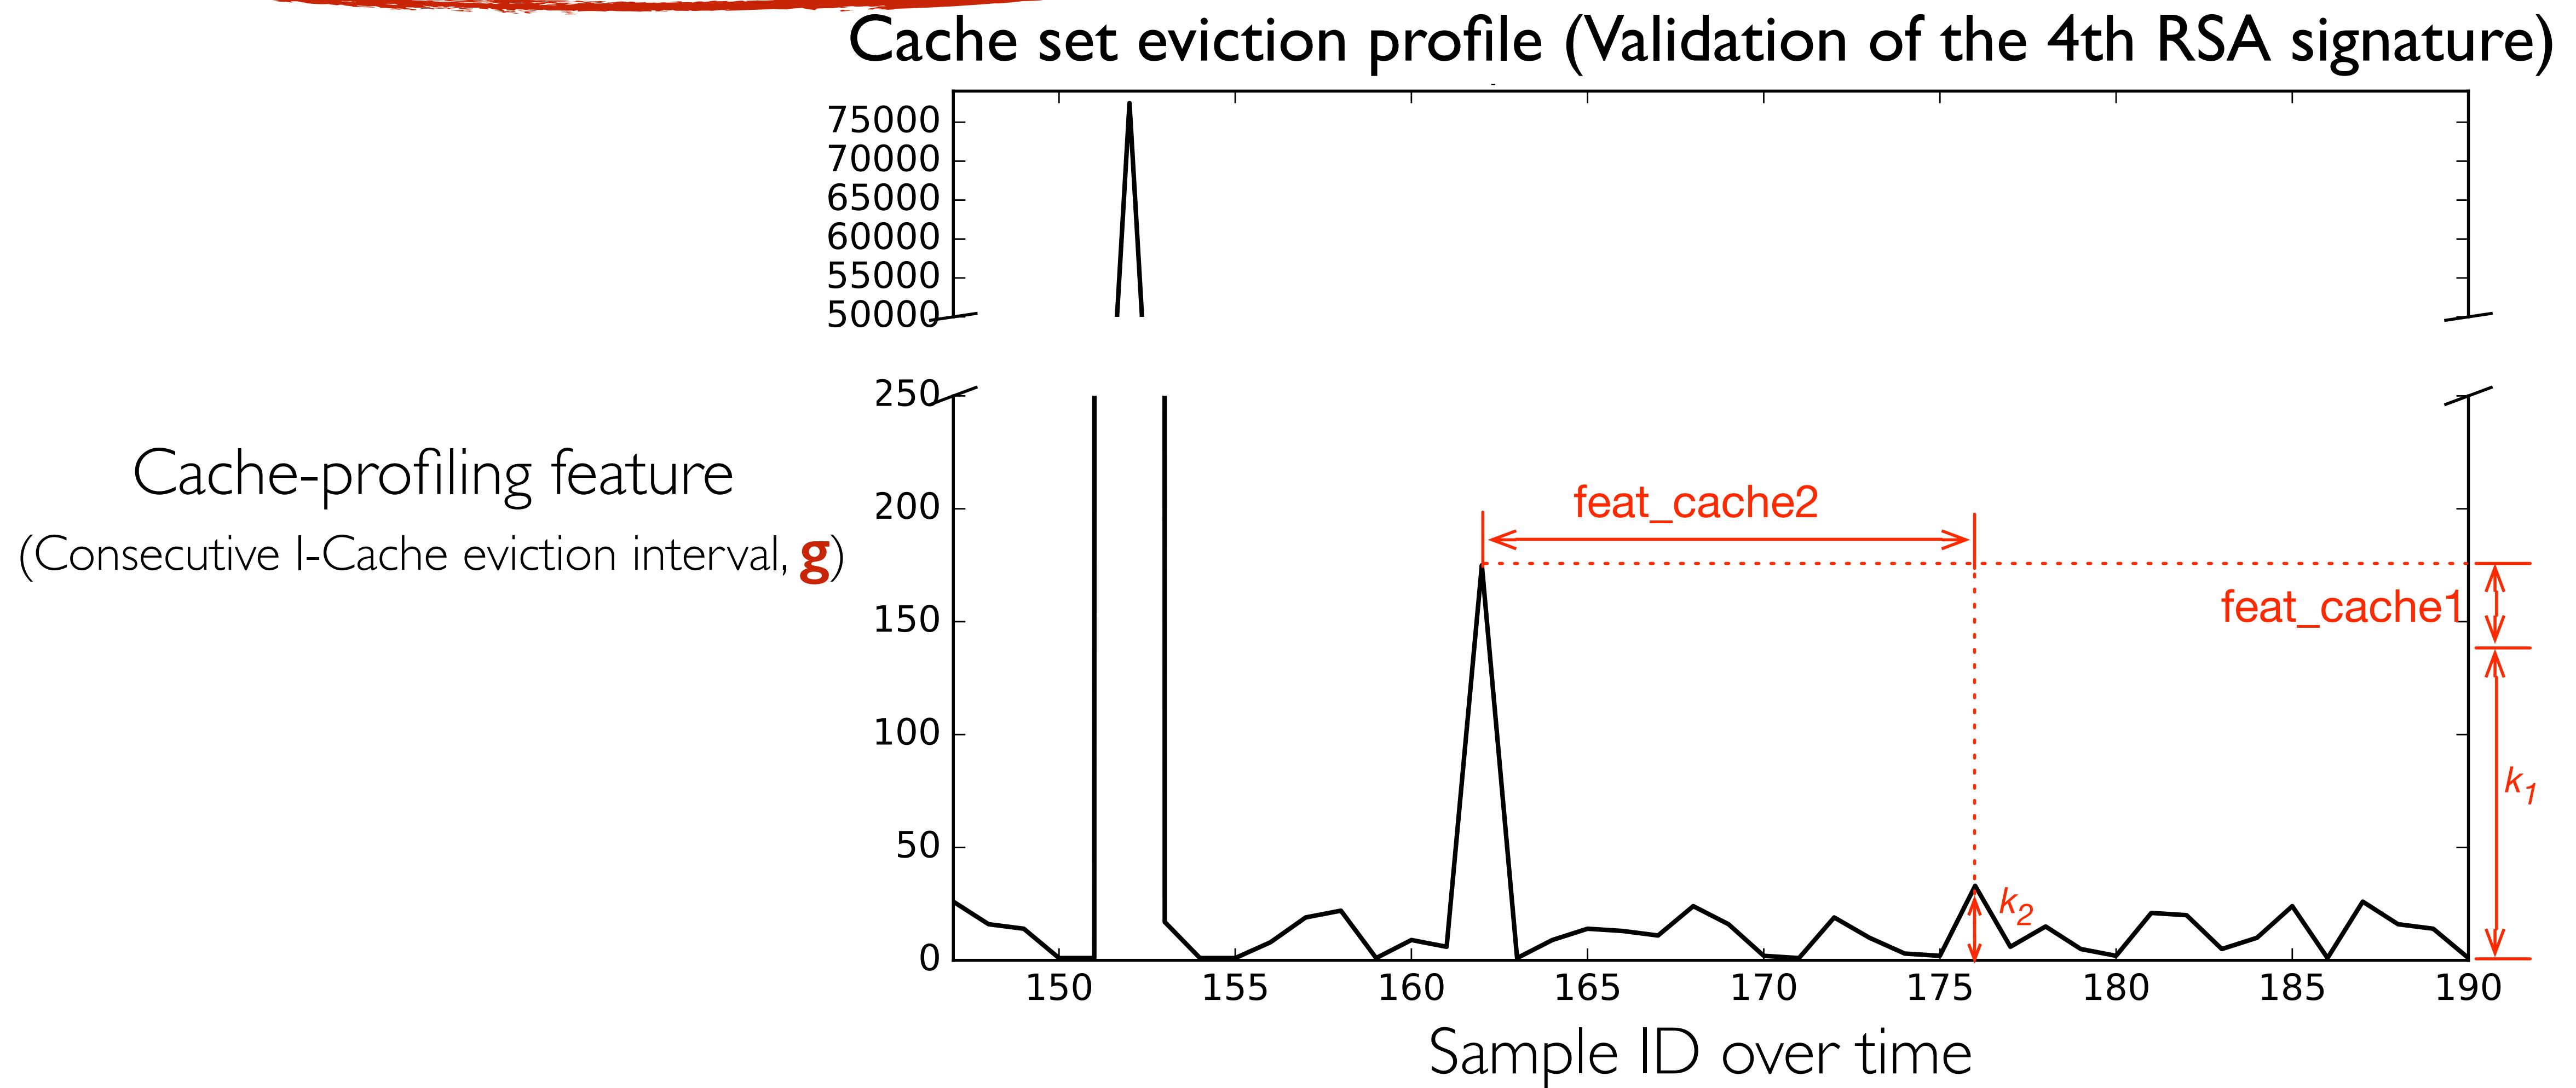
# Timing-Based Code Execution Profiling

4) CERT[1].modulus => .mdt file hashes



# Timing-Based Code Execution Profiling

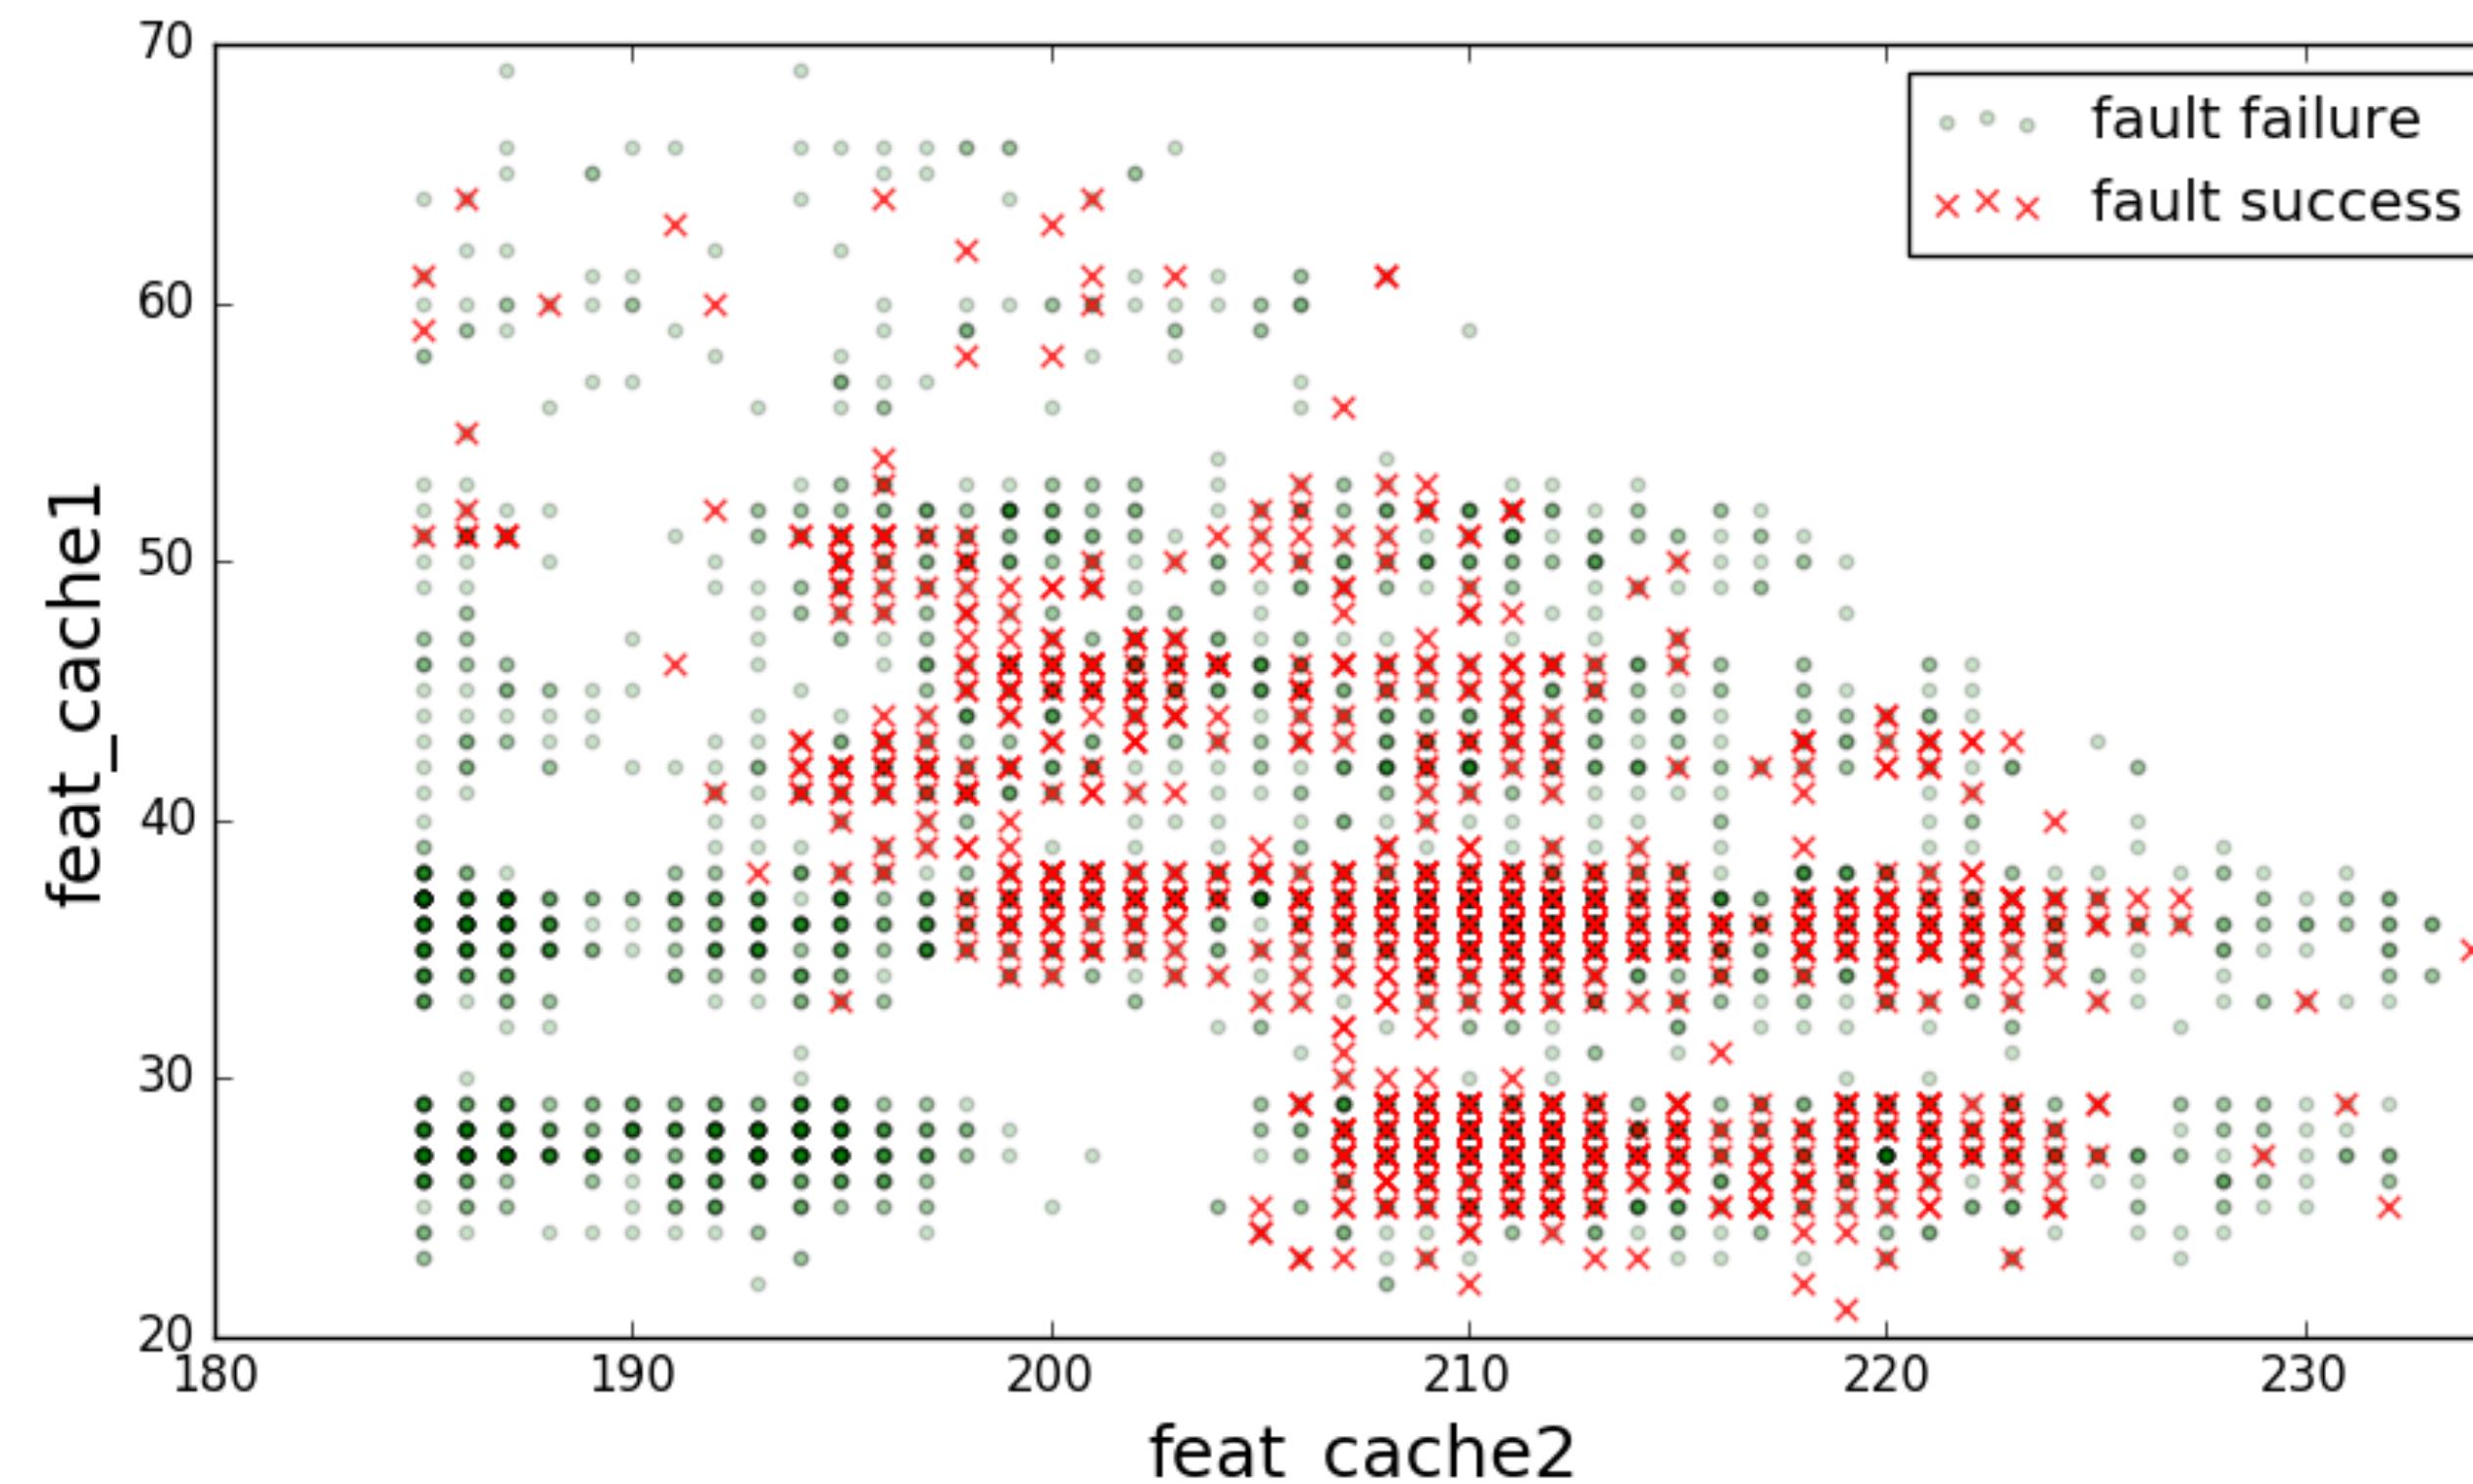
4) CERT[1].modulus => .mdt file hashes



# Timing-Based Code Execution Profiling

Track a “Fault success” as successfully corrupting targeted N modulus

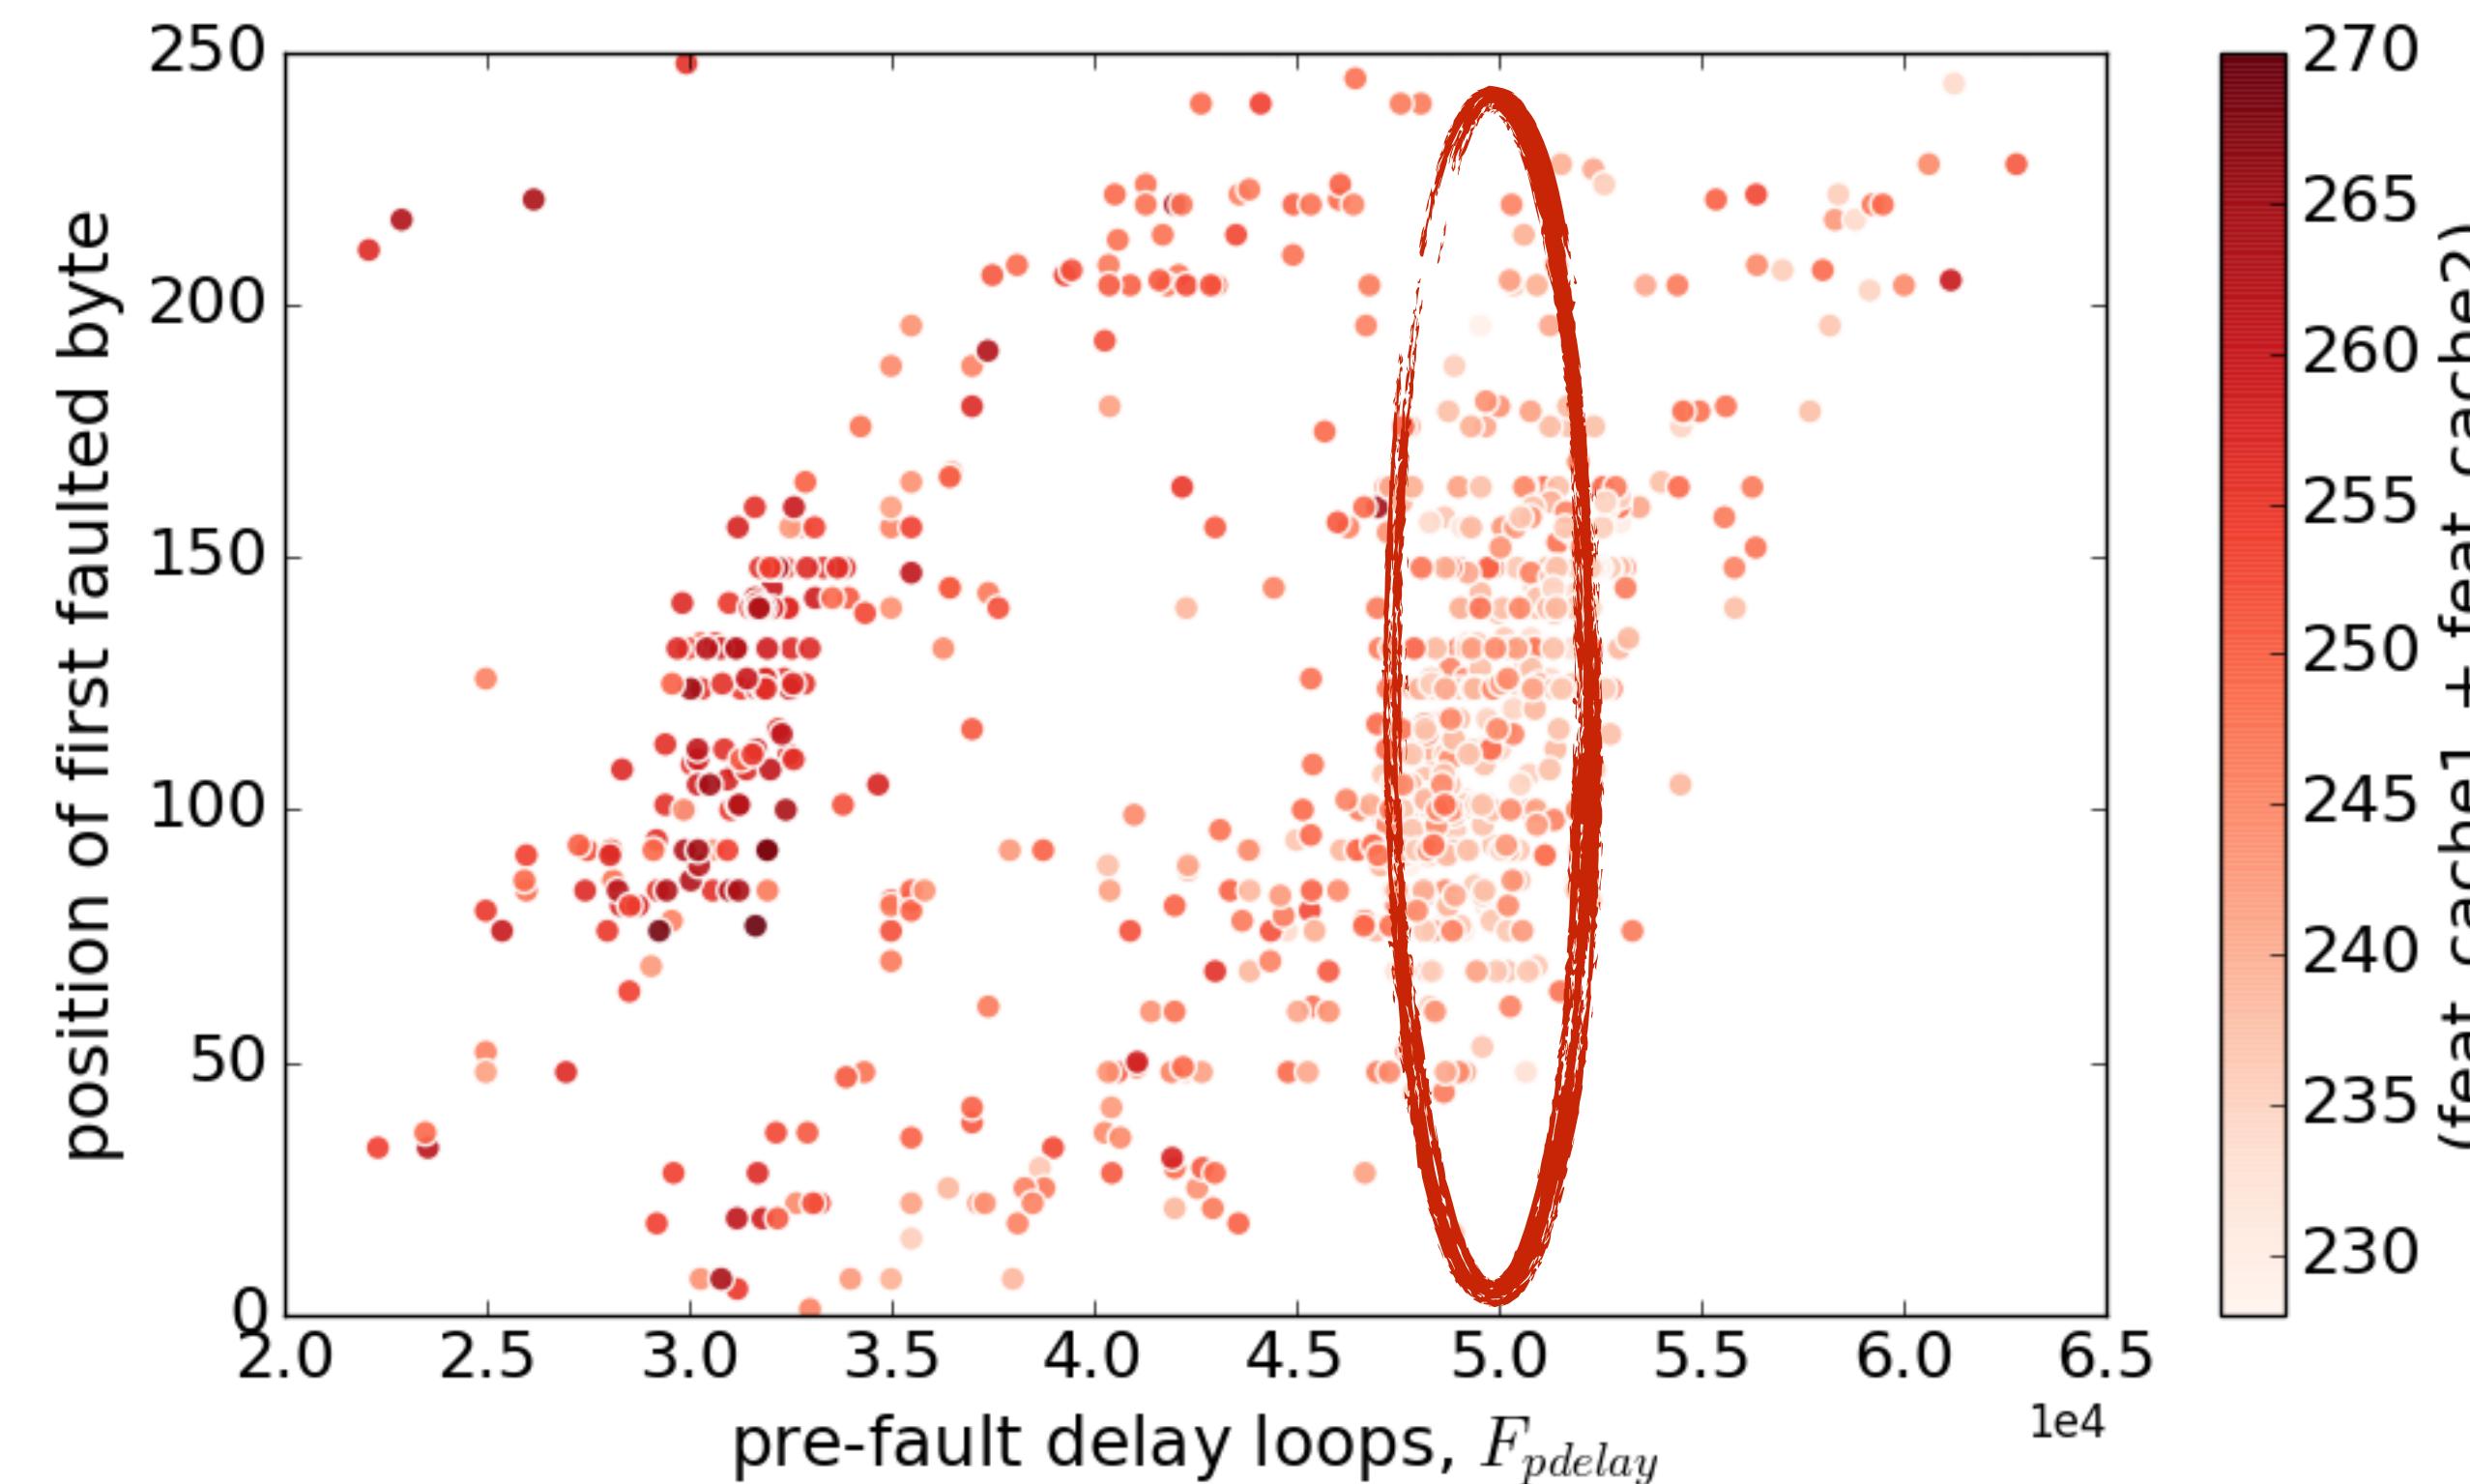
Both `feat_cache1` and `feat_cache2` can influence success rate of faults



# Timing-Based Code Execution Profiling

But these features alone are insufficient!

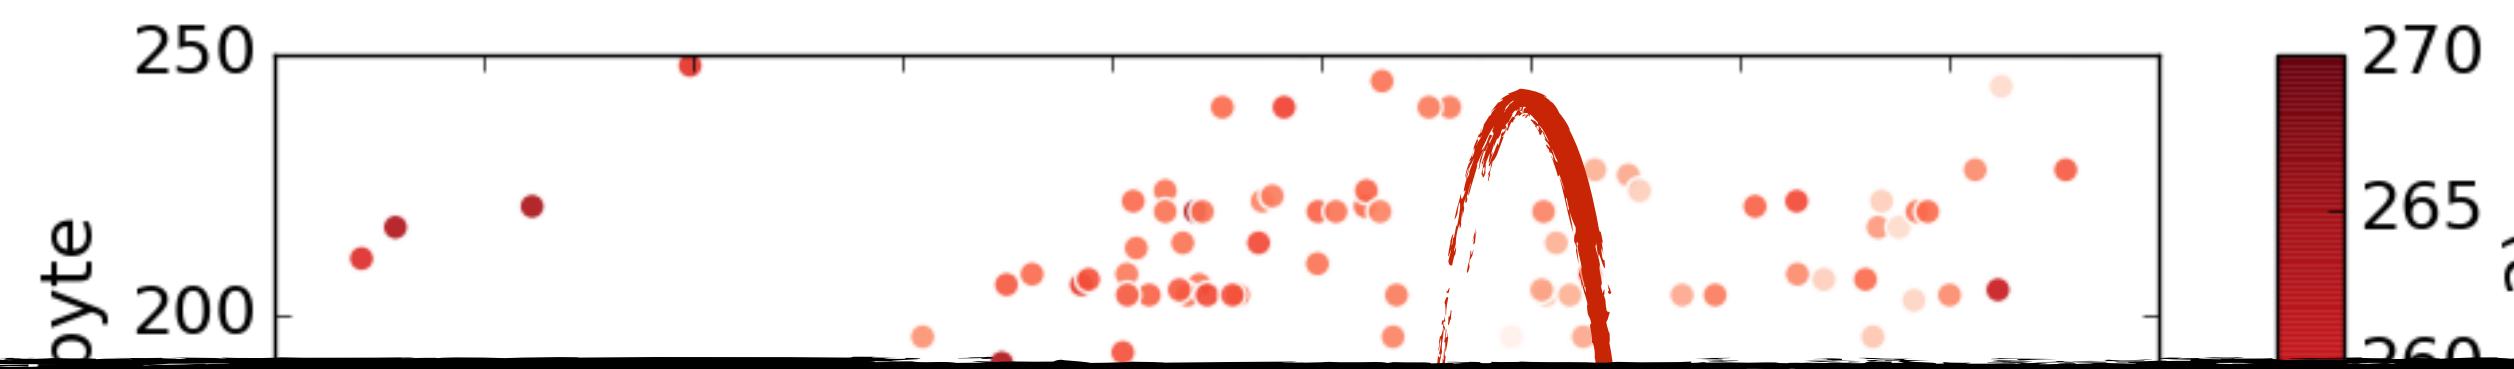
Too much variability given any value of pre-fault delay loops,  $F_{pdelay}$



# Timing-Based Code Execution Profiling

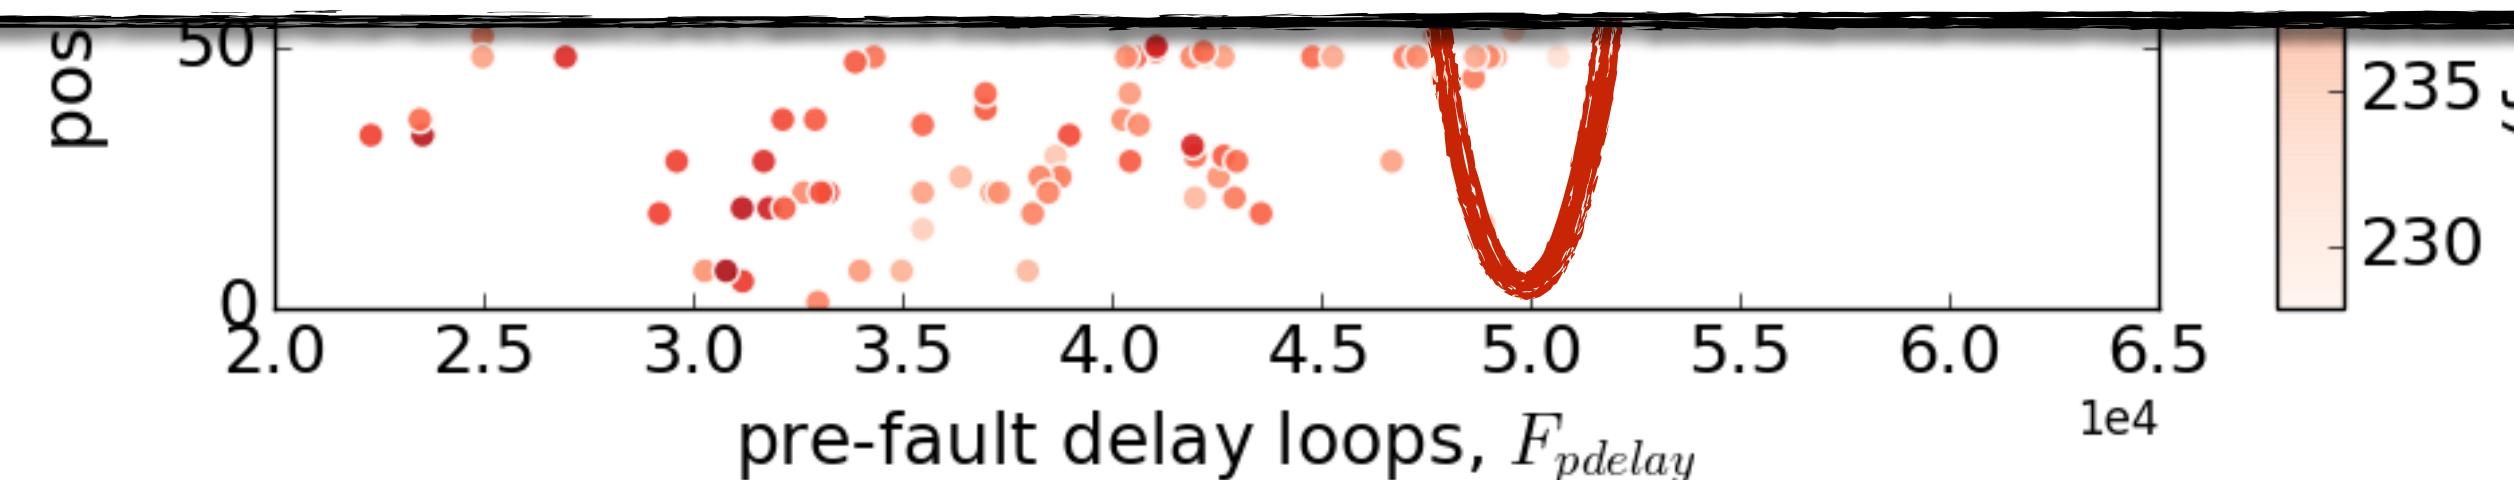
But these features alone are insufficient!

Too much variability given any value of pre-fault delay loops,  $F_{pdelay}$



IDEA

Instead of a fixed  $F_{pdelay}$ , devise an adaptive  $F_{pdelay}$  to target a specific position within  $N$



# Timing-Based Code Execution Profiling

---

Sample lots of faulting parameters and resulting faulted buffer position,  $F_{\text{pos}}$

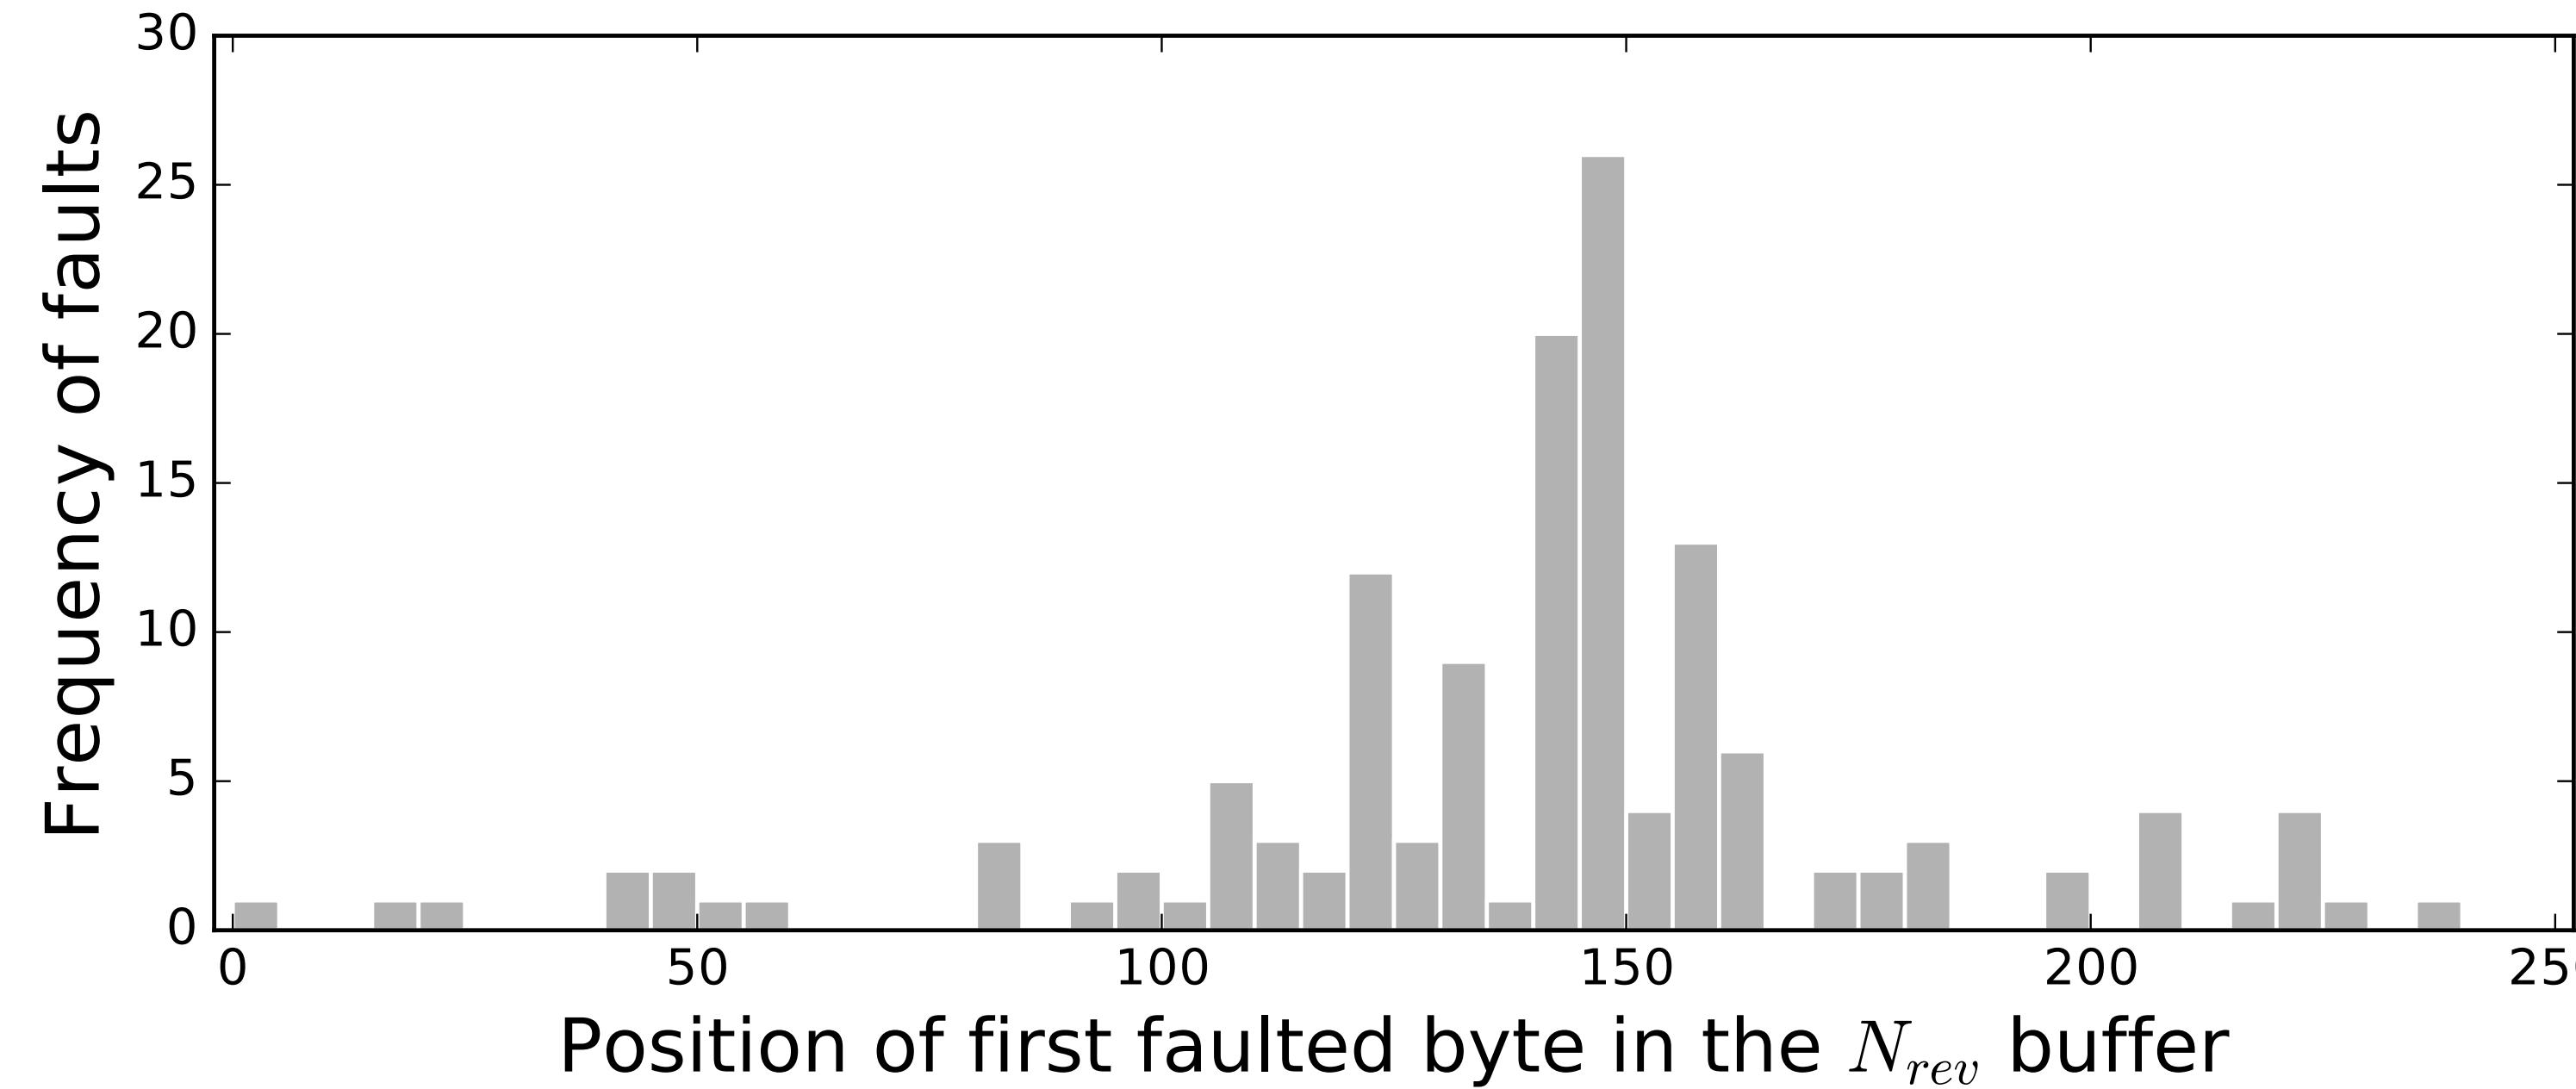
Create a **linear regression model** based on the empirical observations:

$$F_{\text{pos}} \sim \text{feat}_{\text{cache1}} + \text{feat}_{\text{cache2}} + F_{\text{pdelay}} + \text{temperature} + \text{intercept}$$

At runtime, we can then adjust our adaptive pre-fault delay loops,  $F_{\text{pdelay}}$

$$F_{\text{pdelay}} = \mathbf{f} (F_{\text{pos}}, \text{feat}_{\text{cache1}}, \text{feat}_{\text{cache2}}, \text{temperature}, \text{intercept})$$

# Putting it together



Statistics:

- ~20% of faulting attempts (1153 out of 6000) result in a successful desired fault in the  $N_{rev}$  buffer we want
- These faults consist of 805 unique values, of which 38 (4.72%) are factorizable
- One instance of the desired fault in ~65 faulting attempts

# Summary of Attack Enablers

---

- I. No hardware safeguard limits in regulators
- II. Large range of possible combinations of freq/volt for fault injection
- III. Cores deployed in different freq/volt domains
- IV. Hardware regulators operate across security boundaries
- V. Execution timing of Trustzone code can be profiled with hardware cycle counter from outside Trustzone
- VII. Trustzone code execution can be profiled using side-channel-based attacks, like Prime+Probe cache attacks

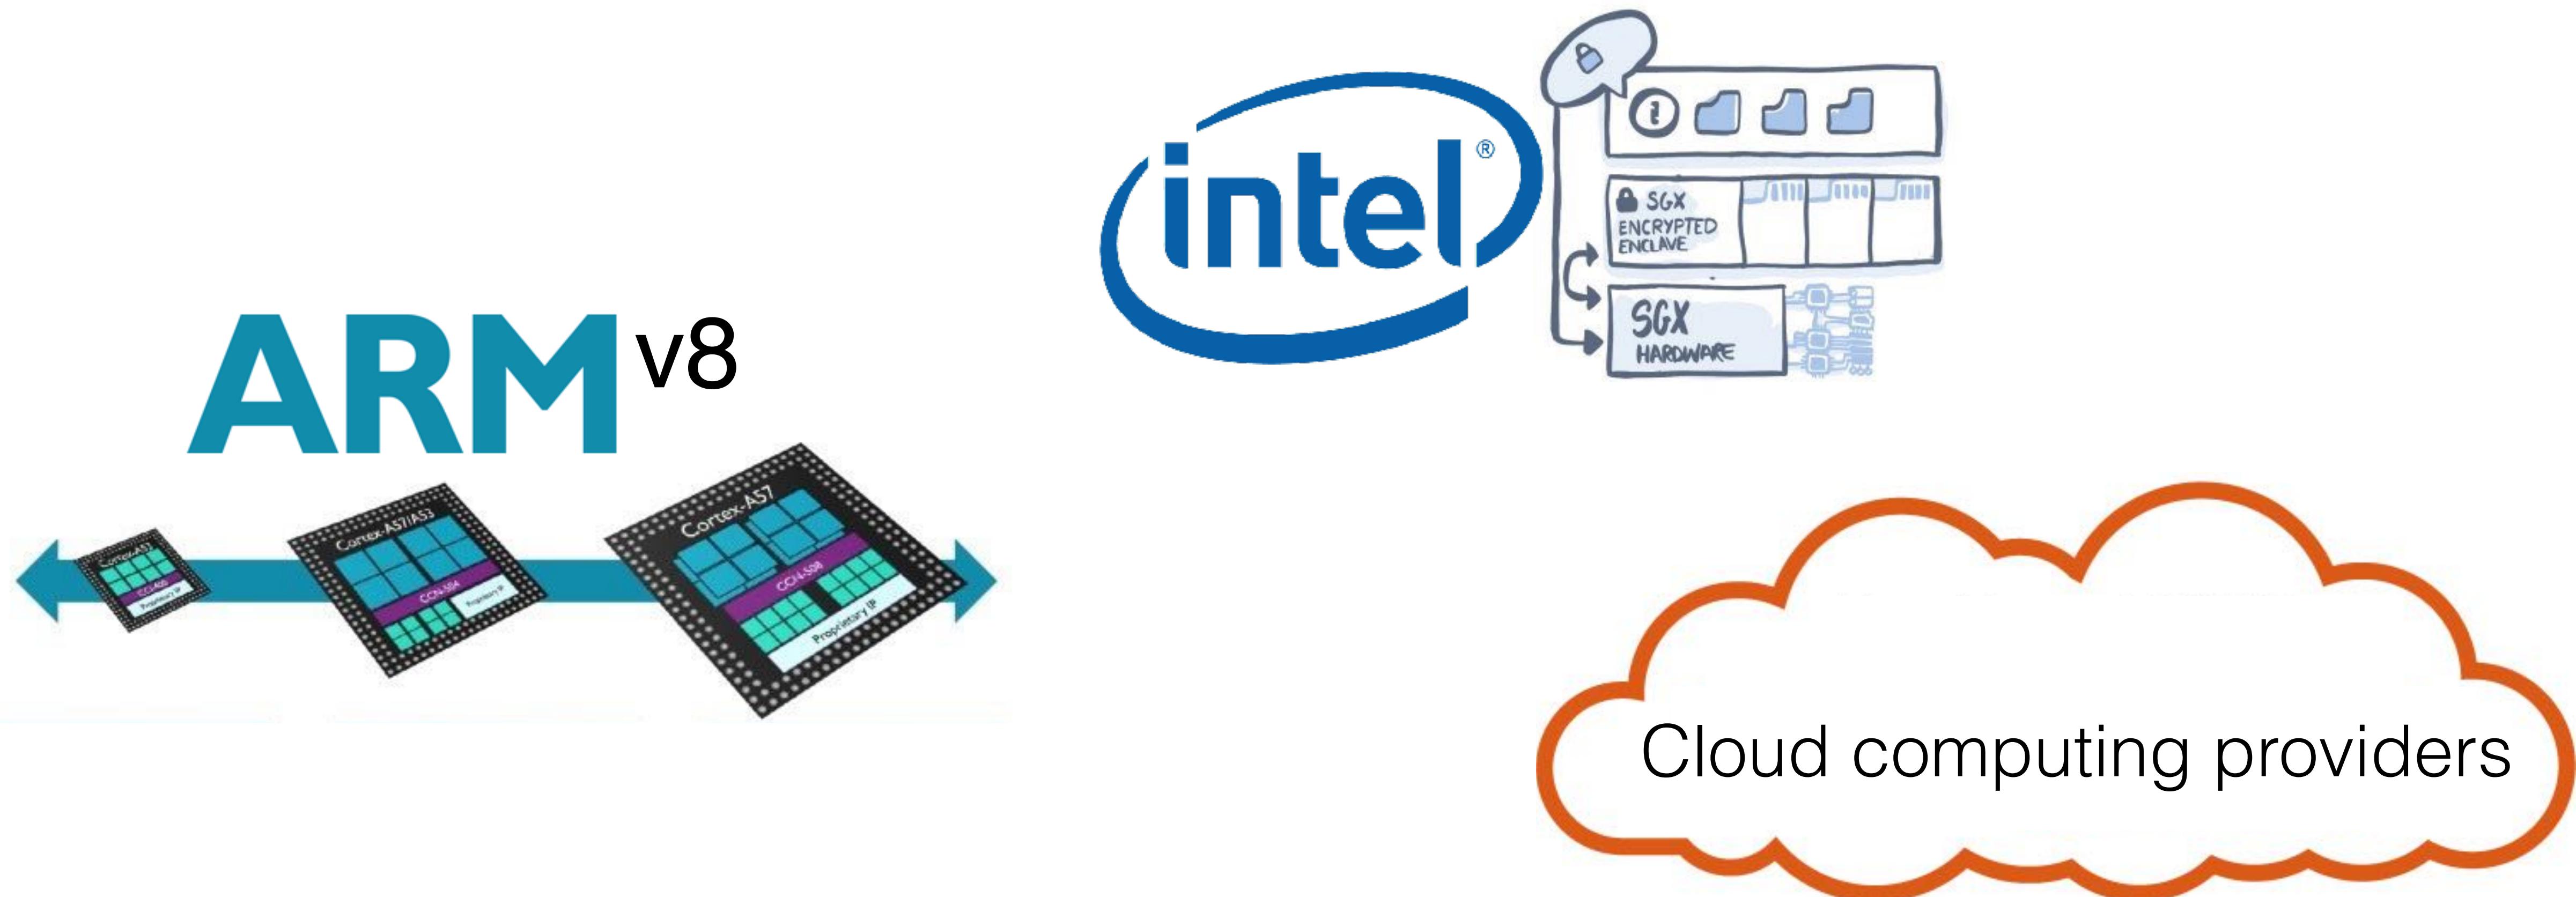
# Outline

---

- I. DVFS and Deep Dive into Hardware Regulators
- II. The CLKSCREW Attack
- III. Trustzone Attack I: Secret AES Key Inference
- IV. Trustzone Attack 2:Tricking RSA Signature Validation
- V. **Concluding Remarks**

# Attack Applicability to Other Platforms

Energy management mechanisms in the industry is trending towards  
finer-grained and increasingly heterogeneous designs



# Possible Defenses

---

## Hardware-Level

Operating limits in hardware

Separate cross-boundary regulators

Microarchitectural Redundancy

## Software-Level

Randomization

Code execution redundancy

# Sound Bytes



New attack surface via energy management software interfaces

Not a hardware or software bug  
Fundamental design flaws in energy management mechanisms

Future energy management designs must take security into consideration



# CLKSCREW

## Exposing the Perils of Security-Oblivious Energy Management

Adrian Tang

Simha Sethumadhavan, Salvatore Stolfo



@0x0atang



<https://github.com/0x0atang/clkscrew>



<https://0x0atang.github.io>