



splunk>

# Defeating the Curse of “First Time” Events

Ignacio Bermudez Corrales | Security Data Scientist @SplunkUBA

[icorrales@splunk.com](mailto:icorrales@splunk.com)

May 2018 | Version 1.0



# Forward-Looking Statements

During the course of this presentation, we may make forward-looking statements regarding future events or the expected performance of the company. We caution you that such statements reflect our current expectations and estimates based on factors currently known to us and that actual events or results could differ materially. For important factors that may cause actual results to differ from those contained in our forward-looking statements, please review our filings with the SEC.

The forward-looking statements made in this presentation are being made as of the time and date of its live presentation. If reviewed after its live presentation, this presentation may not contain current or accurate information. We do not assume any obligation to update any forward-looking statements we may make. In addition, any information about our roadmap outlines our general product direction and is subject to change at any time without notice. It is for informational purposes only and shall not be incorporated into any contract or other commitment. Splunk undertakes no obligation either to develop the features or functionality described or to include any such feature or functionality in a future release.

Splunk, Splunk>, Listen to Your Data, The Engine for Machine Data, Splunk Cloud, Splunk Light and SPL are trademarks and registered trademarks of Splunk Inc. in the United States and other countries. All other brand names, product names, or trademarks belong to their respective owners. © 2018 Splunk Inc. All rights reserved.

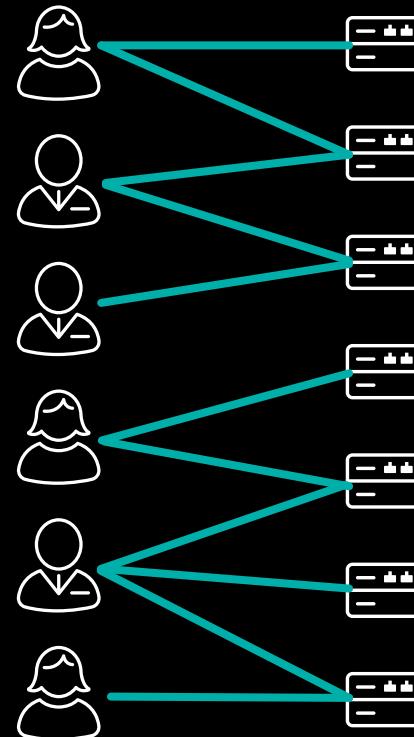
# What are “First Time Events”

And why we should care about them

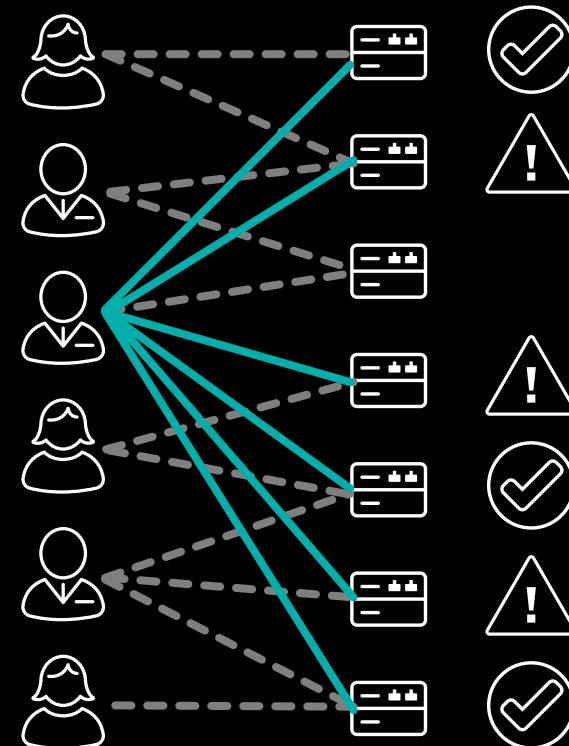
# Machine Accesses

# Do all first time accesses look suspicious?

## Historical Accesses



## First Time Accesses



# Normal Operations?

# Suspicious Accesses?

# Other Examples of First Time Events in Real Life

# Some are normal, some can be risky

## ► Cloud Storage File-sharing

- Employee is planning to leave the company and downloads some unprotected sensitive files
  - A newly made presentation is shared for review among some group of employees

## ► Private Code Repositories

- An account is compromised and cloned projects with intellectual property
  - Cloned project made by a colleague that I'm planning to collaborate on

## ► Badge Accesses

- A worker enters the office servers room during off work hours
  - Same worker enters some other campus' meeting room during a business trip

# A Needle in a Haystack

# The risk of first time events

- ▶ Normal operations trigger a “sea” of first time events
  - ▶ Collecting evidence for suspicious activity may take too long
  - ▶ Suspicious activity can happen only once



# Solution?

# First time event handling

# Conservative

# Relaxed

# Rules

# Machine Learning

All first time events  
are interesting and  
worth investigating

# First time events are not interesting

# Low coverage, high precision

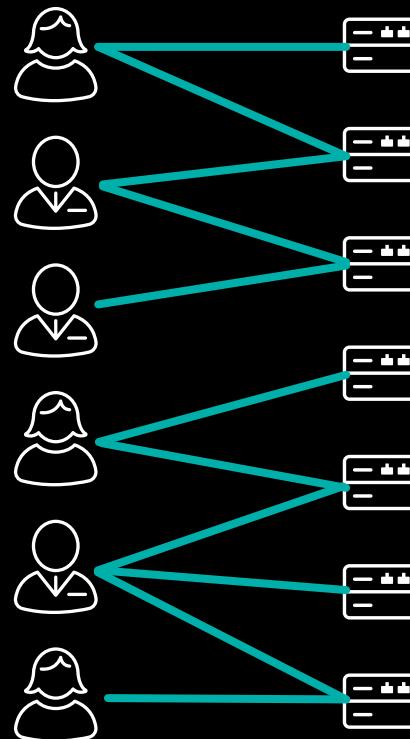
Finds the subtle,  
but learns  
exclusively from  
evidence

# Let's Play with Numbers

## Foundations

# Representation of the Problem

**So a machine learning algorithm can learn**



# Analogous Representations

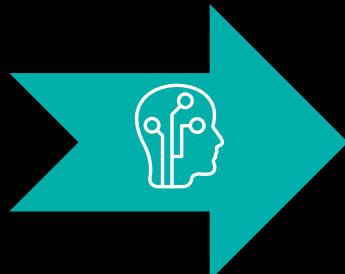
	1	1					
		1	1				
			1				
				1	1		
					1	1	1
							1

# Machine Learning Generalizes

## Filling the gaps

## Historical Data - Evidence

	1	1					
		1	1				
			1				
				1	1		
					1	1	1
							1



# ML generalization capabilities allow Predictions of First Time Events

## ML Reconstruction of Data

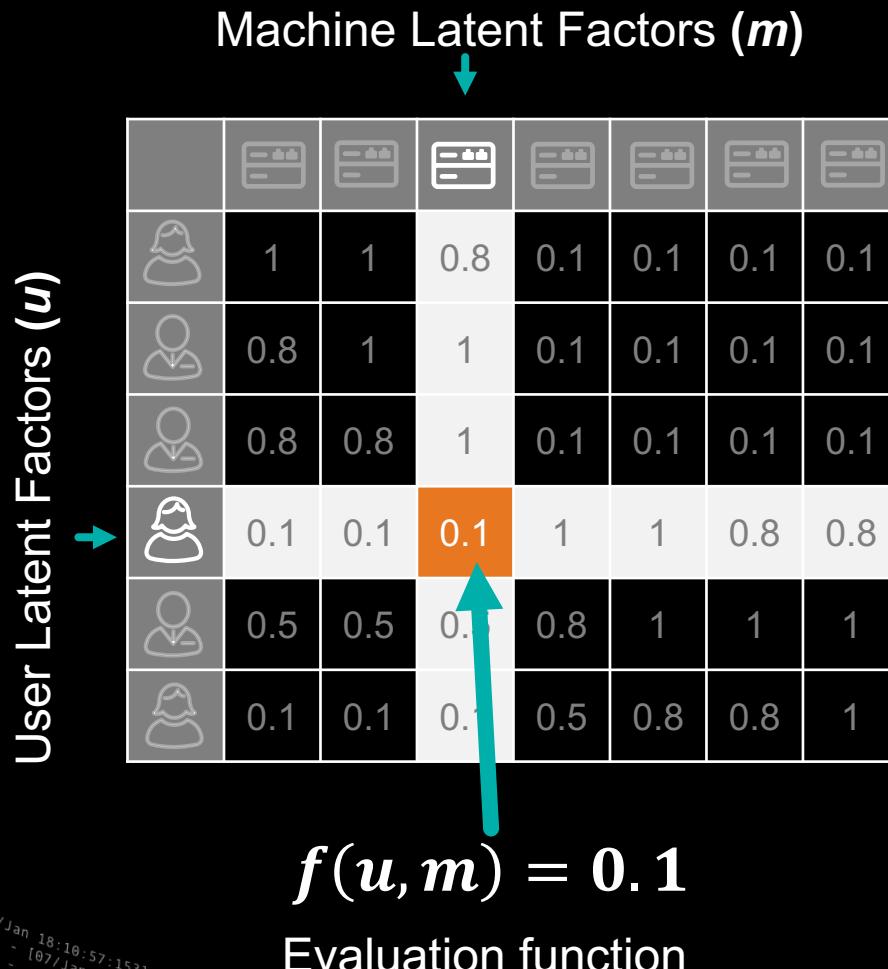


# Suspicious

## Benign

# How Data is Reconstructed

## Latent factor models



Latent factors ( $u, m$ ) are real valued vectors which encode behavior and properties, that evaluated with a function allow reconstruction of the evidence.

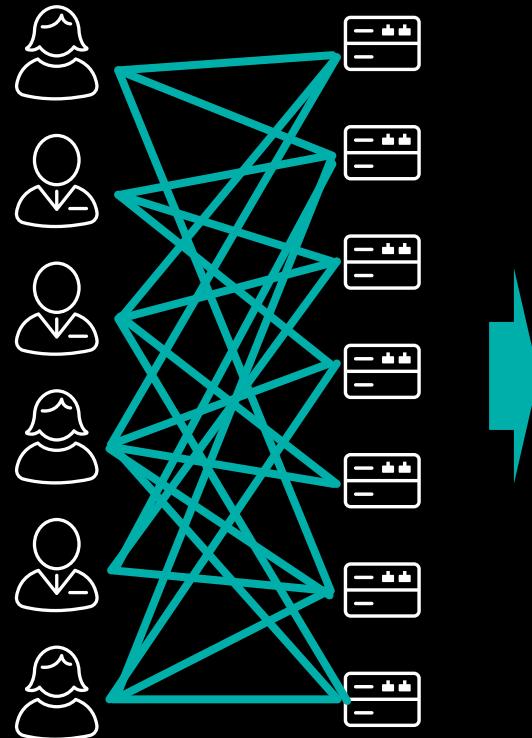
## Netflix's recommendation system

Latent factors may be indicators of  
Of age, income, gender of watchers;  
And for movies the genre, duration, release date

# Taming ML with Rules

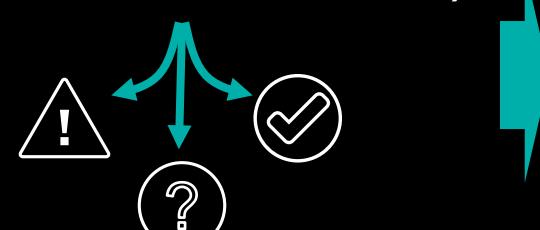
## Passing down knowledge to machine learning algorithm

# Synthetic Generated Events



## Rules

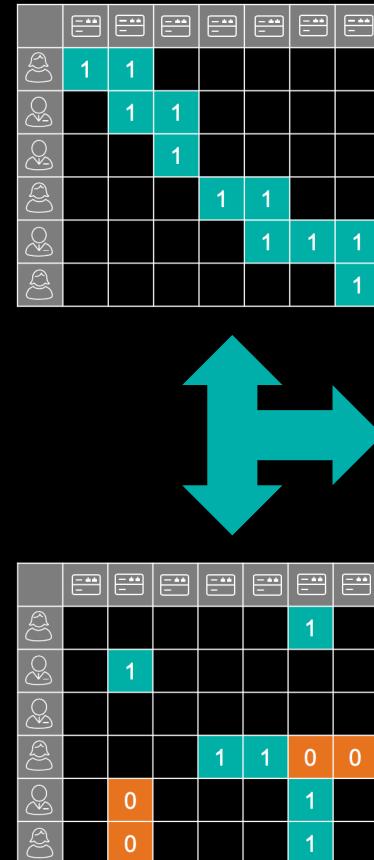
*if(event == condition)*



# Coverage of Rules

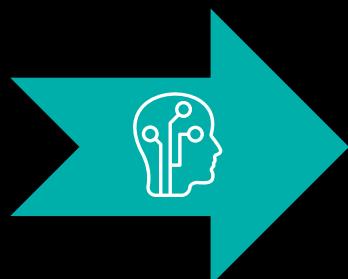

# Connecting Dots

## Tamed machine learning embeds knowledge from rules and evidence



# Historical Data + Rules

	1	1					1
		1	1				
			1				
				1	1	0	0
		0			1	1	1
		0				1	1



# ML Reconstruction of Data

	1	1	0.8	0.1	0.5	0.8	0.5
	0.8	1	1	0.1	0.5	0.5	0.5
	0.5	0.5	1	0.1	0.1	0.1	0.1
	0.1	0.1	0.5	1	1	0.1	0.1
	0.1	0.1	0.5	0.8	1	1	1
	0.1	0.1	0.5	0.5	0.8	0.8	1

# Connecting Dots

**Richer predictive model can describe more complexities**

# Before

	1	1	0.8	0.1	0.1	0.1	0.1
	0.8	1	1	0.1	0.1	0.1	0.1
	0.8	0.8	1	0.1	0.1	0.1	0.1
	0.1	0.1	0.1	1	1	0.8	0.8
	0.5	0.5	0.5	0.8	1	1	1
	0.1	0.1	0.1	0.5	0.8	0.8	1

# Just evidence

# After

	1	1	0.8	0.1	0.5	0.8	0.5
	0.8	1	1	0.1	0.5	0.5	0.5
	0.5	0.5	1	0.1	0.1	0.1	0.1
	0.1	0.1	0.5	1	1	0.1	0.1
	0.1	0.1	0.5	0.8	1	1	1
	0.1	0.1	0.5	0.5	0.8	0.8	1
	0.1	0.1	0.5	0.5	0.8	0.8	1

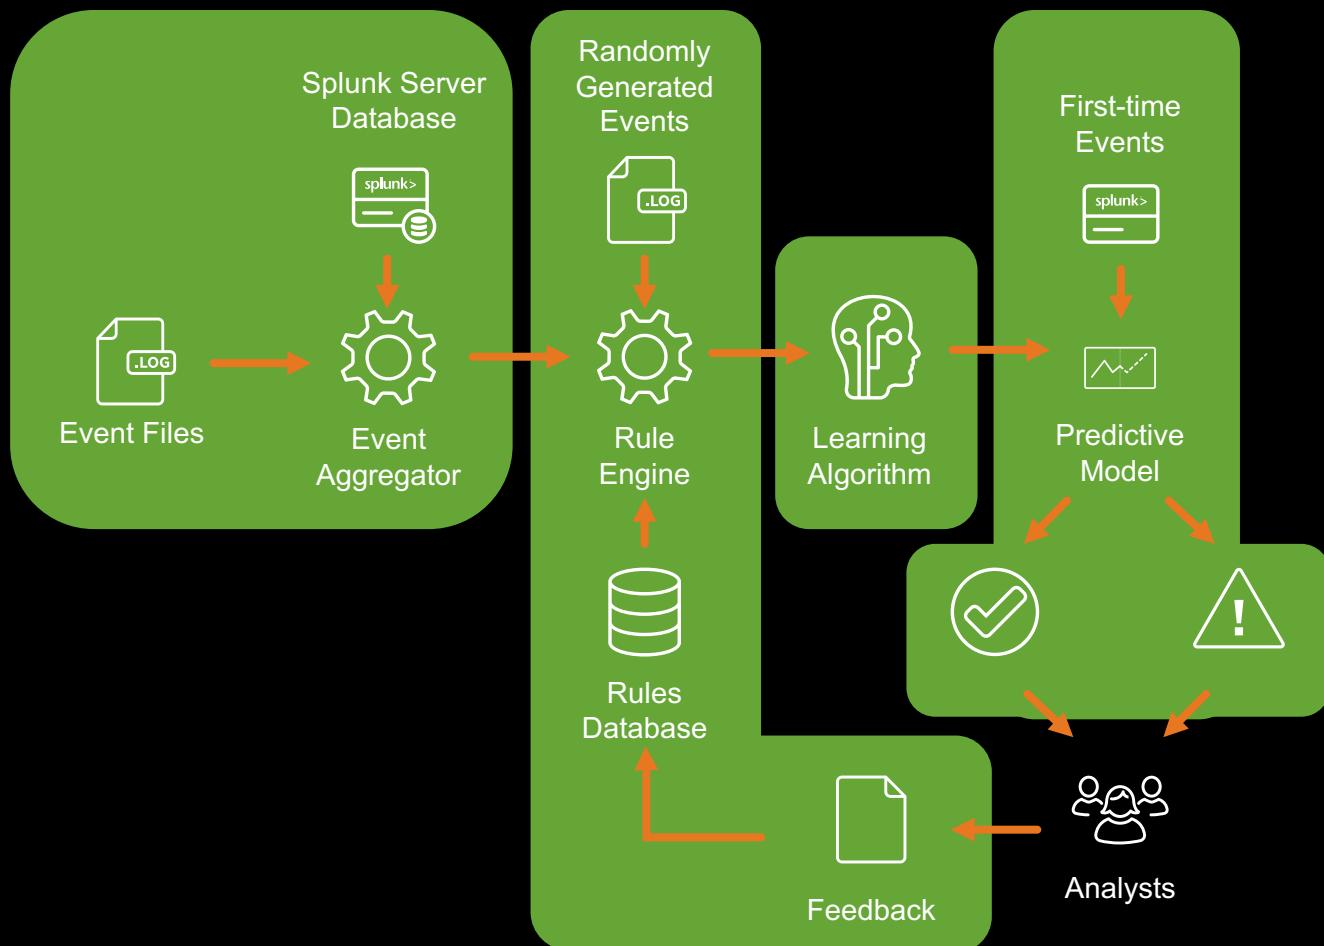
# Rules + Evidence

# Implementation

## Building suspicious first event detector

# Architecture

## 4 steps



# Example Case

## Two employees access two devices

### User Properties

Peergroup: 2  
OU: Finance



Protocol: IPP  
Duration: 2Min



### Device Properties

Device Type: Printer  
Domain: iot



Peergroup: 16  
OU: engineering



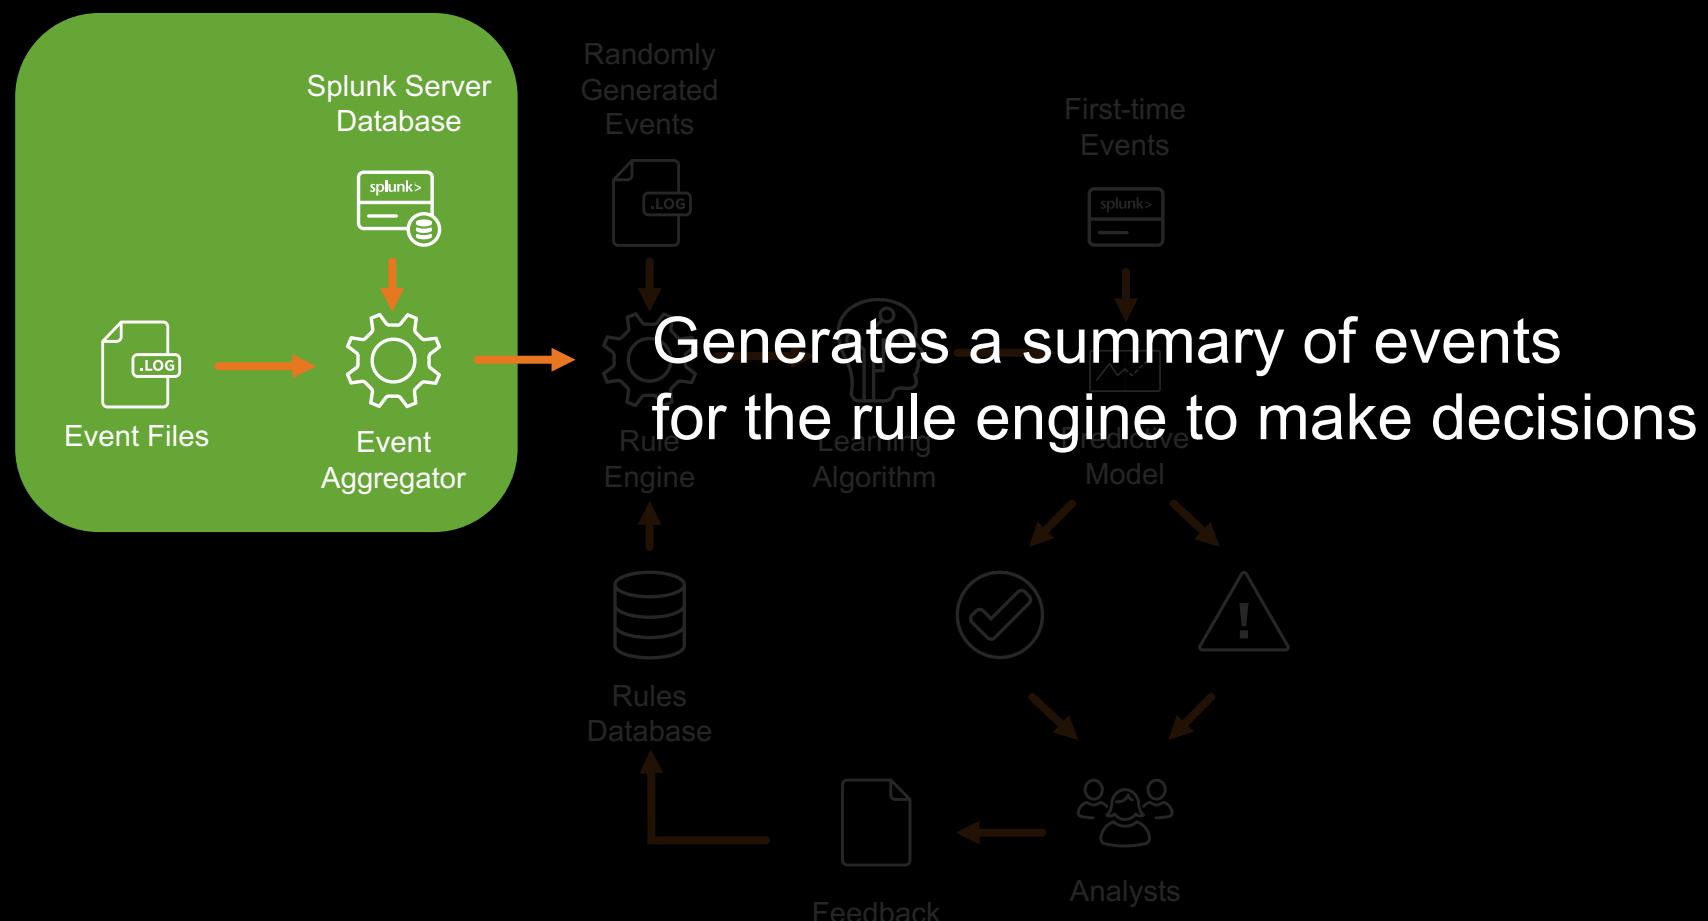
Protocol: SSH  
Duration: 2Hours



Device Type: VM  
Domain: dev

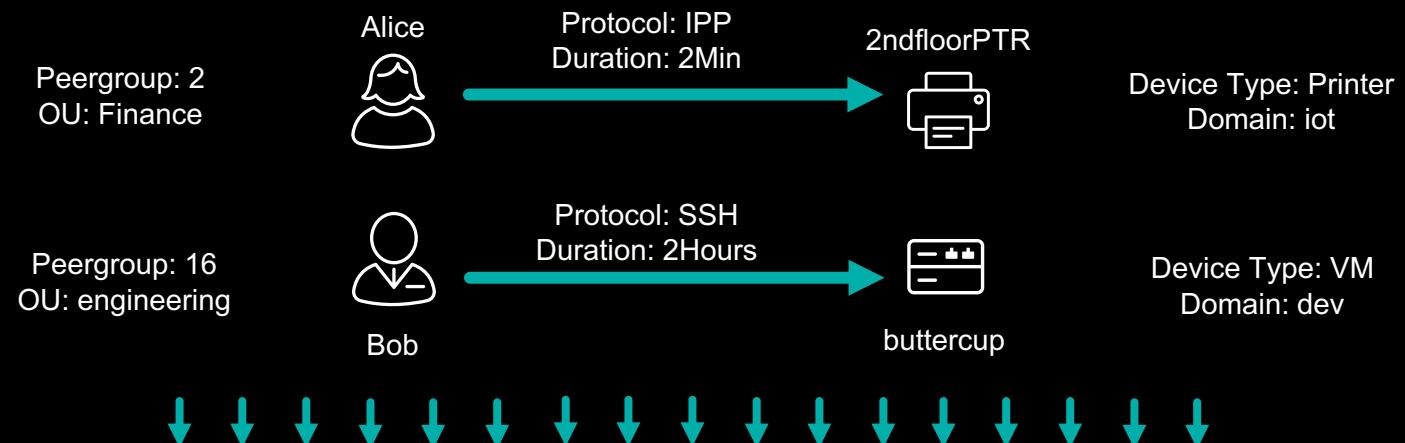
# Event Processing

## **Step 1: Build a summary of information for the rule engine**



# Event Processing

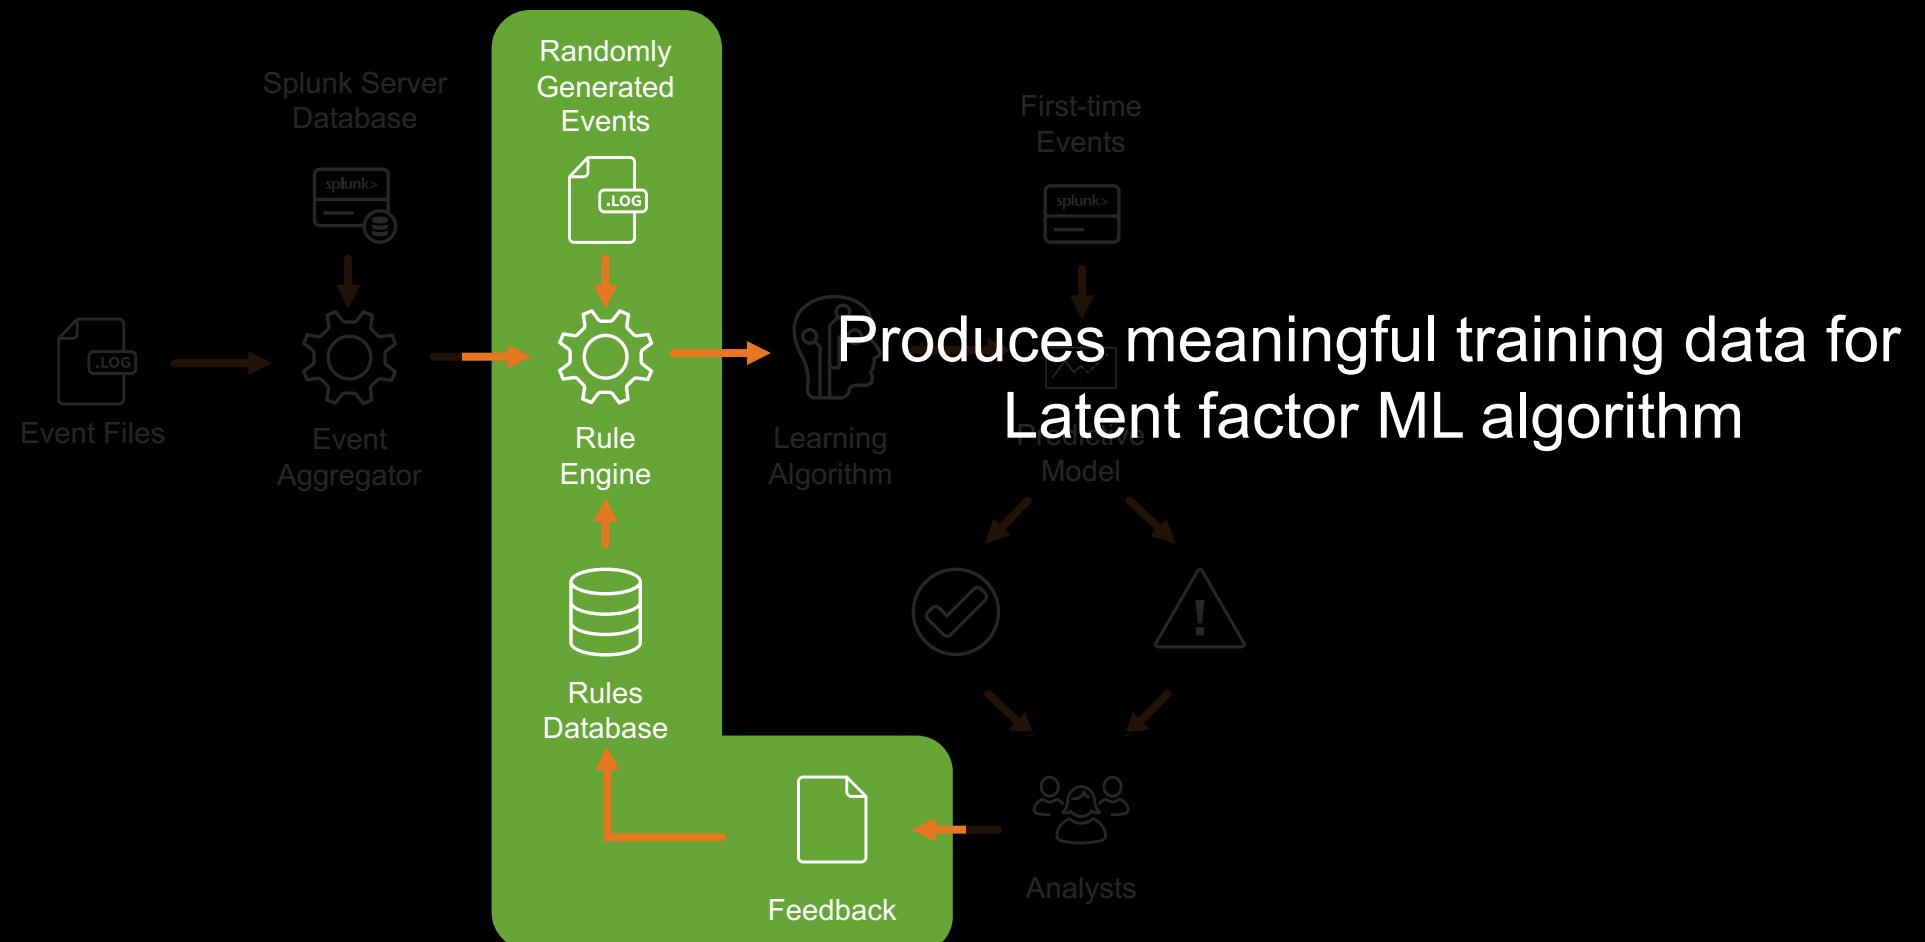
## Step 1: Converting single event into an annotated graph with properties



Link ID	User Properties	User Name	Events Summary	Device Name	Device Properties
1	Peer group: 1 OU: Engineering	Bob	SSH, 2h	Buttercup	Type: VM Domain: dev
2	Peer group: 10 OU: Finance	Alice	IPP, 2m	2ndfloorPRT	Type: Printer Domain: iot

# Rule Engine

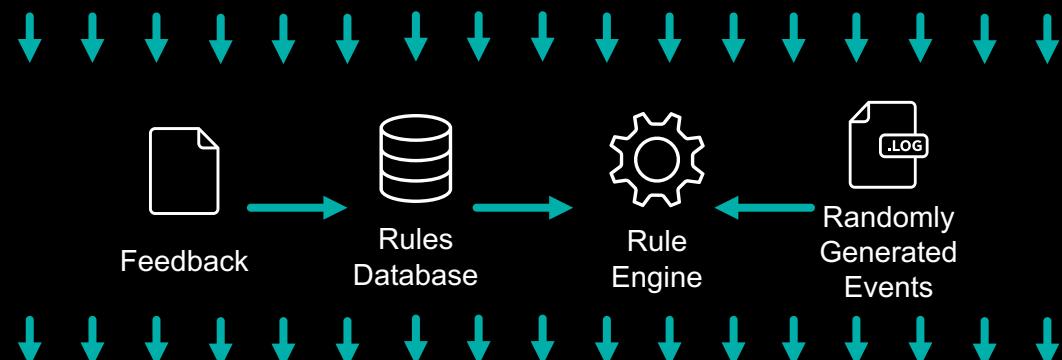
## Step 2: Convert event summaries into numerical signals



# Rule Engine

## Step 2: Converting links into numerical signals

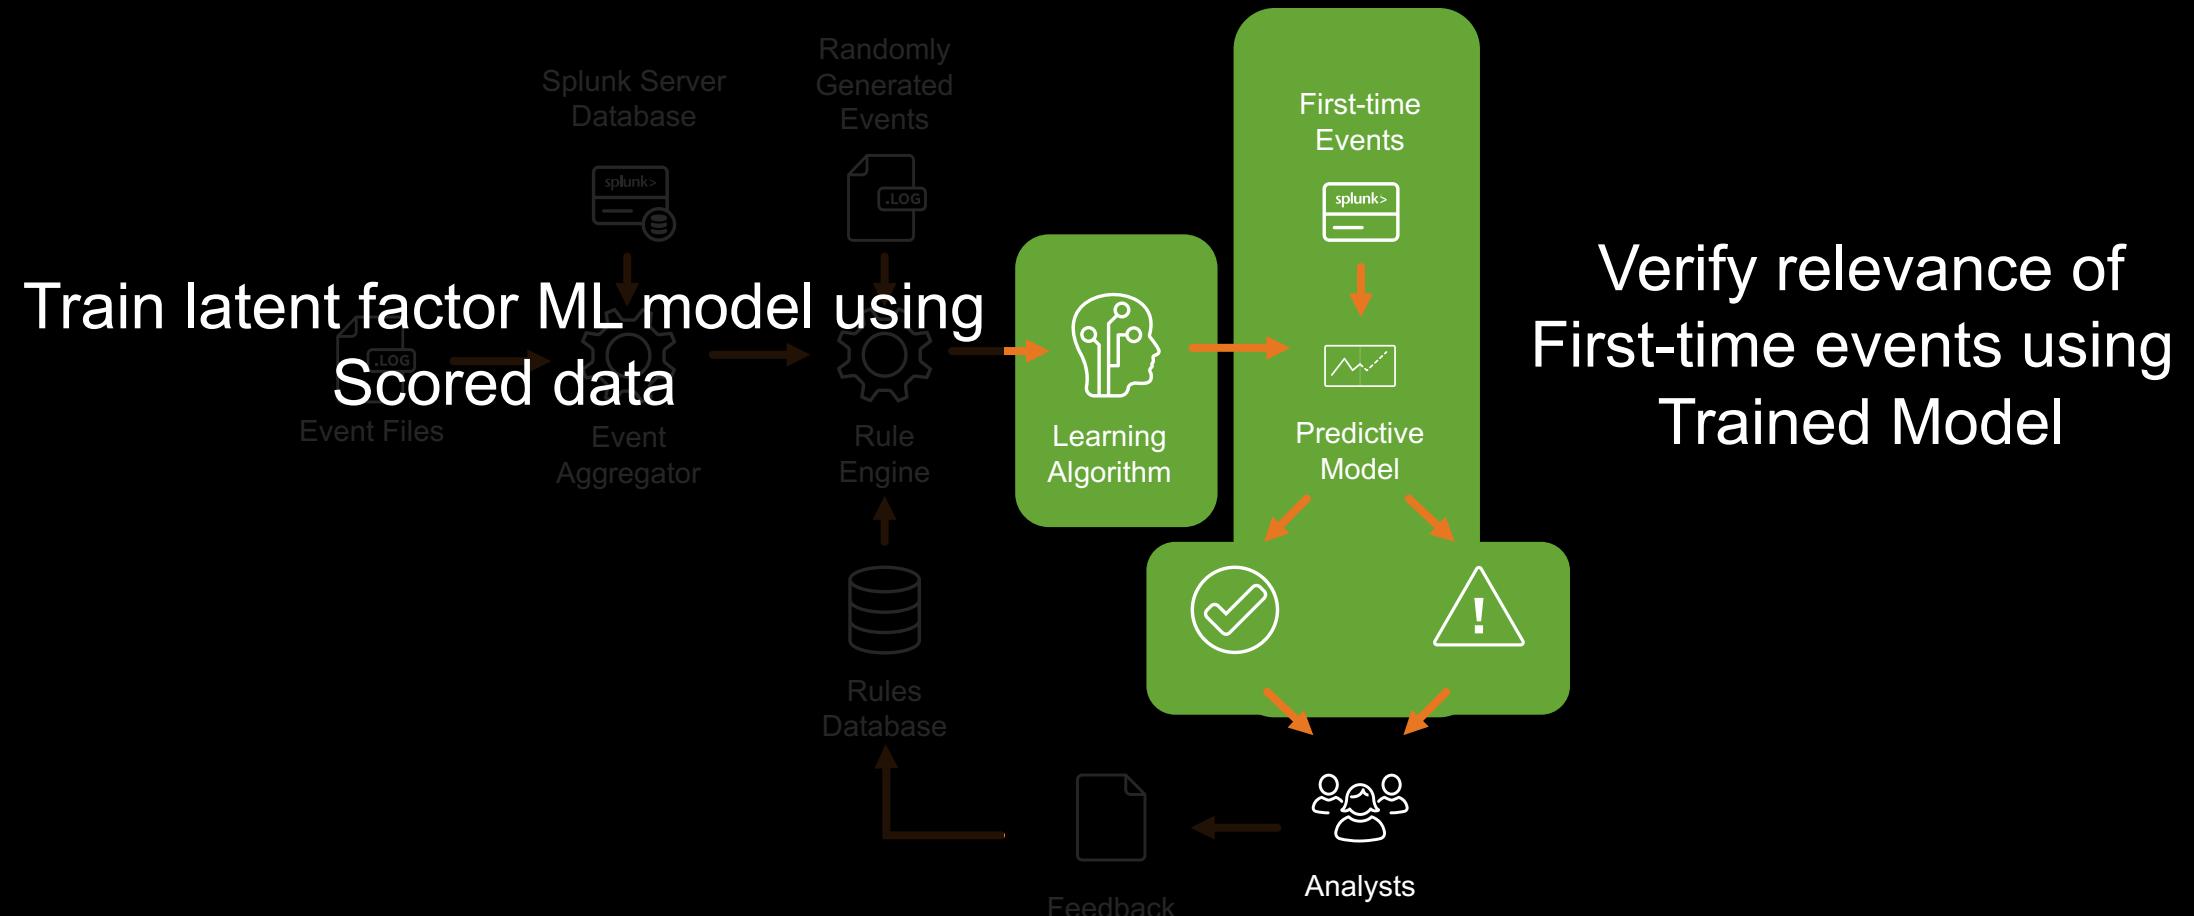
Link ID	User Properties	User Name	Events Summary	Device Name	Device Properties
1	Peergroup: 1 OU: Engineering	Bob	SSH duration 2h	Buttercup	Type: VM Domain: dev
2	Peergroup: 10 OU: Finance	Alice	IPP	2ndfloorPRT	Type: Printer Domain: iot



Link ID	User Name	Rule Score	Device Name
1	Bob	100	Buttercup
2	Alice	3	2ndfloorPRT

# Machine Learning and Prediction

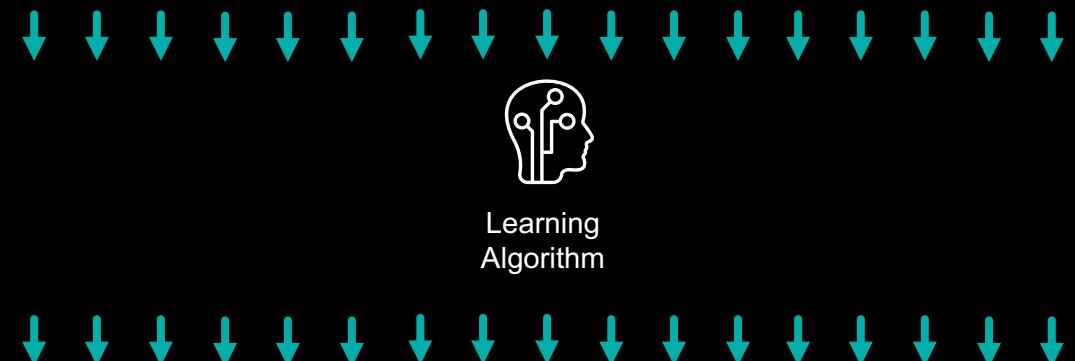
Step 3 and 4: find model that reconstruct data and use it to predict first time events



# Machine Learning

## **Step 3: Find a model that can reconstruct the data**

Link ID	User Name	Rule Score	Device Name
1	Bob	100	Buttercup
2	Alice	3	2ndfloorPRT



User Name	Latent Factors
Bob	[0.0, 0.5, 0.5]
Alice	[0.8, 0.2, 0.0]

$$f(u,v)$$

Device Name	Latent Factors
Buttercup	[0.0, 0.0, 1.0]
2ndfloorPRT	[0.2, 0.8, 0.0]

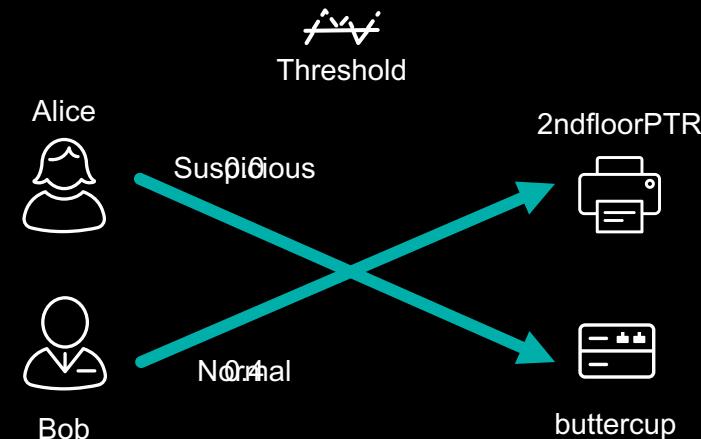
# Predictions

# Step 4: Predict first time events

User Name	Latent Factors
Bob	[0.0, 0.5, 0.5]
Alice	[0.8, 0.2, 0.0]

Device Name	Latent Factors
Buttercup	[0.0, 0.0, 1.0]
2ndfloorPRT	[0.2, 0.8, 0.0]

$$f(u, m) = u \cdot m = u_0 m_0 + u_1 m_1 + u_2 m_2$$



The diagram illustrates Bob's session flow. It starts with a user icon labeled "Normal" above it, followed by the word "Bob". A teal arrow points from Bob to a second user icon labeled "bot" below it. The session log on the left shows a sequence of requests from Bob, including category views, product details, and a purchase attempt. The session log on the right shows a sequence of requests from the bot, including category views, product details, and a purchase attempt.

# Results

Applying this technique on data

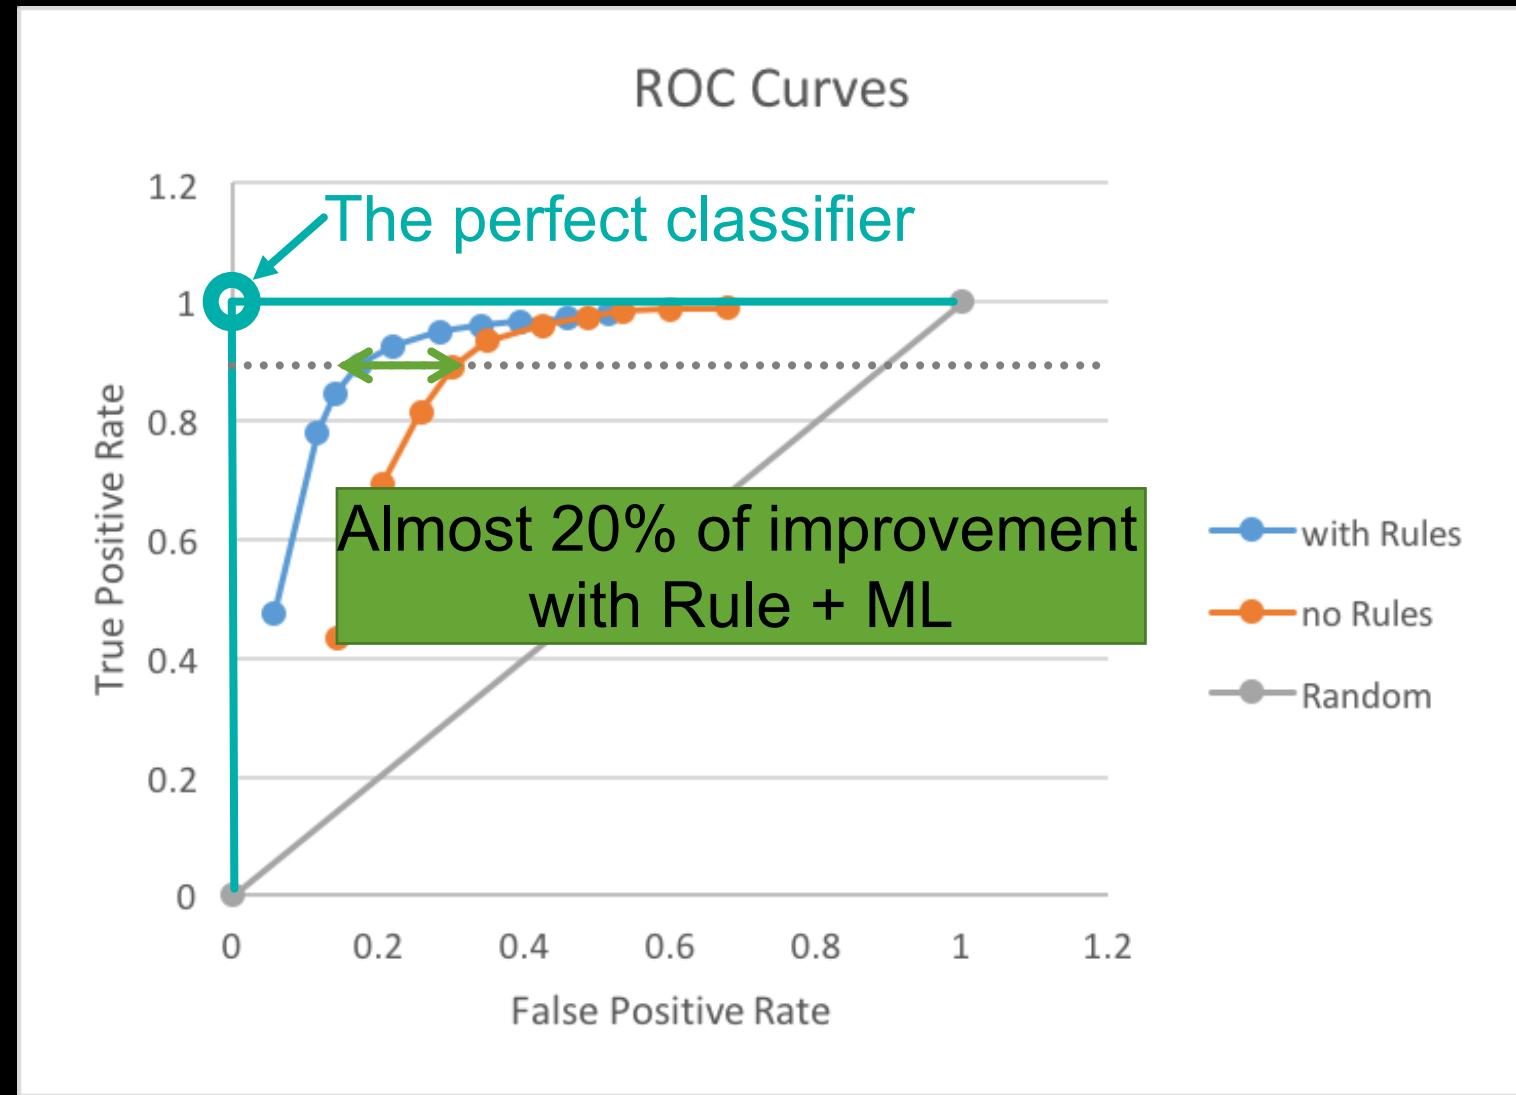


# Device Logins

- ▶ Login events into devices
    - Users have different roles that define their login behaviors
    - Heterogeneous set of device types
    - Other properties such as employee department, device subdomain are available
    - Distribution of device types, and user roles are biased
  - ▶ Suspicious accesses happen with low probability
    - Accesses out-of-profile behavior

# Pure ML vs Rule + ML

Performance of approach in terms of True and False Positives



# Final Words



# Conclusion

## Better than rules or machine learning alone

- ▶ Reduces the volume of first time events for analysts to look at
  - ▶ Allows early detection of suspicious activity
  - ▶ Customizable ML through rules and feedback
  - ▶ Learns natural behavior not foreseen by rules

# Beyond Prediction

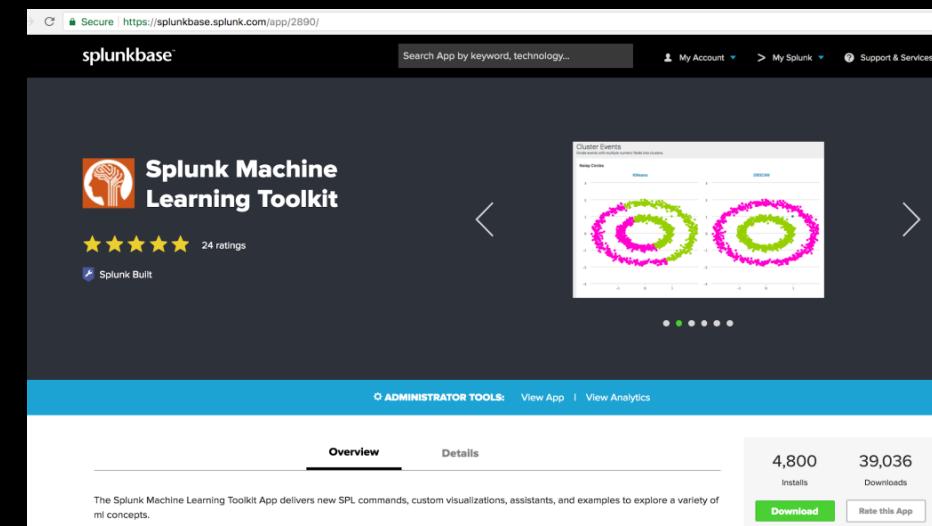
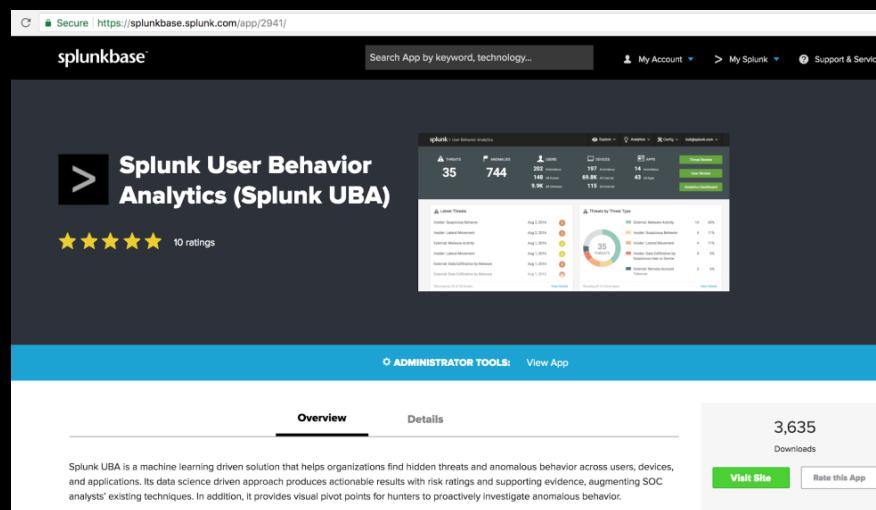
# Not just predictions of suspicious events

- ▶ Latent factors are represented by a vector
    - It encodes behavior and characteristics pretty much like a DNA for entities
  - ▶ These can be clustered using Splunk MLTK
    - Find behavioral peer-groups of users and devices
    - Define user-device affinity groups
    - And many other interesting insights of your data

# Hands on: Where to go from Here!

Get the power of rules and collaborative filtering algorithms

## Splunk User Behavior Analytics Splunk-UBA



## Splunk Machine Learning Toolkit Splunk-MLTK

# Q&A

**Ignacio Bermudez Corrales | icorrales@splunk.com**

# First Time Events in Real Life

**Some are normal, some can be risky**

## ► Machine Accesses

- IT workers frequently update machines. One updates a QA testing VM, is this risky?

## ► Cloud Storage File-sharing

- Files are download or previewed once, some may be read by the “wrong” people

## ► Private Code Repositories

- An employee's account got compromised and cloned projects with sensitive intellectual property

## ▶ Badge Accesses

- Some worker entered the servers room of his office out of working hours

# We Want Early Detection

# The risk of first time event and why you should care

- ▶ Event linked to suspicious activity can happen only once
    - Attacker spear downloads single confidential file to local hard drive
    - Suspicious activity happens once and there are not after events to make a decision
      - Volume of data download, number of devices accessed, number of repos cloned, etc.
  - ▶ Detecting suspicious activity based on post-events can be risky
    - Collecting post-events can take time, can we make an early decision?
  - ▶ Many normal operations happen for a first time
    - Employee access some colleague's dev machine, architect downloads confidential document regarding product design

# First Time Events in Real Life

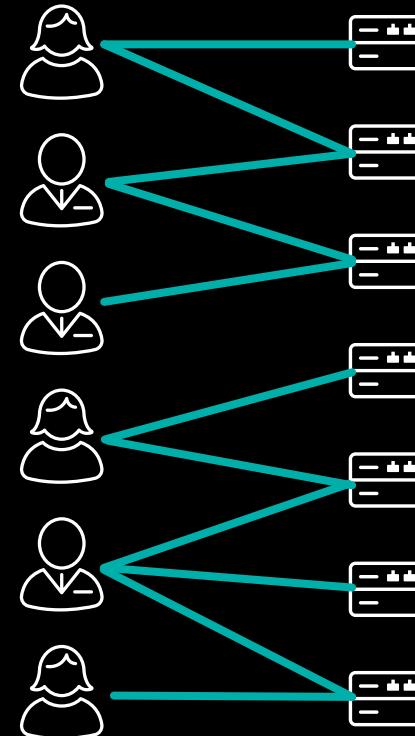
**Some are normal, some can be risky**

- ▶ Users access similar set of machines on a daily basis
    - IT workers frequently update machines. One updates a QA testing VM, is this risky?
  - ▶ Employees share files in some cloud storage system
    - Files are download or previewed once, some may be read by the “wrong” people
  - ▶ A git repository where all developers have read access
    - An employee’s account got compromised and cloned projects with sensitive intellectual property
  - ▶ Company keeps record of badge accesses
    - Some worker entered the servers room of his office out of working hours

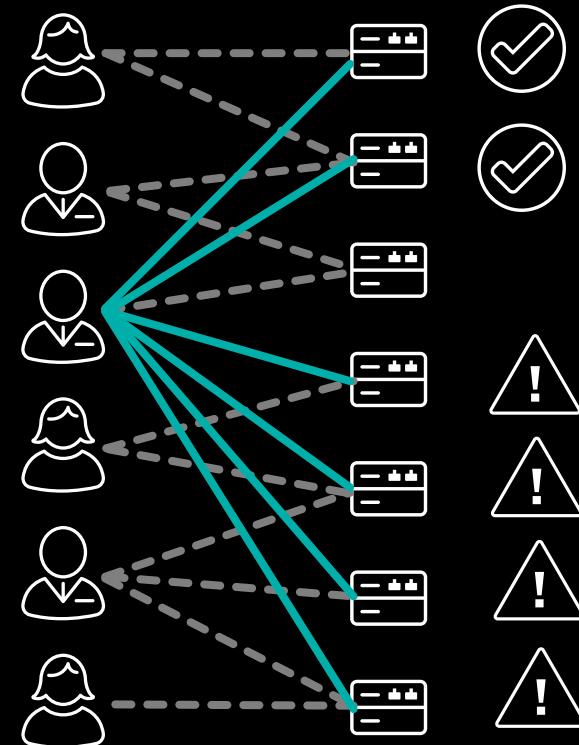
# Machine Accesses

Do all first time accesses look suspicious?

Historical Accesses



First Time Accesses



Normal Operations?

Suspicious Accesses?

```

138.60.4.128.241.220.82 - [07/Jan/18:10:57:153] "GET /category.screen?category_id=GIFTS&SESSIONID=SD15LAFF10ADFF10 HTTP 1.1" 404 720 "http://buttercup-shopping.com/cart.do?action=view&itemId=EST-6&product_id=EST-5&category_id=EST-5&sw=0" "Mozilla/5.0 (Windows NT 5.1; SV1; .NET CLR 1.1.4322; .NET CLR 2.0.50727; .NET CLR 3.0.4506.2152; .NET CLR 3.5.30729; .NET4.0E) AppleWebKit/534.30 (KHTML, like Gecko) Chrome/12.0.742.112 Safari/534.30"
138.60.4.128.241.220.82 - [07/Jan/18:10:57:153] "GET /product.screen?product_id=FL-DSH-01&SESSIONID=SD55L7FF6ADFF9 HTTP 1.1" 404 332 "http://buttercup-shopping.com/cart.do?action=printless&itemId=EST-26&product_id=EST-26&category_id=EST-26&sw=0" "Mozilla/5.0 (Windows NT 5.1; SV1; .NET CLR 1.1.4322; .NET CLR 2.0.50727; .NET CLR 3.0.4506.2152; .NET CLR 3.5.30729; .NET4.0E) AppleWebKit/534.30 (KHTML, like Gecko) Chrome/12.0.742.112 Safari/534.30"
138.60.4.128.241.220.82 - [07/Jan/18:10:57:153] "GET /oldlink?item_id=EST-26&SESSIONID=SD55L9FF1ADEF3 HTTP 1.1" 200 1318 "http://buttercup-shopping.com/cart.do?action=changequantity&itemId=EST-18&product_id=EST-18&category_id=EST-18&SESSIONID=SD55L9FF1ADEF3" "Mozilla/5.0 (Windows NT 5.1; SV1; .NET CLR 1.1.4322; .NET CLR 2.0.50727; .NET CLR 3.0.4506.2152; .NET CLR 3.5.30729; .NET4.0E) AppleWebKit/534.30 (KHTML, like Gecko) Chrome/12.0.742.112 Safari/534.30"
138.60.4.128.241.220.82 - [07/Jan/18:10:57:153] "GET /oldlink?item_id=SURPRISE&SESSIONID=SD85LBF2ZADFF9 HTTP 1.1" 200 385 "http://buttercup-shopping.com/cart.do?action=oldlink?item_id=SURPRISE&SESSIONID=SD85LBF2ZADFF9&product_id=SURPRISE&category_id=SURPRISE&sw=0" "Mozilla/5.0 (Windows NT 5.1; SV1; .NET CLR 1.1.4322; .NET CLR 2.0.50727; .NET CLR 3.0.4506.2152; .NET CLR 3.5.30729; .NET4.0E) AppleWebKit/534.30 (KHTML, like Gecko) Chrome/12.0.742.112 Safari/534.30"
138.60.4.128.241.220.82 - [07/Jan/18:10:57:153] "GET /oldlink?item_id=SURPRISE&SESSIONID=SD85LBF2ZADFF9 HTTP 1.1" 200 385 "http://buttercup-shopping.com/cart.do?action=oldlink?item_id=SURPRISE&SESSIONID=SD85LBF2ZADFF9&product_id=SURPRISE&category_id=SURPRISE&sw=0" "Mozilla/5.0 (Windows NT 5.1; SV1; .NET CLR 1.1.4322; .NET CLR 2.0.50727; .NET CLR 3.0.4506.2152; .NET CLR 3.5.30729; .NET4.0E) AppleWebKit/534.30 (KHTML, like Gecko) Chrome/12.0.742.112 Safari/534.30"
  
```

# Examples of First Time Events

- ▶ First time ...
    - A user access a machine
    - A user enters a room in a building
    - A user opens a file in a private cloud
    - A user pulls code from a git repository
    - A user logins from a particular location
    - An employee mails another employee
    - And many other examples
  - ▶ First time events can be pretty common

# First Time Event Handling

# How to deal with first time events?

## ► Conservative Approach

- All first time events are suspicious and worth looking at one by one
  - In scenarios where the number of first time events is huge this requires excessive resources

## ► Laissez-faire Approach

- We do not care about first time events
  - Easy to implement, but won't catch any concerning first time event

## ► Use Rules

- They can cover low hanging fruits, but can't cover corner cases

## ► Use Machine Learning

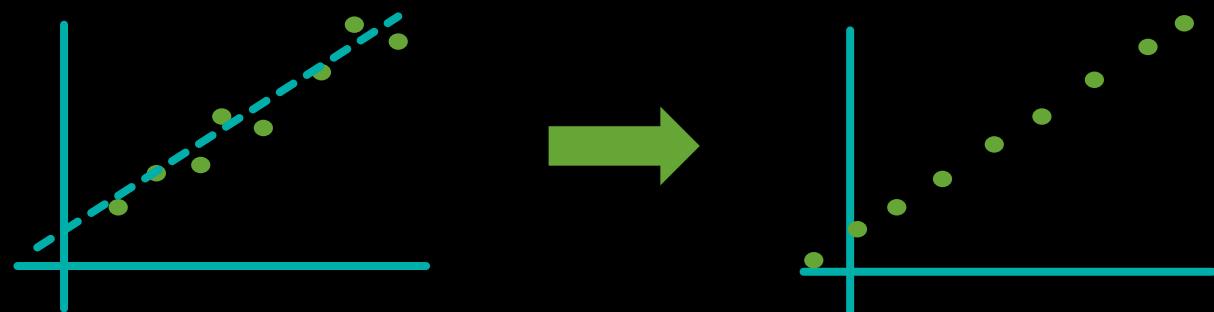
- Learns from historical patterns in the data, but can't learn from what haven't happened

► The answer resides between the last two!

# Some Machine Learning Background

## Things we should keep in mind through the presentation

- ▶ Machine learning algorithms approximate numerical data with a function
    - History of events have to be represented as numbers to learn
    - Linear regression: Given data points  $(x, y)$  find  $a$  and  $b$  such that  $y' = a x + b$  can **reconstruct** the evidence



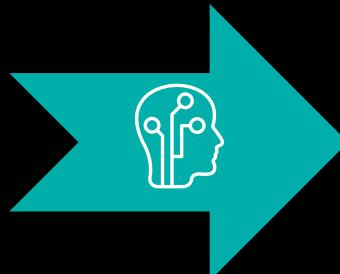
We need to model our problem **numerically** and we will find **parameters** that can reconstruct the evidence

# Machine Learning Generalizes

**ML allows reconstruction of the data**

## Historical Data - Evidence

	1	1					
		1	1				
			1				
				1	1		
					1	1	1
							1



# ML Reconstruction of Data

	1	1	0.8	0.1	0.1	0.1	0.1
	0.8	1	1	0.1	0.1	0.1	0.1
	0.8	0.8	1	0.1	0.1	0.1	0.1
	0.1	0.1	0.1	1	1	0.8	0.8
	0.5	0.5	0.5	0.8	1	1	1
	0.1	0.1	0.1	0.5	0.8	0.8	1

The objective of the ML algorithm is to reconstruct the evidence  
With the minimum error possible using an **evaluation function**

# Taming ML with Rules - Review

## Passing down knowledge to machine learning algorithm

## ► Rules:

- Embed human knowledge
  - Map event properties into numerical signal
  - Consume external properties of event entities
  - May have low coverage
  - Complements what can be seen in the history
  - Don't generalize

## ► Passing them down

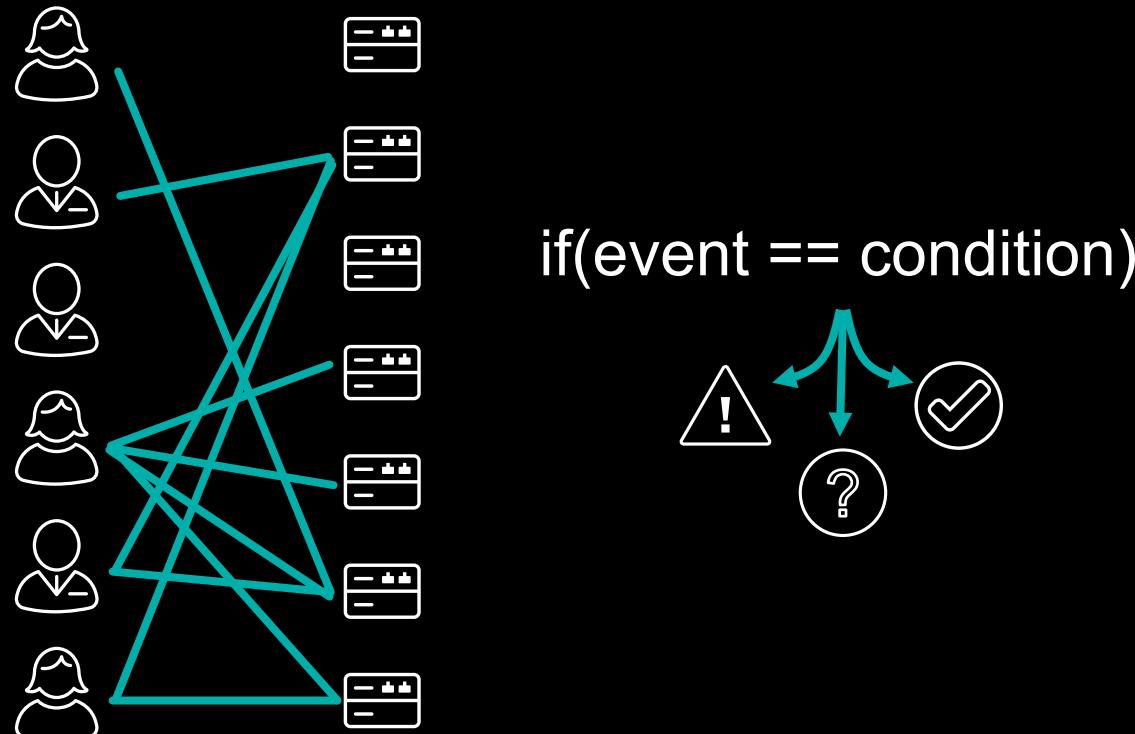
- Generate random events
  - If random event hit a rule add them to history
  - Control the number of random events to avoid overfitting

## Coverage of Rules

							<b>1</b>	
		<b>1</b>						
					<b>1</b>	<b>1</b>	<b>0</b>	<b>0</b>
		<b>0</b>					<b>1</b>	
		<b>0</b>					<b>1</b>	

# Taming ML with Rules - Review

# Passing down knowledge to machine learning algorithm



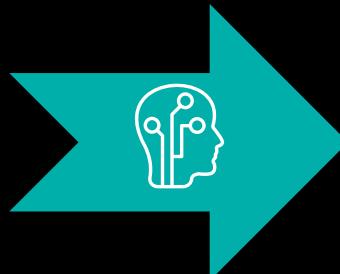
## Coverage of Rules

								1
		1						
					1	1	0	0
		0					1	
		0					1	

# Machine Learning Generalizes

## **ML allows reconstruction of the data**

## Historical Data - Evidence



# ML Reconstruction of Data

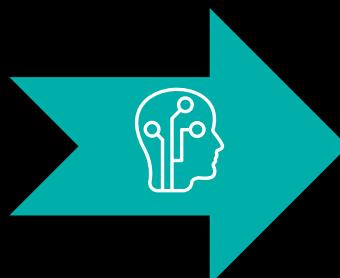
	1	1	0.8	0.1	0.1	0.1	0.1
	0.8	1	1	0.1	0.1	0.1	0.1
	0.8	0.8	1	0.1	0.1	0.1	0.1
	0.1	0.1	0.1	1	1	0.8	0.8
	0.5	0.5	0.5	0.8	1	1	1
	0.1	0.1	0.1	0.5	0.8	0.8	1

This is analogous to the linear regression problem when  $y' = a x + b$   
Reconstructs the Data

# Machine Learning Generalizes

## How a machine learning model predicts first time events

## Historical Data - Evidence



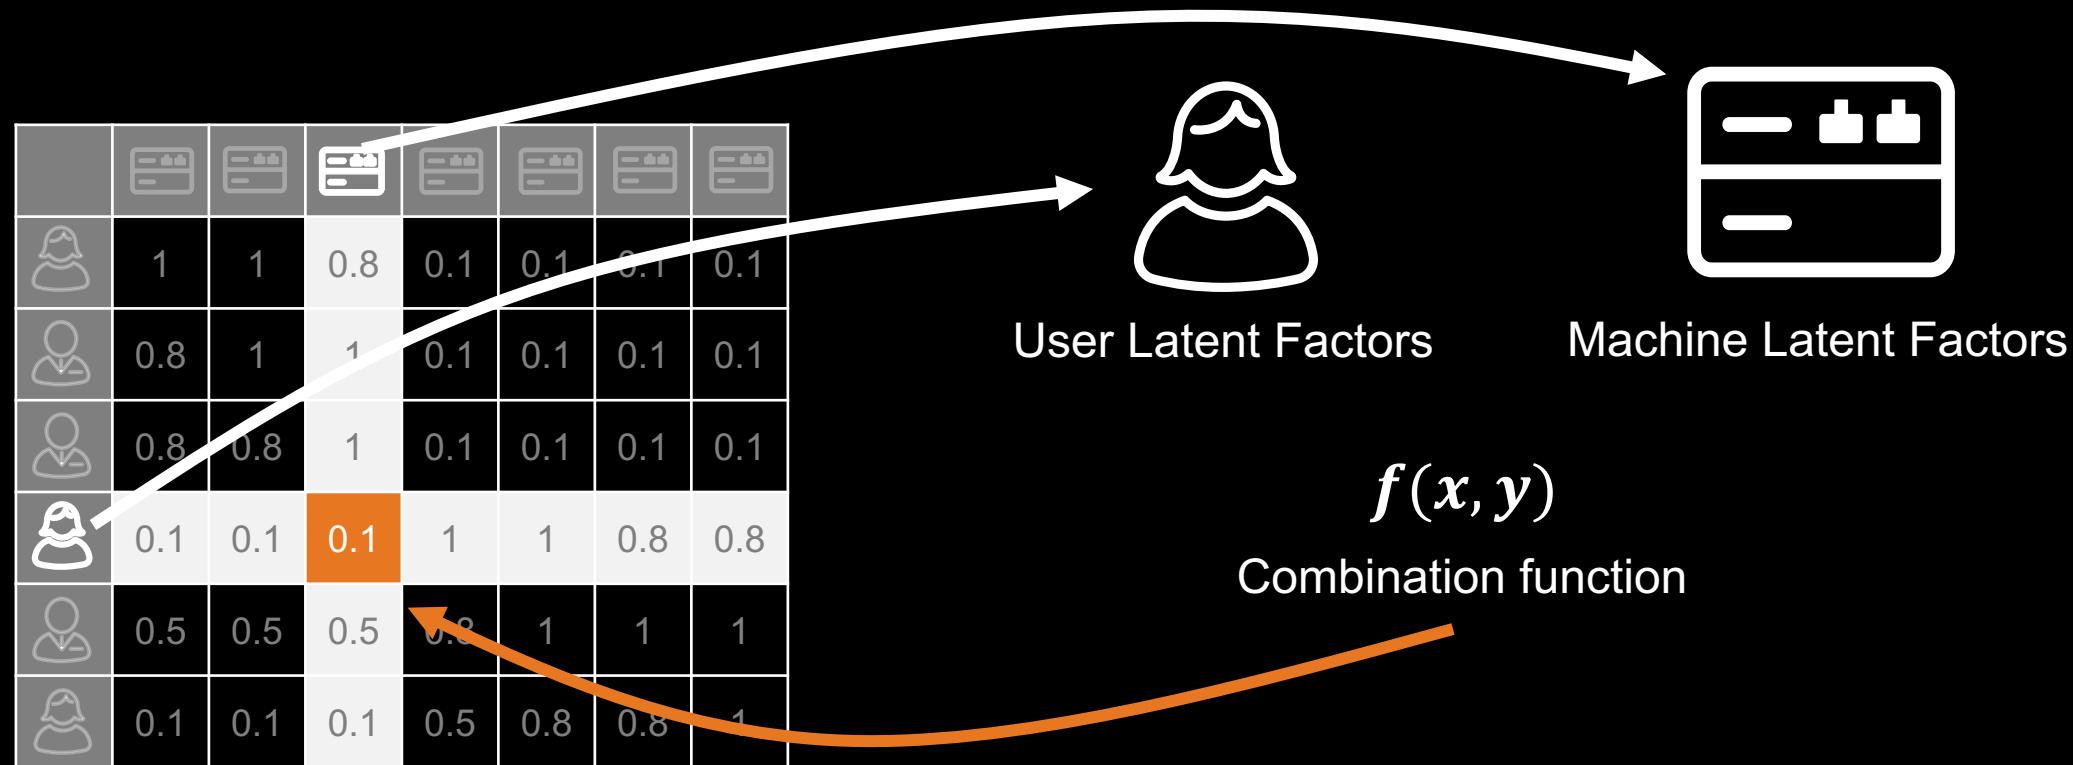
# ML Reconstruction of Data

	1	1	0.8	0.1	0.1	0.1	0.1
	0.8	1	1	0.1	0.1	0.1	0.1
	0.8	0.8	1	0.1	0.1	0.1	0.1
	0.1	0.1	0.1	1	1	0.8	0.8
	0.5	0.5	0.5	0.8	1	1	1
	0.1	0.1	0.1	0.5	0.8	0.8	1

This is analogous to the linear regression problem when  $y' = a x + b$   
Reconstructs the Data

# Machine Learning Generalizes

# How a machine learning model predicts first time events



Analogous to the linear regression problem  $y' = a x + b$

**Latent factors** are the parameters ( $a, b$ ) and the function combines them

# How Data is Reconstructed

## Latent factor models

## Machine Latent Factors (y)

User Latent Factors (x)		Card A	Card B	Card C	Card D	Card E	Card F	Card G
Profile 1	1	1	0.8	0.1	0.1	0.1	0.1	0.1
Profile 2	0.8	1	1	0.1	0.1	0.1	0.1	0.1
Profile 3	0.8	0.8	1	0.1	0.1	0.1	0.1	0.1
Profile 4	0.1	0.1	0.1	1	1	0.8	0.8	0.8
Profile 5	0.5	0.5	0.5	0.8	1	1	1	1
Profile 6	0.1	0.1	0.1	0.5	0.8	0.8	0.8	1

$$f(x, y) = 0.1$$

Evaluation function

In the machine learning training phase we find **latent factors** parameters, that combined with a **function**, we can recreate the values of the matrix (remember, similar to  $y' = a x + b$ )

# How Data is Reconstructed

## Latent factor models

## Machine Latent Factors ( $m$ )

User Latent Factors ( $u$ )

	Card 1	Card 2	Card 3	Card 4	Card 5	Card 6	Card 7	Card 8
User 1	1	1	0.8	0.1	0.1	0.1	0.1	0.1
User 2	0.8	1	1	0.1	0.1	0.1	0.1	0.1
User 3	0.8	0.8	1	0.1	0.1	0.1	0.1	0.1
User 4	0.1	0.1	0.1	1	1	0.8	0.8	0.8
User 5	0.5	0.5	0.5	0.8	1	1	1	1
User 6	0.1	0.1	0.1	0.5	0.8	0.8	0.8	1

Latent factors ( $u, m$ ) are real valued vectors  
that minimize the reconstruction error  
made by the evaluation function

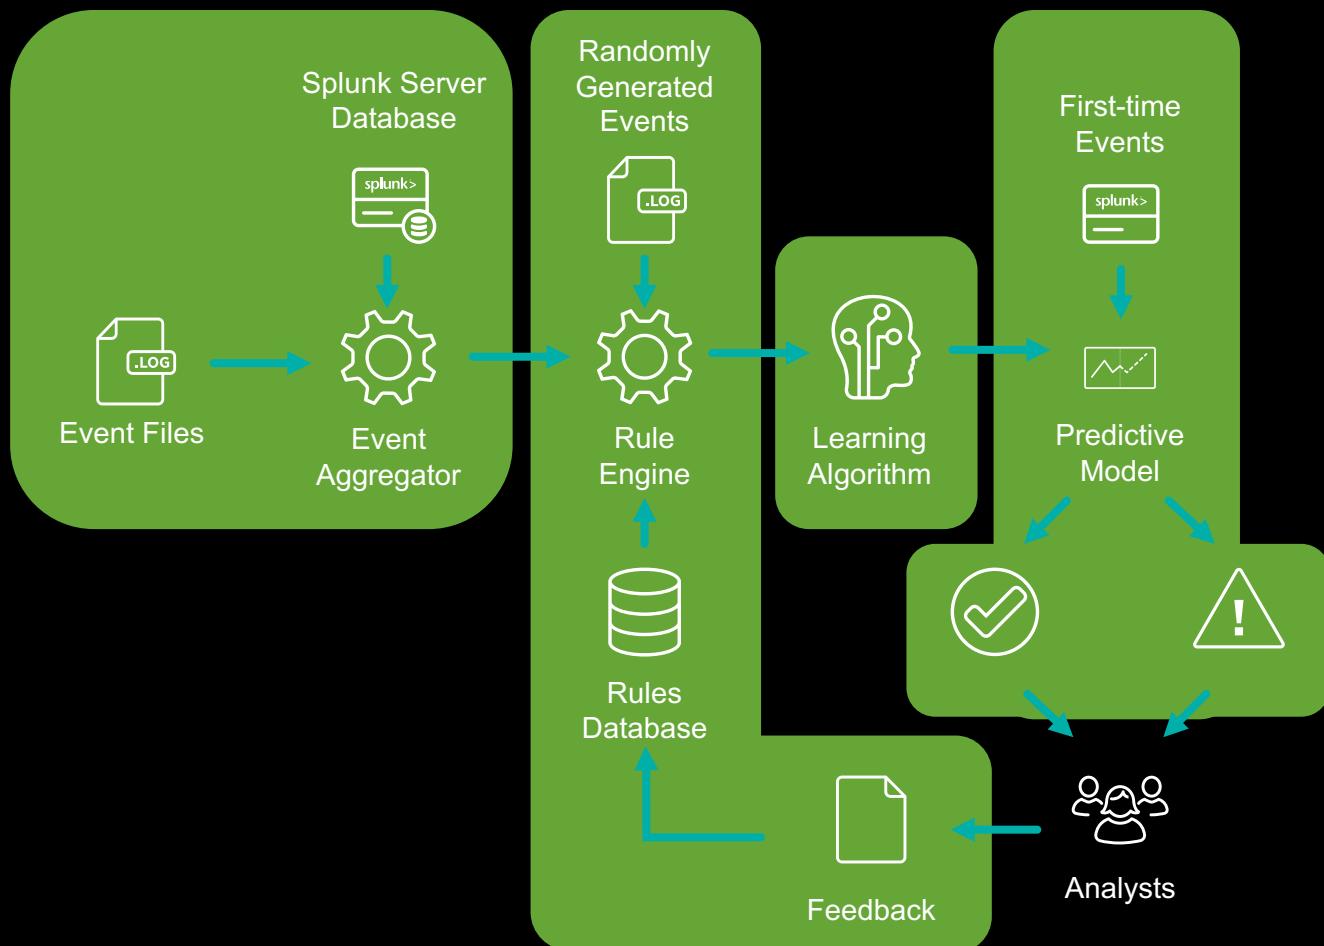
$$f(u, m) = 0.1$$

Evaluation function

This is analogous to linear regression where we find parameters ( $a$ ,  $b$ ) that minimize the error of  $y' = a x + b$

# Architecture

## 4 steps



# Event Processing

## Step 1

# Input

- Event Logs
  - External information about entity properties

# Goal

Convert single event logs  
into an annotated graph  
with event properties

## Output

# Table of aggregated events

# Event Processing

## Step 1

- ▶ Goal
  - Convert single event logs into an annotated graph with event properties
- ▶ Input
  - Event logs
  - External information about the entities
- ▶ Output
  - Table where each row is an aggregation of events between two entities
  - For each row there are properties attached to them

Link ID	User Properties	User Name	Events Summary	Device Name	Device Properties
1	Peergroup: 1 OU: Engineering	Bob	SSH duration 2h	Buttercup	Type: VM Domain: dev
2	Peergroup: 10 OU: Finance	Alice	IPP	2ndfloorPRT	Type: Printer Domain: iot

# Rule Engine

## ► Goal

- Convert links into numerical signals (scores)

## ► Input

- Table of aggregated events
- Randomly generated event aggregations
- Rules database
- Previous analyst feedback

## ► Output

- Table with scored links

Link ID	User Name	Rule Score	Device Name
1	Bob	100	Buttercup
2	Alice	3	2ndfloorPRT

# Machine Learning Algorithm

## ► Goal

- Generate **latent factors** for entities that allow reconstruction of the data

## ▶ Input

- Real valued links with labeled entity identifiers

## ▶ Output

- Evaluation function that can operate on latent factors (given by the algorithm)
  - Latent features for each entity

User Name	Latent Factors
Bob	[0.0, 0.5, 0.5]
Alice	[0.8, 0.2, 0.0]

Device Name	Latent Factors
Buttercup	[0.0, 0.0, 1.0]
2ndfloorPRT	[0.2, 0.8, 0.0]

# Predictions

## Step 4

## ► Goal

- Predict anomalous first time events

## ▶ Input

- Database of entity latent factors
  - Function that operates on latent factors
  - Anomaly threshold

# ▶ Output

- Labeled anomalous first time events

# Example

User factors lookup table

User Name	Latent Factors
Bob	[0.0, 0.5, 0.5]
Alice	[0.8, 0.2, 0.0]

Device factors lookup table

Device Name	Latent Factors
Buttercup	[0.0, 0.0, 1.0]
2ndfloorPRT	[0.2, 0.8, 0.0]

Evaluation function

$$f(x, y) = x \cdot y = x_0y_0 + x_1y_1 + x_2y_2$$

Predicted scores

	Buttercup	2ndfloorPRT
Bob	0.5	0.4
Alice	0.0	0.8

Threshold



Predictions

	Buttercup	2ndfloorPRT
Bob	Normal	Normal
Alice	Suspicious	Normal

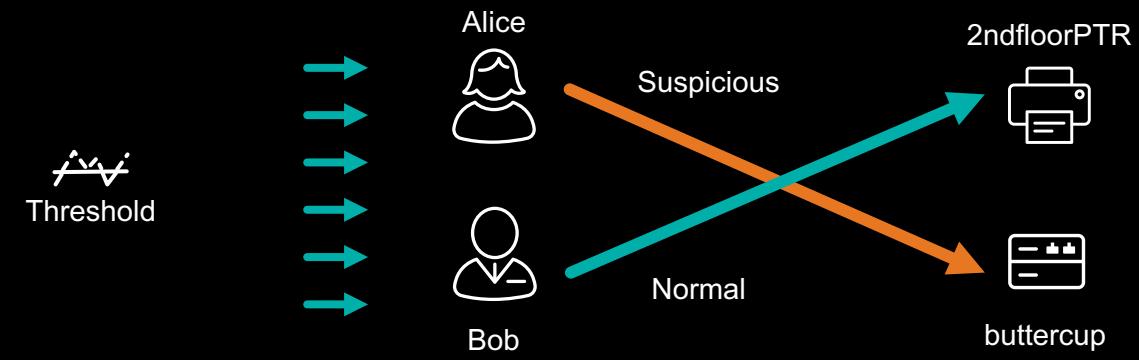
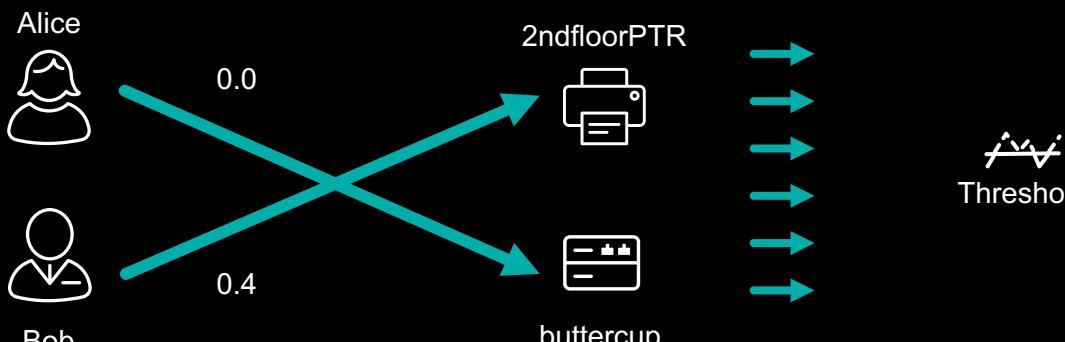
# Predictions

## Step 4: Predict first time events

User Name	Latent Factors
Bob	[0.0, 0.5, 0.5]
Alice	[0.8, 0.2, 0.0]

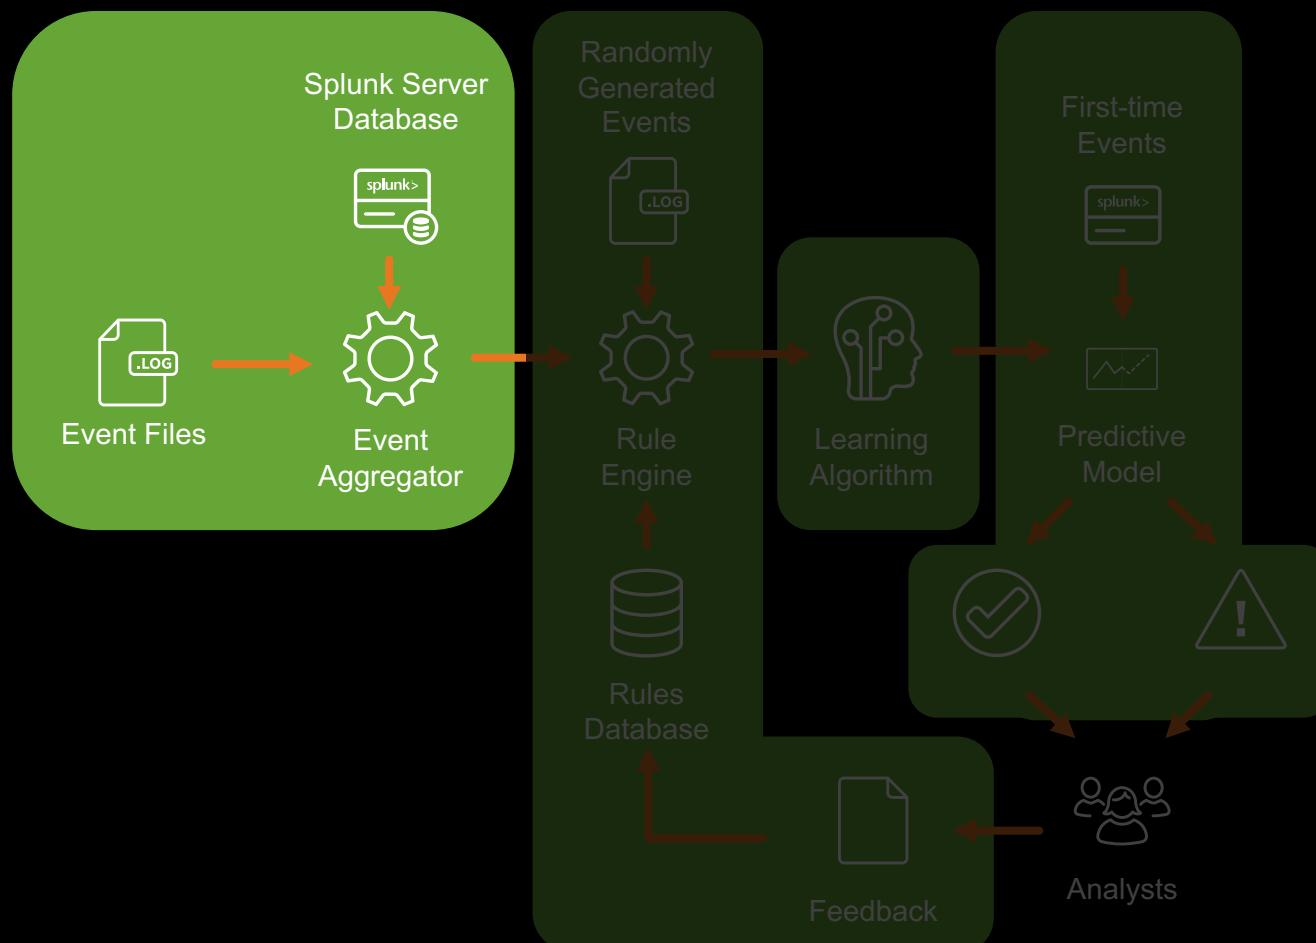
Device Name	Latent Factors
Buttercup	[0.0, 0.0, 1.0]
2ndfloorPRT	[0.2, 0.8, 0.0]

$$f(x, y) = x \cdot y = x_0y_0 + x_1y_1 + x_2y_2$$



# Event Processing

## Step 1: Build a summary of information for the rule engine



# Event Processing

## Step 1

# Input

- Event Logs
  - External information about entity properties

# Goal

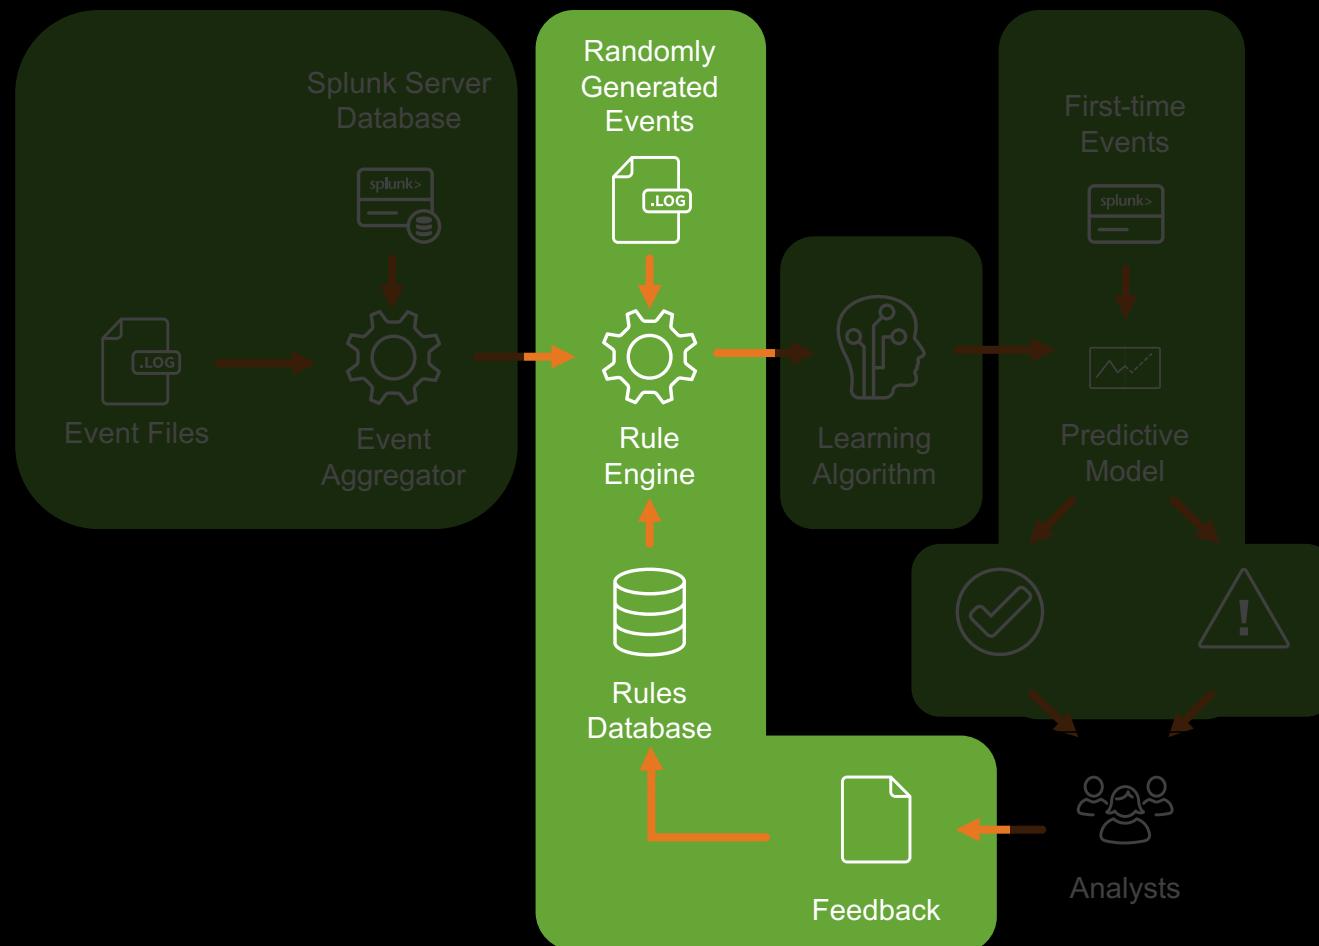
Provide as much information as possible to the rules to make decisions

## Output

## Table of aggregated events

# Rule Engine

## Step 2: Convert event summaries into numerical signals



# Rule Engine

## Step 2

# Input

- Table of aggregated events
  - Randomly generated events
  - Rules database
  - Previous feedback

# Goal

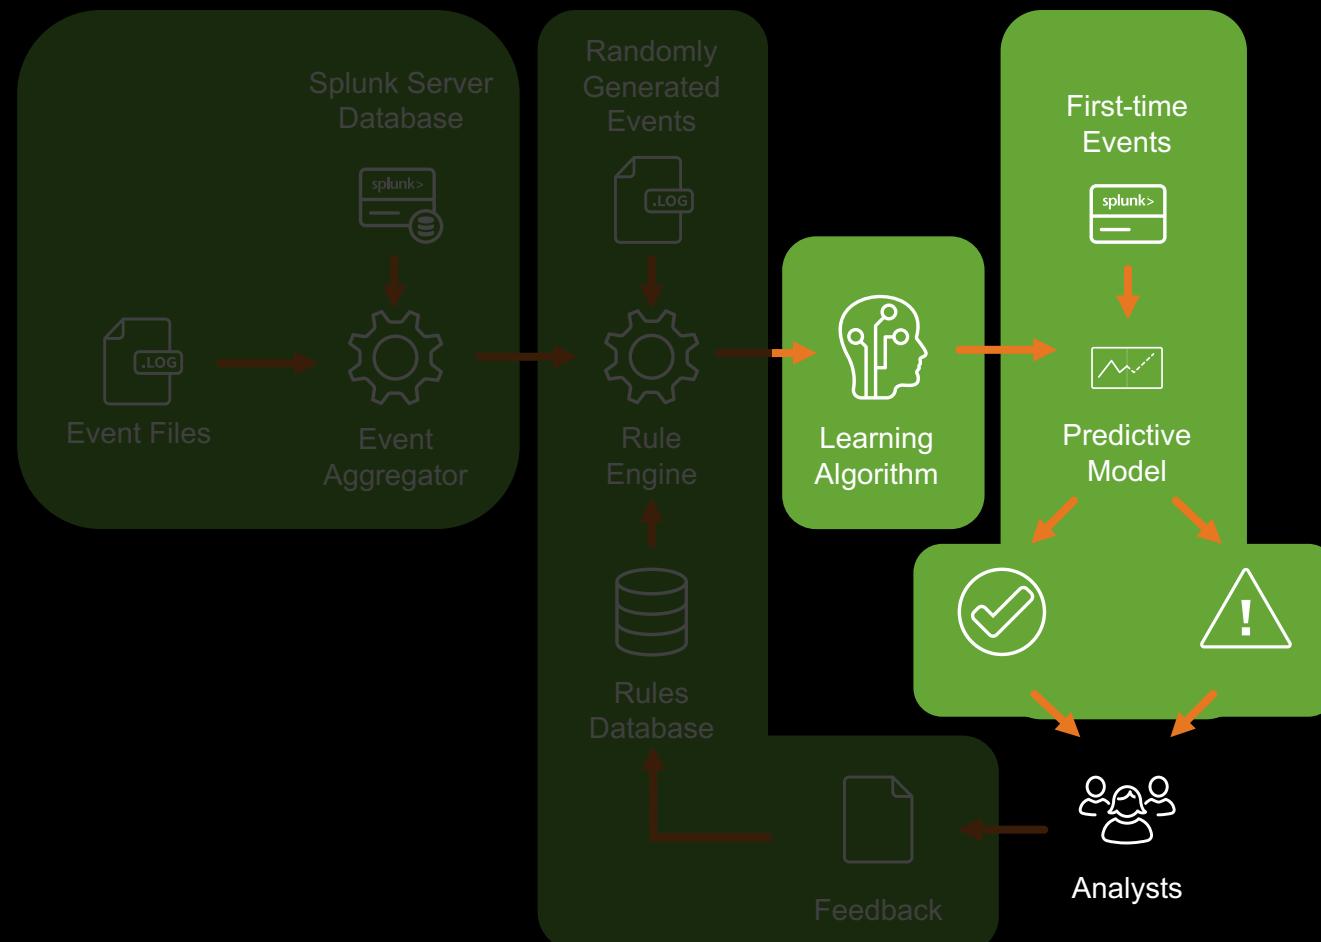
Convert aggregations into numerical signals, so they can be learnt by a Machine Learning algorithm

# Output

## Table of scored links

# Machine Learning and Prediction

Step 3 and 4: find model that reconstruct data and use it to predict first time events



# Machine Learning

## Step 3

# Input

## Table of scored links

# Goal

Find a model that can reconstruct the data

# Output

- Evaluation function
  - Latent factors for each entity

# Predictions

## Step 4

# Input

- Table of aggregated events
  - Randomly generated events
  - Rules database
  - Previous feedback

# Goal

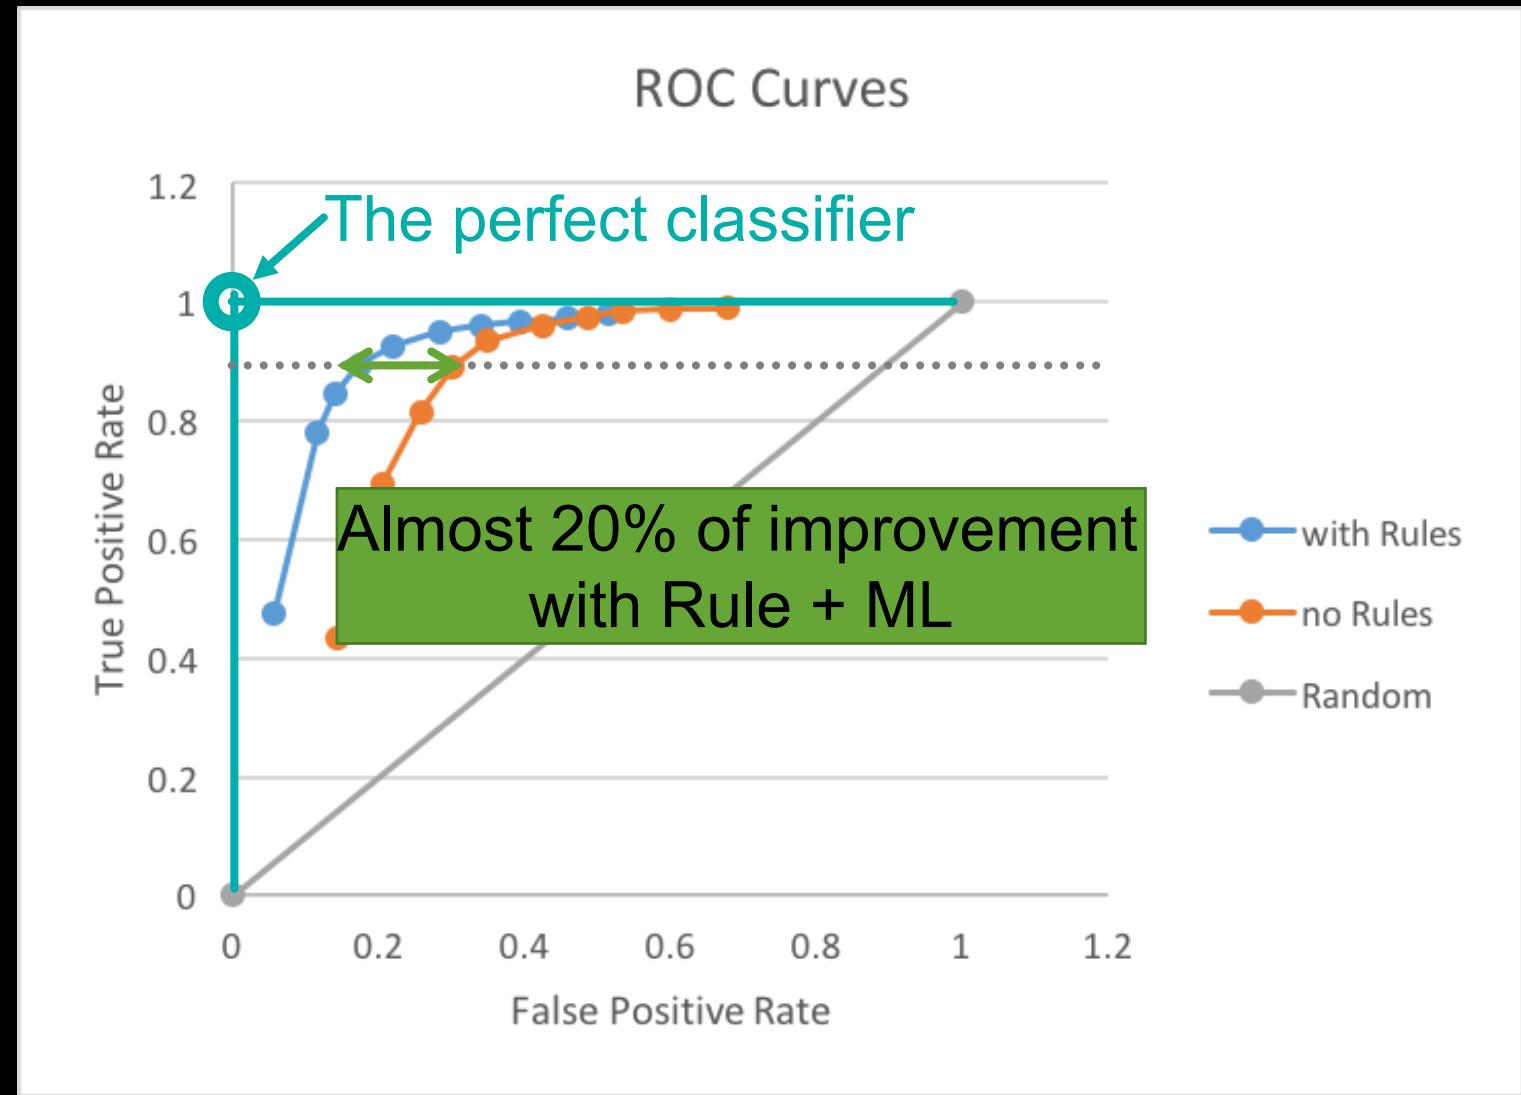
# Predict anomalous first time events

# Output

## Labeled first time events

# Pure ML vs Rule + ML

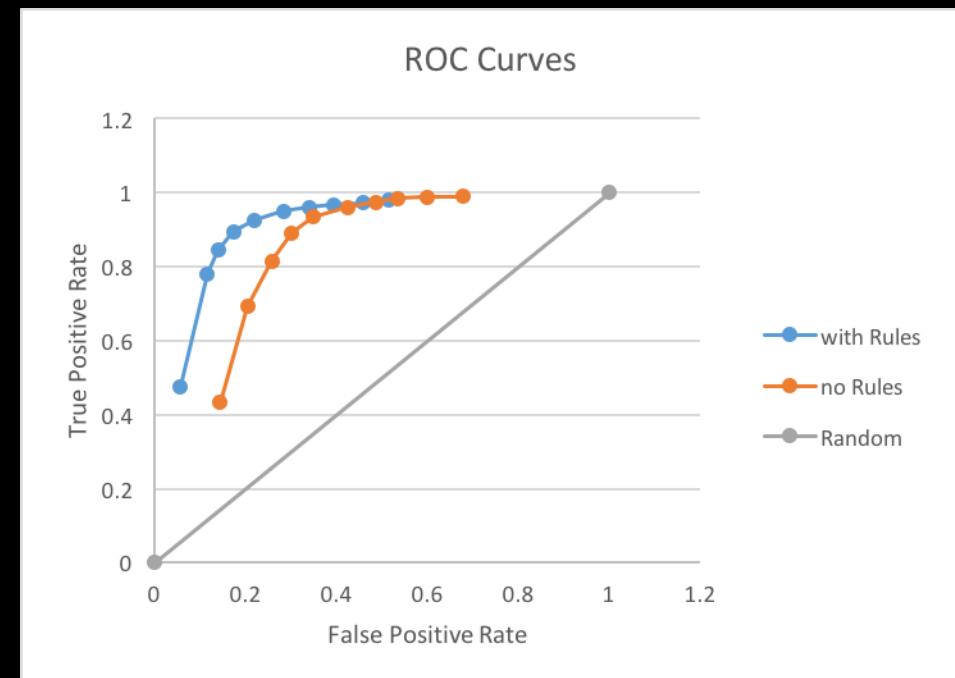
Performance of approach in terms of True and False Positives



# Pure ML vs Rule + ML

## Performance of approach in terms of True and False Positives

- ▶ Compare approach using and not using rules
  - ▶ Rate of TP/FP by moving the detection threshold
  - ▶ Rules improve overall performance
  - ▶ Same TP rate lesser FP
    - TP = 0.9, FP = 0.2 with rules
    - TP = 0.9, FP = 0.35 without rules



# Conclusion ( make 2 slides)

**How this is implemented on Splunk. Where you can find it in splunk. A second slide is how this is gonna help you, compare to other products if possible (competitors).**

- ▶ This can be implemented using Alternating Least Squares (ALS) on Spark
    - Can also be implemented in Python with some specific packages
  - ▶ UBA suspicious first Box accesses use this approach
    - Captures when a user access a resource for the first time
    - Uses Active Directory extracted peer-groups
    - Rules based on file properties
  - ▶ Not just prediction of suspicious first time event
    - Latent factors as “DNA” for entity behavior or characteristic
    - Latent factors ... explain .. Like a DNA
    - Enables further behavioral clustering of entities
  - ▶ What this all means for a business (the impact, why this is important)
    - Save analyst time
    - It learns iteratively, custom to user needs

# Conclusion

## Better than rules or machine learning alone

- ▶ Allows early detection of suspicious activity
  - Reduces the volume of first time events the analyst has to look at
- ▶ Customizable through rules and feedback
  - Analysts can impact ML model by providing feedback on specific events or by writing rules
- ▶ Learns natural behavior not foreseen by rules
  - Collaborative filtering extrapolates from historical interactions

# Hands on

## Get the power of rules and collaborative filtering algorithms

# ▶ Splunk UBA

- Currently used for detecting suspicious access to resources in Box
  - Leverages Spark ML collaborative filtering API

## ► Splunk Machine Learning Toolkit (MLTK)

- Adding more algorithms from Spark ML soon

- ▶ Not just predictions of suspicious events

- Latent factors represent a vector encoding behavior, pretty much like a DNA for entities
  - They can be used to cluster entities based in their behavior to find interesting insights

# Hands on

## Get the power of rules and collaborative filtering algorithms

# ▶ Splunk UBA

- Currently used for detecting suspicious access to resources in Box
  - Leverages Spark ML collaborative filtering API

## ► Splunk Machine Learning Toolkit (MLTK)

- Adding more algorithms from Spark ML soon