

RSA® Conference 2019

San Francisco | March 4–8 | Moscone Center



SESSION ID: CSV-W10

Treating Cloud-Specific Threats with Automatic Remediation

Jay Ball

Head of Application Security
Zocdoc

Josh Barons

Head of Information Security
Zocdoc



#RSAC

Agenda & Goals

- Zocdoc's transformation to the cloud
- Provide examples of threats and issues we encountered both external and internal
 - AWS
 - S3 Bucket Encryption
 - VPC Port Inspection and Closure
 - Ownership
 - Instance Termination
 - GitHub
 - Public Repos
 - Repo ownership
- Offer solutions
- Get buy-in

RSA®Conference2019

Zocdoc's move to the Cloud

Zocdoc's path to the cloud

- Zocdoc is a mission driven healthcare and technology company that needed to move faster
- Innovation was being slowed down by traditional infrastructure barriers
- Moving to the cloud was identified as the most efficient way to remove IT bottlenecks and development barriers

RSA® Conference 2019

The Cloud





Cloud

- Cloud is here - It's already been here and probably in your organization also
- Who is responsible - Providers have a shared responsibility model
- Visibility - How can you see into your environments
- Problems at scale - Minor issues easily become major at scale
- Internal ownership - Who owns the relationship?
- Cloud specific threats - Public exposure - Unencrypted data - Open Security groups
 - Compromised Instance - Visibility and auditing

RSA®Conference2019

Encryption at Rest

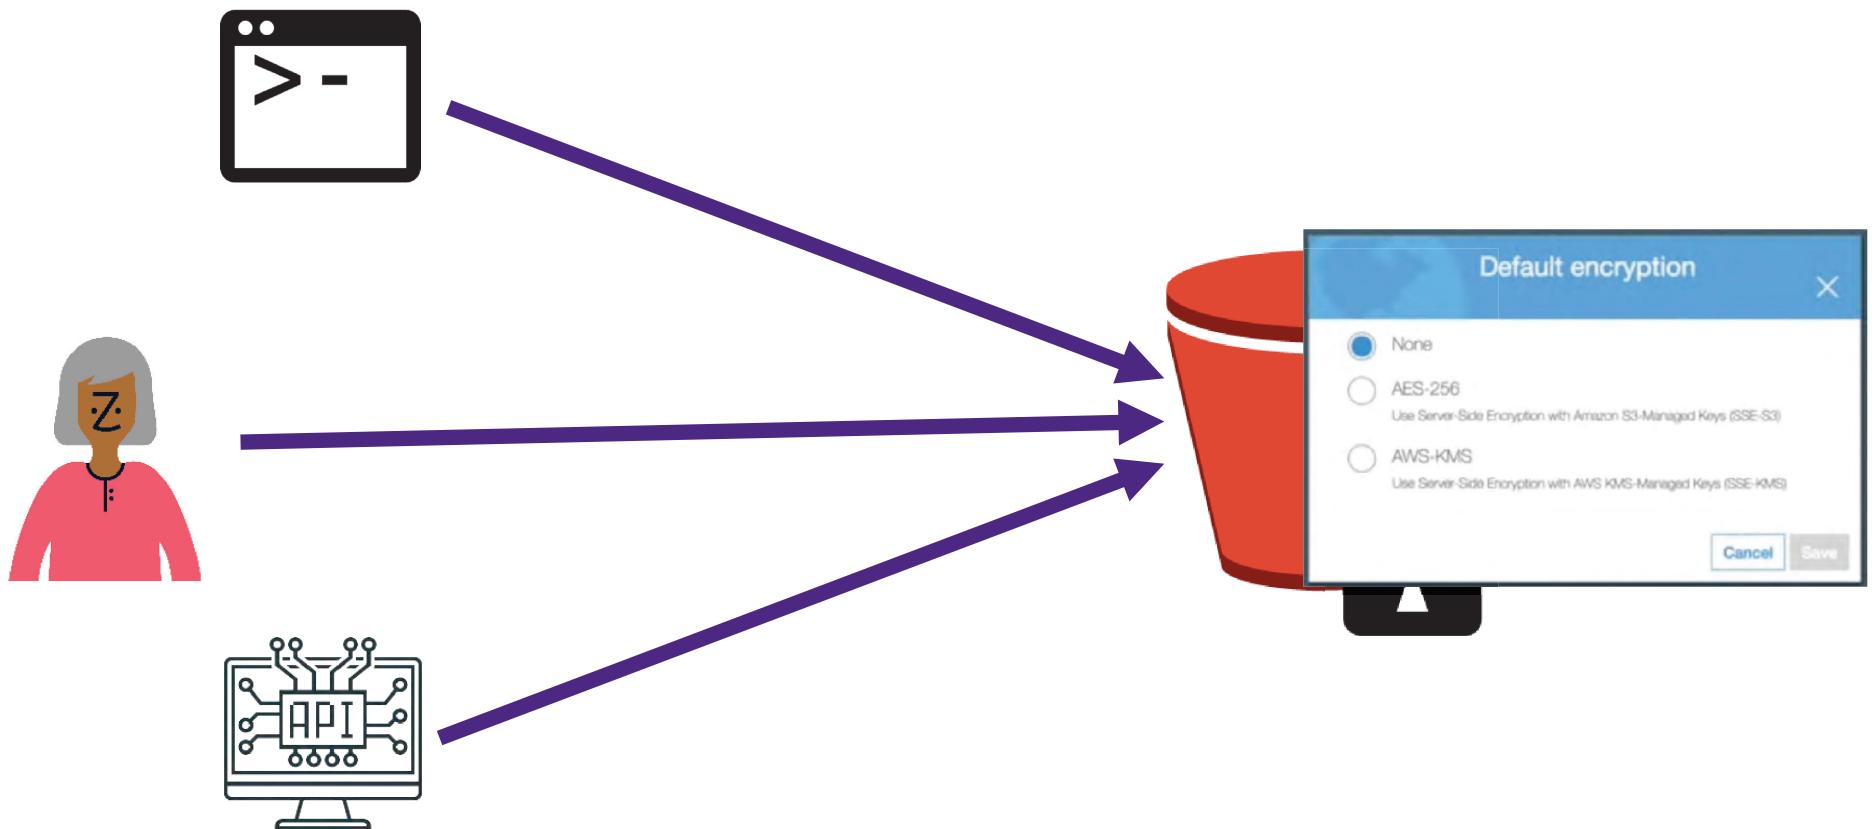
Automated S3 Bucket Encryption

Threat: S3 Bucket Data Exposure

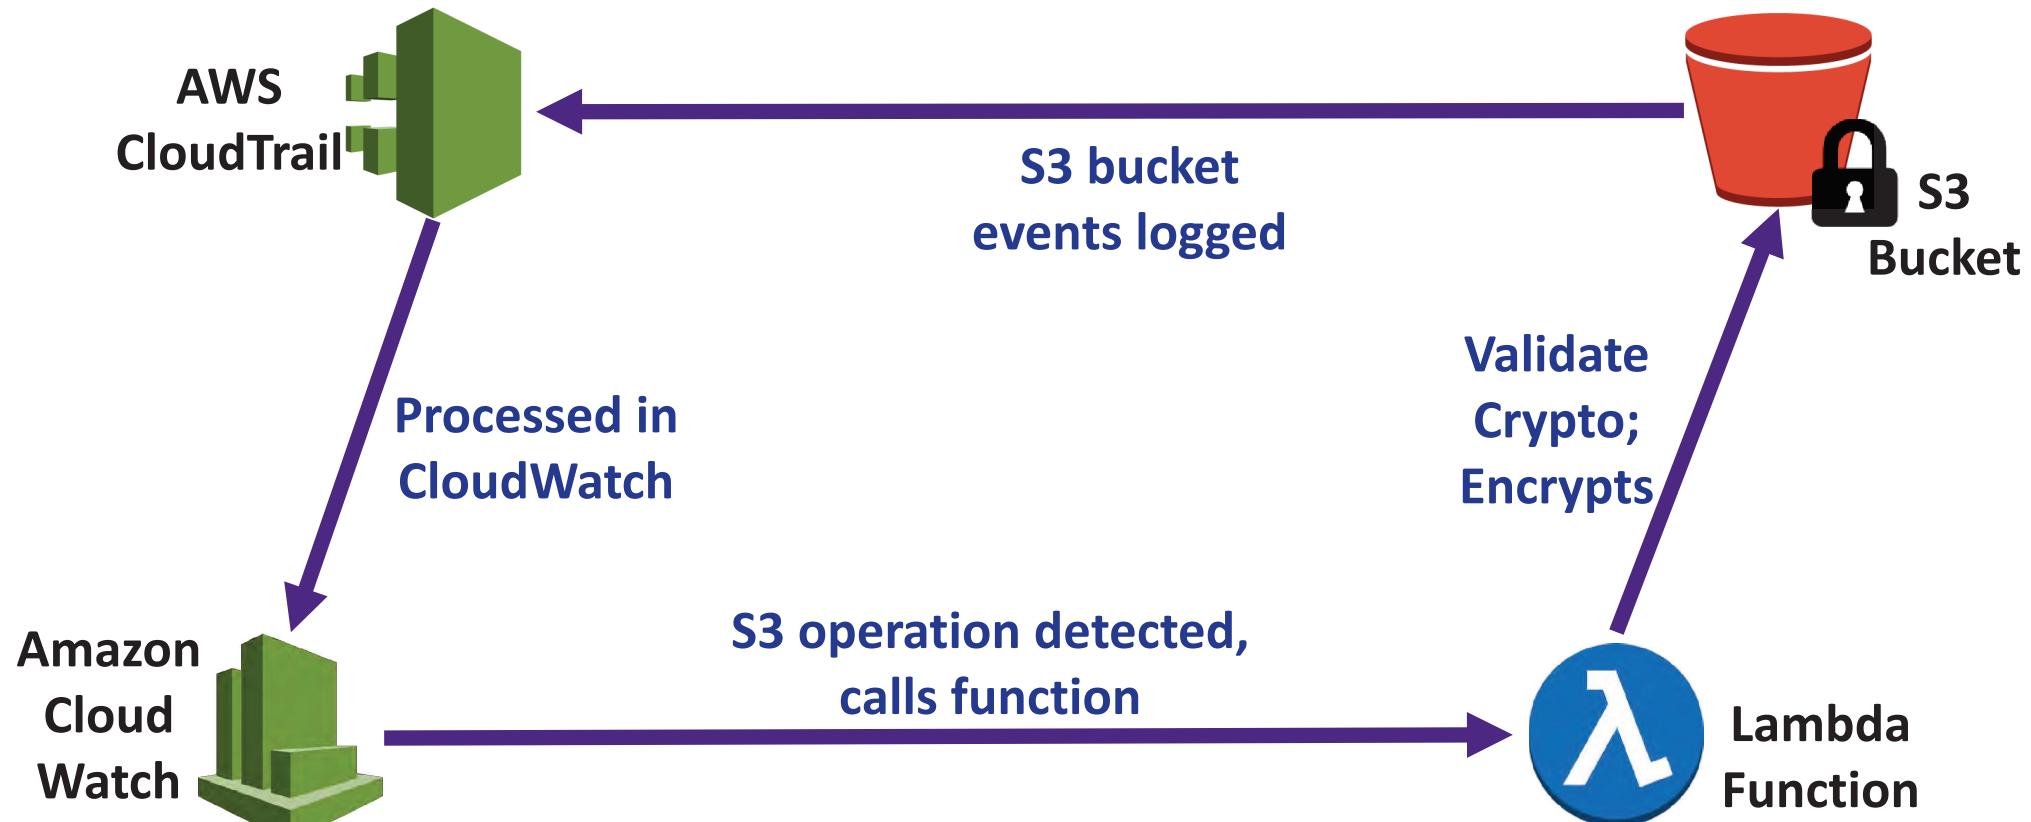
- Pre-cloud, policies and processes enforce encryption at rest
 - We had system build processes
 - Hardened machine images
 - Transparent encryption on RAIDs, SANs, and Servers
- In-cloud, we have these also but data can still be stored unencrypted by default and processes are harder to enforce
- Problem in Cloud:
 - How do we make sure S3 buckets are encrypted
- Solution in Cloud:
 - Automate the bucket encryption when it's lacking



Unencryption Scenarios on S3 Buckets



Automated Encryption of AWS S3 Buckets



RSA®Conference2019

Recorded Demo

Automated S3 Bucket Encryption



CloudWatch Configuration

```
{  
  "source": [ "aws.s3" ],  
  "detail-type": [ "AWS API Call via CloudTrail" ],  
  "detail": {  
    "eventSource": [ "s3.amazonaws.com" ],  
    "eventName": [  
      "CreateBucket",  
      "PutBucketAcl",  
      "PutBucketPolicy",  
      "PutBucketEncryption",  
      "DeleteBucketEncryption"  
    ]  
  }  
}
```

Source Code Highlights

```
s3 = boto3.client('s3')
sns = boto3.client('sns')

def lambda_handler(event, context):
    bucket_name      = event['detail']['requestParameters']['bucketName']
    bucket_creator  = event['detail']['userIdentity']['principalId']
    bucket_user      = bucket_creator.split(":",1)

try:
    currEncrypt = s3.get_bucket_encryption (Bucket=bucket_name)
```

Source Code Highlights

```
except ClientError as e:  
    toEncrypt = s3.put_bucket_encryption (  
        Bucket=bucket_name,  
        ServerSideEncryptionConfiguration = {  
            'Rules':[ {  
                'ApplyServerSideEncryptionByDefault':  
                    { 'SSEAlgorithm': 'AES256' }  
            }]  
        } )  
    response = sns.publish(  
        TopicArn= 'arn:aws:sns:us-east-1: ...  
                  000012345678:unEncryptedS3BucketCreated',  
        Message= 'Encryption Removed by '+ bucket_user[1] )  
else:  
    return currEncrypt
```

Rinse & Repeat; Embrace & Extend

- Design pattern can be replicated for other S3 concerns
- Also implemented “Automated Make Buckets Private”
 - An event causes a bucket to become public
 - Our lambda makes it private if not in the whitelist
 - Amazon finally does this now
- File scanners
 - New file; scan for data leakage compliance, steganography, etc.
- Intracloud duplication
 - Copy all uploads to another cloud provider

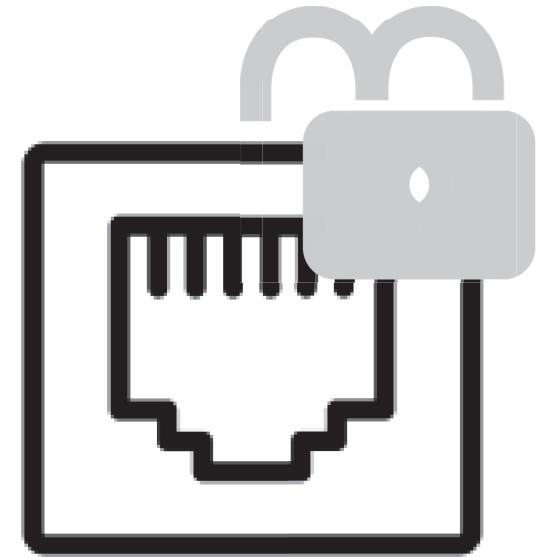
RSA®Conference2019

Sensitive Ports open to Internet

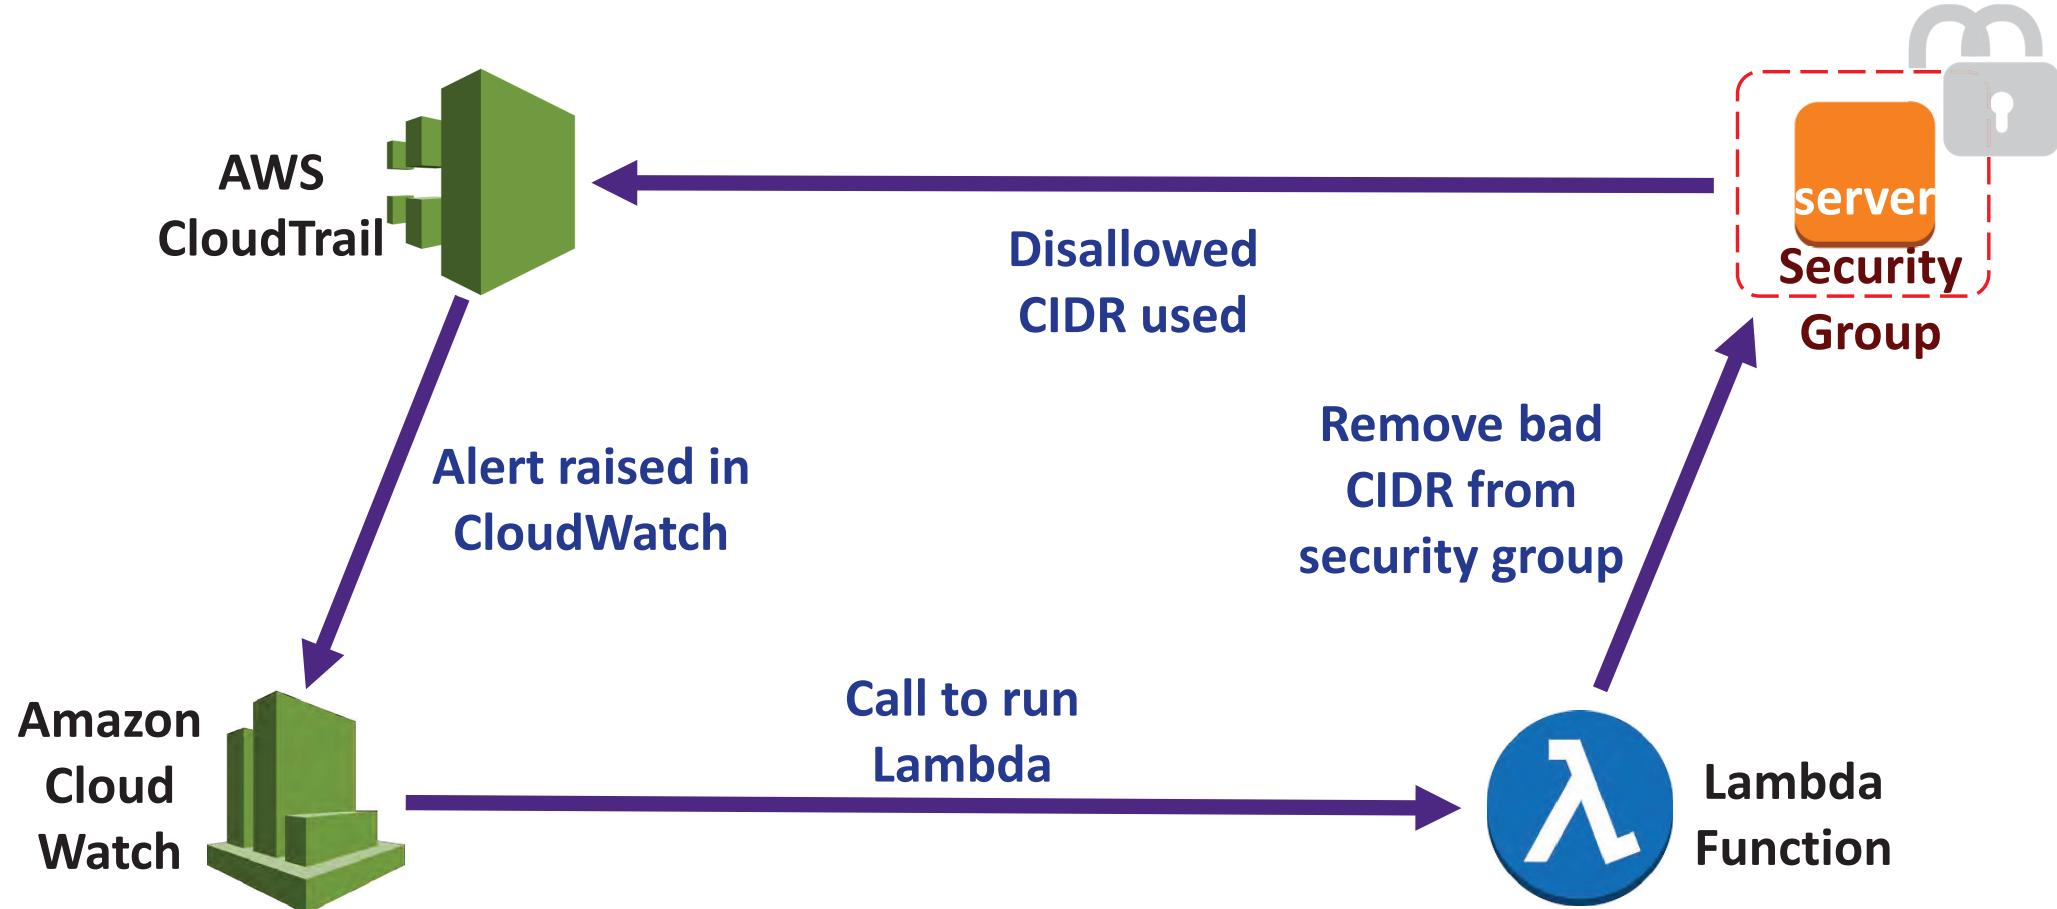
Automated AWS Security Group/VPC Updates

Threat : Ports open to the Internet

- Pre-cloud, we relied on Firewalls and Network security.
 - Deny by default
 - Network change management
- In-cloud, anyone can spin up a system
 - And chose a new or existing security group
- Problems in Cloud:
 - Someone opens sensitive ports to the entire Internet
- Solutions in Cloud:
 - Lock the ports down by auto-correcting the security group



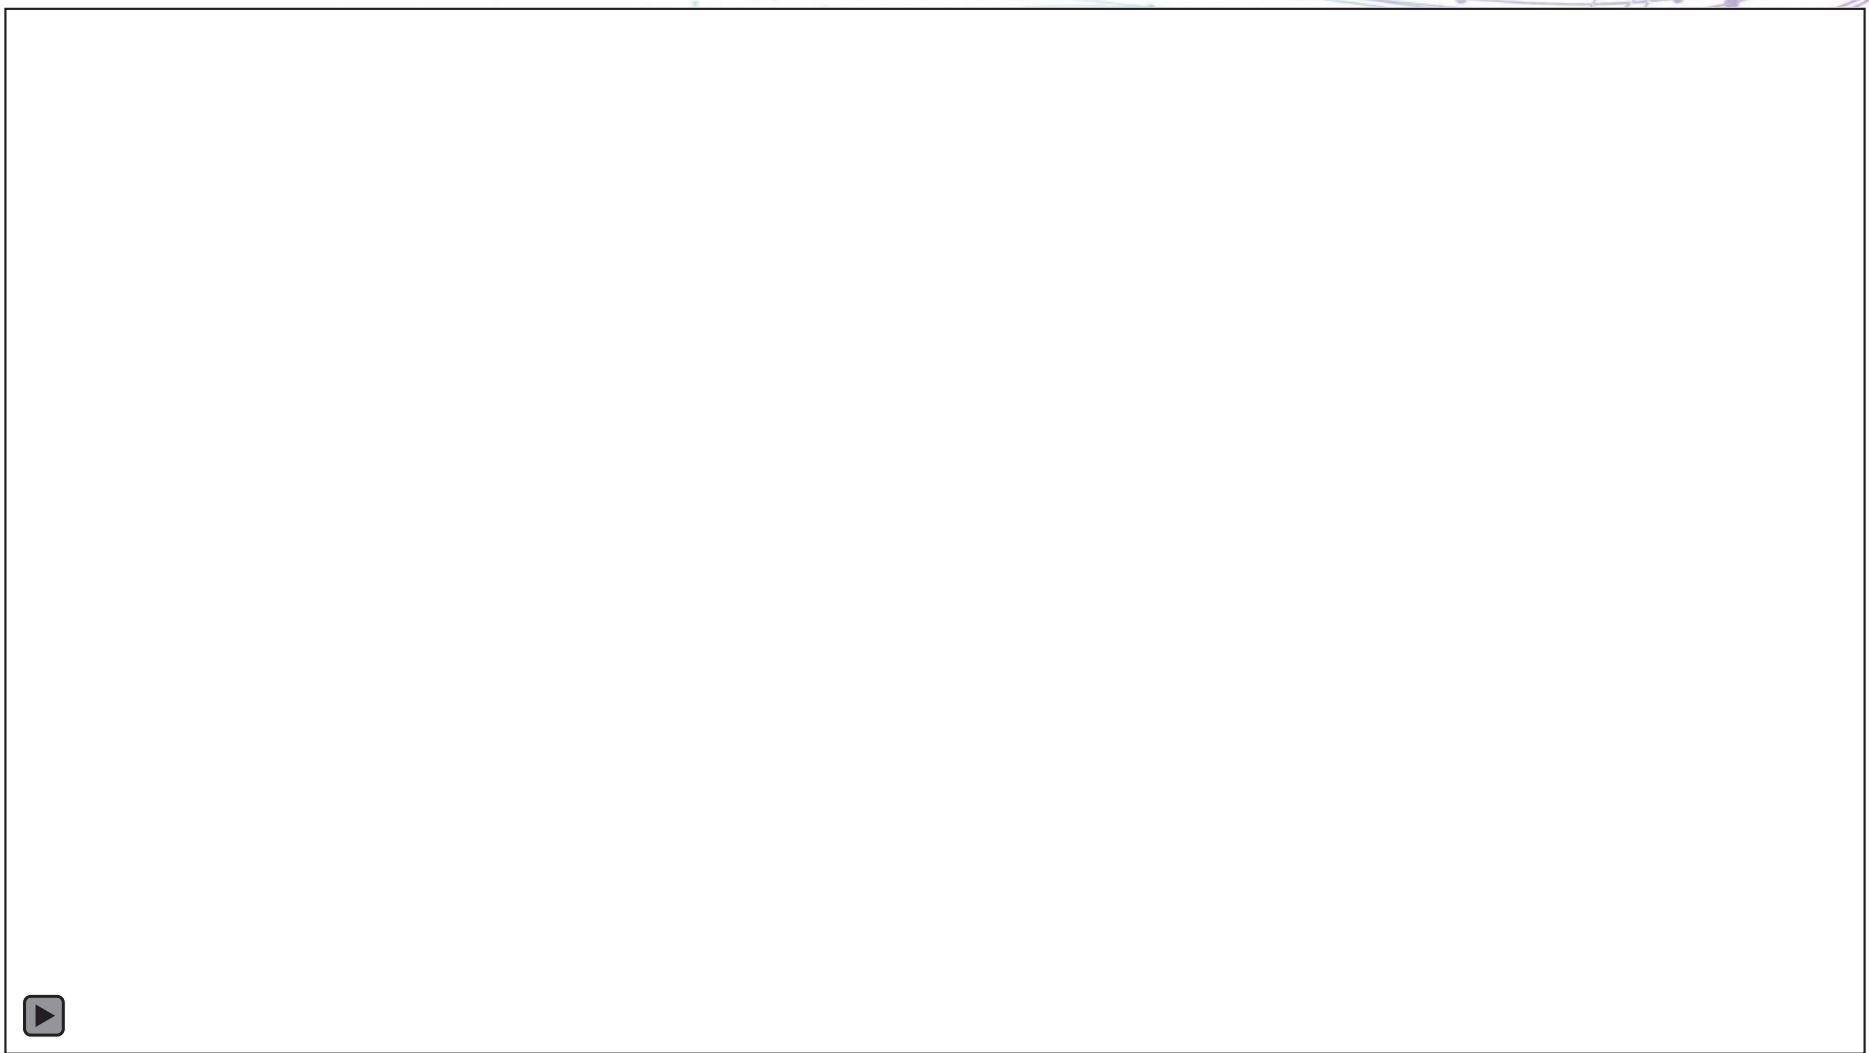
Automated Port Updates in AWS



RSA®Conference2019

Recorded Demo

Automated Ports, No More -1



CloudWatch Configuration

```
{  
  "source":      [ "aws.cloudtrail" ],  
  "detail-type": [ "AWS API Call via CloudTrail" ],  
  "detail": {  
    "eventSource": [ "ec2.amazonaws.com" ],  
    "eventName":   [  
      "CreateSecurityGroup",  
      "AuthorizeSecurityGroupIngress"  
    ]  
  }  
}
```

Code: Initial Sanity Checks

```
def lambda_handler(event, context):
    source = event['detail']['eventSource']
    if source != "ec2.amazonaws.com":
        return      # wrong caller, just silently return

    allowed_event_list = [
        'CreateSecurityGroup',
        'AuthorizeSecurityGroupIngress'
    ]

    event_name = event['detail']['eventName']

    if not(event_name in allowed_event_list):
        return      # wrong caller, just silently return
```

Code: GetSecurityGroupId

```
resp = event['detail']['responseElements']

if (resp["_return"] != True):
    return      # event not successful, so ignore it

SG_id = 'invalid'
if event_name == 'CreateSecurityGroup':
    SG_id = resp['groupId']
elif event_name == 'AuthorizeSecurityGroupIngress':
    SG_id = event['detail']['requestParameters']['groupId']
```

Code: Group Sensitivity Training

```
ec2 = boto3.resource('ec2')
security_group = ec2.SecurityGroup(SG_id)
sensitive_ports = [ 22, 3389, 54321 ]
ingress_list = security_group.ip_permissions
for perm in ingress_list:
    fromport=0 ; toport=0 ; ipprot=0 ; sensitive=False
for each permission in ingress_list:
    if permission.port is a sensitive_ports
        or permission.portrange contains sensitive_ports
    if permission.protocol == -1    ;# aka "all ports"
then ipprot = perm['IpProtocol']
    if ipprot == "1":
        sensitive = True
    if fromport > 0:
        for p in sensitive_ports:
            if fromport <= p and p <= toport:
                sensitive = True
```

Code: Test for Zero CIDR

```
if sensitive:
    for r in perm['IpRanges']:
        if r['CidrIp'] == "0.0.0.0/0":
            print("Removing Ingress Rule: ", json.dumps(perm))
            try:
                security_group.revoke_ingress(
                    CidrIp      = r['CidrIp'],
                    IpProtocol = perm['IpProtocol'],
                    FromPort   = fromport,
                    ToPort     = toport      )
            except ClientError as e:
                print("Error: ", e)
                raise
```

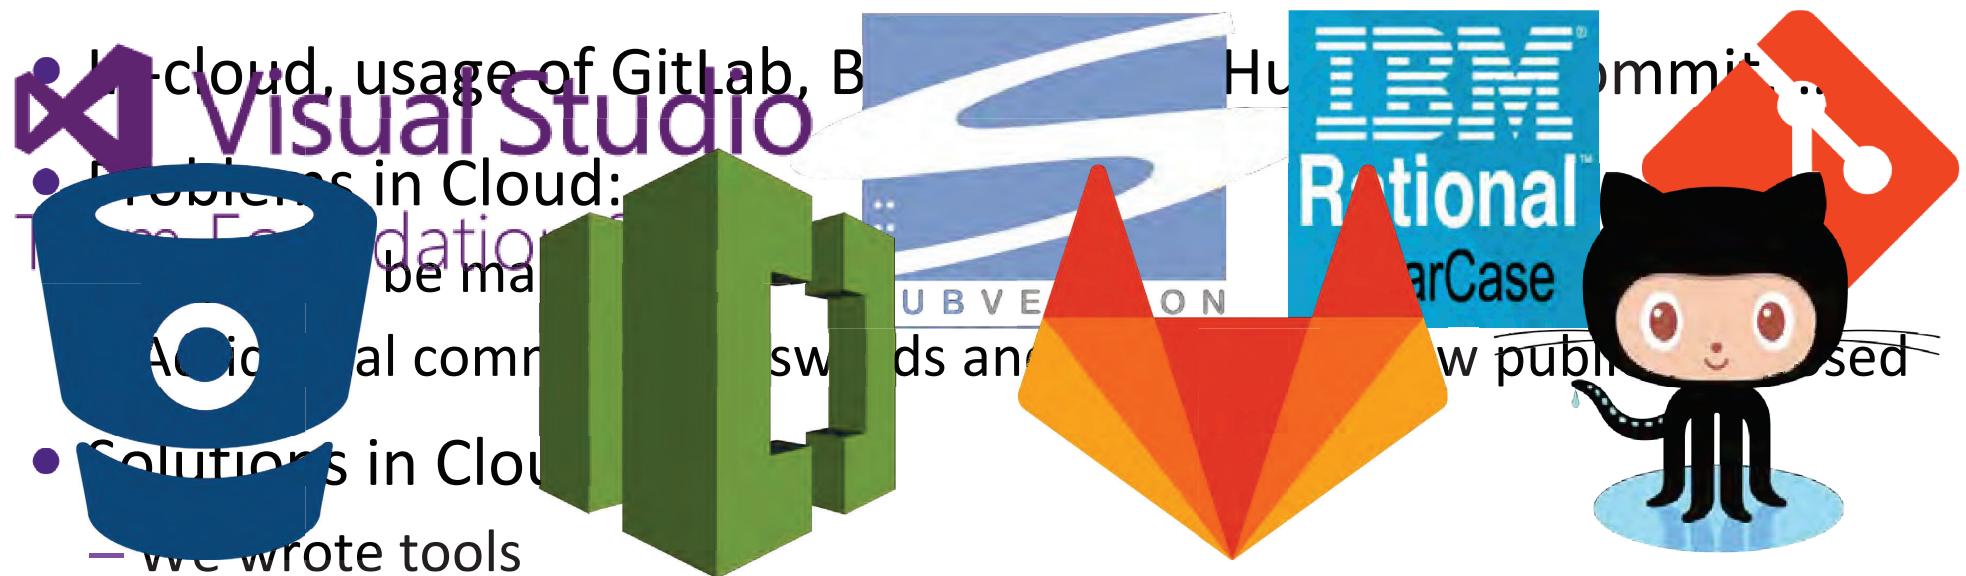
RSA®Conference2019

Source Code Repositories Exposed on Internet

Automate repo protection on GitHub

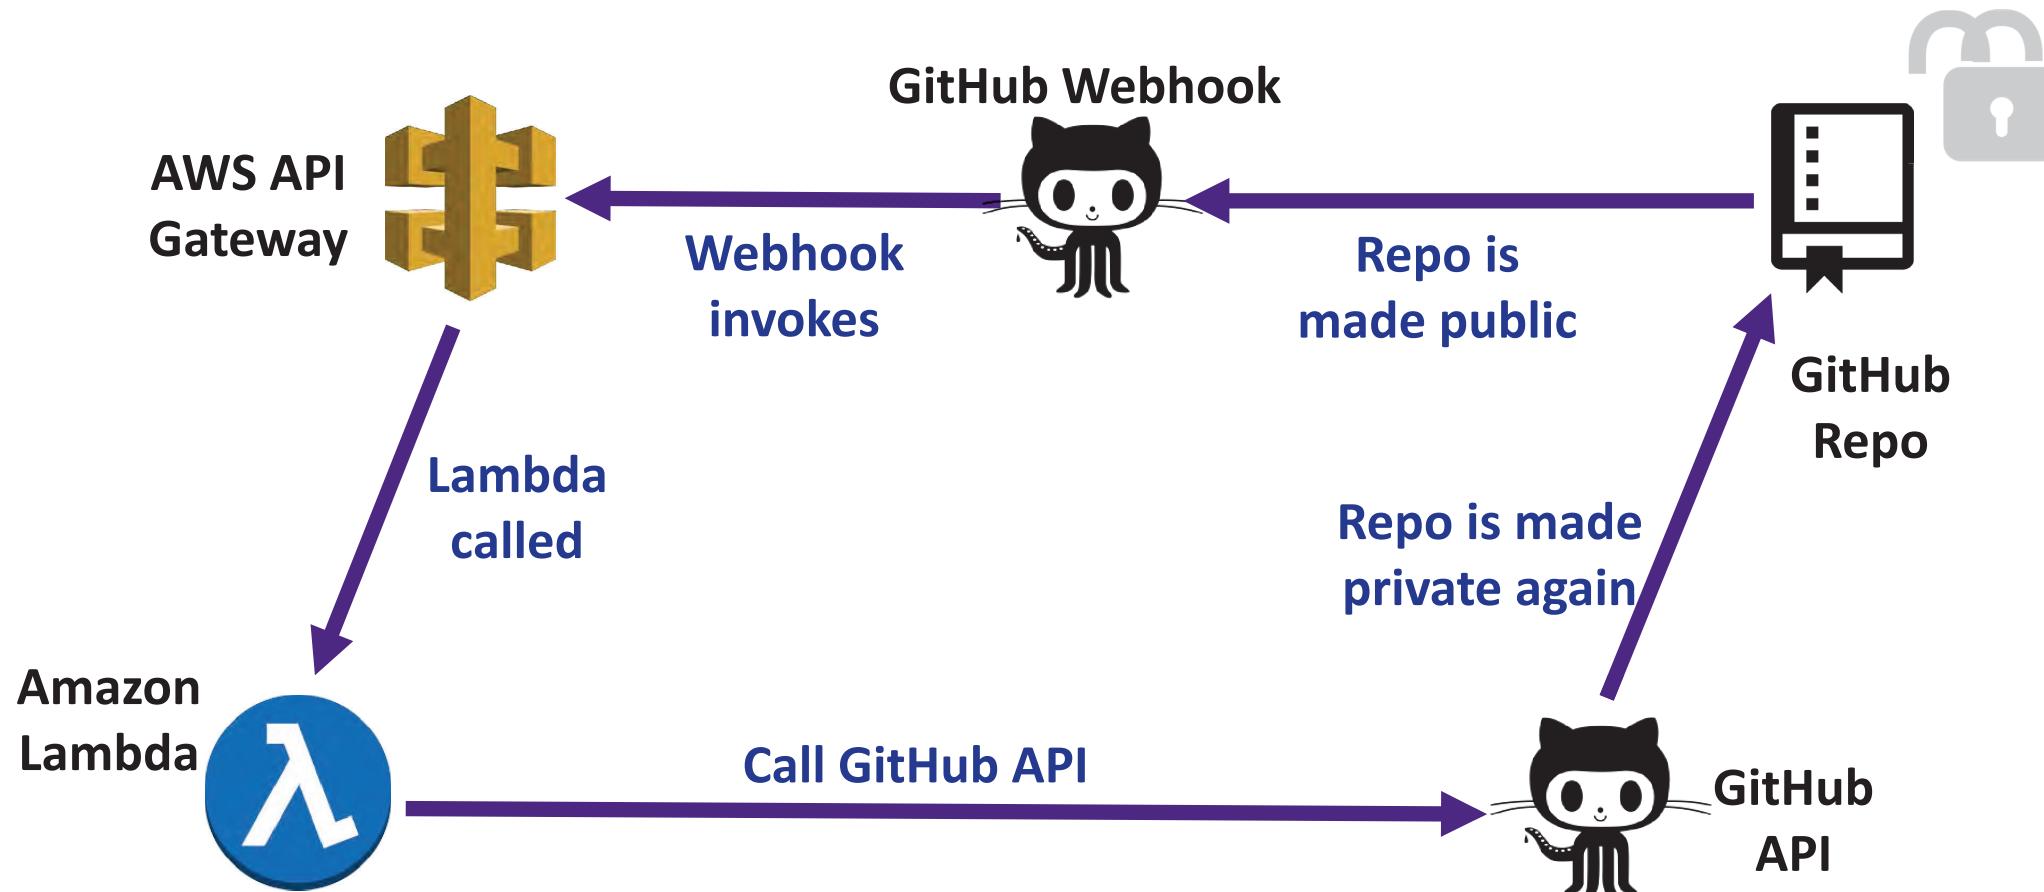
Threat: Public Code Repositories

- Pre-cloud, on-prem code repos, such as TFS, Clearcase, git, ...



- Solutions in Cloud
 - we wrote tools

Public to Private Flow



RSA®Conference2019

Live Demo (with Recorded Backup)

Automate repo protection on GitHub



GitHub WebHook Configuration

```
async function add_hook(organization) {  
  var config = {  
    url: `https://${awshost}/stagename/webhook`,  
    content_type: "json",  
    secret: decrypted_secret_key  
  };  
  const hook = await octokit.orgs.createHook( {  
    org: organization,  
    name: 'web',  
    config,  
    events: ['push','public','repository']  })  
}
```

Code Highlights: Spellbound Invocations

```
const Octokit = require('@octokit/rest');

async function list_topic_and_action(){
  const topics =
    await octokit.repos.listTopics({
      owner: `${body.organization.login}`,
      repo: `${body.repository.name}`,
      headers: {
        accept: 'application/vnd.github+json'
      }
    });
}
```

Code Highlights: Some Privacy Please

```
if ( body.repository.private === false &&
    topics.data.names.includes('open-source')
    === false
) {
  await octokit.repos.update({
    owner: `${body.organization.login}`,
    repo: `${body.repository.name}`,
    name: `${body.repository.name}`,
    private: true } )
```

Code Highlights: Complain to Management

```
var params = {  
  Message: `Repo: ${body.repository.full_name} exposed  
            by user:      ${body.sender.login}  
            at time:       ${body.repository.updated_at}  
            Private was:   ${body.repository.private}  
            Other Repo Topics (none if blank):  
              ${topics.data.names}  
  
            This was corrected, repo was made private.`,  
  TopicArn: `arn:aws:sns:reg:012345:github-repo-monitor`,  
};  
AWS.SNS({apiVersion: '2010-03-31'}).publish(params);
```

RSA® Conference 2019

Repository Ownership Auditing

Automate repo ownership reporting on GitHub



Repository Visibility and Auditing

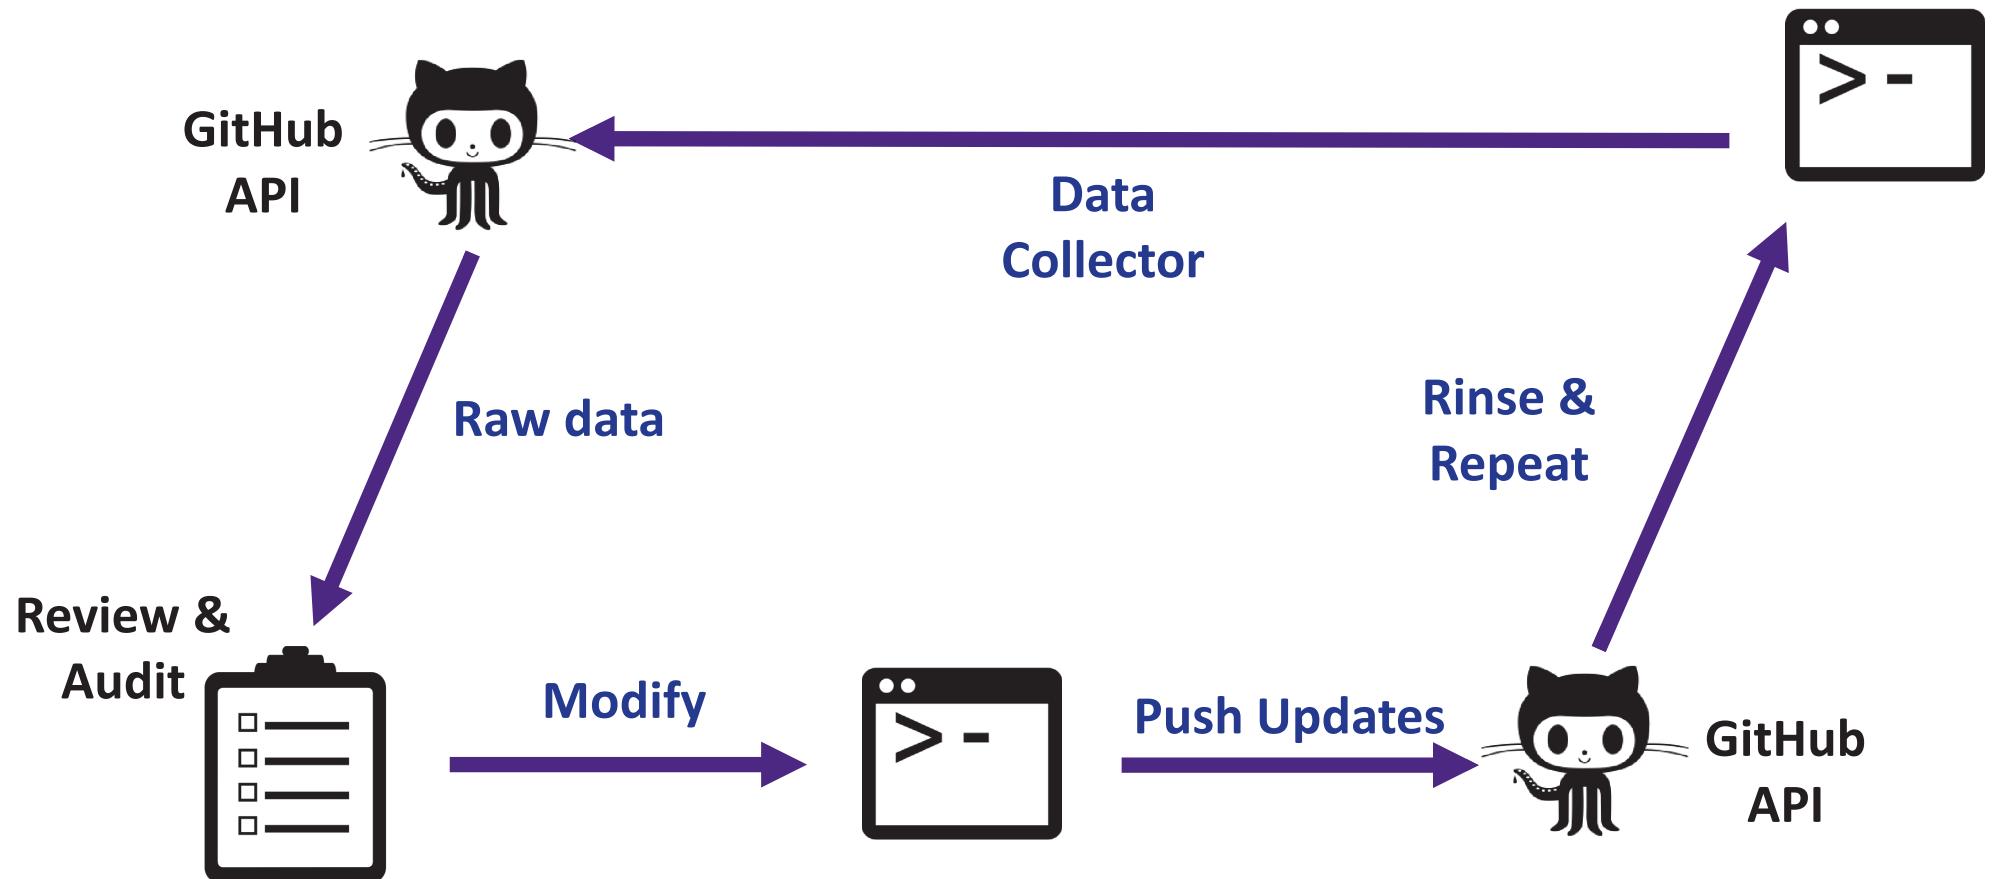
- Pre-cloud, on-prem code repos with built-in auditing, such as Subversion
- In-cloud, we use other things
- Problems in Cloud:
 - Repos are easy to create by anybody, we pay for it
 - Management visibility is hard
- Solutions in Cloud:
 - We wrote tools



Repo Ownership & Responsibility

- Repos at Zocdoc are required to have responsible team
 - Allows both Infosec and Engineering to assign tickets
- We do not use GitHub's team mechanism for repo access control
 - Nearly all repos are owned by the Engineering Organization
 - Developers move between teams and work on many projects
 - Defining ownership and AC using GitHub's mechanism adds overhead
- The ownership is tracked by overloading GitHub topics
 - Topics named “team-NAME” show ownership

Repo Ownership Auditing in GitHub



Excel at Security with Audits

node_id	Name	Description	Archived	Private	Languages	Teams
146348261	ZocSec.GitHub	Repo with GitHub security auditin	FALSE	TRUE	Python JavaScript Shell	infosec/admin
135461720	ZocSec.SecurityAsCo	AWS code from the Zocdoc Inform	FALSE	FALSE	Python	inf

- Basic Info: node_id, Name, Description

- Repo Status: Archived, Private

- Languages

Teams	Topics	Topic_add	Topic_del	Last Committer	Last Commit Date
root Shell	infosec/admin			jay-ball-zocdoc	2019-02-18 16:02:57
infosec/admin	team-infosec open-source			web-flow	2018-10-26 19:02:53

- Teams – GitHub’s Teams and Access Control Permissions
- Topics – All of the topics, including our “teams via topics”
- Update Columns – Topic_Add and Topic_Del
- Last Committer, Last Commit Date

RSA®Conference2019

Recorded Demo

Automate repo ownership

Calling Convention

```
$ python3 github_tasks.py --help  
usage: github_tasks.py [-h] -t TOKEN [-a] [-u UPDATE]
```

This tool allow user to audit Enterprise Repos.

optional arguments:

- h, --help show this help message and exit
- t TOKEN, --token TOKEN GitHub User Private Token
- a, --audit Audit github Repo
- u UPDATE, --update UPDATE Import Topics to Repo

Code Highlights: Startup

```
from github import Github  
from openpyxl import Workbook  
  
if __name__ == "__main__":  
    parser = argparse.ArgumentParser( ... )  
    args = vars(parser.parse_args())  
    user = Github(args["token"])
```

Code Highlights: We Excel at This

```
if args["audit"]:    # audit mode
    wb = Workbook()
    ws = wb.active
    ws.title = "Github Inventory"
    ...
    wb.create_sheet('Repos Deploy keys')
    deploy_keys_sheet = wb['Repos Deploy keys']
    wb.create_sheet('Summary')
    summary_sheet = wb['Summary']

org = user.get_organization("Zocdoc")
```

Code Highlights: Getting Loopy

```
for repo in user.get_user().get_repos():
    languages = ""; teamlist = ""; topics = ""

    for lang in repo.get_languages():
        languages = languages + lang + " "
    for team in repo.get_teams():
        teamlist = teamlist + team.name + "/" +
                  team.permission + " "
    for topic in repo.get_topics():
        topics = topics + " "+ topic
```

Code Highlights: Put a Ring on It

```
try:  
    for commits in repo.get_commits():  
        if commits.committer:  
            github_inventory['k%d' % row] =  
                commits.committer.login  
        if ...endswith("-zocdoc"):  
            human += 1  
    break  
except GithubException as e:  
    github_inventory['k%d' % row] =  
        e.args[1]['message']
```

Code Highlights: More Excel-lent Things

```
github_inventory['a%d' % row] = repo.id
github_inventory['b%d' % row] = repo.name
github_inventory['c%d' % row] = repo.description

github_inventory['d%d' % row] = repo.archived
if repo.archived == True:
    github_inventory['d%d' % row].fill = grassFill

github_inventory['e%d' % row] = repo.private
if repo.private == False:
    github_inventory['e%d' % row].fill = redFill

github_inventory['f%d' % row] = languages
github_inventory['g%d' % row] = teamlist
github_inventory['l%d' % row] = repo.pushed_at
```

Code Highlights: Unloved Repo

```
if topics.find("team-") > 0:  
    github_inventory['h%d' % row] = topics  
else:  
    num_of_topic += 1  
    github_inventory['h%d' % row].fill = redFill  
    github_inventory['h%d' % row] = topics  
    github_inventory['i%d' % row].fill = redFill
```

Code Highlights: Keys and Stats

```
for key in repo.get_keys():
    deploy_keys_sheet[...] = repo.full_name
    deploy_keys_sheet[...] = str(key)
# end of main loop

# Repos with a team & percentage
summary_sheet['a2'] = num_of_topic
summary_sheet['b2'] = num_of_topic/count*100
wb.save(filepath) ;# And save...
```

Code Highlights: Transdermal Topical Applications

```
if args["update"]:  
    for each spreadsheet row:  
        if row.topic_add != "":  
            topics = row.topics.split()  
            topics.extend(row.topic_add)  
            user.get_repo("Zocdoc/%s" % row.repo_name)  
                .replace_topics(topics)  
        if row.topic_del != "":  
            # ... etc.
```

RSA®Conference2019

Ownership & Responsibility of AWS Resources

Determine AWS Resource ownership gaps leveraging the billing statements

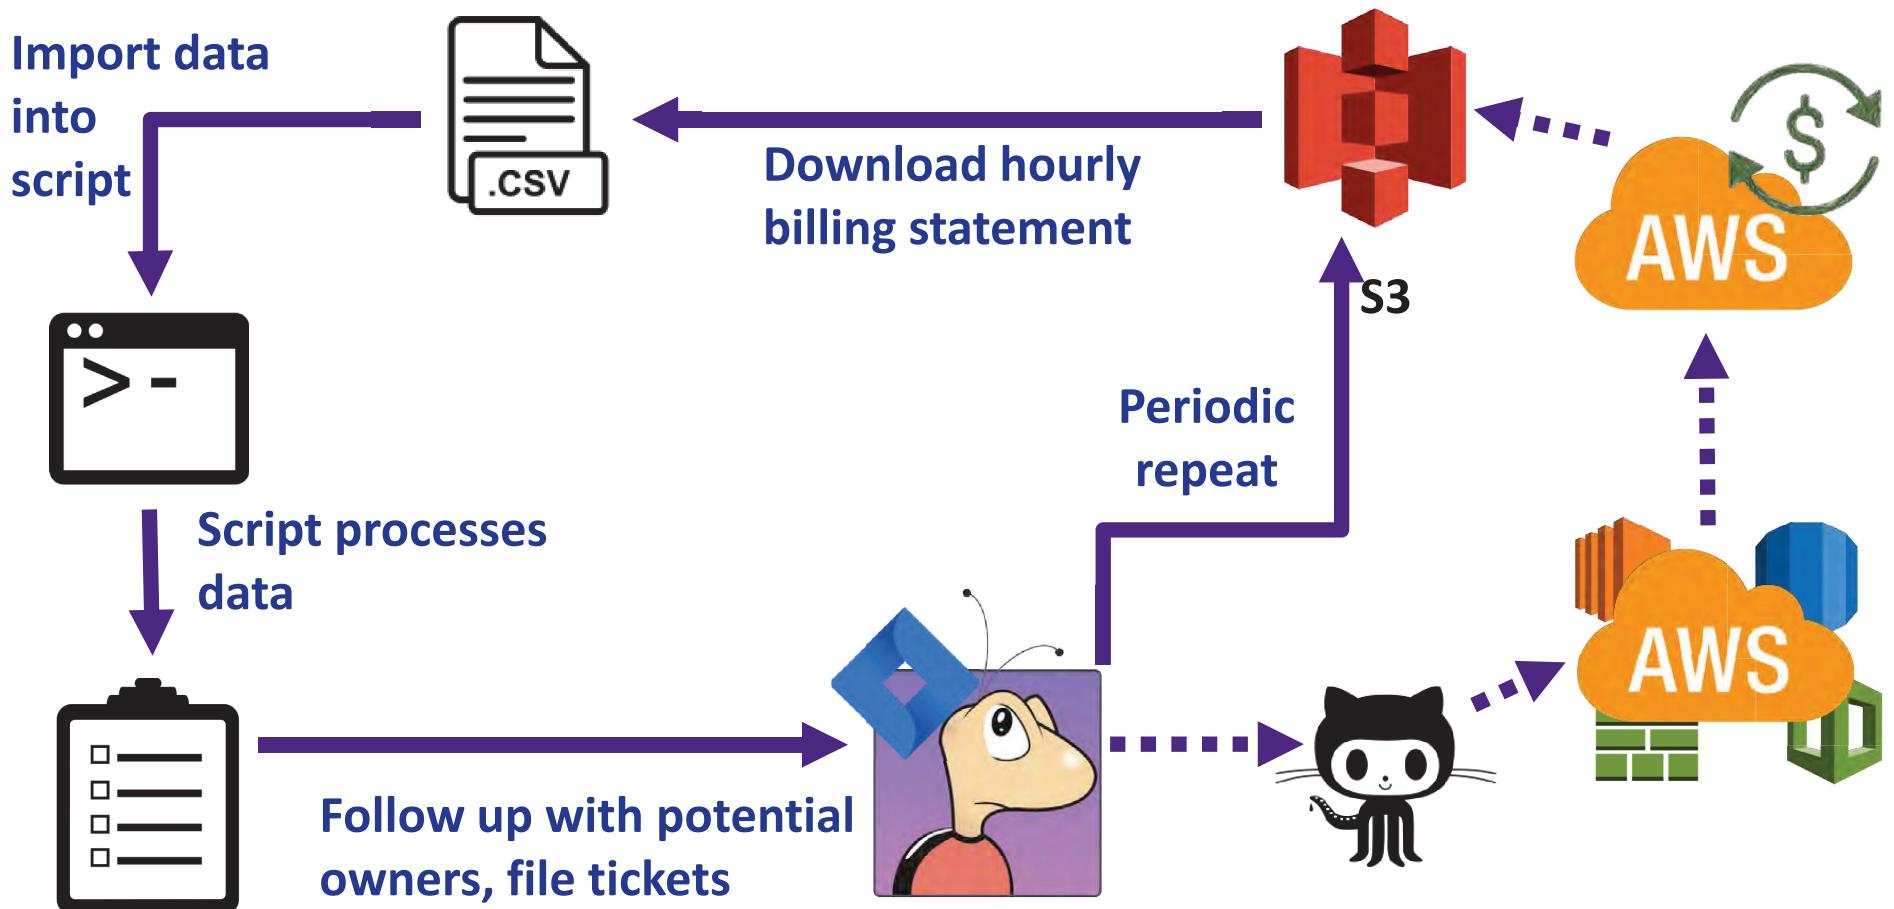
Services in the Cloud

- Pre-cloud, you know who owns servers and storage
 - Since you need hardware, it's easy for IT or Datacenter ops to track
- In-cloud, one can spin up a server or allocate storage without telling anybody
- Problems in Cloud:
 - We're paying for all of this
 - Infosec needs to know ownership during an incident
- Solutions in Cloud:
 - We wrote tools

Assigning Ownership

- Departmental ownership is tracked via AWS tags
- Most deployment is done via Ansible scripts
 - Tags are embedded in script or passed via environment variables
- Infrastructure spun up and torn down as need dictates
- Unlike GitHub, we cannot write easy auto-correction software
 - No point in fixing a tag for an interstitial item
- Owner updates via issue tracker and approved code pulls

AWS Billing and Ownership Flow



RSA®Conference2019

Recorded Demo

Determine AWS Resource ownership gaps via billing statements



Source Code Highlights: I sed a pun, I'm Bourne to do it

```
#!/bin/bash  
input=$1
```

```
# convert csv titles to valid sql column names  
sed -e '1 s/[[:\:\/\]-]/_/g' -e '1 s/$/2/g' $1 \  
      > temp.csv
```

```
sqlite3 < processing.sql
```

Source Code Highlights: Pun Wars II, the Sequel

```
.open cloudhealth-db.sqlite3
.mode csv
.import temp.csv cloudhealthtable

.mode tabs
.header on
.output s3_bucket_owners.txt

select distinct
    lineItem_UsageAccountId,
    lineItem_ResourceId,
    resourceTags_user_Owner
from cloudhealthtable
where
    lineItem_ProductCode = 'AmazonS3' ;
```

Source Code Highlights: Pun Wars III, this is getting silly

```
.output ec2_owners.txt
select distinct
    lineItem_UsageAccountId,
    lineItem_ResourceId,
    resourceTags_user_Owner,
    resourceTags_user_Creator,
    resourceTags_user_Name,
    resourceTags_user_Project
        --- more useful columns...
from cloudhealthtable
where
    product_productFamily = 'Compute Instance' ;
```

RSA®Conference2019

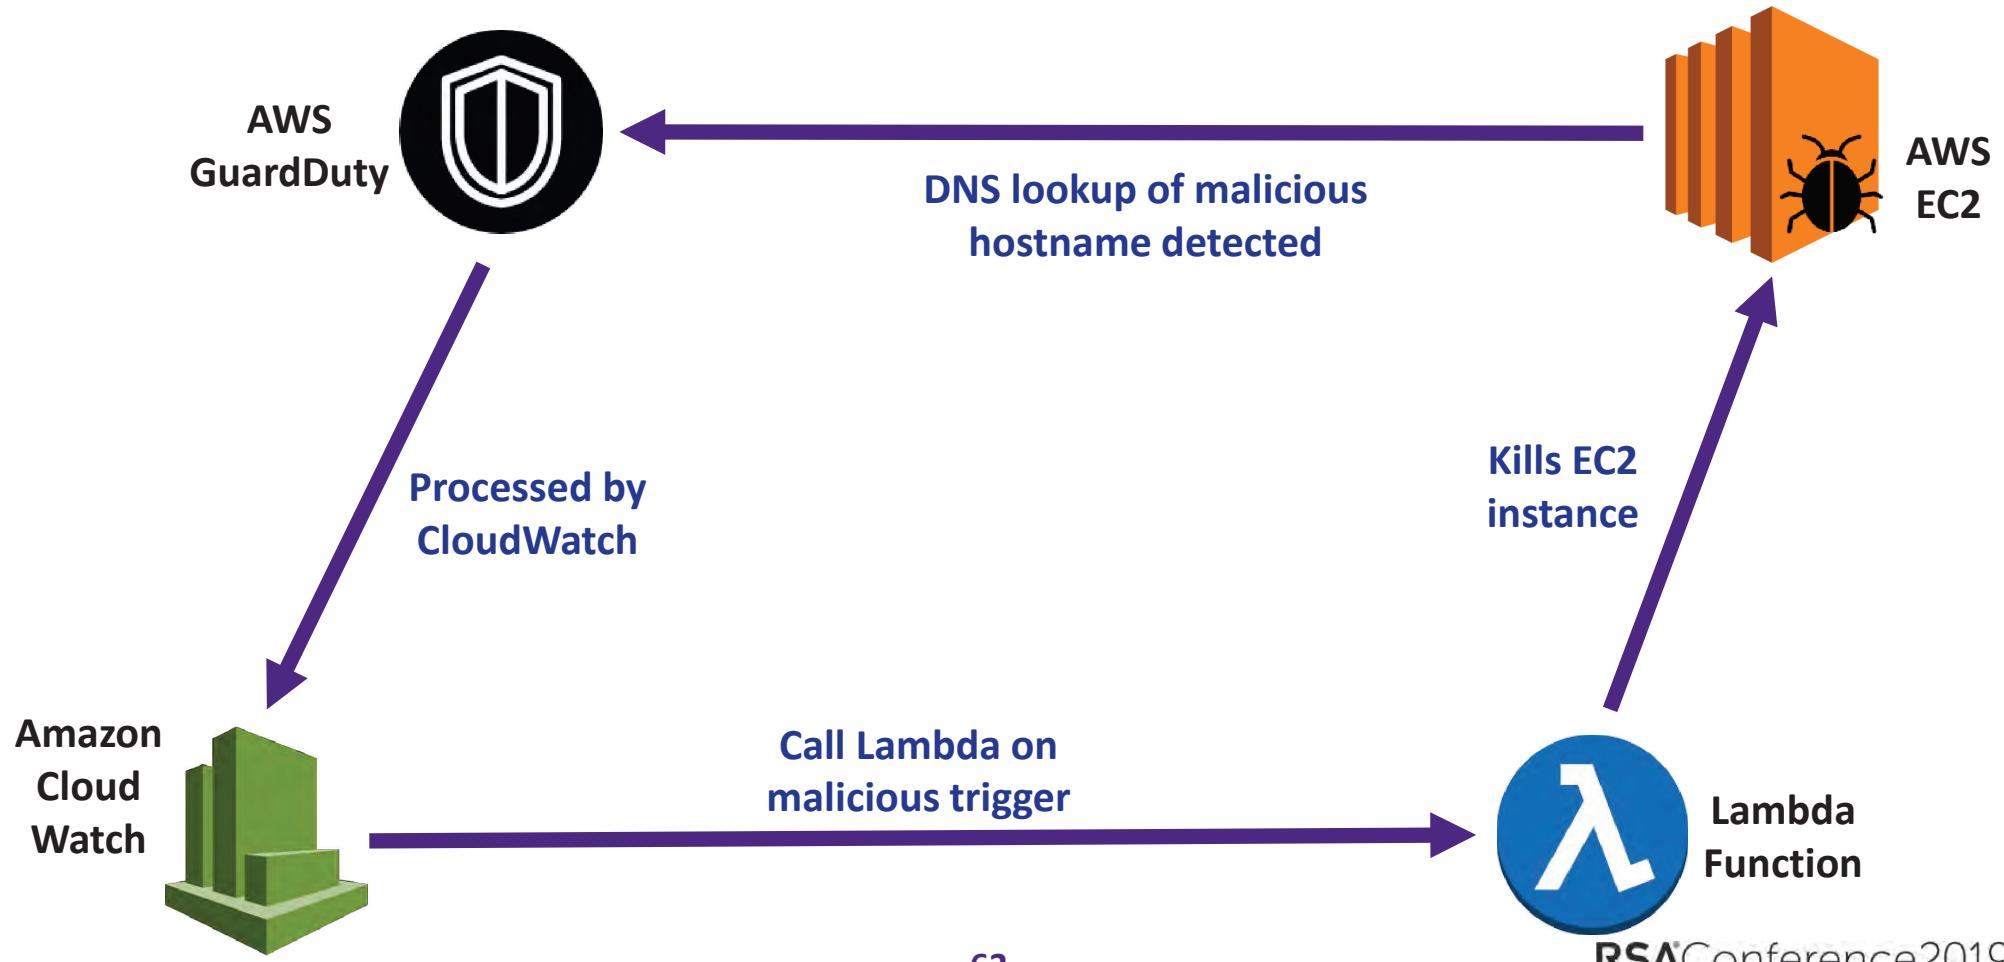
Compromised instance

Just kill it

Compromised systems in the Cloud

- Pre-cloud, leveraged network and security tools or manual processes
- In-cloud, visibility and tools are different
- Problems in Cloud:
 - Visibility
 - Time
 - Scale
- Solutions in Cloud:
 - We wrote a tool

Automated Detection & Response



RSA®Conference2019

Recorded Demo

Et tu, ec2?



Cloud Watch Configuration

```
{  
  "source":      [ "aws.guardduty"      ],  
  "detail-type": [ "GuardDuty Finding" ],  
  "detail": {  
    "type": [  
      "Backdoor:EC2/C&CActivity.B!DNS"  
    ]  
  }  
}
```

Source Code Highlights: Insanity Checks

```
def lambda_handler(event, context):
    source = event['source']
    if source != 'aws.guardduty':
        return ;# wrong caller, silently return

    detail_type = event['detail-type']
    if detail_type != 'GuardDuty Finding':
        return ;# wrong caller, silently return

    event_type = event['detail']['type']
    bad_event_list = [ 'Backdoor:EC2/C&CActivity.B!DNS' ]
    if not(event_type in bad_event_list):
        return ;# wrong event, silently return
```

Source Code Highlights: More Insanity

```
instance_id =  
    event['detail']['resource']['instanceDetails']['instanceId']  
  
# only shutdown certain tags:  
taglist =  
    event['detail']['resource']['instanceDetails']['tags']  
  
for nvp in taglist:  
    if "security_guillotine" == nvp['key'].lower():  
        runme=True  
  
if not(runme):  
    # We are not part of auto-shutdown EC2 group  
    return
```

Source Code Highlights: Euthanize It

```
print("Shutting down instance ", instance_id)

ec2 = boto3.client('ec2')
try:
    response = ec2.stop_instances(
        InstanceIds=[ instance_id ] )

except ClientError as e:
    print('Error', e)
```

Eventful Situations

- GuardDuty alerts on many other events types
 - from mundane to pwned
- We can kill the instance in these critical situations
- Update CloudWatch config and `bad_event_list` variable

'Backdoor:EC2/XORDDOS',
'Backdoor:EC2/Spambot',
'Backdoor:EC2/C&CActivity.B!DNS',
'CryptoCurrency:EC2/BitcoinTool.B!DNS',
'Trojan:EC2/BlackholeTraffic',
'Trojan:EC2/DropPoint',
'Trojan:EC2/BlackholeTraffic!DNS',
'Trojan:EC2/DriveBySourceTraffic!DNS',
'Trojan:EC2/DropPoint!DNS',
'Trojan:EC2/DGADomainRequest.B',
'Trojan:EC2/DGADomainRequest.C!DNS',
'Trojan:EC2/DNSDataExfiltration',
'Trojan:EC2/PhishingDomainRequest!DNS'

RSA® Conference 2019

Getting buy-in

How do you convince management

Cases to build buy-in

- S3 Bucket Encryption – Compliance / Fear
- VPC Port Inspection and Closure – Parity to non-cloud model
- Public Repos in GitHub – Data leakage and exposure
- GitHub Repo ownership – Auditability and closing vulns
- AWS Ownership – Auditability, billing, and monitoring
- Instance Termination – Incident Response, Containment, and Eradication

Apply What You Have Learned Today

- Determine where there is less cloud control versus on-prem
- How can you close this gap
- We give some tools, you can write others
- Small process improvements add up especially at scale

RSA® Conference 2019

Code Available

*Bonus: GitHub remediation
code drops today*

github.com/Zocdoc/ZocSec.SecurityAsCode.AWS

github.com/Zocdoc/ZocSec.SecurityAsCode.GitHub





Special thanks to

Mary Gu

**Information Security
Program Manager**

Linked-In: [mary-gu-ua](#)

Gary Tsai

**Application Security
Engineer**

**Linked-In: [garymalaysia](#)
GitHub: [garymalaysia](#)**



RSA® Conference 2019

San Francisco | March 4–8 | Moscone Center



SESSION ID: CSV-W10

Thank you!

Jay Ball

Head of Application Security
Zocdoc
Linked-In: veggiesspam
Twitter: @veggiesspam
Security Blog: veggiesspam.com
GitHub: veggiesspam



Josh Barons

Head of Information Security
Zocdoc
Linked-In: jbarons
GitHub: CactusSec

#RSAC