

**RSA**® Conference 2020

San Francisco | February 24 – 28 | Moscone Center

HUMAN  
ELEMENT

SESSION ID: HT-R07

# Hunting Linux Malware for Fun and Flags



**Marc-Etienne M.Léveillé**

Senior Malware Researcher

ESET

@marc\_etienne\_

#RSAC

# About this presentation

- This presentation is an introduction to Linux malware analysis and incident response
  - Using commonly available tools
- There is a whole sandbox available where you can test your skills
  - This is where the fun and flags are!
  - Real-world scenarios
    - Trainees are granted root access
    - Use real (defused) malware
    - Real network interactions

# Workshop overview

Hunting Linux Malware for Fun and Flags  
RSAC 2020 Edition

**Session information**

You are assigned **sandbox03-299**

Keep your voucher code ( [REDACTED] ). If you need to connect again or connect from another browser, it can be reused to get to this page.

Status of sandbox: **Running**

| Download OpenVPN configuration file | Slides | Cheat sheet

**Flags**

flag{1234} or 10.3.4.5	Submit
Server 1	0 / 7
Server 2	0 / 9
Server 3	0 / 11

How to connect (Windows) | How to connect (Linux, OS X)

```
% ssh -lroot server1.compromised.network
Warning: Permanently added '10.47.30.42' (ECDSA) to the list of known hosts.
root@10.47.30.42's password:
Welcome to Ubuntu 16.04.6 LTS (GNU/Linux 4.15.0-74-generic x86_64)

| Hunting Linux Malware for Fun and Flags
| Server 1
| https://compromised.network

Last login: Thu Feb  6 20:53:56 2020 from 10.66.0.2
root@trk1-web:~# ps aux
USER      PID %CPU %MEM    VSZ   RSS TTY      STAT START  TIME COMMAND
root        1  0.0  0.1  37180  5300 ?        Ss   20:32  0:00 /sbin/init
root       42  0.0  0.2  35272  8856 ?        Ss   20:32  0:02 /lib/systemd/systemd
root       58  0.0  0.0  41592  3268 ?        Ss   20:32  0:00 /lib/systemd/systemd
syslog    138  0.0  0.0 182660  3204 ?        Ssl  20:32  0:00 /usr/sbin/rsyslogd
root      140  0.0  0.0  29008  3008 ?        Ss   20:32  0:00 /usr/sbin/cron -f
root      257  0.0  0.0  15752  2304 console  Ss+  20:32  0:00 /sbin/agetty --noclear
root      261  0.0  0.1  65512  6176 ?        Ss   20:32  0:00 /usr/sbin/sshd -D
mysql     271  0.0  5.1 1439284 209148 ?        Ssl  20:32  0:01 /usr/sbin/mysqld
root      330  0.0  0.6 298604 25996 ?        Ss   20:32  0:00 /usr/sbin/apache2
```

You only need a web browser, an OpenVPN client and an SSH client to start hunting



**Why this presentation and  
workshop?**

# Dissecting Linux/Moose

The Analysis of a Linux Router-based Worm  
Hungry for Social Networks

Olivier Bilodeau  
& Thomas Dupuy

May 2015



## Unboxing Linux/Mumblehard

Muttering spam from your servers

Marc-Etienne M.Léveillé

April 2015

# OPERATION WINDIGO

The vivisection of a large Linux server-side  
credential stealing malware campaign



ESET Research White papers // December 2018  
Romain Dumont, Marc-Etienne M.Léveillé, Hugo Porcher

# THE DARK SIDE OF THE FORSSHE

A landscape of OpenSSH  
backdoors

# Why malware on Linux servers?

- Servers have a lot of bandwidth
- Servers have a high uptime
- As a result, they make good targets for
  - Sending spam
  - Reverse proxy
  - Open proxy
  - Traffic redirection
  - Hosting services (e.g. DNS) and web pages (e.g. phishing)

# Why care?

- Bad IP reputation
  - Prevents sending legitimate email messages
- Slows down legitimate software and services
- Servers often host the most critical data in the enterprise<sup>1</sup>
- Risk of data exfiltration
  - Passwords
  - E-mail addresses
  - Credit card numbers
  - Etc.

<sup>1</sup> Jon Amato, Mario de Boer, *Gartner Technical Professional Advice – Solution Criteria for Endpoint Protection Platforms*, 2020-01-09

# Why understand them?

- If you don't find out how you were compromised, it might come back by the same door
- If you don't clean everything, it might come back using another backdoor
- Understand what is at risk
- Explain the behavior of very sneaky malware



# Incident response artifacts

# Artifacts

- Filesystem
  - Logs
  - Malware persistence (if any)
- Memory
  - Process memory and state
  - Kernel memory
- Network
  - Configuration
  - Packet capture (in-band and out-of-band)



# Filesystem

# Common file metadata

- Name
- Size
- Type
  - Regular file
  - Directory
  - Symbolic link
  - Special (device)
- Owner
  - User
  - Group
- Access rights
  - Read, write and execute
  - Owner, group and others
- Timestamps
  - Access
  - Last modification
  - Last metadata modification
  - *Creation date*

# Basic filesystem

- Finding new files
- `ls -alt | head`
  - List files that were recently modified in the current directory

```
# ls -lat | head
total 44
-rw----- 1 root root 4322 Oct 29 11:21 .bash_history
drwx----- 1 root root    22 Oct 28 23:52 .aptitude
-rw----- 1 root root     81 Oct 28 22:44 .lessht
drwx----- 1 root root   240 Oct 28 21:59 .
-rw----- 1 root root 5726 Oct 28 21:59 .viminfo
```

# Basic filesystem

- **stat \$FILE**
  - Full file details



All timestamps can  
be tampered with!

```
# stat .viminfo
  File: '.viminfo'
  Size: 5726          Blocks: 16          IO Block: 4096
regular file
Device: 11h/17d  Inode: 63052      Links: 1
Access: (0600/-rw-----)  Uid: (0/root)  Gid: (0/root)
Access: 2015-10-28 21:59:36.283368614 -0400
Modify: 2015-10-28 21:59:36.283368614 -0400
Change: 2015-10-28 21:59:36.283368614 -0400
 Birth: -
```

# Basic filesystem

- **find / -newermt 2019-10-28**
  - Find files that were modified after October 28<sup>th</sup>
  - Based on the same metadata that can be tampered with

```
# find /home/james -newermt 2020-02-01
/home/james
/home/james/wwwroot/index.php
/home/james/wwwroot/static/css/plugins/isimg/css.php
/home/james/.lessht
/home/james/.bash_history
```

# Basic filesystem

- **file \$FILE**
  - Identify file type

```
# file .viminfo
ELF 64-bit LSB executable, x86-64, version 1 (SYSV),
dynamically linked, interpreter /lib64/ld-linux-x86-
64.so.2, for GNU/Linux 2.6.32, stripped
```

# Package integrity

- **debsums**
  - Dpkg-based distributions (Debian, Ubuntu)
- **rpm -Va**
  - RPM-based distributions (RHEL, CentOS, Fedora)

# Malicious or not?

```
# rpm --verify keyutils-libs  
(no error)  
# rpm -qi keyutils-libs  
Name        : keyutils-libs  
Version     : 1.4  
Release     : 4.el6  
Install Date: Mon 27 Jan 2014 06:08:43  
Group       : System Environment/Base  
Size        : 59320  
Signature   : RSA/SHA1, Sun 24 Jun 2012 06:18:51, Key ID 21efc4bf71fbfe7b  
URL         : http://people.redhat.com/~dhowells/keyutils/  
Summary     : Key utilities library  
Description :  
This package provides a wrapper library for the key management facility  
system  
calls.  
  
Relocations: (not relocatable)  
Vendor: CentOS  
Build Date: Fri 22 Jun 2012 02:20:38  
Build Host: c6b10.bsys.dev.centos.org  
Source RPM: keyutils-1.4-  
4.el6.src.rpm  
License: GPLv2+ and LGPLv2+
```

# Logs

- /var/log
  - auth.log
  - HTTP logs
  - messages, syslog, etc.
- systemd's journalctl
- auditd log

# Using auditd

- *The Linux audit framework provides an auditing system that reliably collects information about any security-relevant (or non-security-relevant) event on a system.*
- Part of the kernel
  - Must be enabled during kernel compilation or have a loadable kernel module
  - Enabled on most distributions
- Logs system calls and other types of events
- Logs can be sent over the network

# Using auditd

- **auditctl**: Define what you want to log
- **ausearch**: Search in log files
  - Logs are text files, so **grep** and other tools works fine too for this task

```
# auditctl -a exit,always -S execve  
[...]  
# ausearch -m EXECVE  
type=EXECVE msg=audit(1373838239.340:4474200): argc=4  
a0="rm" a1="-f" a2="-f" a3="/tmp/q"  
type=EXECVE msg=audit(1373838239.341:4474201): argc=4  
a0="touch" a1="-r" a2="/etc/ssh/sshd_config"  
a3="/etc/ssh/ssh_config"
```

# Offline filesystem

- If you don't have access to the live system but have an image of the partition.
- Capture: `dd if=/dev/sda3 of=$IMAGE_FILE`
  - Works over SSH too!
- Browse: `mount -o loop,ro $IMAGE_FILE /mnt`



# System memory

# Analyzing a live process

- Identify running processes

- ps auxw
- top, htop

```
# ps auxw
root      7673  0.0  0.0  55164  5292 ?    Ss   Oct19  0:01 /usr/sbin/sshd -D
root      7718  0.0  0.3 200676 127160 ?    Ss   Oct21  2:29 /lib/systemd/systemd
root      7948  0.0  0.0 248844 25032 ?    Ss   Oct20  0:18 /usr/sbin/apache2 -k
webuser   7953  0.0  0.0 141732  4188 ?    S    Oct20  0:06 /usr/sbin/apache2 -k
webuser   5023  6.2  0.0  39764  8936 ?    Ss   Oct28 66:08 /tmp/.ICE-A5BF7
[...]
```

# Analyzing a live process

- List open files and network streams
  - lssof -p \$PID

sshd	3642	sshd	cwd	DIR	0,17	166	256	/
sshd	3642	sshd	rtd	DIR	0,17	166	256	/
sshd	3642	sshd	txt	REG	0,17	787080	2231	/usr/sbi...
sshd	3642	sshd	mem	REG	0,16		36265	/lib/x86...
sshd	3642	sshd	mem	REG	0,16		36267	/lib/x86...
[...]								
sshd	3642	sshd	2u	CHR	1,3	0t0	1028	/dev/null
sshd	3642	sshd	3u	IPv4	146293912 158.69.117.51:ssh->182.100.67.59:41000	0t0	TCP (ESTABLISHED)	
sshd	3642	sshd	4u	unix	0xffff88072c5cbc00	0t0	146289562	socket

# procfs

- *procfs* provides a lot of useful details
- Mounted at /proc
- Contains one directory per process at /proc/\$PID

```
$ ls /proc/3537
attr          cwd      map_files   oom_adj      schedstat    syscall
autogroup     environ   maps        oom_score    sessionid    task
auxv          exe       mem         oom_score_adj setgroups    timers
cgroup         fd        mountinfo  pagemap      smaps        timerslack_ns
clear_refs    fdinfo    mounts     patch_state  smaps_rollup uid_map
cmdline        gid_map   mountstats personality  stack        wchan
comm           io        net         projid_map  stat
coredump_filter limits   ns          root        statm
cpuset         loginuid  numa_maps  sched       status
```

# *procfs exe magic link*

- Find the path of the executed file
  - `ls -l /proc/$PID/exe`
- Retrieve the executable file **even if it was deleted**
  - `cp /proc/$PID/exe malware.elf`

```
# ps aux | grep 25465
web  25465  6.6  0.0  39764   936 ? Ss Oct29 157:52 crond
# ls -l /proc/25465/exe
lrwxrwxrwx 1 web www-data 0 Oct 29 04:09
/proc/25465/exe -> /tmp/.ICE-684c
```

# *procfs environ*

- `/proc/$PID/environ` contains the environment variables of a process separated by null bytes

```
# tr '\0' '\n' < /proc/1179/environ
MAIL_CONFIG=/etc/postfix
MAIL_LOGTAG=postfix
LANG=C
SSH_CONNECTION=10.0.2.2 58505 10.0.2.15 22
GENERATION=1654316
```

# Process stalling

- Stop a process without destroying its resources.
  - `kill -SIGSTOP $PID`
- Resume a process previously stopped with **SIGSTOP**
  - `kill -SIGCONT $PID`

# Process memory dump

- Acquisition
  - `gcore $PID and cp /proc/$PID/exe malware.elf`
- Alternative acquisition tool
  - `memfetch` from <http://1camtuf.coredump.cx>
- Analysis (simple)
  - `strings`
- Analysis (in-depth)
  - `gdb malware.elf $PID.core`

# Kernel memory

- Acquisition
  - VM snapshot
  - LiME (Linux Memory Extractor)
- Analysis
  - Volatility Framework
- Only helpful if kernel is compromised via malicious kernel module (rootkit)

# RSA® Conference 2020

## Network

# Network configuration

- Dump iptables rules
  - `iptables-save`
  - `ip6tables-save`

```
# iptables-save
[...]
-A POSTROUTING -s 0.0.0.0/0 --dport 8080 -o eth0 -j
SNAT --to-source 89.4.205.9
-A PREROUTING -i eth0 -p tcp -m tcp --dport 8080 -j
DNAT --to-destination 58.48.66.108:80
```

# Network capture

- Acquisition
  - `tcpdump -i eth0 -s 0 -w capture.pcap`
- Analysis
  - `tshark -r capture.pcap`
  - Wireshark
  - `bro -r capture.pcap`



# Malware analysis

# Two approaches

- Script-based malware
  - PHP
  - Perl
  - Python
- Compiled malware
  - ELF executables



# Script-based malware

# Script-based malware

- Can be obfuscated
  - Removed whitespace
  - Variables renamed

```
<?php function  
PZN1YnN0ci($a,$b){$c=array(139,164,40,72);if($b==62)  
{$d=substr($a,$c[0]+$c[1],$c[2]);}elseif($b==12){$d=substr($a,$c[0],$c[1]);}  
}elseif($b=92){$d=trim(substr($a,$c[0]+$c[1]+$c[2]));}return$d;} ?>
```

# Reversing script-based malware

- Most programming languages have a tool to *tidy* code
  - Perl -> `perltidy`
  - Python -> PythonTidy
  - PHP -> `php-cs-fixer`
  - Etc.
- Rename variables with search and replace

# Script-based malware

- Strings and literals can be packed

- 43 ^ 0x20 + 30
  - \x42\x56

```
$stg="ba"."\x73\x65"."64_d".strrev("edo\x63e");eval($st  
g("JHNlcnZlc191c2VyX2FnZW  
JWRVJbJ0hUVFBfVVNFU19BR0VOVCdd0w0KJHNlcnZlc19yZWZlc  
mVyI  
CAgICAgID0gQCRfU0VSvkvSwy  
[...]  
dIVFRQX1JFRkVSRVIInXTsNCiRzZXJ2ZXJfZm9yd2FyZGVkX2Zvc  
iA9IEAkX1N")
```

# Reversing script-based malware

- Always work in an isolated environment
- Use interactive prompts to evaluate parts of the code
  - Perl -> perl -de1
  - Python -> ipython
  - PHP -> php -a
  - Etc.
- Replace eval with print



# Compiled malware

# Compiled malware

- ELF executable in the native architecture of the system
- More challenging to understand
- Can also be packed

# Reverse engineering compiled malware

- Statically
  - strings
  - radare2
  - IDA Pro (\$)

# Reverse engineering compiled malware

- Dynamically
  - strace
  - ltrace for dynamically linked binaries
  - gdb, or any other debugger you like
  - gcore
- **Always work in an isolated environment when playing with malware**

# RSA®Conference2020

**Apply what we've learned**

# This week you should

- Get hands-on experience using the Hunting Linux Malware for Fun and Flags workshop

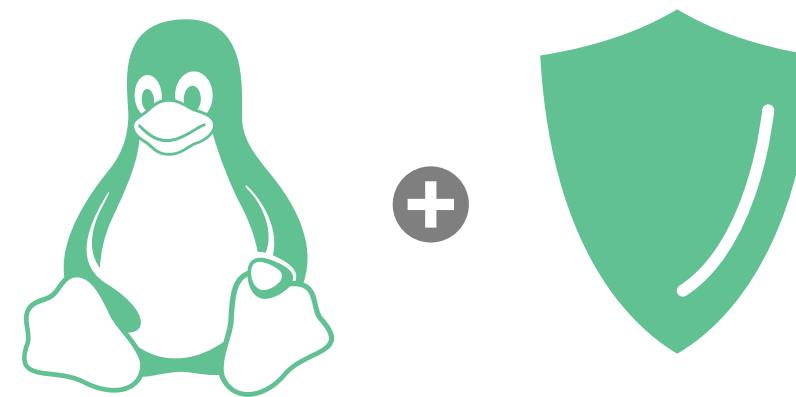
The screenshot shows a web browser window titled "localhost:5000" displaying the "Hunting Linux Malware for Fun and Flags" workshop. The page header includes "RSAC 2020 Edition". Below the header, there's a "Session information" section with a note about being assigned "sandbox03-299" and a placeholder for a voucher code. The status of the sandbox is "Running". A green button allows download of an OpenVPN configuration file. Below this are three tabs: "Slides", "Video", and "Cheat sheet". The main area is titled "Flags" and contains a text input field with "flag{1234} or 10.3.4.5" and a "Submit" button. Below the input field are three sections for "Server 1", "Server 2", and "Server 3", each showing a progress bar at 0/7, 0/9, and 0/11 respectively. At the bottom, there are two "How to connect" sections: "How to connect (Windows)" with the instruction "1. Install OpenVPN from the [openvpn.net download page](#)." and "How to connect (Linux, OS X)" with the instruction "OpenVPN must be installed: use your package manager or".

# Within three months you should

- Identify your Linux assets
  - Both on-premise and rented
  - Who has access to them?
  - From where are they reachable?
- Read articles and papers about Linux malware

## Next you should

- Enable 2FA on all Linux servers
- Consider deploying security product (EPP and/or EDR) on your Linux servers<sup>1</sup>
  - *Malware and attacks impact systems, whether they are workstations, servers, mobile devices or assets in the cloud*



<sup>1</sup> Jon Amato, Mario de Boer, *Gartner Technical Professional Advice – Solution Criteria for Endpoint Protection Platforms*, 2020-01-09



<https://www.eset.com> • <https://www.welivesecurity.com> • @ESETResearch